MUCO: A Mani-Unit-Structured Manipulation Framework for Compositional Long-Horizon Robot Tasks

Xiangyi Wei, Hanzhe Yao, Jingwen Shang, Zhenhong Wu, Yang Li, Changbo Wang
School of Computer Science and Technology,
East China Normal University, Shanghai 200062, China
{52285901014, 51295901117, 10235102513, 10235102482}@stu.ecnu.edu.cn
{yli, cbwang}@cs.ecnu.edu.cn
Corresponding authors: Yang Li, Changbo Wang

arXiv:2610.19913v1 [cs.RO] 14 Oct 2026

Abstract

Learning and executing long-horizon robot manipulation as complete task trajectories raises data-collection costs, limits capability reuse, and causes local failures to cascade into later task stages. We introduce manipulation units (MUs) as reusable operations, each defined by a local goal, an initiation condition, and observable success and failure criteria. MUs are learned and verified individually. During execution, consecutive MUs remain coupled through shared environment dynamics. We formulate MU-Structured Manipulation (MUSM) and realize it through MU-Centric Orchestration (MUCO). Given language and vision, MUCO constructs an MU sequence, matches each MU to compatible policy models in a frozen capability registry, and dispatches the model with the highest calibrated success rate. Visual verification checks MU outcomes against success and failure criteria and governs transitions, enabling retry or fallback to another model from the current state without re-executing verified MUs. Canonical keys enable checkpoint reuse for equivalent MUs across tasks without new demonstrations or updates. We evaluate MUCO on RLBench, RoboTwin 2.0, and five real-world tasks. MUCO achieves higher complete-task success and verified progress, stronger failure containment, and lower error propagation than baselines. Results demonstrate cross-task MU reuse and heterogeneous policy composition without joint training, shared backbones, or common action representations.

Introduction

A robot that can grasp a container, open its lid, and pour its contents should be able to recombine those capabilities when asked to prepare a different drink. Under the dominant trajectory-centric paradigm, however, a new ordering or combination is still treated as a new end-to-end task, typically requiring another complete demonstration or another task-level policy. The limitation is not simply that current policies are too small or their action horizons too short. The deeper problem is that the trajectory is asked to play too many roles at once. It is the unit of data collection and policy learning, as well as the unit of execution, evaluation, and recovery. For long-horizon manipulation, however, that unit is too coarse-grained.

Vision-language-action (VLA) models and related multimodal policies have substantially broadened the behaviors that robots can express by coupling perception, language, and continuous action generation (Brohan et al. 2023; Kim et al. 2024; Physical Intelligence et al. 2025; Driess et al. 2023; Jiang et al. 2023; Open X-Embodiment Collaboration 2024). Yet greater policy capacity does not by itself delineate where one operation ends, whether its intended state change was achieved, or which existing controller should handle the next operation. These questions are of particular importance in laboratory procedures, household activities, and assembly tasks, where an early error changes the physical context faced by every later action (Mees et al. 2022; Jiang et al. 2023). Long-horizon competence therefore requires not only a policy that can act, but also an organization of operation that makes intermediate progress reusable and inspectable.

The trajectory-centric formulation couples semantically distinct operations into one data collection, learning, and execution problem. At the data level, collecting demonstrations for a lengthy task requires complete successful rollouts, even when most constituent operations have already appeared in other tasks (Open X-Embodiment Collaboration 2024; Khazatsky et al. 2024). A change in task order or final objective can therefore require collecting another full trajectory. At the capability level, the trajectory-level formulation does not expose the operations learned by a policy model as units that can be selected, verified, and recombined. Therefore, knowing how to grasp, pour, and insert does not by itself enable a robot to execute those operations in a new long-horizon ordering (Mees et al. 2022; Jiang et al. 2023). At execution time, an early failure changes the state encountered by all later operations. Without explicit Mani-Unit boundaries, it is difficult to check whether each operation succeeded or to retry only the failed operation (Hu et al. 2026; Zhang and Shah 2026). This problem arises because learning and execution are organized around complete trajectories rather than reusable Mani-Units.

To address these problems, we introduce manipulation units (MUs) as reusable operations, each defined by a local goal, an initiation condition, and observable success and failure criteria. An MU specifies what state change should be achieved, while the selected policy model determines how to achieve it through continuous robot actions. Building on MUs, we formulate MU-Structured Manipulation (MUSM), which represents long-horizon tasks as ordered MU compositions. In MUSM, individually learned and verified MUs are recombined into different task sequences. During execution, the outcome of each MU determines the initial state

of the next, so consecutive MUs remain coupled through shared environment dynamics. Data are collected and policy models are trained separately for each MU. The same MU can then be reused across different task sequences and executed by different compatible policy models.| Compared with language-guided planners that primarily organize execution through high-level skill selection or program generation (Ahn et al. 2022; Liang et al. 2023; Huang et al. 2022), MUSM does more than generate an execution sequence. MUSM represents long-horizon tasks as ordered compositions of reusable and verifiable MUs, and different tasks are formed by recombining the same MUs. Each MU retains its own data provenance and observable success and failure criteria. MUSM therefore changes both how task sequences are constructed and the granularity at which data are collected, policies are trained, outcomes are verified, and learned capabilities are reused across tasks. We realize MUSM through **MU-Centric Orchestration (MUCO)**, which comprises four components: an MU constructor that converts task instructions and visual observations into ordered MU records; a matcher backed by a frozen capability registry that identifies compatible policy models and ranks them using calibrated success rates; model-specific adapters that translate observations and MU records into each policy model's native input; and a visual verifier that evaluates success and failure criteria and governs MU transitions and recovery. During execution, MUCO dispatches the highest-ranked compatible model for the active MU. After verified success, MUCO advances to the next MU. After failure, MUCO retries the current model or falls back to another compatible model from the current environment state without re-executing verified MUs. Our experiments test whether this Mani-Unit-centric organization addresses the limitations of trajectory-centric learning described above. We first examine whether a learned Mani-Unit can be reused in a new task. We then test whether assigning different policy models to different Mani-Units improves task success. Finally, we evaluate whether Mani-Unit verification and retry prevent an early error from affecting the rest of a long-horizon task. Our contributions are:

- We identify trajectory granularity as a shared bottleneck in data collection, capability reuse, and failure handling, introduce manipulation units (MUs) as reusable and verifiable operations, and formulate MUSM to represent long-horizon tasks as ordered compositions of MUs that can be recombined across task sequences.
- We introduce MUCO, a concrete realization of MUSM that integrates an MU constructor, a matcher backed by a frozen capability registry, model-specific adapters, and a visual verifier for transition control and failure recovery, without jointly retraining the underlying policy models.
- We establish an evaluation protocol that separates individual policy-model competence from the central claims of cross-task MU reuse without additional demonstrations, heterogeneous policy composition, and robustness to MU-level failures.

Related Work

Vision-language-action models. VLA models and related multimodal policies learn language-conditioned visuomotor behavior from robot demonstrations and large-scale vision-language data (Brohan et al. 2023; Kim et al. 2024; Physical Intelligence et al. 2025; Driess et al. 2023; Jiang et al. 2023; Open X-Embodiment Collaboration 2024). Scaling these models improves the breadth and generality of individual executors, but does not prescribe how a learned capability should be reused across task compositions or replaced after a local failure. MUSM is complementary to such scaling: a VLA can be registered once as an executor for every Mani-Unit it supports, reused across tasks without task-specific architectural changes, and replaced by a specialist only where calibrated evidence indicates that the specialist is more reliable. The contribution is therefore at the learning-and-deployment interface around policies, rather than in a new action decoder.

Hierarchical VLA orchestration. Hierarchical VLA agents pair a high-level VLM with a low-level VLA and study planner choice, controller steerability, termination, observation representation, and memory (Hu et al. 2026). This line establishes that orchestration choices matter even when a strong controller is available. MUSM is complementary but changes the object being orchestrated. A conventional hierarchy typically issues transient subgoals to one low-level controller; MUSM maintains a Mani-Unit as a persistent entity with its own data provenance, admissibility and outcome criteria, and a set of independently trained compatible executors. MUCO can consequently assign different executors to different occurrences of the same Mani-Unit and reuse the Mani-Unit record across task compositions.

Language-guided task organization. Language models have been used to select skills according to affordances, synthesize executable programs, and replan from environmental feedback (Ahn et al. 2022; Liang et al. 2023; Huang et al. 2022). These approaches establish the value of semantic reasoning above robot control. MUSM focuses on the learning and deployment unit beneath that reasoning. Rather than treating a skill call as an atomic symbolic action whose implementation is hidden from the planner, MUSM retains each Mani-Unit as a stochastic feedback-control problem with Mani-Unit-level data, explicit outcome criteria, and multiple possible executors. The organizer sequences Mani-Units, but does not abstract away their continuous execution or its uncertainty.

Task planning and hierarchical reinforcement learning. Classical task planners reason over symbolic action schemas, predicates, and precondition–effect structures (Kaelbling, Littman, and Cassandra 1998); hierarchical reinforcement learning and options introduce temporal abstraction through initiation conditions, intra-option policies, and termination rules (Sutton, Precup, and Singh 1999; Dietterich 2000). MUSM can be embedded in this formalism by treating each Mani-Unit–executor pair as an option, so our claim is not a new temporal-abstraction calculus. The distinction is the factorization MUSM makes explicit. A semantic Mani-Unit is separated from the policy that executes it; one Mani-Unit may map to multiple heterogeneous executors; a shared

observation-based verifier judges all of them; and Mani-Unit data and capability estimates are reused across tasks. Options provide a language for temporally extended control, but do not by themselves prescribe this Mani-Unit-level organization of data, verification, calibration, and cross-task deployment.

Problem Formulation

We consider a language-conditioned manipulation task in a shared partially observed controlled process,

$$\mathcal{M} = (\mathcal{S}, \mathcal{A}, \mathcal{O}, P, \rho_0), \tag{1}$$

where $s_t \in \mathcal{S}$ is the unobserved environment state, $o_t \in \mathcal{O}$ is the robot observation, $a_t \in \mathcal{A}$ is a continuous robot action, $P(s_{t+1} \mid s_t, a_t)$ is the shared transition dynamics, and ρ_0 is the initial-state distribution. The robot receives a language instruction ℓ and must satisfy its task-level completion condition. MUSM organizes trained executors within this shared environment without factorizing P or assuming that the state produced by one MU is irrelevant to the next.

Within this shared controlled process, MUSM associates the task with an ordered sequence of MUs $\mathbf{u} = (u_1, \ldots, u_K)$. An MU is defined as

$$u_i = (g_i, \mathcal{S}_i, \mathcal{A}_i, \mathcal{I}_i, v_i, \Pi_i), \tag{2}$$

where g_i is a local goal, $\mathcal{S}_i \subseteq \mathcal{S}$ and $\mathcal{A}_i \subseteq \mathcal{A}$ specify the state and action components relevant to that goal, $\mathcal{I}_i$ is an initiation condition, v_i returns $\{\textsc{success}, \textsc{failure}, \textsc{continue}\}$, and Π_i is a set of compatible stochastic execution policies. Policies in Π_i may differ in architecture, training data, or control precision, provided that they can pursue g_i through an interface compatible with the MU. The evolution of u_i remains governed by the global dynamics P over the full state, rather than by an independent MU-specific kernel. Each fixed pair (u_i, π_k) is representable as an option; further discussion is provided in the supplementary material.

Boundary-Level Organization. Let $i_m \in \{1, \ldots, K\}$ denote the position of the active MU at the mth decision boundary τ_m, and let h_{τ_m} contain all earlier MU–policy attempts, verification outcomes, and boundary observations. Separating i_m from m is important: a failed attempt creates a new decision boundary without advancing the ordered MU plan. Let $\Pi_{i_m}^{(m)} \subseteq \Pi_{i_m}$ be the policies still eligible after the preceding attempts. At each boundary, an organizer ω selects an executor for the *active* MU using the current observation, task instruction, and history:

$$k_m \sim \omega(\cdot \mid u_{i_m}, o_{\tau_m}, \ell, h_{\tau_m}), \qquad \pi_{k_m} \in \Pi_{i_m}^{(m)}. \tag{3}$$

The active MU must satisfy the explicit admissibility condition

$$\mathcal{I}_{i_m}(o_{\tau_m}, \ell, h_{\tau_m}) = 1. \tag{4}$$

The selected executor generates stochastic continuous controls while the MU remains active,

$$\begin{aligned} a_t \quad &\sim \quad \pi_{k_m}(\cdot \mid o_t, u_{i_m}), \\ &\qquad \tau_m \le t < \tau_{m+1}. \end{aligned} \tag{5}$$

where τ_{m+1} is the first time at which v_{i_m} returns $\textsc{success}$ or $\textsc{failure}$. On success, $i_{m+1} = i_m + 1$. On failure, the organizer retains $i_{m+1} = i_m$ and updates $\Pi_{i_m}^{(m)}$ to implement retry or policy fallback from the current global state; it aborts when no eligible policy remains. This formulation preserves the MU order supplied by the instruction rather than assuming unconstrained task replanning. It introduces explicit verification and executor-reassignment points into a global, state-dependent continuous-control problem.

The task succeeds only when every MU in the instruction reaches a verified $\textsc{success}$ boundary:

$$\mathrm{Succ}(\ell) = \mathbf{1}\left[\bigwedge_{i=1}^{K} \exists m :\ i_m = i \ \wedge\ v_i(o_{\tau_{m+1}}) = \textsc{success}\right]. \tag{6}$$

Let $\mathcal{D}_{\mathrm{task}}$ denote the distribution over language-conditioned task instances. The organizer's task-level objective is to maximize long-horizon task success:

$$\begin{aligned} \omega^* \quad &\in \quad \arg\max_{\omega} J(\omega), \\ J(\omega) \quad &= \quad \mathrm{E}_{\omega}[\mathrm{Succ}(\ell)] \\ &= \quad \Pr_{\omega}[\mathrm{Succ}(\ell) = 1]\,. \end{aligned} \tag{7}$$

The expectation is over $\ell \sim \mathcal{D}_{\mathrm{task}}$, $s_0 \sim \rho_0$, the shared dynamics P, and stochastic executor rollouts. MU-level retry and executor fallback are mechanisms for improving this task-level objective rather than separate objectives.

Method

MUCO realizes MUSM through four components. A Mani-Unit constructor converts task instructions into shared Mani-Unit records, an offline registry links Mani-Units to compatible executors and calibrated success rates, model-specific adapters translate shared inputs into each executor's native inputs, and a visual verifier controls Mani-Unit transitions and recovery. Executors act at their native control frequencies, while MUCO intervenes only at Mani-Unit boundaries and every Δ seconds for verification. This design coordinates heterogeneous policies without joint training, a shared backbone, or a common action representation.

Instruction-Grounded Mani-Unit Records

At time t, MUCO receives the three-view RGB observation $\mathbf{I}_t = (I_t^{\mathrm{left}}, I_t^{\mathrm{front}}, I_t^{\mathrm{right}})$, consisting of images from the left-wrist, front, and right-wrist cameras, respectively. This observation forms the visual component of o_t. Given a task instruction ℓ and the initial three-view observation $\mathbf{I}_0$, the Mani-Unit constructor produces an ordered sequence of implementation records:

$$\begin{aligned} F_{\mathrm{unit}}(\ell, \mathbf{I}_0) \quad &= \quad (\xi_1, \ldots, \xi_K), \\ \xi_i \quad &= \quad (z_i, \ell_i, c_i^+, c_i^-, T_i). \end{aligned} \tag{8}$$

Here, $z_i \in \mathcal{Z}$ identifies a canonical Mani-Unit, ℓ_i describes the required local state change, c_i^+ and c_i^- specify observable success and failure conditions, and T_i sets the maximum duration. The three-view observation $\mathbf{I}_0$ grounds each record in the current scene, while z_i links equivalent Mani-Unit descriptions to the same capability entry.

We implement the semantic part of F_{unit} using a frozen multimodal language model with a fixed output schema.

Given $(\ell, \mathbf{I}_0)$ and $\mathcal{Z}$, the model returns the canonical key, operation description, and visual success and failure criteria for each Mani-Unit; the constructor appends the task-configured T_i. Restricting z_i to $\mathcal{Z}$ allows the registry, dispatcher, and verifier to use the same Mani-Unit record consistently.

In the MUSM formulation, ℓ_i specifies g_i, while (c_i^+, c_i^-, T_i) defines v_i. The remaining elements are instantiated by the evaluator, registry, and executor adapters. At decision boundary τ_m, the evaluator computes

$$\begin{aligned} \iota_m &= F_{\text{init}}(\mathbf{I}_{\tau_m}, \ell, \xi_{i_m}, h_{\tau_m}) \\ &= \mathcal{I}_{i_m}(o_{\tau_m}, \ell, h_{\tau_m}) \in \{0, 1\}. \end{aligned} \tag{9}$$

Using the same frozen model with a binary output schema, F_{init} checks whether the active Mani-Unit can begin. MUCO dispatches the Mani-Unit only when $\iota_m = 1$. After a failed attempt, it evaluates the latest observation again so that recovery starts from the observed state after failure.

Offline Mani-Unit–Executor Capability Registry

Let $\mathcal{E} = \{e_k\}_{k=1}^M$ be a library of independently trained executors, where e_k implements stochastic policy π_k. Each registry entry stores an identifier, a capability and interface descriptor b_k, a launcher, a model-specific adapter ψ_k, and a frozen checkpoint. The descriptor b_k contains canonical Mani-Unit keys, language aliases, and required observation and action fields. This common description covers policies that use only sensor observations as well as VLAs that also receive ℓ_i.

Executors that support the same Mani-Unit may differ in reliability. MUCO therefore calibrates each compatible pairing of a Mani-Unit and an executor before evaluation. For pair (z, k), the registry records N_{zk} trials, S_{zk} successes, and the empirical success rate

$$\widehat{q}_{zk} = \frac{S_{zk}}{N_{zk}}. \tag{10}$$

Simulation calibration uses privileged environment success checks. For physical calibration, human reviewers correct the model's visual judgments when needed. These outcomes are used only to build the offline registry and are never available to the online organizer. The registry remains frozen throughout evaluation.

For a Mani-Unit record ξ_i, a matcher compares the canonical key, local description, and interface requirements with each descriptor b_k. The compatible executor set is

$$\mathcal{E}_i = \{e_k \in \mathcal{E} : F_{\text{match}}(z_i, \ell_i, b_k) = 1\}. \tag{11}$$

The policies associated with $\mathcal{E}_i$ instantiate the compatible set Π_i in the MUSM formulation. We implement F_{match} with the language branch of the same frozen model and a schema that requires a binary output. The matcher uses the Mani-Unit description rather than exact string matching. Compatibility determines which executors can be used, while $\widehat{q}_{zk}$ determines their ranking.

Capability-Aware Dispatch and Adaptation

At decision boundary τ_m, let $B_{i_m}^{(m)} \subseteq \mathcal{E}_{i_m}$ contain the executors whose retry budgets for the active Mani-Unit have been exhausted. The policies associated with $\mathcal{E}_{i_m} \setminus B_{i_m}^{(m)}$ instantiate the remaining eligible set $\Pi_{i_m}^{(m)}$ in the MUSM formulation. When this set is nonempty, MUCO selects the eligible executor with the highest calibrated success rate:

$$k_m^* = \arg \max_{k: e_k \in \mathcal{E}_{i_m} \setminus B_{i_m}^{(m)}} \widehat{q}_{z_{i_m} k}. \tag{12}$$

The dispatcher retrieves the selected entry's launcher, adapter, and checkpoint. It starts the executor, and the adapter converts the current observation and Mani-Unit record into the executor's native input:

$$a_t \sim \pi_{k_m^*}\left(\cdot \mid \psi_{k_m^*}(o_t, \xi_{i_m})\right). \tag{13}$$

Selection is therefore based on measured reliability for the active Mani-Unit. The adapter allows policies with different observation formats, action chunk sizes, and control rates to retain their native interfaces without parameter sharing or retraining.

The selected executor remains fixed during an attempt, so organizer decisions do not interrupt fast feedback control. After a verified failure, MUCO records the executor, outcome, and boundary observation before retrying or choosing the eligible executor with the next highest score. This fallback handles individual outcomes that differ from the calibrated rates and requires no learned meta policy.

Table 1: Mani-Unit transitions controlled by verification. Recovery never replays an already verified Mani-Unit.

Signal	Condition	Transition
CONTINUE	any	keep current executor
SUCCESS	$i < K$	record boundary; activate $i+1$
SUCCESS	$i = K$	declare task success
FAILURE	retry available	retry current executor
FAILURE	fallback available	dispatch next executor
FAILURE	otherwise	declare task failure

Verification-Gated Transition and Local Recovery

During attempt m, the selected policy runs continuously from its native observation stream. Every Δ seconds, MUCO samples the current three-view RGB observation $\mathbf{I}_t$ and evaluates the active Mani-Unit

$$\begin{aligned} y_{i_m,t} &= V(\mathbf{I}_t, c_{i_m}^+, c_{i_m}^-, t - \tau_m, T_{i_m}) \\ &\in \{\text{CONTINUE}, \text{SUCCESS}, \text{FAILURE}\}. \end{aligned} \tag{14}$$

The verifier returns SUCCESS when $c_{i_m}^+$ is visually satisfied, FAILURE when $c_{i_m}^-$ is satisfied or the elapsed time $t - \tau_m$ exceeds T_{i_m}, and CONTINUE otherwise. It uses the same three-view RGB interface in simulation and on the physical robot. Simulator success signals are used only for offline calibration and final evaluation, never for online transitions.

At every decision boundary τ_m, the active Mani-Unit starts from the physical boundary state s_{τ_m} produced by the complete preceding execution history, including verified Mani-Units and any failed attempts. Although this state is only

partially observed, retries and executor fallback both continue from the current global state rather than from a reset or local Mani-Unit initial state. Mani-Unit execution and recovery therefore remain governed by the shared transition dynamics P of the controlled process in Eq. (1), preserving physical coupling across Mani-Unit boundaries.

Table 1 summarizes the resulting state machine. On SUCCESS, MUCO records the verified boundary in $h_{\tau_{m+1}}$, advances to $i_{m+1} = i_m + 1$, and checks whether the next Mani-Unit can begin. On FAILURE, it records the boundary, retains $i_{m+1} = i_m$, and checks the active Mani-Unit again using the latest observation and history. If the Mani-Unit can still begin and the executor has retries remaining, MUCO restarts it from the current state. Otherwise, the executor is added to $B_{i_m}^{(m+1)}$ and the dispatcher selects the best remaining alternative. The task fails when the active Mani-Unit can no longer begin or no eligible executor remains.

Recovery applies only to the active Mani-Unit, so previously verified Mani-Units are not replayed. Every recovery attempt must nevertheless pass the initiation check again. The realized outcomes $y_{i_m,\tau_{m+1}}$ instantiate $v_{i_m}(o_{\tau_{m+1}})$ in the task-success event of Eq. (6).

Experiments

Evaluation Questions and Protocol

Our experiments test three hypotheses that follow directly from the Mani-Unit–executor factorization. **H1 (Mani-Unit organization):** with low-level policy capacity controlled, Mani-Unit decomposition, explicit verification, and local retry or fallback improve long-horizon task completion. **H2 (Mani-Unit reuse):** an executor trained for a Mani-Unit can be reused, without additional Mani-Unit demonstrations or checkpoint updates, when that Mani-Unit appears in a new task composition. **H3 (heterogeneous composition):** selecting a different specialist for each Mani-Unit outperforms forcing one policy model to execute every Mani-Unit. We additionally test whether these conclusions are robust to the multimodal agent used by MUCO.

Terminology. We reserve *Mani-Unit* for the reusable manipulation unit defined in Section 3 and *executor* for a concrete policy checkpoint that can realize it. We use *Mani-Unit-level policy* only when emphasizing training granularity. This avoids overloading *skill*, *option*, or *action*, which refer to different abstractions in prior work.

Metrics and statistical protocol. Task success requires every Mani-Unit to reach a verified SUCCESS boundary; normalized progress is the fraction of Mani-Units verified before termination. We also report per-Mani-Unit success, low-level actions, error propagation $\Pr(F_{i+1} \mid F_i, i+1 \text{ attempted})$, and failure containment, the fraction of failed first attempts recovered without replaying a verified prefix. Unless an official protocol differs, each configuration uses the same 20 initial conditions; success rates include 95% Wilson intervals and paired differences.

Simulation Experiments

RLBench: Standard Task Competence

Role of the benchmark. RLBench (James et al. 2020) provides a standardized comparison with strong VLA policies, but its ten tasks are predominantly short horizon. We therefore use it to establish that Mani-Unit-structured execution preserves standard manipulation competence, rather than as the primary evidence for long-horizon composition. We evaluate Close box, Close laptop lid, Toilet seat down, Sweep to dustpan, Close fridge, Phone on base, Take umbrella out, Frame off hanger, Wine at rack, and Water plants. Following LaST0 and DeepVision-VLA, the robot is a Franka Panda in CoppeliaSim with a single front-view RGB observation. The reported protocol uses 100 waypoint-generated demonstrations per task, 20 rollout trials per task, and three random seeds.

Comparison methods. The comparison spans OpenVLA (Kim et al. 2024), SpatialVLA (Qu et al. 2025), CogACT (Li et al. 2024), CoT-VLA (Zhao et al. 2025), $\pi_{0.5}$ (Physical Intelligence et al. 2025), HybridVLA (Liu et al. 2026a), LaST0 (Liu et al. 2026b), and DeepVision-VLA (Luo et al. 2026). Prior RLBench values in Table 2(a) are transcribed from LaST0 and DeepVision-VLA rather than re-evaluated in our codebase; the exact $\pi_{0.5}$ update count is not disclosed by those reports.

RoboTwin: Bimanual and Compositional Simulation RoboTwin 2.0 (Chen et al. 2025) complements RLBench with bimanual manipulation and strong domain randomization. We use a fixed four-task subset for Mani-Unit composition and heterogeneous dispatch, with 20 matched rollouts per method. Exact task names, decompositions, and demonstration budgets will be stated in the final protocol. Table 2(b) compares the executor library, MUCO, and X-VLA (Zheng et al. 2026), which provides an official RoboTwin 2.0 checkpoint.

Real-World Experiments

Platform, Mani-Unit Library, and Tasks All real-world configurations use the same bimanual robot platform, camera placement, object set, and low-level control interface. These factors are held fixed across methods so that comparisons change only the data and execution organization stated below. Each executor is trained independently on demonstrations of the Mani-Units it supports; no executor is jointly trained with the MUCO organizer.

The canonical Mani-Unit vocabulary contains six entries: A, TRANSFER, grasps an object and moves it to the work area; B, POUR, transfers contents between containers; C, TITRATE, dispenses reagent with a burette; D, PRESS, activates an apparatus; E, CARRY, moves an object to a target table; and F, INSERT, aligns and inserts an object into a receptacle or fixture. INSERT provides a general contact-rich operation beyond liquid handling. The library contains independently trained executors e_1, e_2, e_3 and, where interfaces are compatible, X-VLA; the final paper will report their model families, supported keys, and demonstration counts per Mani-Unit in the implementation details rather than a separate table.

Before formal evaluation, every compatible Mani-Unit–executor pair is tested in $N_{\text{cal}} = 20$ independent calibration trials. Calibration initial states are disjoint from both training

Table 2: Simulation task success rate (%). (a) RLBench prior values are reported by LaST0 (Liu et al. 2026b) and DeepVision-VLA (Luo et al. 2026). (b) RoboTwin 2.0 results are placeholders. Higher is better.

(a) RLBench Method	Box	Laptop	Toilet	Sweep	Fridge	Phone	Umbrella	Frame	Wine	Water	Mean
OpenVLA	60	35	75	55	85	20	30	15	20	5	40
SpatialVLA	80	70	85	20	80	15	25	40	15	30	46
CogACT	90	80	95	50	85	50	55	45	30	25	61
CoT-VLA	95	75	100	80	65	50	40	50	55	50	66
$\pi_{0.5}$	90	95	85	75	100	5	10	80	75	35	65
HybridVLA	85	95	100	90	100	50	50	70	50	50	74
QwenVLA-OFT	95	95	100	15	95	65	30	70	65	50	69
LaST0	95	95	100	80	85	75	75	70	85	60	82
DeepVision-VLA	100	95	100	95	95	75	65	70	85	50	83
MUCO (ours)	–	–	–	–	–	–	–	–	–	–	–

(b) RoboTwin 2.0 Method	RT-1	RT-2	RT-3	RT-4	Mean
X-VLA	–	–	–	–	–
Executor e_1	–	–	–	–	–
Executor e_2	–	–	–	–	–
Executor e_3	–	–	–	–	–
MUCO (ours)	–	–	–	–	–

Table 3: H1 results. SR is complete-task success, Prog is normalized verified progress, EPC is error propagation, and FCR is failure containment. All entries are placeholders.

Method	Short SR	Task-A SR	Task-B SR	General SR	Mean SR	Prog	EPC	FCR
Flat-Same	–	–	–	–	–	–	–	–
Trajectory-E2E	–	–	–	–	–	–	–	–
MUCO-Same	–	–	–	–	–	–	–	–

demonstrations and formal test configurations. The resulting $\widehat{q}_{zk}$ values and compatibility masks are frozen before any long-horizon trial.

We evaluate four compositions: Task-Short, $A \to C$; Task-A, $A \to B \to C \to E \to D$; Task-B, $A \to C \to B \to D \to E$; and Task-General, $A \to F \to D \to E$. No complete trajectory from Task-A, Task-B, or Task-General is used to train an executor for a Mani-Unit.

H1: Mani-Unit Organization versus Trajectory Coupling This experiment isolates Mani-Unit organization from low-level policy capacity. A single comparison cannot simultaneously use an identical checkpoint and different training granularity, so we report two controls. FLAT-SAME uses the same executor checkpoint as MUCO; this checkpoint is trained at Mani-Unit granularity, but Flat-Same disables boundary verification, retry, and fallback. This is the organization-only ablation. TRAJECTORY-E2E trains one task-level policy from complete trajectories using a matched interaction-frame budget; this compares the full trajectory-centric and Mani-Unit-centric formulations. MUCO-SAME forces every Mani-Unit to use the same executor, preserving Mani-Unit construction, visual verification, and retry with $R = 3$ while removing heterogeneous selection.

Table 4: Independent Mani-Unit success rate (%).

Executor	A	B	C	D	E	F
e_1	–	–	–	–	–	–
e_2	–	–	–	–	–	–
e_3	–	–	–	–	–	–
X-VLA	–	–	–	–	–	–

Independent Mani-Unit Capability We next execute each Mani-Unit from a standardized admissible initial state, without preceding or subsequent Mani-Units. Each compatible (z, k) pair is evaluated for 20 trials. The matrix in Table 4 serves two roles: it measures whether individual executors are reliable on their intended Mani-Units, and it provides the frozen capability evidence used to interpret heterogeneous selection. Non-compatible pairs are marked N/A rather than counted as failures.

For keys $(z_1, \ldots, z_K)$, the product $\widehat{U}_{\text{ind}} = \prod_i \max_k \widehat{q}_{z_i k}$ is an *independence estimate*, not an upper bound, because shared-state dependencies correlate Mani-Unit outcomes.

H2–H3: Reuse and Heterogeneous Composition To test H3, all configurations retain the same constructor, veri-

Table 5: Heterogeneous composition and cross-task reuse. (a) H3 task success rate under different executor rules. (b) H2 reuses the same checkpoint with zero new Mani-Unit demonstrations. All result entries are placeholders.

(a) H3: heterogeneous versus homogeneous execution

Method	Rule	A SR	B SR	Mean	V/R
MUCO-e_1	always e_1	–	–	–	Yes
MUCO-e_2	always e_2	–	–	–	Yes
MUCO-e_3	always e_3	–	–	–	Yes
MUCO-Het	$\arg\max_k \widehat{q}_{z_i k}$	–	–	–	Yes

(b) H2: cross-task Mani-Unit reuse

Mani-Unit	Exec.	A SR	B SR	$\lvert\Delta\rvert$	New
Transfer (A)	–	–	–	–	0
Titrate (C)	–	–	–	–	0

Table 6: Robustness to the multimodal agent. All entries are placeholders.

Agent	Mani-Unit F1	Key acc.	Ver. acc.	Task SR
Proprietary	–	–	–	–
Open-source	–	–	–	–

fier, retry budget, and Mani-Unit data. MUCO-Het selects $\arg\max_k \widehat{q}_{z_i k}$, whereas each homogeneous ablation forces one executor across all compatible Mani-Units. Table 5(a) therefore isolates heterogeneous dispatch from Mani-Unit organization.

H2 is established primarily by protocol: Table 5(b) invokes the same checkpoint for the same canonical key in different task compositions, with zero additional Mani-Unit demonstrations and no task-trajectory update. The success-rate difference Δ measures context sensitivity; similar rates support stability but are not required for reuse because different prefixes may produce different physical states.

We replace circular "selection accuracy" with selection regret, $\max_k \widehat{q}_{z_i k} - \widehat{q}_{z_i k_i^*}$, and separately count invalid Mani-Unit keys or compatibility decisions.

Agent Robustness Finally, we replace the frozen multimodal agent while holding the Mani-Unit vocabulary, executor checkpoints, capability registry, prompts, retry budget, and test configurations fixed. At least one proprietary and one open-source multimodal model are evaluated on Task-A. We compare the generated Mani-Unit keys and criteria with manually defined references, and compare visual verification with independently annotated ground truth.

References

Ahn, M.; Brohan, A.; Brown, N.; Chebotar, Y.; et al. 2022. Do As I Can, Not As I Say: Grounding Language in Robotic Affordances. arXiv:2204.01691.

Brohan, A.; Brown, N.; Carbajal, J.; Chebotar, Y.; Chen, X.; Choromanski, K.; Ding, T.; Driess, D.; Dubey, A.; Finn, C.; et al. 2023. RT-2: Vision-Language-Action Models Transfer Web Knowledge to Robotic Control. arXiv:2307.15818.

Chen, T.; Chen, Z.; Chen, B.; Cai, Z.; Liu, Y.; Li, Z.; Liang, Q.; Lin, X.; Ge, Y.; Gu, Z.; et al. 2025. RoboTwin 2.0: A Scalable Data Generator and Benchmark with Strong Domain Randomization for Robust Bimanual Robotic Manipulation. arXiv:2506.18088.

Dietterich, T. G. 2000. Hierarchical Reinforcement Learning with the MAXQ Value Function Decomposition. *Journal of Artificial Intelligence Research*, 13: 227–303.

Driess, D.; Xia, F.; Sajjadi, M. S. M.; et al. 2023. PaLM-E: An Embodied Multimodal Language Model. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, 8469–8488.

Hu, J.; Shridhar, M.; Lu, C.; Shah, D.; Chiang, H.-T. L.; Tan, J.; and Xie, A. 2026. What Matters in Orchestrating Robot Policies: A Systematic Study of Hierarchical VLA Agents. arXiv:2606.10267.

Huang, W.; Xia, F.; Xiao, T.; Chan, H.; Liang, J.; Florence, P.; Zeng, A.; Tompson, J.; Mordatch, I.; Chebotar, Y.; Sermanet, P.; Brown, N.; Jackson, T.; Luu, L.; Levine, S.; Hausman, K.; and Ichter, B. 2022. Inner Monologue: Embodied Reasoning through Planning with Language Models. In *Proceedings of the Conference on Robot Learning*, volume 205, 1769–1782.

James, S.; Ma, Z.; Arrojo, D. R.; and Davison, A. J. 2020. RLBench: The Robot Learning Benchmark and Learning Environment. *IEEE Robotics and Automation Letters*, 5(2): 3019–3026.

Jiang, Y.; Gupta, A.; Zhang, Z.; Wang, G.; Dou, Y.; Chen, Y.; Fei-Fei, L.; Anandkumar, A.; Zhu, Y.; and Fan, L. 2023. VIMA: Robot Manipulation with Multimodal Prompts. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, 14975–15022.

Kaelbling, L. P.; Littman, M. L.; and Cassandra, A. R. 1998. Planning and Acting in Partially Observable Stochastic Domains. *Artificial Intelligence*, 101(1–2): 99–134.

Khazatsky, A.; Pertsch, K.; Nair, S.; Balakrishna, A.; et al. 2024. DROID: A Large-Scale In-The-Wild Robot Manipulation Dataset. In *Proceedings of Robotics: Science and Systems*.

Kim, M. J.; Pertsch, K.; Karamcheti, S.; Xiao, T.; Balakrishna, A.; Nair, S.; Rafailov, R.; Foster, E.; Lam, G.; Sanketi, P.; et al. 2024. OpenVLA: An Open-Source Vision-Language-Action Model. arXiv:2406.09246.

Li, Q.; Liang, Y.; Wang, Z.; Luo, L.; Chen, X.; Liao, M.; Wei, F.; Deng, Y.; Xu, S.; Zhang, Y.; Wang, X.; Liu, B.; Fu, J.; Bao, J.; Chen, D.; Shi, Y.; Yang, J.; and Guo, B. 2024. CogACT: A Foundational Vision-Language-Action Model for Synergizing Cognition and Action in Robotic Manipulation. arXiv:2411.19650.

Liang, J.; Huang, W.; Xia, F.; Xu, P.; Hausman, K.; Ichter, B.; Florence, P.; and Zeng, A. 2023. Code as Policies: Language Model Programs for Embodied Control. In *Proceedings of the IEEE International Conference on Robotics and Automation*, 9493–9500.

Liu, J.; Chen, H.; An, P.; Liu, Z.; Zhang, R.; Gu, C.; Li, X.; Guo, Z.; Chen, S.; Liu, M.; Hou, C.; Zhao, M.; Zhou,

K.; Heng, P.-A.; and Zhang, S. 2026a. HybridVLA: Collaborative Diffusion and Autoregression in a Unified Vision-Language-Action Model. In *International Conference on Learning Representations*.

Liu, Z.; Liu, J.; Chen, H.; Guo, Z.; Hou, C.; Gu, C.; Yu, J.; Mi, X.; Zhang, R.; Che, Z.; Tang, J.; Heng, P.-A.; and Zhang, S. 2026b. LaST0: Latent Spatio-Temporal Chain-of-Thought for Robotic Vision-Language-Action Model. arXiv:2601.05248.

Luo, Y.; Chen, H.; Wu, Z.; Sui, B.; Liu, J.; Gu, C.; Liu, Z.; Feng, Q.; Yu, J.; Gu, S.; Jia, P.; Heng, P.-A.; and Zhang, S. 2026. Look Before Acting: Enhancing Vision Foundation Representations for Vision-Language-Action Models. arXiv:2603.15618.

Mees, O.; Hermann, L.; Rosete-Beas, E.; and Burgard, W. 2022. CALVIN: A Benchmark for Language-Conditioned Policy Learning for Long-Horizon Robot Manipulation Tasks. *IEEE Robotics and Automation Letters*, 7(3): 7327–7334.

Open X-Embodiment Collaboration. 2024. Open X-Embodiment: Robotic Learning Datasets and RT-X Models. In *Proceedings of the IEEE International Conference on Robotics and Automation*.

Physical Intelligence; Black, K.; Brown, N.; Darpinian, J.; Dhabalia, K.; Driess, D.; Esmail, A.; Equi, M.; Finn, C.; Fusai, N.; et al. 2025. $\pi_{0.5}$: A Vision-Language-Action Model with Open-World Generalization. arXiv:2504.16054.

Qu, D.; Song, H.; Chen, Q.; Yao, Y.; Ye, X.; Ding, Y.; Wang, Z.; Gu, J.; Zhao, B.; Wang, D.; and Li, X. 2025. SpatialVLA: Exploring Spatial Representations for Visual-Language-Action Model. arXiv:2501.15830.

Sutton, R. S.; Precup, D.; and Singh, S. 1999. Between MDPs and Semi-MDPs: A Framework for Temporal Abstraction in Reinforcement Learning. *Artificial Intelligence*, 112(1–2): 181–211.

Zhang, M.; and Shah, D. 2026. Visual Verification Enables Inference-Time Steering and Autonomous Policy Improvement. In *Proceedings of Robotics: Science and Systems*.

Zhao, Q.; Lu, Y.; Kim, M. J.; Fu, Z.; Zhang, Z.; Wu, Y.; Li, Z.; Ma, Q.; Han, S.; Finn, C.; Handa, A.; Liu, M.-Y.; Xiang, D.; Wetzstein, G.; and Lin, T.-Y. 2025. CoT-VLA: Visual Chain-of-Thought Reasoning for Vision-Language-Action Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*.

Zheng, J.; Li, J.; Wang, Z.; Liu, D.; Kang, X.; Feng, Y.; Zheng, Y.; Zou, J.; Chen, Y.; Zeng, J.; et al. 2026. X-VLA: Soft-Prompted Transformer as a Scalable Cross-Embodiment Vision-Language-Action Model. In *International Conference on Learning Representations*.