One Agent per Zone: Decentralised Active Inference under Partial Observability for Multi-Zone HVAC Control on Edge Hardware

Panagiotis Kasnesis, Waldiez P.C., Maroussi, Greece
Lazaros Toumanidis, Waldiez P.C., Maroussi, Greece
Christos Chatzigeorgiou, Waldiez P.C., Maroussi, Greece
Amalia Contiero Syropoulou, ThinGenious P.C., Maroussi, Greece
Miguel De Prado, PRAESC AI, Biel, Switzerland

arXiv:2610.19790v1 [cs.MA] 14 Oct 2026

***Abstract*—Multi-zone HVAC control is a shared-resource problem: zones on a floor share one fan and cooling coil, so each zone's demand affects its neighbours and the building's peak power. One controller per zone scales well and isolates faults, but each controller then observes only its own zone, not the shared air loop. In this work, we present a decentralised controller that assigns each zone an independent active-inference agent, which plans by minimising expected free energy and learns its zone dynamics online. The agents exchange no messages, so none depends on peers or a coordinator; each accounts for the unobserved coupling through a latent disturbance and a congestion signal derived from the metered power of the other floors. Fifteen agents are evaluated in EnergyPlus on a Sinergym office building against three rule-based controllers and three MADDPG fleets, one given the same floor meters. On five noisy weather years unseen during training, the fleet attains the lowest mean building peak ($88.3\,\mathrm{kW}$, 3 to 17 % below the baselines) and staggers the floor peaks better than any baseline, with a coincidence factor (building peak over the sum of the floor peaks) of 0.929 versus 0.936 to 0.994. Compared with a reactive rule at similar comfort (0.912 versus 0.916), it lowers the sum of monthly peaks, on which demand charges are billed, by $33\,\mathrm{kW}$ (4 %) in every year, at the cost of about 10 % more energy. It costs 9% less per year than the proactive rule (at lower comfort) and 1% less than the MADDPG fleet given the same floor meters, in every test year. The whole fleet runs on a single low-cost Raspberry Pi 5, closing the observation-to-action loop over MQTT in $442\,\mathrm{ms}$ at p95, far below the 15-minute control step.**

***Index Terms*—edge AI, generative AI, active inference, multi-agent systems, building energy management**

I. Introduction

Heating, ventilation and air conditioning (HVAC) systems account for a large share of the electricity consumed by commercial buildings, and their role is no longer limited to keeping the indoor temperature within a comfort band. Many commercial tariffs now charge for the highest power drawn in each billing month as well as for the energy consumed, and grid operators are increasingly interested in when buildings draw power, not only in how much they consume. As a consequence, a modern building controller has to maintain occupant comfort while also limiting and shaping the demand profile of the building.

For a building served by several air loops, the building peak depends not only on how large the peak of each loop is, but also on whether these peaks occur at the same time. This is measured by the coincidence factor, the ratio of the building peak to the sum of the individual loop peaks: a value of 1 means that all loops peak simultaneously, and lower values mean that their peaks are staggered in time. Reducing coincidence therefore lowers the building peak without necessarily reducing the energy used by each loop.

In a multi-zone building, the shared HVAC infrastructure induces a multi-agent control problem: zones served by the same air handler compete for limited fan and cooling capacity, so the actions of one zone affect the thermal conditions of the others. A single controller that decides all setpoints jointly could account for this coupling, but it scales poorly with the number of zones, is a single point of failure, and has to be redesigned whenever the zoning changes. We therefore consider a decentralised architecture in which each zone is controlled by an independent agent that observes only its own zone, without access to the full building state or communication with neighbouring agents. Variables such as the supply air conditions and the demand of the other zones still influence the local dynamics, but they are not observable by a zone agent and are therefore treated as latent disturbances.

This partial observability matters for two reasons. First, a zone agent cannot directly distinguish changes caused by its own thermal load from changes induced by the other zones on the same air handler; it has to infer their effect from its local observations. Second, peak demand is a collective phenomenon: the building peak is set when several zones call for cooling at the same time, which no zone can see from its own measurements. Effective decentralised control therefore requires agents that reason about hidden influences and, as we show, a minimal aggregate signal that reveals how busy the rest of the building is.

The dominant learning-based approach to this problem has been multi-agent deep reinforcement learning (DRL) under centralised training with decentralised execution (CTDE) [1], [2], in which a centralised critic observes the whole building during training while each deployed actor sees only its own zone. Although this paradigm works, it inherits the properties of its training regime: the coupling between zones is absorbed into the network weights rather than being represented explicitly, and the resulting policy is difficult to inspect at the moment it matters.

In this paper, we take a different route and assign each zone to an independent *active-inference* (AIF) agent [3]. Each agent maintains beliefs over a small set of discrete latent states, corrects them with Bayes' rule as new observations arrive, learns its own transition model online, and selects actions by evaluating how well each candidate action satisfies its comfort and energy preferences over the next two hours. The generative model also contains a hidden disturbance factor that is never observed: when a zone systematically runs warmer or cooler than its model predicts, the agent's belief about this factor shifts accordingly, which is how an agent accounts for an air-loop coupling that it cannot measure. In addition, every agent reads the metered power of the other two air loops as a congestion signal and becomes more reluctant to draw power when the rest of the building is already busy. AIF has recently been applied to building energy management with one agent per building and a supervising community-level agent [4]; however, to the best of our knowledge, this is the first work that applies decentralised AIF at the zone level to multi-zone HVAC control over a shared air loop, and evaluates it against both rule-based control (RBC) and DRL controllers on demand-related metrics.

The main contributions of this work are summarised as follows:

- We present a decentralised AIF controller for shared-resource HVAC systems, in which 15 zone agents stagger their demand across floors without exchanging a single message. Evaluated on five noisy weather years that were not seen during training, the AIF fleet achieves the lowest mean annual peak ($88.3\,\mathrm{kW}$) among the nine controllers that we tested and a lower coincidence factor (0.929) than any baseline.
- We identify why the fleet coordinates. Our ablations show that two ingredients are needed together: a real-time signal of what the rest of the building is drawing, and agents that respond to it differently from one another. Replacing the live signal with a constant of the same mean, or with the signal of the previous day, removes the effect, which rules out a higher energy price or a memorised daily pattern as explanations. Giving the same meters to a MADDPG fleet also lowers its peak, which supports the conclusion that the live load information, and not active inference alone, is what enables the staggering.
- We provide a transparent comparison with well-tuned RBCs, reporting the metrics on which simple rules still prevail and quantifying the demand charge at which the demand savings of the AIF fleet would compensate for its additional energy use.
- We show that all 15 agents run concurrently on a Raspberry Pi 5, with a p95 observation-to-action latency of $442\,\mathrm{ms}$ measured over the full MQTT round trip.

The remainder of this paper is organised as follows. Section II provides the necessary background on AIF, while Section III reviews related work. Section IV introduces the simulation testbed, the evaluation protocol and the controllers. Section V defines the evaluation metrics, Section VI presents the obtained results, and Section VII discusses their implications and limitations. Finally, Section VIII concludes the article. All code, trained models and benchmark harnesses are released under Apache 2.0 at https://github.com/waldiez/aif-sinergym.

II. Background

A. Active Inference in Discrete State Spaces

AIF [3] describes an agent that holds a probabilistic generative model of how its observations are produced by hidden states of the world, and of how these states evolve under its own actions. In the discrete formulation used in this work, the model is specified by a likelihood mapping $\mathbf{A}$, which gives the probability of each observation given the hidden state, a transition model $\mathbf{B}$, which gives the probability of the next state given the current state and the selected action, and a preference distribution $\mathbf{C}$ over observations, which encodes the outcomes that the agent seeks. When the state space is factored, each factor (e.g., temperature, occupancy or a latent disturbance) has its own likelihood and transition terms, which keeps the model small enough to be updated in real time.

At every time step, the agent first infers the posterior over hidden states by combining its prediction from the previous step with the likelihood of the new observation. It then evaluates each candidate policy through its expected free energy, which can be decomposed into a *pragmatic* term, rewarding predicted outcomes that match the preferences in $\mathbf{C}$, and an *epistemic* term, rewarding actions that are expected to reduce uncertainty about the hidden states. As a result, information seeking is part of the objective itself rather than an exploration heuristic added on top of it. Finally, the transition model can be learned online by accumulating Dirichlet counts over observed transitions, so the agent refines its own model of the plant while it controls it.

III. Related Work

Reinforcement learning for HVAC control has been studied extensively on simulation benchmarks built over EnergyPlus [5], among which Sinergym [6], [7] is the framework adopted in this work. Multi-agent formulations typically follow MADDPG [1] with the stabilisation techniques of TD3 [8], which is also the DRL baseline that we implement. Single-agent deep RL has been benchmarked extensively on such testbeds. Campoy-Nieves et al. [7] compared SAC, TD3 and PPO with default and RBC in a two-zone data centre, where PPO learned the most stable policies and saved energy mainly in the intermediate seasons. Zhan et al. [9] compared six algorithms across office and data-centre buildings, climates and reward functions, and reported energy reductions of about 9 to 10 % over RBC, with an exponential comfort penalty reducing temperature fluctuations compared with a linear one. In a national-scale energy setting, Bajrami et al. [10] found SAC faster and more stable than PPO, whereas PPO produced more assertive actions. Sample efficiency remains a central obstacle for these methods: Qu et al. [11] augmented a SAC heating controller with an LSTM surrogate and branched

rollouts, more than halving its training time, while Kadamala et al. [12] pre-trained agents by behavioural cloning and fine-tuned them across buildings with different observation spaces by padding the input features. All of these works control the building with a single agent; the coupling between zones that share an air loop, which is central to our problem, is not modelled explicitly.

AIF has so far been applied mainly to low-dimensional control problems, and `pymdp` [13] provides a reference implementation for discrete state spaces. Its appeal for building control stems from two properties: the epistemic value is part of the objective rather than an exploration rule added afterwards, and the generative model can be inspected at decision time, which makes the behaviour of each agent easier to interpret. Applications of AIF to energy systems are still scarce compared to DRL. Nazemi et al. [4] proposed a dual-layer architecture in which a continuous AIF agent per building sets the HVAC airflow and supply-air temperature while inferring occupancy and infiltration as hidden states, and a discrete AIF agent coordinates storage, photovoltaics and market transactions at the community level. The building agent matched a full-information optimisation benchmark and compared favourably with a DDPG controller. In that work, however, each building is modelled as a single zone and coordination is achieved hierarchically through the community agent. Our agents instead operate at the zone level, share a physical air loop, and coordinate without any supervising agent or message exchange.

IV. Testbed and Controllers

A. Building and Simulation

We use Sinergym [6], [7] driving EnergyPlus 24.1.0 on the ASHRAE 90.1 OfficeMedium prototype building [14] (Denver model, $4979.6\,\mathrm{m}^2$, three floors), but with a custom environment rather than one of the predefined Sinergym scenarios: we expose per-zone setpoints as actions, extend the observation space, and add one electricity meter per air loop, as described below. In particular, of its 18 thermal zones, 15 are conditioned (3 cores and 12 perimeters), while the remaining 3 are return plenums. Each floor is served by one packaged air loop, in which five zones share one fan, one two-speed cooling coil and one outdoor-air system, while each zone has its own damper, electric reheat coil and dual-setpoint thermostat. In addition, we added one electricity meter per air loop, covering the fan, the cooling coil and the reheat coils of the corresponding floor; such a meter per air handler is a modest addition in a commercial building.

The observation vector consists of 95 values: time, outdoor conditions, 18 zone temperatures and humidities, 15 occupancy counts, 30 setpoints, the whole-building HVAC demand and energy, and the three floor meters. The action vector contains a heating and a cooling setpoint for each conditioned zone (30 values), bounded to $15.0\,^{\circ}\mathrm{C}$ to $22.5\,^{\circ}\mathrm{C}$ for heating and $22.5\,^{\circ}\mathrm{C}$ to $30.0\,^{\circ}\mathrm{C}$ for cooling, as specified by our experimental setup using the sinergym simulator. It should be noted that these bounds do not by themselves keep the two setpoints apart, which is why we report deadband violations for every controller. The simulation timestep is 15 minutes, which results in 35 040 steps per simulated year.

B. Evaluation Protocol

Every controller is evaluated under three weather conditions. The first one, denoted as W0, is the deterministic "mixed" TMY3 year (New York JFK), and it is the only year on which the learning controllers are trained or adapted. The second one, W1, consists of the same year perturbed by Ornstein-Uhlenbeck noise on dry-bulb temperature ($\sigma = 1.5\,^{\circ}\mathrm{C}$) and humidity, generated with five different seeds; these five years constitute our main test set. For a given seed, every controller is exposed to exactly the same weather, which we verify by comparing the logged outdoor temperatures, so that all comparisons on W1 can be paired by year. Finally, W2 comprises the hot (Davis-Monthan, Arizona) and cool (Port Angeles, Washington) TMY3 years provided with Sinergym, which serve as unseen climates. All learned controllers are tested with learning switched off, so that no controller adapts to its test years. All controllers target, and are scored against, the same seasonal comfort band: $20.0\,^{\circ}\mathrm{C}$ to $23.5\,^{\circ}\mathrm{C}$ in winter and $23.0\,^{\circ}\mathrm{C}$ to $26.0\,^{\circ}\mathrm{C}$ in summer (1 June to 30 September).

C. Controllers

All controllers share the same interface: every 15 minutes they receive the observation vector and return the 30 setpoints. The learning controllers are deployed in a decentralised manner, with one agent per zone.

1) Rule-based controllers: The two main RBCs are implemented in a single code base and differ in exactly two behavioural switches, so that the pair isolates the value of anticipation and structural awareness rather than comparing unrelated heuristics. The *reactive RBC* applies a seasonal threshold rule on the temperature of each zone and on an occupancy calendar, with a $2\,^{\circ}\mathrm{C}$ setback when the building is unoccupied and without any anticipation. The *proactive RBC* additionally performs a pre-conditioning phase between 04:00 and 06:00 when the outdoor temperature is low, and applies per-zone offsets derived from the core or perimeter position, the orientation and the floor of each zone. The proactive RBC has one setting that controls how hard it cools in summer, the cooling margin. We run it with margins from $0.5\,^{\circ}\mathrm{C}$ to $2.5\,^{\circ}\mathrm{C}$: a small margin gives more comfort at a higher cost, and a large margin the opposite. This shows the full range of comfort and cost that this rule can reach (Fig. 1), so that the AIF fleet is not compared with a single tuning only.

Moreover, we implemented a *demand-limiting RBC*, which extends the proactive RBC with two textbook peak-management measures. First, the three floors start their pre-conditioning 40 minutes apart; second, whenever the building HVAC demand exceeds 90 % of $92.4\,\mathrm{kW}$ (the annual peak of the proactive RBC), all cooling setpoints are raised and all heating setpoints are lowered by $1\,^{\circ}\mathrm{C}$ until the demand falls again. We also tested other thresholds, a larger shed and a rotating shed across floors, none of which lowered the peak.

2) MADDPG: As a DRL baseline, we use a multi-agent actor-critic controller [1] with TD3 stabilisers [8], trained with a centralised twin critic and deployed as 15 independent actors. Each actor receives a 29-dimensional feature vector describing its zone, which contains time, weather, whole-building HVAC demand and energy, and local zone features including a fixed zone identity. The actor consists of a single-layer GRU with 64 hidden units applied over the last eight steps, followed by a small multilayer perceptron, for a total of about 40k parameters. The networks were trained on the shaped reward described in Section V for 50 simulated years, corresponding to roughly 3.5 million gradient updates. We evaluate three fleets: a cold start, a start whose replay buffer was seeded with one year of the reactive RBC, and a cold start whose actors additionally receive the three floor meters (own floor, other floors and whole building), so that the latter has access to the same information as the AIF agents.

3) Active inference: Each zone agent maintains a factored discrete generative model with four factors: the zone temperature T (19 bins, $0.5\,^{\circ}\mathrm{C}$ wide near the comfort band), the occupancy O (2 states, derived from the weekday 06:00 to 20:00 calendar), the outdoor temperature W (5 bins), and a hidden disturbance H (5 bins spanning $\pm 0.8\,^{\circ}\mathrm{C}$) that is never observed. The agent selects among 25 setpoint pairs, formed by five heating values (15.0 to $22.5\,^{\circ}\mathrm{C}$) and five cooling values (22.5 to $30.0\,^{\circ}\mathrm{C}$).

The transition model $\mathbf{B}$ is initialised from a simple physical prior, in which the HVAC drives the zone towards the commanded band and heat is lost towards the outdoor temperature at a rate that depends on whether the zone is a core, a perimeter or located on the top floor. The agent then refines this model from its own experience by means of Dirichlet count updates. Its preferences $\mathbf{C}$ reward temperatures inside the comfort band, relaxed by $4\,^{\circ}\mathrm{C}$ when the zone is unoccupied, and penalise HVAC effort.

At every time step, the agent updates its beliefs and then scores each of the 25 actions by assuming that the action is held for the next eight steps (two hours):

$$G_i(a) = \sum_{k=1}^{8} \Big(\mathbb{E}[C(T_k)] - \big(w_i^E + \lambda_c\, c_i\big)\, \mathbb{E}[D(T_k, a)] \Big) + \beta\, I(T_1; H) - \lambda_{db}\, P_{db}(a), \tag{1}$$

where C is the comfort preference, D the expected HVAC effort, w_i^E the energy weight of agent i, c_i its congestion signal, $I(T_1; H)$ the expected information gain about the hidden disturbance, and P_{db} a penalty for heating and cooling setpoints that are closer than $2.5\,^{\circ}\mathrm{C}$. We use $\lambda_c = 5$, $\beta = 0.2$ and $\lambda_{db} = 8$. Since G is the negative of the expected free energy, the agent selects the action with the highest G. Finally, a reactive safety rule takes over whenever a zone is more than $2\,^{\circ}\mathrm{C}$ outside its comfort band, which occurred in about 4.4 % of the decisions.

Coordination without communication. No agent sends any message to another agent; nevertheless, two terms of Eq. (1) give rise to coordinated behaviour. First, each agent reads the power drawn by the *other two* floors from the meters, divides it by a fixed reference of $92.4\,\mathrm{kW}$ (the annual peak of the incumbent controller, which a facility manager can read from the previous year's bill), and smooths the result into a congestion signal c_i. In this way, drawing power becomes more expensive for an agent when the rest of the building is busy. Second, each agent adapts its own energy weight w_i^E within $[0, 1]$, increasing it while its zone is comfortable and decreasing it otherwise. Since all agents of a floor receive the same congestion signal, these heterogeneous weights are what prevent the five agents from reacting in lockstep.

We refer to this configuration as *AIF (fixed reference)*. For comparison, we also report a variant, *AIF (running reference)*, in which each agent divides the load of the other floors by the highest total floor power that it has observed so far instead of a fixed value. This running maximum is taken over the total power of the three floors and starts from the first load observed after deployment, which is far below the eventual peak; it then only grows over time, so during the first weeks of operation the agents perceive the building as more congested than it actually is. Moreover, its value depends on the history of each deployment, so the same model can scale the signal differently depending on how long it has been running and which loads it has observed. The fixed reference avoids both issues, and we therefore use it as our main configuration, while the running reference serves as a sensitivity check.

We implemented the agents from scratch in PyTorch, so that the whole fleet is expressed as a set of batched tensor operations, which runs on a GPU during simulation and on a CPU at the edge. Importantly, the batch of 15 agents contains no coupling across its batch dimension: we verified that each row of the batch produces exactly the same decisions as a separate single-agent model fed with the same inputs, which is what allows us to deploy 15 truly independent actors from a single model. Each fleet adapts for five years on W0 and is then frozen for testing.

V. Evaluation Metrics

All controllers are scored by the same code against the comfort band of Section IV-B, with comfort measured on weekdays between 07:00 and 19:59.

Comfort. The comfort score CS_{occ} assigns each occupied zone and step a value of 1 when the temperature lies inside the band, and $1 - \min(v, 1)^2$ when it lies $v\,^{\circ}\mathrm{C}$ outside it; these values are then averaged over zones and steps. In addition, the zone comfort rate (ZCR) reports the share of zone-steps that lie strictly inside the band.

Demand. We report the annual peak of the building HVAC demand, the sum of the twelve monthly peaks (which is the quantity actually charged by the tariff), and the coincidence factor

$$\mathrm{CF} = \frac{\max_t \sum_f P_f(t)}{\sum_f \max_t P_f(t)}, \tag{2}$$

where P_f denotes the power drawn by floor f. As explained in Section I, a CF equal to 1 means that the three floors peak

at the same moment, whereas lower values indicate that their peaks are spread out in time.

Cost. We adopt a simplified tariff modelled on PG&E Schedule B-19 [15], with time-of-use energy prices of 0.22, 0.32 and 0.45 \$/kWh (off-peak, mid-peak and on-peak, respectively) and a single demand charge of \$26 per kW of monthly peak.

Finally, the MADDPG fleets are trained on a shaped reward that is shared by all agents and computed at every step t as

$$r_t = -\Big(w_E\, m_t\, \rho(P_t) + w_C\, \kappa_t\, \phi_t + w_{DB}\, \delta_t + w_P\, \rho(\Delta P_t^{pk})\Big) + w_M\, \mu_t\, \psi_t, \tag{3}$$

where P_t is the building HVAC demand and $\rho(x) = \min\big(\log(1 + x/P_{max})/\log 2,\ 2\big)$ is a logarithmic normalisation with $P_{max} = 300\,\mathrm{kW}$. The electricity term is scaled by the time-of-use multiplier m_t, which equals 1, 1.5 and 3 in the off-peak, mid-peak and on-peak periods, respectively, and is increased by a further 10 % when the building is unoccupied. The comfort penalty ϕ_t averages over the zones the band violation v_z of each zone, normalised as $\min(v_z/s, 1)$; during occupied hours $s = 4$ and violations in the direction of the seasonal load (too warm in summer, too cold in winter) are weighted by 1.3, while before occupancy $s = 5$ and after occupancy the band is additionally widened by $3\,^\circ\mathrm{C}$. The occupancy weight κ_t equals 1 during occupied hours (07:00 to 19:59 on weekdays), 0.4 before occupancy (05:00 to 07:00), 0.15 after occupancy (20:00 to 21:59) and 0 otherwise. The deadband penalty $\delta_t = \frac{1}{15}\sum_z \min\big(\max(h_z - c_z + 2, 0)/3,\ 1\big)$ penalises heating and cooling setpoints h_z and c_z that are closer than $2\,^\circ\mathrm{C}$, and the peak term penalises the increase $\Delta P_t^{pk} = \max(P_t - \hat{P}_t, 0)$ over the running monthly peak $\hat{P}_t$, so that it is non-zero only when a new monthly peak is set. Finally, the midpoint bonus ψ_t averages $\max(1 - |T_z - T_{mid}|/2.5,\ 0)$ over the zones, where T_{mid} is the middle of the comfort band, with μ_t equal to 1 during occupied hours, 0.5 before occupancy (with a scale of 4 instead of 2.5) and 0 otherwise. We use $w_E = 0.14$, $w_C = 0.40$, $w_{DB} = 0.05$, $w_P = 0.17$ and $w_M = 0.10$.

VI. Results

A. Main Comparison on the Noisy Test Years

Table I presents the results obtained on the five noisy test years. Since every controller was exposed to the same weather in each year, we also compare the controllers year by year in Table II, reporting the mean paired difference together with the number of years in which the AIF fleet performed better.

Peak demand and coincidence. The AIF fleet achieves the lowest mean annual peak (88.3 kW, against 90.9 to 106.0 kW for the baselines) and a lower coincidence factor than any baseline (0.929). Its peak is lower than that of the reactive RBC and of every MADDPG fleet in every single test year. Against the proactive RBC, however, the annual peak is lower in only three of the five years, so we do not claim a robust peak advantage over that controller.

Comfort and cost. The two RBCs span the range of comfort and cost, and the AIF fleet lies between them. Together with the demand-limiting RBC, the proactive RBC is both the most comfortable controller (CS_{occ} of 0.946) and the most expensive one (about \$61k), whereas the reactive RBC is the cheapest (\$53.7k) and among the least comfortable. The AIF fleet costs \$5.7k less than the proactive RBC in every year, at the price of lower comfort. Compared with the reactive RBC, it reaches a similar comfort level (0.912 against 0.916) with a markedly lower and better-staggered demand, but it consumes about 10 % more energy; under the evaluated tariff, the demand savings do not compensate for this additional energy, so the AIF fleet costs \$1.6k more per year, as discussed further in Section VII.

The cooling-margin sweep. Fig. 1 summarises the comfort, cost and demand trade-offs of all controllers. Changing the summer cooling margin of the proactive RBC trades comfort for cost, but its peak remains between 91 and 93 kW on the noisy years. Interestingly, on the deterministic year one setting ($2.5\,^\circ\mathrm{C}$) happened to miss the design-day peak and reached 82.8 kW, whereas on the noisy years the same setting averaged 92.8 kW. This observation illustrates the risk of judging controllers on a single year: an annual peak is a single 15-minute value, and one favourable day can change it substantially.

Demand limiting. The demand-limiting RBC left the peak unchanged on the deterministic year and even raised it on the noisy years (97.2 kW, against 90.9 kW for the plain proactive RBC). This can be explained by the fact that the annual peak occurs on hot afternoons, when the zones are already above their setpoints, so raising the setpoint by one or two degrees does not reduce the cooling call; moreover, the rule reacts only after the demand has already spiked. A tighter limit even raised the peak further, most likely through a rebound when the shed was released.

Behaviour on the design day. Fig. 2 shows the power drawn by each floor on 5 August, the day of the annual peak for most controllers. Under the proactive RBC, all three floors increase their demand together at about 15:30, and the building reaches 91.6 kW. Under the AIF fleet, the afternoon load remains flatter, and when the bottom floor rises at 15:00 the middle floor stays at about 20 kW instead of following it, which limits the building peak to 83.8 kW. It should also be noted that the AIF fleet starts conditioning earlier in the morning, which appears as a short spike at around 04:30.

B. Deterministic Year and First-Year Behaviour

Table III reports the results on the deterministic year, which is also the year on which the learning controllers were trained. Here, the AIF fleet reaches 84.7 kW and a CF of 0.876, and it already outperforms every baseline in its first year of operation, starting from the physical prior (86.0 kW, CF of 0.855). Two further results of this table are worth mentioning. First, the proactive RBC violates the deadband in 10.5 % of the zone steps, since in winter some core zones receive identical heating and cooling setpoints. Second, even without

TABLE I

RESULTS ON THE FIVE NOISY TEST YEARS (W1), MEAN ± STD. LEARNED CONTROLLERS WERE TRAINED OR ADAPTED ON THE DETERMINISTIC YEAR ONLY AND ARE TESTED WITH LEARNING SWITCHED OFF. LOWER IS BETTER FOR DEMAND, ENERGY AND COST. HIGHER IS BETTER FOR COMFORT.

Controller	Peak (kW)	Σ monthly peaks (kW)	CF	CS_{occ}	ZCR (%)	Energy (MWh)	Cost (k$)
AIF (fixed reference)	**88.3** ± 2.8	765 ± 14	0.929 ± 0.030	0.912 ± 0.003	83.6 ± 0.4	117.4 ± 2.7	55.4 ± 1.1
AIF (running reference)	88.5 ± 3.6	**763** ± 14	**0.927** ± 0.028	0.908 ± 0.003	82.9 ± 0.5	116.7 ± 2.6	55.0 ± 1.1
AIF, no congestion signal	94.1 ± 2.3	784 ± 16	0.967 ± 0.015	0.938 ± 0.001	88.7 ± 0.2	121.2 ± 2.8	57.8 ± 1.2
RBC reactive	93.8 ± 0.5	799 ± 14	0.974 ± 0.005	0.916 ± 0.003	85.6 ± 0.5	**106.2** ± 2.9	**53.7** ± 1.2
RBC proactive	90.9 ± 3.4	767 ± 17	0.958 ± 0.014	**0.946** ± 0.001	**90.6** ± 0.1	139.2 ± 2.8	61.1 ± 1.3
RBC demand-limiting	97.2 ± 2.7	781 ± 19	0.983 ± 0.007	**0.946** ± 0.001	90.5 ± 0.1	138.0 ± 2.8	61.2 ± 1.3
MADDPG cold start	106.0 ± 2.1	787 ± 5	0.994 ± 0.005	0.924 ± 0.001	83.1 ± 0.3	120.3 ± 2.4	56.5 ± 0.8
MADDPG warm start (reactive)	103.5 ± 4.0	798 ± 16	0.977 ± 0.015	0.931 ± 0.002	84.3 ± 0.3	121.0 ± 2.6	57.0 ± 1.2
MADDPG + floor meters	97.0 ± 2.0	774 ± 15	0.936 ± 0.015	0.926 ± 0.002	83.4 ± 0.4	119.6 ± 2.6	55.9 ± 1.1

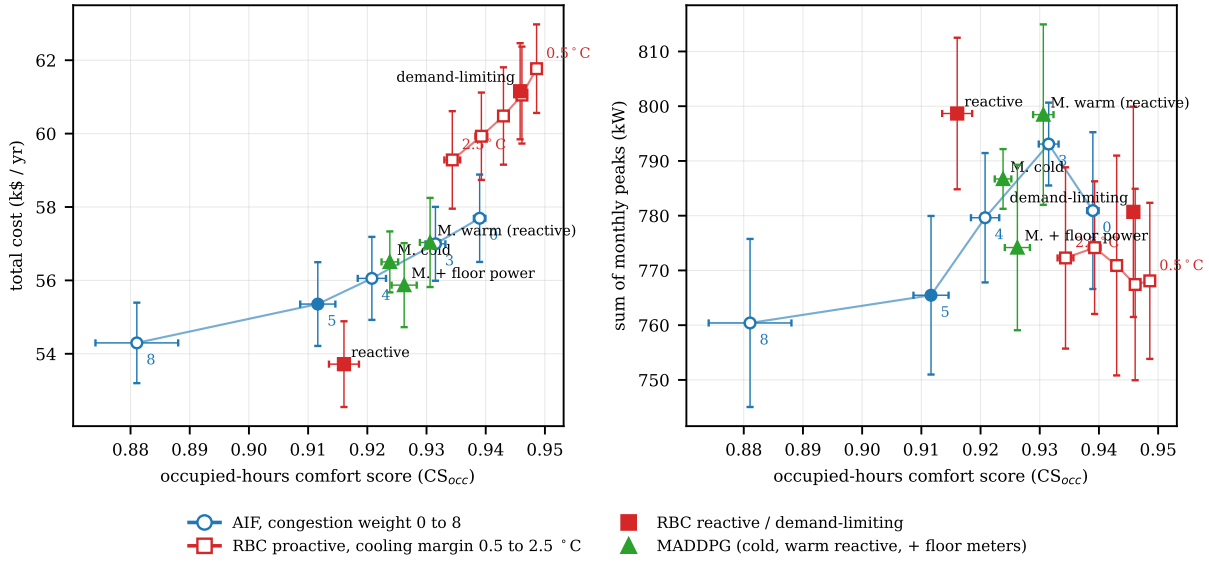

Fig. 1. Comfort against total cost (left) and against the sum of monthly peaks (right) on the five noisy test years (mean ± std over years). Blue: AIF with adaptive energy weights and congestion weights of 0, 3, 4, 5 and 8 (numbers next to the points); the filled point is AIF (fixed reference) with weight 5, while the other points use the running reference. Red: the proactive RBC with summer cooling margins from 0.5 to 2.5 °C, together with the reactive and demand-limiting RBCs (filled squares). Green triangles: MADDPG fleets.

TABLE II

PAIRED DIFFERENCES, AIF (FIXED REFERENCE) MINUS THE OTHER CONTROLLER, MEAN OVER THE FIVE W1 YEARS. THE COUNT SHOWS IN HOW MANY YEARS THE AIF FLEET WAS BETTER.

	RBC reactive	RBC proactive	MADDPG + meters
Annual peak (kW)	−5.5 (5/5)	−2.6 (3/5)	−8.7 (5/5)
Σ monthly peaks	−33 (5/5)	−2 (3/5)	−9 (3/5)
CF	−0.045 (4/5)	−0.029 (4/5)	−0.007 (4/5)
CS_{occ}	−0.004 (0/5)	−0.034 (0/5)	−0.015 (0/5)
Energy (MWh)	+11.2 (0/5)	−21.8 (5/5)	−2.1 (5/5)
Total cost (k$)	+1.6 (0/5)	−5.7 (5/5)	−0.5 (5/5)

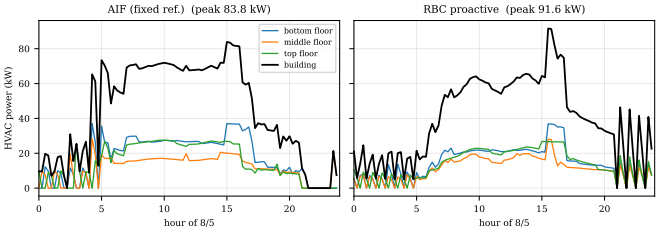

Fig. 2. HVAC power per floor on 5 August of the deterministic year. Left: AIF fleet (last adaptation year). Right: proactive RBC.

the congestion signal, learning the transition model alone lowers the peak of the AIF fleet from 100.1 kW in its first year to 92.9 kW.

These numbers are better than those obtained on W1, because the learning controllers are evaluated on the year they were trained on; for this reason, we regard W1 as our main result. In addition, the strong first year does not fully carry over to the noisy years: an AIF fleet starting from scratch on a noisy year beat the first-year peak of the AIF baseline in four of five years with the fixed reference (and in three of five with the running reference), although it was cheaper in all five.

TABLE III
DETERMINISTIC MIXED YEAR (W0). AIF: MEAN OF THE LAST THREE OF FIVE ADAPTATION YEARS. DB: DEADBAND VIOLATIONS (% OF ZONE STEPS).

Controller	Peak	CF	CS_{occ}	DB	MWh	k$
RBC reactive	93.2	0.987	0.918	0.0	103.8	52.2
RBC proactive	92.4	0.980	0.947	10.5	137.2	60.2
RBC demand-limiting	92.4	0.979	0.947	9.7	136.0	60.1
MADDPG cold	104.7	0.997	0.928	1.6	118.6	55.3
MADDPG warm (reactive)	99.0	0.958	0.934	1.6	119.1	55.8
MADDPG + floor meters	91.9	0.913	0.930	1.9	117.4	54.6
AIF, no congestion	92.9	0.974	0.939	0.0	119.4	56.6
AIF (running reference)	84.8	0.880	0.917	0.0	114.3	53.6
AIF (fixed reference)	84.7	0.876	0.923	0.0	115.4	54.1

TABLE IV
UNSEEN CLIMATES (W2), ONE YEAR EACH.

	Hot			Cool		
Controller	kW	CS_{occ}	k$	kW	CS_{occ}	k$
AIF (fixed reference)	102.0	0.798	80.1	48.5	0.928	27.9
AIF, no congestion	103.8	0.874	83.1	55.2	0.935	28.9
RBC reactive	99.9	0.840	78.7	50.5	0.921	25.6
RBC proactive	95.2	0.884	80.6	48.5	0.947	32.5
RBC demand-limiting	110.0	0.884	81.9	48.5	0.947	32.2
MADDPG cold	109.4	0.848	82.2	61.9	0.906	28.1
MADDPG warm (reactive)	107.8	0.860	82.8	61.5	0.922	28.5
MADDPG + floor meters	99.1	0.848	81.2	57.7	0.899	27.6

C. Unseen Climates

Table IV presents the results on the hot and cool climates. In the cool climate, the AIF fleet matches the peak of the proactive RBC ($48.5\,\mathrm{kW}$) while being \$4.6k cheaper. In the hot climate, however, it performs poorly: its comfort drops to a CS_{occ} of 0.798, the lowest among all controllers, and its peak is higher than that of both main RBCs. When the cooling plant operates at full capacity for long periods, the only way to reduce the peak is to sacrifice a considerable amount of comfort. Allowing the agents to keep learning online recovers a large part of the comfort (CS_{occ} of 0.846, against 0.798 when frozen), but the peak then rises to $115.6\,\mathrm{kW}$. Moreover, in this climate the fixed calendar of the summer comfort band penalises every controller, since on 1 October the band switches to the winter values while the outdoor temperature is still high.

D. AIF component contribution

Table V isolates the contribution of each component of the AIF controller on the deterministic year. As it is shown, neither the congestion signal nor the adaptive energy weights lower the peak on their own; only their combination does, and only from a congestion weight of 5 onwards, since at a weight of 4 the peak remains at $92\,\mathrm{kW}$. The peak therefore behaves like a switch rather than a slope, which is consistent with the coarse action set: on the design day, the fleet either keeps its cooling setpoint or gives up one whole step.

It is also observed that a constant signal with the same mean value, or the same signal delayed by 24 hours, never reaches the low-peak regime, even with a congestion weight of 8. Consequently, the effect is neither a higher energy price in disguise nor a learned daily pattern; instead, the agents react to what the building is drawing at that moment. Signals derived from the agent's own floor or from the whole building also work, but they require a weight of 8 and cost about 0.03 more in comfort, whereas the signal from the other floors reaches the same peak at a lower comfort cost.

TABLE V
AIF ABLATIONS AND MECHANISM CONTROLS ON THE DETERMINISTIC YEAR (MEAN OF THE LAST THREE ADAPTATION YEARS). ALL ROWS USE THE RUNNING REFERENCE UNLESS STATED OTHERWISE.

Configuration	Peak (kW)	CF	CS_{occ}
no congestion, fixed weights	92.9	0.974	0.939
adaptive weights only	92.9	0.976	0.939
congestion 5 only	92.3	0.934	0.930
congestion 4 + adaptive weights	92.0	0.923	0.924
congestion 5 + adaptive weights	**84.8**	**0.880**	0.917
fixed reference	84.7	0.876	0.923
earlier inference code	84.8	0.880	0.921
horizon 4 / 16 steps	87.1 / 85.0	0.891 / 0.929	0.919 / 0.904
stochastic actions	89.7	0.901	0.927
Signal source, congestion weight 5/8:			
own floor	91.9 / 85.0	0.919 / 0.877	0.921 / 0.893
whole building	92.6 / 83.7	0.928 / 0.874	0.921 / 0.892
24 h old	92.8 / 92.9	0.931 / 0.963	0.926 / 0.912
constant	92.9 / 92.9	0.961 / 0.959	0.933 / 0.930

Since every agent of a floor receives the same congestion signal, agents with identical weights react in lockstep. When we replace the adaptive weights with fixed random weights per agent, the fleet reaches the low-peak regime in 4 of 8 random draws, whereas the adaptive weights find a working combination every time, without anyone having to choose them. In practice, the weights rarely settle: most of the time each agent sits at one of the two bounds and switches frequently, so that at any given moment some agents save energy while others protect comfort.

Providing the MADDPG actors with the three floor meters lowered their peak by $9.0\,\mathrm{kW}$ and their CF by 0.057 in every noisy year, while slightly improving comfort. This result independently supports the explanation above, namely that live load information is what allows independent agents to spread their demand. Given the same information, the AIF fleet still achieves a lower peak (by $8.7\,\mathrm{kW}$), a lower cost and a lower energy consumption in every year, but its CF advantage shrinks to 0.007, which lies within the year-to-year variability. It should also be noted that MADDPG required 50 training years and a centralised critic to reach this performance, whereas the AIF fleet adapted for five years without any centralised training.

Furthermore, the additional signal did not change what MADDPG learned in terms of the metrics tracked during training: both cold-start runs reach the same final reward and comfort, and their comfort stops improving after about 30 to 40 training years, while the shaped reward still increases slowly at year 50.

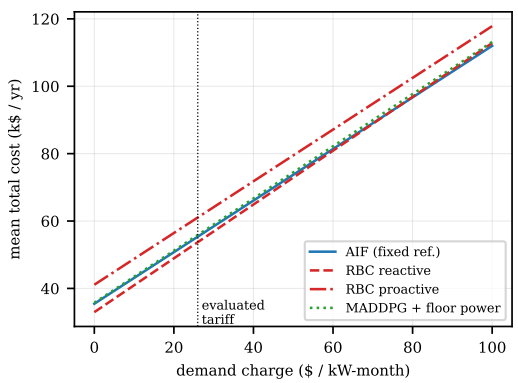

Fig. 3. Mean yearly cost on the noisy test years as a function of the demand charge. The dotted line marks the evaluated tariff ($26/kW-month).

E. Deployment on Edge Hardware

For the deployment we used the open-source `wactorz` actor runtime where each zone agent runs as an independent process and communicates only over MQTT, subscribing to the observation stream and publishing the setpoints of its zone. Decision latency was measured with all 15 AIF agents planning concurrently on a Raspberry Pi 5 while driving the simulator over MQTT. An instrumented bridge timestamps the full observation-to-action round trip, so the reported figures include the messaging overhead rather than computation alone. Over 90 timed decisions (after discarding 10 warm-up decisions), we measured a mean latency of 268 ms, a median of 238 ms, a 95th percentile of 442 ms and a maximum of 822 ms, which is negligible compared with the 15-minute control step.

VII. Discussion and Limitations

The AIF fleet trades energy for demand. Compared with the reactive RBC at similar comfort, the AIF fleet lowers the sum of monthly peaks by 33 kW per year and spreads the demand of the three floors, but it consumes 11 MWh more energy. Whether this trade-off pays off depends on the tariff: with the demand charge of $26/kW-month that we evaluated, it does not, and the fleet would break even with the reactive RBC only at a demand charge of about $75/kW-month (Fig. 3). Nevertheless, lower and better-staggered demand can also be valuable beyond the electricity bill, for instance under coincident-peak charges or local grid constraints, which our tariff does not price.

A well-tuned RBC is a strong baseline. On this building, the proactive RBC is the most comfortable controller, while the reactive RBC is the cheapest, and neither learning approach outperforms both of them on comfort and cost simultaneously. What the AIF fleet adds is lower and better-staggered demand, a single setting (the congestion weight) that covers the range between the two RBCs, and a decision mechanism that can be inspected.

Single-year peaks are fragile. Since an annual peak is a single 15-minute value, results obtained on a single year can be misleading: on the deterministic year, one RBC setting appeared better than the AIF fleet, and this advantage disappeared on the noisy years. Testing on several noisy years, comparing controllers year by year, and reporting the sum of monthly peaks changed several of our conclusions, and we recommend the same practice for HVAC benchmarks.

Scope and limitations. Our study has the following limitations that should be taken into account when interpreting the results:

- We evaluated a single building and a single tariff; the noisy years vary the weather, but not the building or the occupancy patterns.
- The AIF fleet is the least comfortable of the strong controllers, and in a hot climate its comfort degrades considerably when it is not allowed to adapt.
- Comfort is measured on weekdays between 07:00 and 19:59, although EnergyPlus also reports occupants outside these hours, for example in the early morning, in the evening and on Saturdays.
- Finally, each MADDPG variant was trained once.

VIII. Conclusion

In this paper, we presented a decentralised controller for multi-zone HVAC systems in which fifteen active-inference agents, one per zone and exchanging no messages, coordinate the demand of a building that shares one air loop per floor. Evaluated on five noisy weather years that were not seen during training, the AIF fleet achieved the lowest mean annual peak among all the controllers that we tested and a lower coincidence factor than any baseline. Our analysis showed that this coordination stems from a live load signal combined with agents that weigh energy differently, and that a deep reinforcement learning controller given the same meters improves for the same reason. The price of this behaviour is an energy consumption about 10 % higher than that of a simple reactive rule, which pays off only where demand is charged heavily. Finally, we showed that all 15 agents run concurrently on a Raspberry Pi 5 with a p95 latency of 442 ms, measured over the full MQTT round trip.

Future work follows directly from the limitations of this study. In particular, we plan to gate the congestion signal by comfort, so that the agents do not sacrifice comfort in hot climates, to refine the action set, and to evaluate the approach on more buildings and tariffs. Moreover, seeding the transition model with the structural knowledge that makes the proactive RBC so comfortable is a promising direction for combining the strengths of both approaches.

Acknowledgment

This work was funded by the European Network of Excellence dAIEDGE under Grant Agreement No 101120726.

References

[1] R. Lowe, Y. Wu, A. Tamar, J. Harb, P. Abbeel, and I. Mordatch, "Multi-agent actor-critic for mixed cooperative-competitive environments," in *Proceedings of the 31st International Conference on Neural Information Processing Systems*, ser. NIPS'17. Red Hook, NY, USA: Curran Associates Inc., 2017, p. 6382–6393.

[2] L. Yu, Y. Sun, Z. Xu, C. Shen, D. Yue, T. Jiang, and X. Guan, "Multi-agent deep reinforcement learning for hvac control in commercial buildings," *IEEE Transactions on Smart Grid*, vol. 12, no. 1, pp. 407–419, 2021.

[3] K. Friston, T. FitzGerald, F. Rigoli, P. Schwartenbeck, and G. Pezzulo, "Active inference: A process theory," *Neural Computation*, vol. 29, no. 1, pp. 1–49, 2017.

[4] S. D. Nazemi, M. A. Jafari, and A. Matta, "Active inference for energy control and planning in smart buildings and communities," arXiv preprint arXiv:2503.18161, 2025.

[5] D. B. Crawley, L. K. Lawrie, F. C. Winkelmann, W. F. Buhl, Y. J. Huang, C. O. Pedersen, R. K. Strand, R. J. Liesen, D. E. Fisher, M. J. Witte, and J. Glazer, "EnergyPlus: Creating a new-generation building energy simulation program," *Energy and Buildings*, vol. 33, no. 4, pp. 319–331, 2001.

[6] J. Jiménez-Raboso, A. Campoy-Nieves, A. Manjavacas-Lucas, J. Gómez-Romero, and M. Molina-Solana, "Sinergym: a building simulation and control framework for training reinforcement learning agents," in *Proceedings of the 8th ACM International Conference on Systems for Energy-Efficient Buildings, Cities, and Transportation*, ser. BuildSys '21. New York, NY, USA: Association for Computing Machinery, 2021, p. 319–323. [Online]. Available: https://doi.org/10.1145/3486611.3488729

[7] A. Campoy-Nieves, A. Manjavacas, J. Jiménez-Raboso, M. Molina-Solana, and J. Gómez-Romero, "Sinergym – a virtual testbed for building energy optimization with reinforcement learning," *Energy and Buildings*, vol. 327, p. 115075, 2025.

[8] S. Fujimoto, H. van Hoof, and D. Meger, "Addressing function approximation error in actor-critic methods," in *Proceedings of the 35th International Conference on Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 80, 2018, pp. 1587–1596.

[9] L. Zhan, G. Li, C. Xu, L. Zhang, Z. Xiao, Q. Li, S. Liu, and W. Li, "Reinforcement learning based optimal control of multiple building HVAC systems and analysis of influencing factors," *Journal of Building Engineering*, vol. 115, p. 114523, 2025.

[10] E. Bajrami, A. Kulakov, E. Zdravevski, and P. Lameski, "A comparative analysis of PPO and SAC algorithms for energy optimization with country-level energy consumption insights," *IFAC Journal of Systems and Control*, vol. 34, p. 100344, 2025.

[11] K. Qu, H. Zhang, X. Zhou, M. Ferrando, and F. Causone, "A model-based reinforcement learning framework for building heating management with branched rollout strategy and time-series prediction model," *Building Simulation*, vol. 18, no. 7, pp. 1697–1716, 2025.

[12] K. Kadamala, D. Chambers, and E. Barrett, "Improving HVAC control with transfer learning: Using padding techniques for cross-building pre-training and fine-tuning," *Energy and AI*, vol. 21, p. 100531, 2025.

[13] C. Heins, B. Millidge, D. Demekas, B. Klein, K. Friston, I. D. Couzin, and A. Tschantz, "pymdp: A Python library for active inference in discrete state spaces," *Journal of Open Source Software*, vol. 7, no. 73, p. 4098, 2022.

[14] M. Deru, K. Field, D. Studer, K. Benne, B. Griffith, P. Torcellini, B. Liu, M. Halverson, D. Winiarski, M. Rosenberg, M. Yazdanian, J. Huang, and D. Crawley, "U.s. department of energy commercial reference building models of the national building stock," National Renewable Energy Laboratory, Golden, CO, Tech. Rep. NREL/TP-5500-46861, 2011.

[15] Pacific Gas and Electric Company, "Electric schedule B-19: Medium general demand-metered TOU service," Tariff book, PG&E, 2026, accessed: Oct. 2026. [Online]. Available: https://www.pge.com/tariffs/assets/pdf/tariffbook/ELEC_SCHEDS_B-19.pdf