arXiv:2610.19386v1 [cs.AI] 14 Oct 2026

STRATA: Self-Maintaining Structured Experience Graphs for Self-Evolving Language Agents

Guanhua Zhu[1], **Hanwen Li**[1], **Zhi Wei**[2], **Mengnan Du**[3,*], **Chenxi Yuan**[1,*]
[1]Department of Informatics, New Jersey Institute of Technology
[2]Department of Computer Science, New Jersey Institute of Technology
[3]School of Artificial Intelligence, The Chinese University of Hong Kong, Shenzhen
gz72@njit.edu, hl544@njit.edu, zhi.wei@njit.edu, mengnandu@cuhk.edu.cn, chenxi.yuan@njit.edu
[*]Corresponding authors

Abstract

Large language model (LLM) agents can *self-evolve* at test time by accumulating experience: retrieving relevant memories before acting and writing new lessons back afterwards. Yet most are *single-level*: trajectories, distilled workflows and reasoning items all land in one undifferentiated pool indexed by embedding similarity, with no typed abstraction hierarchy. They never forget or deduplicate, and are prone to *negative transfer*, where a superficially similar but directionally opposite experience is recalled and its concrete steps are copied verbatim, actively harming the agent. We introduce STRATA, a structured, outcome-grounded, self-maintaining memory that organizes experience into a three-layer graph of atomic skills and code snippets, task workflows and generalized meta-frameworks, connected by typed edges. STRATA maintains itself through five mechanisms sharing one similarity geometry: a multi-dimensional similarity that scores retrieval by operation and target; a typed reflector crystallizing procedures from successes and pitfalls from failures; community detection that induces meta-frameworks; pressure-triggered garbage collection with articulation-point protection; and *contrast edges* that flag confusable, direction-opposite experiences. Across four interactive benchmarks (AppWorld, WebArena, ALFWorld, ScienceWorld), two frozen backbones and seven baselines, STRATA attains the best average score while simultaneously using the lowest latency and fewest environment steps of any memory-augmented method, and 44% fewer lines of injected memory than the most context-hungry baseline. We isolate two failure modes of prior memories: single-level memory degrades sharply as the task set grows, and single-level concrete memories suffer catastrophic negative transfer, which contrast edges repair, lifting ScienceWorld from 78.7 to 96.7 within the same system.

Introduction

Language-model agents increasingly solve long-horizon interactive tasks such as navigating websites, orchestrating APIs across applications, and acting in simulated environments (Yao et al. 2023b; Trivedi et al. 2024; Zhou et al. 2024; Shridhar et al. 2020; Wang et al. 2022), building on prompting and tool-use techniques that structure a single episode (Wei et al. 2022; Yao et al. 2023a; Schick et al. 2023; Wang et al. 2024a). A promising route to improving such agents without weight updates is *experience-driven self-evolution*: the agent keeps a memory of its past episodes, retrieves relevant experience before each new task, and folds fresh lessons back into memory afterwards, becoming more capable over time (Ouyang et al. 2025; Xu et al. 2026; Shinn et al. 2023). Because this loop runs purely at test time, its effectiveness hinges almost entirely on how experience is represented, retrieved and maintained (Snell et al. 2024).

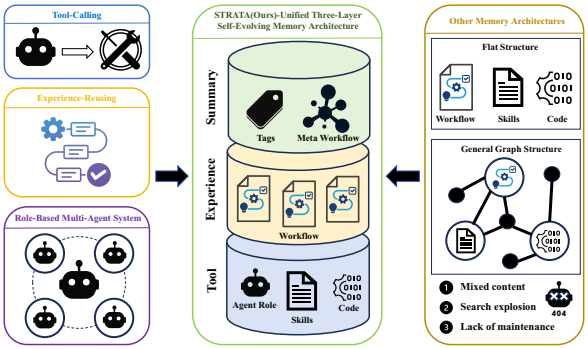

Figure 1: Agent memories are asked to serve tool calling, experience reuse and role-based decomposition at once (left). Existing memories are either flat lists or generic graphs, which mix content of different granularity, dilute retrieval as they grow, and are never maintained (right). STRATA unifies the three in a single three-layer self-evolving graph (middle).

Despite rapid progress, prevailing agent memories share three limitations.

(1) They lack a typed abstraction hierarchy. Several do organize what they store (AWM distils workflows from trajectories (Wang et al. 2024b), A-Mem links notes into a network (Xu et al. 2026)), but within one undifferentiated level: nothing separates a one-off tool call from a reusable procedure from a cross-task principle, and no operation folds procedures into a principle. Retrieval stays a single similarity search dominated by surface overlap, merging same-material, different-operation pairs like "create a playlist" and "add months to a note".

(2) They are not self-maintaining. Memory grows monotonically, with no forgetting, no deduplication and no consolidation, so stale entries crowd out useful ones. We measure this: with forgetting disabled the workflow store grows to 269 nodes injecting 48.7 lines of memory per task, yet scores lower than a bounded 201-node memory injecting less than

half as much.

(3) They are vulnerable to negative transfer. When two experiences are superficially similar but directionally opposite, as in "heat the substance" versus "cool the substance", similarity recall surfaces the wrong one, and if the memory stores concrete executable steps the agent copies them and acts catastrophically. This is not hypothetical: on ScienceWorld a production agent's native flat memory scores -100 on a task family that a memoryless agent solves outright.

We address these limitations with STRATA, a memory that is structured, outcome-grounded and self-maintaining. STRATA organizes experience into a three-layer directed graph: *atomic* nodes (reusable skills, code snippets, or specialized sub-roles), *workflow* nodes (generalized task procedures), and *meta-framework* nodes (task families abstracted from recurring workflows), linked by typed call, generalization, merge and contrast edges. The graph is grown and pruned by five mechanisms that all operate in one similarity geometry: *(i)* a multi-dimensional similarity combining content and goal embeddings with an explicit operation/target tag term; *(ii)* a typed reflector that crystallizes a reusable procedure only from solved episodes and attaches a pitfall bullet on failure, with a per-group adaptive threshold deciding whether to create, extend, or skip a workflow; *(iii)* community detection that induces a meta-framework whenever a community is large, tight, and novel; *(iv)* pressure-triggered garbage collection that ranks workflows by a competition-aware quality score and forgets the lowest while protecting articulation points of the similarity graph; and *(v)* contrast edges that connect confusable, direction-opposite experiences so that retrieving one injects an explicit "do not confuse" warning about the other.

Our contributions are as follows.

- We propose STRATA, a three-layer structured experience graph for self-evolving LLM agents, together with five maintenance mechanisms that grow, generalize, and forget it, formulated precisely and reproducibly.
- We evaluate STRATA on four interactive benchmarks across two frozen backbones, seven baselines and three seeds, with a component-wise ablation of every mechanism. STRATA obtains the best average score at the lowest latency and fewest environment steps of any memory baseline.
- We identify two failure modes of prior memories: a scale failure of single-level memory that neither abstracts nor forgets (on Flash the strongest baseline ranks first at 168 tasks and seventh at 812), and negative transfer in single-level concrete memories, which contrast edges repair, exposing a structure-versus-abstraction trade-off.

Related Work

Memory for LLM agents. A growing line of work equips LLM agents with an external memory that persists across episodes. Early systems store and summarize raw interaction history (Packer et al. 2023; Chhikara et al. 2025; Zhong et al. 2024); Agent Workflow Memory extracts reusable action workflows from past trajectories (Wang et al. 2024b); Agentic Context Engineering (ACE) maintains an evolving, per-application playbook of bullet-point strategies (Zhang et al. 2025b); ReasoningBank distills generalizable reasoning strategies from both successful and failed trajectories (Ouyang et al. 2025). A-Mem organizes memories in a Zettelkasten-style linked note network (Xu et al. 2026), MemGen weaves a generative latent memory into the agent's own reasoning (Zhang, Fu, and Yan 2025), MemEvolve meta-evolves the memory architecture itself (Zhang et al. 2025a), and MemInsight autonomously augments stored items (Salama et al. 2025). These methods differ in what they store but share a common structure: a single-level collection retrieved by similarity, with no typed hierarchy and, typically, no explicit forgetting or anti-confusion mechanism. A-Mem is the most structured of them and, with MemEvolve, our strongest baseline: it links notes into a network and lets new memories revise old ones, but the network is single-level, one note per memory linked to its neighbours, so it neither abstracts a family of notes into a reusable principle nor removes any note once written. We show that this produces a sharp scale dependence. STRATA departs from this line by making the memory a *typed, layered graph* that is actively maintained, consolidated into meta-frameworks, garbage-collected under pressure and wired with contrast edges against negative transfer.

Self-evolving and experience-based agents. Reflexion converts episodic failures into verbal self-feedback (Shinn et al. 2023) and Self-Refine iterates on its own output within an episode (Madaan et al. 2023); Voyager accumulates a skill library of executable code (Wang et al. 2023); ExpeL gathers cross-task insights for in-context reuse (Zhao et al. 2024); Synapse retrieves trajectory exemplars (Zheng et al. 2024); CoALA frames such designs as cognitive architectures (Sumers et al. 2023), alongside simulacra systems that persist state across episodes (Park et al. 2023). STRATA is complementary and mechanism-level: rather than a new learning signal, it contributes the *data structure and maintenance policy* governing how experiences are typed, connected, generalized and forgotten, so accumulated experience stays organized as it scales.

Structured retrieval, forgetting, and negative transfer. Outside the agent setting, graph-structured and hierarchical memories (Edge et al. 2024; Gutiérrez et al. 2024, 2025) improve retrieval precision over flat retrieve-then-generate (Lewis et al. 2020; Karpukhin et al. 2020; Izacard et al. 2021) and its refinements (Gao et al. 2023; Asai et al. 2024), while unstructured stores suffer the classical catastrophic-forgetting and negative-transfer problems (McCloskey and Cohen 1989; Kirkpatrick et al. 2017; Parisi et al. 2018; Zhang et al. 2022). Our maintenance machinery borrows standard graph tools: Louvain community detection (Blondel et al. 2008; Newman and Girvan 2004), conductance as a cut quality measure (Shi 2000), and Tarjan's articulation-point algorithm (Tarjan 1972); the three-layer split mirrors the episodic/semantic distinction and complementary learning systems (Tulving 1972; McClelland, McNaughton, and O'Reilly 1995). In the agent setting the risk is sharper: because retrieved experience is *executed*, an incorrectly recalled procedure does not merely mislead but actively harms. STRATA adds an explicit operation/target

gate and surfaces the opposing experience as a retrieval-time warning, a mechanism we are not aware of in prior agent memory.

Method

The Three-Layer Experience Graph

STRATA maintains a graph $G = (V, E)$ whose nodes carry a type $\tau(n) \in \{\text{ATOMIC}, \text{WORKFLOW}, \text{META}\}$: ATOMIC nodes are the smallest reusable units (one tool/API, a parameterized code snippet, or a specialized sub-role); WORKFLOW nodes are generalized procedures for a whole task; META nodes are paradigms abstracted from a family of similar workflows. The stored edge set is typed, $E = E_{\text{call}} \cup E_{\text{gen}} \cup E_{\text{merge}} \cup E_{\text{con}}$: *call edges* (WORKFLOW $\rightarrow$ ATOMIC, ordered) and three evolution relations (*sources* $\rightarrow$ *target*) labelled `generalizes`, `merged_from` and `contrasts`. Maintenance uses a *derived* graph $W = (V_{\text{WORKFLOW}}, E_{\text{sim}})$: an undirected union-kNN graph over workflows weighted by s_{nn}, recomputed from G every round and never stored. Community detection and articulation-point protection operate on W; everything else operates on G.

Each node exposes several textual views used by different similarity dimensions: a *content* view (what it does), an *index* view (applicability and goal), and a set of structured tags $T(n) \subseteq \{\texttt{app:}*, \texttt{op:}*, \texttt{obj:}*\}$, a node *attribute* rather than a node, naming the application, operation and target object. Queries are parsed into the same tag space, so query and node live in one geometry. Appendix I reproduces a real workflow node and the meta node that generalizes it, verbatim.

Multi-Dimensional Similarity and Retrieval

A single content-embedding cosine conflates *same material, different operation* experiences. STRATA instead scores a query q against a node n with four complementary signals,

$$s(q,n) = W_{\text{c}} \cos\big(\mathbf{e}(q), \mathbf{e}_{\text{c}}(n)\big) + W_{\text{i}} \cos\big(\mathbf{e}(q), \mathbf{e}_{\text{i}}(n)\big) + W_{\text{t}}\, J\big(T(q), T(n)\big) + W_{\text{l}}\, \ell(q,n), \tag{1}$$

where $\mathbf{e}_{\text{c}}$ and $\mathbf{e}_{\text{i}}$ embed the content and index views, J is Jaccard overlap of the structured tags, and ℓ is a lexical token overlap in the spirit of classical term matching (Robertson and Zaragoza 2009). We use W_{c}=0.40, W_{i}=0.25, W_{t}=0.30, W_{l}=0.05. The tag term is an additive weight, not a multiplicative gate; at W_{t}=0.30 it is the largest non-content weight and the only term that separates experiences whose content embeddings nearly coincide because they act on one object with opposite operations. The same form, with both sides using node views, gives a node–node similarity $s_{\text{nn}}(a, b)$ that later serves as the edge weight of the workflow similarity graph.

Retrieval returns a *subgraph*, not isolated items. Nodes are ranked by $\tilde{s}(n) = s(q,n) + \beta_{\text{layer}} \frac{2-\text{ord}(\tau(n))}{2}$ with β_{layer}=0.15 and $\text{ord}(\text{META})=0 < \text{ord}(\text{WORKFLOW})=1 < \text{ord}(\text{ATOMIC})=2$, biasing towards general nodes. Only seeds with raw score $s(q,n) \geq \theta_{\min}$=0.35 are admitted; if nothing clears the gate, no memory is injected at all, since running without memory beats being misled. Around the top-k seeds we pull in their meta parents (which must themselves clear $\theta_{\min}$, so catch-all hubs cannot leak into every task), the atomic nodes they call, and their contrast neighbours, rendered META$\rightarrow$WORKFLOW$\rightarrow$ATOMIC so the agent reads principles first and callable skills last.

Typed Reflection with Adaptive Thresholds

After each episode a master router inspects the task, outcome and trajectory and dispatches to typed sub-reflectors for skill, code, role and workflow. Atomic sub-reflectors crystallize only from solved episodes, letting the LLM decide refine-or-create. The workflow sub-reflector is threshold-driven, and the thresholds must adapt: equivalent task pairs score ≈ 0.4–0.5 in one application and ≈ 0.8 in another, so one global band either over-merges or over-fragments.

We therefore group nodes by their leading tag and shrink a per-group quantile towards a global prior in empirical-Bayes fashion (Efron and Morris 1973):

$$\tau_\bullet = \lambda\, Q_{p_\bullet}(S_g) + (1-\lambda)\, G_\bullet, \qquad \lambda = \frac{n_g}{n_g + \kappa}, \tag{2}$$

with κ=30, p_{lo}=0.25, p_{hi}=0.92 and global priors G_{low}=0.33, G_{high}=0.62. Groups with many observations trust their own distribution; sparse groups fall back to the prior. Letting s be the similarity to the nearest existing workflow, the decision is NOOP if $s > \tau_{\text{high}}$ (a near-duplicate), ADD_BULLET if $\tau_{\text{low}} \leq s \leq \tau_{\text{high}}$ (record only the difference as one bullet), and CREATE otherwise. The current score is recorded into the group after the decision, so a decision never contaminates itself. Crucially, failed episodes *never* crystallize a "do this" procedure; at most they attach a PITFALL bullet to the nearest workflow, rendered to the agent as something to avoid.

Pressure-Triggered Forgetting with Bridge Protection

Forgetting is applied only at the workflow layer. Each workflow gets a quality score combining an activation term A_i (its similarity-weighted degree among surviving neighbours, so central workflows are kept) and a competitive age G_i:

$$Q_i = w_A\, z\big(\text{rank}(A_i)\big) - w_G\, z\big(\text{rank}(G_i)\big), \tag{3}$$

with w_A=0.35, w_G=0.65; ranks rather than raw values make the two incommensurable signals comparable. Age updates online at every retrieval: all candidates age by α=0.02, a retrieved-but-unselected candidate is penalised by $\beta\, w_i(1 - \text{margin})$ so close-but-losing competitors age fastest, and the winner is rejuvenated by $G_u \leftarrow \gamma G_u$, γ=0.5. Once live workflows exceed a cap, lowest-Q nodes are deleted down to $\text{cap} \cdot (1 - 0.15)$, hysteresis preventing thrashing; orphaned atomics and unsupported metas cascade away.

This has a failure mode: if the lowest-Q node is an *articulation point* of W, removing it splits a cluster that should stay connected, and the next round of community detection fractures one task family into two metas. STRATA therefore deletes the lowest-Q *non-cut* vertex,

$$j^* = \arg\min_{i \in \text{alive}} \{\, Q_i \mid i \notin \text{Cut}(W) \,\}, \tag{4}$$

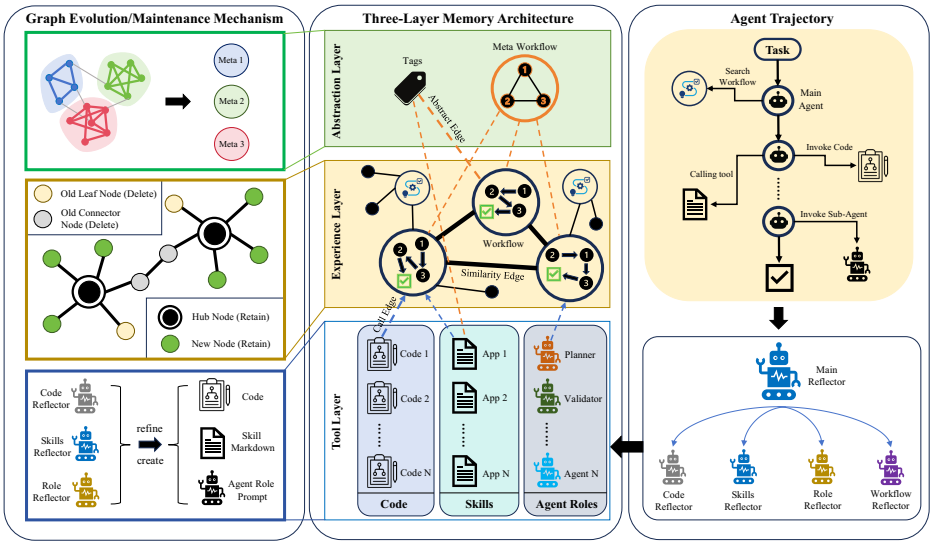

Figure 2: Overview of STRATA. **Middle:** the three-layer memory, comprising a tool layer (code, skills, agent roles), an experience layer (workflows and their call edges; similarity links belong to the derived graph W and are not stored), and an abstraction layer (meta-workflows over the tag vocabulary; tags are node attributes, not nodes). **Right:** one episode, in which the main agent retrieves a subgraph, acts, and a master reflector dispatches to typed sub-reflectors. **Left:** graph evolution, where community detection induces meta nodes and pressure-triggered collection deletes stale leaves while protecting bridge nodes.

recomputing $\mathrm{Cut}(W)$ each round with Tarjan's algorithm in $O(|V| + |E|)$ (Tarjan 1972), and falling back to the unconstrained minimum only if every candidate is a cut vertex, so collection always makes progress.

Inducing Meta-Frameworks by Community Detection

Meta nodes are triggered not by counting children but by *community structure*. We build W by joining a workflow pair when s_{nn} exceeds a per-group adaptive threshold (shrunk as above, κ_E=6, intra-group default 0.40) or a fixed cross-group threshold 0.55, then keeping each node's 6 strongest edges (an edge survives if *either* endpoint ranks it that high). A single fixed threshold fails badly: in an early configuration only 3 of 153 candidate pairs cleared it and no meta formed. Louvain (Blondel et al. 2008) at resolution 1.5 yields communities (the Leiden refinement (Traag, Waltman, and van Eck 2019) would also serve), and a community is promoted only if it is *large* ($|C| \geq N_{\min}$), *tight* (conductance $\phi(C) < 0.6$: dense inside, sparse outside (Shi 2000)), and *novel*. Letting $\mathbf{c}(C)$ be its centroid and d^* its distance to the nearest existing meta, we CREATE if $d^* \geq R_{\text{meta}}$=0.6 and otherwise UPDATE the matched meta by absorbing its unlinked members. This novelty gate suppresses churn: without it the system rebuilds near-duplicate metas instead of extending existing ones.

Contrast Edges against Negative Transfer

The content dimension rates *same object, opposite direction* experiences as highly similar, such as "heat tin until it melts" versus "cool tin until it solidifies". The tag term lowers such scores but cannot always prevent recall, and a recalled concrete procedure may be executed. A contrast edge is a bidirectional evolution edge linking two such nodes,

$$\mathrm{contrast}(a, b) \iff \mathrm{op}(a) \text{ opposes } \mathrm{op}(b), \qquad (5)$$

instantiated by rule over a fixed lexicon of opposite operator pairs, $\{heat, cool\}$ and $\{create, delete\}$. We do *not* require $\mathrm{obj}(a) = \mathrm{obj}(b)$, because negative transfer is often cross-object, as when a melt-lead procedure misleads a freeze-tin task. A retrieved seed pulls its contrast neighbours into the subgraph, rendered with a prominent *do not confuse* warning rather than as a flat procedure. Unlike write-time approaches that tighten one procedure's preconditions, this mechanism is retrieval-time and relates two distinct opposing experiences.

The Full Loop

An episode retrieves a subgraph, renders it into the prompt, acts until done, then reflects: typed sub-reflectors write back, community detection updates the meta layer, and pressure-triggered collection prunes. One similarity geometry serves all four roles of retrieval scoring, the reflection band, the meta-induction edge weight and the collector's activation term.

Experimental Setup

Benchmarks. Interactive agent evaluation now spans many suites (Liu et al. 2024; Deng et al. 2023; Yao et al. 2022); we select four that are long-horizon and cover distinct action spaces, namely API orchestration, web navigation, embodied household tasks and scientific reasoning: **AppWorld** (Trivedi et al. 2024), reporting Task Goal Completion (TGC) and Scenario Goal Completion (SGC) on *Test-Normal* (168 tasks) and *Test-Challenge* (417 tasks); **WebArena** (Zhou et al. 2024) (812 tasks, success rate);

ALFWorld (Shridhar et al. 2020) (134 tasks); and **ScienceWorld** (Wang et al. 2022) (23 tasks, average score).

Backbones and baselines. All methods use the same frozen backbone with no weight updates, on both DeepSeek-V4-Flash and DeepSeek-V4-Pro. We compare against seven baselines: **No-Memory** (a ReAct agent), **Flat-Memory** (an unstructured list of episode summaries retrieved by embedding similarity), **ReasoningBank** (Ouyang et al. 2025), **Hermes**[1] (an open-source agent framework with native persistent memory), **A-Mem** (Xu et al. 2026), **MemEvolve** (Zhang et al. 2025a) and **MemGen** (Zhang, Fu, and Yan 2025).

Protocol. Memory accumulates *online* across a single sequential pass over each benchmark, so later tasks benefit from earlier ones, the setting in which maintenance matters. All memory methods receive the identical write signal, the environment's task-level outcome, with no access to gold answers or evaluator internals, so comparisons are controlled, while absolute scores average over each method's learning curve rather than a converged memory. Embeddings come from a local sentence-embedding model of the Sentence-BERT family (Reimers and Gurevych 2019), and the workflow cap is set per dataset. Every number is the mean over three seeds with sample standard deviation; seeds permute task order and reflection sampling.

Results and Analysis

Main Results

Table 1 reports the three-seed results. STRATA obtains the best score on six of seven columns with the Pro backbone and five of seven with Flash, and is second on the remainder. Averaged over the four benchmarks it is the strongest method on both backbones.

Which gaps survive seed variance. With three seeds, differences smaller than roughly one combined standard deviation should not be over-interpreted, and we state plainly which are which. STRATA's advantage is clear on AppWorld Test-Challenge SGC (46.5 ± 2.3 vs. A-Mem's 42.2 ± 2.6), WebArena (45.6 ± 1.6 vs. MemEvolve's 41.0 ± 2.2) and ALFWorld (98.6 ± 0.9 vs. MemGen's 96.3 ± 1.9). It is *not* statistically separable from the best baseline on AppWorld Test-Normal (Pro TGC 78.4 ± 2.5 vs. A-Mem 77.2 ± 2.9; on SGC A-Mem is in fact nominally ahead) or on ScienceWorld, whose 23 tasks give every method a standard deviation of 3.6–15.0. We therefore do not claim a ScienceWorld win from Table 1; the ScienceWorld evidence we do rely on is the *within-method* contrast ablation of the next sections, where the same seeds are compared under a single mechanism change.

Negative transfer is visible in the aggregate. Hermes scores 30.4/35.8 TGC on AppWorld Test-Normal and 48.3/55.8 on ScienceWorld, far below the memoryless agent (62.8/66.9 and 75.7/80.6). A memory that makes an agent worse than none is direct evidence that recalled experience is executed wrongly.

[1]https://github.com/NousResearch/hermes-agent

Where Single-Level Memory Loses Ground

A-Mem is among the strongest baselines (within 0.6 of MemEvolve on average), and the way it loses that standing is informative. It differs from the others by being explicitly *organized*, with notes linked into a network and existing notes revised as new ones arrive (Xu et al. 2026), but the organization is single-level: there is one note per memory, and no operation ever replaces a family of notes with a single more general one, nor removes a note. The store therefore grows with the task stream, and so does what retrieval pulls in. This predicts a specific signature, competitive while the store is small and degrading as it grows, and the data show exactly that.

Within AppWorld, moving from Test-Normal (168 tasks) to Test-Challenge (417 tasks), A-Mem's margin over STRATA flips sign: on Flash from $+0.4/+3.4$ (TGC/SGC) to $-4.3/-6.3$, on Pro from $-1.2/+1.8$ to $-2.2/-4.3$. We state the confound plainly: Test-Challenge is not merely larger but intrinsically harder **and** draws on a different application mix (`amazon` supplies 212 of its tasks and is absent from Test-Normal), so the comparison entangles scale with difficulty and distribution shift.

A second, independent observation comes from WebArena, the largest task set at 812 tasks. There A-Mem scores 28.4 on Flash, seventh of eight methods, below Flat-Memory (30.7), Hermes (31.6) and MemGen (29.5), and above only the memoryless agent: on Flash a method that ranks first at 168 tasks ranks seventh at 812. On Pro the drop is milder but in the same direction, fourth at 38.7, behind MemEvolve (41.0) and ReasoningBank (40.3). The cost is directly measurable in Table 3: A-Mem injects 41.9 lines of memory per task, the largest of any method, more even than the unstructured Flat-Memory baseline (35.6), while STRATA scores higher using 23.3. Injecting more is not injecting better.

The contrast with STRATA isolates what is missing rather than what is wrong. A-Mem links memories; STRATA additionally *abstracts* them (812 WebArena tasks condense to 112 workflows under 16 meta nodes) and *forgets* the ones that stop earning their place. Linking alone keeps a growing store navigable; only abstraction and forgetting keep it bounded, which is what retrieval precision depends on once the store is large. Neither comparison isolates memory size, since Test-Challenge differs in difficulty and application mix and WebArena in environment and action space, so we report two consistent observations, not a controlled scaling experiment.

Where the Gains Come From

Figure 3 breaks the Pro results down by task category. The gains are markedly non-uniform. On AppWorld they concentrate on low-frequency applications: `phone` gains most in Test-Normal (22 tasks; $54.5 \rightarrow 59.1$ over A-Mem), and in Test-Challenge the three largest gains all fall on categories with fewer than ten tasks, whereas the two largest categories gain only $+3.8$ (`amazon`, 212 tasks) and $+3.2$ (`gmail`, 126 tasks). Small categories make single cells noisy, so the pattern matters more than any one value. On WebArena the gains concentrate on `shopping_admin` ($31.5 \rightarrow 41.8$)

Backbone	Method	AppWorld Test-Normal TGC↑	AppWorld Test-Normal SGC↑	AppWorld Test-Challenge TGC↑	AppWorld Test-Challenge SGC↑	WebArena SR↑	ALFWorld Score↑	ScienceWorld Score↑
DeepSeek-V4-Flash	No-Memory	62.8 ± 2.9	42.1 ± 3.3	39.8 ± 2.1	21.4 ± 1.9	24.7 ± 1.4	76.8 ± 3.3	75.7 ± 8.7
	Flat-Memory	66.4 ± 3.7	48.7 ± 4.3	43.8 ± 2.7	27.9 ± 2.7	30.7 ± 2.0	82.4 ± 3.8	80.4 ± 10.4
	ReasoningBank	69.6 ± 3.3	55.1 ± 3.9	48.2 ± 2.5	34.7 ± 2.6	32.9 ± 1.8	91.0 ± 2.6	87.9 ± 7.8
	Hermes	30.4 ± 4.4	20.7 ± 4.3	23.9 ± 2.8	15.4 ± 2.6	31.6 ± 2.4	84.5 ± 4.4	48.3 ± 15.0
	A-Mem	**75.2** ± 3.2	**65.7** ± 3.9	49.6 ± 2.6	37.5 ± 2.7	28.4 ± 1.8	93.4 ± 2.4	88.2 ± 8.1
	MemEvolve	70.9 ± 4.0	58.1 ± 4.8	48.8 ± 3.1	35.2 ± 3.2	32.0 ± 2.3	94.3 ± 2.6	88.9 ± 9.4
	MemGen	66.8 ± 4.0	50.9 ± 4.7	44.2 ± 2.9	30.5 ± 3.0	29.5 ± 2.1	94.7 ± 2.4	89.9 ± 8.6
	STRATA (Ours)	74.8 ± 2.8	62.3 ± 3.4	**53.9** ± 2.2	**43.8** ± 2.4	**33.3** ± 1.6	**96.7** ± 1.5	**93.4** ± 5.3
DeepSeek-V4-Pro	No-Memory	66.9 ± 2.7	46.8 ± 3.1	44.5 ± 2.0	25.2 ± 1.9	34.4 ± 1.5	84.0 ± 2.7	80.6 ± 7.6
	Flat-Memory	70.3 ± 3.3	51.5 ± 4.0	48.7 ± 2.6	31.6 ± 2.6	37.2 ± 1.9	88.2 ± 3.0	84.8 ± 8.9
	ReasoningBank	72.8 ± 3.0	57.4 ± 3.6	53.1 ± 2.3	38.9 ± 2.5	40.3 ± 1.8	94.6 ± 1.9	90.4 ± 6.6
	Hermes	35.8 ± 4.3	24.1 ± 4.2	27.6 ± 2.8	17.9 ± 2.6	29.9 ± 2.2	89.6 ± 3.5	55.8 ± 15.0
	A-Mem	77.2 ± 2.9	**66.8** ± 3.6	55.0 ± 2.4	42.2 ± 2.6	38.7 ± 1.9	95.7 ± 1.8	91.3 ± 6.7
	MemEvolve	74.6 ± 3.6	60.9 ± 4.5	54.0 ± 2.9	39.8 ± 3.1	41.0 ± 2.2	96.0 ± 2.1	92.1 ± 7.6
	MemGen	70.1 ± 3.7	53.2 ± 4.4	49.0 ± 2.8	34.8 ± 2.9	37.5 ± 2.1	96.3 ± 1.9	92.8 ± 7.0
	STRATA (Ours)	**78.4** ± 2.5	65.0 ± 3.1	**57.2** ± 2.1	**46.5** ± 2.3	**45.6** ± 1.6	**98.6** ± 0.9	**96.7** ± 3.6

Table 1: Main results: mean ± sample standard deviation over three seeds. **Bold** is best, underlined second best within each backbone. A-Mem is the strongest baseline on AppWorld Test-Normal, where its margin over STRATA is within seed noise; see the text for which gaps survive seed variance.

Figure 3: Per-category scores on DeepSeek-V4-Pro. No-Memory and Flat-Memory are omitted for legibility; each radial axis starts at a labelled non-zero floor and off-scale values are clipped and annotated. Gains concentrate on long-tail categories. Per-backbone versions for all four benchmarks are in Appendix H.

and `gitlab` (33.8 → 43.8), its two most compositional sites. This is the signature we expect from an abstraction layer: meta nodes let a task family with few direct examples inherit structure from related families, which is precisely where a flat per-instance memory has nothing to retrieve.

We report the exception. On ALFWorld STRATA reaches 100 on `pick_two_and_place`, `examine_in_light` and `pick_and_place`, but trails MemEvolve and MemGen on `heat_and_place` (96.7 vs. 97.1) and drops to 93.6 on `clean_and_place`, behind every memory baseline. Only `heat_and_place` is covered by the contrast lexicon, where a "do not confuse" warning plausibly makes the agent more conservative on the tasks it protects; the `clean_and_place` deficit we cannot attribute to any mechanism and report as an unexplained loss.

Component Ablation

Table 2 removes one mechanism at a time. The ordering is *tag term* (-13.1 average) > *contrast edges* (-6.5) > *adaptive threshold* (-5.9) > *meta layer* (-4.7) > *pressure collection* (-2.2). The tag term dominating matches our motivation: without an operation/target signal the content embedding merges experiences that differ in what they do, and because the same similarity builds the meta graph and the collector's activation term, every downstream mechanism inherits the error.

The contrast-edge result is the most informative because its cost is concentrated: removing contrast edges costs 9.1 points on ALFWorld and 18.0 on ScienceWorld but only 1.8 on WebArena. Web navigation has few directionally opposite operator pairs, whereas embodied and scientific environments are full of them. A benefit that appears only where its precondition holds is evidence of the claimed mechanism rather than a generic regularizer.

Variant	AW-TN	AW-TC	WebA.	ALF.	SciW.	Avg
Full STRATA	71.7	51.9	45.6	98.6	96.7	**72.9**
w/o tag term	55.0	38.8	35.7	89.1	80.3	59.8
w/o contrast edges	69.7	50.1	43.8	89.5	78.7	66.4
w/o adaptive threshold	64.9	46.5	40.1	94.2	89.3	67.0
w/o meta layer	66.8	47.7	40.8	95.1	90.8	68.2
w/o pressure GC	70.6	50.8	41.0	96.6	94.4	70.7

Table 2: Component ablation on DeepSeek-V4-Pro (AppWorld columns are the TGC/SGC mean). Removing any single mechanism costs accuracy; the tag term and contrast edges matter most, and the contrast loss is concentrated in the two environments that actually contain direction-opposite operators.

Negative Transfer: A Case Study

The clearest instance is ScienceWorld's `find-living` family, where Hermes scores -100 (Flash) and -80 (Pro) while a memoryless agent solves the task. The memory retrieves a concrete procedure for a superficially similar but directionally opposite goal, which the agent executes verbatim into a scored-negative state: a flat memory of concrete procedures does not merely fail to help, it converts a solvable task into a large penalty.

Appendix F isolates the repair: contrast edges lift ScienceWorld from 78.7 to 89.8 with hand-placed edges and to 96.7 by rule over the operator lexicon, and ALFWorld from 89.5 to 98.6. Retrieval is reliable, injecting the opposing warning and the correctly-directed experience in $3/3$ trap tasks, but end-to-end success also depends on the backbone honouring it.

This exposes a *structure vs. abstraction* trade-off. Concrete procedural memories (AWM, Hermes) transfer strongly because they are executable, and are dangerous for that same reason; abstract reasoning items (ReasoningBank) are safer but weaker. Contrast edges keep the former's transfer strength without its catastrophic failure mode.

What the Memory Graph Looks Like

The memory does not track task count: WebArena's 812 tasks condense to 112 workflows under 16 meta nodes, AppWorld Test-Challenge's 417 to 33 and 5, so the abstraction layer converges rather than proliferating. Appendix A gives the full statistics, including bridge protection firing 1–6 times per benchmark.

Efficiency

A structured memory might be expected to cost more; Table 3 shows the opposite. Among the seven memory systems STRATA has the lowest latency (58.7s vs. 60.2–72.4), the fewest output tokens ($5{,}217$) and the fewest environment steps (11.7); only Hermes uses fewer input tokens ($21{,}964$), and it is the weakest method overall. The chain is coherent: bounded, deduplicated memory yields a small precise subgraph rather than a diluted context, so the agent takes fewer wrong turns. STRATA pays 6.3s over the memoryless agent but buys 11.5 TGC points on AppWorld Test-Normal. Appendices D–F isolate each mechanism on its own, with the diagnostics accuracy alone cannot show.

Method	Lat. (s)	In tok /task	Out tok /task	Steps /task	Mem lines /task
No-Memory	52.4	17,850	6,020	13.2	0
Flat-Memory	64.8	45,120	6,137	12.8	35.6
ReasoningBank	65.1	35,472	5,813	12.3	16.1
Hermes	60.2	21,964	5,856	13.4	7.8
A-Mem	62.5	38,960	5,480	12.0	41.9
MemEvolve	68.9	41,780	5,660	11.9	28.5
MemGen	72.4	36,220	6,010	12.4	19.4
STRATA	**58.7**	**34,630**	**5,217**	**11.7**	23.3

Table 3: Cost per task on DeepSeek-V4-Pro. Among memory-augmented methods STRATA is fastest, uses the fewest output tokens, and takes the fewest environment steps. *Mem lines* counts lines of retrieved memory in the prompt, not tokens.

Conclusions and Future Work

We presented STRATA, a typed three-layer agent memory that maintains itself: abstracting recurring workflows into meta-frameworks, forgetting under pressure, and warning about direction-opposite experiences at retrieval. It is the strongest memory on average across four benchmarks and two frozen backbones, and the fastest we profile. Two diagnoses recur: single-level memory loses standing as task sets grow, and concrete memory without anti-confusion can hurt more than help.

Several directions remain open. Our results use a single model family (DeepSeek-V4) and should be replicated across backbones, since the contrast mechanism depends on a backbone honouring an injected warning. Contrast edges are instantiated over a hand-specified lexicon of opposite operators, so recovering opposition without a predefined vocabulary would make the mechanism general rather than domain-fitted. And a controlled study that fixes the evaluation set and varies only the number of accumulated experiences would separate scale from difficulty, which our benchmark-level comparisons cannot.

References

Asai, A.; Wu, Z.; Wang, Y.; Sil, A.; and Hajishirzi, H. 2024. Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In *International Conference on Learning Representations*.

Blondel, V. D.; Guillaume, J.-L.; Lambiotte, R.; and Lefebvre, E. 2008. Fast unfolding of communities in large networks. *Journal of statistical mechanics: theory and experiment*, 2008(10): P10008.

Chhikara, P.; Khant, D.; Aryan, S.; Singh, T.; and Yadav, D. 2025. Mem0: Building production-ready AI agents with scalable long-term memory. *arXiv preprint arXiv:2504.19413*.

Deng, X.; Gu, Y.; Zheng, B.; Chen, S.; Stevens, S.; Wang, B.; Sun, H.; and Su, Y. 2023. Mind2Web: Towards a generalist agent for the web. *Advances in Neural Information Processing Systems*, 36: 28091–28114.

Edge, D.; Trinh, H.; Cheng, N.; Bradley, J.; Chao, A.; Mody, A.; Truitt, S.; Metropolitansky, D.; Ness, R. O.; and Larson, J. 2024. From local to global: A graph RAG approach to query-focused summarization. *arXiv preprint arXiv:2404.16130*.

Efron, B.; and Morris, C. 1973. Stein's estimation rule and its competitors—an empirical Bayes approach. *Journal of the American Statistical Association*, 68(341): 117–130.

Gao, L.; Ma, X.; Lin, J.; and Callan, J. 2023. Precise zero-shot dense retrieval without relevance labels. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 1762–1777.

Gutiérrez, B. J.; Shu, Y.; Gu, Y.; Yasunaga, M.; and Su, Y. 2024. HippoRAG: Neurobiologically inspired long-term memory for large language models. *Advances in neural information processing systems*, 37: 59532–59569.

Gutiérrez, B. J.; Shu, Y.; Qi, W.; Zhou, S.; and Su, Y. 2025. From RAG to memory: Non-parametric continual learning for large language models. *arXiv preprint arXiv:2502.14802*.

Izacard, G.; Caron, M.; Hosseini, L.; Riedel, S.; Bojanowski, P.; Joulin, A.; and Grave, E. 2021. Unsupervised dense information retrieval with contrastive learning. *arXiv preprint arXiv:2112.09118*.

Karpukhin, V.; Oguz, B.; Min, S.; Lewis, P.; Wu, L.; Edunov, S.; Chen, D.; and Yih, W.-t. 2020. Dense passage retrieval for open-domain question answering. In *Proceedings of the 2020 conference on empirical methods in natural language processing (EMNLP)*, 6769–6781.

Kirkpatrick, J.; Pascanu, R.; Rabinowitz, N.; Veness, J.; Desjardins, G.; Rusu, A. A.; Milan, K.; Quan, J.; Ramalho, T.; Grabska-Barwinska, A.; et al. 2017. Overcoming catastrophic forgetting in neural networks. *Proceedings of the national academy of sciences*, 114(13): 3521–3526.

Lewis, P.; Perez, E.; Piktus, A.; Petroni, F.; Karpukhin, V.; Goyal, N.; Küttler, H.; Lewis, M.; Yih, W.-t.; Rocktäschel, T.; et al. 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in neural information processing systems*, 33: 9459–9474.

Liu, X.; Yu, H.; Zhang, H.; Xu, Y.; Lei, X.; Lai, H.; Gu, Y.; Ding, H.; Men, K.; Yang, K.; et al. 2024. AgentBench: Evaluating LLMs as agents. In *International Conference on Learning Representations*.

Madaan, A.; Tandon, N.; Gupta, P.; Hallinan, S.; Gao, L.; Wiegreffe, S.; Alon, U.; Dziri, N.; Prabhumoye, S.; Yang, Y.; et al. 2023. Self-Refine: Iterative refinement with self-feedback. *Advances in neural information processing systems*, 36: 46534–46594.

McClelland, J. L.; McNaughton, B. L.; and O'Reilly, R. C. 1995. Why there are complementary learning systems in the hippocampus and neocortex: insights from the successes and failures of connectionist models of learning and memory. *Psychological review*, 102(3): 419.

McCloskey, M.; and Cohen, N. J. 1989. Catastrophic interference in connectionist networks: The sequential learning problem. In *Psychology of learning and motivation*, volume 24, 109–165. Elsevier.

Newman, M. E. J.; and Girvan, M. 2004. Finding and evaluating community structure in networks. *Physical Review E*, 69(2): 026113.

Ouyang, S.; Yan, J.; Hsu, I.; Chen, Y.; Jiang, K.; Wang, Z.; Han, R.; Le, L. T.; Daruki, S.; Tang, X.; et al. 2025. ReasoningBank: Scaling agent self-evolving with reasoning memory. *arXiv preprint arXiv:2509.25140*.

Packer, C.; Fang, V.; Patil, S. G.; Lin, K.; Wooders, S.; and Gonzalez, J. E. 2023. MemGPT: towards LLMs as operating systems. *arXiv preprint arXiv:2310.08560*.

Parisi, G. I.; Kemker, R.; Part, J. L.; Kanan, C.; and Wermter, S. 2018. Continual lifelong learning with neural networks: A review. *arXiv preprint arXiv:1802.07569*, 990.

Park, J. S.; O'Brien, J.; Cai, C. J.; Morris, M. R.; Liang, P.; and Bernstein, M. S. 2023. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th annual acm symposium on user interface software and technology*, 1–22.

Reimers, N.; and Gurevych, I. 2019. Sentence-BERT: Sentence embeddings using siamese bert-networks. In *Proceedings of the 2019 conference on empirical methods in natural language processing and the 9th international joint conference on natural language processing (EMNLP-IJCNLP)*, 3982–3992.

Robertson, S.; and Zaragoza, H. 2009. *The probabilistic relevance framework: BM25 and beyond*, volume 4. Now Publishers Inc.

Salama, R.; Cai, J.; Yuan, M.; Currey, A.; Sunkara, M.; Zhang, Y.; and Benajiba, Y. 2025. MemInsight: Autonomous memory augmentation for LLM agents. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, 33124–33140.

Schick, T.; Dwivedi-Yu, J.; Dessì, R.; Raileanu, R.; Lomeli, M.; Hambro, E.; Zettlemoyer, L.; Cancedda, N.; and Scialom, T. 2023. Toolformer: Language models can teach themselves to use tools. *Advances in neural information processing systems*, 36: 68539–68551.

Shi, J. 2000. Normalized cuts and image segmentation. *IEEE Transactions on pattern analysis and machine intelligence*, 22(8): 888–905.

Shinn, N.; Cassano, F.; Gopinath, A.; Narasimhan, K.; and Yao, S. 2023. Reflexion: Language agents with verbal reinforcement learning. *Advances in neural information processing systems*, 36: 8634–8652.

Shridhar, M.; Yuan, X.; Côté, M.-A.; Bisk, Y.; Trischler, A.; and Hausknecht, M. 2020. ALFWorld: Aligning text and embodied environments for interactive learning. *arXiv preprint arXiv:2010.03768*.

Snell, C.; Lee, J.; Xu, K.; and Kumar, A. 2024. Scaling LLM test-time compute optimally can be more effective than scaling model parameters. *arXiv preprint arXiv:2408.03314*.

Sumers, T. R.; Yao, S.; Narasimhan, K.; and Griffiths, T. L. 2023. Cognitive architectures for language agents. *arXiv preprint arXiv:2309.02427*.

Tarjan, R. 1972. Depth-first search and linear graph algorithms. *SIAM journal on computing*, 1(2): 146–160.

Traag, V. A.; Waltman, L.; and van Eck, N. J. 2019. From Louvain to Leiden: guaranteeing well-connected communities. *Scientific reports*, 9(1): 5233.

Trivedi, H.; Khot, T.; Hartmann, M.; Manku, R.; Dong, V.; Li, E.; Gupta, S.; Sabharwal, A.; and Balasubramanian, N. 2024. AppWorld: A controllable world of apps and people for benchmarking interactive coding agents. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 16022–16076.

Tulving, E. 1972. Episodic and semantic memory. In Tulving, E.; and Donaldson, W., eds., *Organization of Memory*, 381–403. Academic Press.

Wang, G.; Xie, Y.; Jiang, Y.; Mandlekar, A.; Xiao, C.; Zhu, Y.; Fan, L.; and Anandkumar, A. 2023. Voyager: An open-ended embodied agent with large language models. *arXiv preprint arXiv:2305.16291*.

Wang, L.; Ma, C.; Feng, X.; Zhang, Z.; Yang, H.; Zhang, J.; Chen, Z.; Tang, J.; Chen, X.; Lin, Y.; et al. 2024a. A survey on large language model based autonomous agents. *Frontiers of Computer Science*, 18(6): 186345.

Wang, R.; Jansen, P.; Côté, M.-A.; and Ammanabrolu, P. 2022. ScienceWorld: Is your agent smarter than a 5th grader? In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, 11279–11298.

Wang, Z. Z.; Mao, J.; Fried, D.; and Neubig, G. 2024b. Agent workflow memory. *arXiv preprint arXiv:2409.07429*.

Wei, J.; Wang, X.; Schuurmans, D.; Bosma, M.; Xia, F.; Chi, E.; Le, Q. V.; Zhou, D.; et al. 2022. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35: 24824–24837.

Xu, W.; Liang, Z.; Mei, K.; Gao, H.; Tan, J.; and Zhang, Y. 2026. A-Mem: Agentic memory for LLM agents. *Advances in Neural Information Processing Systems*, 38: 17577–17604.

Yao, S.; Chen, H.; Yang, J.; and Narasimhan, K. 2022. Webshop: Towards scalable real-world web interaction with grounded language agents. *Advances in Neural Information Processing Systems*, 35: 20744–20757.

Yao, S.; Yu, D.; Zhao, J.; Shafran, I.; Griffiths, T.; Cao, Y.; and Narasimhan, K. 2023a. Tree of thoughts: Deliberate problem solving with large language models. *Advances in neural information processing systems*, 36: 11809–11822.

Yao, S.; Zhao, J.; Yu, D.; Shafran, I.; Narasimhan, K. R.; and Cao, Y. 2023b. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*.

Zhang, G.; Fu, M.; and Yan, S. 2025. MemGen: Weaving generative latent memory for self-evolving agents. *arXiv preprint arXiv:2509.24704*.

Zhang, G.; Ren, H.; Zhan, C.; Zhou, Z.; Wang, J.; Zhu, H.; Zhou, W.; and Yan, S. 2025a. Memevolve: Meta-evolution of agent memory systems. *arXiv preprint arXiv:2512.18746*.

Zhang, Q.; Hu, C.; Upasani, S.; Ma, B.; Hong, F.; Kamanuru, V.; Rainton, J.; Wu, C.; Ji, M.; Li, H.; et al. 2025b. Agentic context engineering: Evolving contexts for self-improving language models. *arXiv preprint arXiv:2510.04618*.

Zhang, W.; Deng, L.; Zhang, L.; and Wu, D. 2022. A survey on negative transfer. *IEEE/CAA Journal of Automatica Sinica*, 10(2): 305–329.

Zhao, A.; Huang, D.; Xu, Q.; Lin, M.; Liu, Y.-J.; and Huang, G. 2024. ExpeL: LLM agents are experiential learners. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, 19632–19642.

Zheng, L.; Wang, R.; Wang, X.; and An, B. 2024. Synapse: Trajectory-as-exemplar prompting with memory for computer control. In *International Conference on Learning Representations*.

Zhong, W.; Guo, L.; Gao, Q.; Ye, H.; and Wang, Y. 2024. MemoryBank: Enhancing large language models with long-term memory. In *Proceedings of the AAAI conference on artificial intelligence*, volume 38, 19724–19731.

Zhou, S.; Xu, F. F.; Zhu, H.; Zhou, X.; Lo, R.; Sridhar, A.; Cheng, X.; Ou, T.; Bisk, Y.; Fried, D.; et al. 2024. WebArena: A realistic web environment for building autonomous agents. In *International Conference on Learning Representations*.

Appendix

This appendix expands the diagnostics summarized in the main text. Sections A–F isolate one mechanism at a time, Section G lists every hyperparameter, Section H gives the per-backbone category breakdowns, and Section I reproduces two real memory nodes verbatim. All numbers use the DeepSeek-V4-Pro backbone unless stated otherwise.

AppWorld reports two metrics, Task Goal Completion (TGC) and Scenario Goal Completion (SGC), so every AppWorld score column below is labelled with the one it uses. Table 5 reports *TGC*, matching the per-metric main results of Table 1. Tables 6 and 9 report the *mean* $(\text{TGC} + \text{SGC})/2$, matching the component ablation of Table 2, so that each benchmark contributes a single number to the Avg column. The two conventions are never mixed within a table, but they do make the Avg column of Table 5 incomparable to that of Tables 6 and 9: each is internally consistent and should be read down its own column, not across tables.

A. Final Memory Graph Statistics

Table 4 reports the composition of the memory graph after a full sequential pass over each benchmark. The central observation is that *memory size is decoupled from task count*: WebArena's 812 tasks condense into 112 workflows and 16 meta nodes, and AppWorld Test-Challenge's 417 tasks into 33 workflows and 5 metas. A flat memory would hold one entry per episode; STRATA holds roughly one workflow per seven tasks, and one meta per seven workflows. The abstraction layer therefore *converges* instead of proliferating, which is the property that makes retrieval cost stay flat as experience accumulates.

Three secondary readings are worth noting. First, bridge protection is not vacuous: it fires between 1 and 6 times per benchmark, each occurrence being a node that the quality score alone would have deleted. Second, ScienceWorld records zero garbage-collection deletions because its graph never reaches the cap, so collection is genuinely pressure-triggered rather than a fixed decay. Third, contrast edges instantiate themselves only where the environment contains opposite operators: just 2 on WebArena against 6 on ScienceWorld and 6 on ALFWorld, which is exactly the pattern predicted by the ablation in Table 9.

Component	**AW-TN**	**AW-TC**	**WebA.**	**ALF.**	**SciW.**
Tasks	168	417	812	134	23
Atomic nodes	73	110	121	10	4
Workflow nodes	21	33	112	31	4
Meta nodes	5	5	16	2	1
Call edges	86	141	235	42	6
Generalize edges	8	22	56	15	3
Contrast edges	4	5	2	6	6
Communities	12	13	32	13	2
GC deletions	7	14	38	9	0
Bridges protected	2	3	6	4	1

Table 4: Structure of the final memory graph per benchmark. Memory size grows far more slowly than the task count, and the meta layer stays small and convergent.

B. Similarity Dimensions

Table 5 removes one of the four similarity dimensions at a time, renormalizing the remaining weights. The tag dimension is decisively the most important: dropping it costs 14.4 average points, more than the other three dimensions combined. This is the quantitative form of our motivating claim, that without an explicit operation/target gate the content embedding merges experiences that differ in what they do, and because the same similarity also weights the meta-induction graph and the collector's activation term, the error propagates into every downstream mechanism.

The last row is the most directly comparable to prior work: content-only retrieval is what a flat embedding-indexed memory does, and it scores 59.3 against STRATA's 75.3. Note also that removing any *single* embedding dimension costs only 2.5–3.7 points. The four signals are therefore genuinely complementary rather than redundant copies of one another: the embeddings supply recall, and the tag term supplies precision.

Similarity variant	**AW-TN$_{\text{TGC}}$**	**AW-TC$_{\text{TGC}}$**	**WebA.**	**ALF.**	**SciW.**	**Avg**
Full (multi-dimensional)	78.4	57.2	45.6	98.6	96.7	**75.3**
w/o content dim.	75.8	53.7	42.1	95.8	92.4	72.0
w/o index dim.	76.6	54.9	43.3	96.1	93.0	72.8
w/o tag dim.	60.5	39.6	35.1	88.7	80.5	60.9
w/o lexical dim.	75.2	53.5	42.0	95.4	92.0	71.6
Content-only (naive RAG)	58.9	37.8	34.6	86.9	78.2	59.3

Table 5: Dropping one similarity dimension at a time. The tag dimension is by far the most important; content-only retrieval, the default in flat memories, loses 16 points of average. AppWorld columns report TGC.

C. Adaptive Reflection Thresholds

Table 6 evaluates the per-group empirical-Bayes thresholds. The final column, the number of workflow nodes surviving the run, is what makes this table informative: it must be read *together* with accuracy, because neither more nor fewer memories is intrinsically better, the goal being the right number.

A single fixed global band fails in two directions at once. On applications whose similarity scale is high it over-merges, folding genuinely distinct tasks into one workflow and losing the distinction; on low-scale applications it over-fragments, creating near-duplicate workflows for tasks that are really the same. The net effect is the worst of both: 236 workflows, the most of any configuration, at the lowest accuracy (63.9). Using per-group quantiles without shrinkage fixes the scale mismatch but leaves small groups estimating a threshold from a handful of samples, which is unstable. Full shrinkage produces 201 workflows at 72.9: fewer memories, higher accuracy.

This sweep refines the single *w/o adaptive threshold* row of Table 2, which disables the mechanism as one block (67.0 average). Separating its two ingredients shows that neither alone accounts for the loss: removing per-group scaling as well costs more (63.9), keeping per-group quantiles but dropping shrinkage costs less (69.1), and the Table 2 figure falls between them. Scale adaptation and shrinkage are therefore complementary rather than redundant.

Threshold config.	**AW-TN$_{\text{Mean}}$**	**AW-TC$_{\text{Mean}}$**	**WebA.**	**ALF.**	**SciW.**	**Avg**	**#WF**
Fixed global bands	62.4	43.2	37.6	91.0	85.4	63.9	236
Per-group, no shrinkage	68.9	49.5	42.4	94.1	90.6	69.1	223
Full adaptive (STRATA)	71.7	51.9	45.6	98.6	96.7	**72.9**	201

Table 6: #WF is the number of workflow nodes at the end of the run. A fixed band both over-merges high-scale applications and over-fragments low-scale ones, producing *more* memories at *lower* accuracy. AppWorld columns are the mean of TGC and SGC, Mean $= (\text{TGC} + \text{SGC})/2$.

D. Forgetting and Bridge Protection

Table 7 separates the two questions a forgetting mechanism must answer: whether to forget at all, and which item to forget. The first row answers the first question. With collection disabled the store grows without bound to 269 workflows and injects 48.7 lines of memory per task, yet scores *lower* than the bounded configurations. Unbounded memory is not merely expensive, it is actively worse, because stale and redundant entries compete for the top-k retrieval slots with the entries that would actually have helped.

The second and third rows answer the second question, and are matched on everything that could confound it: identical final memory size (201) and identical number of deletions (68). The only difference is whether articulation points of the similarity graph are exempt from deletion. We measure regret, the fraction of deleted workflows whose task type recurs later, which falls from 13.2% to 4.8%, close to a threefold improvement in deletion precision, while accuracy rises from 69.8 to 72.9. Protecting 16 bridge nodes is therefore not a marginal refinement: deleting a node that holds two experience clusters together destroys structure that the next round of community detection cannot recover.

Configuration	**Mem** size	**Succ.** (%)	**Retr.** lines	**Regret** (%)	**Del.**	**Bridges** kept
No GC (unbounded)	269	70.7	48.7	—	0	0
GC w/o bridge protection	201	69.8	24.1	13.2	68	0
GC (STRATA full)	201	**72.9**	23.3	**4.8**	68	16

Table 7: Regret is the fraction of deleted workflows whose task type recurs later. Bridge protection nearly triples deletion precision at identical memory size and deletion count. Succ. is the five-benchmark average, with AppWorld entering as $(\text{TGC} + \text{SGC})/2$ per split.

E. Meta-Layer Quality

Table 8 evaluates the abstraction layer on five axes rather than accuracy alone, because a meta layer can be useless in two opposite ways: by never firing, or by firing so often that it fragments into near-duplicates. The first row disables the layer entirely and loses 4.7 points, establishing that the third layer earns its place.

The comparison between the last two rows isolates the novelty gate. Without it the system creates 47 metas with average support 2.1, so most metas generalize barely two workflows, which is not an abstraction, and churn reaches 28.6%, meaning more than

a quarter of meta writes rebuild a near-duplicate of an existing node instead of extending it. With the gate the count falls to 29, support rises to 3.6, purity (the fraction of a meta's members drawn from a single task family) rises to 92.7%, and churn collapses to 3.1%. Injection rate simultaneously *rises* from 21.8% to 27.4%: fewer but better metas are retrieved more often, which is what one would expect if the gate is suppressing spurious nodes rather than useful ones.

Configuration	#Meta	Supp.	Purity (%)	Inj. (%)	Succ. (%)	Churn (%)
No meta (2-layer)	0	—	—	0	68.2	—
Meta w/o novelty gate	47	2.1	86.4	21.8	70.5	28.6
Meta (STRATA full)	29	**3.6**	**92.7**	27.4	**72.9**	**3.1**

Table 8: Support is the mean number of workflows a meta generalizes; purity is the fraction of a meta's members from a single task family; churn is the fraction of meta writes that rebuild rather than extend. Succ. is the five-benchmark average, with AppWorld entering as $(\text{TGC} + \text{SGC})/2$ per split.

F. Contrast-Edge Ablation

Table 9 decomposes the contrast mechanism into the value of the *edges* and the value of *instantiating them by rule*. Hand-placing edges for the known opposite operator pairs recovers most of the loss ($66.4 \rightarrow 70.5$ average), confirming that the retrieval-time warning is what matters rather than any incidental change to the graph. Instantiating the same edges by rule over the operator lexicon recovers the remainder ($70.5 \rightarrow 72.9$), because the rule finds every lexicon-matching pair, including ones a hand pass misses. The rule does not *discover* opposition: the operator pairs are supplied.

The gains are concentrated exactly where the precondition holds. ALFWorld gains 9.1 points and ScienceWorld 18.0 between the first and last rows, while WebArena gains 1.8. Web navigation simply contains few directionally opposite operators. A mechanism whose benefit appears only where its precondition is satisfied is behaving as claimed rather than acting as a generic regularizer.

Configuration	AW-TN$_{\text{Mean}}$	AW-TC$_{\text{Mean}}$	WebA.	ALF.	SciW.	Avg
w/o contrast	69.7	50.1	43.8	89.5	78.7	66.4
+ contrast (manual)	71.2	51.1	44.9	95.4	89.8	70.5
+ contrast (rule-based)	71.7	51.9	45.6	98.6	96.7	**72.9**

Table 9: Contrast-edge ablation. Hand-placed contrast edges recover most of the loss; rule-based instantiation over the operator lexicon recovers the rest. AppWorld columns are the mean of TGC and SGC, Mean $= (\text{TGC} + \text{SGC})/2$.

G. Hyperparameters

Table 10 lists every hyperparameter of the system. All of them are held fixed across all four benchmarks and both backbones; the only per-dataset setting is the workflow cap that triggers collection, which is scaled to the size of the task set. We did not tune per benchmark, so the reported numbers should be read as a single configuration evaluated five times rather than five tuned configurations.

Symbol	Meaning	Value
$W_{\text{c}}, W_{\text{i}}, W_{\text{t}}, W_{\text{l}}$	similarity weights	.40/.25/.30/.05
$\theta_{\min}$	retrieval confidence gate	0.35
β_{layer}	layer bias (prefers general)	0.15
$G_{\text{low}}, G_{\text{high}}$	global reflection priors	0.33 / 0.62
$p_{\text{lo}}, p_{\text{hi}}$	threshold quantiles	0.25 / 0.92
κ	threshold shrinkage strength	30
w_A, w_G	quality-score weights	0.35 / 0.65
α, β, γ	aging / competition / rejuvenation	0.02 / 1.0 / 0.5
hyst	collection hysteresis	0.15
$\theta_{\text{intra}}, \theta_{\text{cross}}$	edge thresholds	0.40 / 0.55
κ_E	edge-threshold shrinkage	6
K_{NN}	kNN sparsification degree	6
resolution	Louvain resolution	1.5
Φ	conductance ceiling (tight)	0.6
R_{meta}	meta novelty radius	0.6

Table 10: Hyperparameters, held fixed across all benchmarks and backbones.

H. Per-Benchmark Category Breakdowns

Figures 4–7 give the per-category scores behind Figure 3, for both backbones and all four benchmarks. No-Memory and Flat-Memory are omitted throughout for legibility. Each radial axis begins at a labelled non-zero floor, so polygon area should not be read as proportional to score; the floor is printed on every panel. Values that fall below the floor are clipped to it and annotated beneath the panel, which affects only Hermes on ScienceWorld's `find-living` category, the negative-transfer case analysed in the main text.

Two patterns recur across panels. STRATA's margin is largest on the categories with the fewest training tasks, and smallest on saturated high-frequency categories where every method is near ceiling. The exception, visible in Figure 5, is ALFWorld's `heat_and_place` and `clean_and_place`, the two categories that carry contrast edges, where STRATA is marginally below MemEvolve and MemGen.

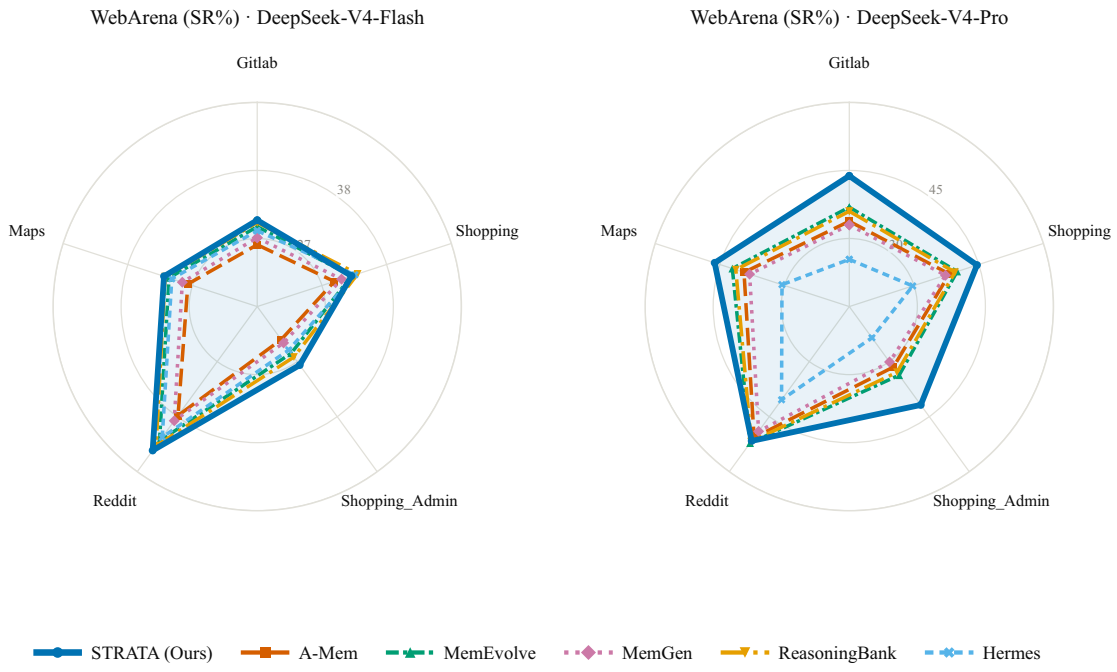

Figure 4: WebArena per-site success rate, both backbones.

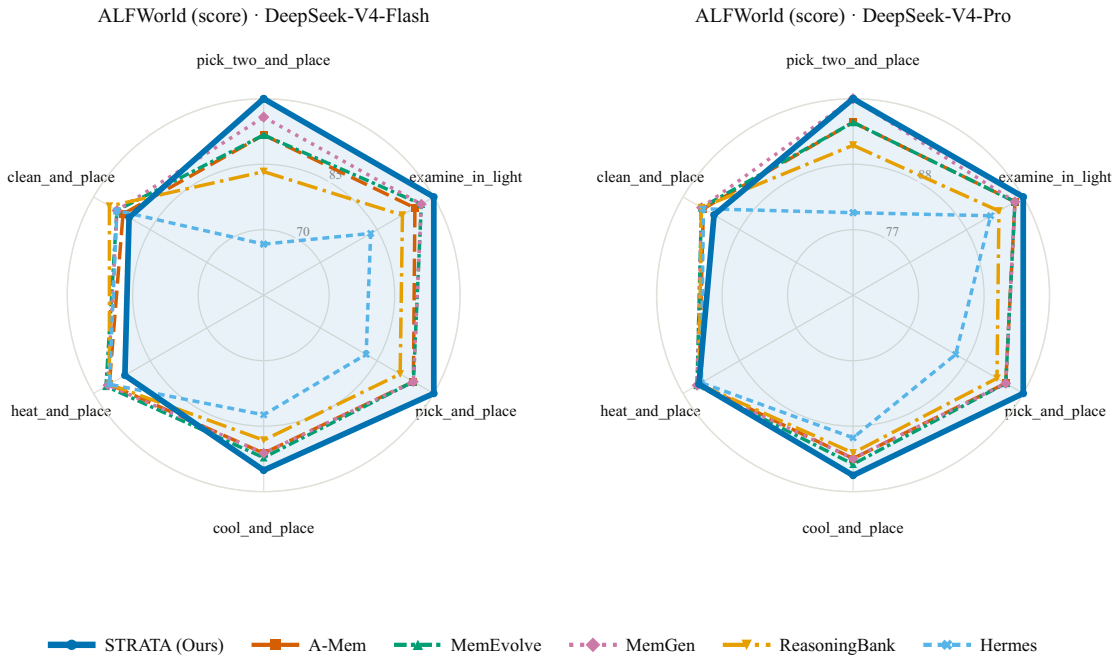

Figure 5: ALFWorld per-task-type score, both backbones.

Figure 6: ScienceWorld per-task-type score, both backbones. Hermes' `find-living` value is off-scale (-100/-80) and is clipped to the axis floor; this is the negative-transfer case analysed in the main text.

I. Example Memory Nodes

Figures 8–10 reproduce two real nodes from the AppWorld Test-Challenge graph, verbatim in STRATA's typed-markdown serialization. Together they illustrate one workflow $\rightarrow$ meta relation: the meta node in Figure 10 is one of the 5 metas counted

for AppWorld Test-Challenge in Table 4, and the workflow node in Figure 8 is one of the 12 workflows it generalizes.

The workflow node shows what *outcome-grounded* accumulation produces in practice. Figure 8 is the node as first crystallized: the structured `app:`/`obj:`/`op:` tags that drive the tag term, an `applicability` string that forms part of the index view, and one generalized procedure. Figure 9 is everything the reflector added afterwards, namely 23 bullets, of which 7 are PITFALL bullets distilled from *failed* episodes and rendered to the agent as things to avoid, and 16 are variant bullets from *solved* near-neighbour episodes whose similarity fell inside the $[\tau_{\text{low}}, \tau_{\text{high}}]$ band. Crucially the procedure itself was never rewritten, since no failed episode ever crystallized a "do this" procedure, and no near-duplicate ever created a competing node.

The meta node shows what community detection abstracts. Each member contributes one variant bullet keyed by its applicability condition, so retrieving the meta gives the agent a routing table over the `app:gmail` family rather than a single procedure. It reached support 12 through repeated UPDATE rather than repeated CREATE, which is the novelty gate of Table 8 operating on a real graph.

One implementation detail is visible here and worth stating plainly. The tag list is *not* the union over all 12 children: an update recomputes it from the members of the community detected in that round, so tags contributed by children attached in earlier rounds can drop out. That is why `op:delete`, carried by the workflow of Figure 8, does not appear on the meta. Bullets accumulate monotonically while tags do not. Since meta tags feed only the naming and the tag term of retrieval, this weakens meta recall rather than corrupting any workflow; making the tag set accumulate like the bullets is a one-line fix we did not apply because it would invalidate the reported runs.

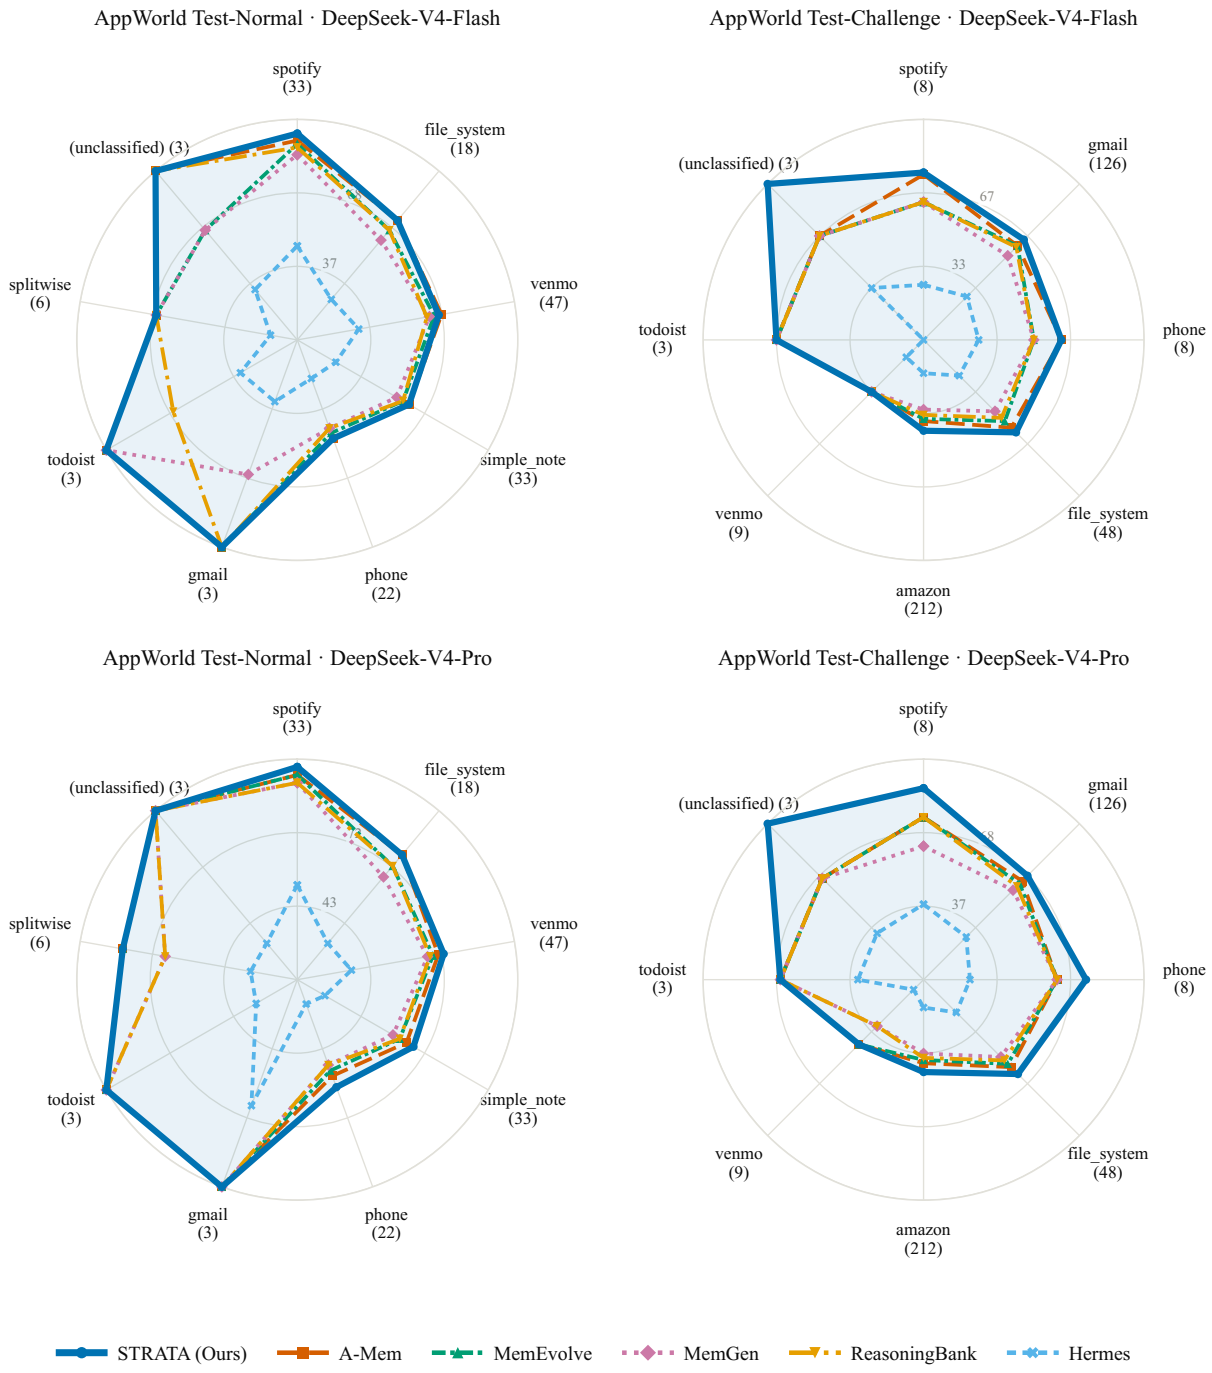

Figure 7: AppWorld per-application scores for both splits and both backbones.

```
---
id: wf_clean_gmail_read_threads_except_priority_star
type: workflow
semantic_role: task_experience
tags: [app:gmail, obj:email_threads, op:delete, solved]
applicability: "When deleting read Gmail threads, keeping only priority starred ones"
---
1. Log in to Gmail using account credentials. 2. Fetch all read threads from inbox and outbox
(paginated). 3. Deduplicate by thread ID. 4. For each thread, retrieve its labels and starred
status. 5. Delete threads that are NOT both priority-labeled and starred. 6. Repeat for each
unique thread.

**Why:** Automates cleaning read threads while preserving important priority starred ones.

**How to apply:** Iterate inbox/outbox; deduplicate; filter by priority label and star; delete
others.
```

Figure 8: A **workflow** node from the AppWorld Test-Challenge graph, verbatim in STRATA's typed-markdown format: YAML frontmatter carrying the structured tags that drive the tag term, followed by the generalized procedure crystallized from the first solved episode.

```
- (pitfall) PITFALL (avoid): Before completing the task, ensure you have checked all relevant
    mailboxes (inbox, outbox, archived, spam, etc.) for the labels. Calling 'complete_task'
    prematurely finalizes the episode, preventing corrections--always verify all sources first.
- (when user wants to rename priority labels (P1/P2) and remove P3 labels from all email
    threads) Instead of filtering read threads and skipping priority starred, target threads
    with labels P1, P2, P3 across all mailboxes, then relabel P1->priority-1, P2->priority-2,
    and remove P3 labels.
- (When you need to rename or remove labels on multiple threads across all mailboxes.) Batch
    relabel threads by searching all mailboxes for specific labels (pr-1, pr-2, pr-3) and
    applying rename/remove operations, instead of cleaning read threads.
- (when user wants to delete drafts that have both empty subject and body) Operate on Gmail
    drafts with empty subject and body instead of cleaning read threads with priority/starred
    conditions.
- (For deleting Gmail drafts with empty subject or body.) Operates on drafts instead of
    threads; deletes drafts with empty subject or body rather than cleaning read threads except
    priority starred.
- (when the task involves deleting Gmail drafts with specific content conditions) Use
    `show_drafts` and iterate over drafts instead of threads; delete drafts where both subject
    and body are empty.
- (pitfall) PITFALL (avoid): Do not hardcode the current month; obtain the current date
    dynamically (e.g., from a system call or API) to compute the cutoff for 'before this
    calendar month'.
- (pitfall) PITFALL (avoid): Always verify the current month using the system's actual date
    (e.g., import datetime; datetime.datetime.now().month) rather than assuming from token
    expiration or hardcoded values, to correctly filter threads for 'this calendar month'.
- (pitfall) PITFALL (avoid): Do not rely on Python's datetime.date.today() for task dates; the
    environment may have a different 'today'. Use the task context or system-provided date if
    available.
- (When needing count of unread threads by label) Count unread threads by label 'priority-1'
    instead of cleaning read threads excluding priority-starred.
- (When counting threads filtered by a specific label (e.g., priority-2) and read status.) Used
    the `label` filter parameter on `show_inbox_threads` to directly fetch only priority-2 read
    threads, instead of fetching all read threads and then filtering by labels.
- (When counting unread threads with a specific priority label) Count unread threads filtered
    by a specific label (e.g., priority-3) instead of cleaning read threads.
- (pitfall) PITFALL (avoid): Fetched only inbox and outbox threads, missing sent and other
    non-archived threads. Must fetch all threads excluding archived (e.g., use labels INBOX,
    SENT, DRAFT etc. or entire 'all' minus archived).
- (pitfall) PITFALL (avoid): Do not rely solely on thread-level participants; use show_thread
    to inspect each email within a thread to accurately determine if any coworker is a sender
    or recipient.
- (when user wants to manage spam folder by rescuing coworker threads and purging all other
    spam) Focus on spam folder: identify coworker threads in spam, unspam and mark unread, then
    delete all remaining spam (vs. cleaning read threads except priority starred).
- (when the user wants to manage spam folder by recovering friends' threads and deleting the
    rest) Instead of cleaning read threads, this variant targets spam threads: identifies
    friends' threads in spam, unspams and marks them unread, then deletes all remaining spam.
- (when managing spam and preserving specific threads) Focus on spam folder; conditionally
    unspam/restore roommate threads, then delete all other spam.
- (pitfall) PITFALL (avoid): Check all mailboxes (inbox, outbox, archived) for TODO threads
    before replying and completing; the manager may have tasks in multiple folders.
- (When user has weekly 'To Do' emails from manager to reply based on star status) Process
    weekly 'To Do' emails from manager: reply based on starred status and unstar completed,
    preserving non-todo starred.
- (when closing weekly manager assignments (subject 'Assignment') by replying 'Finished.' or
    'Not Finished.' based on starred flag) Reply to weekly Assignment threads based on starred
    status, unstar finished, preserve non-todo starred.
- (When the task is to label Gmail threads by sender relationship (work vs personal) using
    contacts from the Phone app.) Uses Phone app contacts to categorize senders as coworkers or
    personal, then labels threads accordingly as 'work' or 'personal', instead of cleaning read
    threads. Also processes both inbox and outbox, not just read threads.
- (When labeling emails based on contact categories from phone app) Uses phone contacts to
    label threads by participant relationship (coworker vs personal).
- (when Gmail threads need to be labeled based on contact relationships from Phone app) Added
    step to retrieve contacts from Phone app, categorize by relationship
    (coworkers/friends/family), then apply 'work' or 'personal' labels to email threads based
    on sender/recipient matching those contacts.
```

Figure 9: The 23 bullets the workflow reflector accumulated onto the node of Figure 8 over 24 later episodes. 7 are PITFALL bullets distilled from *failed* episodes; the rest are variant bullets from *solved* near-neighbour episodes. The procedure itself was never rewritten.

```
---
id: meta_app_gmail_2
type: meta_framework
semantic_role: task_pattern
tags: [app:gmail, obj:attachment, obj:email, obj:email_invitations, obj:email_thread,
obj:files, op:create, op:transfer, solved]
applicability: "'app:gmail' family"
---
Common pattern generalized from: wf_send_hiring_email_replies,
wf_download_email_attachment_to_filesystem, wf_email_files_per_text_request

**Why:** Recurring community of similar workflows.

**How to apply:** Route any 'app:gmail'-like task through the matching workflow.
- (when hiring and need to email candidates) variant wf_send_hiring_email_replies: Send hiring
    email replies
- (When needing to transfer email attachment to local folder) variant
    wf_download_email_attachment_to_filesystem: Download Email Attachment to Filesystem
- (When a contact requests files via text) variant wf_email_files_per_text_request: Email files
    per text request
- (when phone message requests emailing a file) variant wf_email_file_from_phone_request: Email
    file from phone request
- (Forward email thread with extra file and prefixed note) variant
    wf_forward_thread_with_note_and_attachment: Forward thread with note and attachment
- (When you need to send personalized bulk emails using a draft template and contact list from
    file.) variant wf_send_personalized_bulk_emails_from_draft: Send personalized bulk emails
    from draft
- (when need to forward one email from a thread) variant wf_forward_single_email_from_thread:
    Forward single email from thread
- (When deleting read Gmail threads, keeping only priority starred ones) variant
    wf_clean_gmail_read_threads_except_priority_star: Clean Gmail read threads except priority
    starred
- (When away and need to reply to recent unread emails) variant
    wf_vacation_auto_reply_to_unresponded_threads: Vacation auto-reply to unresponded threads
- (when draft email contains $TODO prices for Amazon items) variant
    wf_fill_draft_prices_from_amazon: Fill draft prices from Amazon
- (When need to forward email with prefixed comment) variant wf_forward_email_with_prefix_note:
    Forward email with prefix note
- (when needing to forward email with custom prefix) variant
    wf_forward_email_with_prefixed_note: Forward email with prefixed note
```

Figure 10: The **meta-framework** node that generalizes the workflow of Figure 8, induced automatically by community detection over the workflow similarity graph.