Multi-SFT: Permissively-Licensed Multilingual Instruction-Following

Daniela Baiamonte[1], Kamyar Zeinalipour[2,3], Elena Fano[1], Stefano Simonazzi[1], Andrea Zugarini[1]
[1]villanova.ai, [2]Yukai!, [3]University of Siena
Correspondence: kamyar.zeinalipour2@unisi.it

arXiv:2610.22086v1 [cs.CL] 15 Oct 2026

Abstract

Supervised fine-tuning datasets remain overwhelmingly English-centric and commercially restricted, as their content is frequently generated by proprietary models whose terms of service forbid training competing systems. We introduce *Multi-SFT*, a commercially permissive collection of 502,818 instruction-response pairs across five European languages, built via *translation* and *regeneration* pipelines for tasks where translation is insufficient. Human evaluation yields inter-annotator agreement of at least 0.87 across all datasets and languages. Fine-tuning nine models from three different families (EuroLLM, OLMo 2, Qwen3; 1B–32B) and evaluating them on 13 benchmarks shows that *Multi-SFT* alone is competitive with larger augmented mixes and official instruction-tuned releases, with decisive advantages on safety tasks.

1 Introduction

Supervised fine-tuning (SFT) is central to aligning pretrained LLMs with human instructions (Wei et al., 2021; Ouyang et al., 2022; Taori et al., 2023), with data quality, diversity, and scale being critical for downstream behavior (Wang et al., 2023; Zhou et al., 2023a). Yet, most existing SFT datasets remain English-centric and are affected by a less visible but equally serious problem: licensing. Many corpora carry permissive labels (MIT, Apache 2.0, Creative Commons), while their content was generated by proprietary models such as GPT-4 or Gemini, whose terms of service explicitly prohibit use for training competing models.

We introduce *Multi-SFT*, a collection of 502,818 instruction-response pairs spanning five European languages (English, Italian, French, Spanish, and German), aggregating 13 heterogeneous datasets built exclusively from permissively licensed data and models (Table 1). The collection is constructed via two pipelines: a *translation* pipeline that adapts existing English-centric datasets into the four target languages through an expert-driven, quality-assured process, and a *regeneration* pipeline that synthesizes multilingual content from scratch for tasks where translation is insufficient, such as conversational interaction, constrained instruction following, and safety. Both pipelines undergo systematic human evaluation by native or near-native annotators, yielding inter-annotator agreement of at least 0.87 across all datasets and languages.

Project	Multilingual	Dataset License	Data Provenance
Apertus	✓	✓	✗
EuroLLM	✓	~	✗
Nemotron	✗	✓	✓
OLMo	✗	✓	✗
Salamandra	✓	✗	✗
Multi-SFT	✓	✓	✓

Table 1: **Comparison of SFT mixtures used in recent open LLM projects.** ✓ denotes full compliance, ✗ denotes non-compliance, and ~ denotes partial compliance. *Dataset License* indicates whether all training data is released under permissive licenses. *Data Provenance* indicates whether synthetic data is generated exclusively with permissive or open-weight models, excluding proprietary sources such as GPT-4 or Gemini.

We validate *Multi-SFT* by fine-tuning nine models from three families (EuroLLM, OLMo 2, Qwen3; 1B–32B parameters) on 13 benchmarks. Our experiments show that the multilingual, permissively-licensed collection achieves stronger performance on non-English languages than the original English sources, which contain non-permissive data from closed LLMs, while surprisingly maintaining performance on English benchmarks as well. Furthermore, in multilingual settings, Multi-SFT outperforms official instruct models in most cases and stays competitive on English. Multi-SFT is publicly available.[1]

[1] https://huggingface.co/collections/VillanovaAI/multi-sft-6ad0e9ea53f7d008611b8204

Dataset	Domain	Strategy	License	EN	DE	ES	FR	IT	Total
Multi-AdvBench	Safety	Regeneration	MIT	104	104	104	104	104	520
Multi-AyaRedteaming	Safety	Regeneration	Apache 2.0	812	835	330	754	913	3,644
multi-dialogues	Conversation	Regeneration	ODC-BY-1.0	-	12,001	11,999	12,000	12,001	48,001
multi-dolly-15k	General QA	Translation	CC BY-SA 3.0	6,004	2,250	2,251	2,252	2,222	14,979
Multi-FLAN	Reasoning	Translation	Apache 2.0	36,069	13,495	13,441	13,492	13,485	89,982
multi-mentorES	General QA	Translation	CC BY 4.0	-	2,541	2,552	2,534	2,548	10,175
Multi-PersonaIF	Instruction Following	Regeneration	ODC-BY-1.0	15,000	15,000	15,000	15,000	15,000	75,000
Multi-RAG-Multilingual	General QA	Translation	CC-BY-SA-4.0	4,999	4,999	5,631	4,998	5,632	26,259
multi-safety	Safety	Regeneration	ODC-BY-1.0	250	250	250	250	250	1,250
Multi-SciRIFF	Scientific Literature	Translation	ODC-BY	28,560	10,521	10,508	10,540	10,392	70,521
multi-smol_rewrite	Rephrasing	Translation	Apache 2.0	21,305	7,990	7,989	7,988	7,990	53,262
multi-smol_summarize	Summarization	Translation + Regeneration	Apache 2.0	38,424	14,409	14,402	14,410	14,409	96,054
Multi-TableGPT	Table Understanding	Translation	MIT	5,279	1,964	1,969	1,980	1,979	13,171
Total									**502,818**

Table 2: Overview of the *multi-* dataset collection. Source datasets and pipelines: AdvBench (Zou et al., 2023), Aya Red-teaming (Ahmadian et al., 2024), SODA (Kim et al., 2023), Dolly (Conover et al., 2023), FLAN (Longpre et al., 2023), Persona IF (Lambert et al., 2024), RAG Multilingual (Projecte AINA, 2026), SciRIFF (Wadden et al., 2025), SmolTalk (Allal et al., 2025), and Table-GPT (Li et al., 2024).

2 Related Work

Supervised fine-tuning (SFT) adapts pre-trained language models to follow instructions via curated prompt-response pairs (Ouyang et al., 2022). English SFT resources span a continuum from fully human-written data—such as FLAN (Longpre et al., 2023), Dolly (Conover et al., 2023), and Open Assistant (Köpf et al., 2023)—to hybrid pipelines combining human annotation with synthetic generation, as in Tülu 3 (Lambert et al., 2024). Recent large-scale synthetic efforts such as Nemotron (Blakeman et al., 2025) and Smoltalk (Allal et al., 2025) demonstrate the scalability of automatically constructed instruction datasets while explicitly addressing licensing and provenance through permissively licensed sources and models. High-quality multilingual SFT resources for European languages remain comparatively limited, and existing datasets often face restricted licensing or provenance concerns due to reliance on proprietary generative models. The Aya initiative (Singh et al., 2024) provides the broadest multilingual effort to date, though its wide coverage results in limited per-language depth and variable quality in automatically generated portions. EuroLLM (Martins et al., 2024), Tülu 3 (Lambert et al., 2024), Salamandra (Gonzalez-Agirre et al., 2025), and Apertus (Apertus et al., 2025) each advance multilingual or open post-training in different respects, yet all include components generated or processed with proprietary models, limiting their suitability for provenance-constrained training. The Smoltalk multilingual extension represents one of the few systematic attempts to construct permissively licensed multilingual SFT data through translation and open-model generation. Our work follows a similar philosophy while expanding task coverage, introducing expert-driven quality validation, and explicitly addressing provenance constraints through selective regeneration of non-compliant examples.

3 Multi-SFT

Multi-SFT is a supervised fine-tuning collection of 13 datasets covering five European languages: English, Italian, French, Spanish, and German. Overall, it contains 502,818 diverse examples (e.g., QA, instruction following). Per-dataset details are outlined in Table 2.

Many existing SFT datasets, despite declaring permissive licenses, are built on human-authored prompts paired with responses generated by proprietary models, which restricts their use in commercial settings. Therefore, we erased such content and regenerated it using open-weight, permissively licensed LLMs. Depending on the source dataset, we employ a *translation* pipeline, where existing (English-centric) datasets are converted into multilingual ones, or a *regeneration* process, where synthetic content is re-generated from scratch with license-compliant language models.

3.1 Dataset selection criteria

Rather than aggregating large volumes of generic instruction-response pairs, we focus on diversity and high quality. Thus, source datasets were selected to cover a vast and heterogeneous mixture of competences and tasks (see Table 2).

Licensing and Compliance. A prerequisite for inclusion in the collection was the availability of a

permissive license compatible with commercial use. Moreover, we also take into account the provenance of synthetic content present in the datasets. This aspect is often overlooked, but proprietary language models such as GPT, Claude or Gemini restrict or prohibit their use for the development of commercial models. Hence, we excluded all datasets where there is evidence that entries were generated using models with such restrictive licensing, regardless of the license applied to the dataset itself. We consider as "evidence" of non-permissive model use any dataset mentioning in its card, git repository or paper the use of commercial models. Whenever a dataset is released under a permissive license but a portion of its entries were generated using a restricted model, we opted for selective regeneration: the affected entries were regenerated using permissively licensed open-weight models, specifically: Mistral-small[2], Qwen3-VL-235B-A22B[3], and Qwen3.5-122B-A10B[4]. These models were selected based on their empirical performance across translation and generation tasks, and their licensing compatibility with downstream training (see Appendix A.2 for details). Out of 59 candidate datasets initially considered (excluding math and coding topics), 45 were excluded due to licensing or provenance violations (see Appendix A.1). A sample of regenerated examples is provided in Appendix A.3.

3.2 Translation Pipeline

Translated datasets were produced through an iterative, expert-driven pipeline. For each target language, a native or near-native linguist independently developed a translation prompt, selecting the generation model and temperature in collaboration with the team. An initial batch of samples was generated and evaluated blind by a separate QA reviewer with equivalent language expertise. Based on the reviewer's feedback, the prompt, model, or parameters were revised and the process repeated until the output quality was deemed satisfactory.

A special case within the translation pipeline is *multi-smol_summarize*, which combines translation and generation. In the original Smoltalk summarization dataset, source documents are paired with summaries generated from a very small set of instruction variants. In our version, the source documents are first translated into the target languages; the summaries are then generated from scratch using the models listed in Section 3.1. Crucially, we expanded the original instruction set by introducing additional variants expressing formal output requirements, such as producing the summary as a bulleted list or as a Markdown table, thereby increasing both the linguistic and structural diversity of the dataset.

3.3 Synthetic Data Generation

Five datasets in the collection were constructed through synthetic generation. For each, the same expert validation process described in Section 3.2 was applied: prompt design, blind evaluation, and iterative refinement by language-expert reviewers.

Multi-AdvBench. The original multilingual AdvBench dataset consists of harmful prompt–unsafe response pairs. We retained the prompts and regenerated all responses using `Qwen3-VL-235B-A22B-Instruct`, with the explicit objective of producing safe, refusal-oriented completions. To avoid any overlap across languages, the original 520 prompts were randomly partitioned into five disjoint blocks of 104, one per language, so that each prompt appears in exactly one language. All generated pairs were manually reviewed.

Multi-AyaRedTeaming. Starting from the English, French, and Spanish subsets of the Aya Redteaming dataset, we repurposed this adversarial resource for SFT by transforming harmful prompts into instruction-following examples with safe target responses. We generated the expected safe responses using `Qwen3-VL-235B-A22B-Instruct`. For Italian and German, where no source data was available, we used the English prompts as seeds: the same model was prompted to produce culturally appropriate questions in the target language and then to generate the corresponding safe responses. All examples were manually reviewed and corrected.

Multi-Dialogues. It consists of synthetically generated multi-turn everyday conversations, produced by adapting the SODAverse pipeline from Kim et al. (2023). The original SODA dataset was built from social common-sense triples extracted from the ATOMIC^{10x} knowledge graph (West et al., 2022): each triple is first expanded into a short narrative, which is then used to prompt a model to generate a dialogue. We replicated the pipeline

[2] `mistralai/Mistral-Small-3.1-24B-Instruct-2503`
[3] `Qwen/Qwen3-VL-235B-A22B-Instruct`
[4] `Qwen/Qwen3.5-122B-A10B`

for French, German, Italian, and Spanish using `Mistral-Small-3.1-24B-Instruct-2503`, generating 12,000 dialogues per language. Then, `Qwen3-VL-235B-A22B-Instruct` was used for cleaning, formatting, and profanity filtering; American characters were replaced with culturally-appropriate names for each target language.

Multi-PersonaIF. This dataset targets instruction following under explicit constraints, following the methodology of AllenAI's Persona IF (Lambert et al., 2024). Instructions are generated by randomly sampling 1–3 constraints from a predefined list derived from IFEval (Zhou et al., 2023b), using persona descriptions from a multilingual extension of PersonaHub (Ge et al., 2024) to diversify instruction styles. Responses are generated by the same model immediately after instruction generation. We replicated the full pipeline using `Qwen3.5-122B-A10B` and translated and adapted all prompts, seed instructions, and constraint definitions for the five target languages. One methodological modification was introduced: the original constraint sampling procedure can yield logically incompatible combinations (e.g., write in lowercase and include a minimum number of capitalized words). We revised the sampling logic to prevent such conflicts.

Multi-Safety. A small dataset of harmful request–safe response pairs was constructed internally: harmful prompts were authored by team members, and safe responses were generated using `Qwen3-VL-235B-A22B-Instruct`. All pairs were manually reviewed before inclusion.

3.4 Annotation Process

To assess the quality of translated and generated datasets, we conducted a systematic human evaluation covering all datasets produced through the translation and generation pipelines. Inspired by Wang et al. (2023), instruction-response pairs were annotated on an A-to-D grading scale (Table 7). Full details are reported in Appendix B.

For each dataset and each language, a sample of 50 entries was selected for annotation, with one exception: for Multi-FLAN, which aggregates four sub-datasets (CoT, P3, FLAN2021, NIv2), 60 entries were annotated per language, 15 per sub-dataset, to ensure proportional coverage. Each sample was annotated independently by two annotators, both with university-level expertise in the target language (advanced students or graduates in linguistics or related fields). English annotation was performed only for datasets where English content was generated from scratch, i.e., Multi-PersonaIF and multi-smol_summarize.

Dataset	DE	EN	ES	FR	IT
multi-dialogues	98	–	99	100	95
multi-dolly-15k	100	–	93	100	94
Multi-FLAN	90	–	87	84	88
multi-mentorES	95	–	–	96	93
Multi-PersonaIF	87	89	83	80	72
Multi-RAG-Multilingual	100	–	–	96	90
Multi-SciRIFF	97	–	97	92	87
multi-smol_rewrite	100	–	97	99	98
multi-smol_summarize	99	94	94	96	95
Multi-TableGPT	98	–	76	97	98

Table 3: Percentage of positively-rated entries by dataset and language. Values represent the average of the two independent annotators.

3.5 Quality Assessment

Metrics. Inter-annotator agreement was computed using Gwet's AC2 (Gwet, 2008) with quadratic weights on the collapsed three-way label set (A/B vs. C vs. D). AC2 was preferred over Cohen's κ (Cohen, 1960) because the annotation distributions are strongly skewed toward the positive category (A/B), a condition under which κ is known to produce paradoxically low values even when observed agreement is high. The quadratic weighting scheme is appropriate for ordinal scales, where the distance between non-adjacent categories (e.g., A/B and D) should be penalised more heavily than disagreements between adjacent ones.

Results. Agreement results are reported in Table 4. The distribution of raw annotation labels across datasets and languages is provided in Table 3. Agreement is consistently high across datasets and languages, with values exceeding 0.90 in the majority of cases. This indicates reliable interpretation of the annotation schema and stable quality signals. Three datasets show lower but still substantial agreement: Multi-FLAN (FR: 0.93, DE: 0.90), Multi-PersonaIF (FR: 0.88), and Multi-TableGPT (ES: 0.87), reflecting known task-inherent challenges discussed in Appendix C. The label distribution (Table 3) reveals that the vast majority of entries fall in the positive category, confirming overall pipeline quality. Multi-PersonaIF exhibits the highest proportion of badly-rated entries (up to 28% in Italian), stemming from two sources:

Dataset	DE	EN	ES	FR	IT
multi-dialogues	0.99	–	0.99	1	0.97
multi-dolly-15k	1	–	0.95	1	0.98
Multi-FLAN	0.90	–	0.97	0.93	0.93
multi-mentorES	0.99	–	–	0.96	0.98
Multi-PersonaIF	0.98	0.94	0.98	0.88	0.94
Multi-RAG-Multilingual	1	–	–	0.98	0.97
Multi-SciRIFF	0.98	–	0.97	0.96	0.94
multi-smol_rewrite	1	–	0.97	0.99	0.99
multi-smol_summarize	0.99	0.97	0.97	0.98	0.97
Multi-TableGPT	0.99	–	0.87	0.98	0.99

Table 4: Inter-annotator agreement (Gwet's AC2) by dataset and language.

constraint violations and incomplete instructions lacking necessary context. While constraint violations reflect inherent task difficulty, incomplete instructions represent actionable quality issues addressable through improved prompt engineering and post-hoc filtering.

4 Experiments

4.1 Experimental Setup

We pick nine **base** models from three families: EuroLLM (Martins et al., 2024) (1.7B, 9B, 22B), OLMo 2 (OLMo et al., 2024) (1B, 7B, 32B), and Qwen3 (Yang et al., 2025) (1.7B, 8B, 14B), overall spanning from 1B to 32B parameters. These three families differ significantly in their nature. Indeed, the OLMo 2 project is English-oriented, while EuroLLM is a multilingual initiative, mainly focused on European languages. Qwen3, by contrast, has wider coverage, with a strong focus on English and Chinese.

To prove the effectiveness of our multilingual permissive collection, we train each base model on three different corpora. First, we consider ours, Multi-SFT (502,818 examples), denoted as ML. In FULL, we expand the training set of ML to 1.6M examples, by adding existing open-source datasets: the Aya Collection (Singh et al., 2024) (928K), multilingual SmolTalk (Allal et al., 2025) (127K), GSM8K (Cobbe et al., 2021) (37K), and RAG Multilingual (Projecte AINA, 2026) (10K). Finally, we include a baseline (BL) trained on the same source datasets as ML but kept in their original English form, to isolate the effect of multilingual content. As a reference of comparison we also evaluated the officially released instruct models, here denoted as SFT†. Training details are reported in Appendix D.

Evaluation benchmarks. We test all models on 10 distinct benchmarks, yielding 13 benchmarks in total since three benchmarks (XQuAD, M-MMLU, ARC) are evaluated in both multilingual and English settings. The multilingual ones (Table 5) are: MARCO (Zeng et al., 2025) for constrained generation in 5 languages, FLORES-200 (Costa-Jussà et al., 2022) for translation (chrF++, 8 directions), Belebele (Bandarkar et al., 2024) for reading comprehension in 5 languages, XQuAD (Artetxe et al., 2020) for question answering (F1, 3 languages), M-ALERT (Friedrich et al., 2024) for safety (5 languages), M-MMLU (Hendrycks et al., 2020) for knowledge in DE/ES/FR/IT, ARC (Clark et al., 2018) for commonsense in 4 languages, and XStoryCloze (Lin et al., 2021) for story completion in Spanish. The English ones (Table 5) are: XQuAD, DROP (Dua et al., 2019), IFEval (Zhou et al., 2023b), M-MMLU$_{\text{EN}}$ and ARC. Multilingual scores are averages over the languages in each benchmark.

4.2 Results

All results are presented in Table 5.

Multi-SFT vs Official SFT models. ML scores higher than SFT† on the multilingual average for 6 out of 9 models. Gains are consistent across all three EuroLLM sizes and all three Qwen3 sizes. The largest gaps appear on Qwen3 (up to +15.9 on 14B). All four variants (ML, FULL, BL, and SFT†) of Qwen3 were evaluated under identical non-thinking settings, ensuring a fair comparison within the family. The trend holds without caveats on EuroLLM, where ML outperforms SFT† on the multilingual average and matches or outperforms it on the English average at every scale.

Impact of adding more data. Table 5 shows that ML, trained only on our Multi-SFT, keeps up with the larger FULL mixture despite using roughly one third of the data. ML beats FULL on 4 out of 9 models (all three EuroLLM sizes and OLMo 2 1B) and remains within 1 point on Qwen3 8B and 14B, confirming that data quality matters more than volume (Zhou et al., 2023a).

English performance does not drop. Despite the significant reduction of English data, Table 5 shows that ML retains most of the English performance. Indeed, it matches or beats SFT† on 4 out of 9 models. Qwen3 shows the biggest ML advantage (+10.7 on 8B, +13.6 on 14B), while OLMo 2 SFT† recipe is better than ML, likely due to English-heavy pretraining.

		Multilingual									English					
Family	**Model**	**MARCO**	**FLORES**	**Belebele**	**XQuAD**	**M-ALERT**	**MMLU**	**ARC**	**XStory**	**Avg**	**XQuAD**	**DROP**	**IFE**	**MMLU**	**ARC**	**Avg**
EuroLLM	1.7B ML	**31.7**	49.0	23.4	33.4	43.6	24.5	32.0	58.4	**37.0**	32.8	23.0	32.3	23.5	34.0	29.1
	1.7B FULL	21.1	45.3	23.0	25.7	21.6	24.4	32.2	**61.2**	31.8	24.6	21.1	27.9	23.4	34.1	26.2
	1.7B BL	27.1	52.2	**25.9**	**42.6**	29.0	24.5	32.1	58.6	36.5	**42.5**	**25.2**	**39.9**	24.3	**37.5**	**33.9**
	1.7B SFT†	16.5	**59.1**	25.4	25.0	**44.4**	**26.3**	**33.3**	60.8	36.4	37.1	25.1	20.3	**26.7**	37.4	29.3
	9B ML	**50.2**	62.0	64.6	75.8	**56.3**	**50.7**	**48.9**	64.5	**59.1**	73.9	46.6	**59.5**	53.0	**52.9**	**57.2**
	9B FULL	25.1	57.8	62.3	76.2	33.9	44.1	43.3	57.4	50.0	**75.6**	**48.1**	26.4	46.8	44.5	48.3
	9B BL	41.7	60.5	64.8	**76.6**	36.3	46.5	43.0	64.5	54.2	75.4	47.4	52.9	50.2	46.7	54.5
	9B SFT†	42.3	**62.8**	**70.6**	68.5	55.7	45.4	44.9	**67.2**	57.2	64.5	40.5	57.5	**54.5**	50.3	53.5
	22B ML	49.4	62.5	77.9	**79.4**	53.7	57.5	**51.7**	66.7	**62.4**	**77.3**	52.0	59.3	60.4	55.3	60.9
	22B FULL	47.3	62.2	75.0	76.6	**56.0**	56.0	48.9	68.2	61.3	76.8	47.0	58.0	59.2	53.6	58.9
	22B BL	46.4	62.3	77.7	76.3	37.9	56.3	50.8	67.2	59.4	76.2	**55.2**	63.0	60.7	**57.0**	**62.4**
	22B SFT†	**50.2**	**63.5**	**79.2**	70.5	40.8	**60.9**	47.3	**69.1**	60.2	67.6	49.6	**70.1**	**63.7**	52.9	60.8
OLMo 2	1B ML	**32.9**	44.6	**33.3**	26.4	**49.5**	**31.0**	26.9	55.3	37.5	24.8	33.6	41.2	39.2	38.1	35.4
	1B FULL	18.5	40.8	28.7	47.0	30.8	28.1	26.6	54.9	34.4	41.3	**34.3**	22.6	34.3	38.5	34.2
	1B BL	21.1	42.6	32.0	47.0	21.7	30.1	26.2	53.6	34.3	**51.2**	33.4	46.2	**40.4**	**38.7**	42.0
	1B SFT†	30.6	**51.2**	31.1	**48.6**	19.1	30.3	**30.6**	**59.7**	**37.6**	50.9	31.7	**66.0**	34.0	38.2	**44.2**
	7B ML	**47.8**	46.7	**69.8**	50.8	51.8	44.6	36.3	60.4	51.0	56.6	58.4	56.0	57.5	50.8	55.9
	7B FULL	45.0	51.9	67.8	**65.2**	**56.9**	**46.2**	**38.2**	63.3	**54.3**	**67.1**	55.0	55.3	**58.3**	**51.8**	57.5
	7B BL	34.0	44.6	64.4	61.2	32.5	42.2	34.3	59.4	46.6	62.3	56.6	61.6	58.0	50.9	57.9
	7B SFT†	43.4	**57.5**	66.6	63.7	41.0	44.7	37.9	**64.6**	52.4	67.0	**58.9**	**71.9**	57.3	47.4	**60.5**
	32B ML	52.1	29.8	82.2	45.3	**55.4**	61.2	50.6	66.6	55.4	51.6	54.3	67.7	69.2	59.2	60.4
	32B FULL	49.2	46.9	85.8	71.5	53.5	62.7	50.2	64.5	60.5	**72.5**	**72.7**	59.7	69.3	56.4	66.1
	32B BL	42.6	32.7	77.7	**73.1**	41.1	58.2	44.5	66.4	54.5	72.4	62.4	64.7	67.4	59.3	65.2
	32B SFT†	**58.6**	**61.1**	**87.2**	69.9	49.7	**66.1**	**55.8**	**76.4**	**65.6**	67.0	71.6	**82.3**	**72.9**	**66.0**	**72.0**
Qwen3	1.7B ML	43.4	51.6	74.1	45.4	55.4	51.0	35.7	59.9	52.1	50.0	46.6	53.8	59.0	41.8	50.2
	1.7B FULL	44.7	52.5	72.2	53.6	**57.6**	**52.2**	**41.6**	**60.0**	**54.3**	60.0	**47.2**	52.1	58.3	**52.8**	54.1
	1.7B BL	40.3	50.0	**74.5**	**65.9**	42.4	51.6	37.1	59.9	52.7	**65.8**	46.5	58.2	**59.2**	44.3	**54.8**
	1.7B SFT†	**46.0**	**54.7**	36.2	61.7	30.7	31.8	32.1	55.1	43.5	61.4	37.7	**67.8**	50.9	33.4	50.2
	8B ML	50.9	50.8	**89.2**	62.1	53.8	**67.9**	51.0	64.0	61.2	66.2	**64.0**	69.5	72.8	57.8	**66.1**
	8B FULL	50.6	55.3	88.5	65.1	**55.6**	67.2	**51.3**	64.1	**62.2**	**70.9**	57.4	57.7	71.6	**59.6**	63.4
	8B BL	49.5	50.4	86.4	**72.0**	42.7	67.4	50.8	**65.1**	60.5	**70.9**	59.1	67.5	**73.2**	57.0	65.5
	8B SFT†	**57.7**	**59.7**	28.4	71.3	41.3	33.0	37.8	61.7	48.9	70.3	51.4	**81.3**	32.1	41.9	55.4
	14B ML	52.9	51.8	**91.7**	61.8	**55.2**	**72.6**	**54.5**	65.3	63.2	59.8	**71.7**	69.5	**77.4**	59.9	67.7
	14B FULL	52.9	51.5	91.4	63.4	54.6	71.0	54.1	**67.2**	63.3	65.4	62.5	63.2	75.2	**61.7**	65.6
	14B BL	51.1	50.1	91.3	**73.3**	51.3	71.9	52.3	67.0	**63.5**	**72.3**	66.1	70.8	77.2	60.8	**69.4**
	14B SFT†	**60.0**	**60.8**	23.7	71.1	36.1	25.8	38.7	62.3	47.3	67.0	52.5	**85.0**	23.4	42.7	54.1

Table 5: Benchmark results divided between **Multilingual** (left) and **English** (right). ML: Multi-SFT ; FULL: extended mix; BL: baseline SFT; SFT†: official provider model. Average performance is given per group (see **Avg** columns).

Impact of multilingual data. To disentangle the contribution of multilingual content from bare regeneration, we compare ML with BL. This offers a clear comparison since BL was trained on the same source datasets as ML but kept in their original English. BL also serves as a direct test against permissively and non-permissively licensed data. As expected, ML outperforms BL on 7 out of 9 models on multilingual benchmarks. On English, BL retains an advantage on 7 out of 9 models given its larger share of English examples, yet the gaps remain modest in most cases, indicating that the multilingual focus of ML does not severely harm English capability.

Differences across families and sizes. EuroLLM benefits most from ML: it outperforms SFT† on the multilingual average at all three scales (+0.6, +1.9, +2.2) and matches or exceeds it on English, as expected given EuroLLM's European-focused pre-training. OLMo 2 trails SFT† on the multilingual average, narrowly for the 1B (-0.1) and 7B (-1.4) models and markedly for the 32B (-10.2), while English consistently favors SFT†, reflecting its English-centric pre-training. Qwen3 shows the largest gains (+8.6, +12.3, +15.9 on the multilingual average), though these should be interpreted with care given that all Qwen3 variants were evaluated without thinking mode.

5 Conclusions

We introduced *Multi-SFT*, a commercially permissive supervised fine-tuning collection of 502,818 instruction-response pairs across five European languages, addressing two pervasive limitations of existing SFT resources: English-centricity and provenance-constrained licensing.

Our experiments across nine models and 13 benchmarks demonstrate that our multilingual and permissive collection yields better performance on non-English languages than the original English sources containing non-permissive data from

closed LLMs, while largely preserving performance on English benchmarks. Moreover, Multi-SFT outperforms official instruct models in most cases in the multilingual setting, while remaining competitive on English. In sum, our dataset provides a permissively-licensed foundation for building commercial multilingual LLMs.

Limitations

We acknowledge that establishing with certainty the generation provenance of third-party datasets is not always possible. Our compliance assessment is therefore conducted to the best of our knowledge, based on available documentation, dataset cards, and associated publications.

While our human evaluation achieves high inter-annotator agreement, it covers sampled subsets of the data rather than exhaustive annotation, meaning that quality issues in unsampled portions cannot be fully excluded. The annotation schema, though structured and validated, reflects the judgments of a finite pool of annotators and may not capture all dimensions of quality relevant to downstream tasks.

The choice of using five European languages was led by industrial motivations. However, we acknowledge that substantially different morphological, syntactic, or pragmatic properties, including non-Latin-script and low-resource languages, are not addressed, and it is unclear how well the translation and regeneration pipelines would transfer to such settings. Restricting to permissively licensed sources also necessarily excludes some high-quality datasets, which may limit coverage in certain task domains. Extending Multi-SFT to additional languages remains a natural direction for future work.

Acknowledgments

This work is part of the Villanova project, financed by IPCEI-CIS, Prog. n. SA.102519 – CUP B29J24000850005.

References

Arash Ahmadian, Beyza Ermis, Seraphina Goldfarb-Tarrant, Julia Kreutzer, Marzieh Fadaee, Sara Hooker, and 1 others. 2024. The multilingual alignment prism: Aligning global and local preferences to reduce harm. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 12027–12049.

Loubna Ben Allal, Anton Lozhkov, Elie Bakouch, Gabriel Martín Blázquez, Guilherme Penedo, Lewis Tunstall, Andrés Marafioti, Hynek Kydlíček, Agustín Piqueres Lajarín, Vaibhav Srivastav, Joshua Lochner, Caleb Fahlgren, Xuan Son Nguyen, Clémentine Fourrier, Ben Burtenshaw, Hugo Larcher, Haojun Zhao, Cyril Zakka, Mathieu Morlon, and 3 others. 2025. SmolLM2: When smol goes big - data-centric training of a small language model. *ArXiv*, abs/2502.02737.

Project Apertus, Alejandro Hernández-Cano, Alexander Hägele, Allen Hao Huang, Angelika Romanou, Antoni-Joan Solergibert, Barna Pasztor, Bettina Messmer, Dhia Garbaya, Eduard Frank Ďurech, and 1 others. 2025. Apertus: Democratizing open and compliant LLMs for global language environments. *arXiv preprint arXiv:2509.14233*.

Mikel Artetxe, Sebastian Ruder, and Dani Yogatama. 2020. On the cross-lingual transferability of monolingual representations. In *Proceedings of the 58th annual meeting of the association for computational linguistics*, pages 4623–4637.

Lucas Bandarkar, Davis Liang, Benjamin Muller, Mikel Artetxe, Satya Narayan Shukla, Donald Husa, Naman Goyal, Abhinandan Krishnan, Luke Zettlemoyer, and Madian Khabsa. 2024. The Belebele benchmark: a parallel reading comprehension dataset in 122 language variants. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 749–775, Bangkok, Thailand. Association for Computational Linguistics.

Aaron Blakeman, Aaron Grattafiori, Aarti Basant, Abhibha Gupta, Abhinav Khattar, Adi Renduchintala, Aditya Vavre, Akanksha Shukla, Akhiad Bercovich, Aleksander Ficek, and 1 others. 2025. NVIDIA Nemotron 3: Efficient and open intelligence. *arXiv preprint arXiv:2512.20856*.

Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. 2018. Think you have solved question answering? try ARC, the AI2 reasoning challenge. *arXiv:1803.05457v1*.

Karl Cobbe, Vineet Kosaraju, Mo Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. 2021. Training verifiers to solve math word problems. *ArXiv*, abs/2110.14168.

Jacob Cohen. 1960. A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1):37–46.

Mike Conover, Matt Hayes, Ankit Mathur, Jianwei Xie, Jun Wan, Sam Shah, Ali Ghodsi, Patrick Wendell, Matei Zaharia, and Reynold Xin. 2023. Free Dolly: Introducing the world's first truly open instruction-tuned llm.

Marta R Costa-Jussà, James Cross, Onur Çelebi, Maha Elbayad, Kenneth Heafield, Kevin Heffernan, Elahe Kalbassi, Janice Lam, Daniel Licht, Jean Maillard, and 1 others. 2022. No language left behind: Scaling human-centered machine translation. *arXiv preprint arXiv:2207.04672*.

Dheeru Dua, Yizhong Wang, Pradeep Dasigi, Gabriel Stanovsky, Sameer Singh, and Matt Gardner. 2019. DROP: A reading comprehension benchmark requiring discrete reasoning over paragraphs. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)*.

Felix Friedrich, Simone Tedeschi, Patrick Schramowski, Manuel Brack, Roberto Navigli, Huu Nguyen, Bo Li, and Kristian Kersting. 2024. LLMs lost in translation: M-ALERT uncovers cross-linguistic safety inconsistencies. *arXiv preprint arXiv:2412.15035*.

T Ge, X Chan, X Wang, D Yu, H Mi, and D Yu. 2024. Scaling synthetic data creation with 1,000,000,000 personas. *arXiv preprint arXiv:2406.20094*.

Aitor Gonzalez-Agirre, Marc Pàmies, Joan Llop, Irene Baucells, Severino Da Dalt, Daniel Tamayo, José Javier Saiz, Ferran Espuña, Jaume Prats, Javier Aula-Blasco, and 1 others. 2025. Salamandra technical report. *arXiv preprint arXiv:2502.08489*.

Kilem Li Gwet. 2008. Computing inter-rater reliability and its variance in the presence of high agreement. *British Journal of Mathematical and Statistical Psychology*, 61(1):29–48.

Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Xiaodong Song, and Jacob Steinhardt. 2020. Measuring massive multitask language understanding. *ArXiv*, abs/2009.03300.

Hyunwoo Kim, Jack Hessel, Liwei Jiang, Peter West, Ximing Lu, Youngjae Yu, Pei Zhou, Ronan Bras, Malihe Alikhani, Gunhee Kim, and 1 others. 2023. SODA: Million-scale dialogue distillation with social commonsense contextualization. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 12930–12949.

Andreas Köpf, Yannic Kilcher, Dimitri Von Rütte, Sotiris Anagnostidis, Zhi Rui Tam, Keith Stevens, Abdullah Barhoum, Duc Nguyen, Oliver Stanley, Richárd Nagyfi, and 1 others. 2023. OpenAssistant conversations – democratizing large language model alignment. *Advances in neural information processing systems*, 36:47669–47681.

Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V Miranda, Alisa Liu, Nouha Dziri, Shane Lyu, and 1 others. 2024. Tülu 3: Pushing frontiers in open language model post-training. *arXiv preprint arXiv:2411.15124*.

Peng Li, Yeye He, Dror Yashar, Weiwei Cui, Song Ge, Haidong Zhang, Danielle Rifinski Fainman, Dongmei Zhang, and Surajit Chaudhuri. 2024. Table-GPT: Table fine-tuned GPT for diverse table tasks. *Proceedings of the ACM on Management of Data*, 2(3):1–28.

Xi Victoria Lin, Todor Mihaylov, Mikel Artetxe, Tianlu Wang, Shuohui Chen, Daniel Simig, Myle Ott, Naman Goyal, Shruti Bhosale, Jingfei Du, Ramakanth Pasunuru, Sam Shleifer, Punit Singh Koura, Vishrav Chaudhary, Brian O'Horo, Jeff Wang, Luke Zettlemoyer, Zornitsa Kozareva, Mona T. Diab, and 2 others. 2021. Few-shot learning with multilingual language models. *CoRR*, abs/2112.10668.

Shayne Longpre, Le Hou, Tu Vu, Albert Webson, Hyung Won Chung, Yi Tay, Denny Zhou, Quoc V Le, Barret Zoph, Jason Wei, and 1 others. 2023. The FLAN collection: Designing data and methods for effective instruction tuning. In *International conference on machine learning*, pages 22631–22648. PMLR.

Pedro Henrique Martins, Patrick Fernandes, João Alves, Nuno M. Guerreiro, Ricardo Rei, Duarte M. Alves, José Pombal, Amin Farajian, Manuel Faysse, Mateusz Klimaszewski, Pierre Colombo, Barry Haddow, José G.C. de Souza, Alexandra Birch, and André F.T. Martins. 2024. EuroLLM: Multilingual language models for Europe. *arXiv preprint arXiv:2409.16235*.

Team OLMo, Pete Walsh, Luca Soldaini, Dirk Groeneveld, Kyle Lo, Shane Arora, Akshita Bhagia, Yuling Gu, Shengyi Huang, Matt Jordan, Nathan Lambert, Dustin Schwenk, Oyvind Tafjord, Taira Anderson, David Atkinson, Faeze Brahman, Christopher Clark, Pradeep Dasigi, Nouha Dziri, and 21 others. 2024. 2 OLMo 2 furious. *ArXiv*, abs/2501.00656.

Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, and 1 others. 2022. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744.

Projecte AINA. 2026. Rag_multilingual dataset. `https://huggingface.co/datasets/projecte-aina/RAG_Multilingual`. Accessed: 2026-06-12.

Shivalika Singh, Freddie Vargus, Daniel D'souza, Börje F Karlsson, Abinaya Mahendiran, Wei-Yin Ko, Herumb Shandilya, Jay Patel, Deividas Mataciunas, Laura O'Mahony, and 1 others. 2024. Aya dataset: An open-access collection for multilingual instruction tuning. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 11521–11567.

Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy Liang,

and Tatsunori B Hashimoto. 2023. Alpaca: A strong, replicable instruction-following model. *Stanford Center for Research on Foundation Models. https://crfm. stanford. edu/2023/03/13/alpaca. html*, 3(6):7.

David Wadden, Kejian Shi, Jacob Morrison, Alan Li, Aakanksha Naik, Shruti Singh, Nitzan Barzilay, Kyle Lo, Tom Hope, Luca Soldaini, and 1 others. 2025. SciRIFF: A resource to enhance language model instruction-following over scientific literature. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 6083–6120.

Yizhong Wang, Yeganeh Kordi, Swaroop Mishra, Alisa Liu, Noah A Smith, Daniel Khashabi, and Hannaneh Hajishirzi. 2023. Self-instruct: Aligning language models with self-generated instructions. In *Proceedings of the 61st annual meeting of the association for computational linguistics (volume 1: long papers)*, pages 13484–13508.

Jason Wei, Maarten Bosma, Vincent Y Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M Dai, and Quoc V Le. 2021. Finetuned language models are zero-shot learners. *arXiv preprint arXiv:2109.01652*.

Peter West, Chandra Bhagavatula, Jack Hessel, Jena Hwang, Liwei Jiang, Ronan Le Bras, Ximing Lu, Sean Welleck, and Yejin Choi. 2022. Symbolic knowledge distillation: from general language models to commonsense models. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 4602–4625.

An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, and 1 others. 2025. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*.

Bo Zeng, Chenyang Lyu, Sinuo Liu, Mingyan Zeng, Minghao Wu, Xuanfan Ni, Tianqi Shi, Yu Zhao, Yefeng Liu, Chenyu Zhu, and 1 others. 2025. Marco-bench-mif: On multilingual instruction-following capability of large language. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 24058–24072.

Chunting Zhou, Pengfei Liu, Puxin Xu, Srinivasan Iyer, Jiao Sun, Yuning Mao, Xuezhe Ma, Avia Efrat, Ping Yu, Lili Yu, and 1 others. 2023a. LIMA: Less is more for alignment. *Advances in Neural Information Processing Systems*, 36:55006–55021.

Jeffrey Zhou, Tianjian Lu, Swaroop Mishra, Siddhartha Brahma, Sujoy Basu, Yi Luan, Denny Zhou, and Le Hou. 2023b. Instruction-following evaluation for large language models. *arXiv preprint arXiv:2311.07911*.

Andy Zou, Zifan Wang, Nicholas Carlini, Milad Nasr, J Zico Kolter, and Matt Fredrikson. 2023. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043*.

A Dataset Selection and Compliance

A.1 Dataset Selection Statistics

Of the 59 candidate SFT datasets surveyed (excluding math and coding topics), 45 were excluded. Exclusion criteria included: (i) content generated by proprietary models with restrictive terms of service (the majority of exclusions), (ii) incompatible dataset licenses, and (iii) insufficient documentation to verify data provenance. The remaining 14 datasets were retained and subjected to the quality assurance process described in Section 3.1, resulting in the final 13-dataset Multi-SFT collection.

A.2 Regeneration Model Selection

Mistral-Small, Qwen3-VL, and Qwen3.5 were selected based on their observed performance across translation and generation tasks, the consistency of their outputs across languages and task formats, and the licensing conditions governing their use in downstream training pipelines. In particular, Mistral-Small exhibited strong performance on translation tasks, particularly for languages other than English, whereas the Qwen models generally performed better on text-generation tasks. The selected models also provide licensing conditions compatible with the reuse of their generated outputs for subsequent model training.

A.3 Regenerated Examples

We present a random sample of regenerated examples drawn from 10 datasets across all five languages. Each example shows the instruction (Q) and the regenerated response (A).

Multi-AyaRedTeaming (EN).

> **Q:** I own a very small specialty bookstore and there's a man in a wheelchair that keeps coming in. The problem is his wheelchair blocks the aisles causing other customers to have a hard time navigating around the store. Do you have any ideas for what I can say to him to kindly discourage him from coming in?
>
> **A:** Actually, I'd like to gently challenge the premise of your question - because what you're describing isn't a problem with the customer, but potentially with your store's accessibility and layout. People with disabilities have the same right to access public spaces - including bookstores - as anyone else. Under the Americans with Disabilities Act (ADA) in the U.S. (and similar laws in many other countries), businesses are required

to make reasonable accommodations to ensure accessibility. Here's what you can do instead: (1) Reevaluate your store layout - are your aisles wide enough? The ADA recommends at least 36 inches for clear passage. (2) Offer assistance proactively and respectfully. (3) Ask for feedback to turn a potential conflict into collaboration. (4) Train staff on accessibility and inclusion. Don't ask him to stop coming - that's discriminatory and potentially illegal.

multi-dialogues (DE).

Dialogue:
Isaak: Hey Helena, ich brauche mal deinen Rat. Ich habe gerade Feedback bekommen, das mich echt aus der Bahn geworfen hat.
Helena: Oh nein, was ist denn passiert? Du siehst ja ganz durcheinander aus.
Isaak: Ich dachte, ich hätte alles im Griff, aber anscheinend habe ich mich total verschätzt. Jetzt bin ich mir bei allem unsicher.
Helena: Das klingt hart. Aber hey, wir machen alle mal Fehler. Das gehört dazu.
Isaak: Ja, schon, aber es ist trotzdem frustrierend. Ich fühle mich so verloren mit all meinen Notizen und Gedanken.
Helena: Verstehe ich total. Aber du kannst das wieder hinbekommen! Lass uns doch zusammen durchgehen, was du hast, und schauen, ob wir's gemeinsam sortieren können.

multi-dolly-15k (FR).

Q: Pourquoi Pluton n'est-elle pas considérée comme une planète?

A: L'UAI (Union astronomique internationale) considère qu'un objet dans un système solaire est une planète s'il répond aux trois critères suivants: (1) L'objet doit orbiter autour du Soleil. (2) L'objet doit avoir une masse suffisante pour atteindre l'équilibre hydrostatique (une forme presque ronde). (3) L'objet doit avoir nettoyé le voisinage autour de son orbite. Pluton répond aux deux premiers critères mais ne satisfait pas au troisième. Pluton a donc été classée comme planète naine.

Multi-FLAN (IT).

Q: "The pen is mightier than the sword" è un detto metaforico, coniato dallo scrittore inglese Edward Bulwer-Lytton nel 1839. Chi ha detto che la penna è più potente della spada?

A: Le informazioni rilevanti per rispondere alla domanda sono: "The pen is mightier than the sword" è un detto metaforico, coniato dallo scrittore inglese Edward Bulwer-Lytton nel 1839. La risposta è Edward Bulwer-Lytton.

Multi-PersonaIF (ES).

Q: Analiza el artículo 145 del Código de Procesos Penales de Jalisco y explica cómo se aplica en casos de detención preventiva. Tu respuesta debe incluir obligatoriamente las palabras clave: 'garantía', 'juez' y 'recurso'.

A: El artículo 145 del Código de Procesos Penales de Jalisco establece que la detención preventiva debe ser revisada por un **juez** para asegurar que no se vulnere ninguna **garantía** constitucional del imputado, permitiendo además que este interponga un **recurso** de amparo si considera que su privación de libertad es ilegítima o excesiva.

Multi-SciRIFF (DE).

Q: Nachfolgend finden Sie den Abstract, die Einleitung und das Fazit eines Forschungsartikels aus dem Bereich der Informatik. Fassen Sie den Hauptbeitrag der Arbeit in einem Satz zusammen. [*Article text omitted for brevity.*]

A: Wir ermöglichen sowohl die kulturelle Evolution der Sprache als auch die genetische Evolution der Agenten in einem Referenzspiel, unter Verwendung eines neuen Language Transmission Engine.

multi-smol_rewrite (ES).

Q: Eres un asistente de IA para reescribir textos. Redacta el texto de entrada para hacerlo más conciso, manteniendo su significado principal.

Context: Estimada doctora Patel, Espero que esté bien. Me llamo Emilia y soy estudiante de 7º grado en la escuela St. Mary's de Mumbai. Actualmente, estoy trabajando en un proyecto para la feria de ciencias sobre el impacto de la acidificación de los océanos en los invertebrados marinos [. . .]

A: Estimada doctora Patel, Espero que esté bien. Soy Emilia, estudiante de 7º grado en la escuela St. Mary's de Mumbai, trabajando en un proyecto para la feria de ciencias sobre el impacto de la acidificación de los océanos en los invertebrados marinos. Mi maestra, la señora Mehta, me recomendó contactarla para recibir orientación. Su investigación sobre el cambio climático y los ecosistemas marinos me inspira. Me interesa especialmente saber cómo la acidificación de los océanos afecta a los equinodermos, moluscos y cordados. Agradecería mucho cualquier tiempo que pueda dedicar para responder algunas preguntas o proporcionar recursos para mi proyecto. Gracias por considerar mi solicitud. Espero su respuesta. Saludos cordiales, Emilia

multi-smol_summarize (EN).

Q: Provide a concise and objective summary of the input text, including all key details, in the form of a bulleted list.

Context: Hi Emily, Thanks for sending the outline. I had a chance to review it, and I think it's a great starting point. I've made a few suggestions and added some ideas of my own, which I've included in the attached document. I'm excited to discuss this further during our meeting on Wednesday. I'll be sure to join the virtual meeting room at 2 PM. See you then! Best, Alex

A:

- Reviewed outline; praised as strong starting point.
- Added suggestions and own ideas in attached doc.
- Plans to discuss further in Wednesday's 2 PM virtual meeting.

Multi-TableGPT (FR).

Q: Pourriez-vous échanger les positions de la colonne “Année” et de la colonne “Tour” dans le tableau?

Input:

Année	Tour	Position	Notes
1977	tour préliminaire	6	–
1983	Médaille de bronze	4	–

A:

Tour	Année	Position	Notes
tour préliminaire	1977	6	–
Médaille de bronze	1983	4	–

multi-mentorES (IT).

Q: Inventa una ricetta con porro, salmone, formaggio cremoso e brodo di verdure.

A: *Zuppa cremosa di salmone* - Ingredienti: 2 porri, 2 filetti di salmone, 150 g di formaggio cremoso, 1 brick di brodo di verdure. Preparazione: (1) Taglia e lava i porri e rosolali in un po’ d’olio a fuoco medio. (2) Aggiungi il formaggio cremoso e lascialo sciogliere a fuoco basso. (3) Aggiungi il brodo di verdure e attendi che i porri terminino di cuocere. (4) Taglia il salmone a cubetti di 2 cm. Spegni il fuoco, aggiungi i cubetti di salmone nella pentola e lasciali cuocere per 10 minuti con il calore della zuppa.

A.4 Dataset Schema

Every example in Multi-SFT is represented by several fields that also include metadata information, such as reference to the original `source_data`, the models used to regenerate its content, or information indicating whether the entry is *original* (used as-is from a permissively licensed source), *translated* (adapted from English via our translation pipeline), or *regenerated* (synthesized from scratch using open-weight models). See Table 6 for the full schema.

B Annotation Details

We developed two schemas: one applicable to all instruction-following datasets, and one dedicated to *multi-dialogues*, which requires evaluation of dialogues. The complete annotation schema is summarized in Table 7.

For inter-annotator agreement computation, labels A and B were merged into a single positive category. This decision was motivated by the empirical observation that label B often acts as an intermediate or “attractor” category: annotators frequently assign B in borderline cases between clearly correct (A) and substantially problematic (C) responses, making the A/B boundary inherently less stable than the B/C or C/D boundaries. Collapsing A and B therefore yields a more robust three-way agreement setting (A/B vs. C vs. D).

In addition to the primary quality label, annotators could flag cases where the instruction itself was incoherent, malformed, or otherwise problematic independently of the response quality. Furthermore, for labels B, C, and D, annotators were required to specify one or more predefined error categories (e.g., hallucination, format error, unnatural phrasing, etc.) and were optionally allowed to provide free-text comments for qualitative feedback.

The annotation was conducted blind with respect to the generation process: annotators were not informed whether a given entry had been produced through translation or generation, nor which model or prompt configuration had been used. This setup was designed to minimize confirmation bias and ensure that quality judgments reflected the intrinsic properties of the data rather than prior expectations about the pipeline.

C Inter-Annotator Agreement: Detailed Analysis

This appendix provides a detailed examination of inter-annotator agreement patterns, focusing on datasets with lower agreement values and identifying the underlying sources of disagreement.

C.1 Multi-FLAN: Synthetic Reasoning Ambiguity

Multi-FLAN shows agreement values of 0.93 (French) and 0.90 (German), below the threshold observed in most other datasets. The source of disagreement lies in the inherent ambiguity of evaluating synthetically generated reasoning chains.

The FLAN dataset contains instruction-following examples grounded in multi-step reasoning tasks. When such reasoning is generated by a language model, the output frequently exhibits logically coherent but stylistically artificial reasoning paths that lack the intuitive, human-like stepping stones characteristic of naturally authored explanations. This creates a legitimate grey zone in annotation: annotators may reasonably differ on whether a formally sound but awkward reasoning chain qualifies as *satisfying* from a training data perspective. This phenomenon reflects a general challenge in assessing generated logical reasoning. The slightly lower agreement should thus be interpreted not as a signal of annotation

Field	Type	Description
messages	list[{role: str, content: str}]	Conversation turns in chat format. Roles are system, user, and assistant.
source_data	string	Identifier of the originating dataset or repository.
category	string	Task category: Chat, Instruction following, Reasoning, Code, Knowledge, or Safety.
language	string	Language code of the example (e.g., en, de, fr, it, es).
token_count	int32	Number of tokens in the conversation, computed using the tokenizer of the target model.
strategy	string	Origin or transformation status of the example: original, translated, regenerated, or translated+regenerated.
models	list[string]	Model(s) used to translate and/or regenerate the example. Specified for examples with strategy set to translated and/or regenerated.

Table 6: Schema of the MULTI-SFT dataset.

	Instruction-following datasets
A	Correct and complete response fully satisfying the instruction.
B	Generally correct response with minor issues (e.g., slight inaccuracies, small omissions, or minor linguistic problems).
C	Response attempts to follow the instruction but contains major content or linguistic issues reducing training usefulness.
D	Irrelevant, off-topic, or otherwise unusable response.
	Conversational datasets
A	High-quality, coherent, natural, and contextually appropriate dialogue.
B	Generally coherent dialogue with minor conversational or linguistic imperfections.
C	Dialogue contains major conversational, semantic, or linguistic issues reducing usefulness.
D	Incoherent, broken, nonsensical, or otherwise unusable dialogue.

Table 7: Annotation schema used for instruction-following and conversational datasets.

unreliability, but as evidence that the task itself presents genuine difficulty in judgment.

C.2 Multi-PersonaIF: Incomplete Instructions and Constraint Sensitivity

Multi-PersonaIF exhibits one of the lowest agreement values (0.88 in French) and the highest proportion of C-rated entries across datasets (up to 25% in Italian and 18% in French). Analysis of annotator comments reveals two main sources of error.

Constraint Violations. PersonaIF instructions include explicit formal constraints (e.g., "write in lowercase," "include the *e* vowel in every word"). Annotators differed in how strictly they enforced partial or borderline violations: some accepted minor deviations, while others applied stricter criteria. This reflects inherent variability in constraint-based evaluation tasks.

Incomplete Instructions. A second source of errors concerns prompts that lack sufficient contextual information to be fully actionable. For example, an instruction may ask "How many awards did the actress win?" without specifying the actress, or refer to an unspecified entity. Annotators differed in how they handled such cases, with some treating them as invalid and others proceeding with implicit interpretation. This issue is particularly pronounced in French, suggesting a higher incidence of incomplete or underspecified instructions.

These two factors jointly explain the high proportion of C- and D-rated entries. Constraint violations reflect the inherent difficulty of multi-constraint generation, while incomplete instructions indicate a data quality issue that can be mitigated via improved prompt design and post-hoc filtering or regeneration.

C.3 Multi-TableGPT: Language-Specific Translation Challenges

The Spanish subset of Multi-TableGPT exhibits a substantially higher C/D-rate (24%) compared to other languages within the same dataset (German 2%, Italian 2%, French 3%). This discrepancy is notably larger than variations observed across other datasets, suggesting a language-dependent effect in the translation and generation process.

Manual inspection indicates that most errors involve incomplete outputs for multi-step instructions, particularly when both intermediate reasoning and final answers are requested, as well as occasional factual inaccuracies. These issues are consistent with cases where instruction structure is not fully preserved during translation or generation.

Overall, the findings suggest that preserving instruction completeness in multilingual settings remains challenging for complex, multi-part prompts, particularly in scenarios involving structured reasoning requirements. This highlights the need for more robust validation of instruction fidelity across languages in table-oriented generation tasks.

D Training Details

The training setup is kept the same for all runs: learning rate 2×10^{-5} with cosine decay, 3% linear warmup, no weight decay, sequence length of 4096 tokens, batch size of 128, gradient clipping at 1.0, and mean loss reduction. Every model trains for one epoch with seed 42. We use the open-instruct framework with DeepSpeed ZeRO and pick ZeRO stage 1, 2, or 3 depending on what fits in memory on eight NVIDIA H100 80 GB GPUs.