QREF: Improving DeepSearch Agents by Aligning Search with Reasoning via Query Refinement and Evidence Filtering

Yujia Chen[1,2*], Yuhui Wang[1,2*], Junzhe Wang[1,2*], Shihan Dou[1], Di Liang[2], Yujiong Shen[1], Yuhao Zhou[1], Zhiheng Xi[1], Ming Zhang[1], Senjie Jin[1], Jiayi Chen[1], Dingwei Zhu[1], Yajie Yang[1], Yuan Hua[2], Chenyao Lu[2], Tao Gui[1†], Qi Zhang[1], Xuanjing Huang[1]

[1]Fudan University [2]Tencent
tgui@fudan.edu.cn

arXiv:2610.21512v1 [cs.AI] 15 Oct 2026

Abstract

DeepSearch agents achieve strong performance on complex information-seeking tasks through iterative search and reasoning. However, they often issue intent-misaligned, under-constrained, or redundant search queries, resulting in poor evidence coverage and accumulating retrieval noise. To mitigate these limitations, we propose **QREF**, a plug-and-play **Q**uery **R**efiner and **E**vidence **F**ilter. Specifically, the Query Refiner rewrites search queries according to the current reasoning intent while ensuring complementary coverage across each search batch. The Evidence Filter retains evidence relevant to both the current search batch and the reasoning objective, reducing distracting retrieval results. We initialize the Query Refiner with supervised fine-tuning and optimize it through reinforcement learning based on interactions with a real-world web search environment, reducing ineffective queries. Comprehensive experiments demonstrate that our approach improves the performance of DeepSearch agents across four challenging benchmarks, achieving average relative Pass@1 improvements of 33.55% on WebSailor and 11.04% on TongyiDeepResearch without modifying architectures or retraining their reasoning policies. Compared with the original DeepSearch backbone without QREF, the RL-trained refiner reduces the fraction of query batches that lack discriminative retrieval anchors by 95%. Compared with the SFT-trained refiner, it also lowers the filter-based ineffective query ratio.

1 Introduction

Large language models (LLMs) have demonstrated remarkable capabilities in mathematical reasoning (Guo et al. 2025), code generation (Li et al. 2022), and logical reasoning (Wei et al. 2022). However, hallucinations and unreliable parametric knowledge can still cause factual errors on knowledge-intensive or domain-specific tasks (Wang et al. 2024). To overcome these limitations, DeepSearch agents interleave reasoning with iterative web search to progressively gather external evidence for complex knowledge-intensive tasks. Existing methods have improved search performance through enhanced planning, memory mechanisms, and reinforcement learning (Li et al. 2025; Wu et al. 2026a; Gao et al. 2025; Tao et al. 2025b). However, the quality of intermediate search queries, which strongly influences retrieval effectiveness at each search step, has received much less attention. During search, the agent must translate its evolving reasoning state into one or more executable queries. These queries serve as the interface between reasoning and retrieval, yet their quality is typically optimized only implicitly as part of the overall agent policy.

To characterize query failures in DeepSearch agents, we analyze 300 trajectories generated by WebSailor-32B (Li et al. 2025) and identify three recurring failure patterns in intermediate search queries: **Missing Discriminative Constraints**, **Low Query Specificity**, and **Redundant Query Diversification**. Missing Discriminative Constraints and Low Query Specificity reflect failures of individual queries to adequately capture the retrieval intent implied by the current reasoning state, whereas Redundant Query Diversification reflects insufficient complementary coverage of retrieval intents across queries within the same search batch. Analysis of 1,992 query batches shows that each of the three failure patterns occurs in more than 20% of batches, while 47.04% of batches contain at least one failure. Moreover, unsuccessful trajectories accumulate nearly three times as many ineffective queries as successful ones on average, with the gap emerging as early as the first search round. These findings identify intermediate query quality as a critical factor for effective retrieval and multi-hop reasoning.

Recent methods have begun to improve intermediate search behavior through process supervision or retrieval-grounded reinforcement learning (Tao et al. 2025a; Zheng et al. 2025a; Wen, Dong, and Dou 2026). Nevertheless, two limitations remain. First, query optimization is generally learned within the search agent's reasoning policy or trajectory-generation process, which limits its reuse with already-trained agents. Second, existing methods typically assess or rewrite queries individually, whereas a DeepSearch step may issue a query batch whose collective coverage matters. Moreover, query reformulation alone cannot prevent real-world search engines from returning irrelevant, redun-

*These authors contributed equally.
†Corresponding author.

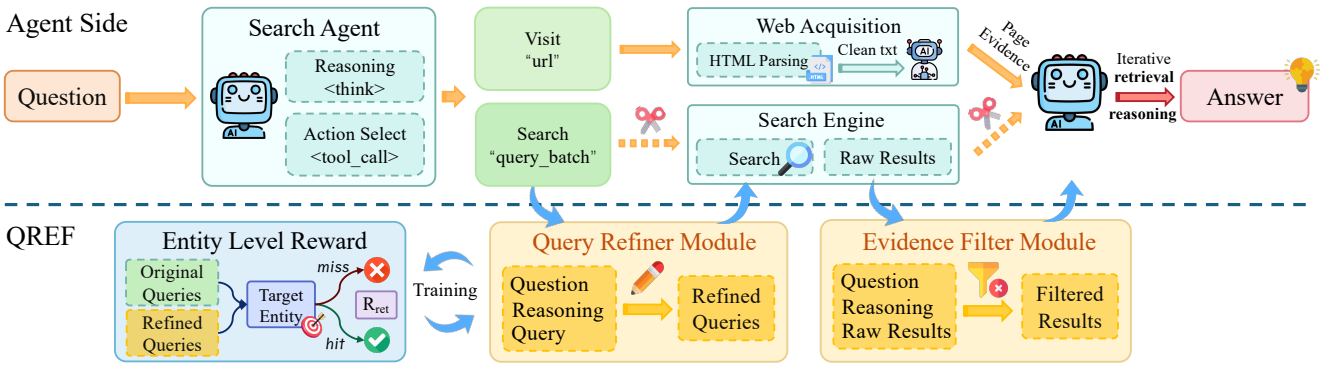

Figure 1: Overview of the QREF integrated into a tool-augmented DeepSearch agent.

dant, or overly broad evidence. Effective interface optimization therefore requires improving both directions of communication: translating reasoning states into complementary retrieval actions before search, and controlling which retrieved evidence is returned to the reasoning loop afterward.

In this work, we present **QREF**, a plug-and-play **Q**uery **R**efiner and **E**vidence **F**ilter for DeepSearch agents. Figure 1 provides an overview of QREF and its integration into the reasoning–search loop of a DeepSearch agent. Specifically, QREF consists of two components. Before retrieval, the **Query Refiner** rewrites intermediate queries into self-contained, specific, and search-effective formulations while preserving the agent's reasoning intent and ensuring complementary coverage across the search batch. After retrieval, the **Evidence Filter** retains evidence relevant to both the current query and the reasoning objective, filtering irrelevant, redundant, or overly general search results to reduce retrieval noise. Our module can be integrated into existing DeepSearch agents without modifying their architectures or retraining their reasoning policies.

We first initialize the Query Refiner by supervised fine-tuning (SFT) from Qwen3-8B (Yang et al. 2025). We then further optimize it with Proximal Policy Optimization (PPO) (Schulman et al. 2017) using an entity-level retrieval reward computed by comparing the retrieval outcomes of refined and original queries in real-world web search. Compared with the SFT model, reinforcement learning further improves Pass@1 across seven of eight backbone–benchmark settings and reduces the ineffective-query ratio across all backbone–benchmark settings.

In summary, our main contributions of this work are as follows:

- We analyze 300 WebSailor-32B trajectories, identifying three recurring failure categories in intermediate queries that each occur in over 20% of query batches and are strongly associated with lower final task accuracy.
- We propose QREF, a plug-and-play Query Refiner and Evidence Filter, enabling constraint-preserving, retrieval-oriented query reformulation and evidence filtering without modifying or retraining the underlying agent.
- We develop a two-stage training strategy for the Query Refiner, combining supervised fine-tuning with PPO under a fine-grained entity-level retrieval reward that directly optimizes evidence recovery in real-world web search.
- Experiments on four benchmarks and two backbone agents demonstrate consistent improvements in retrieval quality and downstream reasoning, while reducing ineffective and redundant queries.

2 Related Works

DeepSearch Agents The paradigm of retrieval-augmented language models has evolved from static, single-turn RAG (Lewis et al. 2020) toward autonomous DeepSearch agents that interleave long-horizon planning with iterative interaction with external search engines (Zheng et al. 2025b; Gao et al. 2025). To strengthen the reasoning–retrieval loop, one line of work leverages reinforcement learning to learn search policies directly from real-world retrieval environments (Zheng et al. 2025b) and to jointly optimize reasoning and search behaviors (Jin et al. 2025; Chen et al. 2026). A second line scales training to support longer and more effective search trajectories through improved data curation and rollout stability (Team et al. 2025a; Gao et al. 2025; Li et al. 2025; Wu et al. 2026a), while another line augments agents with structured knowledge (Lu et al. 2025), planning and memory mechanisms (Tao et al. 2025b), or hierarchical retrieval interfaces (Du et al. 2026). Despite these advances, existing methods primarily pursue *agent-level* optimization of reasoning, planning, or search policies, leaving intermediate query generation tightly coupled with the backbone agent and difficult to reuse across already-trained systems. Our work instead decouples query optimization from the backbone by introducing a backbone-agnostic, plug-and-play refinement module that can be seamlessly integrated into existing DeepSearch agents.

Query Refinement Early query refinement addressed vocabulary mismatch through linguistically motivated expansion (Selvaretnam and Belkhatir 2016) and conversational ambiguity through semantic de-contextualization (Jiang et al. 2025b). With the advent of LLMs, research shifted toward semantic-level rewriting, where the Rewrite-Retrieve-Read framework established explicit rewriting criteria and trained rewriting policies with reinforcement learning (Ye et al.

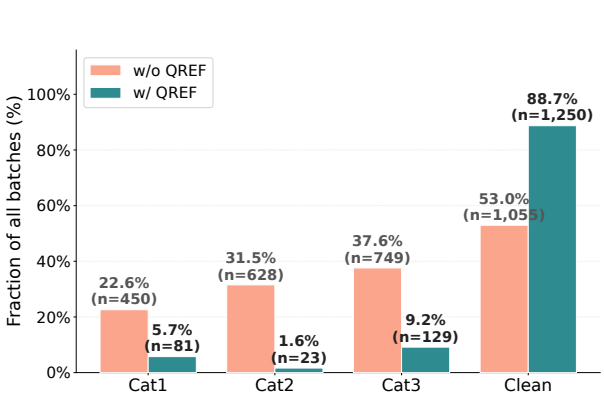

(a) Failure category distribution.

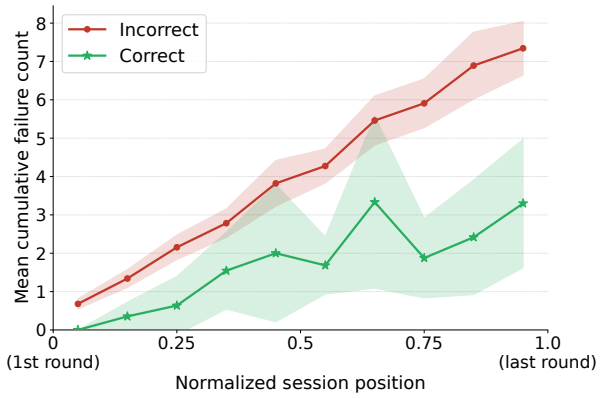

(b) Failure accumulation.

Figure 2: Analysis of ineffective queries on WebSailor. (a) Occurrence rates of the three ineffective-query categories before/after query refinement. (b) Cumulative ineffective-query counts for successful and unsuccessful trajectories.

2023), followed by preference alignment and intent-aware optimization (Zhao, Yu, and Xu 2025; Cao, Li, and Zhu 2026). These methods, however, target single-turn retrieval and do not account for queries whose intent evolves with the agent's reasoning state. More recent work extends query optimization to search agents through planning-aware query decomposition (Zhang, Feng, and Zhang 2025), query productivity estimation (Li et al. 2026), process-level supervision of intermediate query generation (Wen, Dong, and Dou 2026), and reinforcement learning that directly optimizes retrieval effectiveness (Jiang et al. 2025a). In parallel, Over-searching (Xie et al. 2026) identifies redundant low-value retrieval as a common failure mode, and DeepDive (Lu et al. 2025) attributes such failures to the semantic gap between reasoning and executable queries. However, these methods generally do not explicitly optimize the complementary coverage of query batches or the alignment of retrieved evidence with the current reasoning objective. In contrast, we cast query optimization as an independent bidirectional reasoning–retrieval interface that optimizes query generation before retrieval and evidence selection afterward, both plug-and-play with respect to the backbone agent.

3 Diagnosing Ineffective Queries

3.1 Analysis Setup

To better understand the prevalence, structure, and downstream impact of ineffective intermediate queries, we conduct a focused analysis on 300 randomly sampled WebSailor-32B trajectories from BrowseComp (Wei et al. 2025). These trajectories contain 1,992 search batches and 7,545 individual queries. The category criteria are as follows:

Cat1. Missing Discriminative Constraints The query batch fails to preserve one or more discriminative constraints already identified in the agent's current reasoning.

Cat2. Low Query Specificity The query batch lacks sufficiently specific retrieval anchors, causing the search to remain overly broad even if it is semantically relevant.

Cat3. Redundant Query Diversification Multiple queries in the batch repeatedly explore the same information need with only superficial wording differences, rather than covering complementary search directions.

Each search batch is labeled with three non-exclusive failure categories, so multiple failure types may co-occur in the same batch. GPT-4.1 serves as an LLM judge and evaluates the formulation quality of each batch conditioned on the original question and the agent's current reasoning trace. Importantly, the judge assesses the search strategy itself rather than whether the returned results already contain the correct answer. On a 100-batch subset annotated by three human raters, inter-annotator agreement is substantial for all three categories, with Fleiss' κ of 0.786, 0.667, and 0.713 for Cat1–Cat3, respectively. GPT-4.1 also agrees well with the human majority vote, achieving κ/F1 of 0.705/0.839 on Cat1, 0.900/0.950 on Cat2, and 0.770/0.909 on Cat3.

3.2 Failures Are Prevalent

We first examine how frequently ineffective query batches occur in WebSailor trajectories. As shown in Figure 2a, all three failure categories occur in more than 20% of query batches. Overall, 47.04% of query batches contain at least one failure, and 86.3% of complete trajectories contain at least one ineffective batch.

These results show that ineffective intermediate queries are not isolated errors, but a recurring failure pattern in WebSailor. Even a strong search agent frequently fails to convert intermediate reasoning into effective retrieval actions.

3.3 Failure Categories Co-occur

We then investigate whether different failure modes occur independently. Based on the binary annotations of the 1,992 query batches, we compute pairwise Pearson correlation coefficients among the three categories.

As illustrated in the supplementary material, all pairwise correlations are positive and statistically significant (all $p < 0.001$), suggesting that these failure behaviors often co-occur

rather than arise independently. In particular, *Low Query Specificity* and *Redundant Query Diversification* exhibit the strongest correlation ($r = 0.537$). This pattern suggests that batches with weak retrieval anchors are also more likely to repeat nearly identical search directions instead of covering complementary ones.

3.4 Failures Are Associated with Lower Final Accuracy

To quantify the association between ineffective queries and downstream reasoning, we compare the final correctness of trajectories with and without query failures using a χ^2 test. The underlying counts are reported in the supplementary material.

The difference is highly significant ($\chi^2 = 15.997$, $p < 0.001$): trajectories with no ineffective query batches are correct on 26.8% of examples (11/41), whereas trajectories with at least one ineffective batch are correct on only 6.2% (16/259). This result indicates that intermediate query quality is strongly associated with end-to-end reasoning performance, rather than merely with retrieval efficiency.

3.5 Failures Accumulate Over Search Rounds

We further analyze how ineffective query batches accumulate during multi-step search. Figure 2b compares cumulative query failures in trajectories that ultimately succeed or fail.

Although both groups exhibit approximately linear growth as search progresses, the two curves diverge from the first search round and continue separating throughout the trajectory. By the end of the search process, unsuccessful trajectories accumulate roughly three times more ineffective batches per trajectory than successful ones.

This observation suggests that ineffective queries rarely occur in isolation. Instead, they accumulate over successive search iterations, progressively increasing retrieval noise, reducing evidence quality, and compounding downstream reasoning errors.

4 Methodology

4.1 QREF

QREF consists of two components: the Query Refiner and the Evidence Filter. The overall framework is illustrated in Figure 1.

Query Refiner The Query Refiner is an LLM-based rewriting module that takes the original question, the current reasoning trace, and the candidate query batch parsed from the agent's `<search>` action. It follows four principles. First, *self-containment*: each rewritten query should be interpretable without omitted context or pronouns. Second, *atomic decomposition*: multi-faceted information needs are split into smaller, retrieval-friendly sub-queries. Third, *no answer leakage*: the refiner must not inject external knowledge or hidden answers beyond the current reasoning state. Fourth, *search-engine alignment*: operators such as exact-match quotes, wildcards, negation, and source restrictions (e.g., `site:`, `filetype:`) are used only when they improve retrieval precision. Beyond rewriting vague queries, the refiner removes redundant candidates and preserves already effective ones, yielding a query batch that is cleaner and more diverse than the original candidate set.

Evidence Filter The Evidence Filter is applied after the search engine returns snippets for each refined query. For every result list $D^{\text{raw}}_{t,j}$, it jointly considers the original question, the current reasoning trace, and the active refined query $\tilde{q}_{t,j}$, and predicts which result IDs should be kept. The filter retains evidence that helps identify entities, verify facts, narrow candidates, or advance the current reasoning step, while discarding results that are off-topic, overly general, redundant, or misaligned with the current query intent.

Figure 3 illustrates this pipeline on a representative BrowseComp question: the Query Refiner introduces the discriminative constraints missing from the agent's original queries, enabling the retrieval of the target evidence, while the Evidence Filter subsequently removes distracting off-topic pages so that the correct entity can be identified.

4.2 QREF Formulation

Given a question x, the goal is to produce the final answer $\hat{y}$ with a tool-augmented web agent. The backbone agent π_θ maintains an interaction history H_t and, at each step t, outputs a reasoning trace τ_t, an action $a_t \in \{\texttt{search}, \texttt{visit}, \texttt{answer}\}$, and the corresponding action content c_t.

If $a_t = \texttt{search}$, we parse from c_t a candidate query batch $Q_t = \{q_{t,1}, \ldots, q_{t,m_t}\}$. The Query Refiner R_ϕ takes the original question, the current reasoning trace, and Q_t as input, and outputs

$$o^{\text{ref}}_t = (\rho_t, \tilde{Q}_t) = R_\phi(x, \tau_t, Q_t), \tag{1}$$

where ρ_t is a short rationale and $\tilde{Q}_t = \{\tilde{q}_{t,1}, \ldots, \tilde{q}_{t,\tilde{m}_t}\}$ is the refined query batch.

The search engine executes each refined query and returns raw results

$$Z^{\text{raw}}_t = \{(\tilde{q}_{t,j}, D^{\text{raw}}_{t,j})\}_{j=1}^{\tilde{m}_t}, \tag{2}$$

where $D^{\text{raw}}_{t,j} = \{d^{\text{raw}}_{t,j,k}\}_{k=1}^{n_{t,j}}$ denotes the returned snippets for $\tilde{q}_{t,j}$, and each snippet contains a title, URL, and short text span. The Evidence Filter F_ψ evaluates each result list conditioned on x, τ_t, and the active refined query $\tilde{q}_{t,j}$, and predicts a kept-index set $I_{t,j}$. The filtered evidence is

$$Z_t = F_\psi(x, \tau_t, \tilde{Q}_t, Z^{\text{raw}}_t) = \{(\tilde{q}_{t,j}, D_{t,j})\}_{j=1}^{\tilde{m}_t}, \tag{3}$$

where $D_{t,j} = \{d^{\text{raw}}_{t,j,k} \mid k \in I_{t,j}\}$.

If $a_t = \texttt{visit}$, the agent extracts a target URL u_t from c_t and obtains page evidence $v_t = \mathcal{V}(u_t)$. The resulting observation is appended to the interaction history and used in the next step. The process continues until the agent outputs $a_t = \texttt{answer}$, from which $\hat{y}$ is extracted.

4.3 Cold-Start

Although LLMs have strong general-purpose query generation ability, direct prompting often yields unstable refinements across reasoning contexts. To obtain a reliable initialization for subsequent reinforcement learning, we first perform SFT on high-quality query refinement demonstrations distilled from GPT-5.2 (OpenAI 2025).

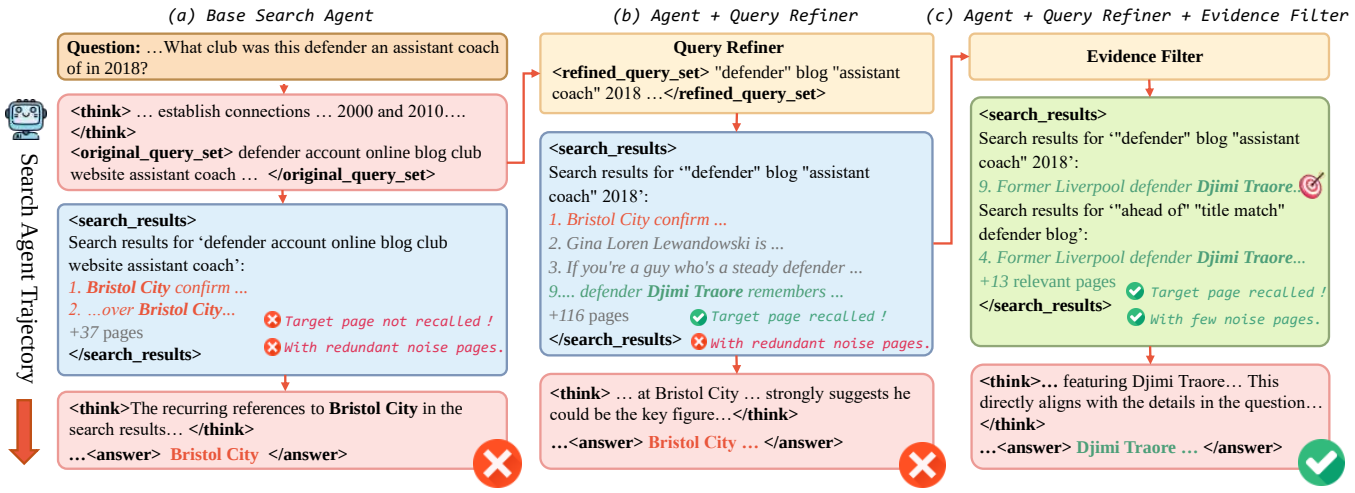

Figure 3: Case study illustrating how the proposed modules refine queries and filter retrieved evidence. (a) Vanilla Search Agent with ineffective intermediate queries. (b) Search Agent with Query Refiner. (c) Search Agent with Query Refiner and Evidence Filter.

Algorithm 1: Tool-Augmented Reasoning with the Query Refinement Module

Require: Question x, agent π_θ, refiner R_ϕ, filter F_ψ
Ensure: Answer $\hat{y}$

$H \leftarrow [m_{\text{sys}}, m_{\text{user}}(x)]$
while agent has not produced an answer **do**
  $(\tau_t, a_t, c_t) \leftarrow \text{Parse}(\pi_\theta(H))$
  if $a_t = \texttt{search}$ **then**
    $Q_t \leftarrow \text{ExtractQueries}(c_t)$
    $(\rho_t, \tilde{Q}_t) \leftarrow R_\phi(x, \tau_t, Q_t)$
    $Z_t^{\text{raw}} \leftarrow \mathcal{S}(\tilde{Q}_t)$
    $Z_t \leftarrow F_\psi(x, \tau_t, \tilde{Q}_t, Z_t^{\text{raw}})$
  else if $a_t = \texttt{visit}$ **then**
    $u_t \leftarrow \text{ExtractURL}(c_t)$
    $Z_t \leftarrow \mathcal{V}(u_t)$
  else if $a_t = \texttt{answer}$ **then**
    $\hat{y} \leftarrow \text{ExtractAnswer}(c_t)$
    return $\hat{y}$
  end if
  $H \leftarrow H \cup \{(\tau_t, c_t), Z_t\}$
end while

Specifically, we construct a query refinement dataset based on DeepResearch-9K (Wu et al. 2026b). For each search step, GPT-5.2 rewrites the candidate query batch given the original question, the current reasoning context, and the parsed queries from the agent tool call. Each training instance corresponds to one search turn in an agent trajectory and is represented as a tuple $(x_i, \tau_i, Q_i, o_i^{\text{ref}})$, where x_i is the original question, τ_i is the reasoning context extracted from the immediately preceding <think> block, Q_i is the candidate query batch parsed from the corresponding <tool_call> search action, and $o_i^{\text{ref}} = (\rho_i, \tilde{Q}_i)$ is the GPT-5.2 supervision target containing a concise rationale and the refined query batch.

The query refinement model is optimized using the standard autoregressive language modeling objective over the serialized target output. This supervised initialization teaches the model effective query refinement behavior and provides a strong starting point for subsequent reinforcement learning. Training on eight NVIDIA H20 GPUs takes approximately two hours.

4.4 Reinforcement Learning

Although supervised fine-tuning enables the model to imitate high-quality query refinement behaviors, it primarily learns the surface patterns of rewriting. It does not receive fine-grained rewards from complex and noisy web-search interactions, nor does it explicitly optimize refinement strategies for real-world retrieval or downstream retrieval effectiveness. Therefore, we further optimize the query refinement policy through reinforcement learning on the InfoSeek dataset (Xia et al. 2025) using PPO. Training on eight NVIDIA H20 GPUs takes approximately five hours and uses a total of about 100,000 queries.

At each search step, the query refinement model receives the original question, the current reasoning context, and the candidate query batch as input, and generates a refined query batch $\tilde{Q}$ that is subsequently issued to the search engine. The retrieval outcomes of $\tilde{Q}$ are compared against those of the original queries to compute a retrieval-grounded reward, encouraging the model to produce refined queries that recover more of the required evidence than the agent's own queries.

Entity-Level Retrieval Reward Instead of relying on a coarse task-success signal, we design a fine-grained reward that directly reflects how well the refined queries retrieve the evidence needed by the current reasoning step. For each training instance, we annotate a set of expected entities $\mathcal{E}$ that the agent must recover from the web at this step. Both

the original query batch Q and the refined query batch $\tilde{Q}$ are executed against the search engine, and for each entity $e \in \mathcal{E}$ we determine whether it is *hit* by a query batch.

Concretely, let $\mathcal{R}(q)$ denote the top-K retrieved snippets of query q. A single query q is said to hit entity e if a sufficient fraction of its results mention e:

$$\mathrm{hit}(q, e) = \mathbb{I}\left[\sum_{r \in \mathcal{R}(q)} \mathbb{I}[e \preceq r] \;\geq\; \max(1, \lfloor \gamma \, |\mathcal{R}(q)| \rfloor)\right], \tag{4}$$

where $e \preceq r$ denotes a normalized match of e in snippet r, and γ is a hit-ratio threshold. A query batch Q_{batch} is said to hit e if at least one $q \in Q_{\mathrm{batch}}$ hits e. We denote by $h_Q(e)$ and $h_{\tilde{Q}}(e)$ the hit indicators of the original and refined query batches, respectively.

Reward Formulation Based on $(h_Q(e), h_{\tilde{Q}}(e))$, each expected entity falls into one of four disjoint categories, capturing how query refinement changes retrieval coverage:

- **Recovered** ($h_Q(e)$=0, $h_{\tilde{Q}}(e)$=1, count N_{nh}): the refined queries retrieve an entity that the original queries miss.
- **Preserved** ($h_Q(e)$=1, $h_{\tilde{Q}}(e)$=1, count N_{hh}): both retrieve the entity.
- **Lost** ($h_Q(e)$=1, $h_{\tilde{Q}}(e)$=0, count N_{hn}): the refined queries lose an entity originally retrieved.
- **Missed** ($h_Q(e)$=0, $h_{\tilde{Q}}(e)$=0, count N_{nn}): neither retrieves the entity.

The retrieval reward is defined as a weighted per-entity average:

$$R_{\mathrm{ret}} = \frac{\alpha_{nh} N_{nh} + \alpha_{hh} N_{hh} - \alpha_{hn} N_{hn} - \alpha_{nn} N_{nn}}{N_{nh} + N_{hh} + N_{hn} + N_{nn}}, \tag{5}$$

where the weights $(\alpha_{nh}, \alpha_{hh}, \alpha_{hn}, \alpha_{nn})$ are set to $(1.0, 0.1, 0.8, 0.2)$.

Invalid-Retrieval Penalty To prevent the policy from exploiting degenerate rewrites that produce empty or malformed search results, we further impose a hard penalty. Let $\rho_\emptyset$ denote the fraction of refined queries in $\tilde{Q}$ that return no results. We clip the reward as

$$R = \begin{cases} -1, & \text{if } \tilde{Q} \text{ is unparseable or } \tilde{Q} = \emptyset, \\ -1, & \text{if } R_{\mathrm{ret}} < 0 \text{ and } \rho_\emptyset \geq \theta_\emptyset, \\ R_{\mathrm{ret}}, & \text{otherwise.} \end{cases} \tag{6}$$

Additionally, if the ground-truth annotation contains no expected entities or the denominator in R_{ret} is zero, the reward is set to 0 so that such instances contribute no gradient signal.

Through this entity-level, retrieval-grounded reward, the policy is directly optimized to enlarge the effective evidence coverage of each search step, rather than merely imitating teacher rewrites or optimizing a coarse end-task signal. In our implementation we use γ=0.3, K=10, and $\theta_\emptyset$=0.5.

5 Experiments

5.1 Evaluation Benchmarks

We evaluate our method on four challenging information-seeking benchmarks, including: (1) **BrowseComp** (1,266 questions) (Wei et al. 2025), an English web browsing benchmark requiring deep multi-hop web exploration; (2) **BrowseComp-ZH** (Zhou et al. 2025)(289 questions), the Chinese counterpart of BrowseComp requiring multi-hop exploration over Chinese web resources; (3) **GAIA** (117 questions, text-only subset) (Mialon et al. 2024), a benchmark evaluating multi-step reasoning and tool-assisted problem solving; and (4) **HLE** (2,158 questions, text-only subset) (Phan et al. 2025) a benchmark covering diverse knowledge-intensive reasoning tasks.

5.2 Backbone Agents

We evaluate the proposed module on two DeepSearch backbones: WebSailor-32B and TongyiDeepResearch (Team et al. 2025b). Both agents are served with SGLang and operate in a multi-turn ReAct-style loop that alternates between reasoning, tool use, and answer generation.

5.3 Auxiliary Models

The Query Refinement Module is instantiated with four refiner configurations: (1) *+Qwen3-8B*, a prompt-based refiner using a locally deployed open-source LLM; (2) *+GPT-5.2*, a prompt-based refiner using a proprietary API model; (3) *+SFT*, a refiner initialized through supervised fine-tuning; and (4) *+RL*, a refiner further optimized from *+SFT* via reinforcement learning. We use Qwen3-8B as the Evidence Filter. For final-answer evaluation, we follow the backbone models' evaluation pipeline and use benchmark-specific prompts.

5.4 Implementation Details

Web Search Tools The agent interacts with the web through two tools: a search tool powered by the Serper API[1] for query execution and a visit tool powered by the Jina Reader API[2] for webpage retrieval and parsing.

Inference Setup The backbone agents use a temperature of 0.6, top-p of 0.95, and a maximum context length of 31K tokens. To reduce contamination risk, we apply a rule-based URL blocklist that excludes known benchmark-hosting websites (e.g., `huggingface`) during retrieval.

5.5 Main Results

Table 1 reports *Pass@1* results. Both the trained Query Refiners (SFT and RL) and GPT-5.2 as the Query Refiner consistently improve performance over the vanilla agents across both backbones. On WebSailor, *+RL* achieves the best performance among trained refiners, improving over the baseline by 6.79, 6.23, 5.99, and 3.11 points on BrowseComp, BrowseComp-ZH, GAIA, and HLE, respectively. Similar gains are observed on TongyiDeepResearch, demonstrating the generality of the proposed module across different DeepSearch agents.

Comparing *+SFT* and *+RL*, reinforcement learning outperforms supervised fine-tuning on both backbones, suggesting that retrieval-based rewards provide a more effective optimization signal for query refinement. In contrast, directly

[1] https://serper.dev/

[2] https://jina.ai/reader/

Backbone	Method	BrowseComp			BrowseComp-ZH			GAIA			HLE		
		P@1 ↑	Turns ↓	IQR ↓	P@1 ↑	Turns ↓	IQR ↓	P@1 ↑	Turns ↓	IQR ↓	P@1 ↑	Turns ↓	IQR ↓
Tongyi DeepResearch	Baseline	39.89	46.33	-	39.66	18.10	-	74.36	10.60	-	27.61	11.53	-
	+Qwen3-8B	30.65	58.49	59.28	38.41	28.18	34.06	70.69	11.36	51.01	25.53	14.01	55.06
	+SFT	44.39	29.57	62.22	48.10	15.59	37.22	74.36	7.77	57.11	**28.08**	8.81	59.10
	+RL(Ours)	45.62	**20.46**	**36.41**	49.48	11.50	**29.36**	77.78	7.51	**31.21**	27.73	6.07	**37.25**
	+GPT-5.2	**46.29**	22.23	52.76	**51.21**	**10.15**	35.26	**79.49**	**5.97**	45.15	27.90	**5.67**	45.57
WebSailor	Baseline	10.82	**6.75**	-	19.03	4.42	-	46.15	**2.44**	-	12.09	**3.63**	-
	+Qwen3-8B	10.43	13.21	40.95	19.03	8.71	30.22	42.74	4.09	44.80	12.33	4.86	**35.61**
	+SFT	15.56	9.34	71.46	22.49	4.51	33.35	49.57	3.22	63.74	12.97	4.71	71.27
	+RL(Ours)	17.61	7.77	**39.33**	25.26	5.02	**24.25**	**52.14**	3.07	**44.16**	15.20	4.44	50.81
	+GPT-5.2	**20.14**	8.03	60.62	**25.95**	**4.39**	31.16	44.44	3.39	61.64	**15.25**	4.03	58.71

Table 1: Main results and search efficiency. P@1 denotes Pass@1. Turns indicates the average number of search turns for WebSailor and search actions for TongyiDeepResearch. IQR denotes the ineffective query ratio (%). **Bold** and underlined numbers denote the best and second-best results, respectively. ↑ indicates higher is better, and ↓ indicates lower is better.

applying an off-the-shelf LLM as the refiner is ineffective: Qwen3-8B consistently underperforms the trained refiners and even the baseline on most benchmarks, indicating that general instruction-following ability alone is insufficient. Although GPT-5.2 achieves the best overall performance, its gains are largely driven by generating substantially more queries, which improves recall at the cost of precision and introduces irrelevant retrievals that reduce search efficiency.

5.6 Further Analysis

Ablation Study We ablate the two inference-time components on WebSailor, as shown in Table 2. Removing either component degrades performance relative to the full model, indicating that both contribute to the overall effectiveness. Using the *Evidence Filter* alone still improves performance on BrowseComp but falls below the baseline on BrowseComp-ZH, suggesting that evidence filtering cannot consistently compensate for low-quality search queries. In contrast, using the *Query Refiner* alone yields consistent improvements over the baseline on both benchmarks. The full model nevertheless achieves the best performance, demonstrating that the two components provide complementary benefits.

Search Effectiveness Let Q_{ref} denote the set of all executed refined queries over the evaluation set. A query $q \in Q_{\text{ref}}$ is considered *ineffective* if, after retrieval and the Evidence Filter, it yields no retained evidence snippet for the current reasoning step. We define the **ineffective query ratio** as

$$R_{\text{ineff}} = \frac{N^{\text{ineff}}}{N^{\text{all}}}, \tag{7}$$

where N^{all} is the total number of executed refined queries and N^{ineff} is the number of queries that produce no retained evidence after filtering.

Table 1 reports the average number of search turns and the ineffective query ratio. On TongyiDeepResearch, query refinement substantially reduces search turns, indicating improved search efficiency, whereas on WebSailor it slightly increases search turns, encouraging the agent to further verify candidate answers through additional exploration. More importantly, the RL-trained refiner achieves a lower ineffective query ratio than the SFT model across all benchmarks while delivering stronger or comparable accuracy, suggesting that reinforcement learning better aligns query refinement with downstream retrieval utility by reducing low-value searches.

Method	BrowseComp-ZH	BrowseComp
WebSailor	19.03	10.82
w/o Query Refiner	15.22	13.51
w/o Evidence Filter	22.84	12.88
Full Method	**25.26**	**17.61**

Table 2: Ablation study of the proposed modules on WebSailor. *w/o Query Refiner* retains only the Evidence Filter, whereas *w/o Evidence Filter* retains only the Query Refiner.

Query Quality Analysis To further understand the effect of query refinement, we analyze the same 300 questions mentioned in Section 3. Figure 2a compares the occurrence rates of the three failure categories after applying the Query Refinement Module.

As shown in Figure 2a, *Low Query Specificity* is reduced by 95%, indicating that refined queries introduce more discriminative retrieval anchors. *Missing Discriminative Constraints* and *Redundant Query Diversification* are both reduced by approximately 3×, suggesting better preservation of reasoning constraints and more complementary search directions. We calculate the pairwise Intersection over Union (IoU) within query batches for the w/o QREF and w/ QREF settings. The average query IoU within batches dropped from 0.41 to 0.29, indicating reduced redundancy and increased diversity.

6 Conclusion

In this work, we propose **QREF**, a plug-and-play Query Refiner and Evidence Filter. The Query Refiner rewrites intermediate search queries into self-contained, constraint-aware, and complementary formulations before retrieval, and the Evidence Filter retains only evidence relevant to both the

refined query and the reasoning objective. Furthermore, we devise a two-stage training strategy that couples supervised fine-tuning with reinforcement learning under an entity-level retrieval reward in a real-world web search environment. Our method consistently improves retrieval and reasoning on four benchmarks with two backbones, delivering improved retrieval quality and downstream reasoning while substantially reducing the proportion of ineffective queries. This work highlights intermediate query quality as a key interface between reasoning and retrieval in DeepSearch agents.

References

Cao, Z.; Li, P.; and Zhu, Q. 2026. Multi-Faceted Self-Consistent Preference Alignment for Query Rewriting in Conversational Search. In *Findings of the Association for Computational Linguistics: ACL 2026*, 13083–13100.

Chen, M.; Sun, L.; Li, T.; Sun, H.; Zhu, C.; Wang, H.; Pan, J.; Zhang, W.; Chen, H.; Yang, F.; et al. 2026. Learning to reason with search for llms via reinforcement learning. *Advances in Neural Information Processing Systems*, 38: 85287–85307.

Du, M.; Xu, B.; Zhu, C.; Wang, S.; Wang, P.; Wang, X.; and Mao, Z. 2026. A-RAG: Scaling Agentic Retrieval-Augmented Generation via Hierarchical Retrieval Interfaces. *arXiv preprint arXiv:2602.03442*.

Gao, J.; Fu, W.; Xie, M.; Xu, S.; He, C.; Mei, Z.; Zhu, B.; and Wu, Y. 2025. Beyond ten turns: Unlocking long-horizon agentic search with large-scale asynchronous rl. *arXiv preprint arXiv:2508.07976*.

Guo, D.; Yang, D.; Zhang, H.; Song, J.; Wang, P.; Zhu, Q.; Xu, R.; Zhang, R.; Ma, S.; Bi, X.; et al. 2025. DeepSeek-R1 incentivizes reasoning in LLMs through reinforcement learning. *Nature*, 645(8081): 633–638.

Jiang, P.; Lin, J.; Cao, L.; Tian, R.; Kang, S.; Wang, Z.; Sun, J.; and Han, J. 2025a. Deepretrieval: Hacking real search engines and retrievers with large language models via reinforcement learning. *arXiv preprint arXiv:2503.00223*.

Jiang, Y.; Shen, L.; Niu, L.; Zhao, S.; Su, W.; and Zheng, B. 2025b. Qagent: A modular search agent with interactive query understanding. *arXiv preprint arXiv:2510.08383*.

Jin, B.; Zeng, H.; Yue, Z.; Yoon, J.; Arik, S.; Wang, D.; Zamani, H.; and Han, J. 2025. Search-r1: Training llms to reason and leverage search engines with reinforcement learning. *arXiv preprint arXiv:2503.09516*.

Lewis, P.; Perez, E.; Piktus, A.; Petroni, F.; Karpukhin, V.; Goyal, N.; Küttler, H.; Lewis, M.; Yih, W.-t.; Rocktäschel, T.; et al. 2020. Retrieval-augmented generation for knowledge-intensive nlp tasks. *Advances in neural information processing systems*, 33: 9459–9474.

Li, K.; Zhang, Z.; Yin, H.; Zhang, L.; Ou, L.; Wu, J.; Yin, W.; Li, B.; Tao, Z.; Wang, X.; et al. 2025. Websailor: Navigating super-human reasoning for web agent. *arXiv preprint arXiv:2507.02592*.

Li, Y.; Choi, D.; Chung, J.; Kushman, N.; Schrittwieser, J.; Leblond, R.; Eccles, T.; Keeling, J.; Gimeno, F.; Dal Lago, A.; et al. 2022. Competition-level code generation with alphacode. *Science*, 378(6624): 1092–1097.

Li, Z.; Solar-Lezama, A.; Yue, Y.; and Zheng, S. 2026. ENCOMPASS: Enhancing Agent Programming with Search Over Program Execution Paths. *Advances in Neural Information Processing Systems*, 38: 105751–105815.

Lu, R.; Hou, Z.; Wang, Z.; Zhang, H.; Liu, X.; Li, Y.; Feng, S.; Tang, J.; and Dong, Y. 2025. Deepdive: Advancing deep search agents with knowledge graphs and multi-turn rl. *arXiv preprint arXiv:2509.10446*.

Mialon, G.; Fourrier, C.; Wolf, T.; LeCun, Y.; and Scialom, T. 2024. GAIA: a benchmark for General AI Assistants. In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024*. OpenReview.net.

OpenAI. 2025. Update to GPT-5 System Card: GPT-5.2. Technical report, OpenAI.

Phan, L.; Gatti, A.; Han, Z.; Li, N.; Hu, J.; Zhang, H.; Zhang, C. B. C.; Shaaban, M.; Ling, J.; Shi, S.; Choi, M.; Agrawal, A.; Chopra, A.; Khoja, A.; Kim, R.; Ren, R.; Hausenloy, J.; Zhang, O.; Mazeika, M.; Yue, S.; Wang, A.; and Hendrycks, D. 2025. Humanity's Last Exam. *CoRR*, abs/2501.14249.

Schulman, J.; Wolski, F.; Dhariwal, P.; Radford, A.; and Klimov, O. 2017. Proximal Policy Optimization Algorithms. *CoRR*, abs/1707.06347.

Selvaretnam, B.; and Belkhatir, M. 2016. A linguistically driven framework for query expansion via grammatical constituent highlighting and role-based concept weighting. *Information Processing & Management*, 52(2): 174–192.

Tao, Z.; Shen, H.; Li, B.; Yin, W.; Wu, J.; Li, K.; Zhang, Z.; Yin, H.; Ye, R.; Zhang, L.; et al. 2025a. Webleaper: Empowering efficiency and efficacy in webagent via enabling info-rich seeking. *arXiv preprint arXiv:2510.24697*.

Tao, Z.; Wu, J.; Yin, W.; Zhang, J.; Li, B.; Shen, H.; Li, K.; Zhang, L.; Wang, X.; Jiang, Y.; et al. 2025b. Webshaper: Agentically data synthesizing via information-seeking formalization. *arXiv preprint arXiv:2507.15061*.

Team, T. D.; Li, B.; Zhang, B.; Zhang, D.; Huang, F.; Li, G.; Chen, G.; Yin, H.; Wu, J.; Zhou, J.; et al. 2025a. Tongyi deep-research technical report. *arXiv preprint arXiv:2510.24701*.

Team, T. D.; Li, B.; Zhang, B.; Zhang, D.; Huang, F.; Li, G.; Chen, G.; Yin, H.; Wu, J.; Zhou, J.; et al. 2025b. Tongyi DeepResearch Technical Report. *arXiv preprint arXiv:2510.24701*.

Wang, Y.; Wang, M.; Manzoor, M. A.; Liu, F.; Georgiev, G. N.; Das, R. J.; and Nakov, P. 2024. Factuality of Large Language Models: A Survey. In Al-Onaizan, Y.; Bansal, M.; and Chen, Y., eds., *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing, EMNLP 2024, Miami, FL, USA, November 12-16, 2024*, 19519–19529. Association for Computational Linguistics.

Wei, J.; Sun, Z.; Papay, S.; McKinney, S.; Han, J.; Fulford, I.; Chung, H. W.; Passos, A. T.; Fedus, W.; and Glaese, A. 2025. BrowseComp: A Simple Yet Challenging Benchmark for Browsing Agents. *CoRR*, abs/2504.12516.

Wei, J.; Wang, X.; Schuurmans, D.; Bosma, M.; Ichter, B.; Xia, F.; Chi, E. H.; Le, Q. V.; and Zhou, D. 2022. Chain-of-Thought Prompting Elicits Reasoning in Large Language

Models. In Koyejo, S.; Mohamed, S.; Agarwal, A.; Belgrave, D.; Cho, K.; and Oh, A., eds., *Advances in Neural Information Processing Systems 35: Annual Conference on Neural Information Processing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November 28 - December 9, 2022*.

Wen, T.; Dong, G.; and Dou, Z. 2026. SmartSearch: Process Reward-Guided Query Refinement for Search Agents. *arXiv preprint arXiv:2601.04888*.

Wu, J.; Li, B.; Fang, R.; Yin, W.; Zhang, L.; Wang, Z.; Tao, Z.; Zhang, D.-C.; Xi, Z.; Tang, R.; et al. 2026a. Webdancer: Towards autonomous information seeking agency. *Advances in Neural Information Processing Systems*, 38: 120957–120985.

Wu, T.; Wang, Y.; Ma, X.; He, X.; Wang, S.; Yin, D.; and Zhao, X. 2026b. DeepResearch-9K: A Challenging Benchmark Dataset of Deep-Research Agent. In Moffat, A.; Scholer, F.; Bast, H.; Najork, M.; and Zhang, M., eds., *Proceedings of the 49th International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 2026, MelbourneVICAustralia, July 20-24, 2026*, 3481–3488. ACM.

Xia, Z.; Luo, K.; Qian, H.; and Liu, Z. 2025. Open Data Synthesis For Deep Research. arXiv:2509.00375.

Xie, R.; Gopinath, D.; Qiu, D.; Lin, D.; Sun, H.; Potdar, S.; and Dhingra, B. 2026. Over-Searching in Search-Augmented Large Language Models. *arXiv preprint arXiv:2601.05503*.

Yang, A.; Li, A.; Yang, B.; Zhang, B.; Hui, B.; Zheng, B.; Yu, B.; Gao, C.; Huang, C.; Lv, C.; Zheng, C.; Liu, D.; Zhou, F.; Huang, F.; Hu, F.; Ge, H.; Wei, H.; Lin, H.; Tang, J.; Yang, J.; Tu, J.; Zhang, J.; Yang, J.; Yang, J.; Zhou, J.; Zhou, J.; Lin, J.; Dang, K.; Bao, K.; Yang, K.; Yu, L.; Deng, L.; Li, M.; Xue, M.; Li, M.; Zhang, P.; Wang, P.; Zhu, Q.; Men, R.; Gao, R.; Liu, S.; Luo, S.; Li, T.; Tang, T.; Yin, W.; Ren, X.; Wang, X.; Zhang, X.; Ren, X.; Fan, Y.; Su, Y.; Zhang, Y.; Zhang, Y.; Wan, Y.; Liu, Y.; Wang, Z.; Cui, Z.; Zhang, Z.; Zhou, Z.; and Qiu, Z. 2025. Qwen3 Technical Report. arXiv:2505.09388.

Ye, F.; Fang, M.; Li, S.; and Yilmaz, E. 2023. Enhancing conversational search: Large language model-aided informative query rewriting. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, 5985–6006.

Zhang, Z.; Feng, Y.; and Zhang, M. 2025. Levelrag: Enhancing retrieval-augmented generation with multi-hop logic planning over rewriting augmented searchers. *arXiv preprint arXiv:2502.18139*.

Zhao, S.; Yu, T.; and Xu, A. 2025. Beyond the limitation of a single query: Train your LLM for query expansion with Reinforcement Learning. *arXiv preprint arXiv:2510.10009*.

Zheng, X.; An, K.; Wang, Z.; Wang, Y.; and Wu, Y. 2025a. StepSearch: Igniting LLMs search ability via step-wise proximal policy optimization. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, 21816–21841.

Zheng, Y.; Fu, D.; Hu, X.; Cai, X.; Ye, L.; Lu, P.; and Liu, P. 2025b. Deepresearcher: Scaling deep research via reinforcement learning in real-world environments. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, 414–431.

Zhou, P.; Leon, B.; Ying, X.; Zhang, C.; Shao, Y.; Ye, Q.; Chong, D.; Jin, Z.; Xie, C.; Cao, M.; Gu, Y.; Hong, S.; Ren, J.; Chen, J.; Liu, C.; and Hua, Y. 2025. BrowseComp-ZH: Benchmarking Web Browsing Ability of Large Language Models in Chinese. *CoRR*, abs/2504.19314.

Appendix

A Case Study

Figure 3 presents a representative example illustrating how our two modules progressively improve retrieval quality and downstream reasoning.

Given the question “Ahead of a title match, a defender who played for one of the teams... What club was this defender an assistant coach of in 2018”, the vanilla DeepSearch agent first generates several vague intermediate queries such as "assistant coach defender blog club website" and "defender account online blog assistant coach". These queries omit critical constraints (e.g., the title match, the defender’s identity, and the associated event), causing the search engine to retrieve numerous irrelevant pages centered around Bristol City. Since the target evidence never appears in the retrieved documents, the agent incorrectly infers that Bristol City is the answer.

After applying the Query Refiner, the intermediate queries become substantially more search-oriented. The refiner decomposes the information need into multiple complementary queries while explicitly preserving important constraints, producing queries such as “"defender" blog "assistant coach" 2018" and “"ahead of" "title match" defender blog”. As a result, the target page discussing Djimi Traoré is successfully retrieved. However, the retrieved evidence still contains many irrelevant pages because the search results include documents matching only partial query intents. Consequently, although the correct evidence is present, the reasoning process can still be distracted by redundant information and arrive at the same incorrect conclusion.

Finally, the Evidence Filter removes documents that are inconsistent with the current reasoning objective while preserving evidence directly supporting the intermediate query. The remaining context is dominated by pages describing Djimi Traoré, allowing the reasoning process to correctly establish the connection between the defender and his role as an assistant coach in 2018, ultimately producing the correct answer.

This example demonstrates that the two modules address complementary bottlenecks in DeepSearch agents. Query Refinement primarily improves target evidence recall by generating more discriminative and search-effective queries, whereas Evidence Filtering further improves the signal-to-noise ratio of the retrieved context, enabling the downstream reasoning module to utilize the recalled evidence more reliably. Only when both modules are combined can the agent consistently retrieve the correct evidence and reason over it successfully.

B Prompt Templates

This section reports the prompts used by the Query Refiner and Evidence Filter. The placeholders enclosed in angle brackets are replaced with the corresponding trajectory information at inference time.

B.1 Query Refiner Prompt

query-refiner-prompt

You are a professional assistant skilled at rewriting complex or unclear questions into simpler, precise, and search-engine-friendly sub-questions.
Your task is to generate a set of search-engine-friendly queries based on the provided input context (Question, Current Reasoning, and Candidate Queries).
OPERATIONAL GUIDELINES:
1. **Self-Containment**:
Each query must be independent and understandable without the original context. Avoid all demonstrative pronouns (e.g., "this", "it").
2. **Atomic Decomposition**:
Break down multi-faceted information needs into distinct, atomic sub-questions.
3. **No Spoilers /Inference**:
Strictly do NOT leak potential answers from the reasoning chain into the queries. Do not use internal knowledge to bridge gaps; let the search engine do the work.
4. **Search Syntax Mastery**:
- Use ‘""’ for exact matching to precisely search for specific entities.
- Use the asterisk ‘*’ as a wildcard for unknown terms in a phrase with ‘""’(e.g., ‘"best * architecture"’).
- Use ‘-’ to disambiguate (e.g., ‘mercury -planet’ if searching for the element).
- Use ‘site:’ or ‘filetype:’ if the reasoning suggests a need for official documentation or specific formats.
5. **Evolutionary Filtering**:
Analyze the "Current candidate query set". Retain high-quality queries, refine vague ones, and eliminate redundant or overlapping ones.
CONSTRAINTS:
- **Conciseness**:
No conversational filler ("Search for...", "Find information about..."). Use information-dense keywords.
- **No Hallucination**:
Do not invent entities or facts not contained in the input.
- **Output Integrity**:
You must output ONLY a valid JSON object.
INPUT CONTEXT:
- Question: {QUESTION}
- Current reasoning: {THOUGHT}
- Current candidate query set: {QUERY}
OUTPUT FORMAT:
{
"reasoning": "Explain the strategy for the refinement, specifically why certain operators were chosen or why the candidate set was modified.",
"refined_query_set": ["sub-query-1", "sub-query-2"]
}

B.2 Evidence Filter Prompt

evidence-filter-prompt

You are a search-result filtering expert for a DeepSearch agent.
Your task is to remove noisy, irrelevant search results and retain only those that genuinely help advance the agent's current information need.
Judge usefulness strictly based on:
- The original question
- The agent's current reasoning and search intent
- The current search query
Keep results that can:
- Identify key entities,
- Align with search intent,
- Verify critical facts,
- Narrow down candidates,
- Or directly advance reasoning.
Discard results that are:
- Off-topic
- Overly general
- Background-only
- Redundant
- Not aligned with the current query intent
[Input]
Original Question: question
Current Reasoning (for internal understanding only): thought
Current Search Query: query
Search Results: results
[Output Format]
Output ONLY the following JSON, without any extra text:
{ "kept_ids": [1, 2, ...] }

C Training Details

We summarize the main training hyperparameters used for supervised fine-tuning (SFT) and reinforcement learning (RL). Machine-specific settings, such as internal paths, proxy variables, cluster startup commands, and system memory configuration, are omitted because they do not affect the training objective.

C.1 Supervised Fine-Tuning

We first perform full-parameter SFT on the Hunyuan backbone. The model is trained for one epoch with a maximum sequence length of 20,480 tokens using bfloat16 precision. We use a learning rate of 1×10^{-5} with a cosine schedule and a warmup ratio of 0.1. DeepSpeed ZeRO-3 with parameter and optimizer offloading is used to reduce memory consumption. FlashAttention-2, the Liger kernel, and gradient checkpointing are enabled for efficient long-context training. Table 3 summarizes the main settings.

Setting	Value
Fine-tuning type	Full-parameter SFT
Maximum length	20,480 tokens
Batch size	2 per device
Gradient accumulation	2 steps
Learning rate	1×10^{-5}
Training epochs	1
Scheduler	Cosine, 0.1 warmup ratio
Precision	bfloat16
Distributed optimization	DeepSpeed ZeRO-3 offload
Efficiency features	FlashAttention-2, Liger kernel, gradient checkpointing

Table 3: Main hyperparameters used for supervised fine-tuning.

C.2 Reinforcement Learning

Starting from the SFT checkpoint, we further optimize the model with proximal policy optimization (PPO) using the VERL training framework. Advantages are estimated with generalized advantage estimation (GAE), and rewards are computed by a task-specific reward function. Training uses a batch size of 64, with a maximum prompt length of 2,048 tokens and a maximum response length of 2,048 tokens. The actor and critic learning rates are 1×10^{-6} and 1×10^{-5}, respectively. We generate one rollout per prompt and train for one epoch. KL regularization is applied through the reward, while an additional actor KL loss is disabled. Table 4 lists the main RL settings.

Setting	Value
RL algorithm	PPO with GAE
Training batch size	64
Prompt / response length	2,048 / 2,048 tokens
Actor / critic learning rate	1×10^{-6} / 1×10^{-5}
PPO mini-batch size	64
Actor / critic micro-batch	1 / 2 per GPU
Rollouts per prompt	1
Training epochs	1
Entropy coefficient	0
KL configuration	KL in reward enabled; actor KL loss disabled
Parallelism	1 node, 8 GPUs; actor TP$=2$, rollout TP$=1$
Pipeline / context parallelism	PP$=1$, CP$=1$
Batching and token budget	Dynamic batching; 4,096 tokens per GPU for actor, critic, reference, and rollout log-prob computation
Rollout engine	vLLM; memory utilization 0.7; 15,000 maximum batched tokens
Memory optimization	Actor, critic, reference, gradient, and optimizer offloading

Table 4: Main hyperparameters used for PPO training. TP, PP, and CP denote tensor, pipeline, and context parallelism, respectively.

D Additional Results and Failure Analysis

This section provides complementary results for the WebSailor-based agent and additional analyses of ineffective query generation. We first report performance under three-attempt evaluation and then examine how ineffective query batches relate to trajectory correctness and to one another.

D.1 WebSailor Results under Three-Attempt Evaluation

Table 5 reports the performance of WebSailor equipped with the Query Refinement Module on BrowseComp-English (BC-EN), BrowseComp-Chinese (BC-ZH), Humanity's Last Exam (HLE), and GAIA. Across all four benchmarks, Pass@3 is higher than Avg@3, indicating that allowing multiple attempts improves the probability of obtaining at least one successful trajectory. The largest Pass@3 values are observed on GAIA, followed by BC-ZH, while HLE remains the most challenging benchmark under both metrics.

Method	BC-EN	BC-ZH	HLE	GAIA
Avg@3	18.07	23.99	15.06	48.43
Pass@3	30.15	41.18	25.81	65.51

Table 5: Additional WebSailor results with the Query Refinement Module, reported using *Avg@3* and *Pass@3*.

D.2 Ineffective Query Batches and Trajectory Correctness

To quantify the relationship between ineffective query generation and final answer correctness, we divide trajectories according to whether they contain at least one ineffective query batch. Table 6 gives the contingency counts used in the χ^2 analysis. Among trajectories without a failure batch, 11 of 41 are correct, whereas only 16 of 259 trajectories with at least one failure batch are correct. This difference shows a strong association between ineffective intermediate search behavior and unsuccessful final trajectories, while not by itself establishing a causal relationship.

Trajectory Group	Incorrect	Correct	Total
No failure batches	30	11	41
At least one failure batch	243	16	259
Total	273	27	300

Table 6: Contingency table for the χ^2 analysis of trajectory correctness versus the presence of ineffective query batches.

D.3 Relationships among Ineffective-Query Categories

We compute pairwise Pearson correlations among the three failure modes using annotations from 1,992 query batches. As shown in Figure 4, all correlations are positive and statistically significant ($p < 0.001$), indicating that the failure modes tend to co-occur. The strongest correlation is observed between *Low Query Specificity* and *Redundant Query Diversification* ($r = 0.537$).

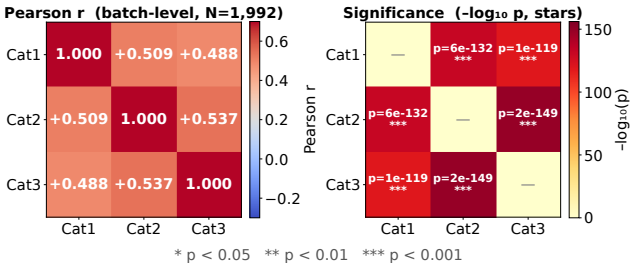

Figure 4: Pairwise Pearson correlation matrix for the three ineffective-query categories. The heatmap summarizes the extent to which different query-failure modes co-occur across trajectories.