ToolOPD: Credit-Aware On-Policy Distillation for Tool-Use Agents

Chu Zhao[1], Lei Tang[2], Yanming Wang[2], Jianzhe Zhao[1], Guibing Guo[1]*, Zhengzong Chen[2]*, Yuanyuan Zhao[2], Fei Huang[2]
[1]Northeastern University
[2]Honor Device Co., Ltd

https://github.com/MagicAgent-Search/ToolOPD

arXiv:2610.25693v1 [cs.AI] 19 Oct 2026

Abstract

On-policy distillation (OPD) enhances the capabilities of small language agents by providing dense, token-level supervision from a teacher model on student-generated trajectories. However, standard OPD implicitly assumes that credit is uniformly distributed across long-horizon trajectories, assigning the same strength of teacher supervision to different actions. Through credit-distribution analysis and outcome-reversal counterfactual experiments, we find that trajectory credit is concentrated on a small number of critical actions. Moreover, even when the trajectory is held fixed, reversing only the final outcome substantially changes the credit assigned to multiple interaction turns. These findings show that the uniform-credit assumption fails to identify critical decisions and ultimately degrades the model's reasoning accuracy. To address these challenges, we propose **ToolOPD**, an outcome-aware on-policy distillation framework designed for long-horizon tool use. ToolOPD first summarizes each completed trajectory and its outcome into a reusable hindsight skill. It then measures how this skill changes the teacher's probability of each student action token and converts the change into outcome-aware credit. The resulting credit preserves supervision that supports successful decisions or discourages failed ones, while suppressing irrelevant and conflicting signals. ToolOPD uses these credits to reweight the action-level distillation loss and jointly optimizes it with reinforcement learning. Extensive experiments across multiple long-horizon tool interaction benchmarks demonstrate that ToolOPD significantly outperforms existing reinforcement learning and self-distillation baselines.

1 Introduction

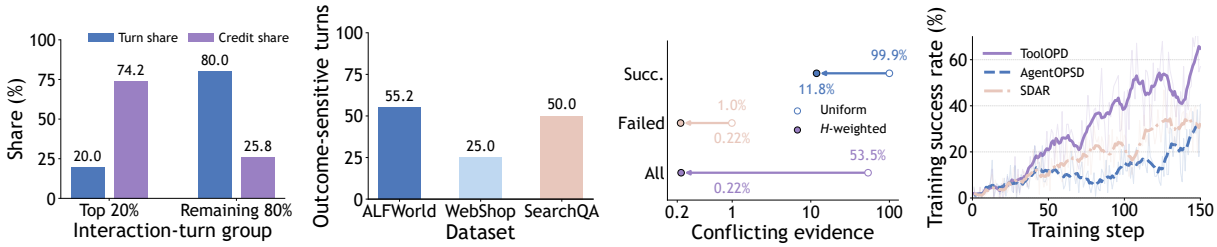

Figure 1: Empirical Study Results.

Large language models (LLMs) (Naveed et al., 2025; Minaee et al., 2024; Chang et al., 2024; Zhao et al., 2026b; Thirunavukarasu et al., 2023) have demonstrated strong capabilities in complex interactive tasks by invoking external tools (Qu et al., 2025; Xu et al., 2025; Qin et al., 2024) and adapting their decisions to environmental feedback. However, long-horizon tool-use tasks (Liu et al., 2026; Li et al., 2026b) often require tens or hundreds of interaction turns while providing only a sparse success-or-failure reward at the end of the trajectory. This makes it difficult to determine which intermediate actions contribute to the final outcome. On-policy distillation (OPD) (Agarwal et al., 2024; Li et al., 2026a; Song & Zheng, 2026; Yang et al., 2026a) addresses sparse supervision by matching a compact student policy to a teacher distribution on trajectories generated by the student itself. Standard OPD implicitly assumes that credit is uniformly distributed across long-horizon trajectories, such that different actions receive equal credit and the same strength of teacher supervision. This uniform credit assignment fails to distinguish critical tool decisions from routine transitional actions, thereby weakening effective supervision for critical decisions.

To systematically characterize this issue, we conduct a series of empirical studies. As shown in Fig. 1(a), on the ALFWorld diagnostic trajectories, the top 20% of interaction turns account for 74.2% of the total

*Corresponding authors.

credit, demonstrating that trajectory credit is highly concentrated rather than uniformly distributed. We further construct an outcome-reversal counterfactual experiment by fixing the complete task and state–action sequence while reversing only the outcome label provided to the teacher. Fig. 1(b) shows that this intervention substantially changes the credit assigned to 55.2%, 25.0%, and 50.0% of interaction turns in ALFWorld, WebShop, and SearchQA, respectively. Thus, the credit of an action depends not only on its local context but also on its relationship with the final outcome. Fig. 1(c) further shows that outcome-conflicting signals account for 53.5% of the total teacher evidence under uniform distillation. Together, these findings establish that effective long-horizon distillation must identify critical actions and determine whether the corresponding teacher signals are consistent with the trajectory outcome. Detailed settings are provided in Appendix H. Recent studies have begun to explore fine-grained credit assignment for long-horizon tasks. ATOD (Tan et al., 2026b) adaptively adjusts the distillation strength of different actions according to the discrepancy between the teacher and student policies, while TurnSight (Qu et al., 2026b) constructs multi-view hindsight information from actual tool executions and selects reliable turn-level supervision through cross-view consistency.

Unlike existing methods that determine distillation weights from local teacher–student divergence, training progress, or interaction difficulty, in this work, we propose **ToolOPD**, an outcome-aware on-policy distillation framework for long-horizon tool use. ToolOPD uses the terminal outcome to determine both the direction and strength of action credit. For each student-generated trajectory, ToolOPD conditions the teacher on the complete interaction process and its outcome to induce a concise hindsight skill that summarizes the effective workflow or failure pattern. It then compares the teacher's action probabilities with and without the hindsight skill to estimate the outcome-relevant credit of each action token. For successful trajectories, ToolOPD retains skill-induced probability increases that reinforce effective actions; for failed trajectories, it retains probability decreases that suppress erroneous actions. The resulting relevance is further calibrated by teacher confidence and experience-based reliability to prevent uncertain or historically inconsistent teacher signals from dominating the update. ToolOPD uses the calibrated credit to dynamically reweight the distillation loss over `<action>` tokens and jointly optimizes it with Group Relative Policy Optimization (GRPO) (Guo et al., 2025). This design concentrates supervision on outcome-relevant decisions while filtering redundant and conflicting signals.

Our main contributions are summarized as follows:

- **Empirical findings.** We show that credit in long-horizon OPD is concentrated on critical actions and depends on trajectory outcomes, while uniform distillation transfers outcome-conflicting supervision.
- **Method.** We propose ToolOPD, which induces outcome-conditioned hindsight skills and converts their effects on teacher predictions into fine-grained, outcome-asymmetric action credit.
- **Experiments.** Across multiple tool-use benchmarks and model scales, ToolOPD consistently outperforms representative reinforcement learning and OPD baselines in task performance, long-context robustness, and optimization stability.

2 Method

2.1 Problem Formulation

We consider multi-turn tool use tasks with a discrete set of executable actions. For a given task instance x_i, the student policy π_θ interacts with the external environment and generates a trajectory:

$$\tau_i = \{(s_{i,t}, a_{i,t}, o_{i,t+1})\}_{t=0}^{T_i-1}, \tag{1}$$

where $s_{i,t}$ denotes the interaction history and the current environmental state at step t, $a_{i,t} \in \mathcal{A}(s_{i,t})$ is the action selected by the student policy, and $o_{i,t+1}$ is the observation feedback returned by the environment. Upon the termination of the trajectory, the environment returns a scalar task reward R_i and a binary outcome indicator $y_i = \mathbb{I}(R_i > 0) \in \{0, 1\}$. In long-horizon tool use scenarios, this setting introduces severe credit assignment challenges: due to the extremely long interaction horizon and highly sparse rewards, the success or failure of a task is typically determined by only a few critical tool decisions (e.g., specific API selections or parameter settings), while the vast majority of the remaining turns consist of routine state queries and transitional actions. How to accurately attribute and assign the terminal global reward to the intermediate specific actions is the core difficulty in optimizing such policies.

We are given a teacher model π_ϕ with frozen parameters. Our objective is to distill the transferable strategy summarized by the teacher based on the complete trajectory and the final outcome into the student, while preserving the exploration capability of on-policy reinforcement learning. The trajectories during training are generated entirely by the current student policy, i.e., $\tau_i \sim \pi_\theta$. Therefore, our method falls under the category of On-Policy Distillation, rather than simple behavior cloning on offline teacher trajectories.

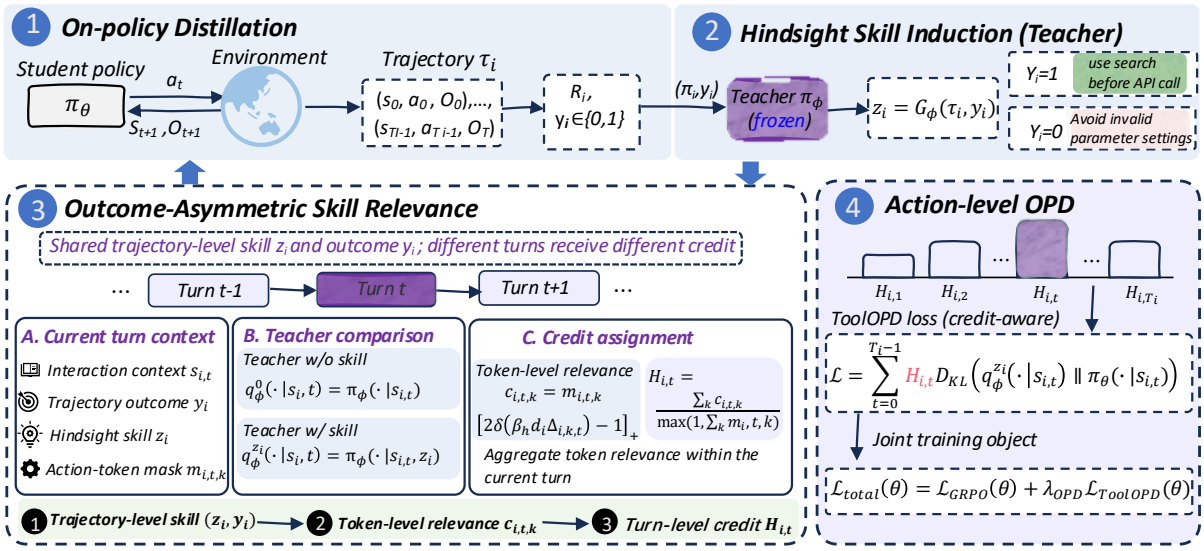

Figure 2: ToolOPD framework overview.

2.2 Overview of ToolOPD

To address the credit assignment dilemma caused by sparse rewards in long-horizon interactions, ToolOPD organically integrates GRPO with outcome-aware on-policy distillation. Its core idea is to transform sparse trajectory-level credit into dense action-level credit. The specific data flow is as follows: First, the student policy interacts with the environment to generate a set of on-policy trajectories, and the intra-group relative advantages for GRPO are computed. Subsequently, the teacher reads the complete trajectory and the final success/failure outcome to induce a concise, reusable hindsight skill. Next, by comparing the probability distributions of the teacher with and without this skill on the student's actions, the outcome relevance of the skill to the current action is calculated, thereby achieving fine-grained credit assignment. Finally, a weighted teacher-to-student OPD is performed exclusively on the action (<action>) tokens. Throughout the training process, the teacher model parameters remain frozen, and all gradients are used solely to update the student model. The framework of ToolOPD is illustrated in Fig. 2.

2.3 Outcome-Conditioned Hindsight Skill Induction

Standard distillation typically directly imitates the local action distribution of the teacher, but the teacher lacks global perception of the final success or failure of the current trajectory in this process. To address this issue, ToolOPD first introduces a hindsight skill induction mechanism, allowing the teacher to distill experience using the complete trajectory and the final outcome:

$$z_i = G_\phi(\tau_i, y_i), \tag{2}$$

where G_ϕ denotes the skill induction process of the teacher. For successful trajectories ($y_i = 1$), the teacher summarizes effective operational procedures or key decisions; for failed trajectories ($y_i = 0$), the teacher identifies the root cause of the error and provides a general rule to avoid it. To facilitate context injection, the skill is constrained to a single concise rule, formatted as: `Rule Name: imperative guidance.`

This design serves a dual purpose: first, it compresses lengthy interaction trajectories into reusable strategy priors, essentially distilling the macroscopic outcome credit at the trajectory level into microscopic strategy signals that can guide specific decisions; second, it prevents the teacher from directly rewriting the complete trajectory, thereby focusing the distillation signal on high-level decision-making patterns. All states within the same trajectory share the corresponding skill z_i for that trajectory. During the inference phase, the student can be deployed independently without the teacher or explicit skill generation, thus incurring no additional inference cost.

2.4 Outcome-Asymmetric Skill Relevance

Hindsight skills do not possess equal explanatory power for every decision within a trajectory, which precisely reflects the highly uneven credit distribution in long-horizon interactions. To measure the guiding value of a skill for the current action (i.e., to perform fine-grained credit assignment), we compare the probability distributions of the teacher on the student's actual sampled actions, both with and without

the skill. We define the base teacher distribution and the skill-conditioned teacher distribution as:

$$q_\phi^0(\cdot \mid s_{i,t}) = \pi_\phi(\cdot \mid s_{i,t}), q_\phi^{z_i}(\cdot \mid s_{i,t}) = \pi_\phi(\cdot \mid s_{i,t}, z_i). \tag{3}$$

For the k-th token $a_{i,t,k}$ in the student's action, we define its skill-induced log-probability shift as:

$$\Delta_{i,t,k} = \log q_\phi^{z_i}(a_{i,t,k} \mid s_{i,t}) - \log q_\phi^0(a_{i,t,k} \mid s_{i,t}). \tag{4}$$

We interpret this change asymmetrically based on the final outcome of the trajectory. Let the direction factor be $d_i = 2y_i - 1 \in \{-1, 1\}$, then the token-level outcome relevance is computed as:

$$c_{i,t,k} = m_{i,t,k} \left[2\sigma\left(\beta_h d_i \Delta_{i,t,k}\right) - 1\right]_+, \tag{5}$$

where $\sigma(\cdot)$ is the sigmoid function, $[\cdot]_+ = \max(\cdot, 0)$, β_h controls the mapping sensitivity, and $m_{i,t,k} \in \{0, 1\}$ is the <action> token mask.

This definition embodies outcome asymmetry and achieves dynamic credit assignment: when the trajectory is successful ($y_i = 1$), positive relevance is obtained only when the skill increases the probability of the successful action ($\Delta > 0$); when the trajectory fails ($y_i = 0$), positive relevance is obtained only when the skill decreases the probability of the failed action ($\Delta < 0$); skill influences inconsistent with the direction of the final outcome are truncated to zero. Averaging the token relevance over the current action yields the decision-level weight:

$$H_{i,t} = \frac{\sum_k c_{i,t,k}}{\max\left(1, \sum_k m_{i,t,k}\right)}. \tag{6}$$

This weight essentially quantifies the credit contribution of the current action to the final outcome. In the subsequent distillation, we directly adopt $w_{i,t} = H_{i,t}$ as the distillation weight for OPD.

2.5 Action-Level On-Policy Distillation

The standard OPD objective is to minimize the KL divergence from the teacher distribution to the student distribution over the state distribution generated by the student. Leveraging the dynamically assigned credit weights, ToolOPD can precisely pinpoint critical decisions in long-horizon trajectories and dynamically adjust the distillation intensity for different decisions:

$$\mathcal{L}_{\text{ToolOPD}}(\theta) = \mathbb{E}_{\tau_i \sim \pi_\theta}\left[\sum_{t=0}^{T_i - 1} H_{i,t} D_{\text{KL}}\left(q_\phi^{z_i}(\cdot \mid s_{i,t}) \parallel \pi_\theta(\cdot \mid s_{i,t})\right)\right]. \tag{7}$$

Since the training samples are drawn from the student policy π_θ, we employ importance sampling to estimate the aforementioned KL divergence. For the action tokens sampled by the student, we define the log-ratio $\ell_{i,t,k} = \log q_\phi^{z_i}(a_{i,t,k} \mid s_{i,t}) - \log \pi_\theta(a_{i,t,k} \mid s_{i,t})$, and compute the clipped importance weight:

$$r_{i,t,k} = \exp\left(\text{clip}(\ell_{i,t,k}, -c, c)\right), \tag{8}$$

where c is the maximum log-ratio clipping value (set to $c = 5$ in this study). The unbiased estimation form of the KL divergence for a single token is $r_{i,t,k} \log r_{i,t,k} - r_{i,t,k} + 1$. Therefore, the practically optimized ToolOPD loss is:

$$\widehat{\mathcal{L}}_{\text{ToolOPD}} = \frac{1}{N_A} \sum_{i,t,k} m_{i,t,k} H_{i,t} \left(r_{i,t,k} \log r_{i,t,k} - r_{i,t,k} + 1\right), \tag{9}$$

where $N_A = \sum_{i,t,k} m_{i,t,k}$ denotes the total number of action tokens in the batch. This design ensures that distillation is applied exclusively to the executable decision tokens within <action>, without interfering with natural language reasoning processes such as <think>, thereby allowing the student to retain its own reasoning expression paradigm.

2.6 Joint Optimization with GRPO

To preserve the exploration and policy improvement capabilities driven by environmental rewards, we adopt GRPO as the foundational reinforcement learning objective. For G trajectories sampled for the same task, the group-relative advantage is computed as:

$$A_i = \frac{R_i - \text{mean}(\{R_j\}_{j=1}^G)}{\text{std}(\{R_j\}_{j=1}^G) + \epsilon}. \tag{10}$$

Table 1: Performance comparison across ALFWorld, Search QA, and WebShop benchmarks.

Method	ALFWorld							Search QA								WebShop	
	Pick	Look	Clean	Heat	Cool	Pick2	**Avg.**	NQ	Triv	Pop	Hotp	2Wk	MuS	Bam	**Avg.**	Score	**Succ.**
Teacher Model: Qwen3-4B-Instruct																	
GRPO	79.3	50.0	16.1	63.2	21.7	30.0	43.4	45.4	64.1	48.2	41.6	46.1	17.7	19.8	40.4	71.4	46.1
Distilled Student Model: Qwen3-0.6B																	
Vanilla	0.0	0.0	0.0	0.0	0.0	5.0	0.8	11.9	16.4	13.3	10.5	20.7	1.3	1.6	10.8	34.8	2.3
OPD	71.4	38.5	14.8	**50.0**	4.0	0.0	29.8	7.2	9.3	12.4	6.2	4.1	0.8	2.0	6.0	3.7	0.8
GRPO	17.1	15.4	29.6	18.8	16.0	8.3	17.5	**39.6**	54.1	**44.2**	**32.8**	**29.5**	**9.6**	11.7	31.6	54.9	10.9
OPSD	4.1	5.5	6.4	4.3	0.0	0.0	3.4	0.0	0.2	0.1	0.1	0.3	0.0	0.0	0.1	36.7	0.7
TCOD	11.4	23.1	7.4	0.0	4.0	4.2	8.3	17.9	19.0	19.6	14.4	12.5	5.3	0.0	12.7	9.0	5.5
SDAR	50.0	33.3	40.6	30.8	28.6	**29.2**	35.4	36.6	49.4	40.4	28.5	28.3	6.6	20.6	30.1	44.3	10.9
Skill-SD	44.4	33.3	28.1	23.1	14.3	12.5	26.0	38.8	53.5	43.3	29.4	27.8	7.8	12.1	30.4	70.6	49.4
AgentOPSD	58.6	50.0	54.8	36.8	13.0	20.0	38.9	32.8	48.1	41.5	21.1	21.4	4.1	3.2	24.6	63.3	24.2
ATOD	**91.2**	53.3	17.4	42.9	25.0	6.7	39.4	32.8	50.3	38.1	27.3	25.4	6.7	8.9	27.1	34.8	23.4
Ours	65.1	**75.0**	**66.7**	33.3	**79.0**	22.2	**56.9**	39.2	**54.6**	43.1	31.8	26.2	8.6	**23.3**	**32.4**	**80.1**	**58.6**
Distilled Student Model: Qwen3-1.7B																	
Vanilla	24.1	33.3	3.2	0.0	0.0	5.0	10.9	29.4	46.9	37.0	23.5	19.6	6.4	10.5	24.8	71.4	46.1
OPD	65.7	38.5	18.5	50.0	20.0	8.3	33.5	8.1	18.7	13.2	6.2	4.6	1.4	5.6	8.3	26.5	2.3
GRPO	70.0	44.4	42.8	47.4	26.3	17.4	41.4	40.0	58.9	43.5	35.4	30.3	12.0	65.7	40.8	72.1	40.6
OPSD	22.2	22.2	3.1	0.0	14.3	4.2	11.0	4.2	8.3	4.6	6.6	15.3	0.7	1.2	5.8	42.1	3.9
TCOD	42.9	38.5	3.7	12.5	0.0	12.5	18.3	23.9	35.7	25.0	19.1	18.8	1.4	0.0	17.7	44.1	5.5
SDAR	73.5	25.0	**76.9**	33.3	40.0	36.8	47.6	39.7	58.9	45.3	**35.9**	**35.3**	12.6	65.3	41.9	74.7	54.4
Skill-SD	52.9	37.5	69.2	42.9	60.0	36.8	49.9	39.1	57.5	**45.4**	34.8	34.1	10.7	64.1	40.8	**81.8**	53.9
AgentOPSD	44.8	**66.7**	29.0	63.2	30.4	35.0	44.8	41.1	55.9	43.2	33.4	28.7	10.7	62.1	39.3	65.4	36.7
ATOD	**86.2**	**66.7**	35.5	**68.4**	8.7	10.0	45.9	32.6	46.6	32.0	23.0	19.3	7.1	18.1	25.5	43.0	35.1
Ours	82.4	45.0	72.7	57.1	**70.6**	**47.6**	**62.6**	**42.1**	**59.3**	**45.4**	**35.9**	33.9	**13.1**	**68.6**	**42.6**	77.6	**55.7**

Defining the token-level probability ratio between the current policy and the sampling policy as $\rho_{i,t,k}(\theta) = \frac{\pi_\theta(a_{i,t,k}|s_{i,t})}{\pi_{\theta_{\text{old}}}(a_{i,t,k}|s_{i,t})}$, the clipped policy loss for GRPO is:

$$\mathcal{L}_{\text{GRPO}} = -\mathbb{E}_{i,t,k}\left[\min\left(\rho_{i,t,k}A_i, \text{clip}(\rho_{i,t,k}, 1-\epsilon_l, 1+\epsilon_h)A_i\right)\right]. \tag{11}$$

Ultimately, we combine the reinforcement learning signal from GRPO with the distillation signal from ToolOPD, defining the overall training objective as:

$$\mathcal{L}_{\text{total}}(\theta) = \mathcal{L}_{\text{GRPO}}(\theta) + \lambda_{\text{OPD}}\widehat{\mathcal{L}}_{\text{ToolOPD}}(\theta), \tag{12}$$

where λ_{OPD} is the distillation loss weight. Under this configuration, GRPO is responsible for learning the foundational task policy from environmental rewards and maintaining exploration, while ToolOPD precisely addresses the fine-grained credit assignment problem in long-horizon tool use through outcome-conditioned teacher supervision. The smaller distillation coefficient ensures that the teacher signal serves as an auxiliary regularizer, guiding the direction of policy optimization without overriding the primary gradients of reinforcement learning. The pseudocode of ToolOPD is provided in Appendix E, and the prompt templates used are presented in Appendix G

3 Experiments

3.1 Experimental Setup

Benchmarks. We evaluate our method on two major categories of benchmarks to comprehensively assess the model's performance across tool-use scenarios of varying complexity. The first category consists of classic agent interaction benchmarks, including the text-embodied environment ALFWorld (Shridhar et al., 2020) , SearchQA (Jin et al., 2025) covering single-hop and multi-hop question answering, and the interactive online shopping environment WebShop (Yao et al., 2022). The second category comprises benchmarks specifically designed for complex tool invocation, including the FTRL Ye et al. (2026) dataset, which covers single-step, sequential, and parallel tool-calling scenarios, and the BFCL (Patil et al., 2025) and ToolHop (Ye et al., 2025) dataset, which focus on evaluating long-horizon reasoning and execution capabilities with multi-step tool dependencies. Detailed descriptions are provided in the Appendix A.

Table 2: Performance of the teacher and distilled student models on FTRL, BFCL, and ToolHop.

Method	FTRL				BFCL							ToolHop	Avg.
	Solve-F1	Solve-R	Solve-P	Avg.	Base	MF	MP	LC	MT.OA	Memory	Avg.	Acc.	
Teacher: Qwen3-4B-Instruct													
GRPO	33.9	33.3	36.7	34.6	39.5	29.5	32.5	27.5	32.3	17.9	25.1	30.0	29.9
Student: Qwen3-0.6B													
Vanilla	9.3	8.5	40.0	19.3	4.0	2.0	3.0	2.5	2.9	9.3	6.1	13.0	12.8
GRPO	13.3	13.7	25.4	17.5	3.5	2.0	2.0	2.5	2.5	5.8	4.2	17.9	13.2
ToolRL	**16.2**	15.8	20.4	17.5	2.5	0.0	1.5	1.5	1.4	8.4	4.9	16.9	13.1
MatchTIR	14.0	14.9	18.9	15.9	4.0	1.5	3.0	**3.5**	3.0	6.5	4.7	17.8	12.8
TurnSight	12.3	12.8	18.7	14.6	1.5	0.5	1.0	0.5	0.9	5.4	3.1	17.5	11.7
SDAR	16.0	**16.7**	36.8	23.2	4.5	3.0	1.5	1.5	2.6	8.4	5.5	17.7	15.5
Ours	**16.2**	15.8	**42.7**	**24.9**	**5.0**	**4.0**	**4.5**	**3.5**	**3.3**	**10.0**	**7.0**	**18.5**	**16.8**
Student: Qwen3-1.7B													
Vanilla	19.9	22.5	21.1	21.2	15.5	11.0	12.0	7.5	11.8	9.3	10.4	20.4	17.3
GRPO	22.6	23.6	25.3	23.8	18.0	15.5	12.0	11.5	14.3	9.9	12.1	23.8	19.9
ToolRL	20.4	22.3	21.0	21.2	16.0	11.0	7.0	8.0	11.1	11.8	11.0	19.9	17.5
MatchTIR	21.6	24.5	22.7	23.0	13.0	14.5	15.0	10.0	13.1	8.8	11.0	20.3	18.1
TurnSight	21.7	**25.6**	21.3	22.9	**24.0**	16.0	**20.0**	13.0	**18.3**	8.2	13.2	23.0	19.7
SDAR	22.4	**25.6**	22.8	23.6	18.0	**17.0**	14.0	10.0	14.9	5.8	10.3	23.1	19.0
Ours	**22.7**	23.9	**28.4**	**25.0**	19.0	13.0	13.0	**16.0**	13.4	**12.4**	**13.6**	**27.8**	**22.2**

Baselines. We compare ToolOPD against three categories of methods: (1) **Training-free**: Vanilla (the instruction-tuned model without task-specific training). (2) **Outcome RL and General (Self-) Distillation**: including GRPO (Guo et al., 2025), ToolRL (Qian et al., 2026), and MatchTIR (Qu et al., 2026a) which optimize outcome rewards, as well as OPD, TCOD (Wang et al., 2026b), OPSD (Zhao et al., 2026a), ATOD (Tan et al., 2026b), and TurnSight (Qu et al., 2026b) which employ general distillation without explicit skill construction. (3) **Skill-Conditioned Self-Distillation**: Skill-SD (Wang et al., 2026a), SDAR (Lu et al., 2026), and AgentOPSD (Wang et al., 2026c), which incorporate skill guidance or trajectory reflection mechanisms. Details of each baseline are provided in the Appendix B.

Training Details. All experiments are conducted on 8 A800 GPUs. We adopt Qwen3-0.6B and Qwen3-1.7B as student models, and a Qwen3-4B model trained with GRPO (150-step checkpoint) as the teacher model. In Table 1, the training and optimization setups for ALFWorld, SearchQA, and WebShop are kept consistent with the AgentOPSD baseline. In Table 2, the models are trained exclusively on the FTRL dataset, with the remaining datasets used for out-of-distribution (OOD) generalization evaluation, and their base parameter settings aligned with TurnSight. Detailed hyperparameter configurations are provided in the Appendix C. The evaluation metrics used in this paper are presented in Appendix I.

3.2 Main Results

Tables 1 and 2 present the performance of ToolOPD across various tool-use benchmarks. Overall, ToolOPD significantly outperforms existing reinforcement learning and self-distillation baselines in all evaluated scenarios, validating its effectiveness in addressing the credit assignment problem in long-horizon tool use.

Credit Assignment Breakthrough in Long-Horizon Interactions (Table 1). In long-horizon interaction tasks such as ALFWorld, SearchQA, and WebShop, task success or failure is typically determined by only a few critical tool decisions. Compared to pure reinforcement learning methods (e.g., GRPO), ToolOPD achieves substantial improvements (e.g., the success rate of Qwen3-0.6B on WebShop jumps from 10.9% to 58.6%), demonstrating that dense and precise token-level supervision can effectively overcome the credit assignment failure caused by terminal sparse rewards. Furthermore, compared to standard OPD, which implicitly assumes uniform weights (tending to distribute gradients across redundant actions), and AgentOPSD, which lacks outcome filtering, ToolOPD precisely quantifies and amplifies the weights of critical actions that substantially contribute to the final outcome through the outcome-asymmetric mechanism ($H_{i,t}$), fundamentally overcoming the dilution of core decision signals.

Complex Tool Dependencies and Generalization Capabilities (Table 2). Across the FTRL, BFCL, and ToolHop benchmarks, ToolOPD not only achieves the highest average performance on the FTRL training set but also demonstrates strong generalization capabilities on out-of-distribution (OOD) test sets such as ToolHop and BFCL. Particularly in the ToolHop task, which features extremely long interaction turns and a large number of ineffective retrievals, ToolOPD significantly outperforms the second-best

Table 3: Ablation study on the ALFWorld benchmark. We progressively add components to GRPO: standard action-level distillation (Vanilla OPD), hindsight skill conditioning (Skill-OPD), and the full ToolOPD. The best results in each column are highlighted in bold.

Method	Pick	Look	Clean	Heat	Cool	Pick2	Scc.
GRPO	17.1	15.4	29.6	18.8	16.0	8.3	14.8
GRPO + Vanilla OPD	55.6	66.7	62.5	30.7	**92.9**	25.0	53.9
GRPO + Skill-OPD	**82.4**	50.0	46.2	19.1	60.0	**52.6**	54.7
ToolOPD (Ours)	65.1	**75.0**	**66.7**	**33.3**	79.0	22.2	**59.4**

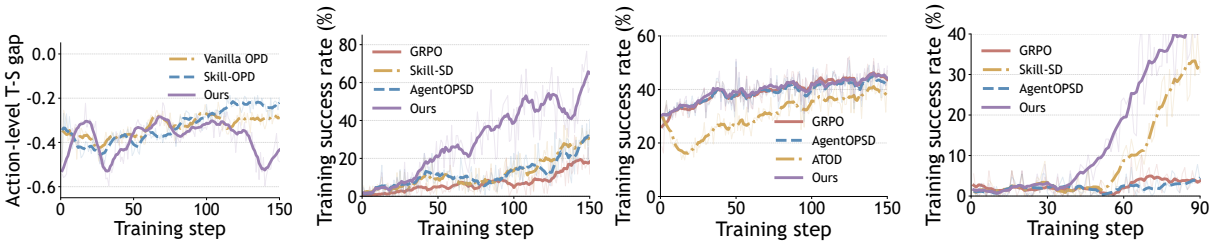

(a) Teacher-student gap. (b) ALFWorld success rate. (c) SearchQA success rate. (d) WebShop success rate.

Figure 3: Training dynamics across different tasks.

baseline, SDAR (improving the accuracy of Qwen3-1.7B from 19.0% to 22.2%). This is attributed to ToolOPD's hindsight skill induction mechanism, which compresses macroscopic trajectory outcomes into microscopic strategy priors, combined with asymmetric relevance to truncate gradients for ineffective actions. This ensures that the model can precisely pinpoint core API calls within long-horizon trajectories, effectively enhancing reasoning and execution capabilities under complex state transitions.

3.3 Ablation Study

To investigate the contribution of each component in ToolOPD to long-horizon credit assignment, we conduct a progressive ablation study on the ALFWorld benchmark, as reported in Table 3. GRPO, which relies solely on sparse trajectory-level rewards, achieves an average success rate of only 14.8%, highlighting the limitation of sequence-level supervision in identifying pivotal decisions during long-horizon tool interactions. Incorporating standard action-level distillation substantially improves the average success rate to 53.9%, demonstrating that dense teacher supervision over action tokens provides effective local learning signals. However, Vanilla OPD assigns relatively uniform distillation weights across interaction steps and therefore cannot adequately distinguish pivotal decisions from redundant actions. Further introducing hindsight skill conditioning yields only a marginal improvement to 54.7%, with performance degradation on several task categories. This result suggests that skill conditioning alone does not guarantee consistent gains, as the model must additionally determine whether the induced skill is aligned with the trajectory outcome and whether the corresponding teacher signal is reliable.

The complete ToolOPD combines outcome-asymmetric skill relevance $H_{i,t}$ with teacher confidence and experience-based reliability to emphasize actions that are more relevant to the final outcome. It achieves the best average success rate of 59.4%, outperforming Vanilla OPD and Skill-OPD by 5.5 and 4.7 percentage points, respectively. The largest gains over Skill-OPD occur on Look, Clean, and Heat, with improvements of 25.0, 20.5, and 14.3 points. These results show that outcome-aware relevance and reliability calibration are critical for robust credit assignment in long-horizon tool interactions.

3.4 Training Dynamics

Figure 3 presents the training dynamics of different methods across multiple tool-interaction benchmarks. ToolOPD achieves faster and more sustained performance improvements on ALFWorld and WebShop, establishing a clear advantage over the baselines in the later stages of training. Although the differences among methods are smaller on SearchQA, ToolOPD still attains the best final performance. These results suggest that its benefits are not specific to a particular environment, but generalize across tasks with different interaction horizons and reward structures. Meanwhile, the teacher–student gap in Figure 3(a) is not monotonically minimized toward zero, but instead changes dynamically throughout training. This behavior indicates that ToolOPD does not uniformly imitate all teacher signals; rather, it adaptively adjusts the strength of distillation according to the current policy state and trajectory outcome.

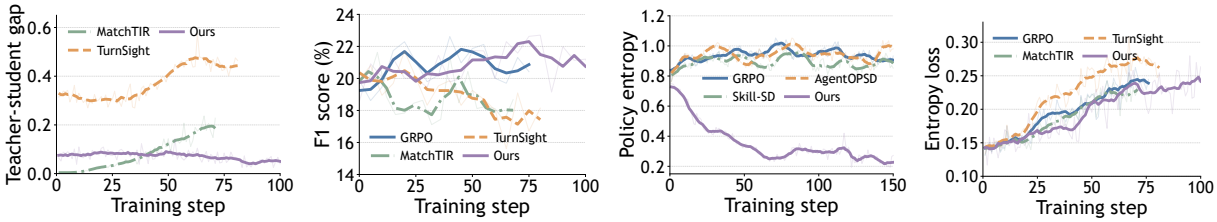

(a) Teacher-student gap. (b) FTRL F1 score. (c) Policy entropy. (d) FTRL entropy loss.

Figure 4: Training dynamics of the model across ALFWorld (c) and FTRL (a, b, and d).

Fig. 4 further characterizes the optimization behavior underlying these performance gains. On FTRL, ToolOPD maintains a consistently small and stable teacher–student discrepancy, whereas the discrepancies of MatchTIR and TurnSight progressively increase during training. Correspondingly, ToolOPD improves its F1 score more rapidly and reaches a higher level in the middle and later stages. On ALFWorld, the policy entropy of ToolOPD decreases substantially while its success rate continues to improve, suggesting that the policy gradually shifts from broad exploration toward more concentrated tool decisions rather than merely suffering from performance degradation. The entropy dynamics on FTRL exhibit a similar pattern: ToolOPD shows a smoother increase in entropy and avoids the pronounced growth in uncertainty observed for TurnSight. Overall, these results indicate that outcome-aware dynamic distillation enables ToolOPD to control teacher–student discrepancy while progressively focusing optimization on critical actions, leading to a more stable and efficient policy learning process.

3.5 Hyperparameter Analysis of ToolOPD

To assess the effect of the ToolOPD coefficient λ_{OPD}, we conduct a sensitivity analysis on SearchQA using both Qwen3-0.6B and Qwen3-1.7B. As shown in Fig. 5, both models exhibit a similar inverted-U-shaped trend: validation success first improves as λ_{OPD} increases to 0.005, and then declines when the coefficient is further increased. Specifically, the 0.6B model achieves its best success rate of 38.58% at $\lambda_{\text{OPD}} = 0.005$, outperforming the settings of 0.001 (38.11%) and 0.01 (37.24%). The 1.7B model follows the same pattern, reaching 42.94% at 0.005, while performance decreases to 41.66% and 39.20% at 0.01 and 0.02, respectively. These results suggest that a moderate action-level distillation weight provides effective dense supervision: a coefficient that is too small is insufficient to leverage ToolOPD for credit assignment, whereas an overly large coefficient can cause teacher supervision to dominate policy optimization. Based on this trend, we use $\lambda_{\text{OPD}} = 0.005$ as the default setting for SearchQA.

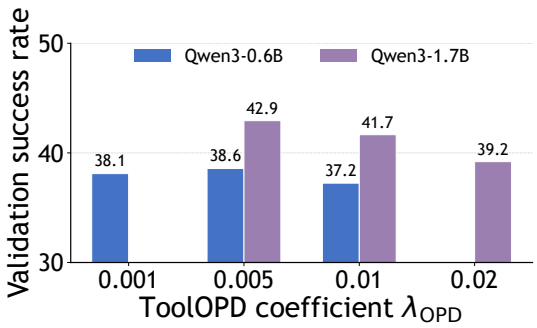

Figure 5: Hyperparameter sensitivity analysis.

3.6 Further Analysis of Credit Assignment

To further examine the credit assignment capability of ToolOPD, we analyze both performance gains and supervision quality. As shown in Fig. 6(a), ToolOPD achieves an average success rate of 59.38%, outperforming Vanilla OPD and Skill-OPD by 5.47 and 4.69 percentage points, respectively. This result indicates that dense distillation or hindsight skills alone are insufficient to achieve optimal performance. Fig. 6(b) further explains the source of this improvement. When considering only the top 10%–20% highest-credit turns, the proportion of outcome-aligned signals reaches approximately 71.3%, substantially higher than random turns (approximately 50.0%) and bottom-credit turns (approximately 37.6%). As the retained proportion increases, the differences gradually diminish and eventually converge. These results demonstrate that ToolOPD effectively ranks turns according to supervision quality and prioritizes

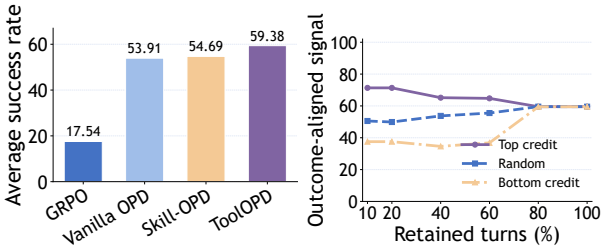

Figure 6: Further analysis of Tool-OPD. **(Left)** Performance under different credit-assignment ablation settings. **(Right)** Quality of training signals with different top-k selections.

distillation signals associated with tool decisions that are more likely to affect the final outcome. Taken together, Fig. 6 (a) and (b) suggest that the performance gains of ToolOPD primarily arise from fine-grained, outcome-aware credit assignment, rather than merely introducing additional teacher supervision.

4 Related Work

4.1 On-policy Distillation

On-policy distillation (OPD) (Song & Zheng, 2026) aligns the student and teacher policies by matching the student policy to the teacher distribution over its self-generated trajectories. By utilizing objectives such as token-level KL divergence, OPD effectively mitigates the state distribution shift inherent in offline distillation. As OPD is extended to long-horizon agent tasks, research has transitioned from uniform token-level matching to step- or turn-level supervision. In Toolintegrated scenarios, existing approaches primarily operate along two dimensions. First, they leverage dynamic skills or hindsight information from completed trajectories to guide self-distillation (including Skill-SD (Wang et al., 2026a), SEED (Wu et al., 2026), and TurnSight Qu et al. (2026b)). Second, they dynamically adjust the distillation weights across interaction turns based on factors including teacher-student divergence, training budgets, and interaction difficulty (including SOD (Zhong et al., 2026), TurnOPD (Zhou et al., 2026), AgentOPSD (Wang et al., 2026c), and HIDIFFTIR (Guo et al., 2026)). These studies demonstrate that online distillation for Toolusing agents must model multi-turn interaction structures, tool invocation difficulty, and hindsight information, rather than relying on uniform token-level supervision.

In contrast to existing methods, this paper focuses on the credit assignment problem for small models executing tool invocations in long-horizon tasks on resource-constrained terminal devices. We propose ToolOPD, which moves beyond the reliance on local training difficulty or teacher-student divergence by directly evaluating the true contribution of tool decisions based on final task outcomes. The core of ToolOPD lies in performing non-uniform credit assignment for distinct tool decisions and dynamically adjusting the distillation intensity of action tokens accordingly, thereby maximizing the distillation efficacy of small models under resource constraints.

4.2 Credit Assignment in Long-Horizon On-Policy Distillation

In long-horizon tool use tasks (Feng et al., 2026; Jiang et al., 2025; Qian et al., 2026), standard OPD typically performs uniform token-level distribution matching across all visited states. However, because different interaction turns contribute unequally to task success, applying supervision with uniform intensity is highly prone to introducing noise. Recent studies have introduced fine-grained credit assignment into long-horizon interactions and policy distillation. Some methods evaluate intermediate action values by quantifying information gain, leveraging hindsight information, or decomposing trajectory-level outcomes (e.g., ΔBelief-RL (Auzina et al., 2026), HCAPO (Tan et al., 2026a), and PBSD (Tian et al., 2026)). Others identify critical action spans and adjust local advantages from the perspectives of supervision reliability, failure correlation, or policy comparison (e.g., Prune-OPD (Yang et al., 2026b), HINT-SD (Yeo et al., 2026), UCOB (Tu et al., 2026), and SGCD (Ding et al., 2026)). Overall, existing research has transitioned from uniform token-level matching to credit-aware selective supervision.

In this work, ToolOPD locates critical tool decisions by quantifying the shift in the teacher's action distribution conditioned on hindsight skills. It then performs multi-dimensional credit calibration by integrating the final outcome, trajectory consistency, and teacher confidence. Ultimately, the calibrated turn-level credit is directly mapped to the action tokens actually sampled by the student, achieving non-uniform on-policy distillation oriented towards long-horizon task outcomes.

5 Conclusion

We introduce ToolOPD, an outcome-aware on-policy distillation framework that addresses credit assignment in long-horizon tool-use tasks. Our empirical analysis shows that the uniform supervision of standard OPD fails to distinguish the contributions of different tool decisions and may introduce teacher signals that conflict with the final outcome. To address this limitation, ToolOPD induces hindsight skills from complete trajectories and their outcomes, and dynamically assigns credit to action tokens through outcome-asymmetric skill relevance. This preserves supervision consistent with the task outcome while suppressing irrelevant or conflicting signals. Experiments across multiple long-horizon tool-interaction benchmarks and model scales demonstrate that ToolOPD consistently outperforms reinforcement-learning and on-policy-distillation baselines while exhibiting more stable optimization over extended interactions.

References

Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos Garea, Matthieu Geist, and Olivier Bachem. On-policy distillation of language models: Learning from self-generated mistakes. In *International Conference on Learning Representations*, volume 2024, pp. 21246–21263, 2024.

Ilze Amanda Auzina, Joschka Strüber, Sergio Hernández-Gutiérrez, Shashwat Goel, Ameya Prabhu, and Matthias Bethge. Intrinsic credit assignment for long horizon interaction. *arXiv preprint arXiv:2602.12342*, 2026.

Yupeng Chang, Xu Wang, Jindong Wang, Yuan Wu, Linyi Yang, Kaijie Zhu, Hao Chen, Xiaoyuan Yi, Cunxiang Wang, Yidong Wang, et al. A survey on evaluation of large language models. *ACM transactions on intelligent systems and technology*, 15(3):1–45, 2024.

Tianyu Ding, Jianhong Xin, and Juan Pablo De la Cruz Weinstein. Keep policy gradient in charge: Sibling-guided credit distillation for long-horizon tool-use agents. *arXiv preprint arXiv:2606.12634*, 2026.

Jiazhan Feng, Shijue Huang, Xingwei Qu, Ge Zhang, Yujia Qin, Baoquan Zhong, Chengquan Jiang, Jinxin Chi, and Wanjun Zhong. Retool: Reinforcement learning for strategic tool use in llms. In *International Conference on Learning Representations*, volume 2026, pp. 37909–37926, 2026.

Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Peiyi Wang, Qihao Zhu, Runxin Xu, Ruoyu Zhang, Shirong Ma, Xiao Bi, et al. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.

Yucan Guo, Xiaohan Wang, Miao Su, Saiping Guan, Zhongni Hou, Jiajun Chai, Wei Lin, Guojun Yin, Xiaolong Jin, Jiafeng Guo, et al. Hidifftir: Hierarchical difficulty-aware policy optimization for multi-turn tool-integrated reasoning. *arXiv preprint arXiv:2608.21863*, 2026.

Dongfu Jiang, Yi Lu, Zhuofeng Li, Zhiheng Lyu, Ping Nie, Haozhe Wang, Alex Su, Hui Chen, Kai Zou, Chao Du, et al. Verltool: Towards holistic agentic reinforcement learning with tool use. *arXiv preprint arXiv:2509.01055*, 2025.

Bowen Jin, Hansi Zeng, Zhenrui Yue, Jinsung Yoon, Sercan Arik, Dong Wang, Hamed Zamani, and Jiawei Han. Search-r1: Training llms to reason and leverage search engines with reinforcement learning. *arXiv preprint arXiv:2503.09516*, 2025.

Yaxuan Li, Yuxin Zuo, Bingxiang He, Jinqian Zhang, Chaojun Xiao, Cheng Qian, Tianyu Yu, Huan-ang Gao, Wenkai Yang, Zhiyuan Liu, et al. Rethinking on-policy distillation of large language models: Phenomenology, mechanism, and recipe. *arXiv preprint arXiv:2604.13016*, 2026a.

Yu Li, Guangfeng Cai, Shengtian Yang, Han Luo, Shuo Han, Xu He, Dong Li, and Lei Feng. Phgpo: Pheromone-guided policy optimization for long-horizon tool planning. *arXiv preprint arXiv:2602.13691*, 2026b.

Jiayu Liu, Qihan Lin, Cheng Qian, Rui Wang, Emre Can Acikgoz, Xiaocheng Yang, Jiateng Liu, Zhenhailong Wang, Xiusi Chen, Heng Ji, et al. Planbench-xl: Evaluating long-horizon planning of llm tool-use agents in large-scale tool ecosystems. *arXiv preprint arXiv:2606.22388*, 2026.

Zhengxi Lu, Zhiyuan Yao, Zhuowen Han, Zi-Han Wang, Jinyang Wu, Qi Gu, Xunliang Cai, Weiming Lu, Jun Xiao, Yueting Zhuang, et al. Self-distilled agentic reinforcement learning. *arXiv preprint arXiv:2605.15155*, 2026.

Shervin Minaee, Tomas Mikolov, Narjes Nikzad, Meysam Chenaghlu, Richard Socher, Xavier Amatriain, and Jianfeng Gao. Large language models: A survey. *arXiv preprint arXiv:2402.06196*, 2024.

Humza Naveed, Asad Ullah Khan, Shi Qiu, Muhammad Saqib, Saeed Anwar, Muhammad Usman, Naveed Akhtar, Nick Barnes, and Ajmal Mian. A comprehensive overview of large language models. *ACM Transactions on Intelligent Systems and Technology*, 16(5):1–72, 2025.

Shishir G Patil, Huanzhi Mao, Fanjia Yan, Charlie Cheng-Jie Ji, Vishnu Suresh, Ion Stoica, and Joseph E Gonzalez. The berkeley function calling leaderboard (bfcl): From tool use to agentic evaluation of large language models. In *Forty-second International Conference on Machine Learning*, 2025.

Cheng Qian, Emre Can Acikgoz, Qi He, Hongru Wang, Xiusi Chen, Dilek Hakkani-Tur, Gokhan Tur, and Heng Ji. Toolrl: Reward is all tool learning needs. *Advances in Neural Information Processing Systems*, 38: 105523–105553, 2026.

Yujia Qin, Shengding Hu, Yankai Lin, Weize Chen, Ning Ding, Ganqu Cui, Zheni Zeng, Xuanhe Zhou, Yufei Huang, Chaojun Xiao, et al. Tool learning with foundation models. *ACM Computing Surveys*, 57 (4):1–40, 2024.

Changle Qu, Sunhao Dai, Xiaochi Wei, Hengyi Cai, Shuaiqiang Wang, Dawei Yin, Jun Xu, and Ji-Rong Wen. Tool learning with large language models: A survey. *Frontiers of Computer Science*, 19(8):198343, 2025.

Changle Qu, Sunhao Dai, Hengyi Cai, Jun Xu, Shuaiqiang Wang, and Dawei Yin. Matchtir: Fine-grained supervision for tool-integrated reasoning via bipartite matching. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 11953–11968, 2026a.

Changle Qu, Sunhao Dai, Hengyi Cai, Yuqi Zhou, Xinran Chen, Jun Xu, et al. Turnsight: Turn-level hindsight self-distillation for tool-integrated reasoning. *arXiv preprint arXiv:2608.04007*, 2026b.

Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté, Yonatan Bisk, Adam Trischler, and Matthew Hausknecht. Alfworld: Aligning text and embodied environments for interactive learning. *arXiv preprint arXiv:2010.03768*, 2020.

Mingyang Song and Mao Zheng. A survey of on-policy distillation for large language models. *arXiv preprint arXiv:2604.00626*, 2026.

Hui-Ze Tan, Xiao-Wen Yang, Hao Chen, Jie-Jing Shao, Yi Wen, Yuteng Shen, Weihong Luo, Xiku Du, Lan-Zhe Guo, and Yu-Feng Li. Hindsight credit assignment for long-horizon llm agents. *arXiv preprint arXiv:2603.08754*, 2026a.

Qitai Tan, Zefang Zong, Yang Li, and Peng Chen. Atod: Annealed turn-aware on-policy distillation for multi-turn autonomous agents. *arXiv preprint arXiv:2606.27814*, 2026b.

Arun James Thirunavukarasu, Darren Shu Jeng Ting, Kabilan Elangovan, Laura Gutierrez, Ting Fang Tan, and Daniel Shu Wei Ting. Large language models in medicine. *Nature medicine*, 29(8):1930–1940, 2023.

Yang Tian, Rui Wang, Xumeng Wen, Junjie Li, Shizhao Sun, Lei Song, Jiang Bian, and Bo Zhao. Pbsd: Privileged bayesian self-distillation for long-horizon credit assignment. *arXiv preprint arXiv:2606.09348*, 2026.

Songjun Tu, Chengdong Xu, Qichao Zhang, Yiwen Ma, Yaocheng Zhang, Linjing Li, Dong Li, Xiangyuan Lan, and Dongbin Zhao. Ucob: Learning to utilize and evolve agentic skills via credit-aware on-policy bidirectional self-distillation. *arXiv preprint arXiv:2606.29502*, 2026.

Hao Wang, Guozhi Wang, Han Xiao, Yufeng Zhou, Yue Pan, Jichao Wang, Ke Xu, Yafei Wen, Xiaohu Ruan, Xiaoxin Chen, et al. Skill-sd: Skill-conditioned self-distillation for multi-turn llm agents. *arXiv preprint arXiv:2604.10674*, 2026a.

Jiaqi Wang, Wenhao Zhang, Weijie Shi, Yaliang Li, and James Cheng. Tcod: Exploring temporal curriculum in on-policy distillation for multi-turn autonomous agents. *arXiv preprint arXiv:2604.24005*, 2026b.

Zi-Han Wang, Zhengxi Lu, Zhiyuan Yao, Jinyang Wu, Jie Wu, Zhengzhou Cai, Yueqing Sun, Ziang Ye, Linji Hao, Qi Gu, et al. Agentopsd: Recursive self-distillation for agentic reinforcement learning. *arXiv preprint arXiv:2608.05987*, 2026c.

Jinyang Wu, Shuo Yang, Zhengxi Lu, Fan Zhang, Yuhao Shen, Lang Feng, Haoran Luo, Zheng Lian, Shuai Zhang, Zhengqi Wen, et al. Seed: Self-evolving on-policy distillation for agentic reinforcement learning. *arXiv preprint arXiv:2607.14777*, 2026.

Weikai Xu, Chengrui Huang, Shen Gao, and Shuo Shang. Llm-based agents for tool learning: A survey: W. xu et al. *Data Science and Engineering*, 10(4):533–563, 2025.

Wenkai Yang, Weijie Liu, Ruobing Xie, Kai Yang, Saiyong Yang, and Yankai Lin. Learning beyond teacher: Generalized on-policy distillation with reward extrapolation. *arXiv preprint arXiv:2602.12125*, 2026a.

Zhicheng Yang, Zhijiang Guo, Yifan Song, Minrui Xu, Yongxin Wang, Yiwei Wang, Xiaodan Liang, and Jing Tang. Prune-opd: Efficient and reliable on-policy distillation for long-horizon reasoning. *arXiv preprint arXiv:2605.07804*, 2026b.

Shunyu Yao, Howard Chen, John Yang, and Karthik Narasimhan. Webshop: Towards scalable real-world web interaction with grounded language agents. *Advances in Neural Information Processing Systems*, 35: 20744–20757, 2022.

Junjie Ye, Zhengyin Du, Xuesong Yao, Weijian Lin, Yufei Xu, Zehui Chen, Zaiyuan Wang, Sining Zhu, Zhiheng Xi, Siyu Yuan, et al. Toolhop: A query-driven benchmark for evaluating large language models in multi-hop tool use. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 2995–3021, 2025.

Junjie Ye, Changhao Jiang, Zhengyin Du, Yufei Xu, Xuesong Yao, Zhiheng Xi, Xiaoran Fan, Qi Zhang, Tao Gui, Xuan-Jing Huang, et al. Feedback-driven tool-use improvements in large language models via automated build environments. In *Findings of the Association for Computational Linguistics: ACL 2026*, pp. 2293–2323, 2026.

Woongyeng Yeo, Yumin Choi, Taekyung Ki, and Sung Ju Hwang. Hint-sd: Targeted hindsight self-distillation for long-horizon agents. *arXiv preprint arXiv:2605.17873*, 2026.

Siyan Zhao, Zhihui Xie, Mengchen Liu, Jing Huang, Guan Pang, Feiyu Chen, and Aditya Grover. Self-distilled reasoner: On-policy self-distillation for large language models. *arXiv preprint arXiv:2601.18734*, 2026a.

Wayne Xin Zhao, Kun Zhou, Junyi Li, Tianyi Tang, Zican Dong, Yupeng Hou, Beichen Zhang, Yingqian Min, Junjie Zhang, Peiyu Liu, et al. A survey of large language models. *Frontiers of Computer Science*, 20 (12):2012627, 2026b.

Qiyong Zhong, Mao Zheng, Mingyang Song, Xin Lin, Jie Sun, Houcheng Jiang, Xiang Wang, and Junfeng Fang. Sod: Step-wise on-policy distillation for small language model agents. *arXiv preprint arXiv:2605.07725*, 2026.

Yuhang Zhou, Kai Zheng, Haoling Li, Dengyun Peng, Can Xu, and Jingjing Chen. Turnopd: Making on-policy distillation turn-aware for efficient long-horizon agent training. *arXiv preprint arXiv:2607.05804*, 2026.

Appendix

Contents

Appendix 13

A Benchmarks 14

B Baselines 14

B.1 Outcome RL and General (Self-) Distillation 14

B.2 Skill-Conditioned Self-Distillation 15

C Implementation Details 15

D Computational Overhead Analysis 16

E Algorithm 17

F Trajectory Examples and Skill Generation 17

G Prompts Used for Training 18

G.1 Environment Interaction Prompts 19

G.2 Outcome-Conditioned Hindsight Skill Induction 20

G.3 Skill-Conditioned Teacher Prompt 21

H Empirical Study: Detailed Setup 22

I Evaluation metrics 22

J Evaluation metrics 22

A Benchmarks

- **ALFWorld** is an interactive benchmark that aligns abstract textual reasoning with concrete embodied environments, requiring agents to translate high-level text policies into multi-step navigation and object manipulation actions in simulated household settings, thereby evaluating cross-modal grounding and long-horizon planning.
- **SearchQA** is a large-scale open-domain question answering dataset comprising over 140,000 question-answer pairs augmented with real-world search engine contexts, designed to assess model robustness in precise machine reading comprehension and information extraction amidst highly redundant and noisy retrieved snippets.
- **WebShop** provides a simulated e-commerce environment grounded in real-world web interactions, featuring over 1.18 million real products and 12,000+ crowdsourced natural language instructions. It requires agents to perform sequential decision-making to search, filter, and purchase target items, systematically evaluating compositional instruction understanding, dynamic query reformulation, and long-horizon multi-step interaction.
- **FTRL** is a foundational evaluation dataset focused on tool-use capabilities, comprising over 2,000 automatically constructed tool interaction scenarios with verifiable execution environments. It typically serves as a basis for in-domain training and evaluation, aiming to systematically measure a model's tool-calling accuracy and sequential decision-making efficacy within standardized environments.
- **BFCL (Berkeley Function Calling Leaderboard)**, proposed by Patil et al., is a comprehensive function-calling evaluation benchmark. Its core dataset encompasses thousands of question-function-answer pairs spanning multiple programming languages and interfaces (e.g., Python, Java, JavaScript, and REST APIs). Through iterative updates, BFCL has expanded to include complex test cases involving multi-turn dialogues, multi-step dependencies, parallel function calls, and real-time dynamic scenarios, primarily assessing a model's instruction-following capability and dynamic decision-making robustness in real-world API interactions.
- **ToolHop** is a query-driven benchmark specifically designed to evaluate the multi-hop tool use capabilities of large language models. The dataset consists of 995 complex user queries and 3,912 associated locally executable tools. By constructing toolchains with strong interdependencies, it rigorously examines a model's comprehensive proficiency in long-horizon logical reasoning, dynamic query reformulation, and multi-step tool orchestration.

B Baselines

B.1 Outcome RL and General (Self-) Distillation

- **GRPO** (Guo et al., 2025) is a reinforcement learning method that eliminates the need for an explicit value model. It samples multiple responses for the same input and estimates advantages based on their relative rewards within the group. Its supervision mainly comes from the final task outcome, such that all decisions along a trajectory typically share the same outcome-level reward.
- **ToolRL** (Qian et al., 2026) is a reinforcement learning framework designed for tool-use tasks. It trains tool-using agents with reward signals associated with task outcomes and tool execution, aiming to improve tool selection and invocation capabilities in complex and unseen tool-use scenarios.
- **MatchTIR** (Qu et al., 2026a) addresses the coarse-grained credit assignment problem of trajectory-level rewards in multi-turn tool-integrated reasoning. It matches generated tool-use trajectories with reference trajectories via bipartite matching to construct turn-level rewards, and combines trajectory-level and turn-level advantages for more fine-grained policy optimization.
- **OPD** lets the student model generate trajectories under its current policy and queries a teacher model to provide token-level distribution supervision on states visited by the student. Compared with offline distillation, OPD reduces the train–test distribution mismatch by directly training on the student's on-policy state distribution.
- **TCOD** (Wang et al., 2026b) focuses on trajectory-level KL instability caused by error accumulation in multi-turn on-policy distillation. It adopts a temporal curriculum that gradually increases the interaction horizon from short trajectories to full-length trajectories, enabling the student to progressively acquire long-horizon interaction capabilities while improving training stability.
- **OPSD** (Zhao et al., 2026a) is an on-policy self-distillation approach that removes the dependence on a separate large teacher model. The same model acts as both student and teacher under different contextual conditions: the student generates on-policy trajectories, while a privileged teacher view with additional task information provides token-level supervision.

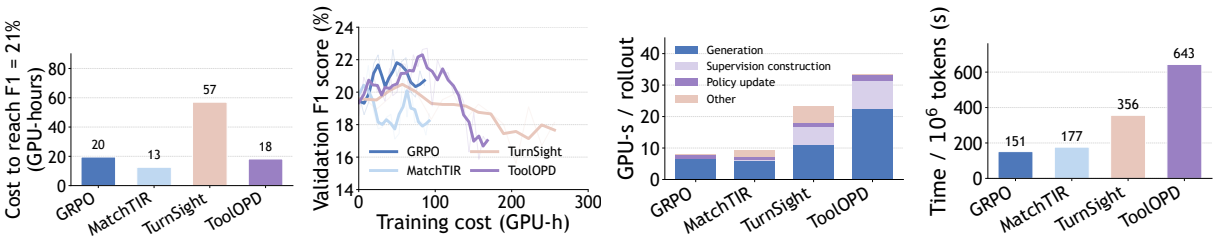

(a) Cost to reach F1 21%. (b) F1 versus GPU hours. (c) Rollout cost breakdown. (d) Time per million tokens.

Figure 7: Computational overhead analysis on FTRL.

- **ATOD** Tan et al. (2026b) is a turn-aware on-policy distillation method for long-horizon multi-turn agents. It combines OPD with GRPO through an annealed training scheme and dynamically adjusts the distillation strength at each turn according to teacher–student disagreement and uncertainty, thereby allocating stronger supervision to interaction steps that remain difficult for the student.
- **TurnSight** (Qu et al., 2026b) introduces turn-level hindsight self-distillation for tool-integrated reasoning. It constructs a hindsight teacher using tool execution outcomes and measures teacher–student distribution discrepancies at individual turns. Multi-step lookahead and group-wise normalization are further used to assign differentiated supervision weights across tool-interaction turns.

B.2 Skill-Conditioned Self-Distillation

- **Skill-SD** (Wang et al., 2026a) is a skill-conditioned self-distillation framework for multi-turn LLM agents. It summarizes completed agent trajectories into compact natural-language skills describing successful behaviors, common mistakes, and reusable workflows, and uses these skills as dynamic privileged information for the teacher branch. The student is then trained to internalize such skill guidance through on-policy distillation, together with an importance-weighted reverse-KL objective and dynamic teacher synchronization for stable training.
- **SDAR** (Lu et al., 2026) integrates on-policy self-distillation into agentic reinforcement learning as a gated auxiliary objective, while retaining RL as the primary optimization signal. It derives token-level teacher–student discrepancy signals and maps them into adaptive gates, strengthening distillation on teacher-endorsed tokens while attenuating unreliable negative guidance, thereby improving training stability in long-horizon multi-turn tasks.
- **AgentOPSD** Wang et al. (2026c) introduces recursive self-distillation for fine-grained credit assignment in long-horizon agentic reinforcement learning. It aggregates token-level teacher–student log-probability gaps into turn-level evidence and recursively updates a Bayesian belief state across interaction turns, using the resulting belief revisions to estimate the importance of individual decisions and reweight policy optimization without introducing an additional critic or extra rollouts.

C Implementation Details

For the long-horizon tasks reported in Table 1, we use GRPO as the base policy optimization algorithm and jointly optimize it with the ToolOPD auxiliary objective. ToolOPD operates exclusively on tool-action tokens. For each completed trajectory, the teacher generates an outcome-conditioned hindsight skill, whose effect on the teacher's action probabilities is converted into fine-grained credit for reweighting the distillation loss. We set the hindsight scaling factor to $\beta_h = 5.0$ and the replay calibration factor to $\alpha = 1.0$, while linearly increasing the evidence weight λ_e from 0.2 to 0.5 during training. The ToolOPD loss coefficient λ_{OPD} is selected from 0.001 to 0.01 according to the task and student model size, and the log-probability ratio is clipped at 5.0 for training stability.

We use Qwen3-0.6B or Qwen3-1.7B as the student policy and a GRPO-trained Qwen3-4B as the frozen teacher. The same teacher performs both hindsight skill induction and action-level supervision, avoiding the need for an additional skill-generation model or external API. During training, four GPUs are allocated to the student policy and four to the independent teacher pool. The environment-specific training configurations are summarized in Table 4.

For tool-use benchmark, we augment vanilla GRPO with a ToolOPD auxiliary loss applied only to tool-action tokens. We use Qwen3-0.6B as the student and a GRPO-trained Qwen3-4B as the frozen teacher. We set $\beta_h = 10.0$, $\alpha = 1.0$, $\lambda_{\text{OPD}} = 0.03$, and $\lambda_e = 0$, while clipping the teacher–student log-ratio at 3.0. Training uses eight rollouts per prompt, a batch size of 32, and a learning rate of 2×10^{-7}.

Table 4: Core ToolOPD training configurations across the three tool-use benchmarks.

Hyperparameter	ALFWorld	WebShop	SearchQA
Training steps	150	150	150
Training batch size	16	16	128
Rollout group size G	8	8	8
Maximum prompt length	2048	4096	4096
Maximum response length	512	512	512
Maximum interaction turns	50	15	4
Train/validation temperature	1.0 / 0.4	1.0 / 0.4	1.0 / 0.4
GPUs (student / teacher)	4 / 4	4 / 4	4 / 4
Rollout tensor-parallel size	2	2	2

Table 5: Core training hyperparameters for ToolOPD on the tool-use benchmarks.

Hyperparameter	Value
Student model	Qwen3-0.6B / Qwen3-1.7B
Teacher model	GRPO-trained Qwen3-4B
Policy optimization algorithm	Vanilla GRPO
Group size G (rollouts per task)	8
Training batch size	32
PPO mini-batch size	32
Optimization epochs per batch	1
Learning rate	2×10^{-7}
Training temperature	1.0
Maximum prompt length	10,000 tokens
Maximum response length	13,000 tokens
Maximum tokens per GPU	23,000
Tensor parallel size	1
KL-loss coefficient	0.0
Entropy coefficient	0.0
Hindsight scaling factor β_h	10.0
Replay calibration factor α	1.0
Evidence weight λ_e	0.0
ToolOPD loss coefficient λ_{OPD}	0.03
Maximum log-ratio threshold	3.0
Hindsight generation length	96 tokens
Teacher maximum prompt length	14,000 tokens
Hindsight maximum prompt length	14,000 tokens
ToolOPD supervision scope	Tool-action tokens only

D Computational Overhead Analysis

To quantify the additional computational cost of ToolOPD, we use a matched eight-GPU setup on FTRL and exclude evaluation and checkpointing time from all measurements. As shown in Fig. 7(d), ToolOPD does not improve raw training throughput: it requires 643 seconds per million training tokens, compared with 151, 177, and 356 seconds for GRPO, MatchTIR, and TurnSight, respectively. The per-rollout cost breakdown in Fig. 7(c) further shows that ToolOPD is dominated by rollout generation (22.5 GPU-s) and outcome-aware supervision construction (8.8 GPU-s), whereas policy updates account for only a small fraction of the total cost (2.1 GPU-s). These results indicate that the additional overhead primarily arises from teacher-driven hindsight skill inference and action-level supervision construction, rather than standard policy optimization.

Despite this overhead, the performance–cost curve in Fig. 7(b) reveals a favorable optimization trade-off in some regimes. In particular, Fig. 7(a) shows that ToolOPD reaches an F1 score of 21% after 18.2 GPU-hours, requiring less computation than GRPO (19.5 GPU-hours) and substantially less than TurnSight (57.0 GPU-hours), although it remains more expensive than MatchTIR (12.5 GPU-hours). Therefore, ToolOPD should not be viewed as a method for improving raw computational efficiency; instead, it trades additional computation for more targeted credit assignment, allowing it to reach a comparable performance level with lower cost than TurnSight.

E Algorithm

Algorithm 1 Outcome-Aware On-Policy Distillation for ToolOPD

Input: Student policy π_θ, old policy $\pi_{\theta_{\text{old}}}$, frozen teacher π_ϕ, group size G, sensitivity β_h, distillation coefficient λ_{OPD}, log-ratio bound c
Output: Optimized student policy π_θ

for each training step **do**
  Sample G trajectories $\{\tau_i, R_i, y_i\}_{i=1}^{G} \sim \pi_{\theta_{\text{old}}}$ for each task
  GRPO advantage: normalize rewards within each trajectory group,
    $A_i \leftarrow \dfrac{R_i - \frac{1}{G}\sum_{j=1}^{G} R_j}{\text{std}(\{R_j\}_{j=1}^{G}) + \epsilon}$
  Hindsight skill induction: for each trajectory τ_i,
    $z_i \leftarrow G_\phi(\tau_i, y_i)$
  Outcome-asymmetric credit assignment: for each decision turn (i, t),
    Compute teacher probabilities $q_\phi^0(\cdot \mid s_{i,t})$ and $q_\phi^{z_i}(\cdot \mid s_{i,t})$
    $\Delta_{i,t,k} \leftarrow \log q_\phi^{z_i}(a_{i,t,k} \mid s_{i,t}) - \log q_\phi^0(a_{i,t,k} \mid s_{i,t})$
    $c_{i,t,k} \leftarrow m_{i,t,k}[2\sigma(\beta_h(2y_i - 1)\Delta_{i,t,k}) - 1]_+$
    $H_{i,t} \leftarrow \dfrac{\sum_k c_{i,t,k}}{\max(1, \sum_k m_{i,t,k})}$
  Action-level on-policy distillation:
    $\ell_{i,t,k} \leftarrow \log q_\phi^{z_i}(a_{i,t,k} \mid s_{i,t}) - \log \pi_\theta(a_{i,t,k} \mid s_{i,t})$
    $r_{i,t,k} \leftarrow \exp(\text{clip}(\ell_{i,t,k}, -c, c))$
    $\widehat{\mathcal{L}}_{\text{ToolOPD}} \leftarrow \dfrac{1}{N_A}\sum_{i,t,k} m_{i,t,k} H_{i,t}\,(r_{i,t,k}\log r_{i,t,k} - r_{i,t,k} + 1)$
  GRPO policy loss:
    $\mathcal{L}_{\text{GRPO}} \leftarrow -\mathbb{E}_{i,t,k}[\min(\rho_{i,t,k}A_i,\ \text{clip}(\rho_{i,t,k}, 1-\epsilon_l, 1+\epsilon_h)A_i)]$
    where $\rho_{i,t,k} = \pi_\theta(a_{i,t,k} \mid s_{i,t}) / \pi_{\theta_{\text{old}}}(a_{i,t,k} \mid s_{i,t})$
  Joint policy update:
    $\mathcal{L}_{\text{total}} \leftarrow \mathcal{L}_{\text{GRPO}} + \lambda_{\text{OPD}}\widehat{\mathcal{L}}_{\text{ToolOPD}}$
    $\theta \leftarrow \theta - \eta\nabla_\theta\mathcal{L}_{\text{total}}$
end for
return π_θ

F Trajectory Examples and Skill Generation

We present representative trajectories from ALFWorld and FTRL to illustrate how ToolOPD induces reusable hindsight skills from completed interactions. We retain all outcome-critical tool actions while omitting repetitive reasoning and redundant observations using [...].

Case 1: Long-Horizon Interaction in ALFWorld

Task. Put a clean cup in the microwave.

Trajectory.
Agent – Step 1: `go to microwave 1`
Environment – Step 1: The agent arrives at `microwave 1`, which is closed.
Agent – Step 2: `open microwave 1`
Environment – Step 2: The microwave is open and contains `cup 1`.
Agent and Environment – Steps 3–11: The agent repeatedly examines the microwave and checks its inventory, while the cup remains inside the microwave. [...]
Agent – Step 12: **`take cup 1 from microwave 1`**
Environment – Step 12: The agent successfully picks up `cup 1`.
Agent – Step 13: `examine cup 1`
Environment – Step 13: The cup is empty, but its required clean state has not yet been established.
Agent and Environment – Steps 14–18: The agent performs several redundant examinations and location transitions that do not change the task state. [...]

Agent – Steps 19–20:
`go to sinkbasin 1`
`clean cup 1 with sinkbasin 1`
Environment – Step 20: The cup is successfully cleaned.

Agent – Steps 21–24:
`go to microwave 1`
`examine cup 1`
`examine microwave 1`
`move cup 1 to microwave 1`
Environment – Step 24: The clean cup is placed in the microwave, completing the task.

Outcome. Success.

Generated Hindsight Skill. Verify and clean objects before placing them in appliances.

Although the trajectory contains numerous examinations and unnecessary transitions, its success is primarily determined by the pivotal action sequence **`take`** → **`clean`** → **`place`**. ToolOPD therefore extracts an outcome-grounded skill from the noisy trajectory and concentrates supervision on actions that causally change the task state.

Case 2: Multi-Step Tool Reasoning on FTRL

Task. What is the result of subtracting 8 from the reverse of the birth year of the star of *Lady in a Cage*?

Trajectory.
Agent – Step 1: The agent identifies the main star of the target film:
```
movie_star_finder(
  movie_title="Lady in a Cage",
  role_type="main_star")
```
Environment – Step 1: The tool returns **`Olivia de Havilland`**.

Agent – Step 2: The agent retrieves the birth year of the identified person:
```
biographical_info_retriever(
  name="Olivia de Havilland",
  information_type="birth_year")
```
Environment – Step 2: The tool returns **`1916`**.

Agent – Step 3: The agent reverses **`1916`** to obtain **`6191`** and performs the required subtraction:
```
advanced_arithmetic_calculator(
  operation="subtract",
  operands=[6191, 8])
```
Environment – Step 3: The tool returns **`6183`**.

Agent – Final Answer: **`<answer>6183</answer>`**

Outcome. Correct answer.

Generated Hindsight Skill. Resolve entities before computation: retrieve each dependent entity and numerical attribute with tools, and then apply the requested transformations in their specified order.

The FTRL trajectory forms a clear dependency chain: identifying the target person and retrieving the corresponding birth year are necessary prerequisites for the final numerical operation. Conditioned on the successful outcome, ToolOPD preserves supervision along this reasoning chain and emphasizes the tool calls that establish essential intermediate variables. Together, the two cases show that ToolOPD induces transferable skills and assigns credit to outcome-relevant decisions across embodied interaction and multi-step tool reasoning.

G Prompts Used for Training

We report the prompts used to collect on-policy trajectories and construct ToolOPD supervision. For environment interaction, each prompt specifies the task, recent interaction history, current observation, admissible actions or tools, and the required response format. For ToolOPD, the completed trajectory and its outcome are additionally provided to the teacher for hindsight skill induction. The generated

skill is subsequently inserted into the teacher context when computing action-level distillation targets. Placeholders enclosed in braces are instantiated online.

G.1 Environment Interaction Prompts

G.1.1 ALFWorld

ALFWorld Interaction Prompt

```
You are an expert agent operating in the ALFRED Embodied Environment.

Your task is to:
{TASK_DESCRIPTION}

Prior to this step, you have already taken {STEP_COUNT} step(s).
Below are the most recent {HISTORY_LENGTH} observations and
the corresponding actions you took:
{ACTION_HISTORY}

You are now at step {CURRENT_STEP}.
Your current observation is:
{CURRENT_OBSERVATION}

Your admissible actions in the current situation are:
[{ADMISSIBLE_ACTIONS}]

Now it is your turn to take one action. First reason step by step
about the current situation. The reasoning must be enclosed within
<think> and </think> tags and contain at most two sentences.

Then select exactly one admissible action and present it within
<action> and </action> tags. Output one action only and then stop.
```

For the first interaction step, the history block is omitted because no previous observation–action pair is available.

G.1.2 WebShop

WebShop Interaction Prompt

```
You are an expert autonomous agent operating in the WebShop
e-commerce environment.

Your task is to:
{TASK_DESCRIPTION}

Prior to this step, you have already taken {STEP_COUNT} step(s).
Below are the most recent {HISTORY_LENGTH} observations and
the corresponding actions you took:
{ACTION_HISTORY}

You are now at step {CURRENT_STEP}.
Your current observation is:
{CURRENT_OBSERVATION}

Your admissible actions in the current situation are:
[
{AVAILABLE_ACTIONS}
]

Now it is your turn to take one action. First reason step by step
about the current situation and determine which admissible action
best advances the shopping goal. The reasoning must be enclosed
within <think> and </think> tags.

Then select one admissible action and present it within
```

```
<action> and </action> tags.
```

As in ALFWorld, the interaction-history block is omitted at the first step.

G.1.3 FTRL

For FTRL, the available tool schemas vary across instances and are dynamically inserted into the system prompt.

FTRL System Prompt

```
# Tools

You may call one or more functions to assist with the user query.

You are provided with function signatures within <tools></tools>
XML tags:

<tools>
{TOOL_SCHEMAS}
</tools>

For each function call, return a JSON object containing the
function name and arguments within <tool_call></tool_call> tags:

<tool_call>
{"name": <FUNCTION_NAME>, "arguments": <ARGUMENTS_JSON>}
</tool_call>
```

FTRL User Prompt

```
Please call the given tools to answer the question. All information
must be obtained by calling tools rather than answering the question
directly.

If a tool call fails, correct the call and continue until you arrive
at an answer. Output only the final answer, as a word, number, or
short phrase, inside the <answer></answer> tag, without explanations
or additional information.

Question:
{QUESTION}
```

G.2 Outcome-Conditioned Hindsight Skill Induction

ToolOPD applies an independent teacher to the complete trajectory after its outcome becomes available. The prompt explicitly distinguishes successful and failed trajectories so that the generated skill captures either an effective workflow or an error-avoidance principle.

G.2.1 ALFWorld and WebShop

Hindsight Skill Induction Prompt

```
You are an independent RL teacher analyzing a completed
long-horizon tool-use trajectory.

Task description:
{TASK_DESCRIPTION}

Complete trajectory:
{COMPLETE_TRAJECTORY}
```

```
Final outcome:
{SUCCESS_OR_FAILURE}

If the trajectory succeeds:
Describe the reusable workflow that made the trajectory succeed.

If the trajectory fails:
Identify the task-specific failed decision pattern and state one
reusable workflow or avoidance principle. Ground every detail in
the task and trajectory. Explicitly identify the target object,
the relevant receptacle or tool, the missing prerequisite operation
or subgoal, and any wrong object selection or repeated no-progress
action loop. For tasks requiring cleaning, cooling, or heating,
distinguish actively performing the operation from merely finding
an object that already has the desired state. Do not invent
unobserved objects or states.

Return ONLY one valid JSON object:
{"skill": "one concise policy-facing skill"}

Do not output turn numbers, importance scores, action corrections,
numeric ratings, or Yes/No judgments.
```

G.2.2 FTRL

FTRL Hindsight Skill Induction Prompt

```
You are an expert tool-use teacher.

Trajectory:
{COMPLETE_TRAJECTORY}

Outcome:
{SUCCESS_OR_FAILURE}

Induce exactly one short, reusable tool-use strategy.

For a successful trajectory, capture the effective workflow or
pivotal decision. For a failed trajectory, capture the root error
and explain how to avoid it.

Use exactly the following format:
<skill>
Rule Name: imperative guidance. Apply when the relevant condition holds.
</skill>

Use at most 50 words. Do not include analysis, scores, or a
replacement trajectory.
```

G.3 Skill-Conditioned Teacher Prompt

After hindsight skill induction, ToolOPD prepends the generated skill to the original user context of the teacher. The remaining environment prompt, tool schemas, and interaction history are kept unchanged.

Skill-Conditioned Teacher Context

```
[Hindsight Skill]
{GENERATED_SKILL}

{ORIGINAL_USER_PROMPT}
```

The student policy receives the original environment prompt, whereas the teacher is conditioned on the generated hindsight skill. ToolOPD compares their action-token distributions and uses the resulting outcome-asymmetric skill relevance to dynamically weight the distillation objective.

H Empirical Study: Detailed Setup

To investigate the connection between performance degradation in long-horizon tool learning and outcome-dependent credit assignment, we design three complementary diagnostic experiments. These experiments examine long-context robustness, outcome sensitivity, and supervision conflicts (Figure 1(a)–(c)). The first evaluates the impact of interaction length on model performance, while the latter two analyze the behavior of ToolOPD credit weights on fixed trajectories. We uniformly use Qwen3-1.7B for the long-context experiments. For the credit analysis, we use Qwen3-4B models trained up to step 150 on each dataset as the teacher and probability evaluator.

Credit concentration. Fig. 1(a) examines how credit is distributed across interaction turns in long-horizon ALFWorld trajectories. For each trajectory, we rank all turns according to their estimated outcome-aware credit $H_{i,t}$ and compute the fraction of total credit assigned to the top 20% of turns:

$$\mathrm{CreditMass}_{20\%} = \frac{\sum_{t \in \mathcal{T}^i_{20\%}} H_{i,t}}{\sum_{t \in \mathcal{T}^i} H_{i,t}},$$

where $\mathcal{T}^i_{20\%}$ denotes the highest-credit 20% of turns in trajectory i. As shown in Fig. 1(a), these turns account for 74.2% of the total credit, whereas the remaining 80% account for only 25.8%. This highly concentrated distribution shows that different interaction turns contribute unequally to the final outcome. Consequently, uniform OPD spreads teacher supervision over many low-credit transitions and weakens the learning signal for pivotal tool decisions.

Outcome-sensitive credit assignment. Fig. 1 (b) employs a controlled outcome counterfactual experiment. We keep the task, state, and action sequence unchanged. We then flip the final outcome between success and failure, and regenerate the factual and counterfactual skills accordingly. Next, we compute the factual credit $H_{i,t}$ and the counterfactual credit $H^{\mathrm{cf}}_{i,t}$ on the same action tokens. The credit calculation follows the outcome-asymmetric definition of ToolOPD:

$$H_{i,t,k} = \left[2\sigma\left(\beta s_i \left(\log \pi_{\mathrm{tea}}(a_{i,t,k} \mid z_i) - \log \pi_{\mathrm{tea}}(a_{i,t,k})\right)\right) - 1\right]_+,$$

where $s_i \in \{1, -1\}$ denotes success and failure, respectively, and we set $\beta = 5$. If a turn satisfies $\left|H_{i,t} - H^{\mathrm{cf}}_{i,t}\right| > 0.01$, we label it as an outcome-sensitive turn. This experiment covers sampled data from ALFWorld (6 trajectories, 29 turns), WebShop (4 trajectories, 12 turns), and SearchQA (4 trajectories, 8 turns). It directly quantifies the dependence of credit assignment on the final outcome.

Conflicting supervision analysis. Fig. 1 (c) analyzes the consistency between the evidence provided by the teacher skill for observed actions and the final outcome. For an action token $a_{i,t,k}$, we define the signed teacher evidence as:

$$e_{i,t,k} = s_i \left[\log \pi_{\mathrm{tea}}(a_{i,t,k} \mid z_i) - \log \pi_{\mathrm{tea}}(a_{i,t,k})\right].$$

When $e_{i,t,k} > 0$, skill conditioning shifts the teacher distribution in a direction consistent with the outcome; otherwise, it conflicts. We compute the proportion of conflicting evidence in both uniform distillation and credit-weighted distillation:

$$R_{\mathrm{uniform}} = \frac{\sum_{i,t,k} [-e_{i,t,k}]_+}{\sum_{i,t,k} |e_{i,t,k}|}, \qquad R_H = \frac{\sum_{i,t,k} H_{i,t} [-e_{i,t,k}]_+}{\sum_{i,t,k} H_{i,t} |e_{i,t,k}|}.$$

Both settings share the same trajectories and teacher outputs. The only difference is whether ToolOPD credit weighting is applied. We conduct this analysis on ALFWorld diagnostic trajectories and report the conflict proportions for successful, failed, and all trajectories.

Additionally, Fig. 1 (d) shows the success rate of Qwen3-0.6B on ALFWorld over 150 continuous training steps. This verifies whether the differences in the aforementioned credit assignment mechanisms translate into actual optimization gains.

I Evaluation metrics

J Evaluation metrics

For ALFWorld, SearchQA, and WebShop in Table 1, we use task success rate as the primary metric, defined as the proportion of evaluation instances in which the agent successfully completes the environment objective or produces the correct final answer. WebShop additionally reports a normalized task score that

Table 6: Task and metric abbreviations used in Table 1.

ALFWorld		SearchQA		WebShop	
Abbr.	**Full name**	**Abbr.**	**Full name**	**Abbr.**	**Full name**
Pick	Pick and Place	NQ	Natural Questions	Score	Average Score
Look	Look at Object in Light	Triv	TriviaQA	Succ.	Success Rate
Clean	Pick, Clean, and Place	Pop	PopQA		
Heat	Pick, Heat, and Place	Hotp	HotpotQA		
Cool	Pick, Cool, and Place	2Wk	2WikiMultiHopQA		
Pick2	Pick Two Objects and Place	MuS	MuSiQue		
		Bam	Bamboogle		

Table 7: Metric abbreviations used in the FTRL, BFCL, and ToolHop evaluation.

FTRL		BFCL		ToolHop	
Abbr.	**Full name**	**Abbr.**	**Full name**	**Abbr.**	**Full name**
Solve-F1	Solve F1 Score	Base	Multi-Turn Base	Acc.	Accuracy
Solve-R	Solve Recall	MF	Multi-Turn Missing Function		
Solve-P	Solve Precision	MP	Multi-Turn Missing Parameter		
		LC	Multi-Turn Long Context		
		MTOA	Multi-Turn Overall Accuracy		
		Memory	Agent Memory Accuracy		

measures partial progress in unsuccessful trajectories. The task abbreviations are summarized in Table 6. The column "Avg." denotes the average performance across the corresponding task categories.

For Table 2, we follow the official evaluation protocol of each benchmark. FTRL reports Solve-P, Solve-R, and Solve-F1, which measure tool-use precision, task completion, and their harmonic mean, respectively. BFCL reports accuracy on multi-turn base, missing-function, missing-parameter, long-context, and memory tasks, with MTOA denoting Multi-Turn Overall Accuracy. ToolHop uses Accuracy (Acc.) to evaluate the correctness of the final answer after multi-hop tool reasoning. Unless otherwise specified, all metrics are reported as percentages, and higher values indicate better performance. The corresponding abbreviations are summarized in Table 7.