arXiv:2610.27516v1 [cs.CV] 20 Oct 2026

HealBench: A Fine-Grained Benchmark for Evaluating Psychological Comfort in Text-to-Image Generation

Junjie Xu[1], Yuzhe Yang[1], Rui Yang[2], Luwei Xiao[3], Wenxin Hu[1], Xingjiao Wu[1], Liang He[1]

[1]East China Normal University
[2]Johns Hopkins University
[3]National University of Singapore

Abstract

Current text-to-image and image-editing models can generate photorealistic images with strong semantic fidelity and high visual quality. Yet an important human-centered capability remains underexplored: whether these systems can create images that users perceive as psychologically comforting. Drawing on cognitive evaluation theory, we conceptualize psychological comfort not as a low-level visual property, but as an affective response shaped by a viewer's appraisal of restorative and emotionally supportive cues. To study this problem, we introduce HealBench, a large-scale, fine-grained benchmark for evaluating perceived psychological comfort in restorative visual stimuli. HealBench contains 30K high-quality image-text pairs spanning 19 high-level categories and 92 subcategories, constructed from 100K candidates through a human-and-expert-in-the-loop pipeline. To preserve the restorative nature of the data, the pipeline integrates VLM-based semantic matching, multiple human preference models, expert filtering, and augmentation guided by highly restorative exemplars. We further construct HealBench-Eval, a rigorously selected evaluation set of 1,095 instances using feature clustering. Using HealBench-Eval, we conduct extensive objective and subjective evaluations of 14 representative image generation models, together with prompt-rewriting generation strategies. Our results reveal a clear gap between conventional visual generation quality and perceived psychological comfort: models that perform well on semantic alignment and aesthetics still struggle to produce images experienced as emotionally restorative. Further reinforcement learning experiments show that HealScorer can serve as a practical reward signal for assessing and improving the emotional and restorative quality of generated images. These findings suggest that psychological comfort is a distinct and underexplored dimension of human-centered image generation, and establish HealBench as a foundation for future work on affect-aware generative models.

Keywords: Text-to-Image, Emotional Comfort, Benchmark and datasets, Image Manipulation, Multi-modal Learning

Introduction

Recently, text-to-image (T2I) generation and instruction-guided image editing have made rapid progress in semantic alignment (Fu et al. 2025b), visual quality, editing fidelity (Yu et al. 2025; Fu et al. 2025a), and human preference modeling (Ma et al. 2025). These advances have enabled generative models to produce images that are increasingly faithful, realistic, and visually appealing. However, as visual generation systems are increasingly used in emotionally sensitive and human-centered scenarios, a more subtle question emerges: can a generated image not only **look correct**, but also **feel right** to the viewer? This distinction matters when images are expected to support emotional experience, restoration, or psychological relief rather than merely satisfy semantic or aesthetic criteria. An image may match the prompt and appear visually polished, yet still fail to provide perceived psychological comfort (Belanche et al. 2025).

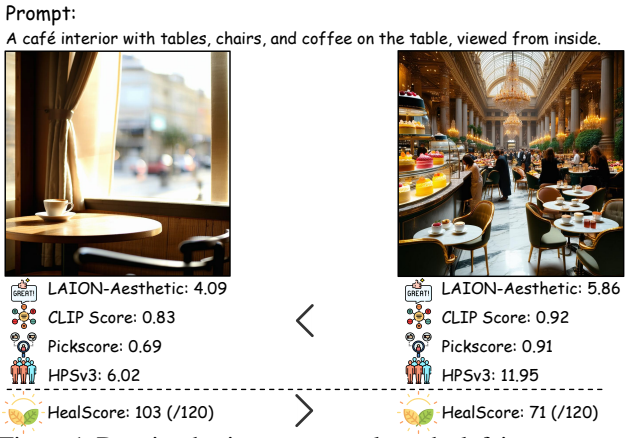

Figure 1: Despite sharing a prompt class, the left image scores poorly on general models (like CLIP) but high in psychological comfort (HealScore). Conversely, the right image has high evaluation scores but low comfort.

We argue that this limitation is not merely a weakness of current models, but also a consequence of how the target has been formulated and evaluated. Existing generation pipelines are typically optimized for text-image alignment, perceptual quality, editing fidelity, or generic human preference. These objectives are valuable, but they do not directly capture the affective and restorative experience required by psychological comfort. In particular, current systems suffer from three mismatches. First, there is a ***target mismatch***: *perceived psychological comfort* depends on affective and

Copyright © 2027, Association for the Advancement of Artificial Intelligence (www.aaai.org). All rights reserved.

Table 1: Comparison with representative benchmarks and evaluators for T2I generation and image editing. Existing resources mainly focus on preference, alignment, fidelity, compositionality, or generic quality assessment, whereas **HealBench** explicitly targets **perceived psychological comfort** with **fine-grained diagnostic dimensions** and **actionable rationale-bearing feedback**.

Benchmark / Evaluator	Venue	Scope	Primary Target	Fine-grained	Comfort-specific	Rationale / Feedback	Optimization
Pick-a-Pic	NeurIPS'23	T2I	Pairwise human preference	No ✗	No ✗	No ✗	Limited
AIGCIQA2023	LNCS'24	T2I	Quality, authenticity, and correspondence	Partial	No ✗	No ✗	No ✗
T2I-CompBench	NeurIPS'23	T2I	Compositional correctness	Yes ✓	No ✗	No ✗	No ✗
HATIE	CVPR'25	Editing	Human-aligned editing quality	Yes ✓	No ✗	No ✗	Limited
NTIRE 2025 T2I QA	CVPR'25	T2I	Fine-grained alignment and structural distortion	Yes ✓	No ✗	Partial	No ✗
ICE-Bench	ICCV'25	T2I + Editing	Unified evaluation for creation and editing	Yes ✓	No ✗	No ✗	No ✗
Task-Decomposed AutoEval	ACL'25	T2I evaluator	Automatic multi-criterion quality assessment	Yes ✓	No ✗	Partial	Limited
HealBench (ours)	**AAAI'27**	**T2I + Editing**	**Perceived psychological comfort**	**Yes ✓**	**Yes ✓**	**Yes ✓**	**Yes ✓**

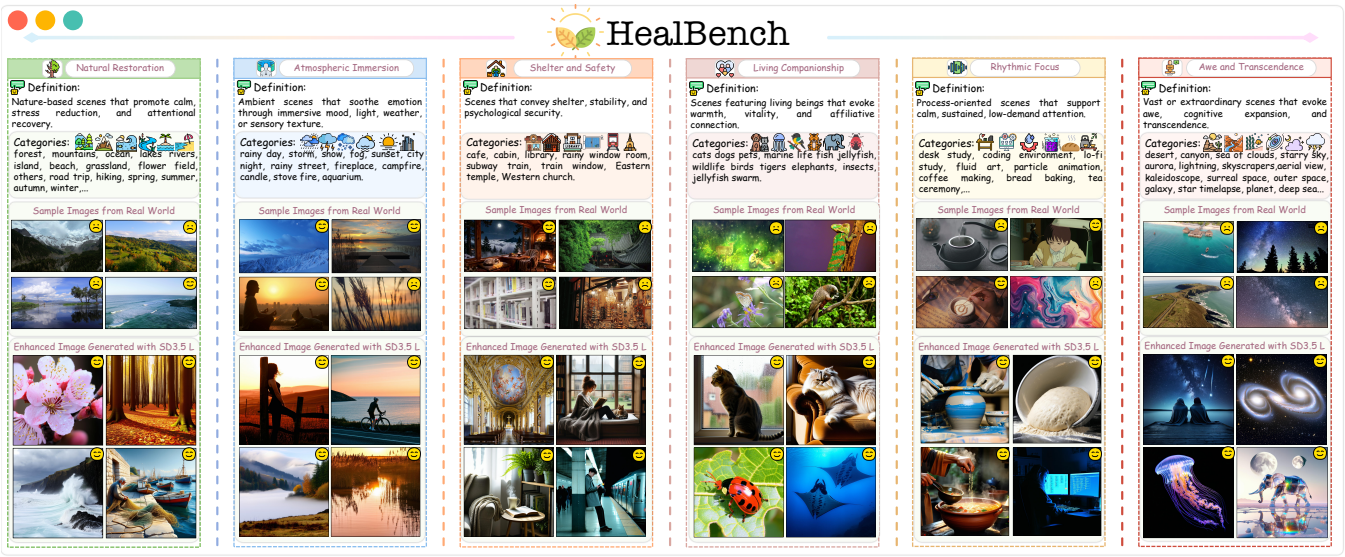

Figure 2: Overview of HealBench. HealBench organizes healing-related visual content into 19 categories and 92 subcategories, which are further grouped into six major classes according to different healing functions. For each category, we present both real-world examples and enhanced images generated via SD 3.5-L, which is also used for dataset augmentation.

restorative cues that are not explicitly captured by existing optimization targets (Shi et al. 2025; Belanche et al. 2025; Ortega, Murai, and Whitney 2025). Second, there is an ***evaluation-granularity mismatch***: comfort-related failure may arise from different sources, including prompt misalignment, incorrect healing semantics, or poor image quality (Gao, Ajith, and Peelen 2025; Ortega, Murai, and Whitney 2025). As illustrated in Figure 1, two images may receive similarly low comfort ratings for fundamentally different reasons, making a single scalar score insufficient for diagnostic evaluation. Third, there is a ***benchmark-and-feedback mismatch***: as summarized in Table 1, existing benchmarks and evaluators mainly target preference, alignment, fidelity, compositionality, or generic quality, but do not explicitly target *perceived psychological comfort* or provide actionable, rationale-bearing feedback. Against this background, we ask: how can perceived psychological comfort be turned into an evaluable and optimization-relevant objective for image generation and image editing? Rather than treating comfort as a vague holistic preference, we formulate healing-oriented generation through three necessary dimensions: **text-image alignment**, **heal accuracy**, and **image quality**. A generated image should be considered healing only when it satisfies all three dimensions; failure in any one dimension is sufficient to disqualify the output. Accordingly, our goal is not merely to rank images from better to worse, but to identify whether, where, and why a generated image fails to provide perceived psychological comfort.

Based on this formulation, we propose **HealBench** (shown in Figure 2), a fine-grained benchmark for evaluating the perceived psychological comfort of generated images. HealBench is constructed through a human-and-expert-in-the-loop pipeline that curates high-quality data from a large-scale candidate image-text pool, performs imitative data augmentation, and builds a rigorously filtered test set for unified comparison across models and prompting strategies. Unlike preference-style benchmarks that primarily produce aggregate scores, HealBench is designed around failure-based diagnosis: it identifies whether an output violates text-image alignment, heal accuracy, or image quality, thereby revealing different ways in which a generated image may fail to be healing. Evaluation alone, however, is insufficient for improving generation systems. To make diagnosis actionable, we intro-

duce a **score-plus-rationale** feedback paradigm, where each result is paired with both a comfort score and an explicit explanation of which dimension fails and why. This structured feedback supports prompt rewriting for both T2I generation and instruction-guided image editing. Since large-scale model iteration cannot rely entirely on costly human evaluation, we further train an automatic evaluator, Tri-dimensional Healing Scorer (**THS**), on top of HealBench to approximate these structured judgments and study its use as an optimization signal for downstream improvement, including RL-based fine-tuning.

Within this work, we treat psychological comfort as a distinct human-centered objective for visual generation, rather than as a by-product of semantic alignment, aesthetics, or generic preference. Our contributions are as follows: (1)We formulate *perceived psychological comfort* as a failure-based evaluation problem for healing-oriented visual generation, defined by three necessary dimensions: **text-image alignment**, **heal accuracy**, and **image quality**. (2)We build **HealBench**, a fine-grained benchmark that diagnoses whether a generated image fails on any of the three dimensions, instead of relying solely on aggregate preference-style scoring. (3)We introduce a **score-plus-rationale** feedback paradigm and **THS**, turning structured diagnosis into an actionable signal for prompt refinement and downstream model optimization.

Related Work

Image Generation and Editing Models. Recent visual generation models have made substantial progress in fidelity, semantic alignment, controllability, and human preference alignment. Diffusion models such as DDPM, latent diffusion, Imagen, and DiT enable high-resolution synthesis and strong prompt following (Ho, Jain, and Abbeel 2020; Rombach et al. 2022; Saharia et al. 2022; Peebles and Xie 2023); autoregressive and token-based models such as DALL·E, Make-A-Scene, and Muse improve complex content modeling and composition (Ramesh et al. 2021; Gafni et al. 2022; Chang et al. 2023); and flow-matching or rectified-flow methods further enhance training and sampling efficiency (Lipman et al. 2023; Liu et al. 2025; Esser et al. 2024). Preference optimization and reinforcement learning methods further align T2I models with human judgments (Black et al. 2024; Fan et al. 2023; Wallace et al. 2024), while text-guided editing methods provide instruction-based control over image content and style (Hertz et al. 2023; Mokady et al. 2023; Brooks, Holynski, and Efros 2023). However, these methods mainly optimize whether images are faithful, realistic, editable, or generally preferred. Existing affective methods, such as Affective Image Filter and EmoEdit, focus on explicit emotion cues or targeted emotional editing (Weng et al. 2023; Yang et al. 2025), leaving subtler outcomes such as psychological comfort, safety, companionship, and healing underexplored.

Text-to-Image Evaluation Benchmarks and Metrics. Existing T2I evaluation primarily targets semantic correctness, visual quality, compositional reasoning, and generic preference. CLIP-based metrics estimate text-image alignment (Radford et al. 2021); reward and preference models such as PickScore, ImageReward, HPS, and HPSv3 capture overall human preference (Kirstain et al. 2023; Xu et al. 2023; Wu et al. 2023; Ma et al. 2025); and aesthetic metrics such as AVA and NIMA assess visual appeal (Murray, Marchesotti, and Perronnin 2012; Talebi and Milanfar 2018). Broader benchmarks, including DrawBench, GenEval, T2I-CompBench, VQAScore / GenAI-Bench, and HEIM, evaluate compositionality, spatial relations, reasoning, robustness, safety, and overall quality (Saharia et al. 2022; Ghosh, Hajishirzi, and Schmidt 2023; Huang et al. 2023; Lee et al. 2023; Lin et al. 2024; Li et al. 2024). Yet these metrics mainly ask whether an image is correct, high-quality, or preferred, rather than whether it is psychologically comforting. Although HRS-Bench introduces emotion-grounded evaluation (Bakr et al. 2023), and HEIM and GenAI-Bench cover broad evaluation dimensions (Lee et al. 2023; Li et al. 2024), deeper affective outcomes such as soothingness and emotional comfort remain insufficiently evaluated. Recent studies further show that metric rankings can be unstable across prompts and protocols (Wiles et al. 2025), and that faithfulness metrics remain unreliable on challenging subsets (Saxon et al. 2024). HealBench addresses this gap by treating perceived psychological comfort as a dedicated, fine-grained, and failure-diagnostic evaluation target.

Method

Overview

HealBench is constructed in four stages: (1) taxonomy design and web-video-based image collection; (2) captioning, hierarchical labeling, and reward-model-based filtering; (3) category-aware evaluation protocol construction and HealScore annotation; (4) tri-dimensional scorer training and data augmentation. Experts participate in taxonomy design, question revision, and iterative quality checks, while VLMs/LLMs are used as scalable assistants.

Data Collection and Filtering

We first build a two-level taxonomy for psychologically comforting imagery. Experts define 18 coarse-grained categories, supplemented with an *Others* class, to cover common restorative visual functions. The coarse-grained category set is $\mathcal{T} = \{T_i\}_{i=1}^{19}$. For each coarse category T_i, we further define a fine-grained subcategory set $\mathcal{S}_i = \{S_{i,j}\}_{j=1}^{n_i}$, which also includes an *Others* subtype. In total, the taxonomy contains 19 coarse-grained categories and 102 fine-grained subcategories. Each image I is assigned to exactly one category pair:

$$c(I) = (T_i, S_{i,j}), \quad T_i \in \mathcal{T},\ S_{i,j} \in \mathcal{S}_i. \tag{1}$$

Based on this taxonomy, we manually design category-aware retrieval templates to crawl web videos and obtain a raw video set ($\mathcal{D}_{\text{video}}$). We then extract frames from these videos, producing a raw image set ($\mathcal{D}_{\text{raw}}$) of about 100K frames. For each image in $\mathcal{D}_{\text{raw}}$, we use QWEN2.5-VL-7B-INSTRUCT to generate a caption, and then use Claude to assign hierarchical labels. Manual inspection removes invalid or ambiguous category pairs, leaving 92 valid pairs for subsequent experiments. The category pairs and the detailed data collection and annotation process are shown in Fig.3. To remove noisy samples caused by semantic drift, poor composition, and low visual

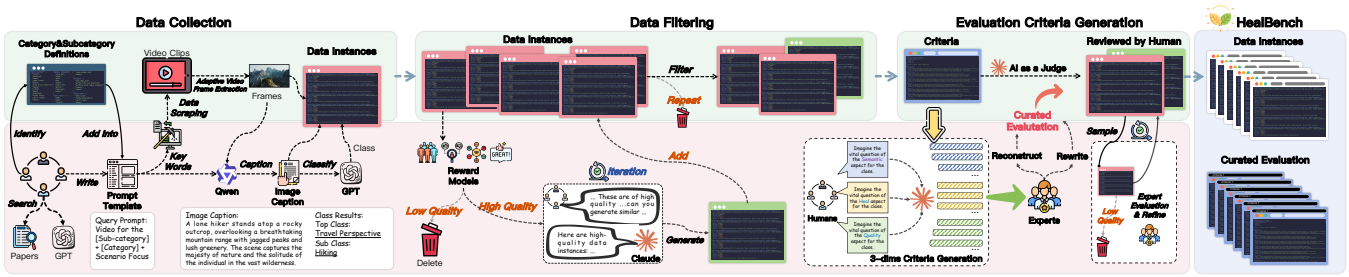

Figure 3: Overview of the HealBench construction pipeline. The benchmark is built through three stages: data collection, reward-model-based filtering and augmentation, and tri-dimensional evaluation-criteria generation. Human and expert review is incorporated throughout the pipeline to ensure the quality of both data instances and evaluation criteria.

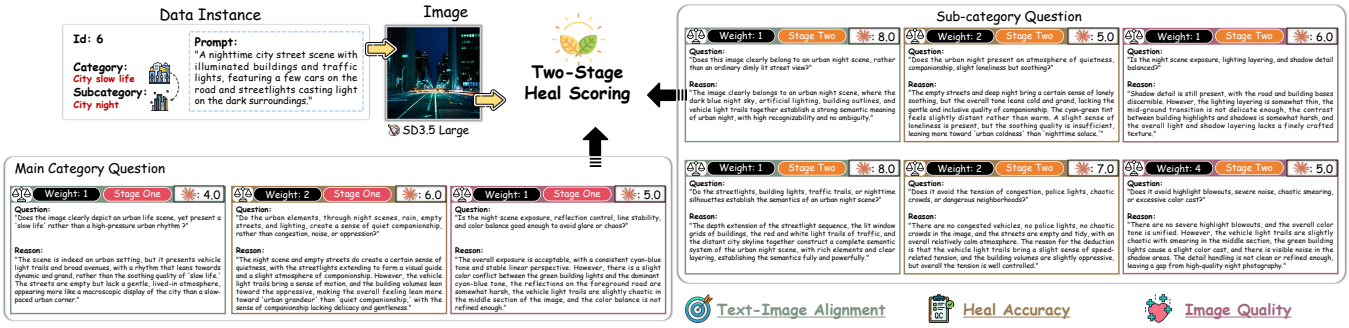

Figure 4: Overview of the two-stage HealScore pipeline. HealScore first evaluates each image with three category-level questions on *text-image alignment*, *heal accuracy*, and *image quality*, and then refines the judgment with six subcategory-level questions. The final score is computed by weighted aggregation of the resulting multi-dimensional responses.

quality, we rank images with four complementary reward signals: CLIP Score, Aesthetic Score, PickScore, and HPSv3. These models serve as coarse quality gates rather than final healing labels. Denote their scores by $r_{\text{clip}}, r_{\text{aes}}, r_{\text{pick}}, r_{\text{hps}}$, respectively. After independent normalization, we compute the overall score as

$$r_{\text{overall}} = \tilde{r}_{\text{clip}} + \tilde{r}_{\text{aes}} + \tilde{r}_{\text{pick}} + \tilde{r}_{\text{hps}}. \tag{2}$$

Within each fine-grained subcategory, we rank images by r_{overall} and retain the top 33%, which removes low-quality candidates while preserving category diversity, yielding the filtered real-image subset $\mathcal{D}_{\text{filter}}$.

Evaluation Criteria and HealScore

Existing image-generation metrics mainly emphasize semantic alignment and aesthetic preference, but are insufficient for evaluating psychological soothing effects. We therefore propose *HealScore*, a structured metric defined on three dimensions: text-image alignment, heal accuracy, and image quality. Here, heal accuracy refers to the perceived healing degree of the image under its target category, not to clinical efficacy. For each category pair $c = (T_i, S_{i,j})$, we build a category-aware question set. An LLM first generates one coarse-level question for the general healing function of T_i and two fine-grained questions for the concrete visual or affective cues of $S_{i,j}$ under each dimension. Experts then revise the questions to remove ambiguous, redundant, or category-mismatched formulations. The resulting question set is

$$\mathcal{Q}_c = \left\{ q^c_{d,k} \mid d \in \{\text{sem}, \text{heal}, \text{qual}\},\ k \in \{0,1,2\} \right\}, \tag{3}$$

where $q^c_{d,0}$ is the coarse-level question and $q^c_{d,1}, q^c_{d,2}$ are fine-grained supplementary questions. For each image $I \in \mathcal{D}_{\text{filter}}$, an LLM scores every question in $\mathcal{Q}_{c(I)}$ on a 1–10 scale. Let the resulting scores be $r_{\text{sem},k}(I)$, $r_{\text{heal},k}(I)$, and $r_{\text{qual},k}(I)$. We aggregate them within each dimension:

$$s_d(I) = \sum_{k=0}^{2} r_{d,k}(I), \quad d \in \{\text{sem}, \text{heal}, \text{qual}\}. \tag{4}$$

We report the three dimension scores for diagnosis and use a compact summary score only for ranking, training, and large-scale comparison. Since psychological comfort is the primary objective of the benchmark, the summary score gives the healing dimension twice the weight of the alignment and quality dimensions:

$$\text{HealScore}(I) = s_{\text{sem}}(I) + 2\, s_{\text{heal}}(I) + s_{\text{qual}}(I). \tag{5}$$

The 1:2:1 design therefore emphasizes the target dimension while retaining the two supporting constraints, rather than masking a failed dimension in diagnostic interpretation. To reduce annotation noise, we manually inspect sampled scored results and iteratively refine the question set by expert feedback. The detailed scoring process of HealScore is shown in Fig.4.

To further enrich the dataset, we select, for each category pair, the top-10 images ranked by HealScore and the top-10 images ranked by the normalized reward-model total score. Their captions are rewritten by Claude and used as prompts for SD3.5-Large to synthesize additional images, producing an augmented set $\mathcal{D}_{\text{augment}}$. We merge $\mathcal{D}_{\text{augment}}$ with $\mathcal{D}_{\text{filter}}$ and remove near-duplicates based on image similarity, obtaining the final dataset $\mathcal{D}_{\text{heal}}$.

Tri-dimensional Heal Scorer (THS)

To approximate rule-based HealScore with a learnable model, we train a *Tri-dimensional Heal Scorer (THS)*. Given an image and its category pair $(T_i, S_{i,j})$, THS predicts three dimension-specific scores for text-image alignment, heal accuracy, and image quality, and combines them into the summary score. We adopt BLIP image-text matching as the backbone. Given image I and category text

$$x = \texttt{"An image of } T_i\texttt{, } S_{i,j}\texttt{"}, \tag{6}$$

we encode image and text separately, project them into a shared embedding space, and obtain normalized embeddings v and t. Their fused representation is defined as $h = [v; t; v \odot t]$, where $[\cdot\,;\,\cdot]$ denotes concatenation and $\odot$ denotes element-wise multiplication. The fused feature is fed into three independent MLP heads to predict the three dimension scores. Let the supervision targets be $y_{\text{sem}}, y_{\text{heal}}$, and y_{qual}, with corresponding predictions $\hat{y}_{\text{sem}}, \hat{y}_{\text{heal}}$, and $\hat{y}_{\text{qual}}$. We first optimize the dimension-wise regression loss:

$$\mathcal{L}_{\text{dim}} = \sum_{d \in \{\text{sem,heal,qual}\}} \|\hat{y}_d - y_d\|^2 . \tag{7}$$

To enforce consistency with the summary HealScore while preserving dimension-wise supervision, we further define

$$\begin{aligned} y_{\text{total}} &= y_{\text{sem}} + 2y_{\text{heal}} + y_{\text{qual}}, \\ \hat{y}_{\text{total}} &= \hat{y}_{\text{sem}} + 2\hat{y}_{\text{heal}} + \hat{y}_{\text{qual}}, \end{aligned} \tag{8}$$

and apply an additional total-score loss: $\mathcal{L}_{\text{total}} = \|\hat{y}_{\text{total}} - y_{\text{total}}\|^2$. The final objective is $\mathcal{L} = \mathcal{L}_{\text{dim}} + \lambda \mathcal{L}_{\text{total}}$, where λ balances the two terms. We optimize THS with AdamW, using different learning rates for the pretrained backbone and the newly initialized heads. During training, we retain only samples with complete annotations on all three dimensions, and select the final model according to the validation error on $\mathcal{L}_{\text{total}}$.

Experiments

Overview

We evaluate HealBench from three perspectives: validity, utility, and scalability, covering benchmark representativeness, tri-dimensional diagnosis, feedback-based improvement, automatic scoring, and downstream optimization.

Table 2: Agreement between HealScore-induced rankings and pairwise preferences from human annotators and VLM-based judges.

Model	Accuracy	Kendall's τ	Spearman's ρ
Haiku	0.6831	0.3546	0.4319
Sonnet 4-6	0.6933	0.3868	0.4711
People	0.8012	0.4487	0.5448

HealBench-Eval Construction and Validation

Construction To enable controlled evaluation without sacrificing diversity, we construct a compact HealBench-Eval from the filtered real-image subset. For each of the 92 valid category pairs, we extract SigLIP embeddings, run K-means clustering, and select representative samples. Figure 5 shows that the resulting subset follows the overall HealBench distribution rather than concentrating in a few dense regions, making it suitable for controlled healing-oriented evaluation.

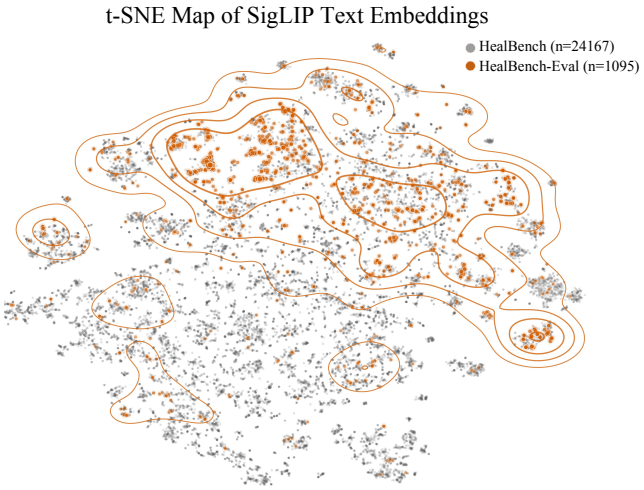

Figure 5: t-SNE visualization of SigLIP embeddings for the full HealBench data and the selected HealBench-Eval.

Validation To test whether the tri-dimensional formulation captures a meaningful healing-oriented signal, we compare HealScore-induced rankings with human and VLM preferences on 1,500 image pairs. Annotators, Claude-Haiku, and Claude-Sonnet 4-6 select the more psychologically comforting image. Table 2 shows that human judgments achieve the highest agreement, while Sonnet slightly outperforms Haiku, indicating that HealScore provides a meaningful approximation of this structured preference.

Table 3: Quantitative comparison of different text-to-image models across multiple evaluation metrics.

Model	LAION Aesthetic ↑	CLIPScore ↑	PickScore ↑	HPSv3 ↑
SD 2.1	4.403	0.885	0.785	3.406
SDXL Base 1.0	4.631	0.885	0.790	3.602
pixArt-XL	4.717	0.946	0.864	12.623
SDXL Turbo	4.095	0.975	0.812	5.543
Playground v2.5	5.164	**0.988**	**0.899**	12.712
Hunyuan-Dit	**5.271**	0.976	0.872	10.991
SD 3 Medium	4.841	0.965	0.863	11.515
Kolors	4.638	0.901	0.823	9.587
FLUX.1 [dev]	4.798	0.938	0.876	**13.585**
FLUX.1 [schnell]	4.682	0.943	0.840	10.624
SD 3.5 Large	4.912	0.971	0.867	11.792
SD 3.5 Medium	4.867	0.974	0.867	11.533
sana16	4.666	0.967	0.870	12.955
Lumina-Image 2.0	4.950	0.984	0.864	11.919
GT	4.954	0.925	0.838	9.797

Benchmarking Baseline T2I Models

We next benchmark representative diffusion and flow-matching T2I models under the original prompts on healing-

Table 4: Baseline summary table. All metrics are positive indicators. For prompt-rewrite models, the orange mark denotes the improvement over the corresponding non-rewrite base model. For image editing models, the blue mark denotes the improvement over the corresponding original model.

Model	Overall	Text-Image Alignment	Heal Accuracy	Image Quality	Natural Restoration	Atmospheric Immersion	Shelter and Safety	Living Companionship	Rhythmic Focus	Awe and Transcendence
Diffusion Models										
SD 2.1	19.390	5.172	4.820	4.578	18.754	19.678	18.484	21.780	19.412	20.147
SDXL Base 1.0	20.723	5.319	5.173	5.057	20.258	21.269	21.915	24.284	19.859	20.481
PixArt-XL	25.080	6.435	6.169	6.308	24.953	24.832	25.115	27.787	25.036	24.844
SDXL Turbo	18.765	5.419	4.517	4.312	17.372	18.689	17.536	19.837	20.549	18.695
Playground v2.5	25.088	6.479	6.094	6.422	24.551	25.304	24.449	28.177	25.162	25.253
Hunyuan-DiT	25.020	6.378	6.178	6.286	24.341	24.849	24.286	28.057	25.427	25.347
Kolors	23.535	6.005	5.866	5.797	23.586	23.924	22.832	26.369	23.267	23.255
Sana-1.6B	24.419	6.386	5.917	6.201	23.748	25.044	24.291	29.624	23.821	24.541
Flow Matching Models										
SD 3 Medium	23.465	6.121	5.731	5.882	22.556	23.691	21.617	28.560	23.462	24.780
FLUX.1 [dev]	24.836	6.289	6.152	6.243	24.211	24.195	23.489	29.759	25.424	25.222
FLUX.1 [schnell]	22.717	5.954	5.561	5.642	22.080	21.632	21.108	27.830	23.225	23.913
SD 3.5 Large	24.571	6.346	6.008	6.210	23.694	24.381	23.160	29.057	25.106	25.154
SD 3.5 Medium	24.164	6.321	5.884	6.075	23.285	24.131	22.391	28.730	24.266	25.620
Lumina-Image 2.0	23.169	6.132	5.627	5.782	22.483	22.855	22.035	28.447	23.433	23.602
Prompt-Rewrite Models										
Hunyuan-DiT (Rewrite)	26.631 ↑1.612	6.617 ↑0.239	6.711 ↑0.533	6.592 ↑0.306	26.841 ↑2.500	25.843 ↑0.994	26.336 ↑2.050	29.546 ↑1.489	27.183 ↑1.757	25.574 ↑0.226
FLUX.1 [schnell] (Rewrite)	24.514 ↑1.797	6.328 ↑0.374	6.051 ↑0.490	6.084 ↑0.442	24.411 ↑2.331	23.090 ↑1.459	23.942 ↑2.835	28.794 ↑0.965	25.043 ↑1.818	24.552 ↑0.638
SD 3.5 Large (Rewrite)	26.149 ↑1.578	6.638 ↑0.293	6.471 ↑0.464	6.568 ↑0.358	25.851 ↑2.157	24.993 ↑0.611	25.825 ↑2.664	31.057 ↑2.000	26.680 ↑1.575	25.836 ↑0.683
SD 3.5 Medium (Rewrite)	25.598 ↑1.434	6.556 ↑0.235	6.322 ↑0.438	6.398 ↑0.323	25.071 ↑1.786	24.683 ↑0.552	24.554 ↑2.163	29.979 ↑1.248	26.179 ↑1.912	26.046 ↑0.426
Image Editing Models										
Pix2Pix (Original)	20.666	5.454	5.081	5.050	20.349	21.401	20.832	23.837	20.832	19.059
Pix2Pix (Rewrite)	22.802 ↑2.136	5.853 ↑0.399	5.664 ↑0.583	5.620 ↑0.570	23.047 ↑2.698	23.193 ↑1.792	23.649 ↑2.817	25.674 ↑1.837	22.719 ↑1.887	20.606 ↑1.547
SDXL-Pix2Pix (Original)	19.307	5.137	4.722	4.727	18.612	19.576	18.351	23.780	19.272	19.733
SDXL-Pix2Pix (Rewrite)	22.053 ↑2.746	5.711 ↑0.574	5.445 ↑0.723	5.453 ↑0.726	22.111 ↑3.499	22.575 ↑2.999	21.464 ↑3.113	26.255 ↑2.475	21.908 ↑2.636	20.917 ↑1.184
Ground Truth										
Real Images	24.528	6.375	5.995	6.163	23.186	24.119	25.205	27.901	25.216	24.503

oriented generation. Table 3 reports standard metrics as reference baselines. Recent models perform strongly on Aesthetic score, CLIPScore, PickScore, and HPSv3, sometimes exceeding the real-image reference on individual metrics. Because these metrics do not target perceived psychological comfort, they should not be treated as evidence of healing-oriented generation quality.

Overall healing-oriented performance. Table 4 shows clear gaps across model families. Among original baselines, *Playground v2.5* (25.088), *PixArt-XL* (25.080), and *Hunyuan-DiT* (25.020) perform best overall, while *SDXL Turbo* (18.765) and *SD 2.1* (19.390) lag behind. Even when advanced generators approach the real-image reference, their scores remain uneven across dimensions and healing functions.

Tri-dimensional analysis. In Table 4, the three dimensions reveal why aggregate scores are insufficient. Strong baselines usually score above 6.0 on *Text-Image Alignment* and *Image Quality*, but are lower and more variable on *Heal Accuracy*: *Playground v2.5* obtains 6.479/6.422 on Alignment/Quality but 6.094 on Heal Accuracy, and *SD 3.5 Medium* obtains 6.321/6.075 but 5.884. Thus, perceived healing degree cannot be inferred from alignment or visual quality alone.

Class-level analysis. As shown in Table 4, Class-wise results further show uneven difficulty. *Living Companionship* is often easiest, reaching 29.624 for *Sana-1.6B*, 29.759 for *FLUX.1 [dev]*, and 29.057 for *SD 3.5 Large*. In contrast, *Shelter and Safety* and *Natural Restoration* show larger model-dependent variation, and no model dominates all six classes. HealBench therefore exposes both dimension-specific and function-specific limitations.

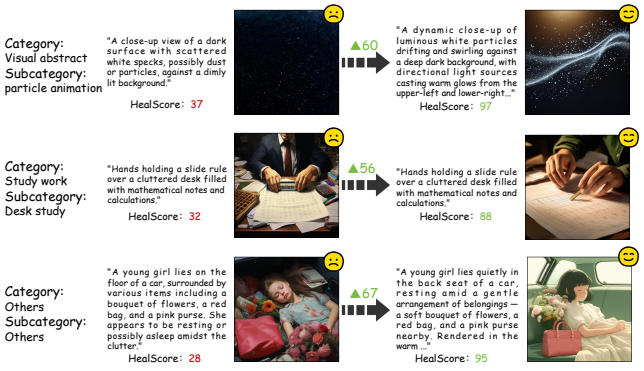

Figure 6: Qualitative examples of our prompt rewriting strategy. Using the same image generation model, targeted prompt rewriting markedly improves generation outcomes, yielding more coherent and safer images with substantially higher HealScores across multiple categories.

Prompt Rewriting as Intervention

We next test whether score-plus-rationale feedback can improve prompts on healing-oriented generation for both text-to-image generation and image editing.

Prompt rewriting improves healing-oriented generation. Table 4 shows that rewriting improves all tested T2I models,

including *Hunyuan-DiT*, *FLUX.1 [schnell]*, *SD 3.5 Large*, and *SD 3.5 Medium*, with the most consistent gains on *Heal Accuracy*.

Prompt rewriting also benefits image editing. For editing, rewritten prompts improve both *Pix2Pix* and *SDXL-Pix2Pix* in overall HealScore and all three dimensions.

Qualitative analysis. Figure 6 shows that rewritten prompts produce more coherent, less visually tense images aligned with the intended healing semantics.

Table 5: User study on personalized healing delivery with 24 subthreshold-depression college students. Participants in the personalized condition could choose among the six HEAL classes, while those in the random condition viewed randomly assigned healing images. Mood change was measured with HAMD-17 after one hour of viewing.

Condition	N	Improved Ratio (%) ↑	HAMD-17 ↓
Random playback	12	45.5	-1.27
Personalized class selection	12	75.0	-3.42
Difference / Gain	–	+29.5	-2.15

User Study on Personalized Healing Delivery

In a preliminary study with 24 college students with subthreshold depression, participants either choose a HEAL class or view randomly assigned healing images. Personalized selection yields a higher improvement ratio (75.0% vs. 45.5%) and larger HAMD-17 reduction (-3.42 vs. -1.27), suggesting that the six-class organization is useful beyond benchmark analysis.

Table 6: THS vs. generic scorers on pairwise healing-preference ranking.

Model	Acc. ↑	Kendall's τ ↑	Spearman's ρ ↑
PickScore	0.5593	0.1525	0.2254
LAION Aesthetic	0.5507	0.1077	0.1595
CLIPScore	0.5413	0.1272	0.1882
HPSv3	0.5013	0.0379	0.0595
THS	**0.8527**	**0.7741**	**0.9223**

THS as an Automatic Healing Scorer

We next examine whether THS can provide a more suitable automatic evaluation signal for healing-oriented generation than generic scoring models. To this end, we compare THS with several widely used baselines, including PickScore, LAION Aesthetic, CLIPScore, and HPSv3, using the same pairwise ranking protocol as in the preference validation experiment. As shown in Table 6, THS substantially outperforms all generic baselines across all three ranking metrics, achieving 0.8527 in Accuracy, 0.7741 in Kendall's τ, and 0.9223 in Spearman's ρ. By contrast, the generic scorers remain far behind, with Accuracy values clustered around 0.50–0.56 and much weaker rank correlation. These results indicate that existing scoring models for semantics, aesthetics, or broad human preference do not adequately capture the specific notion of perceived psychological comfort targeted in this work. Results justify the need for a task-specific automatic scorer. Rather than serving only as a descriptive metric, THS provides a scalable evaluation signal that is better aligned with healing-oriented preference and can therefore support both large-scale benchmarking and downstream optimization.

Table 7: Main results of the base SD 3.5 Medium model and its Flow-GRPO fine-tuned variant. Higher values are highlighted in blue.

Model	Overall	Semantic	Heal	Quality
SD 3.5 Medi (base)	24.16	6.32	5.88	6.08
SD 3.5 Medi + Flow-GRPO	**25.67**	**6.48**	**6.45**	**6.28**

Exploratory RL Fine-Tuning with THS Reward

Finally, we conduct a preliminary study to examine whether THS can serve not only as an evaluator, but also as a reward signal for direct model optimization. To this end, we fine-tune *SD 3.5 Medium* with Flow-GRPO using a reward defined by a weighted combination of THS and CLIPScore. As shown in Table 7, the Flow-GRPO fine-tuned model improves over the base *SD 3.5 Medium* across all four metrics, including *Overall* (23.47 → 25.67), *Semantic* (6.12 → 6.48), *Heal* (5.73 → 6.45), and *Quality* (5.88 → 6.28). Notably, the largest gain appears on the *Heal* dimension, which is consistent with the role of THS as a healing-oriented reward signal. This suggests that the proposed scorer is not only useful for benchmark evaluation, but can also provide an effective optimization signal for improving healing-oriented generation. We emphasize that this experiment is exploratory rather than a fully optimized RL recipe. Nevertheless, the consistent gains across all dimensions provide initial evidence that THS can support downstream alignment toward psychologically comforting image generation.

Conclusion

We present **HealBench**, the first fine-grained benchmark for evaluating perceived psychological comfort in text-to-image generation and image editing. By explicitly defining healing-oriented generation along three complementary dimensions (*text-image alignment*, *heal accuracy*, and *image quality*), we move psychological comfort from an implicit notion of preference to a structured and measurable objective. Building on this formulation, we develop a large-scale human-and-expert-in-the-loop benchmark, introduce the rule-based **HealScore**, and further propose **THS**, a learnable automatic evaluator. Our experiments reveal a clear gap between strong performance on conventional generation metrics and true effectiveness under healing-oriented evaluation. Beyond benchmarking, we show that structured score-plus-rationale feedback can directly support generation improvement, and that the HEAL taxonomy provides a useful basis for personalized healing-content delivery. In conclusion, our results establish psychological comfort as a distinct evaluation and optimization target in generative vision, and position HealBench as a strong foundation for future work on affect-aware, human-centered multimodal generation.

References

Bakr, E. M.; Sun, P.; Shen, X.; Khan, F. F.; Li, L. E.; and Elhoseiny, M. 2023. Hrs-bench: Holistic, reliable and scalable benchmark for text-to-image models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 20041–20053.

Belanche, D.; Ibáñez-Sánchez, S.; Jordán, P.; and Matas, S. 2025. Customer reactions to generative AI vs. real images in high-involvement and hedonic services. *International Journal of Information Management*, 85: 102954.

Black, K.; Janner, M.; Du, Y.; Kostrikov, I.; and Levine, S. 2024. Training Diffusion Models with Reinforcement Learning. In *The Twelfth International Conference on Learning Representations*.

Brooks, T.; Holynski, A.; and Efros, A. A. 2023. Instructpix2pix: Learning to follow image editing instructions. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 18392–18402.

Chang, H.; Zhang, H.; Barber, J.; Maschinot, A.; Lezama, J.; Jiang, L.; Yang, M.-H.; Murphy, K. P.; Freeman, W. T.; Rubinstein, M.; et al. 2023. Muse: Text-To-Image Generation via Masked Generative Transformers. In *International Conference on Machine Learning*, 4055–4075. PMLR.

Esser, P.; Kulal, S.; Blattmann, A.; Entezari, R.; Müller, J.; Saini, H.; Levi, Y.; Lorenz, D.; Sauer, A.; Boesel, F.; et al. 2024. Scaling rectified flow transformers for high-resolution image synthesis. In *Forty-first international conference on machine learning*.

Fan, Y.; Watkins, O.; Du, Y.; Liu, H.; Ryu, M.; Boutilier, C.; Abbeel, P.; Ghavamzadeh, M.; Lee, K.; and Lee, K. 2023. Dpok: Reinforcement learning for fine-tuning text-to-image diffusion models. *Advances in Neural Information Processing Systems*, 36: 79858–79885.

Fu, F.; Zhang, L.; Huang, M.; and Mao, Z. 2025a. Feededit: Text-based image editing with dynamic feedback regulation. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, 2661–2670.

Fu, M.; Wang, G.-H.; Cao, L.; Chen, Q.-G.; Xu, Z.; Luo, W.; and Zhang, K. 2025b. CHATS: Combining Human-Aligned Optimization and Test-Time Sampling for Text-to-Image Generation. In *Forty-second International Conference on Machine Learning*.

Gafni, O.; Polyak, A.; Ashual, O.; Sheynin, S.; Parikh, D.; and Taigman, Y. 2022. Make-a-scene: Scene-based text-to-image generation with human priors. In *European conference on computer vision*, 89–106. Springer.

Gao, C.; Ajith, S.; and Peelen, M. V. 2025. Object representations drive emotion schemas across a large and diverse set of daily-life scenes. *Communications Biology*, 8(1): 697.

Ghosh, D.; Hajishirzi, H.; and Schmidt, L. 2023. Geneval: An object-focused framework for evaluating text-to-image alignment. *Advances in Neural Information Processing Systems*, 36: 52132–52152.

Hertz, A.; Mokady, R.; Tenenbaum, J.; Aberman, K.; Pritch, Y.; and Cohen-Or, D. 2023. Prompt-to-Prompt Image Editing with Cross-Attention Control. In *The Eleventh International Conference on Learning Representations*.

Ho, J.; Jain, A.; and Abbeel, P. 2020. Denoising diffusion probabilistic models. *Advances in neural information processing systems*, 33: 6840–6851.

Huang, K.; Sun, K.; Xie, E.; Li, Z.; and Liu, X. 2023. T2i-compbench: A comprehensive benchmark for open-world compositional text-to-image generation. *Advances in Neural Information Processing Systems*, 36: 78723–78747.

Kirstain, Y.; Polyak, A.; Singer, U.; Matiana, S.; Penna, J.; and Levy, O. 2023. Pick-a-pic: An open dataset of user preferences for text-to-image generation. *Advances in neural information processing systems*, 36: 36652–36663.

Lee, T.; Yasunaga, M.; Meng, C.; Mai, Y.; Park, J. S.; Gupta, A.; Zhang, Y.; Narayanan, D.; Teufel, H.; Bellagente, M.; et al. 2023. Holistic evaluation of text-to-image models. *Advances in Neural Information Processing Systems*, 36: 69981–70011.

Li, B.; Lin, Z.; Pathak, D.; Li, J.; Fei, Y.; Wu, K.; Xia, X.; Zhang, P.; Neubig, G.; and Ramanan, D. 2024. Evaluating and improving compositional text-to-visual generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 5290–5301.

Lin, Z.; Pathak, D.; Li, B.; Li, J.; Xia, X.; Neubig, G.; Zhang, P.; and Ramanan, D. 2024. Evaluating text-to-visual generation with image-to-text generation. In *European Conference on Computer Vision*, 366–384. Springer.

Lipman, Y.; Chen, R. T.; Ben-Hamu, H.; Nickel, M.; and Le, M. 2023. Flow Matching for Generative Modeling. In *The Eleventh International Conference on Learning Representations*.

Liu, H.; Wang, J.; Huang, R.; Liu, Y.; Lu, H.; Zhao, Z.; and Xue, W. 2025. Flashaudio: Rectified flow for fast and high-fidelity text-to-audio generation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 13694–13710.

Ma, Y.; Wu, X.; Sun, K.; and Li, H. 2025. Hpsv3: Towards wide-spectrum human preference score. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 15086–15095.

Mokady, R.; Hertz, A.; Aberman, K.; Pritch, Y.; and Cohen-Or, D. 2023. Null-text inversion for editing real images using guided diffusion models. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 6038–6047.

Murray, N.; Marchesotti, L.; and Perronnin, F. 2012. AVA: A large-scale database for aesthetic visual analysis. In *2012 IEEE conference on computer vision and pattern recognition*, 2408–2415. IEEE.

Ortega, J.; Murai, Y.; and Whitney, D. 2025. Integration of affective cues in context-rich and dynamic scenes varies across individuals. *Nature Communications*, 17: 786.

Peebles, W.; and Xie, S. 2023. Scalable diffusion models with transformers. In *Proceedings of the IEEE/CVF international conference on computer vision*, 4195–4205.

Radford, A.; Kim, J. W.; Hallacy, C.; Ramesh, A.; Goh, G.; Agarwal, S.; Sastry, G.; Askell, A.; Mishkin, P.; Clark, J.; et al. 2021. Learning transferable visual models from

natural language supervision. In *International conference on machine learning*, 8748–8763. PmLR.

Ramesh, A.; Pavlov, M.; Goh, G.; Gray, S.; Voss, C.; Radford, A.; Chen, M.; and Sutskever, I. 2021. Zero-shot text-to-image generation. In *International conference on machine learning*, 8821–8831. Pmlr.

Rombach, R.; Blattmann, A.; Lorenz, D.; Esser, P.; and Ommer, B. 2022. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 10684–10695.

Saharia, C.; Chan, W.; Saxena, S.; Li, L.; Whang, J.; Denton, E. L.; Ghasemipour, K.; Gontijo Lopes, R.; Karagol Ayan, B.; Salimans, T.; et al. 2022. Photorealistic text-to-image diffusion models with deep language understanding. *Advances in neural information processing systems*, 35: 36479–36494.

Saxon, M.; Jahara, F.; Khoshnoodi, M.; Lu, Y.; Sharma, A.; and Wang, W. Y. 2024. Who evaluates the evaluations? objectively scoring text-to-image prompt coherence metrics with t2iscorescore (ts2). *Advances in Neural Information Processing Systems*, 37: 85630–85657.

Shi, J.; Zhang, N.; Sun, Y.; Jiang, J.; Duan, H.; and Gao, W. 2025. The impact of virtual images of coastal landscape features on stress recovery based on EEG. *Scientific Reports*, 15(1): 17369.

Talebi, H.; and Milanfar, P. 2018. NIMA: Neural image assessment. *IEEE transactions on image processing*, 27(8): 3998–4011.

Wallace, B.; Dang, M.; Rafailov, R.; Zhou, L.; Lou, A.; Purushwalkam, S.; Ermon, S.; Xiong, C.; Joty, S.; and Naik, N. 2024. Diffusion model alignment using direct preference optimization. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 8228–8238.

Weng, S.; Zhang, P.; Chang, Z.; Wang, X.; Li, S.; and Shi, B. 2023. Affective image filter: Reflecting emotions from text to images. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 10810–10819.

Wiles, O.; Zhang, C.; Albuquerque, I.; Kajic, I.; Wang, S.; Bugliarello, E.; Onoe, Y.; Papalampidi, P.; Ktena, I.; Knutsen, C.; et al. 2025. Revisiting text-to-image evaluation with Gecko: on metrics, prompts, and human rating. In *The Thirteenth International Conference on Learning Representations*.

Wu, X.; Sun, K.; Zhu, F.; Zhao, R.; and Li, H. 2023. Human preference score: Better aligning text-to-image models with human preference. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2096–2105.

Xu, J.; Liu, X.; Wu, Y.; Tong, Y.; Li, Q.; Ding, M.; Tang, J.; and Dong, Y. 2023. Imagereward: Learning and evaluating human preferences for text-to-image generation. *Advances in Neural Information Processing Systems*, 36: 15903–15935.

Yang, J.; Feng, J.; Luo, W.; Lischinski, D.; Cohen-Or, D.; and Huang, H. 2025. Emoedit: Evoking emotions through image manipulation. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, 24690–24699.

Yu, Q.; Chow, W.; Yue, Z.; Pan, K.; Wu, Y.; Wan, X.; Li, J.; Tang, S.; Zhang, H.; and Zhuang, Y. 2025. Anyedit: Mastering unified high-quality image editing for any idea. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, 26125–26135.