NeSySpec: Neuro-Symbolic Technical Intelligence for Auditing Engineering Specifications with an Application to Aircraft

Pier Mattia Basciano[1,2,*,†] Gianvito Antenori[1,2] Cristiano Andolfi[2]
Giovanni Farina[1] Salvatore Monteleone[1]

[1] *Department of Engineering, Niccolò Cusano University, Rome, Italy*

[2] *SOMA Research Division, Rome, Italy*

{piermattia.basciano, gianvito.antenori, giovanni.farina, salvatore.monteleone}@unicusano.it

{piermattia.basciano, gianvito.antenori, cristiano.andolfi}@somatwin.com

[*] Corresponding author. [†] Lead author, who carried out the large majority of the work.

arXiv:2610.35403v1 [cs.AI] 26 Oct 2026

Abstract

Engineers who audit a declared specification of an engineered system need to know which physical constraints it respects, which it violates, which field was altered and how to correct it. NeSySpec is an integrated and interpretable neuro-symbolic audit framework, built on Logic Tensor Networks, that answers these questions in one pipeline. Its axioms encode the physical relations of the domain as differentiable graded predicates whose tolerances are learned from a fleet of real records; from the same degrees, it diagnoses the violated constraint, identifies the altered field by abduction, proposes a minimum-change correction and ties image evidence to the parts it describes, abstaining exactly where evidence is missing, and returns all of it in one verification report. We validate NeSySpec on 107 aircraft types, with 16 test aircraft of 9 families held out of training, on specification records paired with ground photographs and satellite imagery and on synthetic errors: generated corruptions, edits labeled from the edit alone and deliberate falsifications. NeSySpec combines all the assessed audit outputs in one pipeline, which no learned, rule-based, solver-based, neuro-symbolic or language-model comparator does; on the test aircraft none is ahead of it in the macro-F1 of its constraint verdicts, in naming the violated constraint and the altered field or in the repair error, and it is competitive in record and image-pair detection. Its diagnosis names the violated constraint in 94.2% of corrupted test records, against 79.5% for classical Fellegi-Holt diagnosis, and on edits labeled from the edit alone, without retraining, abduction finds the altered field in 59.3% of records, against 30.6% for minimum-cardinality localization with random tie-breaking. Only the knowledge base, the visual predicates, the corruption generator and the data are specific to the domain.

Keywords: Engineering Specification Auditing; Neuro-Symbolic Artificial Intelligence; Logic Tensor Networks; Constraint-Based Diagnosis; Explainable Anomaly Detection; Reverse Engineering; Technical Intelligence; Multimodal Consistency

License. This work is licensed under a Creative Commons Attribution 4.0 International License (CC BY 4.0, `https://creativecommons.org/licenses/by/4.0/`). Third-party photographs reproduced in the figures are excluded from this license.

1 Introduction

Certification, maintenance and procurement engineers work with engineered systems they did not design and reverse-engineer them from declared specifications of masses, dimensions, ratings and performance. A record can be incomplete, inconsistent, valid for another variant or misreported, and public knowledge bases hold wrong numerical values [1]; the errors of a record propagate into the performance models it feeds [2]. Auditing such records falls within the scope of civil technical intelligence, which gathers and assesses information about technologies and technical products [3, 4]. Technical intelligence increasingly draws on open sources [5, 6]. For aircraft, photographs and satellite imagery can

Table 1: Acronyms.

Acronym	Meaning	Acronym	Meaning
AUROC	area under the receiver operating characteristic curve	MLP	multilayer perceptron
BCE	binary cross-entropy	MLP-SHAP	grouped SHAP attributions of the residual MLP
CAD	computer-aided design	MTOW	maximum take-off mass
CAM	class activation map	NeSy	neuro-symbolic artificial intelligence
CFD	computational fluid dynamics	OEW	operating empty mass
FAA	Federal Aviation Administration	PGD	projected gradient descent
FEM	finite element method	PSL	probabilistic soft logic
FPR	false-positive rate	RRR	Right for the Right Reasons
ISA	International Standard Atmosphere	SAM	Segment Anything Model
LEN	Logic Explained Networks	SHAP	Shapley additive explanations
LTN	Logic Tensor Networks	SMT	satisfiability modulo theories
MCAR	missing completely at random	SVM	support vector machine
MICE	multivariate imputation by chained equations	TPR	true-positive rate

also help verify the visually observable attributes reported in the specifications [7, 8]. We use reverse engineering in a first-approach sense: inferring which physical relations a declared specification and, when available, its image respect or violate, which field departs from them and which value restores them. Table 1 lists the acronyms.

Learned methods, from anomaly detectors to tabular foundation models and language models, score each record [9–11], with field-level attributions [12]. Statistical data editing and consistency-based diagnosis identify violated constraints and the fields implicated in those violations, based on bounds set by experts [13, 14]. Neuro-Symbolic Artificial Intelligence (NeSy) integrates learning with explicit knowledge and reasoning [15, 16], also surveyed for advanced air mobility [17], and can learn such bounds from data together with the constraints they delimit. To our knowledge, no previous work provides and evaluates together the outputs that such an audit needs, from the violated constraint and the altered field to part-grounded image evidence and a report (Table 2).

We present NeSySpec, an integrated and interpretable neuro-symbolic framework for auditing engineering specifications, built on Logic Tensor Networks (LTN) [19]: trained on fleet data, it determines through its logical axioms which physical constraints a declared record respects and which it violates, and returns a verification report (Figures 1 and 5). Its graded predicates measure each physical, definitional or regulatory relation by its residual against bands learned from the training fleet, and cross-modal axioms compare the fields with what a vision module reads. Masked quantification gives NeSySpec exact abstention on missing evidence: each verdict reduces to the constraints that the available evidence makes checkable, and the report lists every undeclared field with the constraints it leaves unchecked. The audit proceeds from detection to diagnosis and repair. It first scores the record using a masked mean of constraint violation degrees. This overall verdict can then be traced back to individual constraints, with the largest violation degree identifying the most violated constraint. An abductive search over single-field changes then localizes the altered field. Finally, a minimum-change repair searches for the smallest adjustment that restores coherence. A grounding axiom encourages each visual head to place its class-activation evidence on the part its predicate describes, so that the report ties image evidence to parts.

We validate NeSySpec on 107 aircraft, with specification records, FGVC-Aircraft photographs [7] and RarePlanes satellite imagery [8], holding 16 test aircraft of 9 families out of training, and evaluate constraint verdicts and diagnosis, field localization and repair, record coherence, missing evidence and absent images, the grounding of the visual axioms and robustness to deliberate falsification, on errors labeled by the corruption generator and on errors labeled from the edit alone (Section 4). The constraint verdicts reach a macro-F1 of 0.795 on generator labels and, without retraining, 0.575 on intervention labels, above a network trained with the DL2 loss on the same residuals and targets, the like-for-like neuro-symbolic test, and find every labeled violation in 0.934 of generator-labeled test corruptions against 0.522 for that network, with micro precision on par on test and lower on the

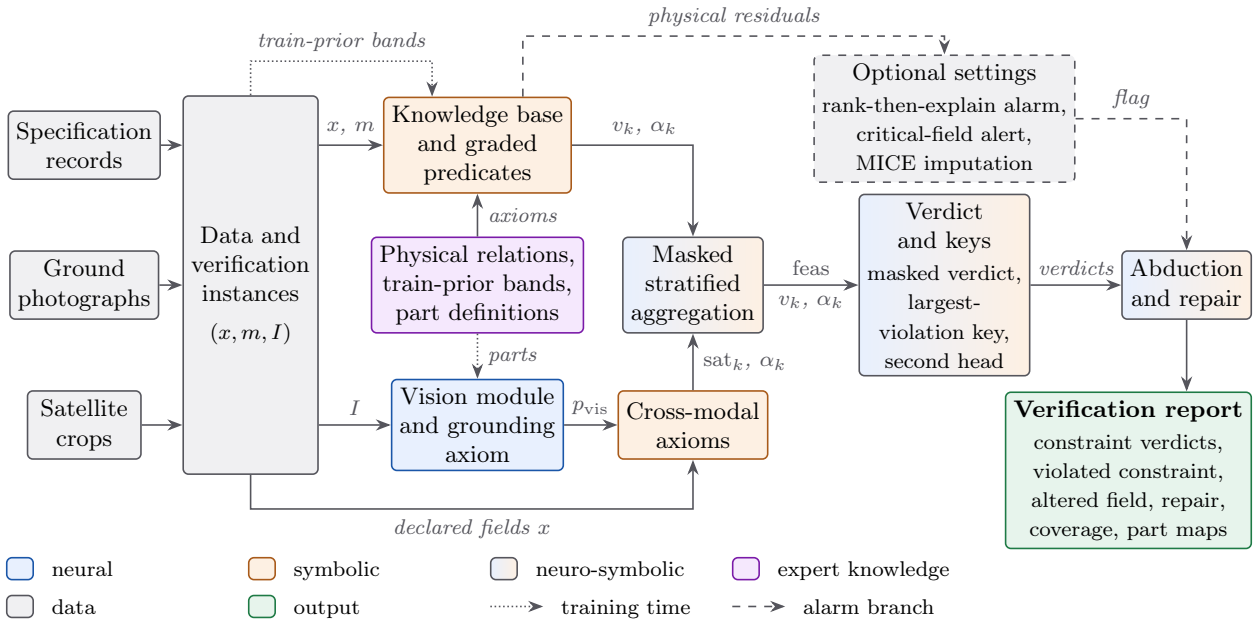

Figure 1: Overview of the NeSySpec pipeline.

unseen class (Section 4.2, Appendix G.1). The largest-violation key identifies the violated constraint in 0.942 of generator-labeled test corruptions, outperforming Fellegi–Holt, solver-based and rule-based diagnosis, DL2 applied to the declared fields, and a prompted language model. Abduction then traces the error to an individual field, correctly identifying the altered field in 0.593 of independently labeled errors. By comparison, minimum-cardinality localization reaches 0.306 with random tie-breaking and 0.453 with oracle tie-breaking; a sizing-regression residual also reaches 0.453 (Figure 6). The record verdict detects incoherent records at a test AUROC of 0.979 on generator labels, and on satellite pairs the multimodal verifier, which names the contradicted axiom, detects incoherent record-image pairs at 0.994 over all pair types, against 0.748 for a prompted vision-language model (Section 4.4). Table 5 lines these results up. No comparator returns every output of the audit, and on the test aircraft no comparator is ahead of NeSySpec in the macro-F1 of the constraint verdicts, in naming the violated constraint and the altered field or in the repair error on the altered field; in record and pair detection, learners on the residuals and late fusions are close to it and in some settings rank higher (Section 4.1).

Among the approaches evaluated here - the DL2 method on the declared fields, the network trained with the DL2 loss on the same residuals read per constraint, a constraint-layer repair ordering, LTN ablations, solver and rule diagnosis and prompted language and vision-language models -, NeSySpec is, to our knowledge, the most suitable neuro-symbolic framework for auditing engineering specifications: it provides and evaluates together the seven outputs of Table 2, from constraint verdicts and field localization to part-grounded image evidence and a verification report. Its logic contributes more to the physical audit than to record classification. Training its knowledge base through *Real Logic* (the differentiable formulation of first-order fuzzy logic underlying LTN [19]) improves detection over the same knowledge base left at its initial values. Detection also improves as more knowledge is encoded, a trend shared by a residual MLP using the same constraints (Section 4.8). Only the knowledge base, visual predicates, corruption generator and data of NeSySpec are aircraft-specific, so another engineered system enters through these four parts (Appendix A).

The contributions of this work are:

- An integrated and interpretable neuro-symbolic audit framework on Logic Tensor Networks whose graded residual-aware axioms, trained on fleet data, give each physical constraint of a declared engineering specification a held, violated or masked verdict and derive from these verdicts the altered field, a repair proposal, an exact accounting of missing evidence, part-grounded visual evidence and a verification report;

- A verification protocol that re-estimates the priors from training records, refits the generator

oracle without test or out-of-distribution aircraft, labels errors independently of the generator, reads the held and violated verdicts against per-constraint labels and stress-tests the verdicts against falsification and human part annotations;

- A comparative study of rule-based and solver-based methods, models trained with constraint losses, language and vision-language models, and residual learners, all using the same physical knowledge base (Table 5). NeSySpec identifies the violated constraint more often than classical diagnosis, DL2 applied to the declared fields, and a prompted language model. It also identifies the altered field more often than minimum-cardinality localization with random tie-breaking and produces a repair proposal from the same violation degrees (Figure 6);
- A deployment map (Table 6) that assigns each audit question to the output of NeSySpec that answers it, including its optional rank-then-explain alarm and critical-field alert.

The remainder of this paper is organized as follows. Section 2 reviews background and related works. Section 3 presents the NeSySpec framework in full. Section 4 reports experimental results. Section 5 concludes with directions for future work. Appendix A describes how to port NeSySpec to other engineered systems.

2 Background and Related Works

2.1 Background

Real Logic and Logic Tensor Networks. LTN [19] implement Real Logic, a first-order language whose symbols are interpreted on data by a parametric grounding $\mathcal{G}_\Theta$: constants are mapped to tensors of real features, variables to sequences of individuals, and predicates to differentiable functions into $[0, 1]$. Fuzzy connectives and aggregators make the truth of every formula of a knowledge base $\mathcal{K}$ differentiable, so that Θ is learnable from data. In NeSySpec an individual is a declared record with its availability mask and an optional image, and the predicates are graded physical predicates and visual attribute heads (Sections 3.2 and 3.3). From the product configuration [19, 21], NeSySpec uses the negation $\neg a = 1 - a$, the product t-norm $a \wedge b = ab$ and the Reichenbach implication $a \to b = I_R(a, b) = 1 - a + ab$. The abovementioned implication passes its consequent the gradient $\partial I_R / \partial b = a$, where the residuated Goguen and Łukasiewicz implications pass none wherever $a \leq b$ [21], so the localization pressure of the grounding axiom $P_k \to \mathrm{Loc}_k$ on a visual head grows with the truth of its own assertion (Section 3.3). The universal quantifier is interpreted by the generalized mean error of the truth values $a_1, \ldots, a_n$ that a formula takes on n individuals [19],

$$\mathrm{pME}_p(a_1, \ldots, a_n) = 1 - \Big(\frac{1}{n}\sum_{i=1}^{n}(1 - a_i)^p\Big)^{1/p}, \qquad p \geq 1, \tag{1}$$

which is the arithmetic mean for $p = 1$ and tends to the minimum as $p \to \infty$. Guarded quantification, $\forall x\,[g(x)]\colon \varphi(x)$, aggregates only over the individuals that satisfy a crisp condition g and passes no gradient to the others [19].

Learning from a knowledge base. Learning maximizes the satisfaction of the closed formulas φ of $\mathcal{K}$, aggregated by an operator such as pME_p into SatAgg [19], that is, it minimizes

$$\mathcal{L}(\Theta) = 1 - \mathrm{SatAgg}_{\varphi \in \mathcal{K}}\, \mathcal{G}_\Theta(\varphi). \tag{2}$$

Labels enter as data axioms over the positive and negative examples of a predicate [19], whose satisfaction, aggregated with the log-product that van Krieken et al. [21] recommend, is the negative binary cross-entropy (BCE) up to normalization. In NeSySpec the verdict of a record is a masked aggregation of its axiom truths that applies this guard per axiom, with availability as the guard (Section 3.5), and both objectives are BCEs of data axioms, the visual one adding the loss of Eq. (2) over the grounding axioms (Sections 3.2 and 3.3).

Real Logic for specification auditing. Our formulation of the audit in *Real Logic* has four properties. First, the bands and slopes of graded predicates are learned through their truth values, jointly with the record verdict. Second, guarded quantification uses evidence availability as its guard: constraints whose required evidence is missing are excluded, and the verdict reduces exactly to that of the available constraints (Proposition 1). Third, physical predicates and visual heads share a common truth scale, allowing cross-modal axioms to enter the same knowledge base. Finally, the grounding axiom uses the Reichenbach implication to scale the localization pressure with the truth of the visual predicate.

In the specification stratum, learned graded predicates are aggregated by a masked mean, with no logical connectives between them. Each axiom therefore retains its own truth value at inference, and the record verdict decomposes into constraint-level verdicts. Crisp rule validation also provides exact abstention (Table 2), and Section 4.8 examines the contribution of each of the four properties.

2.2 Related Works

Specification auditing draws on several lines of research, from detecting and repairing errors in structured records to combining domain knowledge with visual evidence. We review work on data editing and diagnosis, learned anomaly detection, differentiable logic and neuro-symbolic reasoning, explanation supervision, and multimodal verification.

Reverse engineering, specification records and data editing. Physics-driven correction flags machine-learning estimates of design parameters, such as those of an aero engine, whose accumulated physical error exceeds a threshold, and corrects them by surrogate-assisted optimization against physical evaluations [22]; NeSySpec works from fleet data without a simulator and names the constraint that fails. Fellegi–Holt editing localizes errors by finding the minimum-weight set of fields whose modification satisfies every edit [13]. Data validation leaves a rule with a missing value unevaluated [23]. Data repair methods, benchmarked by Ni et al. [24], correct cells through probabilistic inference under integrity constraints in HoloClean [25] or using a few example repairs in Baran [26]. The satisfiability modulo theories (SMT) solver Z3 [27] returns conflict sets for consistency-based diagnosis [14] as unsatisfiable cores. It also supports minimum-cardinality localization through MaxSMT [28]. Rule engines and language models check building designs against codes [29, 30]. Their bounds, where present, are crisp, set by experts or mined from data [31], and they take no image evidence; NeSySpec grades every constraint against bands learned from the fleet, ranks the violated constraints by the graded degrees of its own verdict, localizes the altered field by abduction with no oracle to break ties and adds image evidence (Sections 4.2 and 4.3). A reference architecture for auditing mechanical engineering drawings checks the entities read from them against a standards knowledge graph [32]. Hybrid fault diagnosis pairs model-based residuals with anomaly classifiers [33], or residuals learned from fault-free data with consistency-based isolation [34], and a neuro-symbolic architecture lets knowledge-based and learned methods support each other in anomaly detection and fault diagnosis [35], a division of labor that NeSySpec offers as its optional rank-then-explain alarm, a residual-learner alarm explained by the violated constraint.

Anomaly detection and learned detectors. TabPFN v2 [10] and AnoLLM [11] score records by in-context learning and language-model likelihood, RegValidator checks the detection ideas of a language model with data-validation rules [36], GIDCL detects and corrects errors with a graph network and cleaning rules generated by a language model [37], and ZeroED detects erroneous values without manual labels through language-model reasoning [38]. NeSySpec returns, for each record, which constraints the record respects and which it violates, the altered field, a repair proposal and the evidence coverage, together with a score; Section 4.4 compares the scores in area under the receiver operating characteristic curve - AUROC.

Differentiable logic and neuro-symbolic frameworks. Neuro-symbolic methods use explicit knowledge to guide learning and reasoning. Frameworks such as LTN, DeepProbLog [42], Scallop [43] and NeurASP [44] combine trainable models with logical reasoning. Knowledge can also enter through

Table 2: Comparison with the state of the art.

Work	Violated constraint	Altered field	Repair proposal	Exact abstention	Learned tolerances	Part grounding	Record report
LTN and its applications [18, 19, 39–41]	(✓)	(✓)	✗	(✓)	(✓)	(✓)	✗
DeepProbLog, Scallop, NeurASP, abductive learning[b] [42–45]	(✓)	(✓)	(✓)	(✓)	(✓)	✗	✗
DL2[a], semantic loss, constraint layers[a] [20, 46, 47]	(✓)	✗	(✓)	✗	✗	✗	✗
KENN, LEN, PSL[b] [48–50]	(✓)	(✓)	(✓)	(✓)	(✓)	✗	(✓)
Statistical data editing and validation[a] [13, 23]	✓	✓	✓	✓	(✓)	✗	✓
SMT diagnosis[a], language model with SMT [14, 28, 51]	✓	✓	✓	(✓)	✗	✗	(✓)
HoloClean, Deequ [25, 31]	(✓)	✓	✓	✗	(✓)	✗	(✓)
Learned and language-model data cleaning [26, 37, 38]	(✓)	✓	✓	✗	(✓)	✗	(✓)
Tabular learners and anomaly scores[a] [10, 11]	✗	✗	✗	(✓)	✗	✗	✗
Language models with rules[a] [36, 52]	(✓)	(✓)	(✓)	(✓)	(✓)	✗	(✓)
Vision-language models with rules[a] [53, 54]	(✓)	✗	✗	(✓)	✗	✗	(✓)
Rule checking of building designs [29, 30]	✓	(✓)	✗	✗	✗	✗	✓
Hybrid fault diagnosis [33, 34]	✓	✓	✗	✗	✓	✗	(✓)
Neuro-symbolic visual compliance checking [55, 56]	✓	✓	✗	✗	✗	(✓)	(✓)
Explanation supervision[a] [57–61]	✗	✗	✗	✗	✗	(✓)	✗
Tabular-image fusion[a] [62, 63]	✗	✗	✗	(✓)	✗	✗	✗
NeSySpec (ours)	✔	✔	✔	✔	✔	✔	✔

✓: provided and evaluated in a cited work; (✓): expressible or partially provided; ✗: not reported.
[a] Benchmarked here (Section 4); constraint layers used for repair ordering (Section 4.3).
[b] Expressible capabilities, not evaluated in the cited works or benchmarked here.

training objectives or mechanisms that constrain or refine predictions, as in abductive learning [45], DL2 [20], the semantic loss [46], constraint layers [47, 64] and Knowledge Enhanced Neural Networks (KENN) [48]. Probabilistic soft logic (PSL) infers truth values for open atoms under weighted rules [50]. Logic Explained Networks (LEN) extract rules from a trained network [49]. These frameworks can compute a truth, a probability or a distance to satisfaction per formula, from which a violated constraint can be read by an argmax or argmin, can carry a tolerance band as a neural or weighted predicate, and marginalize, infer or leave free what is unobserved (Table 2, note b). Encoding NeSySpec's knowledge base in DeepProbLog, Scallop or PSL would change its axiomatic and logical structure, with probabilistic facts and proofs in place of graded truths and masked quantification, so that comparison is left to future work (Section 5). NeSySpec joins these per-constraint truths with tolerance bands learned from the fleet, abstention under missing evidence (Proposition 1), a search for the altered field with its repair through the same degrees, part-grounded image evidence and a per-record report, provided and evaluated together, which no cited work does; Sections 4.2 and 4.8 compare it with the DL2 method on the declared fields and with a network trained with the DL2 loss on the same residuals, which returns per-constraint probabilities. LTN has grounded part-of relations, fault signatures, process traces [39,65, 66], fault-diagnosis rules for gearboxes and rotating machinery [40,41] and medical rules with learnable thresholds [18]; an answer set program names the missing, misplaced or wrongly connected component of electrical control panels after a separately trained detector [55], and GuardEn executes compiled safety rules on scene graphs of an image [56]; NeSySpec adds graded constraints against bands learned from the fleet and visual predicates grounded in the parts they describe (Section 4.6).

Explanation supervision and part-level evidence. Explanation-guided learning adds supervision on model explanations [67]. Right for the Right Reasons (RRR) penalizes input gradients in regions annotated as irrelevant [57], while GradMask penalizes them outside the annotated region [58]. RES learns from imprecise human masks [59]. The Energy loss of Rao et al. [60] increases the share of positive attribution inside annotated bounding boxes, and Pukdee et al. [61] learn from explanation constraints satisfied in expectation. NeSySpec expresses part-level supervision as an axiom on each head's class-activation map, gated by the head's assertion and guarded by part visibility. Section 4.6 compares this approach with RRR, GradMask and Pukdee et al. under a common protocol.

Multimodal consistency and verifiers. TIP pre-trains on images and masked tables for incomplete tabular data [62], and AutoGluon-Multimodal (AutoGluon-MM) fuses tables and images with

foundation models [63]. The fused verdict of NeSySpec aggregates its strata without labels and reduces exactly to the available constraints, and Section 4.4 compares it with these fusions; its cross-modal axioms read the attributes predicted by the visual heads, as concept bottleneck models predict a label from predicted concepts [68]. Open vision-language models [53, 69] can judge whether a record and an image agree and a language model [52] whether a record is coherent, and vision-language models have been benchmarked on engineering rules, computer-aided design renderings and drawings, where they still struggle [54]; NeSySpec returns with its judgment a verdict for every constraint (Section 4.4).

Table 2 compares the reviewed approaches on seven outputs needed for specification auditing: identification of the violated constraint, localization of the altered field, a repair proposal, exact abstention under missing evidence, learned tolerances, part-grounded image evidence, and a per-record verification report. To our knowledge, none of the reviewed works provides and evaluates all seven outputs together. Statistical data editing and validation come closest, with crisp bounds set by experts or fitted to data and no image evidence. The neuro-symbolic formulation of NeSySpec uses graded constraint violation degrees for the largest-violation key and for abductive field localization, without an oracle to break ties (Sections 4.2 and 4.3).

3 Proposed Methodology

NeSySpec audits a declared record in one integrated, interpretable NeSy framework: through its logical axioms it gives a held or violated verdict on each physical relation that binds the record's fields and, when an image is paired with it, on whether the image agrees. The framework is organized into seven modules (Figure 1), built on Real Logic (Section 2.1). Guarded quantification allows each axiom to contribute only when its required evidence is available (E1), while graded predicates express how well each constraint is satisfied (E2). A grounding axiom ties visual predictions to the relevant image parts (E3), and stratification preserves a separate verdict for each evidence source (E4). A monotone second head complements the main verdict with an additional score and diagnosis key (E5).

- **Data and Verification Instances** encode records and images, then generate corrupted examples and controls;
- **Knowledge Base and Graded Predicates** grade each constraint by a violation degree, with bands and slopes learned from the fleet (E2);
- **Vision Module and Grounding Axiom** encourage each head to place its evidence on its part (E3);
- **Cross-Modal Axioms** compare visual posteriors with declared fields;
- **Masked Stratified Aggregation** guards each axiom by its availability, so the verdict reduces exactly to the available evidence (E1), with one stratum per evidence source (E4);
- **Verdict and Diagnosis Keys** return the verdict and name the violated constraint, with a monotone second head (E5);
- **Verification Report and Optional Checks** give the altered field, a repair proposal and the report, with three optional settings of the audit.

The verdict is the masked mean of the constraint degrees, so it decomposes into constraint-level verdicts, from which the violated constraint, the altered field, the repair, the accounting of missing evidence and the report follow; detection is one of these outputs. Algorithm E.1 in Appendix E transcribes the forward pass from the code. In our framework, we used the LTNtorch library [70]. Appendix B describes the organization of the implementation and gives an example of a verification report.

Domain-specific and domain-independent parts. The domain enters NeSySpec through the knowledge base, the visual predicates, the corruption generator and the data; masked aggregation, graded predicates, grounding axiom, keys, report and protocol are written to apply unchanged to

engineered systems whose declared fields obey physical, definitional or regulatory relations with statable tolerances; fitted parameters and numbers concern aircraft. Appendix A lists what a port to another engineered system requires.

3.1 Data and Verification Instances

Data acquisition. A *specification record* (Appendix D) for each of 107 aircraft types of 71 families in 7 classes declares 16 numeric fields in SI units: maximum take-off and operating empty mass $MTOW$ and OEW, wingspan b, wing area S, aspect ratio AR, fuselage length L_{fus}, cruise speed V_{cr} and altitude h_{cr}, range R, ceiling h_{ceil}, total static thrust F_{tot} or shaft power P_{tot}, engine count n_{eng}, usable fuel mass m_{fuel}, stall speed V_{s} and passenger limit n_{pax}, with the class c and the engine type e. *Ground photographs* of FGVC-Aircraft [7], 6,800 paired with the record of their type or family (Appendix D.1), carry variant, family and manufacturer labels, from which per-variant attribute labels are assigned. The real release of RarePlanes [8] annotates 14,707 aircraft in *satellite* tiles; a diamond through nose, wingtips and tail and the attributes form a record of 9 numeric fields: wingspan b, length L, diamond area $A_{\diamond}$, n_{eng}, vertical-stabilizer count n_{fin}, and codes for the airplane design group of the Federal Aviation Administration (FAA)[1], canards, wing type t and wing position, with the role r and the propulsion π.

Encoding. In an instance (x, m, I), the record $x \in (\mathbb{R} \cup \{\varnothing\})^{n_x}$ holds $n_x = 16$ or 9 declared numeric fields and the categorical attributes, $\varnothing$ marking an undeclared field; $m_j = \mathbf{1}[x_j \neq \varnothing]$, and $I = \varnothing$ without an image. Axiom k is available, $\alpha_k(x, I) = 1$, when its fields are declared and the record falls within its scope. In Figure 2(a), a wide-body turbofan record, outside the scope of the propeller axiom A4P, declares no stall speed, so A6 is masked ($\alpha_k = 0$ for both), and an undetected landing gear masks the main-gear axiom AX3; panel (c) is the 23-field frame of the cross-dataset transfer (Appendix D.1).

Corruption generator and labels. The specification benchmark splits the records by family: the business-jet families form the unseen class, six random families the unseen families, and the others fill training, validation and test at 60/20/20 (Appendix D.1). Each clean record yields up to two instances of each of five corruption families: an impossible single value and geometric, performance and cruise inconsistencies multiply or divide one field by $1 + \epsilon$, with $\epsilon = 0.05$ and 0.10, and a cross-class block swap (Frankenstein) copies a coherent block of fields from another class in the same split. The generator oracle, a crisp check of the same relations against class bands [0.88 min, 1.12 max] of each quantity over all 107 clean aircraft of every split, keeps an instance only when it flags the target constraint (any constraint for a swap); otherwise ϵ doubles, for at most six attempts (Table D.1; realized instances in Figure 2(d)). The verifier is fitted on training-split instances only: the oracle's bands reach it through their labels and per-constraint targets, which thus carry the class extremes of all 107 aircraft, and no field value of another split enters training. Two noisy-but-coherent controls per record perturb each continuous field by a Gaussian factor with relative standard deviation 0.02, with AR rescaled according to the induced change in b^2/S; both controls must remain coherent. The *intervention benchmark* (2,582 rows, same aircraft and splits) draws the same operators once per magnitude, without escalation or oracle certification, and labels each row from the edit alone. A row is labeled incoherent when a continuous field differs from its clean source by more than three noise standard deviations or an integer field changes. These tolerances are fixed before generation, and the label applies regardless of whether a physical relation is violated. An available constraint receives a positive label when its scope includes such an altered field. Noisy copies and edits within tolerance count as unaltered. These labels evaluate the constraint verdicts, the diagnosis key, abduction and detection. The satellite and multimodal benchmarks likewise include corrupted examples and controls. In the multimodal benchmark, each image is paired with its matching record, a mismatched record, or an altered version of its matching record. Mismatches involve aircraft from another class or role, or a sister variant. Occluded and degraded photographs serve as controls..

[1] FAA Advisory Circular 150/5300-13B, Airport Design, March 31, 2022.

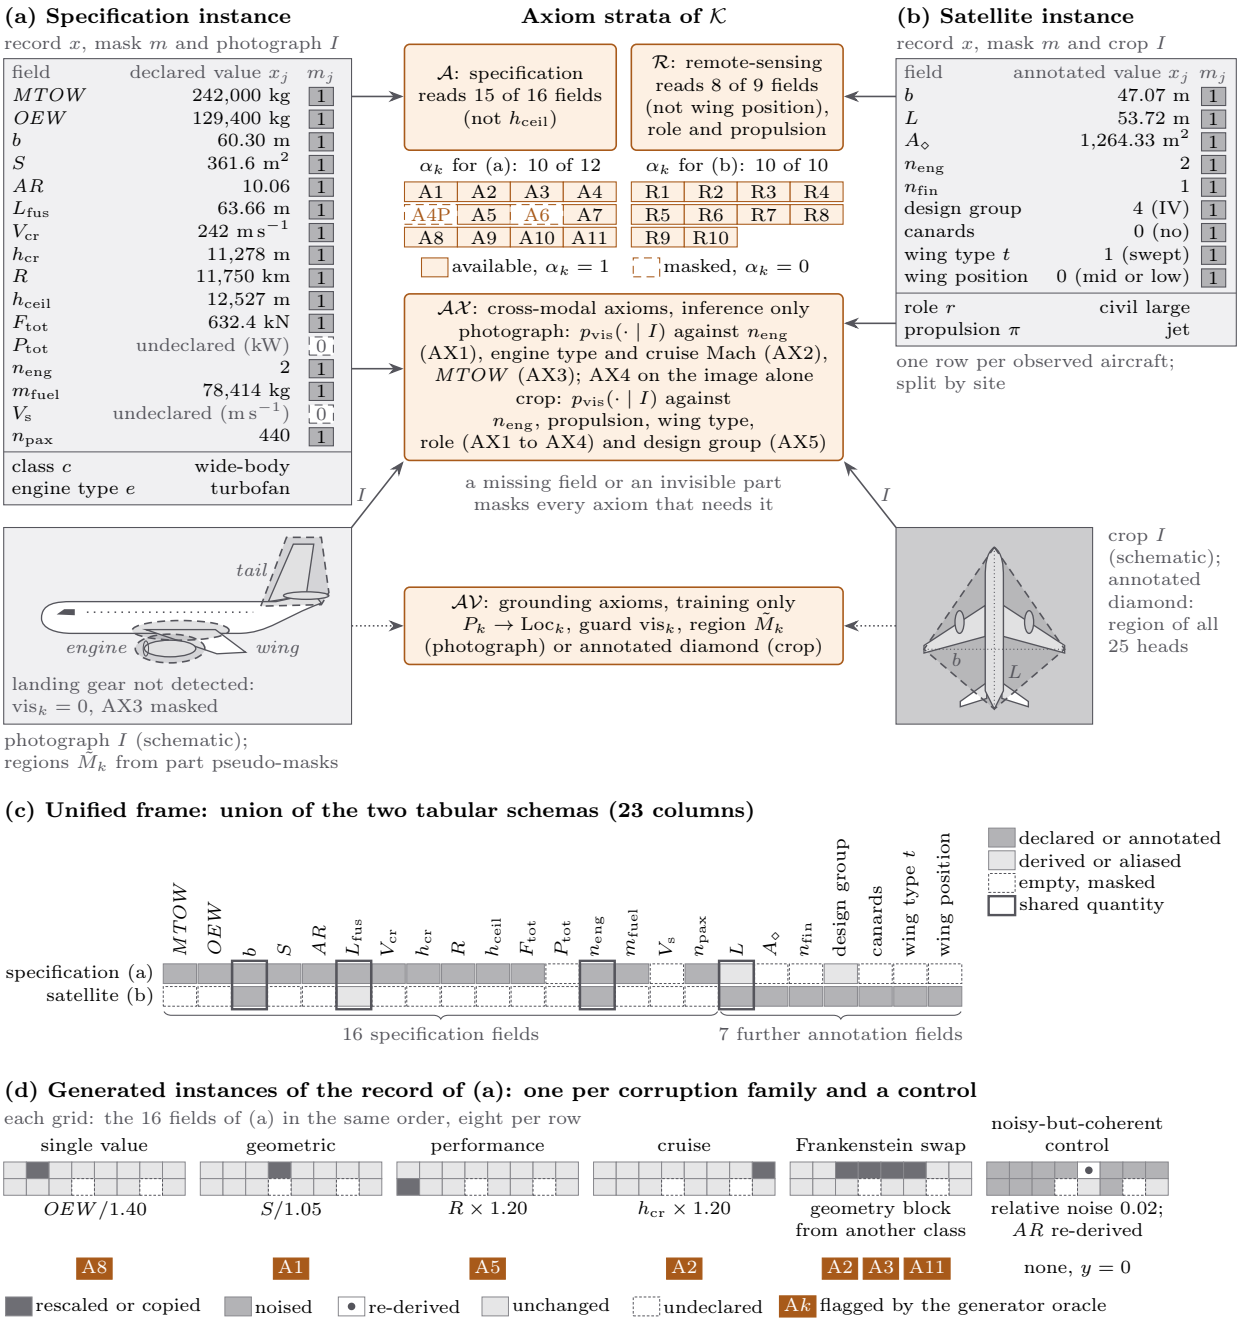

Figure 2: Data and verification instances.

3.2 Knowledge Base and Graded Predicates

The knowledge base $\mathcal{K}$ has four strata (Table 3; visual part in Table E.1): the specification axioms $\mathcal{A}$, 12 relations on records; the remote-sensing axioms $\mathcal{R}$, 10 relations on annotations; the grounding axioms $\mathcal{AV}$, in training only; and the cross-modal axioms $\mathcal{AX}$, at inference only. The tabular strata contain no neural network (Figure 3): each axiom computes an exact, differentiable quantity q of the declared fields, with atmospheric properties from the International Standard Atmosphere (ISA)[2] (Appendix C), reads it in units of the scatter of clean training records, $z = (\log_{10} q - m_g)/\sigma_g$ for a band and $z = \xi/\sigma_g$ for the residual ξ of an identity or the mass-balance margin, and grades it with a residual-aware predicate (E2),

$$v = \mathrm{softplus}\left(\kappa(\ell_g - z)\right) + \mathrm{softplus}\left(\kappa(z - h_g)\right) + \beta\left(\sqrt{1 + (z - c_g)^2} - 1\right), \qquad \mathrm{sat} = \exp(-v), \tag{3}$$

[2]ISO 2533:1975, Standard Atmosphere.

Table 3: Axiom catalog.

ID	Graded predicate and quantity	Fields read	Available ($\alpha_k = 1$) when	Learned parameters
Specification axioms $\mathcal{A}$ (records; 336 trainable scalars, 126 of them latents; $w_k = 1$)				
A1	graded identity of $(AR - \hat{A})/\hat{A}$, $\hat{A} = b^2/S$	b, S, AR	all declared	tolerance (8); κ, β (2)
A2	graded band of $C_L = 2MTOW\,g/\big(\rho(h_{\mathrm{cr}})V_{\mathrm{cr}}^2 S\big)$	$MTOW$, V_{cr}, h_{cr}, S	all declared	band (16); κ, β (2)
A3	graded band of $MTOW/S$	$MTOW$, S	all declared	band (16); κ, β (2)
A4	graded band of $F_{\mathrm{tot}}/(MTOW\,g)$	F_{tot}, $MTOW$	all declared; turbofan	band (16); κ, β (2)
A4P	graded band of $MTOW/P_{\mathrm{tot}}$	P_{tot}, $MTOW$	all declared; propeller	band (16); κ, β (2)
A5	graded identity of $\mathrm{clip}_3 \ln(R/\hat{R})$, $\hat{R} = \frac{V_{\mathrm{cr}}}{g\,c_T}\frac{L}{D}\ln\frac{MTOW}{W_1}$ for a turbofan, $\hat{R} = \frac{\eta_p}{g\,c_P}\frac{L}{D}\ln\frac{MTOW}{W_1}$ for a propeller aircraft; $W_1 = \max(MTOW - 0.94\,m_{\mathrm{fuel}},\, 0.05\,MTOW)$	R, m_{fuel}, $MTOW$; V_{cr} (turbofan)	R, V_{cr}, m_{fuel}, $MTOW$ and AR declared	tolerance (8); κ, β (2); L/D, c_T, c_P, η_p per pair (84)
A6	graded identity of $\mathrm{clip}_2 \ln(V_{\mathrm{s}}/\hat{V}_{\mathrm{s}})$, $\hat{V}_{\mathrm{s}} = \sqrt{2MTOW\,g/(\rho_0 S C_{L,\max})}$	V_{s}, $MTOW$, S	all declared	tolerance (8); κ, β (2); $C_{L,\max}$ per pair (21)
A7	lower hinge of the margin $(MTOW - OEW - f_u\,m_{\mathrm{fuel}} - 95\,n_{\mathrm{pax}})/MTOW$	OEW, m_{fuel}, n_{pax}, $MTOW$	all declared	edge offsets (7); κ (1); f_u per pair (21)
A8	graded band of $OEW/MTOW$	OEW, $MTOW$	all declared	band (16); κ, β (2)
A9	graded band of $F_{\mathrm{tot}}/n_{\mathrm{eng}}$ if F_{tot} is declared, else of $P_{\mathrm{tot}}/n_{\mathrm{eng}}$	n_{eng}; F_{tot} or P_{tot}	n_{eng}, and F_{tot} or P_{tot}, declared	2 bands (32); κ, β (4)
A10	graded band of $M = V_{\mathrm{cr}}/a(h_{\mathrm{cr}})$, per engine type e	V_{cr}, h_{cr}	all declared	band per e (8); κ, β (2)
A11	sum of the graded bands of L_{fus}/b and $n_{\mathrm{pax}}/L_{\mathrm{fus}}$	L_{fus}, b, n_{pax}	all declared	2 bands (32); κ, β (4)
Remote-sensing axioms $\mathcal{R}$ (annotations; 57 trainable scalars; $w_k = 1$ except $w_{\mathrm{R9}} = 0.5$)				
R1	graded band of b for the declared design group, edges fixed by regulation	b, design group	all declared	κ, β (2)
R2	graded band of b per role r; civil bands fixed by definition, military bands learned	b, r	all declared	military bands (8); κ, β (4)
R3	graded band of L/b per r	L, b, r	all declared	band per r (14); κ, β (2)
R4	graded band of n_{eng} per r	n_{eng}, r	all declared	band per r (14); κ, β (2)
R5	$-\ln$ of $[\pi = \text{unpowered}]\,\zeta + [\pi \neq \text{unpowered}]\,(1-\zeta)$, $\zeta = \exp\big(-(n_{\mathrm{eng}}/\tau)^2\big)$	n_{eng}, π	all declared	τ (1)
R6	$-\ln$ of p_r, $1 - p_r$ or $\frac{1}{2}$ for a jet, propeller or unpowered π; $p_r = \sigma(\theta_r)$	π, r	all declared	θ_r per r (7)
R7	$-\ln$ of $I_R\big([t \in \{\text{delta, variable-swept}\}], [r \text{ military}]\big)$	t, r	all declared	none
R8	$-\ln$ of $I_R\big([\text{canards}], [r \in \{\text{bomber, fighter}\}]\big)$	canards, r	all declared	none
R9	$-\ln$ of $I_R\big([n_{\mathrm{fin}} \geq 2], [r = \text{fighter}]\big)$	n_{fin}, r	all declared	none
R10	graded identity of $\ln A_\diamond - \ln(bL/2)$	$A_\diamond$, b, L	all declared	tolerance (1); κ, β (2)

Field symbols: Section 3.1; latents and graded predicates: Section 3.2, Eq. (3). Formulas in SI units. Graded band: two edge softplus hinges plus the in-band term; identity: same, with symmetric tolerance around zero; lower hinge: lower edge only. Degrees capped at 10^4; logical axioms use $-\log\max(t, 10^{-6})$, with truth t. Specification banks: fleet value plus class offset scaled by $n_g/(n_g+5)$ (engine type for A10); remote-sensing banks: flat. $\rho(h), a(h)$: International Standard Atmosphere density and sound speed; ρ_0: sea-level density; g: standard gravity; pair: (class, engine type); $[\cdot]$: indicator; clip_a: clipping to $[-a, a]$; propeller: turboprop or piston. Initial values and bounds: Appendix C..

a violation degree v and its truth. The mean m_g and scatter $\sigma_g \geq 10^{-3}$ of group g (class; engine type for A10; fleet for A1 and A2) come from its n_g clean training rows (132 for records: 44 clean aircraft and 88 noisy copies), shrunk toward the fleet with weight $n_g/(n_g+5)$, σ_g being a root-mean-square around 0 for an identity or the margin; $[\ell_g, h_g]$ is the band, $[-t_g, t_g]$ for an identity, with center c_g; each derived quantity has a slope κ and a weight $\beta \geq 0$ of the in-band term, which grades a value inside the band by its distance from the center. The mass balance A7 keeps its lower hinge only, A9 grades thrust or power per engine as declared, A11 sums its two quantities, every degree is capped at 10^4, and $\mathrm{sat}_k = \exp(-v_k)$ is never clamped. Bounded latents supply quantities that records rarely declare: cruise lift-to-drag ratio L/D, thrust- and power-specific fuel consumptions c_T and c_P, propeller efficiency η_p, maximum lift coefficient $C_{L,\max}$ and fuel utilization f_u, each a sigmoid of one parameter per pair of class c and engine type e (21 pairs), rescaled to conceptual-design ranges [71]; shared by a class and engine type, they check a record against its class.

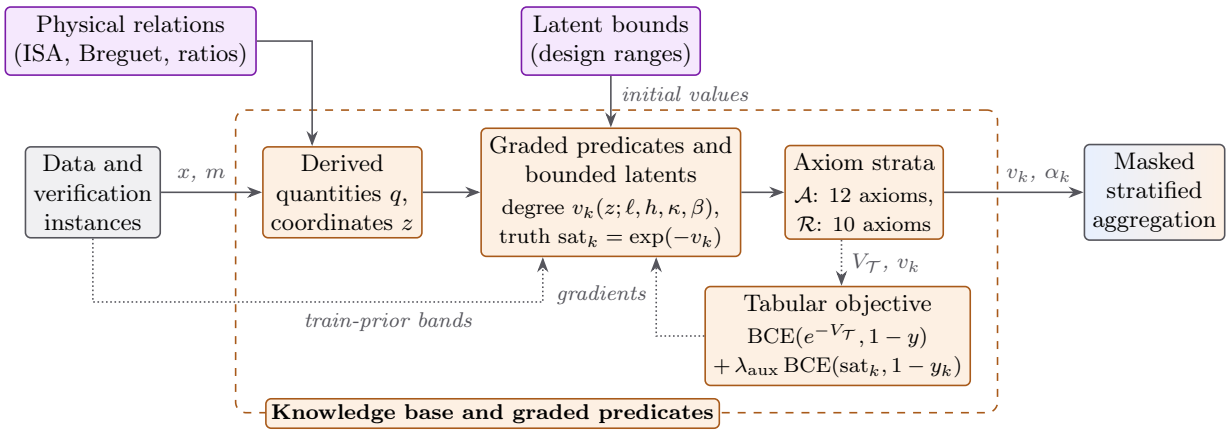

Figure 3: Knowledge base and graded predicates.

The train prior fits every data-dependent band on the 44 clean training aircraft and widens it to [0.88 min, 1.12 max], as the oracle does. The bands start from the train prior and are learned with their slopes from the fleet and its labeled corruptions, so they act as learned tolerances: 45.7% of them keep both edges inside the training fleet range widened in the same way, and the others, all wing-loading bands among them, extend beyond it after training (Appendix G.9). The learned slopes act as implicit per-constraint weights. In hierarchical banks each band or tolerance parameter of a class is a fleet value plus a class offset with the same weight, so a class without training aircraft, such as the business jets, is read with fleet statistics and parameters. The remote-sensing axioms follow the same recipe: six graded slots (R1, R2 for civil and for military roles, R3, R4 and R10) read the residual in standard units of the clean training rows, and the logical axioms R5 to R9 enter as the negative logarithm of their truth. Each tabular stratum's settings (the power p of a masked power mean of the degrees, the in-band term, flat or hierarchical banks) are chosen on validation alone among eight candidates trained with eight seeds each; both strata select $p = 1$, the masked mean of Eq. (9), and on generator-oracle labels the specification stratum keeps the in-band term and hierarchical banks.

Training objective. With instance labels y_i, $y = 1$ for incoherence, and per-constraint targets y_{ik}, the generator oracle's verdict on constraint k for a corruption and zero otherwise, the tabular stratum $\mathcal{T} \in \{\mathcal{A}, \mathcal{R}\}$ minimizes over a mini-batch $\mathcal{B}$

$$\mathcal{L}_{\mathrm{tab}} = \frac{1}{|\mathcal{B}|}\sum_{i\in\mathcal{B}} \mathrm{BCE}\big(e^{-V_\mathcal{T}(x_i)},\, 1-y_i\big) + \lambda_{\mathrm{aux}}\,\big\langle \mathrm{BCE}\big(\mathrm{sat}_k(x_i),\, 1-y_{ik}\big)\big\rangle_{\alpha_{ik}}, \qquad \langle q\rangle_\chi = \frac{\sum_j \chi_j\, q_j}{\max\big(1, \sum_j \chi_j\big)}, \tag{4}$$

with $V_\mathcal{T}$ of Eq. (9), both cross-entropies computed in log space from the unclamped degrees, $\lambda_{\mathrm{aux}} = 0.5$ and $\langle q\rangle_\chi$ the mean of q where $\chi_j = 1$. The residual learners of the comparison receive no per-constraint target. AdamW [72] trains bands, slopes and latents for at most 100 epochs, with early stopping on validation AUROC (Table F.2).

3.3 Vision Module and Grounding Axiom

For each image, the vision module predicts the attributes that the cross-modal axioms compare with the declared record. To let an analyst inspect the evidence behind these predictions, a grounding axiom encourages each head to concentrate its evidence on the part it describes, and the report displays the corresponding activation map (Section 3.4).

The module uses a backbone with one sigmoid head per class of each visual predicate: 25 heads covering eight predicates for photographs, and 25 covering seven predicates for satellite crops (Figure E.2 in Appendix E). On photographs a ResNet-50 [73] pretrained on ImageNet [74] is read at its third stage, a 28×28 grid at 448 pixels. Head k returns $P_k(I) = \sigma(s_k(I))$, with s_k linear in the

spatially averaged feature maps, and its class activation map (CAM) [75] normalized to unit mass,

$$H_k(I) = \mathrm{ReLU}\Big(\textstyle\sum_c w_{kc}\, A_c(I)\Big), \qquad \hat{H}_k(I) = \frac{H_k(I)}{\sum_{u,v} H_k(I)[u,v]+\varepsilon}, \tag{5}$$

where $A_c(I)$ is channel c of the feature map, w_{kc} the weight of head k on it, $[u, v]$ a grid cell and $\varepsilon = 10^{-6}$.

On photographs the region $\tilde{M}_k(I)$ of head k is the union of the automatic pseudo-masks of its parts, OWLv2 [76] boxes of confidence at least 0.15 segmented by the Segment Anything Model (SAM) [77], dilated by 2 native pixels and averaged over the grid cells, none annotated or verified by hand. The head is visible, $\mathrm{vis}_k(I) = 1$, when a part is detected with a confidence of at least 0.20 (Appendix E). On satellite crops every region is the annotated diamond. The in-mask share ρ_k and the localization predicate Loc_k,

$$\rho_k(I) = \frac{\sum_{u,v} \hat{H}_k(I)[u,v]\, \tilde{M}_k(I)[u,v]}{\sum_{u,v} \hat{H}_k(I)[u,v] + \varepsilon}, \qquad \mathrm{Loc}_k(I) = \sigma\big(\kappa_k(\rho_k(I) - \ell_k)\big), \quad \ell_k = \sigma(\tilde{\ell}_k), \quad \kappa_k = \exp(\tilde{\kappa}_k), \tag{6}$$

read "at least a fraction ℓ_k of the attribution lies in the region", with a learned threshold ℓ_k and slope κ_k per head. The grounding axiom of head k (E3) is the guarded Reichenbach implication

$$\mathcal{AV}_k : \ \forall I\,[\mathrm{vis}_k(I)] : \ P_k(I) \to \mathrm{Loc}_k(I), \qquad \mathrm{sat}_{\mathcal{AV}_k}(I) = I_R\big(P_k(I), \mathrm{Loc}_k(I)\big) = 1 - P_k(1 - \mathrm{Loc}_k). \tag{7}$$

Its *gate* P_k asks a head that asserts its class to localize its evidence, with no label, and its *guard* masks a conjunct whose part is not visible.

Training objective. With per-head labels t_{ik}, the one-hot attribute labels, and their availability ω_{ik}, the vision module minimizes

$$\begin{aligned} \mathcal{L}_{\mathrm{vis}} &= \big\langle \mathrm{BCE}\big(P_k(I_i),\, t_{ik}\big)\big\rangle_{\omega_{ik}} + \lambda_{\mathcal{AV}}\Big(1 - \big\langle \mathrm{sat}_{\mathcal{AV}_k}(I_i)\big\rangle_{\mathrm{vis}_k(I_i)}\Big) + \lambda_{\mathrm{fl}}\,\Phi, \\ \Phi &= \Big\langle \mathrm{sg}\,\big[P_k(I_i)\big]\, \max\big(0, 1 - S_k(I_i)/S_0\big)\Big\rangle, \end{aligned} \tag{8}$$

with $\lambda_{\mathcal{AV}} = 1$ and $\lambda_{\mathrm{fl}} = 0.1$. The *floor* Φ (S_k the total activation of H_k, $S_0 = 1$, sg a stop-gradient) stops a head from asserting its class through its bias with a negligible map; the cross-modal axioms read the frozen posteriors (Table F.2). Explanation supervision [57,59,61] suggests instead an auxiliary loss, the Energy loss [60], here the mean of $1 - \rho_k$ over the visible label-positive heads, with about 4.3 times the axiom's localization weight per conjunct at the same loss weight (Section 4.6).

We train and compare multiple visual backbones. Satellite experiments use ResNet-50, ConvNeXt-T and ViT-B/16, while photographic grounding is also evaluated across multiple backbones (Table F.2; Appendix G.6).

3.4 Cross-Modal Axioms

The cross-modal axioms (Table E.1; center of Figure 2) compare declared attributes with the visual posteriors $p_{\mathrm{vis}}(\cdot \mid I)$ predicted from the image. For each predicate, the posterior is obtained by a softmax over its heads, or a sigmoid for a binary predicate. Axiom truths are computed from these posteriors, expected truths of implications, or products of such terms, without additional trainable parameters. An axiom is available when its declared value is covered by the corresponding visual head and the relevant parts are visible.

The report displays each grounded head's activation map and in-mask share ρ_k alongside the cross-modal verdicts (Figure E.1), allowing an analyst to inspect which image regions support a reported contradiction.

3.5 Masked Stratified Aggregation

A tabular stratum $\mathcal{T} \in \{\mathcal{A}, \mathcal{R}\}$ averages the degrees of its available axioms (E1), and the cross-modal stratum aggregates their truths by the masked p-mean error,

$$V_{\mathcal{T}}(x) = \frac{\sum_{k\in\mathcal{T}} w_k\,\alpha_k\,v_k}{\sum_{k\in\mathcal{T}} w_k\,\alpha_k}, \qquad \mathrm{feas}_{\mathcal{AX}}(x,I) = 1 - \left(\frac{\sum_{k\in\mathcal{AX}} w_k\,\alpha_k\,(1-\mathrm{sat}_k)^4}{\sum_{k\in\mathcal{AX}} w_k\,\alpha_k}\right)^{1/4}, \tag{9}$$

each where its denominator is positive, with $w_k = 1$ except 0.5 for R9 and the satellite AX5. The feasibility $\mathrm{feas}_{\mathcal{T}} = \exp(-V_{\mathcal{T}})$ is the weighted geometric mean of the available truths, and $V_{\mathcal{T}} = \ln 2$ (feasibility 0.5) without available axioms; with unit weights $\mathrm{feas}_{\mathcal{AX}}$ is the masked aggregator of LT-Ntorch, and without available axioms the cross-modal stratum drops out. The strata present combine by a weighted mean (E4),

$$\mathrm{feas}(x,I) = \frac{\lambda_{\mathcal{T}}\,\mathrm{feas}_{\mathcal{T}}(x) + \lambda_{\mathcal{AX}}\,\gamma_{\mathcal{AX}}(x,I)\,\mathrm{feas}_{\mathcal{AX}}(x,I)}{\lambda_{\mathcal{T}} + \lambda_{\mathcal{AX}}\,\gamma_{\mathcal{AX}}(x,I)}, \qquad \gamma_{\mathcal{AX}} = \Big[\textstyle\sum_{k\in\mathcal{AX}} w_k\alpha_k > 0\Big], \tag{10}$$

with $\lambda_{\mathcal{T}} = \lambda_{\mathcal{AX}} = 1$; when $\gamma_{\mathcal{AX}} = 0$ the implementation returns $\mathrm{feas}_{\mathcal{T}}$ itself. Photographic pairs are scored by the cross-modal stratum alone, which reads no specification truth.

Proposition 1 (Exact reduction) *Let the verdict be computed with Eqs.* (9) *and* (10)*. (a) For every stratum $\mathcal{S}$, instance (x, I) and fixed availability pattern α (which, for cross-modal axioms, depends on the declared value and on visibility), let $\mathcal{S}' = \{k \in \mathcal{S} : \alpha_k(x, I) = 1\}$. If $\mathcal{S}'$ is nonempty, the stratum verdict, $V_{\mathcal{S}}$ or $\mathrm{feas}_{\mathcal{AX}}$, coincides with Eq.* (9) *on $\mathcal{S}'$ with every axiom available; otherwise a tabular stratum takes $V_{\mathcal{S}} = \ln 2$ and the cross-modal stratum is absent ($\gamma_{\mathcal{AX}} = 0$). Either way it depends neither on the degree or truth of a masked axiom nor on a field read only by masked axioms, and has zero derivative with respect to them. (b) With a tabular stratum ($\lambda_{\mathcal{T}} > 0$), if $I = \varnothing$, or every cross-modal axiom is masked, then $\mathrm{feas}(x, I) = \exp\big(-V_{\mathcal{T}}(x)\big)$ exactly; otherwise $\mathrm{feas}(x, I)$ is the verdict of the knowledge base restricted to the strata present (including the tabular stratum at feasibility $1/2$ when none of its axioms is available), with their weights renormalized.*

The proof is in Appendix E.1. The reduction holds because only available axioms enter the mean, a structural property of E1 that also holds in training, where a masked axiom receives zero gradient; its practical value is an exact abstention, which crisp rule validation shares (Table 2) and learned heads do not keep (Section 4.5); stratification (E4) keeps one verdict per evidence source (Section 4.8).

3.6 Verdict and Diagnosis Keys

NeSySpec's verdict is $V = -\ln \mathrm{feas}(x, I)$, the masked mean degree $V_{\mathcal{T}}$ for a record without image. Bands and slopes are learned for each derived quantity, with no additional weights learned for aggregation. The resulting verdict serves as the record score. NeSySpec converts this score into a binary decision using the threshold $t^\star$ that maximizes Youden's index on the validation set [78]. A specification record with no available constraint scores $V = \ln 2$, above the validation thresholds of the specification verifier (0.310 to 0.558 across seeds), so it is flagged for lack of evidence, with an evidence coverage of 0 in its report (Section 3.7). For a flagged specification record with at least one available constraint, *constraint diagnosis* identifies the constraint with the largest violation degree. This gives the *largest-violation key*: the argmax of the degrees entering the verdict, with no additional parameter to fit. We also retain a *secondary key*, identifying the constraint that contributes most to the second head's score. The second head (E5) is a logistic regression with non-negative weights on the degrees $e_k = \alpha_k v_k$ of the available axioms,

$$\hat{p}_{\mathrm{m}}(x, I) = \sigma\Big(\phi_0 + \textstyle\sum_k \big(\phi_k\, e_k + \phi^{\mathrm{a}}_k\, \alpha_k\big) + \phi_{\mathrm{s}}\, V\Big), \qquad \phi_k,\ \phi_{\mathrm{s}} \geq 0, \tag{11}$$

with availability terms ϕ^{a}_k free in sign, fitted by L_2-regularized maximum likelihood with a bound-constrained quasi-Newton method [79] and selected on validation AUROC; on satellite pairs it reads $V_{\mathcal{R}}$ in place of V and also the cross-modal truths, with non-positive weights, and the cross-modal score

$1 - \text{feas}_{\mathcal{AX}}$ (0 without an image), with a non-negative weight. The contribution $c_k = \text{logit}\,\hat{p}_{\text{m}}(x, I) - \text{logit}\,\hat{p}_{\text{m}}(x, I \mid v_k{:=}0)$, with V (or $V_{\mathcal{R}}$) recomputed, gives the secondary key, a learned second ranking of the same degrees that the report prints beside the largest-violation key (Section 4.2, Figure 6b).

3.7 Verification Report and Optional Checks

Every instance receives a verification report (Figure E.1 in Appendix E; example in Figure 5) whose per-constraint degrees state, by one rule, which constraints the record violates and which it respects: an available constraint is *violated* when its truth $\exp(-v_k)$ is below 0.5, that is $v_k > \ln 2$, reached at the learned band edge, where a hinge of Eq. (3) alone equals $\ln 2$, or inside the band where the learned in-band term adds to it (A1, A2); it is *held* otherwise, and a masked constraint is not evaluated. The rule has no tuned parameter: it reads at 0.5 the truths that the per-constraint term of Eq. (4) trains toward 1 on constraints the generator leaves held and toward 0 on those it violates. Its evidence coverage counts the available constraints among those whose scope contains the record, so a masked axiom is an abstention, and the report gives the class of the nearest clean training record and its fleet similarity $\exp(-d/\tau)$, with d the root-mean-square difference from it over the fields both records declare, in log-scaled standardized coordinates, and τ the median leave-one-out nearest-neighbor d in the clean training fleet, so that a value below e^{-1} marks a record farther from the fleet than a typical clean training record.

For a flagged specification record, *field localization* names the altered field by abduction, a search over single-field changes for the one that restores consistency; a field whose change alone restores it is a single-fault diagnosis in the sense of consistency-based diagnosis [14]. For every field f read by an available constraint, the verifier evaluates the verdict when f alone moves by d in log space,

$$\psi_f(d) = V_{\mathcal{A}}\big(x^{(f,d)}\big), \qquad x_f^{(f,d)} = x_f \exp(d), \quad x_j^{(f,d)} = x_j \ (j \neq f), \qquad d \in \{-3, -2.995, \ldots, 3\}, \quad (12)$$

and ranks first the fields whose change alone can bring the verdict to $V \leq -\ln 0.9 \approx 0.1054$, by the smallest such change, then the others by the lowest verdict they reach, a rule and a threshold selected on validation and independent of the diagnosis. *Repair proposals* then solve

$$d^\star = \arg\min_d \ \|d\|_2^2 \quad \text{s.t.} \quad \text{feas}_{\mathcal{A}}\big(x \odot \exp(d)\big) \geq \tau_r, \qquad d_j = 0 \text{ for } j \notin F_{k_r}(x), \quad (13)$$

the smallest log-space change of the k_r top-ranked fields $F_{k_r}(x)$ that brings the verdict to the threshold, a counterfactual over the declared fields [80], by an augmented-Lagrangian method [81] with Adam [82] steps, keeping its best-effort change when the threshold is not reached; validation chose $\tau_r = 0.9$ ($V \leq -\ln 0.9$), a target stricter than the decision threshold, and $k_r = 1$ for the repair error (Section 4.3).

NeSySpec's audit has three optional settings. With *rank-then-explain*, the residual MLP, an MLP on the 14 physical quantities of the oracle (Table F.1), raises the alarm at its validation Youden threshold and the largest-violation key explains each flagged record; the flagged-and-diagnosed rate, the share of corrupted records flagged with a violated top-ranked constraint, scores this setting. The *critical-field alert* flags a record that leaves undeclared a field required by an applicable axiom and declared by at least 95% of the training records of its class - or of its engine type for classes with fewer than five training records -, reading only the undeclared fields, the class, the engine type and a table of the training records; with it enabled, NeSySpec flags a record when its verdict or the alert fires. With *MICE imputation* of the missing fields enabled, NeSySpec scores imputed values in place of abstaining on them; Section 4.5 compares this setting with learners given the same imputer.

4 Evaluation and Results

4.1 Experimental Setup

The vision module and the tabular verifier are fitted separately, as described in Section 3 (Table F.2, Appendix F); where a result uses one of NeSySpec's optional settings, rank-then-explain, the critical-field alert or MICE imputation, the text names it. The configuration was selected on validation data and held-out training classes. The reported results use a visibility threshold calibrated on a subset

Table 4: Benchmarks and labels.

Benchmark	Labels	Independent of the oracle	Sections	Size
Specification records	generator labels: corruptions kept when the oracle flags them, with per-constraint labels; noisy-but-coherent controls; fields hidden at random	no	4.2 to 4.5, 4.8	107 aircraft of 71 families; test: 16 aircraft of 9 families, 204 rows with 32 controls
Intervention-labeled records	from the edit alone, with a tolerance fixed before generation	yes	4.2 to 4.4, 4.8	2,582 rows, 1,810 altered; same aircraft and splits
Falsification	the attacker's edits of genuine records	yes	4.7	736 attack problems on the 46 genuine held-out records
Satellite pairs (RarePlanes)	generator labels; image swaps	no	4.4, 4.5	5,250 pairs; 10 test and 8 unseen sites
Photographic pairs (FGVC-Aircraft)	constructed mismatches, attribute labels, automatic pseudo-masks	yes	4.4, 4.6	4,543 pairs; 800 test photographs
Human part masks	PASCAL-Part [84], PartImageNet [85] and UDA-Part [86] annotations, for evaluation only	yes	4.6	535 aircraft in 506 PASCAL-Part images; 40 UDA-Part items

including test and out-of-distribution images, and generator labels from an oracle fitted on all 107 aircraft (Appendix F). Sensitivity analyses at the end of this section repeat threshold selection and oracle fitting without test or out-of-distribution data.

Benchmarks and labels. Table 4 maps each benchmark to its labels and to the sections that use it; every error in this study is synthetic, an edit or a constructed mismatch of public records and images. Two label families recur. *Generator labels* come from the generator oracle of Section 3.1, which judges the same residual quantities as the knowledge base against class bands fitted on all 107 aircraft; its single-field corruptions are escalated from 5 or 10% until the oracle flags them, mostly to 20 to 160% (geometric ones stay at 5 or 10%; Table D.1). *Intervention labels* cover single-field edits drawn once at 5 to 80% and cross-class block swaps: a field counts as altered beyond three noise standard deviations (0.06 in natural-log units for continuous fields, 0.134 for the aspect ratio, any change of an integer field), so edits within this tolerance, e.g., the 5% edits of continuous fields, count as unaltered, and a constraint is labeled when its scope contains an altered field. Such a constraint is involved in the edit, which need not violate it: an edit can keep the record inside the physical envelope of the fleet, and 613 of the 1,810 altered rows, about a third, break no constraint even under bands fitted on all 107 aircraft, so that no physical check can flag them (Appendix G.4). Generator labels follow the same relations as the axioms and were NeSySpec's per-constraint training targets; intervention labels, set from the edit alone, are independent of them. Minimum-cardinality localization (Fellegi-Holt and MaxSMT [28], all optimal solutions) breaks ties at random unless the oracle tie-break is named.

Statistics. Intervals are 95% two-level cluster-bootstrap percentile intervals [87], reported as [aircraft] · [family] for records. The ± notation denotes standard deviation across seeds. Residual MLP and NeSySpec runs are paired by seed. A method is higher or lower than another, ahead or below, when both intervals of their difference exclude zero, and on par with it when both include zero; when only one excludes zero, when a bound prints 0.000 or lies within 0.003 of zero, or when the reading changes in one of 20 alternative bootstrap streams, the text gives the levels in plain words (close, comparable, slightly). The text gives levels, Tables G.1 to G.4 and Appendix G.1 the differences and intervals, and Appendix G the spread over seeds and the readings across the alternative streams; no multiplicity correction is applied.

Comparators, metrics and cost. Learners on the raw fields or on the residuals, the 14 physical quantities that the oracle checks, score records (the residual MLP, TabPFN v2 [10], LightGBM [88], a one-class support vector machine, SVM); class standardization and the grouped Shapley additive explanations (SHAP) [12] of a second model (MLP-SHAP, LightGBM-SHAP) rank the residuals; Fellegi-Holt, Z3 and rule checkers return crisp violated sets, a network trained with the DL2 loss [20] on the same residuals a probability per constraint, and language and vision-language models given the rules

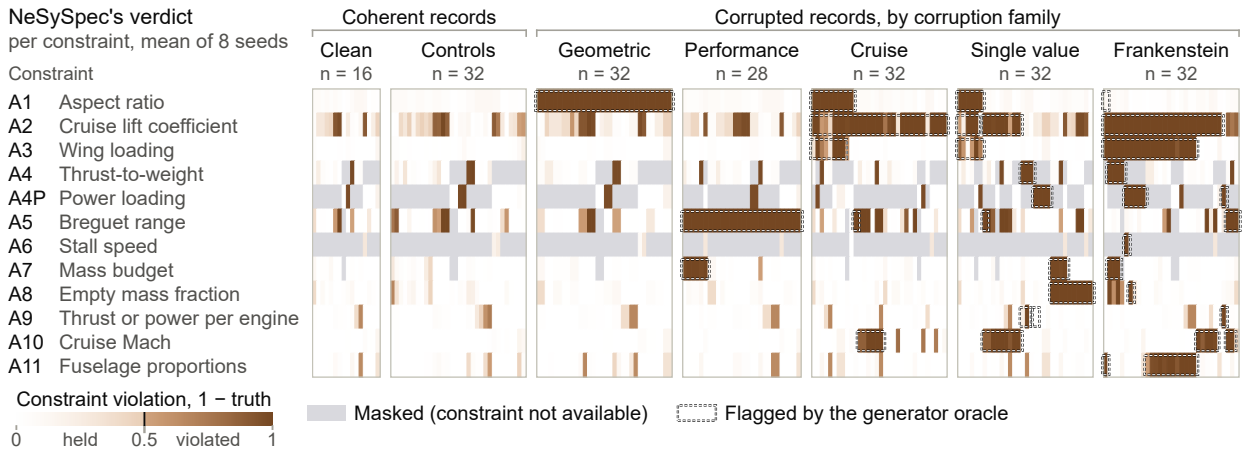

Figure 4: Constraint-level verdicts on test records.

a judgment (Table F.1). We call V the record verdict and the held or violated outputs per constraint the constraint verdicts. The constraint verdicts use exact-set agreement per record, macro-F1 over the constraints with at least five labeled violations and the false-violation rate on clean records. Diagnosis and localization use top-1 rates, repair the log-space error and grounding the pointing game [89]; detection scores the record verdict by AUROC and the false-positive rate (FPR) on the 32 controls at a true-positive rate (TPR) of 0.90. Experiments ran on an AMD Ryzen 9 9950X3D with 192 GB of RAM, an NVIDIA RTX 5090 and an RTX 4090. Per record, the verdict and the key take one pass through the knowledge base, abduction evaluates the verdict at 1,201 log-space steps for each field read by an available constraint, and the repair solves one augmented-Lagrangian problem; training curves are in Appendix G.10.

Figure 4 shows, for each test record, the violation $1 - \mathrm{sat}_k = 1 - \exp(-v_k)$ of each constraint (mean of 8 seeds) with the rule of Section 3.7 drawn as a step at 0.5: dark cells are violated constraints, gray cells masked ones, framed cells the oracle's flags.

The audit at a glance. Table 5 lines up NeSySpec with its comparators on every output of the audit. Learners on the physical residuals return a record score and need a second model, a head per constraint or attributions, for a verdict or a diagnosis; solvers return a violated set, an edit or a repair; language and vision-language models return a judgment. NeSySpec is the only method that returns every output, and without an image its fused verdict alone returns the record verdict exactly. On the test aircraft no comparator is ahead of it in the macro-F1 of the constraint verdicts, in naming the violated constraint and the altered field or in the repair error on the altered field; the network trained with the DL2 loss keeps fewer false violations (Section 4.2). In record and pair detection, learners on the residuals and late fusions are close to NeSySpec and in some settings rank records or pairs higher (Sections 4.4 to 4.5). Learners on the residuals do so on the unseen class, after retraining on intervention labels (Table G.3) and with 10% of the fields hidden, the MLP trained with masks on test and the residual MLP on unseen families; late fusions do so on satellite pairs after learning from labeled pairs, as does Qwen3-VL-32B on photographs of unseen families. With the axiom's 8 seeds, the auxiliary loss points higher on test pseudo-masks (Section 4.6, Appendix F).

Robustness to the label oracle and the visibility threshold. Two further oracles exclude test and out-of-distribution aircraft: the first is fitted on the 44 training aircraft with the paper's tolerances, the second on the 61 training and validation aircraft with tolerances chosen by leaving one family out. On intervention labels, under both oracles, NeSySpec keeps its leads over Fellegi-Holt and Z3 in diagnosis and over minimum-cardinality localization with random tie-break on every split. It keeps its macro-F1 lead over the DL2-loss network on unseen families under both oracles, on test under the first and slightly under the second, and on the unseen class under the first only. Its narrower leads over the oracle tie-break, the crisp edits and the Class I residual become on par or lower on some splits. Its own levels on test are lower (macro-F1 0.488 and 0.537, abduction 0.532 and 0.491, key 0.732 and, slightly,

Table 5: The audit at a glance: NeSySpec and its comparators on each output.

Method	Verdicts macro-F1	Violated constraint	Altered field	Repair error	Exact, no image	Part pointing, PASCAL-Part	Records AUROC	Half fields hidden, unseen class	Satellite pairs
Learners on the physical residuals									
Residual MLP	·	·	·	·	·	·	0.990	0.536	·
TabPFN v2 with native missing values	·	·	·	·	·	·	0.957	0.539	·
LightGBM trained with masks	·	·	·	·	·	·	-	0.566	·
MLP-SHAP	0.693/0.479[a]	0.942	·	·	·	·	0.990	-	·
LightGBM-SHAP	0.595/0.418[a]	0.871	·	·	·	·	0.957	-	·
Network with the DL2 loss, one head per constraint	0.661/0.426[b]	0.752	·	·	·	·	0.965	-	·
Class-standardized residual	0.714/0.480[b]	0.929	0.403/-	·	·	·	-	-	·
Solvers, statistics and imputation									
Fellegi-Holt	0.829/0.505	0.795	-	0.241	·	·	0.856	-	·
Z3	0.829/0.505	0.825	-	-	·	·	0.867	-	·
Minimum-cardinality localization (Fellegi-Holt, MaxSMT), random tie-break	-	-	0.367/0.306	-	·	·	-	-	·
Same localization, oracle tie-break	-	-	0.629/0.453	-	·	·	-	-	·
Class I sizing regressions	·	·	0.597/0.453	0.160	·	·	-	-	·
MICE at the Class I field	·	·	·	0.120	·	·	·	·	·
Constraint-layer ordering (PiShield style)	·	·	·	0.218	·	·	·	·	·
Knowledge given as a loss or in words									
DL2 method on the declared fields	0.425/0.259	0.443	·	·	·	·	0.703	-	·
Language model given the constraints in words	0.153/-	0.128	·	·	·	·	0.620	-	·
Images									
Auxiliary localization loss	·	·	·	·	·	0.739	·	·	·
RRR, GradMask, Pukdee et al.	·	·	·	·	·	lower	·	·	·
End-to-end fusions trained on labeled pairs	·	·	·	·	·	·	·	·	lower[c]
Learned fusions and heads on satellite pairs, NeSySpec's second head included	·	·	·	·	0 of 207	·	·	·	close[d]
Qwen3-VL-8B prompted with the rules	·	·	·	·	·	·	·	·	0.748
NeSySpec (ours)	**0.795/0.575**	**0.942**	**0.742/0.593**	**0.133**	**9 of 9**	**0.737**	**0.979**	**0.577**	**0.994**

Test unless stated otherwise; slash: generator/intervention labels. Diagnosis: top-1 accuracy; repair: log-space error on the altered field (lower is better); detection: AUROC (Z3 score: violated-edit count). Field masking is random; satellite AUROC pools all pair types. Exactness: satellite cells (test and unseen sites) whose no-image verdict matches the record verdict bit for bit. Pointing: 3 runs per method, loss weights 1/0.3 (ours/auxiliary). Fellegi–Holt/Z3 verdicts, MICE/Class I repairs and auxiliary-loss pointing are on par with NeSySpec. ·: output not provided or no image input; -: not reported. [a]Corrupted records only; corresponding NeSySpec macro-F1: 0.832/0.601. [b]Validation-tuned thresholds: network 0.746/0.541; residual 0.560 on intervention labels. [c]ResNet-50 for both; test and unseen sites. [d]Some fusions on knowledge-base truths, trained on labeled pairs, achieve higher AUROC on test swaps, unseen sites or with 30% of images missing (Section 4.4; Appendix G.5). Comparisons and intervals: Sections 4.2 to 4.6; Appendix H.

0.738, against 0.575, 0.593 and 0.786), and its false violations rise only under the paper's tolerances, whose bands flag 38 of the 63 clean held-out aircraft and narrow the bands it learns (Appendix G.9). Re-selected on validation alone, the part detector's visibility threshold becomes 0.25. Retrained at 0.25, the photograph module keeps its comparisons against the networks without a localization term or with permuted masks and against RRR, GradMask and Pukdee et al. on PASCAL-Part, all kept at 0.20, against the auxiliary loss also retrained at 0.25 and against the vision-language models. Its pointing is slightly lower on test pseudo-masks (0.774 against 0.780) and on PASCAL-Part human masks (0.711 against 0.737), where the auxiliary loss falls as well (0.708 against 0.739), and its cross-modal AUROC is slightly lower on unseen families and the unseen class (Appendix F).

4.2 Constraint Verdicts and Diagnosis

For every available constraint of a record, NeSySpec returns a graded degree and, by the fixed rule of Section 3.7, a held or violated constraint verdict (Figures 4 and 5); these verdicts reach a macro-F1

(a) Declared record (test split)

class; engine type	narrow-body; turbofan	*MTOW* (kg)	68,000	wingspan (m)	34.1
engine count	2	*OEW* (kg)	39,500	wing area (m^2)	122.6
total thrust (kN)	**160.17** *(clean 192.2)*	usable fuel (kg)	19,465	aspect ratio	9.48
not declared: shaft power, stall speed		passenger limit	132	fuselage length (m)	31.44

cruise speed (m/s)	230
cruise altitude (m)	11,278
range (km)	5,750
ceiling (m)	12,131

(b) Constraint-level verdicts

constraint		*verdict*	v_k	*degree*
A1	aspect ratio	held	0.082	
A2	lift coefficient	held	0.096	
A3	wing loading	held	0.053	
A4□	thrust to weight	**violated**	10.001	
A4P	power loading	*out of scope*		*propeller axiom*
A5	Breguet range	**violated***	1.053	*class latents*
A6	stall speed	*masked*		*no stall speed*
A7	mass balance	held	0.094	
A8	empty-mass fraction	held	0.386	
A9□	thrust per engine	held	0.384	
A10	cruise Mach number	held	0.005	
A11	fuselage ratios	held	0.013	
				ln 2
verdict V, mean degree			1.217	

decision: flagged (validation threshold 0.310, dashed)
clean record: V 0.191, not flagged; *A5 violated there too
largest-violation key (seed 0): A4, then A5
evidence coverage: 10 of the 11 constraints in scope
nearest clean record: narrow-body, fleet similarity 0.274, below e^{-1}: farther from the fleet than typical
□ oracle flag (benchmark only) ▱ value off scale

(c) Abductive field search and repair

field changed alone	*lowest V*	*log change*
total thrust	0.181	+0.245
usable fuel	1.103	−0.200
range	1.112	+0.200
as declared	1.217	

altered field: total thrust; alone it brings V below the decision threshold (0.181 against 0.310)
violated bands (seed 0): A4 $T/W = 0.240$ *(clean 0.288)*, below the learned band 0.278 to 0.332; A5 range 5,750 km, below the learned edge 5,830 km;
repair of one field, stricter target $V \leq -\ln 0.9 \approx 0.1054$: not reached: A5, which does not read the thrust, and the in-band degrees keep V above it; proposal 204.59 kN *(clean 192.2 kN)*, a value to cross-check

Figure 5: Verification report of one declared record.

of 0.795 on generator-label test and, without retraining, 0.575 on intervention-label test. Against a network trained with the DL2 loss on the same residuals and per-constraint targets, which returns a probability per constraint, they reach a higher macro-F1 on every split and both label families (the network at 0.661 and 0.426 on the two tests; Figure 6a), and the same key read from the network ranks the violated constraint first less often on every generator-label split (0.752 against 0.942). They also find every labeled violation of a corrupted record more often than the network on every split and both label families (0.934 against 0.522 on generator-label test; micro recall 0.959 against 0.635), with micro precision on par on test and lower on the unseen class (0.458 against 0.577 on generator and 0.517 against 0.639 on intervention labels; Appendix G.1). With the network's thresholds tuned on validation the macro-F1 lead holds on generator-label test and is slight on intervention-label test (0.575 against 0.541; Table G.5). The network reports fewer false violations on clean records (0.020 against 0.074 on generator-label test, also on intervention labels), and NeSySpec also returns the altered field and a repair proposal from the same degrees.

Crisp consistency edits on validation-selected bands (Fellegi-Holt and Z3) are on par with the constraint verdicts in macro-F1 on generator labels, and on intervention labels the constraint verdicts lead them on unseen families (0.694 against 0.569; Appendix G.1); the graded degree ranks the violations, and the largest-violation key names the violated constraint first more often than their complete violated set (below). With thresholds from the specification benchmark's validation, the constraint verdicts are ahead of the class-standardized residual, MLP-SHAP and LightGBM-SHAP in macro-F1 on test with both label families; with the residual's thresholds chosen on the intervention benchmark's validation, the residual is on par with them in macro-F1 on intervention-label test (0.560 against 0.575) and below them in exact-set agreement, with more false violations (0.174 against 0.084). The DL2 method on the declared fields and a language model given the constraints in words are below the constraint verdicts in macro-F1 on every split and label family (the language model on the specification benchmark only).

On clean and control test records the constraint verdicts report a violation in 0.074 of the checks (Figure 4), mostly on A2 and A5 (0.255 and 0.208 of their checks), where A2 evaluates the cruise lift coefficient at maximum take-off mass and A5 uses class-level latents; the report shows which relation the declared values strain and by how far (Figure 5c), and the key and abduction rank the degrees without this rule.

The largest degree, the largest-violation key, names the violated constraint in 0.942 of corrupted test records (generator labels; on each of 8 seeds), above Fellegi-Holt, Z3 and rule diagnosis (0.795 to

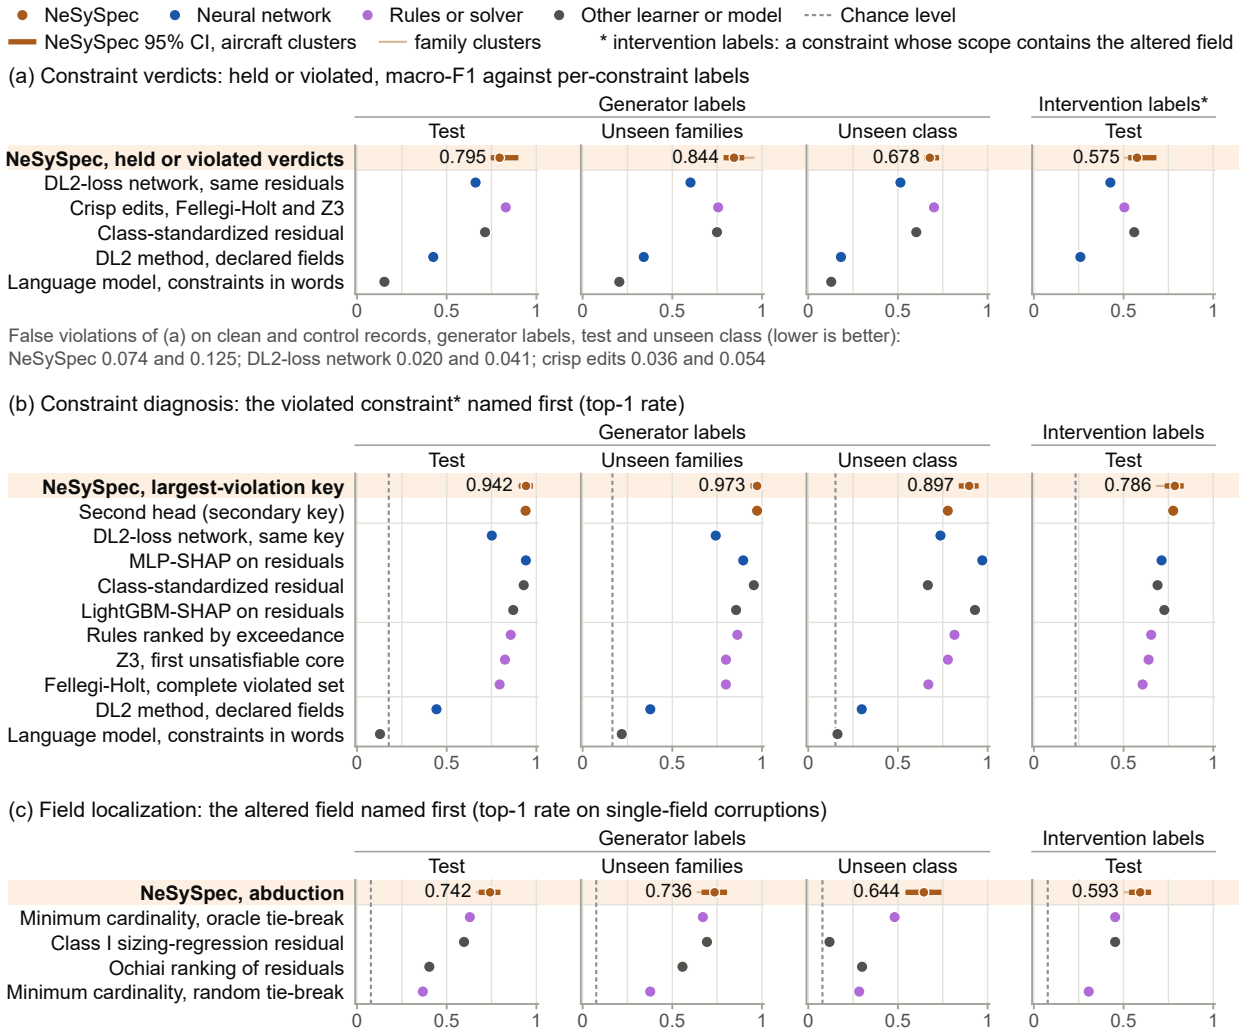

Figure 6: Constraint verdicts, diagnosis and field localization.

0.857; Figure 6b); 49 to 85% of these margins come from the 16 test records that the crisp edits leave unflagged (key 0.812, chance 0.106), and on the 140 they flag it is on par with them. On the 81 test records with a single violated constraint the key reaches 0.963, above the class-standardized residual, a z-score of the same residuals, and MLP-SHAP, which are on par with it over all corrupted test records; it comes with a held or violated verdict against learned tolerances, the altered field with its repair and exact abstention under missing evidence, all read from NeSySpec's own degrees, and the residual's field ranking is below abduction (Section 4.3). On unseen families the key reaches 0.973 on each of 8 seeds, above Fellegi-Holt, Z3, MLP-SHAP and LightGBM-SHAP and on par with the class-standardized residual, with the rules at 0.863 (cells without an alternative-stream check; Appendix G.2), and on the unseen class 0.897 ± 0.007, above Fellegi-Holt, Z3, the rules and the class-standardized residual and on par with LightGBM-SHAP, with MLP-SHAP, a second model on the residuals, at 0.970 (Table G.6). On intervention labels without retraining, the key ranks first a constraint whose scope contains the altered field in 0.786 ± 0.006 of corrupted test records (chance 0.232), above Fellegi-Holt, Z3 and rule diagnosis, with most of these margins from the 76 records that the crisp edits leave unflagged, and on par with LightGBM-SHAP, the strongest attribution on these labels; on unseen families NeSySpec's second head, its secondary key, reaches 0.848 (largest-violation key 0.828; Appendix G.2). With the optional rank-then-explain setting, NeSySpec flags and correctly diagnoses 0.901 ± 0.013 of corrupted test records and 0.849 ± 0.043 on the unseen class (5 seed pairs), above Fellegi-Holt and Z3 behind the same alarm on test and on par with MLP-SHAP, with the violated constraint, the altered field and a repair from its own degrees; with its own alarm, the record verdict and the key reach 0.929 on test (Appendix G.2).

Figure 5 shows the report of the first single-field corruption of the test split whose altered field an available constraint reads, chosen by a fixed rule. Its total thrust is altered: the report names A4, the thrust-to-weight ratio below the band learned from the fleet, as the largest violation, abduction finds the thrust, and changing the thrust alone brings V below the decision threshold (0.181 against 0.310). A9, flagged by the oracle's crisp band, is held, its thrust per engine inside the wider learned band; A5, which does not read the thrust, is violated as on the clean record (asterisk in panel b), so the repair, short of its stricter target, is a best-effort value to cross-check.

4.3 Field Localization and Repair

Abduction names the altered field more often than minimum-cardinality localization on every split, on generated and on independently labeled errors, and needs no oracle to break ties (Figure 6c). On single-field test corruptions (generator labels) it finds the altered field in 0.742 ± 0.011, above minimum-cardinality localization (0.367) and the Ochiai ranking of the class-standardized residuals; with the oracle tie-break, the setting most favorable to it, minimum-cardinality localization reaches 0.629, on par with abduction. On single-field records that no oracle selected or certified, abduction through the verifier fitted on generator labels, without retraining, finds the field moved beyond tolerance in 0.593 of test records, above minimum-cardinality localization (0.306) and its oracle tie-break (0.453); 78% of the margin over the oracle tie-break comes from the 70 records that the crisp edits leave unflagged, and on the 175 of 245 that they flag abduction stays above random tie-breaking and on par with the oracle tie-break. With the selected abduction rule (Appendix G.9), abduction is also ahead of the Class I sizing-regression residual, the field farthest from the conceptual-design (Class I) weight regressions on maximum take-off mass fitted on the training aircraft [71], on test (0.597), on the unseen class (0.119) and on intervention labels (0.453, the level of the oracle tie-break as well), and on par with it on unseen families (0.693), with the violated constraint and a repair from the same verdict.

The repair proposal of Eq. (13) comes from NeSySpec's own abduction and verdict, with the field and the violated constraint in the same report. It cuts the log-space error on the altered field from 0.366 to 0.133 and on the whole record to 0.200, and is closer on the record than the Fellegi-Holt, MaxSMT, nearest-neighbor and regression repairs and on the altered field than a constraint-layer repair ordering in the style of PiShield [47]; at the field of the Class I residual, MICE [83] (0.120 and 0.161) and missForest are on par with it on both errors, as is the regression on the altered field (Table G.6). The solver reaches its target, a verdict $V \leq -\ln 0.9 \approx 0.1054$ stricter than the decision threshold, on 0.555 of the single-field test records and otherwise returns its best-effort value (Figure 5c); the report issues a proposal for each flagged record, as a value to cross-check.

4.4 Record Coherence

On generator labels, with 16 test aircraft of 9 families held out of training, the record verdict tells coherent from incoherent records at a test AUROC of 0.979 ± 0.002, above learners on the raw fields (0.748 to 0.792), and each of its flags arrives with the violated constraint, the altered field and a repair proposal, read from the degrees of Figure 4. In distribution TabPFN v2 is on par with it and the other learners on the residuals are close to it, the residual MLP at 0.990 ± 0.003 and TabPFN v2, LightGBM and the one-class SVM at 0.957 to 0.958 (Appendix G.4). At a TPR of 0.90 the record verdict's false alarms on the 32 controls (0.086) are fewer than those of the Mahalanobis distance (0.656), each naming the constraint that raised it (Table G.3).

On the unseen class of business jets, the key and abduction lead classical diagnosis and minimum-cardinality localization (Sections 4.2 and 4.3). In detection there the record verdict reaches 0.934 and five learners on the residuals rank higher (residual MLP 0.978, TabPFN v2 up to 0.984), and NeSySpec's optional rank-then-explain setting lets the residual MLP raise the alarm and NeSySpec explain it, with false alarms on 0.205 of the 38 controls and 0.168 of the 19 pristine records at validation thresholds. On intervention labels, whose altered single fields move beyond the tolerance, by up to 80%, the key and abduction lead the classical comparators without retraining (Sections 4.2 and 4.3); in detection there, the residual MLP retrained on these labels and TabPFN v2, whose statistics include the scored rows, rank test records higher (Table G.3).

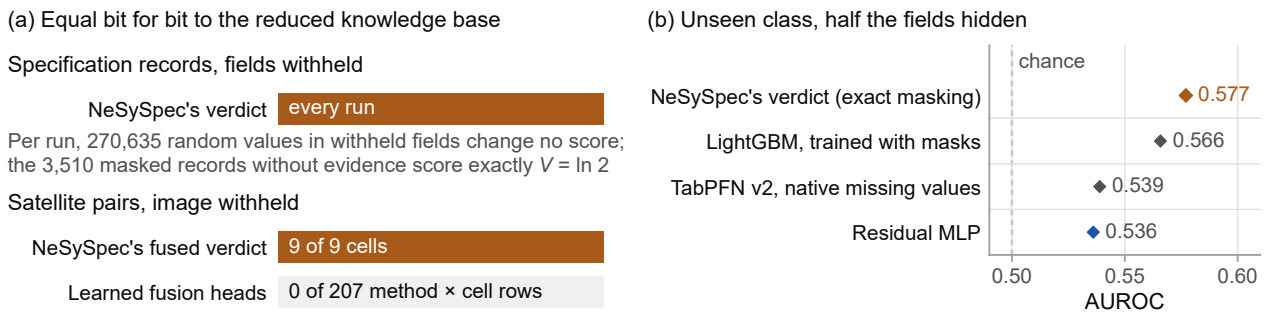

Figure 7: Verdicts under missing evidence.

Records with images. With a paired image, NeSySpec also checks the record against the image and names the cross-modal axiom that the image contradicts (Section 3.4). On satellite pairs, the multimodal verifier of Eq. (10) achieves an AUROC of 0.994±0.002 over all pair types with a ConvNeXt-T backbone [90] (8 seeds). Qwen3-VL-8B [53], prompted with rules R1-R10, reaches 0.748; the verifier's advantage holds across all evaluated backbones.

With ResNet-50 used for both, the verifier also outperforms end-to-end fusions trained on labeled pairs on test and unseen sites. Some learned fusions achieve higher AUROC on image swaps and unseen-site pairs. Without an image, the verifier returns the record verdict bit for bit. Appendix G.5 details the comparison protocols and their asymmetries. On photographs the cross-modal stratum is on par with Qwen3-VL-8B and Qwen3-VL-32B prompted with the same axioms on test (0.692 against 0.668 and 0.712) and prints the part map of its evidence beside the contradicted cross-modal axiom; on the families held out of NeSySpec's training, most of which the 32-billion-parameter model names from the photograph alone, that model reaches 0.827 and the stratum 0.733. This advantage may partly reflect the visual and semantic priors acquired through extensive multimodal pretraining [53]. These families are held out from our task-specific training; pretraining exposure is not controlled across the compared systems. Further comparisons should distinguish prior recognition from reasoning over the supplied physical constraints.

4.5 Missing Evidence and Absent Images

By Proposition 1, masked quantification makes every verdict a function of the available evidence only, and every run keeps this property (Figure 7a): every masked record scores bit for bit as its sub-knowledge-base, random values written into withheld fields change nothing, and masked records without evidence score exactly $V = \ln 2$; without an image, NeSySpec's fused verdict returns the tabular verdict bit for bit in 9 of 9 satellite cells, against 0 of 207 learned method × cell rows.

With fields hidden completely at random (MCAR) [91], a residual with a hidden input is missing for the learners; with half the fields hidden on the unseen class, where 26.4% of the records keep no available constraint, NeSySpec reaches 0.577, above the residual MLP (0.536) and TabPFN v2 with native missing values (0.539) and on par with LightGBM trained with masks (0.566; Figure 7b), while MICE imputation at scoring yields a slightly higher AUROC (+0.033). The masked verdict reads no filled value, and its report states which constraints remain checkable. With 10% of the fields hidden in distribution it reaches 0.848 on test and 0.824 on unseen families, where the MLP trained with masks (0.891 on test), the residual MLP (0.916 on unseen families) and, with 10 to 30% hidden on unseen families, MICE imputation of every missing field at scoring rank higher; with no field hidden, the masked verdict is ahead of MICE at scoring (Appendix G.3).

With MICE imputation enabled, NeSySpec keeps its constraint verdicts and largest-violation key and ranks records at 0.914 with 10% of the fields hidden on test, on par with the MLP trained with masks given the same imputer (0.905); at its threshold it flags 0.210 of the clean-labeled records, against 0.137 for the masked verdict, NeSySpec's default.

With 30% of satellite images hidden, the fused verdict keeps the exact reduction that no learned head keeps and ranks test pairs at 0.947, and heads trained with modality dropout [92], such as CatBoost on knowledge-base truths (0.978), rank them higher; across sensors without labels the record

verdict is the fallback (Appendix G.5).

4.6 Grounding of the Visual Axioms

When a head asserts an attribute, the grounding axiom encourages it to place its evidence on the corresponding part. The Reichenbach gate scales this soft penalty with the strength of the head's assertion, without using attribute labels. The report then displays the evidence map for inspection.

We compare methods using per-example penalties, keeping the batch's normalization constants fixed and stopping gradients through those statistics. Loss weights are selected on validation using one seed per candidate value (Appendix H). Under this common protocol, the network trained with the grounding axiom localizes the annotated parts substantially better than RRR [57] and GradMask [58]. Every cluster-bootstrap interval of its difference from these two and from Pukdee et al. [61], run both with the calibration of the protocol, which flattened their teacher toward the uniform prediction, and with a recalibrated arm, lies above zero on the FGVC pseudo-masks and on the human part masks of PASCAL-Part, PartImageNet [85] and UDA-Part [86] (Appendix H). On test pseudo-masks it points to the part in 0.780 ± 0.017 of maps (8 seeds), against 0.360 ± 0.066 with permuted masks and 0.103 ± 0.017 without a localization term (3 runs each), so localization depends on anatomically correct masks, and it points to the annotated part in 0.737 ± 0.015 (3 runs at weight 1) of PASCAL-Part human masks outside the FGVC domain (Figure G.1 in Appendix G.6).

Across loss weights 0.3, 1 and 3, the grounding axiom keeps at most one collapsed head, compared with up to 5.7 on average for the auxiliary localization loss. Its pointing accuracy rises from 0.512 at weight 0.3 to 0.790 at 1 and 0.803 at 3, with three runs per weight. Localization also holds across all evaluated backbones with per-head maps (Appendix G.6). At the validation-selected weights (1 for the axiom and 0.3 for the auxiliary loss), the axiom has a small accuracy advantage out of distribution (Appendix H). Localization is on par on test, out of distribution and on PASCAL-Part when comparing three runs per method. Including all eight axiom runs gives the auxiliary loss a higher pointing accuracy on test pseudo-masks (Appendix F).

Figure 8 shows maps of NeSySpec's visual module with the grounding axiom and without a localization term: on ground photographs (a) each map prints its in-mask share ρ of Eq. (6); on satellite crops (b), outlined by the annotated diamond on which the axiom trains, the maps show where the heads look. Each row names the cross-modal axiom its head feeds, and the report prints the map beside that axiom's verdict (part maps in Figure 1; layout in Figure E.1), so the analyst sees on which part the evidence of a contradicted axiom lies.

4.7 Robustness to Deliberate Falsification

Under a black-box falsification protocol of bare lies, cover-ups that edit up to five further fields, and omissions (Table 4 and Appendix G.7), a lie flagged by the record verdict is covered less often than one flagged by the residual MLP: with up to five further edits 0.936 of cover-ups pass against 1.000, and 0.247 of its flagged bare lies cannot be covered, against 0.000 for the MLP; a minimal cover of its flagged bare lies takes 2.69 further edited fields on average, against 1.38, 1.37 and 1.63 for the MLP, LightGBM and TabPFN v2 on the residuals (point estimates, censored at six; Appendix G.7). At that budget the record verdict is also evaded less often than TabPFN v2, the one-class SVM, LightGBM on the raw fields and an isolation forest, while the share of its flagged bare lies beyond cover is 0.247, against 0.494 and 0.325 for the rules with hierarchical pooling and a Mahalanobis distance on the declared fields, whose minimal covers take more edits (3.62 and 3.45); each of its flags comes with the field found by abduction and a repair proposal. On bare lies the record verdict (evasion 0.740 ± 0.038) is evaded less often than the rules with hierarchical pooling and LightGBM on the raw fields and on par with the MLP (0.754); LightGBM on the residuals (0.685) and TabPFN v2, evaded less often on bare lies, leave almost none of their flagged lies beyond cover (0.004 and 0.024, against 0.247).

By Proposition 1, a withheld field masks the axioms that require it, and the report lists every undeclared field with the constraints it leaves unchecked; withholding up to two fields lets 0.964 of lies pass the record verdict alone. NeSySpec's optional critical-field alert addresses this channel, and each of its flags names the undeclared critical field or the violated constraint: against an attacker who

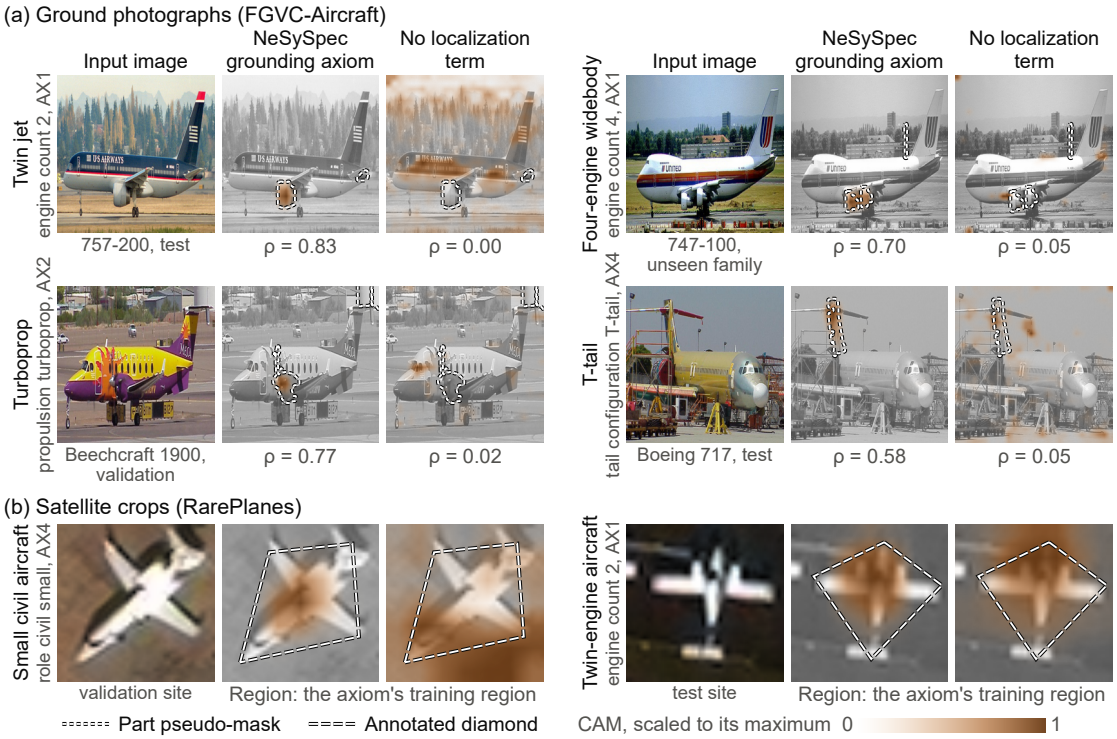

Figure 8: Evidence heatmaps.

knows the alert and withholds up to two fields, 0.729 of lies pass NeSySpec with the alert, on par with the residual MLP given the same alert (0.756) with all splits pooled. With the alert, false alarms on the 32 test controls are 0.148, below the Mahalanobis distance (0.656).

4.8 What the Logic Adds: Other Encodings of the Same Knowledge

The logical formulation provides the constraint-level verdicts used for diagnosis and repair. For record-level detection, learners on the same residuals perform similarly to NeSySpec in distribution and better on the unseen class (Section 4.4). Against the network trained with the DL2 loss on the same residuals and targets, which learns one head per constraint in place of the graded predicates and is their like-for-like test, NeSySpec's graded predicates trained through Real Logic give constraint verdicts with a higher macro-F1 and a better key, with the field and a repair from the same degrees (Section 4.2). Of the other Real Logic properties of Section 2.1, guarded quantification gives the exact abstention that none of 207 learned method and cell rows keeps without an image (Section 4.5), the shared truth scale the fused verdict (Section 4.4) and the gate grounding stable across loss weights (Section 4.6). Stratification (E4) keeps one verdict per evidence source, so the report says whether the record or the image raised the flag (Appendix G.8). Training the knowledge base through the differentiable truths of Real Logic improves detection over keeping it at its initial values: AUROC rises from 0.907 to 0.979 on test and from 0.845 to 0.934 on the unseen class. On generator labels, the full model outperforms the flat-bank ablation on every split. Removing in-band grading lowers detection on the unseen class and slightly on test and unseen families. The ablation using a generalized mean at $p = 4$ is on par with the full model (Table G.8). In detection on test the record verdict (0.979) is slightly above the DL2-loss network (0.965) and a per-constraint monotone model on the same residual magnitudes without LTN (0.966), and on par with the network out of distribution; NeSySpec reads these magnitudes as graded predicates in physical units (Figure 5c), with the key, abduction and repair from the same degrees. The generator's per-constraint targets (Section 4.2) and the graded predicates act together: the ablation trained with the binary record label only equals the generator labels in 0.081 of test records and is below the crisp edits on all three measures. Given as a constraint loss on the declared fields or as rules in words, the same knowledge detects less than NeSySpec's verdict: the DL2 method on the declared fields reaches 0.703 ± 0.030 and a language model prompted with the constraints A1-A11 in words

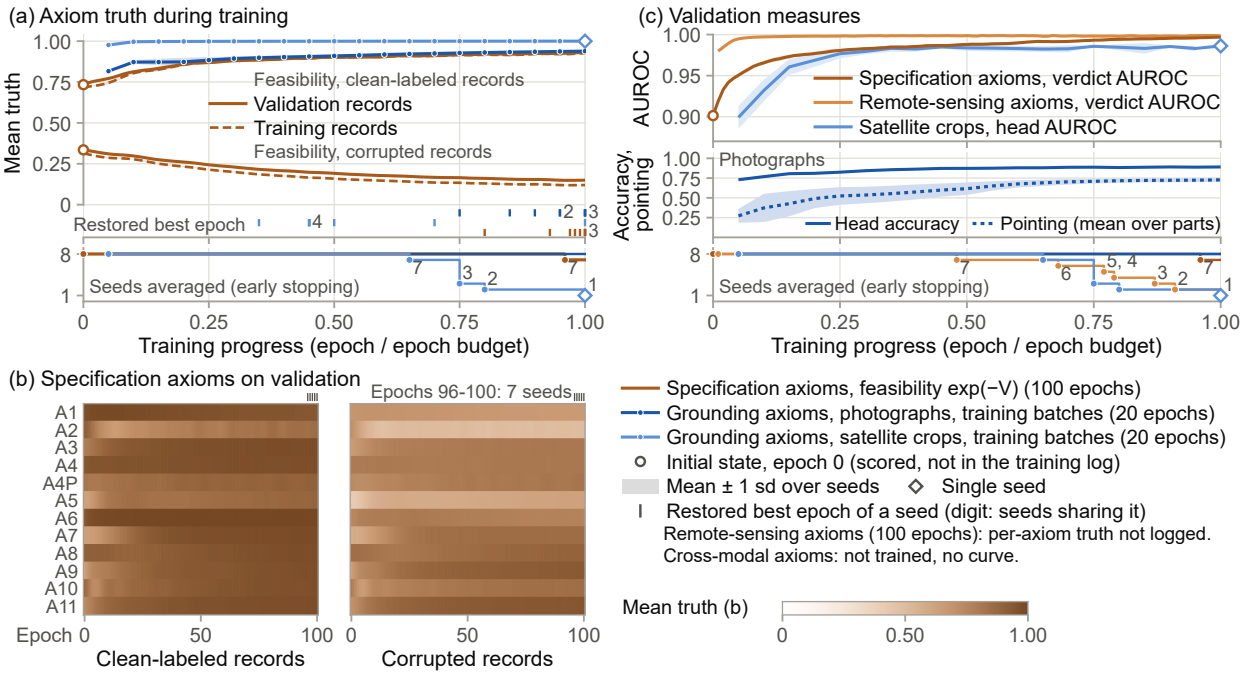

Figure 9: Training dynamics of the axioms.

(Qwen3-8B [52], anonymized records) 0.620, below the record verdict on every split (Table G.7), and both are below its key and verdicts (Section 4.2).

Figure 9 traces the training of the axioms over eight seeds per stratum. For the specification axioms, a re-run of NeSySpec's training whose only addition is a copy of the state at each epoch reproduces every frozen checkpoint and its predictions bit for bit, together with every logged loss and validation AUROC; each state is scored at the end of its epoch on the training and validation records. From the initial state, shared by all seeds, to the restored best epoch, the mean feasibility $\exp(-V)$ on validation rises from 0.734 to 0.930 on clean-labeled records and falls from 0.335 to 0.151 on corrupted ones, while the verdict's validation AUROC rises from 0.901 to 0.997. Panel b averages each axiom's truth $\exp(-v_k)$, which the constraint verdicts compare with 0.5, over the records where the axiom is available; the feasibility of Eq. (9) in panel a is the geometric mean of a record's available truths. The satisfaction of the grounding axioms, averaged over the training batches of each epoch while the weights move, rises from 0.816 at the first epoch to 0.936 at the restored best epoch on photographs and from 0.976 to 0.999 on satellite crops. The training logs of the remote-sensing axioms hold no per-axiom truth, so panel c shows their verdict's validation AUROC beside the photographs' head accuracy and pointing; the cross-modal axioms have no trainable parameter, and late epochs average fewer seeds where early stopping ended a run.

Detection also improves with the knowledge the base encodes (Figure 10). Retrained with subsets of its 12 specification constraints, the record verdict's intervention-label AUROC, averaged over the three held-out splits, rises from 0.535 with one constraint to 0.723 with all twelve (Spearman 0.93). It rises with diminishing returns in the number of constraints (panel a) and linearly with the share of edits that alter a field some of them read (R^2 0.88, panel b); the residual MLP on the same subsets gains as well (0.540 to 0.706), with no trend in their gap. Removing one constraint (panel c) lowers detection for A5, the Breguet range, and A2, the cruise lift coefficient, and slightly for A4, A7 and A8, the thrust-to-weight, mass-balance and empty-weight-fraction checks, and raises it for none; it also lowers abduction for 7 of the 12 constraints and slightly for 3 more (Appendix G.8).

4.9 Deployment Map

Table 6 assigns each regime to the output of NeSySpec that serves it, with the workflow of an audit and the strongest comparator of each regime.

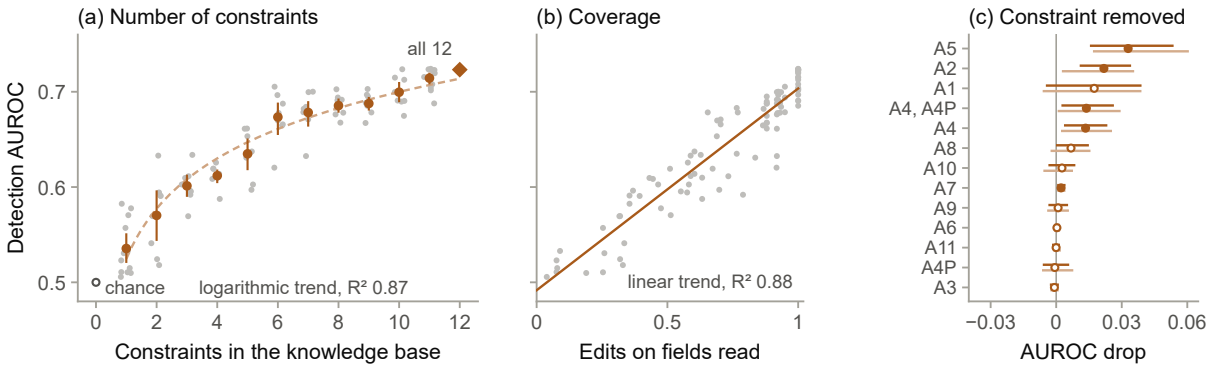

Figure 10: Detection and the knowledge encoded.

4.10 Limitations

Each limitation is bounded by the protocol and opens a line of future work (Section 5). *Domain and use.* NeSySpec is validated on aircraft, whose public civil records and relations a reader can verify; only its knowledge base, visual predicates, corruption generator and data are domain-specific (Appendix A), and its verdicts support a human analyst, with aircraft detection, tracking and identification outside its scope because of the dual-use risks of such capabilities [93]. *Synthetic errors and scale.* Every error is generated or edited, so its type, magnitude and field are known and each verdict is checked against them, and the intervention labels are independent of the generator; they mark the constraints an edit involves, about a third of the altered rows break none of them physically, and expert per-constraint physical labels will score the verdicts against violations (Section 5); the 107 aircraft and one published split hold every test family and satellite site out of training. Naturally occurring errors, expert red teams, larger fleets and a replication on a new fleet will extend this evidence. *Class-level relations.* A2 evaluates the cruise lift coefficient at maximum take-off mass and A5 uses latents shared by a class and engine type, estimable from 44 training aircraft; most violations reported on clean records come from these two relations (Section 4.2), which per-type latents and a cruise-weight form of A2 will refine. *Classes without training aircraft.* There, residual learners rank higher in record-level detection (Section 4.4); the optional rank-then-explain setting lets such a learner raise the alarm while NeSySpec names the constraint, the field and a repair, and larger fleets will reduce the classes read with fleet statistics alone. *Omission.* A withheld field masks the axioms that read it, the exact abstention of Proposition 1; the optional critical-field alert flags the omission, and class-level axioms will make the omission of a critical field a violation. *Masks, statistics and encodings.* Automatic pseudo-masks supervise the grounding axiom, so the protocol is reproducible from public data, and hand-annotated part masks will extend it; no multiplicity correction is applied, which the new-fleet replication will address; encoding the knowledge base in DeepProbLog, Scallop or PSL, with probabilistic facts and proofs, changes its axiomatic and logical structure and is left to a dedicated comparison.

5 Conclusion and Future Works

NeSySpec is an integrated and interpretable neuro-symbolic framework for auditing engineering specifications: its graded axioms on Logic Tensor Networks, trained on fleet data, determine which physical constraints a declared specification respects and which it violates, and from the degree behind each verdict it derives the altered field, a repair proposal, exact abstention on missing evidence, part-grounded image evidence and a verification report. Validated on 107 aircraft with generated, intervention-labeled and falsified errors, it provides and evaluates these outputs together.

Using a fixed rule to classify each constraint as held or violated, NeSySpec reaches a test macro-F1 of 0.795 with generator labels and, without retraining, 0.575 with intervention labels. On every split evaluated with generator labels, it outperforms three comparators in macro-F1: a network trained with the DL2 loss on the same residuals; DL2 applied to the declared fields; and a language model given

Table 6: Deployment map.

Regime	What NeSySpec provides	Strongest comparator
Workflow: run NeSySpec first and read its report; on a class without training aircraft, enable its optional rank-then-explain setting, in which a residual learner raises the alarm and NeSySpec explains it; cross-check a repair value against a statistical imputer such as MICE; against omission, enable the optional critical-field alert.		
Constraint verdicts and diagnosis (Section 4.2)		
classical diagnosis and other encodings	the largest-violation key, from the verdict's own degrees, above Fellegi-Holt, Z3 and the rules (also on intervention labels), the DL2-loss network's key, the DL2 method on the declared fields and a language model; 0.812 where the crisp edits flag nothing	the crisp edits, on par where they flag the record (generator labels)
held or violated, test and unseen families	the constraint verdicts, macro-F1 0.795 on generator-label test, above the DL2-loss network, with the key, the field and a repair	crisp edits, on par in macro-F1; the DL2-loss network, with fewer false violations (0.020 against 0.074)
residual rankings	the key, above both on 81 single-violation records and above the residual on the unseen class, with a held or violated verdict, a repair and exact abstention	the class-standardized residual and MLP-SHAP, on par in distribution; MLP-SHAP on the unseen class
Field localization and repair (Section 4.3)		
field	abduction, above random tie-breaking on every split and label family and above the Class I sizing-regression residual on the unseen class and intervention labels, with no oracle to break ties	an oracle tie-break, on par in distribution
value, in distribution	the repair proposal, with its field and violated constraint in the same report, issued for flagged records (the proposals change 0.562 of clean and control records); the solver reaches its target on 0.555 of single-field test records	MICE at the field of the Class I sizing-regression residual, on par, the recommended value-recovery step
Record coherence (Section 4.4)		
in distribution	the record verdict, which with the key flags and correctly diagnoses 0.929 of corrupted test records, with the constraint, field and repair of each flag	the residual MLP, close in detection, and TabPFN v2, on par
class without training aircraft	rank-then-explain (optional): false alarms on 0.205 of controls and 0.168 of 19 pristine records, with the constraint, field and repair behind each alarm	the residual MLP and TabPFN v2, in detection
intervention labels	the key and abduction, without retraining, naming the constraint and the field	TabPFN v2 (statistics from the scored rows) or the MLP retrained on them, in detection
record with an image	the multimodal verifier (satellite), above a vision-language model given the rules and end-to-end fusions, and the cross-modal stratum (photographs), each naming the contradicted axiom, with the exact reduction	vision-language models given the axioms, on par on test photographs; learned fusions on image swaps and unseen-site pairs; the largest vision-language model on unseen families
Missing evidence and absent images (Section 4.5)		
10% of fields hidden, test or unseen families	the coverage, for the audit; NeSySpec with MICE imputation enabled, to detect	the MLP trained with masks (test) and the residual MLP (unseen families), against the masked verdict; the MLP with masks and the same imputer, on par with the MICE setting (test)
half hidden, unseen class	the masked verdict, above the residual MLP and TabPFN v2, with the checkable constraints and an exact verdict	LightGBM trained with masks, on par; MICE imputation at scoring, close
image absent or partly missing	the fused verdict, which reduces exactly to the record verdict	on labeled pairs, heads trained with modality dropout, such as CatBoost on knowledge-base truths, to rank
Grounding of the visual axioms (Section 4.6)		
any backbone	the grounding axiom, which localizes the annotated parts above RRR and GradMask under one protocol, with at most one collapsed head across loss weights	an auxiliary localization loss at its validation weight, on par with 3 runs each and higher on test pseudo-masks against the axiom's 8 seeds
Robustness to deliberate falsification (Section 4.7)		
cover-ups, up to five further edits	the record verdict, which a cover-up passes less often than the residual MLP, with the altered field and a repair for each flag	rules with hierarchical pooling or a Mahalanobis distance, on par
omission of up to two fields, alert-aware	the optional critical-field alert, naming the undeclared field or violated constraint of each flag	the residual MLP with the same alert, on par with all splits pooled

On intervention labels without retraining, on unseen families, NeSySpec's second head, its secondary key, reaches 0.848 (largest-violation key 0.828). Satellite cells: comparators not retrained or retuned, and fusions on knowledge-base truths read the truths of the saturating-envelope variant of NeSySpec's remote-sensing stratum they were fitted on; asymmetries in Appendix G.5.

the constraints in words.

The largest-violation key identifies the violated constraint in 0.942 of corrupted test records with generator labels, outperforming Fellegi–Holt, solver-based diagnosis and rule-based diagnosis. On independently labeled errors, abduction identifies the altered field in 0.593 of records without retraining. Minimum-cardinality localization reaches 0.306 with random tie-breaking and 0.453 with oracle tie-breaking. The Class I sizing-regression residual also reaches 0.453. The record verdict detects inco-

herent records at a test AUROC of 0.979 on generator labels, and a missing field or image leaves each verdict exactly that of the available constraints.
An image lets the audit check the declared record against visible attributes and identify the axiom behind a disagreement. On satellite pairs, the multimodal verifier achieves an AUROC of 0.994 across all pair types, compared with 0.748 for a prompted vision-language model.
The report also shows where each visual head places its evidence. With the grounding axiom, these maps point to the annotated part in 0.737 of PASCAL-Part human masks. Localization is substantially better than RRR and GradMask on human masks from PASCAL-Part, PartImageNet and UDA-Part, with wider intervals on the 40 UDA-Part items (Appendix H).

Among the approaches evaluated here (the DL2 method on the declared fields, the network trained with the DL2 loss on the same residuals read per constraint, a constraint-layer repair ordering, LTN ablations, solver and rule diagnosis and prompted language and vision-language models), NeSySpec is, to our knowledge, the most suitable neuro-symbolic framework for auditing engineering specifications: it provides and evaluates together the outputs of Table 2, and its key and abduction lead classical diagnosis, the DL2 method, a prompted language model and minimum-cardinality localization with random tie-breaking. No comparator returns all the outputs of the audit, and on the test aircraft none is ahead of NeSySpec in the macro-F1 of the constraint verdicts, in naming the violated constraint and the altered field or in the repair error on the altered field (Table 5). With oracles fitted without test or out-of-distribution aircraft, its leads on intervention labels over Fellegi-Holt and Z3 in diagnosis and over minimum-cardinality localization with random tie-breaking hold on every split, at lower levels (Section 4.1). Its logic serves the physical audit more than record classification: in record-level detection, learners on the same residuals are close to it and in some settings rank records higher (Section 4.4). Its detection grows with the knowledge it encodes, following the share of edits that alter a field its constraints read, with diminishing returns in their number.

This study is a first approach to specification auditing, evaluated on aircraft and constructed errors. Broader physical knowledge, richer input evidence and larger, more diverse datasets may improve generalization to unseen systems, to be tested through independent replications. The logical formulation helps interpret anomalies by exposing which relations fail and which fields are implicated. Future work will extend NeSySpec along the following directions:

- *Other neuro-symbolic encodings.* A version of the knowledge base in DeepProbLog, Scallop and PSL, with probabilistic facts and proofs in place of graded truths and masked quantification, will be compared with NeSySpec, together with the semantic loss and the monotone model read per constraint.

- *Aircraft-level latents and refined relations.* Per-type latents in A5 and a cruise-weight formulation of A2 will refine the constraint verdicts. Per-constraint thresholds selected on validation will calibrate these verdicts, and expert physical labels for each constraint will be used to evaluate them. The report will display every learned band in physical units. A new relation involving the ceiling will cover the only declared field not currently checked by any constraint. Additional relations can improve detection by checking different physical relationships among fields already covered elsewhere, as A2 does.

- *Other engineering applications.* Visual recognition of components and geometry could extend the audit to mechanical assemblies, industrial equipment and civil structures, with adapted knowledge bases, visual predicates and corruption generators.

- *Poorly characterized observations.* The framework could also support the analysis of observations whose identity or physical behaviour remains uncertain. Physical constraints could help compare candidate explanations and identify the additional measurements needed to distinguish them.

- *Generative design and simulation.* Exporting structured parameter vectors, units and constraints could connect the audit to agentic pipelines that generate parametric CAD models from text and images. Candidate models could then undergo finite element method (FEM) simulations, computational fluid dynamics (CFD) and other numerical analyses, feeding their results back into the audit.

- *Training and hyperparameters.* Further ablations will assess how supervision, loss balancing and joint versus separate module training affect diagnosis and generalization. Hyperparameter sensitivity will be evaluated on validation data.
- *Inference efficiency.* Profiling will measure the inference cost of each module. Batching, caching and adaptive search budgets will be evaluated in terms of latency, diagnostic accuracy and repair quality.
- *Naturally occurring errors.* Cross-source discrepancies, variant mix-ups, with block abduction, and uncorroborated fields will bring the audit to real errors, and expert red teams will extend the falsification protocol.
- *Evidence, images and declarations.* Imputation channels chosen with block-wise hiding will extend NeSySpec with MICE imputation enabled, hand-annotated FGVC-Aircraft part masks and visual strata adapted to new sensors will extend the grounding axiom, and class-level axioms will make the omission of a critical field a violation.

Declaration of Generative AI and AI-Assisted Technologies in the Manuscript Preparation Process

During the preparation of this work, the authors used AI tools to improve the language, readability, and grammar of the manuscript, as well as to assist with debugging and bug hunting.

Data Licenses

FGVC-Aircraft [7] is provided for non-commercial research only and its photographs are not redistributed, apart from the four photographs of Figure 8, reproduced for illustration with credit to their photographers (airliners.net); RarePlanes [8] is distributed under the Creative Commons Attribution-ShareAlike 4.0 International license; the two satellite crops of Figure 8 are adapted from it (cropped, resized and overlaid with activation maps) and shared under the same license, under which the tables and crop indices derived from it will be released. The human part masks used for evaluation are not redistributed.

References

[1] Dominik Wienand and Heiko Paulheim. Detecting incorrect numerical data in DBpedia. In *The Semantic Web: Trends and Challenges (ESWC 2014)*, volume 8465 of *Lecture Notes in Computer Science*, pages 504–518. Springer, 2014.

[2] Junzi Sun, Jacco M. Hoekstra, and Joost Ellerbroek. OpenAP: An open-source aircraft performance model for air transportation studies and simulations. *Aerospace*, 7(8):104, 2020.

[3] Eckhard Lichtenthaler. Technology intelligence processes in leading European and North American multinationals. *R&D Management*, 34(2):121–135, 2004.

[4] Yi Zhang, Douglas K. R. Robinson, Alan L. Porter, Donghua Zhu, Guangquan Zhang, and Jie Lu. Technology roadmapping for competitive technical intelligence. *Technological Forecasting and Social Change*, 110:175–186, 2016.

[5] Michael Glassman and Min Ju Kang. Intelligence in the internet age: The emergence and evolution of open source intelligence (OSINT). *Computers in Human Behavior*, 28(2):673–682, 2012.

[6] Javier Pastor-Galindo, Pantaleone Nespoli, Félix Gómez Mármol, and Gregorio Martínez Pérez. The not yet exploited goldmine of OSINT: Opportunities, open challenges and future trends. *IEEE Access*, 8:10282–10304, 2020.

[7] Subhransu Maji, Esa Rahtu, Juho Kannala, Matthew Blaschko, and Andrea Vedaldi. Fine-grained visual classification of aircraft. arXiv:1306.5151, 2013.

[8] Jacob Shermeyer, Thomas Hossler, Adam Van Etten, Daniel Hogan, Ryan Lewis, and Daeil Kim. RarePlanes: Synthetic data takes flight. In *Proceedings of the IEEE Winter Conference on Applications of Computer Vision (WACV 2021)*, pages 207–217, 2021.

[9] Songqiao Han, Xiyang Hu, Hailiang Huang, Minqi Jiang, and Yue Zhao. ADBench: Anomaly detection benchmark. In *Advances in Neural Information Processing Systems 35 (NeurIPS 2022), Datasets and Benchmarks Track*, pages 32142–32159, 2022.

[10] Noah Hollmann, Samuel Müller, Lennart Purucker, Arjun Krishnakumar, Max Körfer, Shi Bin Hoo, Robin Tibor Schirrmeister, and Frank Hutter. Accurate predictions on small data with a tabular foundation model. *Nature*, 637(8045):319–326, 2025.

[11] Che-Ping Tsai, Ganyu Teng, Phillip Wallis, and Wei Ding. AnoLLM: Large language models for tabular anomaly detection. In *Proceedings of the 13th International Conference on Learning Representations (ICLR 2025)*, pages 7562–7592, 2025.

[12] Scott M. Lundberg and Su-In Lee. A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems 30 (NIPS 2017)*, pages 4765–4774, 2017.

[13] I. P. Fellegi and D. Holt. A systematic approach to automatic edit and imputation. *Journal of the American Statistical Association*, 71(353):17–35, 1976.

[14] Raymond Reiter. A theory of diagnosis from first principles. *Artificial Intelligence*, 32(1):57–95, 1987.

[15] Artur d'Avila Garcez and Luís C. Lamb. Neurosymbolic AI: The 3rd wave. *Artificial Intelligence Review*, 56(11):12387–12406, 2023.

[16] Laura von Rueden, Sebastian Mayer, Katharina Beckh, Bogdan Georgiev, Sven Giesselbach, Raoul Heese, Birgit Kirsch, Julius Pfrommer, Annika Pick, Rajkumar Ramamurthy, Michal Walczak, Jochen Garcke, Christian Bauckhage, and Jannis Schuecker. Informed machine learning: A taxonomy and survey of integrating prior knowledge into learning systems. *IEEE Transactions on Knowledge and Data Engineering*, 35(1):614–633, 2023.

[17] Kamal Acharya, Iman Sharifi, Mehul Lad, Liang Sun, and Houbing Song. Integrating neurosymbolic AI in advanced air mobility: A comprehensive survey. In *Proceedings of the Thirty-Fourth International Joint Conference on Artificial Intelligence (IJCAI 2025), Survey Track*, pages 10362–10370, 2025.

[18] Semanto Mondal, Antonino Ferraro, Fabiano Pecorelli, and Giuseppe De Pietro. A Logic Tensor Network-based neurosymbolic framework for explainable diabetes prediction. *Applied Sciences*, 15(21):11806, 2025.

[19] Samy Badreddine, Artur d'Avila Garcez, Luciano Serafini, and Michael Spranger. Logic Tensor Networks. *Artificial Intelligence*, 303:103649, 2022.

[20] Marc Fischer, Mislav Balunović, Dana Drachsler-Cohen, Timon Gehr, Ce Zhang, and Martin Vechev. DL2: Training and querying neural networks with logic. In *Proceedings of the 36th International Conference on Machine Learning (ICML 2019)*, volume 97 of *Proceedings of Machine Learning Research*, pages 1931–1941. PMLR, 2019.

[21] Emile van Krieken, Erman Acar, and Frank van Harmelen. Analyzing differentiable fuzzy logic operators. *Artificial Intelligence*, 302:103602, 2022.

[22] Ruiyuan Kang, Tingting Mu, Panagiotis Liatsis, and Dimitrios Kyritsis. Physics-driven ML-based modelling for correcting inverse estimation. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*, pages 15639–15659, 2023.

[23] Mark P. J. van der Loo and Edwin de Jonge. Data validation infrastructure for R. *Journal of Statistical Software*, 97(10):1–31, 2021.

[24] Wei Ni, Xiaoye Miao, Xiangyu Zhao, Yangyang Wu, Shuwei Liang, and Jianwei Yin. Automatic data repair: Are we ready to deploy? *Proceedings of the VLDB Endowment*, 17(10):2617–2630, 2024.

[25] Theodoros Rekatsinas, Xu Chu, Ihab F. Ilyas, and Christopher Ré. HoloClean: Holistic data repairs with probabilistic inference. *Proceedings of the VLDB Endowment*, 10(11):1190–1201, 2017.

[26] Mohammad Mahdavi and Ziawasch Abedjan. Baran: Effective error correction via a unified context representation and transfer learning. *Proceedings of the VLDB Endowment*, 13(12):1948–1961, 2020.

[27] Leonardo de Moura and Nikolaj Bjørner. Z3: An efficient SMT solver. In *Tools and Algorithms for the Construction and Analysis of Systems (TACAS 2008)*, volume 4963 of *Lecture Notes in Computer Science*, pages 337–340. Springer, 2008.

[28] Nikolaj Bjørner, Anh-Dung Phan, and Lars Fleckenstein. νZ: An optimizing SMT solver. In *Tools and Algorithms for the Construction and Analysis of Systems (TACAS 2015)*, volume 9035 of *Lecture Notes in Computer Science*, pages 194–199. Springer, 2015.

[29] C. Eastman, Jae-min Lee, Yeon-suk Jeong, and Jin-kook Lee. Automatic rule-based checking of building designs. *Automation in Construction*, 18(8):1011–1033, 2009.

[30] Odin Iversen and Lizhen Huang. Leveraging large language models for BIM-based automated compliance checking. *Automation in Construction*, 182:106707, 2026.

[31] Sebastian Schelter, Dustin Lange, Philipp Schmidt, Meltem Celikel, Felix Biessmann, and Andreas Grafberger. Automating large-scale data quality verification. *Proceedings of the VLDB Endowment*, 11(12):1781–1794, 2018.

[32] Xin Chen and Kai Cheng. AI-assisted auditing review of mechanical engineering drawings: A Perception-Cognition-Collaboration reference architecture and its implementation perspectives. *International Journal of Computer Integrated Manufacturing*, pages 1–22, 2026. Published online, doi:10.1080/0951192X.2026.2714117.

[33] Daniel Jung, Kok Yew Ng, Erik Frisk, and Mattias Krysander. Combining model-based diagnosis and data-driven anomaly classifiers for fault isolation. *Control Engineering Practice*, 80:146–156, 2018.

[34] Arman Mohammadi, Mattias Krysander, and Daniel Jung. Consistency-based diagnosis using data-driven residuals and limited training data. *Control Engineering Practice*, 159:106283, 2025.

[35] Tim Bohne, Anne-Kathrin Patricia Windler, and Martin Atzmueller. A domain-agnostic neuro-symbolic architecture for multimodal human-in-the-loop anomaly detection and complex fault diagnosis. *IEEE Access*, 13:210201–210236, 2025.

[36] Haoliang Huang, Zihuang Cai, Zhuo Tang, Yifan Liu, Chen Tian, Kenli Li, and Changjian Chen. Integrating data validation with large language models for regulation-guided tabular anomaly detection. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6559–6581. Association for Computational Linguistics, 2026.

[37] Mengyi Yan, Yaoshu Wang, Yue Wang, Xiaoye Miao, and Jianxin Li. GIDCL: A graph-enhanced interpretable data cleaning framework with large language models. *Proceedings of the ACM on Management of Data*, 2(6):236:1–236:29, 2024.

[38] Wei Ni, Kaihang Zhang, Xiaoye Miao, Xiangyu Zhao, Yangyang Wu, Yaoshu Wang, and Jianwei Yin. ZeroED: Hybrid zero-shot error detection through large language model reasoning. In *2025 IEEE 41st International Conference on Data Engineering (ICDE)*, pages 3126–3139, 2025.

[39] Ivan Donadello, Luciano Serafini, and Artur d'Avila Garcez. Logic Tensor Networks for semantic image interpretation. In *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence (IJCAI 2017)*, pages 1596–1602, 2017.

[40] Eduard Hogea, Darian M. Onchiş, Ruqiang Yan, and Zheng Zhou. LogicLSTM: Logically-driven long short-term memory model for fault diagnosis in gearboxes. *Journal of Manufacturing Systems*, 77:892–902, 2024.

[41] Eduard Hogea, Darian M. Onchiş, Ruqiang Yan, and Zheng Zhou. Rule guided transformers for dynamic knowledge adaptation in rotating machinery fault diagnosis. *Advanced Engineering Informatics*, 72:104444, 2026.

[42] Robin Manhaeve, Sebastijan Dumančić, Angelika Kimmig, Thomas Demeester, and Luc De Raedt. Neural probabilistic logic programming in DeepProbLog. *Artificial Intelligence*, 298:103504, 2021.

[43] Ziyang Li, Jiani Huang, and Mayur Naik. Scallop: A language for neurosymbolic programming. *Proceedings of the ACM on Programming Languages*, 7(PLDI):1463–1487, 2023.

[44] Zhun Yang, Adam Ishay, and Joohyung Lee. NeurASP: Embracing neural networks into answer set programming. In *Proceedings of the Twenty-Ninth International Joint Conference on Artificial Intelligence (IJCAI 2020)*, pages 1755–1762, 2020.

[45] Wen-Chao Hu, Wang-Zhou Dai, Yuan Jiang, and Zhi-Hua Zhou. Efficient rectification of neuro-symbolic reasoning inconsistencies by abductive reflection. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 17333–17341, 2025.

[46] Jingyi Xu, Zilu Zhang, Tal Friedman, Yitao Liang, and Guy Van den Broeck. A semantic loss function for deep learning with symbolic knowledge. In *Proceedings of the 35th International Conference on Machine Learning (ICML 2018)*, volume 80 of *Proceedings of Machine Learning Research*, pages 5502–5511. PMLR, 2018.

[47] Mihaela Cătălina Stoian, Alex Tatomir, Thomas Lukasiewicz, and Eleonora Giunchiglia. PiShield: A PyTorch package for learning with requirements. In *Proceedings of the Thirty-Third International Joint Conference on Artificial Intelligence (IJCAI 2024), Demo Track*, pages 8805–8809, 2024.

[48] Alessandro Daniele and Luciano Serafini. Knowledge enhanced neural networks. In *PRICAI 2019: Trends in Artificial Intelligence*, volume 11670 of *Lecture Notes in Computer Science*, pages 542–554. Springer, 2019.

[49] Gabriele Ciravegna, Pietro Barbiero, Francesco Giannini, Marco Gori, Pietro Liò, Marco Maggini, and Stefano Melacci. Logic explained networks. *Artificial Intelligence*, 314:103822, 2023.

[50] Stephen H. Bach, Matthias Broecheler, Bert Huang, and Lise Getoor. Hinge-loss Markov random fields and probabilistic soft logic. *Journal of Machine Learning Research*, 18(109):1–67, 2017.

[51] Yung-Shen Hsia, Fang Yu, and Jie-Hong Roland Jiang. Neuro-symbolic compliance: Integrating LLMs and SMT solvers for automated financial legal analysis. In *2025 2nd IEEE/ACM International Conference on AI-powered Software (AIware)*, pages 1–10. IEEE, 2025.

[52] An Yang et al. Qwen3 technical report. arXiv:2505.09388, 2025.

[53] Shuai Bai et al. Qwen3-VL technical report. arXiv:2511.21631, 2025.

[54] Anna C. Doris, Daniele Grandi, Ryan Tomich, Md Ferdous Alam, Mohammadmehdi Ataei, Hyunmin Cheong, and Faez Ahmed. DesignQA: A multimodal benchmark for evaluating large language models' understanding of engineering documentation. *Journal of Computing and Information Science in Engineering*, 25(2):021009, 2025.

[55] Vito Barbara, Massimo Guarascio, Nicola Leone, Giuseppe Manco, Alessandro Quarta, Francesco Ricca, and Ettore Ritacco. Neuro-symbolic AI for compliance checking of electrical control panels. *Theory and Practice of Logic Programming*, 23(4):748–764, 2023.

[56] Jisoo Kim, TaeYoon Kwack, Jinwoo Jang, and Honguk Woo. Visual compliance via executable safety rule entailment. arXiv:2609.18328, 2026. Accepted to EMNLP 2026.

[57] Andrew Slavin Ross, Michael C. Hughes, and Finale Doshi-Velez. Right for the right reasons: Training differentiable models by constraining their explanations. In *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence (IJCAI 2017)*, pages 2662–2670, 2017.

[58] Becks Simpson, Francis Dutil, Yoshua Bengio, and Joseph Paul Cohen. GradMask: Reduce overfitting by regularizing saliency. arXiv:1904.07478, 2019.

[59] Yuyang Gao, Tong Steven Sun, Guangji Bai, Siyi Gu, Sungsoo Ray Hong, and Liang Zhao. RES: A robust framework for guiding visual explanation. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 432–442. ACM, 2022.

[60] Sukrut Rao, Moritz Böhle, Amin Parchami-Araghi, and Bernt Schiele. Studying how to efficiently and effectively guide models with explanations. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV 2023)*, pages 1922–1933, 2023.

[61] Rattana Pukdee, Dylan Sam, J. Zico Kolter, Maria-Florina Balcan, and Pradeep Ravikumar. Learning with explanation constraints. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*, pages 49883–49926, 2023.

[62] Siyi Du, Shaoming Zheng, Yinsong Wang, Wenjia Bai, Declan P. O'Regan, and Chen Qin. TIP: Tabular-image pre-training for multimodal classification with incomplete data. In *Computer Vision – ECCV 2024*, volume 15073 of *Lecture Notes in Computer Science*, pages 478–496. Springer, 2024.

[63] Zhiqiang Tang, Haoyang Fang, Su Zhou, Taojiannan Yang, Zihan Zhong, Cuixiong Hu, Katrin Kirchhoff, and George Karypis. AutoGluon-Multimodal (AutoMM): Supercharging multimodal AutoML with foundation models. In *Proceedings of the Third International Conference on Automated Machine Learning (AutoML 2024)*, volume 256 of *Proceedings of Machine Learning Research*, pages 15/1–35. PMLR, 2024.

[64] Mihaela Cătălina Stoian, Salijona Dyrmishi, Maxime Cordy, Thomas Lukasiewicz, and Eleonora Giunchiglia. How realistic is your synthetic data? Constraining deep generative models for tabular data. In *Proceedings of the 12th International Conference on Learning Representations (ICLR 2024)*, pages 32167–32204, 2024.

[65] Maximilian-Peter Radtke and Jürgen Bock. Expert knowledge induced Logic Tensor Networks: A bearing fault diagnosis case study. *PHM Society European Conference*, 7(1):421–431, 2022.

[66] Devashish Gaikwad, Wil M. P. van der Aalst, and Gyunam Park. Neuro-symbolic process anomaly detection. In *Advanced Information Systems Engineering (CAiSE 2026)*, volume 16559 of *Lecture Notes in Computer Science*, pages 240–255. Springer, 2026.

[67] Yuyang Gao, Siyi Gu, Junji Jiang, Sungsoo Ray Hong, Dazhou Yu, and Liang Zhao. Going beyond XAI: A systematic survey for explanation-guided learning. *ACM Computing Surveys*, 56(7):188:1–188:39, 2024.

[68] Pang Wei Koh, Thao Nguyen, Yew Siang Tang, Stephen Mussmann, Emma Pierson, Been Kim, and Percy Liang. Concept bottleneck models. In *Proceedings of the 37th International Conference on Machine Learning (ICML 2020)*, volume 119 of *Proceedings of Machine Learning Research*, pages 5338–5348. PMLR, 2020.

[69] Weiyun Wang et al. InternVL3.5: Advancing open-source multimodal models in versatility, reasoning, and efficiency. arXiv:2508.18265, 2025.

[70] Tommaso Carraro, Luciano Serafini, and Fabio Aiolli. LTNtorch: PyTorch implementation of Logic Tensor Networks. arXiv:2409.16045, 2024.

[71] Daniel Raymer. *Aircraft Design: A Conceptual Approach*. American Institute of Aeronautics and Astronautics, Inc., sixth edition, 2018.

[72] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *Proceedings of the 7th International Conference on Learning Representations (ICLR 2019)*, 2019. OpenReview Bkg6RiCqY7.

[73] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR 2016)*, pages 770–778, 2016.

[74] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. ImageNet: A large-scale hierarchical image database. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR 2009)*, pages 248–255, 2009.

[75] Bolei Zhou, Aditya Khosla, Agata Lapedriza, Aude Oliva, and Antonio Torralba. Learning deep features for discriminative localization. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR 2016)*, pages 2921–2929, 2016.

[76] Matthias Minderer, Alexey Gritsenko, and Neil Houlsby. Scaling open-vocabulary object detection. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*, pages 72983–73007, 2023.

[77] Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, Tete Xiao, Spencer Whitehead, Alexander C. Berg, Wan-Yen Lo, Piotr Dollár, and Ross Girshick. Segment anything. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV 2023)*, pages 3992–4003, 2023.

[78] W. J. Youden. Index for rating diagnostic tests. *Cancer*, 3(1):32–35, 1950.

[79] Richard H. Byrd, Peihuang Lu, Jorge Nocedal, and Ciyou Zhu. A limited memory algorithm for bound constrained optimization. *SIAM Journal on Scientific Computing*, 16(5):1190–1208, 1995.

[80] Riccardo Guidotti. Counterfactual explanations and how to find them: Literature review and benchmarking. *Data Mining and Knowledge Discovery*, 38(5):2770–2824, 2024.

[81] Jorge Nocedal and Stephen J. Wright. *Numerical Optimization*. Springer Series in Operations Research and Financial Engineering. Springer, New York, second edition, 2006.

[82] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *Proceedings of the 3rd International Conference on Learning Representations (ICLR 2015)*, San Diego, 2015. arXiv:1412.6980.

[83] Stef van Buuren and Karin Groothuis-Oudshoorn. mice: Multivariate imputation by chained equations in R. *Journal of Statistical Software*, 45(3):1–67, 2011.

[84] Xianjie Chen, Roozbeh Mottaghi, Xiaobai Liu, Sanja Fidler, Raquel Urtasun, and Alan Yuille. Detect what you can: Detecting and representing objects using holistic models and body parts. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR 2014)*, pages 1979–1986, 2014.

[85] Ju He, Shuo Yang, Shaokang Yang, Adam Kortylewski, Xiaoding Yuan, Jie-Neng Chen, Shuai Liu, Cheng Yang, Qihang Yu, and Alan Yuille. PartImageNet: A large, high-quality dataset of parts. In *Computer Vision – ECCV 2022*, volume 13668 of *Lecture Notes in Computer Science*, pages 128–145. Springer, 2022.

[86] Qing Liu, Adam Kortylewski, Zhishuai Zhang, Zizhang Li, Mengqi Guo, Qihao Liu, Xiaoding Yuan, Jiteng Mu, Weichao Qiu, and Alan Yuille. Learning part segmentation through unsupervised domain adaptation from synthetic vehicles. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR 2022)*, pages 19118–19129, 2022.

[87] C. A. Field and A. H. Welsh. Bootstrapping clustered data. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 69(3):369–390, 2007.

[88] Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu. LightGBM: A highly efficient gradient boosting decision tree. In *Advances in Neural Information Processing Systems 30 (NIPS 2017)*, pages 3146–3154, 2017.

[89] Jianming Zhang, Sarah Adel Bargal, Zhe Lin, Jonathan Brandt, Xiaohui Shen, and Stan Sclaroff. Top-down neural attention by excitation backprop. *International Journal of Computer Vision*, 126(10):1084–1102, 2018.

[90] Zhuang Liu, Hanzi Mao, Chao-Yuan Wu, Christoph Feichtenhofer, Trevor Darrell, and Saining Xie. A ConvNet for the 2020s. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR 2022)*, pages 11966–11976, 2022.

[91] Donald B. Rubin. Inference and missing data. *Biometrika*, 63(3):581–592, 1976.

[92] Natalia Neverova, Christian Wolf, Graham Taylor, and Florian Nebout. ModDrop: Adaptive multi-modal gesture recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 38(8):1692–1706, 2016.

[93] Miles Brundage et al. The malicious use of artificial intelligence: Forecasting, prevention, and mitigation. arXiv:1802.07228, 2018.

[94] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. In *Proceedings of the 32nd International Conference on Machine Learning (ICML 2015)*, volume 37 of *Proceedings of Machine Learning Research*, pages 448–456. PMLR, 2015.

A Porting to Other Engineered Systems

This appendix guides the port of NeSySpec to the declared specifications of another engineered system whose fields are linked by physical, definitional or regulatory relations with statable tolerances; Table A.1 lists what the new system supplies, what is fitted again and what is reused.

The aircraft study suggests this ordered checklist.

1. *Scope.* NeSySpec audits the coherence of a declaration with physics and with its evidence, and its report supports human review (Section 4.10).

2. *Fields and relations.* List the fields with their units and one relation per constraint, and let every constraint that reads an undeclared field abstain: on the records as declared, filling the undeclared fields at scoring lowers detection, with global means on every split and with MICE on test and unseen families, while with fields hidden at random MICE at scoring ranks unseen-family records higher (Appendix G.3).

3. *Priors and labels.* Fit initial bands on clean training records only, and fix an intervention tolerance before generation, so that the intervention labels also reach errors that the generator oracle misses: on aircraft, the oracle flags no constraint on 613 of the 1,810 rows labeled as altered (33.9%; Appendix G.4).

Table A.1: Porting to a new engineered system.

Component	Content	Aircraft study
Supplied by the new system		
Knowledge base	one exact, differentiable quantity per relation, each with a graded predicate, an availability condition and a scope; bounded latents for rarely declared quantities; initial bands fitted on clean training records; regulatory bands frozen	12 constraints on records and 10 on satellite annotations (Table 3); 336 trainable parameters on records, 126 of them latents, and 57 on satellite annotations; no neural network
Visual predicates	one sigmoid head per class, with the parts where its evidence must lie; cross-modal axioms with the same form	25 heads per dataset (Table E.1)
Generator and data	corruptions that break one relation or swap a block of fields; noisy-but-coherent controls; an oracle with per-constraint labels; an intervention tolerance fixed before generation; sourced records with units, split by family or site	five corruption families, controls and intervention labels; 107 records of 71 families (Appendix D)
Fitted again		
Bands and visual modules	tolerances learned from the training fleet; visual heads specific to their sensor	45.7% of the learned bands with both edges inside the widened training fleet range (Appendix G.9); with a visual stratum trained on photographs, the fused verdict reaches an AUROC of 0.724 on satellite pairs, against 0.796 for the record verdict alone (Appendix G.5)
Reused as Section 3 defines it		
Verdict and report	masked verdict with slopes and bands learned per derived quantity, and its exact reduction for any knowledge base; a diagnosis key read from its degrees and a monotone second head; abduction, repair proposals and the report; stratified aggregation; grounding axiom with gate, guard and floor	Sections 3.2 to 3.7; Appendix B
Protocol	priors built without held-out aircraft, intervention labels, black-box falsification within an edit budget, cluster-bootstrap intervals, comparators grouped by the knowledge they receive (Table F.1), statistical imputers among them	Section 4.1; Appendix F

4. *Masks and sensors.* Measure automatic part masks against public human masks, and train visual heads on images of the sensor in use (Sections 4.5 and 4.6; Appendix G.5).

5. *Comparators.* Include learners on the same physical residuals: on aircraft records a multilayer perceptron on them serves as the alarm of the optional rank-then-explain setting (Section 4.4).

6. *Critical fields.* NeSySpec abstains on every constraint that reads an undeclared field; on aircraft, NeSySpec's optional critical-field alert brings the evasion of lies that withhold fields back to the level of bare lies, which withhold none (Section 4.7).

B Implementation and Verification Report

This appendix summarizes how the implementation of NeSySpec is organized and shows an example of a verification report. The implementation follows the seven modules of Section 3. Every prediction file, of NeSySpec and of each comparator, has one format (instance identifier, split and score, with per-constraint outputs where the method returns them), so a single evaluation script scores every method in the same way.

B.1 Example of a Verification Report

Figure 5 (Section 4.2) shows a report assembled from the saved outputs of NeSySpec (train prior, seed 0) for a record chosen by a rule fixed before any value was read: the first single-field corruption of the test split, in file order, whose altered field an available constraint reads. The corruption divides the total thrust of a narrow-body turbofan by 1.2, and the generator oracle flags A4 and A9. The verdict flags the record ($V = 1.217$, above the validation threshold $t^\star = 0.310$ of this seed), and the report lists the degree of every available constraint, the largest being A4 (10.001), A5 (1.053), A8 (0.386) and A9 (0.384). By the rule of Section 3.7, A4 and A5 exceed $\ln 2$ and are violated, A5 also on the clean

record, while A8 and A9 are held. The largest-violation key names A4, which the oracle flags; it ranks A4 first in every seed, followed by A5 in this seed and by A8 in 5 of the 8 seeds.

C Physical Relations and Constants

The specification axioms evaluate closed-form relations of conceptual aircraft design [71] (Table 3). Table C.1 gives the constants of the relations and of the International Standard Atmosphere (ISO 2533:1975) with the initial values of the graded degrees, and Table C.2 the remote-sensing constants with the Airplane Design Groups of the Federal Aviation Administration (FAA)[3]. The train prior widens each per-class range of a derived quantity over the 44 clean training aircraft to [0.88 min, 1.12 max], as the oracle does; these initial bands are then learned with their slopes (Section 3.2), and business jets, absent from training, start from the band of all training classes. The learned bands act as tolerances learned from the fleet. On specification records, 45.7% of them keep both edges inside the training fleet range widened in the same way, and the others extend beyond it (Appendix G.9). In the remote-sensing stratum, 103 of the 144 learned group bands (the 18 bands of R2 for military roles, R3 and R4 in each of the 8 seeds) extend beyond the clean training range widened on each side by 12% of its width. The initial bands of R2 for military roles, R3 and R4 (Table C.2) were set from the annotations and from knowledge of the aircraft types to contain every clean annotation, those of test and unseen sites included, and the generator oracle of the satellite benchmark was fitted on the clean annotations of every site, so both may carry held-out information; the satellite benchmark was not regenerated, nor the remote-sensing stratum retrained, without them.

D Data and Benchmark Details

This appendix describes the specification records and the benchmarks built from the records and images of Section 3.1: the corruption operators with their realized magnitudes (Table D.1), the realized instances per type and split (Appendix D.1), the frame of the cross-dataset runs and the rule that chose the examples of Figure 8. Every generator is deterministic given its seed. The 107 specification records cover seven classes: wide-body (17 aircraft of 8 families), narrow-body (18 of 6), regional jet (13 of 6), regional turboprop (11 of 8), business jet (19 of 17, the unseen class), general-aviation turboprop (12 of 11) and general-aviation piston (17 of 15). They were compiled from the type-certificate data sheets of the aviation authorities of the United States, the European Union, Canada, the United Kingdom and France and from manufacturer documents (airport-planning manuals, specification sheets and flight-manual excerpts), with other public sources used for cross-checks. The 67 jets declare thrust and the 40 propeller aircraft shaft power, and 18 records, all general-aviation aircraft, declare the stall speed; every record passed a loose screening with the verifier's physics module, so the clean records were checked against the relations that the verifier evaluates; the intervention labels, set from the edit alone, bypass this dependence. Table D.1 lists the corruption operators. A candidate that the oracle does not flag is generated again with its magnitude doubled, for at most six attempts (Section 3.1), so the realized magnitudes reflect the oracle tolerances and extend beyond the declared grid of 5 and 10%. They extend furthest for performance corruptions, whose Breguet tolerance must absorb the largest residual among the clean records.

D.1 Benchmarks

Table D.2 gives the realized instances per type and split of four benchmarks: the specification records, the photographic mismatch pairs, the satellite annotations and the satellite multimodal pairs. Of the 100 FGVC-Aircraft variants, 68 are paired with a specification record, 39 with the record of their type and 29 with that of their family; their 6,800 photographs follow the split of the paired family (2,500 training, 1,300 validation, 800 test, 1,500 unseen-family and 700 unseen-class photographs), and the 32 unpaired variants enter no split. The satellite annotations start from 3,000 clean RarePlanes annotations, a sample stratified by role with a floor of 200 annotations per role (all annotations of

[3]FAA Advisory Circular 150/5300-13B, Airport Design, March 31, 2022.

Table C.1: Constants of the physical relations.

Standard atmosphere		Relations and graded degrees	
g	$9.80665\ \mathrm{m\,s^{-2}}$	units	SI, except R (km), F_{tot} (kN) and P_{tot} (kW)
R_{air}; γ_{air}	$287.05287\ \mathrm{J\,kg^{-1}\,K^{-1}}$; 1.4	A5	$W_1 = \max\{MTOW - 0.94\,m_{\mathrm{fuel}},\ 0.05\,MTOW\}$, a 6% fuel reserve; log residual clipped to $[-3, 3]$; available only when AR is also declared and, for propeller aircraft, V_{cr}
T_0; ρ_0	288.15 K; $1.225\ \mathrm{kg\,m^{-3}}$	A6	log residual clipped to $[-2, 2]$; wide $C_{L,\max}$ bounds, because published stall speeds mix clean and landing configurations
Γ	$6.5 \times 10^{-3}\ \mathrm{K\,m^{-1}}$ (troposphere)	A7	95 kg per passenger with baggage; $f_u \in [0.3, 1]$
h_{tr}; T_{tr}	11,000 m; 216.65 K, isothermal above; density continuous	unused	an estimate of maximum L/D from AR (Oswald factor 0.80, zero-lift drag coefficient 0.020) is implemented but read by no axiom; A5 takes L/D from its latent
$\bar{h}$	h_{cr} clamped to [0, 20,000] m	coordinates	$z = (\log_{10} q - m_g)/\sigma_g$ for a band, $z = \xi/\sigma_g$ for the residual ξ of A1, A5, A6 and A7; m_g, $\sigma_g \geq 10^{-3}$ from the clean training rows of group g (σ_g a root-mean-square around 0 for ξ), shrunk toward the fleet with weight $n_g/(n_g + 5)$, the scatter through its square; a group without rows takes the fleet values
$a(\bar{h})$	$\sqrt{\gamma_{\mathrm{air}} R_{\mathrm{air}} T(\bar{h})}$	initial values	tolerances τ/σ_g with $\tau = 0.05$ (A1), 0.25 (A5) and 0.15 (A6); slope $\kappa = 8/\tilde{w}$, with $\tilde{w}$ the median initial band width in standard units (A7: $\kappa = 25\,\sigma$, with σ the fleet scatter of its margin, and lower edge at 0); in-band weight 0.05
		latent bounds	sigmoid of one parameter per pair, rescaled to these bounds and starting at their midpoint; L/D: 15-21 (wide-body), 14-19 (narrow-body), 12-17 (regional jet), 11-17 (regional turboprop, business jet), 9-15 (general-aviation turboprop), 7-14 (general-aviation piston); $C_{L,\max}$: 2.2-3.2 (wide-body, narrow-body), 2.0-3.0 (regional jet), 1.9-3.0 (regional turboprop), 1.6-2.8 (business jet), 1.4-2.4 (general-aviation turboprop), 1.2-2.2 (general-aviation piston); c_T (turbofan): 1.2×10^{-5} to $2.3 \times 10^{-5}\ \mathrm{kg\,N^{-1}\,s^{-1}}$; c_P: 7.5×10^{-8} to 1.1×10^{-7} (turboprop), 7.0×10^{-8} to 1.0×10^{-7} (piston) $\mathrm{kg\,W^{-1}\,s^{-1}}$; η_p: 0.75-0.88
		degrees	capped at 10^4; truth $\exp(-v)$ never clamped; a masked axiom contributes no value and no gradient; $V = \ln 2$ without an available axiom

g: standard gravity; R_{air}, γ_{air}: gas constant and ratio of specific heats of air; T_0, ρ_0: sea-level temperature and density; Γ: temperature lapse rate; h_{tr}, T_{tr}: altitude and temperature of the tropopause; $\bar{h}$: altitude at which the atmosphere is evaluated; a: speed of sound; q: derived quantity of an axiom; n_g: number of clean training rows of group g.

Table C.2: Constants of the remote-sensing axioms.

Design group	Wingspan (m)
I	below 14.9352
II	14.9352-24.0792
III	24.0792-35.9664
IV	35.9664-52.1208
V	52.1208-65.2272
VI	65.2272-79.8576

Role	R2 span (m), initial	R3 L/b, initial	R4 n_{eng}, initial	R6 P(jet), initial
Civil small	0-15.005 (frozen)	0.55-1.70	0-2	0.25
Civil medium	15.005-36.0 (frozen)	0.70-1.45	1-4	0.55
Civil large	36.0-90.0 (frozen)	0.75-1.35	2-4	0.98
Military transport	7-75	0.70-1.60	1-4	0.55
Military bomber	7-60	0.60-2.20	2-4	0.90
Military fighter	5-20	0.90-2.70	1-2	0.98
Military trainer	5-15	0.80-2.00	1-2	0.60

Initial tolerance 0.3 (R5) and 0.10 (R10); weight of R9 0.5; frozen: R1 and civil R2 bands
Graded slots R1, R2 (civil and military), R3, R4 and R10; 57 trained parameters

a rarer role), and each operator draws one candidate from each of them, so the kept and discarded candidates of a type sum to 3,000. In the cross-dataset runs of Appendix G.5, a knowledge base fitted on one source scores the other through a 23-field unified frame (Figure 2(c)) with only the axioms R2 to R4, which both sources can evaluate; a specification record takes its role and design group from its clean wingspan. On the records of each split, the report prints the class of the nearest clean training record (Section 3.7; example in Figure 5): this proximity names the class of a record from an unseen family and attributes each business jet, a class absent from training, to its nearest training class.

Table D.1: Corruption types and magnitudes.

Operator	Fields altered	Kept if the oracle flags	Kept	Disc.	Realized magnitude ϵ 5%	10%	20%	40%	80%	160%	320%
Single value	one of S, F_{tot} or P_{tot}, OEW, V_{cr}	the targeted envelope (A3, A4 or A4P, A8, A10)	212	2	0	0	53	105	46	8	0
Geometric	one of b, S, AR	A1	208	6	104	104	0	0	0	0	0
Performance	R or m_{fuel}	A5	189	25	1	2	13	35	55	66	17
Cruise	one of V_{cr}, h_{cr}, S	A2	206	8	0	8	59	85	46	7	1
Frankenstein	one block of a record of another class	any constraint	176	38				-			
Noisy-but-coherent control	every declared non-integer field (relative standard deviation 0.02); AR re-derived	no constraint ($y = 0$)	214	0				-			

Constraints violated per corrupted instance	1	2	3	4-6	Total
Corrupted instances	553	253	125	60	991

Kept: instances kept; Disc.: candidates discarded, because the operator does not apply to the record (a missing field, or for a Frankenstein swap no record of another class in the same split) or because the condition of the third column fails in all six attempts. Each operator draws two candidates per clean record, 214 in all; the 38 discarded Frankenstein swaps are those of the unseen class, a single class. Constraints violated: flagged by the generator oracle. Disc. has the same meaning in Table D.2.

Examples of Figure 8. The figure shows, for seed 0, the network trained with the grounding axiom and the network trained without a localization term. A rule fixed in advance chose its examples: each is the first candidate on which both networks give a non-degenerate map (maximum above 0) and the axiom-trained network predicts the target class with an in-mask share ρ of at least 0.15. Photographs for engine count, propulsion and tail configuration are taken in identifier order from validation and test, and from unseen families only when validation and test offer no such candidate. Satellite crops of civil aircraft, for the role small civil aircraft and for an engine count of 2, are taken in identifier order from validation and test sites, the second from a site other than the first's, with the annotated diamond as region. Both networks are drawn at their selected checkpoints; on photographs the network without a localization term is at its first epoch. Pointing is higher with the axiom in every seed pair at the selected and at the last checkpoints, and the automatic pseudo-masks include some background or off-part regions, visible in the figure. The examples illustrate the maps; pointing measures localization (Section 4.6).

Table D.2: Benchmarks.

Type	Label	Train	Val.	Test	Unseen fam.	Unseen class	Total
Specification records							
Clean	coherent	44	17	16	11	19	107
Single value	incoherent	86	34	32	22	38	212
Geometric	incoherent	88	32	32	22	34	208
Performance	incoherent	81	29	28	22	29	189
Cruise	incoherent	86	32	32	22	34	206
Frankenstein	incoherent	88	34	32	22	0	176
Noisy-but-coherent control	coherent	88	34	32	22	38	214
Total		561	212	204	143	192	1,312
Photographic mismatch pairs							
Matched	coherent	788	499	412	250	51	2,000
Inter-class	mismatched	227	108	70	138	0	543
Intra-family	mismatched	227	168	137	68	0	600
Attribute	mismatched	232	158	132	63	15	600
Occlusion	coherent control	161	115	69	52	3	400
Degradation	coherent control	163	87	88	50	12	400
Total		1,798	1,135	908	621	81	4,543

Type	Label	Train	Val.	Test	Unseen sites	Disc.
Satellite annotations						
Clean	coherent	1,903	262	357	478	-
Wingspan rescaling	incoherent	1,903	262	357	478	0
Length rescaling	incoherent	1,139	149	224	306	1,182
Design-group change	incoherent	1,903	262	357	478	0
Role change	incoherent	1,903	262	357	478	0
Engine-count swap	incoherent	1,233	170	254	415	928
Propulsion change	incoherent	1,890	256	357	478	19
Frankenstein	incoherent	1,899	262	356	478	5
Noisy-but-coherent control	coherent	1,903	262	357	478	0
Total		15,676	2,147	2,976	4,067	2,134
Satellite multimodal pairs						Total
Clean	coherent	950	144	171	235	1,500
Tabular corruption	incoherent	956	152	162	230	1,500
Image swap	incoherent	971	134	170	225	1,500
Geometric corruption	incoherent	483	75	80	112	750
Total		3,360	505	583	802	5,250

Photographic pairs: matched, a photograph with its own record; inter-class, with the record of an aircraft of another class; intra-family, with the record of another variant of its family; attribute, with its own record altered in engine count, engine type or maximum take-off mass; occlusion and degradation, with its own record and a part of the photograph occluded or the photograph degraded (blur, compression or low light). Satellite multimodal pairs: clean, an annotation with its own crop; tabular corruption, the annotation altered in engine count, propulsion or role, with its own crop; image swap, the clean annotation with the crop of an aircraft of another role from the same split; geometric corruption, the annotation with its wingspan or length rescaled. Disc.: as in Table D.1; in the multimodal panel the last column gives the totals.

E Method Details

This appendix completes Section 3. Table E.1 adds the visual predicates and the cross-modal axioms to the catalog of Table 3; Appendix E.1 proves Proposition 1 and reports its numerical check; Figures E.1 and E.2 zoom into the verification report and the vision module; Appendix E.2 defines the localization terms of the matched grounding study; and Algorithm E.1 transcribes the forward pass for one declared record.

E.1 Proof of the Exact Reduction

Proof. (a) Fix α. A masked axiom has $\alpha_k = 0$, and the implementation sets its degree (for a cross-modal axiom, its truth) to 0 by selection before any sum; its term and its weight therefore vanish from the numerator and from the denominator of Eq. (9), which reduce to the sums over $\mathcal{S}'$. If $\mathcal{S}'$ is empty, a tabular stratum takes $V_{\mathcal{S}} = \ln 2$ by definition and the cross-modal stratum has $\gamma_{\mathcal{AX}} = 0$. Every degree and its derivatives are finite, because undeclared fields carry a finite placeholder and degrees are capped at 10^4, so the selection passes no value and no gradient from the degree v_k, or the truth sat_k, of a masked axiom. The stratum verdict therefore has zero derivative with respect to them, and a field read only by masked axioms does not affect it. (b) Without an image, no cross-modal axiom

Table E.1: Axiom catalog: visual predicates and cross-modal axioms.

Predicate	Classes (one sigmoid head each)	Part region	Read by
Photographs (8 predicates, 25 heads; regions from the pseudo-masks)			
EngineCount	1, 2, 3, 4	engine	AX1
EnginePlacement	underwing, rear, overwing	engine	AX4
Propulsion	jet, turboprop, piston	propeller and engine	AX2
WingPosition	high, mid, low	wing	AX5 (not implemented)
TailConfig	conventional, T-tail, cruciform, V-tail	tail	AX4
WingletPresence	present (binary)	none (conjuncts masked)	none
MainGearWheelCount	2, 4, 6 to 10, 12 or more (no attribute labels)	landing gear	AX3
FuselageSlenderness	stubby, standard, slender (no attribute labels)	none (conjuncts masked)	none
Satellite crops (7 predicates, 25 heads; region of every head: the annotated diamond)			
role	civil small, medium, large; military transport, bomber, fighter, trainer	diamond	AX4, AX5
wing type	straight, swept, delta, variable-swept	diamond	AX3
wing position	mid or low, high	diamond	none
propulsion	jet, propeller, unpowered	diamond	AX2
engine count	0, 1, 2, 3, 4	diamond	AX1
canards	no, yes	diamond	none
tail fins	1, 2	diamond	none

ID	Photographs: truth (availability)	Satellite crops: truth (availability)
AX1	$p_{\text{vis}}(n_{\text{eng}} \mid I)$ (declared $n_{\text{eng}} \in \{1, \dots, 4\}$; engine visible)	$p_{\text{vis}}(n_{\text{eng}} \mid I)$ (declared $n_{\text{eng}} \in \{0, \dots, 4\}$)
AX2	$p_{\text{vis}}(\hat{e} \mid I)\,\big(1 - p_{\text{vis}}(\text{propeller} \mid I)\,\sigma(25(M - 0.62))\big)$, each factor entering only when its declared fields are present (engine type, or V_{cr} and h_{cr}, declared; propeller or engine visible)	$p_{\text{vis}}(\pi \mid I)$ (declared π)
AX3	$\sum_j p_{\text{vis}}(j \mid I)\,\mu_j(\log_{10} MTOW)$, band edges at $\log_{10} MTOW = 4$, 5.05 and 5.5, slope 16 ($MTOW$ declared; landing gear visible)	$p_{\text{vis}}(t \mid I)$ (declared t)
AX4	$I_R\big(p_{\text{vis}}(\text{rear} \mid I),\, p_{\text{vis}}(\text{T-tail} \mid I) + p_{\text{vis}}(\text{cruciform} \mid I)\big)$ (engine and tail visible; no declared field)	$p_{\text{vis}}(r \mid I)$ (declared r)
AX5	wing position against class: planned, not implemented	$1 - \sum_{r'} p_{\text{vis}}(r' \mid I)\,\iota_{r'}$, where $\iota_{r'} = 1$ when r' is a civil role whose wingspan band does not overlap the band of the declared design group; weight 0.5 (declared civil r and design group)

$p_{\text{vis}}(\cdot \mid I)$: posterior of a predicate, the softmax over its heads (a sigmoid for a binary predicate); $\hat{e}$: visual class of the declared engine type (turbofan to jet); propeller: turboprop or piston; M: cruise Mach number $V_{\text{cr}}/a(h_{\text{cr}})$ of the declared fields, as in A10; μ_j: soft membership of the declared mass in the wheel band j; I_R: Reichenbach implication; $[\cdot]$: indicator. Every head of a satellite crop shares the diamond region, visible whenever it is annotated. The photographic score reads the cross-modal axioms alone; the satellite verdict combines them with $\mathcal{R}$ (Section 3.5). Settings: Table F.2.

is available. Then, as when every cross-modal axiom is masked, $\gamma_{\mathcal{AX}} = 0$, and since $\lambda_{\mathcal{T}} > 0$, Eq. (10) reduces to $\text{feas}_{\mathcal{T}} = \exp(-V_{\mathcal{T}})$; the implementation returns $\text{feas}_{\mathcal{T}}$ itself in this case, so the equality holds bit for bit. Otherwise, the cross-modal stratum enters both sums of Eq. (10) with weight $\lambda_{\mathcal{AX}}\gamma_{\mathcal{AX}}$ and the tabular stratum keeps $\lambda_{\mathcal{T}}$, which is the stated renormalization over the strata present. □

In float32, the masked verdict equals the verdict of the knowledge base restricted to the available constraints bit for bit in 16 of 16 runs (Figure 7a), eight seeds under each of the two tolerance priors: the train prior and the textbook prior, which takes handbook envelopes in place of the bands fitted on the training fleet (Table G.7, note f). It agrees with a float64 recomputation up to float32 rounding, within 2.2×10^{-7} relative, and the reduction without an image also holds in all 180 runs scored with newly drawn masks that hide a nominal 30% of the satellite images.

E.2 Method Components

Figure E.1 zooms into the module that assembles the verification report and the optional settings of Section 3.7.

Figure E.2 zooms on the vision module and the grounding axiom of Section 3.3.

Table F.2 gives the settings of the components. The matched grounding study isolates the gate of the grounding axiom: on photographs, it trains localization terms $\lambda\langle g_{ik}(1 - \text{Loc}_k(I_i))\rangle_{\text{vis}}$ with different gates g_{ik}, at the loss weights λ of 0.3, 1 and 3 used in Section 4.6, with three runs per weight. With

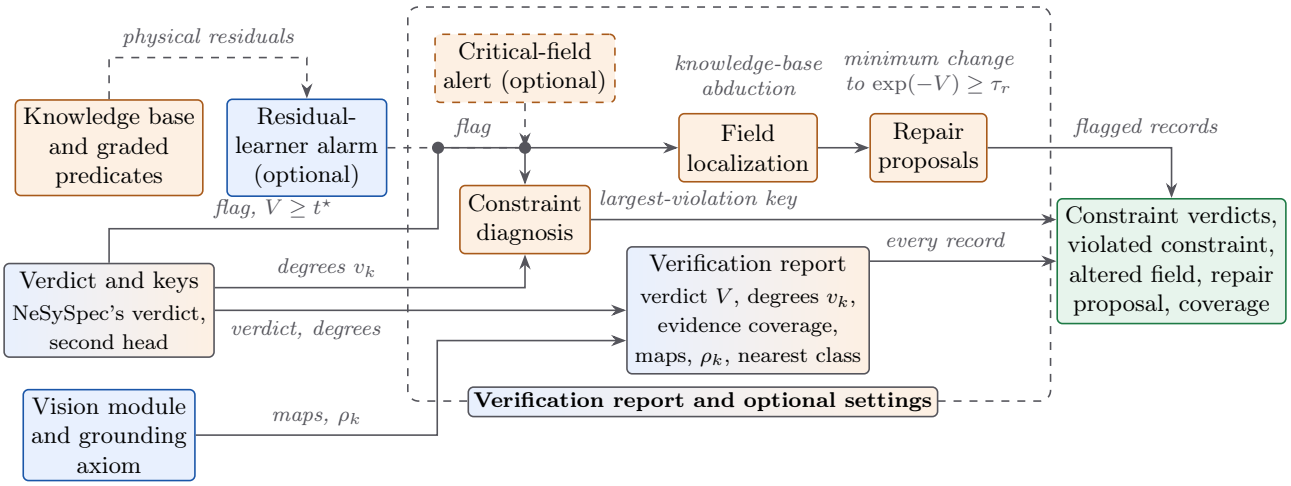

Figure E.1: Verification report and optional settings.

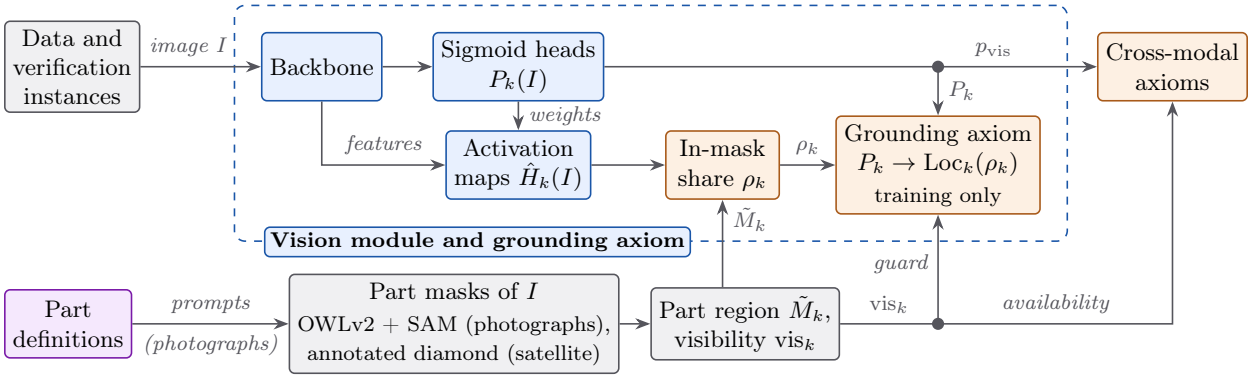

Figure E.2: Vision module and grounding axiom.

$g_{ik} = P_k(I_i)$ the term is the grounding axiom; with g_{ik} equal to the positive label of a supervised head and the normalizer of the axiom, only the gate changes (gate only); with $g_{ik} = 1$ the gate is removed (ungated term). A fourth variant gates by the label and averages over the visible label-positive heads, which changes the normalizer together with the gate. The auxiliary loss of Section 3.3 is trained at the same weights as a reference; Section 4.6 and Appendix H report its comparison with the axiom.

On photographs, the engine-count and engine-type axioms AX1 and AX2 carry the instance-specific signal, since the gear heads read by AX3 have no attribute label and AX4 reads no declared field; a pair without an available axiom scores 0. On satellite crops, AX5 is available on 95.2% of the pairs.

Algorithm E.1 transcribes from the code the forward pass of Figure 1 for one declared record.

F Protocol Details

This appendix completes the setup of Section 4.1: it lists the comparators and the configuration, states how the configuration was selected and which choices used data beyond the training split, and reports how the vision results change when the part detector's visibility threshold is re-selected on validation data alone. Table F.1 orders the comparators by the knowledge they receive; each shares the partitions, the labels and the evaluation code of NeSySpec, with settings selected on validation data within a budget fixed before its runs, and one planned arm, the language model in its thinking mode, has no result: in a pilot run too few of its answers could be parsed, below the parse rate fixed before the runs. Table F.2 gives the configuration of every module of NeSySpec, of its optional settings and of the baseline detectors; bands, slopes, latents, visual heads, read-outs, operating thresholds and checkpoints are fitted or selected on data, and every other value was set once. The class bands of the generator oracle are fitted on all 107 aircraft, test and out-of-distribution aircraft included, and reach training through the generator labels and per-constraint targets (Section 3.1); Appendix G.9 refits the

Algorithm E.1: NeSySpec verification of a declared record.

Input: x, m, optional image I, stratum $\mathcal{T} \in \{\mathcal{A}, \mathcal{R}\}$, threshold $t^\star$
Output: verification report

foreach *axiom* $k \in \mathcal{T}$ **do** availability α_k, degree v_k by Eq. (3) ($-\ln$ of its truth for R5 to R9), truth $\mathrm{sat}_k = \exp(-v_k)$
$V_\mathcal{T} \leftarrow$ Eq. (9), or $\ln 2$ without available axioms; $\quad \mathrm{feas} \leftarrow \exp(-V_\mathcal{T})$
if $I \neq \varnothing$ **then**
  foreach *axiom* $j \in \mathcal{AX}$ **do** availability α_j, truth sat_j
  if $\gamma_{\mathcal{AX}} = 1$ *and* I *is a satellite crop* ($\mathcal{T} = \mathcal{R}$) **then** feas $\leftarrow$ Eq. (10)
  // *photographs: score* $1 - \mathrm{feas}_{\mathcal{AX}}$*, or 0 if* $\gamma_{\mathcal{AX}} = 0$
$V \leftarrow -\ln \mathrm{feas}$; $\quad$ flagged $\leftarrow [V \geq t^\star]$; $\quad e_k \leftarrow \alpha_k v_k$; $\quad \hat{p}_\mathrm{m}$ and c_k by Eq. (11)
report $\leftarrow$ degrees, truths, coverage, maps and the nearest clean training record
if *a specification record with at least one available constraint is flagged by its verdict or by an optional setting (the residual-learner alarm of rank-then-explain or the critical-field alert)* **then**
  diagnosis $\arg\max_{k:\,\alpha_k=1} v_k$ (secondary key $\arg\max_k c_k$); $\quad$ field ranking by Eq. (12); $\quad$ repair by Eq. (13)
return *report*

oracle without test or out-of-distribution aircraft.

Selection of the configuration. NeSySpec's configuration was selected on validation data and on held-out training classes; during development, earlier configurations were also evaluated on the test splits. Validation data selected the configuration of both tabular strata, the largest-violation key, the abduction and repair rules and the optional settings, and held-out training classes also selected the key and the critical-field alert (Appendix G.9). The rule that reads each constraint as held or violated has no tuned parameter. The part detector's visibility threshold was moved from 0.25 to 0.20 on a 240-image stratified subset including test and out-of-distribution images, by an informal trade-off, with no numeric criterion and no saved tuning output. The results of this paper use 0.20.

Visibility threshold re-selected on validation. Re-selected on 240 stratified validation images alone, under a criterion declared beforehand that keeps propeller detections on jets, which have no propeller, at or below 10% of jet images, the threshold becomes 0.25, as in 197 of 200 resampled subsets: at 0.20 those detections reach 17% on validation, against 7% on the original subset. Weighing part visibility and those detections equally keeps 0.20. Without retraining, 0.25 turns 12.1% of the part flags of training, validation and test images to not visible, 2,794 flags, 1,618 of them on the wing.

Vision module retrained at 0.25. Retrained at 0.25 with the same code, configuration and 8 seeds (even seeds on the RTX 5090, which trained the original module, and odd seeds on the RTX 4090; the even seeds alone give the same accuracy and pointing readings), the photograph module keeps its macro head accuracy (0.916 on test and 0.944 on unseen families, as at 0.20) and at most one collapsed head per seed on test. It also keeps its pointing lead over the network without a localization term (+0.671) and with permuted masks (+0.414), its lead over RRR, GradMask and Pukdee et al. on PASCAL-Part (+0.568, +0.676 and +0.469, every interval above zero) and the nine overall readings of the cross-modal stratum against Qwen3-VL-8B, InternVL3.5-8B and Qwen3-VL-32B on test, unseen families and the unseen class. Of the 21 readings by pair type, 4 change by one step under each of the two bootstrap streams of that comparison, 3 of them to the disadvantage of the stratum. The networks without a localization term or with permuted masks, RRR, GradMask and Pukdee et al. were trained with the flags at 0.20 and not retrained. At 0.25 the cross-modal AUROC is 0.693 on test (0.692 at 0.20) and slightly lower on unseen families (0.716 against 0.733) and on the unseen class (0.827 against 0.844); the axiom on engine placement and tail configuration (AX4) is available on 0.784 of the test pairs scored for detection, instead of 0.956. Pointing on pseudo-masks is slightly lower on test (0.774 against 0.780) and on unseen families (0.778 against 0.784): there the seed-bootstrap and the image-and-seed bootstrap intervals include zero, while the intervals over paired seeds, $[-0.010, -0.001]$ and $[-0.010, -0.004]$, do not. On PASCAL-Part human masks pointing is slightly lower with 3 runs (0.711 against 0.737, -0.026 $[-0.051, -0.002]$) and lower over 8 seeds (0.708 against 0.729, -0.020

[−0.037, −0.005]). The auxiliary loss at weight 0.3, retrained at 0.25 with its 3 seeds, falls on PASCAL-Part as the axiom does (0.708 against 0.739, −0.031 [−0.055, −0.005]). With both at 0.25 the two stay on par in localization on test pseudo-masks (−0.009), out of distribution (−0.027) and on PASCAL-Part (+0.003 [−0.028, +0.032]), while the axiom keeps its out-of-distribution accuracy lead (+0.016 [+0.010, +0.021]). With all 8 seeds of the axiom, the axiom's pointing on test pseudo-masks is slightly lower than that of the auxiliary loss at 0.25 (−0.019) and lower at 0.20 (−0.026). A first check against the auxiliary loss still at 0.20, added after the results, had shown the axiom lower on PASCAL-Part (−0.028); the like-for-like comparison attributes that difference to the threshold.

G Additional Results and Figures

The evaluation at a glance. NeSySpec is tested against 14 groups of comparators (Table F.1), from knowledge-free detectors, residual learners and tabular foundation models to solvers, constraint losses and prompted language and vision-language models, on six benchmarks, four of them labeled independently of the generator oracle (Table 4). Specification records carry two label families, generator and intervention labels, and are evaluated on three splits (test, unseen families and the unseen class), satellite pairs on two (test and unseen sites); each stratum of NeSySpec is trained with eight seeds (seeds 0 to 7, Table F.2) and the satellite backbone grid with three seeds per backbone. Unless the text or a table note names another scheme, every interval on specification records rests on 5,000 cluster-bootstrap resamples at each of the two cluster levels (seed 20261002); the removals of Table G.9 use 2,000, and the curve over the number of constraints in Appendix G.8 resamples the subsets of each size and the seeds of each drawn subset (5,000 draws). Tables G.1 to G.4 print 168 contrasts with both intervals, and Table G.8 prints 41, of which 36 are not in Table G.3. Every primary diagnosis contrast of the largest-violation key keeps its reading in all 20 alternative bootstrap streams, except two on intervention labels, which keep it in 19 and 17 of the 20. So do 11 ablation contrasts, the unseen-class contrasts with half the fields hidden, also under each of 10 further bootstrap seeds of that analysis, and every primary localization and repair contrast except two on the record repair error.

This appendix gives the complete levels behind Section 4, the audit tasks first: constraint verdicts, diagnosis, localization and repair, and missing evidence, then record coherence, records with images, grounding, falsification and the ablations; it ends with the selected components and the sensitivity analyses, including the generator oracles refitted without test or out-of-distribution aircraft (Appendix G.9), and with the training checkpoints (Appendix G.10). Tables G.1 to G.4 give, for the comparisons of Sections 4.2 to 4.5 and 4.7, the difference with its intervals, written $[a, b] \cdot [c, d]$, [aircraft] · [family] on specification records; those of Section 4.6 are in Appendices G.6 and H.

G.1 Constraint Verdicts

NeSySpec's constraint verdicts reach a macro-F1 of 0.795 on generator-label test and, without retraining, 0.575 on intervention-label test (Table G.5), from a rule fixed by the semantics (a degree above $\ln 2$) with no tuned parameter; comparator thresholds were chosen on validation rows only. Against the network trained with the DL2 loss on the same residuals and per-constraint targets, their macro-F1 is higher in all six split and label-family cells, with every interval above zero: on test by +0.134 [+0.086, +0.242] · [+0.099, +0.264] with generator and +0.149 [+0.090, +0.268] · [+0.095, +0.276] with intervention labels, on unseen families by +0.243 and +0.250 and on the unseen class by +0.164 and +0.121. In all six cells they also find every labeled violation of a corrupted record more often (on generator-label test 0.934 against 0.522, +0.413 [+0.367, +0.460] · [+0.371, +0.510]) and reach a higher micro recall (0.959 against 0.635 on generator-label test). Their micro precision is on par on test and lower on the unseen class (0.458 against 0.577 on generator and 0.517 against 0.639 on intervention labels), and their exact-set agreement is higher on unseen families (0.721 against 0.508 and 0.373 against 0.221) and on par on test and the unseen class. Over the 8 seeds, NeSySpec's generator-label test macro-F1 spans 0.767 to 0.808 (sd 0.013) and its exact-set agreement 0.446 to 0.549, the network's macro-F1 0.442 to 0.838; the intervals average the seeds. The network reports fewer false violations on clean-labeled and control records on test and on the unseen class (0.020 against 0.074 on generator-label test, 0.041 against 0.125 on the unseen class), and the crisp edits are on par in

Table F.1: Comparators.

Knowledge	Comparators	Input	Output	Data
None	Isolation Forest, Mahalanobis, MLP on normalized fields	declared fields	error score	records
	LightGBM, XGBoost, CatBoost, TabPFN v2 on raw fields	declared fields, availability mask, one-hot categories	error score	records
Crisp rules	train-fitted rule checker; rules with hierarchical pooling	band exceedances of the residuals	score, violated rules	records
Classical diagnosis	Fellegi-Holt, first unsatisfiable core of Z3, MaxSMT	crisp constraints	violated set, minimum-cardinality localization	records
Sizing statistics	Class I sizing regressions on maximum take-off mass, fitted on the training aircraft	declared fields	largest class-I residual, its field	records
Residual learners	MLP, LightGBM, XGBoost, CatBoost, logistic regression, TabPFN v2, TabICL with its v2 checkpoint, local outlier factor, one-class SVM, autoencoder; class-standardized residual and its Ochiai ranking; SHAP attributions	physical residuals with availability	score, attribution	records
Differentiable logic	DL2 method (network on declared fields, projected-gradient-descent (PGD) adversary); network with the DL2 loss and the per-constraint targets on the physical residuals	declared fields; physical residuals with availability	score, per-constraint probabilities	records
Knowledge in words	Qwen3-8B with the constraints A1-A11	anonymized record	verdict, score, first violated constraint	records
Localization losses	auxiliary localization loss and the matched ablations of the grounding axiom	photograph, pseudo-masks in training	attribute heads, evidence maps	photographs
Explanation supervision	RRR, GradMask, RES-L, Pukdee et al.	photograph, part masks in training	attribute heads, evidence maps	photographs, human masks
Sequential pipeline	cross-modal stratum (vision trained alone, axioms on its posteriors); independent concept-bottleneck model with LightGBM	photograph or crop with the record	mismatch score	photographs, satellite pairs
Late fusion	logistic regression, LightGBM, CatBoost, TabPFN on the record alone, on the record with the visual posteriors, and on these with the image embeddings or with the knowledge-base truths; best fusion; monotone read-out	record, visual posteriors, embeddings, knowledge-base truths	pair score	photographs, satellite pairs
End-to-end fusion	AutoGluon-MM 1.6.3 (out of the box), TIP-style (re-implementation)	raw crop and declared fields	pair score	satellite pairs
Vision-language models	Qwen3-VL-8B, InternVL3.5-8B, Qwen3-VL-32B (4-bit NormalFloat, NF4); Qwen3-VL-8B with the rules R1-R10	photograph with the declared specification; crop with the declared vector and the rules	verdict, score	photographs, satellite pairs
Imputers	MICE, missForest, nearest-neighbor imputation	record with missing fields	imputed values	records
NeSySpec (ours)	specification and remote-sensing axioms with graded predicates, grounding axiom, cross-modal axioms	declared record with availability mask; optional photograph or satellite crop	verdict, per-constraint degrees (constraints respected and violated), evidence coverage and maps, largest-violation key, field localization, repair proposals, report	records, photographs, satellite pairs

macro-F1 on generator-label test (exact-set agreement, macro-F1 and false-violation rate 0.569, 0.829 and 0.036 against 0.504, 0.795 and 0.074; macro-F1 -0.034 $[-0.114,\ +0.075] \cdot [-0.132,\ +0.064]$) and below the constraint verdicts in macro-F1 on intervention-label unseen families (0.569 against 0.694, $+0.125$ $[+0.053,\ +0.187] \cdot [+0.035,\ +0.186]$); beyond the verdicts, NeSySpec also returns the altered field and a repair from the same degrees.

Table F.2: Configuration.

Setting	Value
Specification stratum	graded degrees in standard units of the clean training rows of each group, shrunk toward the fleet with weight $n_g/(n_g+5)$; in-band grading; hierarchical banks; truth $\exp(-v_k)$, degrees capped at 10^4; masked mean verdict ($p=1$), $V=\ln 2$ without available axioms; 336 trainable parameters, 126 of them bounded latents; chosen on validation among eight candidates with eight seeds each
Remote-sensing stratum	six graded slots (R1, R2 for civil and military roles, R3, R4, R10) in standard units of the clean training rows; R5 to R9 as the negative logarithm of their truth; masked weighted mean verdict ($p=1$); in-band grading; flat banks; 57 trained parameters; chosen on validation among eight candidates
Training of the strata	AdamW [72], learning rate 5×10^{-3} multiplied by 5, that is 2.5×10^{-2}, for every trained parameter of a stratum (bands, tolerances, slopes, in-band weights, latents and the parameters of R5 and R6), weight decay 0, batch 256; at most 100 epochs, early stopping on validation AUROC with patience 15, best state restored; $\lambda_{\text{aux}}=0.5$; seeds 0-7
Fusion of the strata	weighted mean between strata with $\lambda_{\mathcal{R}}=\lambda_{\mathcal{AX}}=1$; masked p-mean error with $p=4$ inside the cross-modal stratum; Reichenbach implication (R7-R9, grounding axiom, photographic AX4, satellite AX5); axiom weight 1, except R9 and satellite AX5 (0.5)
Second head	monotone logistic regression on the degrees, the availabilities and the verdict (on satellite pairs $V_{\mathcal{R}}$, the cross-modal truths and $1-\text{feas}_{\mathcal{AX}}$), inverse regularization in $\{10^{-4},\dots,10^{3}\}$ and, on specification records, the inclusion of the verdict selected on validation AUROC; operating thresholds: maximum of Youden's J [78] on validation
Abduction and repair	single-field log changes on a grid from -3 to 3 in steps of 0.005; smallest-change rule; feasibility threshold 0.9 ($V\le-\ln 0.9\approx0.1054$) for both; repair on the $k_r=1$ top-ranked field by an augmented-Lagrangian method with Adam steps
Optional settings	rank-then-explain: residual MLP alarm at its validation Youden threshold; critical-field alert: undeclared fields required by an applicable axiom and declared by at least 95% of the training records of the class, or of the engine type for classes with fewer than five training records; MICE imputation: chosen among nine configurations on validation
Photographs	ResNet-50 with ImageNet weights up to stage 3 (28×28 grid at 448 pixels), 25 sigmoid heads, class-activation maps; bounding box enlarged by 10% per side, 448×448 pixels; AdamW, learning rate 3×10^{-4} for heads and localization parameters and 3×10^{-5} for the backbone, weight decay 10^{-4}, gradient clipping at 5, batch 32, mixed precision; 20 fixed epochs without early stopping, checkpoint at the first epoch with the best validation macro head accuracy
Grounding	$\lambda_{\mathcal{AV}}=1$; floor weight 0.1 and mass floor $S_0=1$; Loc_k with ℓ_k initialized at 0.5 and κ_k at 10; visibility guard at 0.5 on the 0/1 visibility flag of each part (photographs: detector confidence at least 0.20, as in the pseudo-masks row; satellite crops: the annotated diamond)
Satellite crops	ResNet-50, ConvNeXt-T and ViT-B/16 (token maps), ImageNet weights, last stage; grounding region: the annotated diamond; aircraft box enlarged by 15% per side, 224×224 pixels; batch 64; at most 20 epochs, early stopping on the validation macro head AUROC with patience 5; other settings as for photographs
Pseudo-masks	OWLv2 (base, patch 16, ensemble), boxes with score at least 0.15, up to 4 for engines and landing gear and 1 otherwise; SAM ViT-B on the kept boxes; visibility: detector confidence at least 0.20; dilation: maximum filter of radius 2 native pixels
Cross-modal axioms	photographs: AX2 Mach cap 0.62 and slope 25; AX3 band edges at 10^4, $10^{5.05}$ and $10^{5.5}$ kg, slope 16
Baseline detectors	perceptron with two hidden layers of 64 units on the normalized fields and the mask (Adam, learning rate 10^{-3}, weight decay 10^{-5}, batch 256, at most 200 epochs, patience 20); Isolation Forest with 200 trees and Mahalanobis distance with Ledoit-Wolf covariance, both fitted on the clean training records

G.2 Diagnosis, Localization and Repair

From the same degrees, NeSySpec's largest-violation key names the violated constraint in 0.942 of corrupted test records and in 0.897 ± 0.007 on the unseen class, abduction finds the altered field in 0.742 ± 0.011 of single-field test corruptions, and the repair proposal lowers the log-space error on the altered field from 0.366 to 0.133 (Table G.6, the generator-label levels behind Figure 6b and c; Table G.5 gives the test levels of panel a). On unseen families the key, 0.973 on each seed, is above Fellegi-Holt ($+0.174$ [$+0.051$, $+0.304$] $\cdot$ [$+0.074$, $+0.365$]), Z3 ($+0.174$ [$+0.047$, $+0.311$] $\cdot$ [$+0.101$, $+0.286$]), MLP-SHAP ($+0.077$ [$+0.020$, $+0.143$] $\cdot$ [$+0.018$, $+0.189$]) and LightGBM-SHAP ($+0.117$ [$+0.052$, $+0.184$] $\cdot$ [$+0.055$, $+0.163$]) and on par with the class-standardized residual ($+0.018$); these unseen-family contrasts were not checked across the 20 alternative bootstrap streams. Class-I localization names the field of the largest residual of the Class I statistical sizing regressions on maximum take-off mass fitted on the training aircraft (the Class I sizing-regression residual of Section 4.3). MLP-SHAP, the attribution of a second model on the residuals, ranks the violated constraint first more often on the unseen class (0.970), and MICE at the field of the class-I residual is on par with the repair on both errors (0.120 and 0.161); NeSySpec's key, abduction and repair come with a held or violated verdict per constraint and the report. With 8 seed pairs instead of 5, the optional rank-then-explain

Table G.1: Contrasts behind the body results: constraint diagnosis.

Comparator	Difference and intervals
Constraint diagnosis (top-1); NeSySpec returns the violated constraint from its own degrees, by the largest-violation key	
Test	
NeSySpec's second head	+0.002 [−0.010, +0.018] · [−0.013, +0.011]
Fellegi-Holt	+0.147 [+0.068, +0.234] · [+0.079, +0.298]
Z3	+0.117 [+0.040, +0.209] · [+0.045, +0.256]
Rule diagnosis	+0.085 [+0.027, +0.146] · [+0.023, +0.150]
LightGBM-SHAP	+0.071 [+0.025, +0.118] · [+0.003, +0.098]
DL2 method	+0.499 [+0.455, +0.539] · [+0.428, +0.538]
Qwen3-8B [a]	+0.814 [+0.759, +0.868] · [+0.742, +0.875]
Class-standardized residual	+0.013 [−0.020, +0.045] · [−0.019, +0.041]
MLP-SHAP	+0.001 [−0.043, +0.043] · [−0.050, +0.028]
Test, 81 records with one violated constraint	
Fellegi-Holt	+0.284 [+0.186, +0.387] · [+0.193, +0.476]
Z3	+0.241 [+0.152, +0.335] · [+0.162, +0.398]
Rule diagnosis	+0.228 [+0.132, +0.329] · [+0.139, +0.391]
LightGBM-SHAP	+0.144 [+0.070, +0.220] · [+0.041, +0.195]
Class-standardized residual	+0.062 [+0.016, +0.105] · [+0.025, +0.104]
MLP-SHAP	+0.046 [+0.005, +0.103] · [+0.009, +0.074]
Test, 140 records the crisp edits flag	
Fellegi-Holt	+0.083 [−0.019, +0.192] · [−0.016, +0.291]
Z3	+0.050 [−0.045, +0.162] · [−0.041, +0.239]
Rule diagnosis	+0.014 [−0.057, +0.087] · [−0.053, +0.129]
Unseen families	
DL2 method	+0.595 [+0.544, +0.643] · [+0.569, +0.656]
Qwen3-8B	+0.755 [+0.691, +0.827] · [+0.675, +0.817]
Class-standardized residual	+0.018 [−0.036, +0.073] · [−0.044, +0.100]
Rule diagnosis	+0.110 [+0.024, +0.190] · [0.000, +0.200]
Unseen class	
NeSySpec's second head	+0.119 [+0.036, +0.219] · [+0.037, +0.221]
Fellegi-Holt	+0.228 [+0.152, +0.302] · [+0.154, +0.310]
Z3	+0.119 [+0.083, +0.155] · [+0.085, +0.156]
Rule diagnosis	+0.082 [+0.046, +0.121] · [+0.047, +0.124]
Class-standardized residual	+0.231 [+0.151, +0.310] · [+0.156, +0.303]
DL2 method	+0.599 [+0.529, +0.672] · [+0.529, +0.674]
Qwen3-8B	+0.734 [+0.647, +0.816] · [+0.646, +0.815]
LightGBM-SHAP	−0.032 [−0.081, +0.017] · [−0.081, +0.016]
MLP-SHAP	−0.073 [−0.125, −0.023] · [−0.124, −0.025]
Unseen class, NeSySpec's second head (its secondary key): the comparisons that the two keys read differently, a reading in plain words for either counting as different	
Z3	−0.000 [−0.085, +0.075] · [−0.084, +0.072]
Rule diagnosis	−0.037 [−0.143, +0.054] · [−0.143, +0.051]
LightGBM-SHAP	−0.152 [−0.283, −0.040] · [−0.285, −0.043]

Comparator	Difference and intervals
Diagnosis on intervention labels (top-1), without retraining; NeSySpec returns the violated constraint by the same key	
Test	
Fellegi-Holt	+0.180 [+0.126, +0.243] · [+0.135, +0.283]
Z3	+0.146 [+0.092, +0.214] · [+0.090, +0.231]
Rule diagnosis	+0.132 [+0.085, +0.181] · [+0.079, +0.182]
MLP-SHAP	+0.074 [+0.022, +0.119] · [+0.005, +0.128]
Z3, records the crisp edits flag	+0.074 [+0.001, +0.164] · [+0.004, +0.215]
Rule diagnosis, same records	+0.054 [−0.004, +0.115] · [−0.000, +0.156]
LightGBM-SHAP	+0.058 [−0.018, +0.129] · [−0.066, +0.126]
Class-standardized residual	+0.096 [+0.043, +0.145] · [−0.014, +0.143]
Test, NeSySpec's second head (its secondary key): the comparisons that the two keys read differently, a reading in plain words for either counting as different	
MLP-SHAP	+0.065 [+0.016, +0.108] · [−0.002, +0.115]
Class-standardized residual	+0.087 [+0.037, +0.131] · [−0.016, +0.128]
Z3, records the crisp edits flag	+0.070 [−0.003, +0.157] · [−0.001, +0.209]
Rule diagnosis, same records	+0.050 [−0.008, +0.110] · [−0.006, +0.152]
Unseen families	
NeSySpec's second head	−0.019 [−0.035, −0.006] · [−0.046, −0.006]
MLP-SHAP	+0.101 [+0.057, +0.150] · [+0.057, +0.155]
LightGBM-SHAP	+0.024 [−0.031, +0.097] · [−0.033, +0.136]
Class-standardized residual	+0.068 [−0.0032, +0.138] · [+0.011, +0.174]
Rank-then-explain, flagged and correctly diagnosed (5 seed pairs); NeSySpec returns the violated constraint behind the residual MLP's alarm	
Test, against NeSySpec's verdict with its own alarm and the same key (a declared sensitivity)	
Flagged and diagnosed	−0.028 [−0.052, −0.006] · [−0.064, −0.013]
False alarms, 32 controls	−0.094 [−0.2125, −0.0125] · [−0.256, 0.000]
Comparators behind the same alarm	
Fellegi-Holt, test	+0.132 [+0.060, +0.211] · [+0.066, +0.274]
Z3, test	+0.101 [+0.031, +0.190] · [+0.028, +0.231]
MLP-SHAP, test	−0.009 [−0.047, +0.026] · [−0.061, +0.021]
MLP-SHAP, unseen class	−0.041 [−0.098, +0.013] · [−0.098, +0.009]

Difference: NeSySpec, in the setting named by the headings above the row, minus the comparator named by the row or by its heading (rank-then-explain minus the comparator in its block); for false alarms, repair errors and evasion, where lower is better, a negative difference favors that setting. Intervals: 95% cluster-bootstrap percentiles at each of the two cluster levels, [aircraft] · [family] on specification records, [annotated aircraft] · [site] on satellite pairs and [image] · [variant] on photographs. Generator-oracle labels unless intervention labels are named; native NaN: TabPFN v2 with native missing values.
[a] Qwen3-8B: the language model given the constraints in words, on anonymized records.

setting stays above Fellegi-Holt (+0.130) and Z3 (+0.100) behind the same alarm on test, on par with the class-standardized residual and below NeSySpec's verdict with its own alarm and the same key (−0.028); strict abstention, which gives no diagnosis when no available degree exceeds $\ln 2$, withdraws the diagnosis of 3 to 5 records per seed, none of them correctly diagnosed.

G.3 Missing Evidence

Masked quantification makes every verdict of NeSySpec a function of the available evidence alone: Figure 7 in Section 4.5 shows the exact reduction on specification records and satellite pairs and the unseen class with half the fields hidden, where the masked verdict (0.577) is ahead of the residual MLP and TabPFN v2 with native missing values and on par with LightGBM trained with masks (Table G.2). With 10% of the fields hidden in distribution, the MLP trained with masks ranks test records higher (−0.0425) and the residual MLP unseen-family records (−0.0925); both read every field, imputed or masked in training, while the masked verdict reads the declared evidence alone and states which constraints remain checkable. With no field hidden beyond those the records already leave undeclared, the masked verdict is ahead of MICE imputation of every missing field at scoring, by +0.051 [+0.031, +0.063] · [+0.031, +0.062] on test and +0.035 [+0.011, +0.062] · [+0.012, +0.064] on unseen

Table G.2: Contrasts behind the body results: field localization, repair and missing evidence.

Comparator	Difference and intervals
Field localization (exact top-1); NeSySpec returns the altered field by abduction, the single-field change that restores consistency	
Test	
Minimum cardinality	+0.375 [+0.286, +0.470] · [+0.293, +0.529]
Ochiai ranking	+0.339 [+0.234, +0.450] · [+0.246, +0.413]
Largest class-I residual	+0.145 [+0.065, +0.232] · [+0.071, +0.203]
Same comparator, NeSySpec with the largest-feasibility rule [a]	+0.095 [−0.006, +0.192] · [−0.020, +0.176]
Oracle tie-break	+0.113 [−0.033, +0.280] · [−0.028, +0.454]
Unseen families	
Minimum cardinality	+0.359 [+0.241, +0.485] · [+0.241, +0.496]
Ochiai ranking	+0.179 [+0.041, +0.305] · [+0.029, +0.302]
Oracle tie-break	+0.065 [−0.125, +0.274] · [−0.089, +0.370]
Largest class-I residual	+0.043 [−0.041, +0.136] · [−0.018, +0.141]
Unseen class	
Minimum cardinality	+0.360 [+0.254, +0.478] · [+0.257, +0.480]
Largest class-I residual	+0.526 [+0.394, +0.653] · [+0.380, +0.646]
Ochiai ranking	+0.344 [+0.220, +0.463] · [+0.225, +0.460]
Oracle tie-break	+0.163 [−0.002, +0.335] · [−0.003, +0.351]
Intervention labels, without retraining, test	
Minimum cardinality	+0.287 [+0.220, +0.359] · [+0.238, +0.411]
Same, records the crisp edits flag	+0.248 [+0.153, +0.338] · [+0.164, +0.406]
Oracle tie-break	+0.140 [+0.041, +0.248] · [+0.041, +0.361]
Largest class-I residual	+0.140 [+0.073, +0.204] · [+0.068, +0.200]
Repair (log-space error, test); NeSySpec returns a repair proposal from its own abduction and verdict	
Error on the altered field	
No repair	−0.233 [−0.262, −0.204] · [−0.272, −0.214]
MICE, class-I localization	+0.013 [−0.022, +0.047] · [−0.033, +0.044]
Regression, class-I localization	−0.027 [−0.065, +0.008] · [−0.089, +0.004]
missForest, class-I localization	−0.028 [−0.068, +0.011] · [−0.085, +0.007]
PiShield-style ordering, class-I localization	−0.085 [−0.119, −0.049] · [−0.117, −0.049]
Error on the whole record, class-I localization	
Nearest neighbor	−0.106 [−0.171, −0.046] · [−0.210, −0.050]
Regression	−0.082 [−0.142, −0.018] · [−0.186, −0.027]
MICE	+0.039 [−0.015, +0.099] · [−0.013, +0.112]
missForest	−0.037 [−0.102, +0.029] · [−0.133, +0.021]
PiShield-style ordering	−0.063 [−0.120, +0.001] · [−0.128, −0.013]
Error on the whole record, other comparators	
No repair	−0.166 [−0.217, −0.107] · [−0.215, −0.110]
Fellegi-Holt	−0.114 [−0.176, −0.042] · [−0.163, −0.047]
MaxSMT	−0.108 [−0.170, −0.035] · [−0.175, −0.042]
MICE, Fellegi-Holt localization	−0.151 [−0.378, −0.001] · [−0.497, +0.024]

Comparator	Difference and intervals
Fields hidden at random (AUROC); NeSySpec returns its verdict on the available constraints and lists the unchecked ones	
Unseen class, half of the fields hidden	
Residual MLP	+0.0410 [+0.0085, +0.0738] · [+0.0074, +0.0734]
TabPFN v2, native NaN	+0.0382 [+0.0047, +0.0736] · [+0.0060, +0.0735]
LightGBM trained with masks	+0.0113 [−0.0366, +0.0655] · [−0.0387, +0.0704]
10% of the fields hidden	
MLP trained with masks, test	−0.0425 [−0.0705, −0.0168] · [−0.0627, −0.0038]
Residual MLP, unseen families	−0.0925 [−0.1266, −0.0635] · [−0.1320, −0.0507]
Satellite images hidden or from another sensor (AUROC), test [b]; NeSySpec returns the fused verdict, the record verdict bit for bit when the image is hidden	
30% of the images hidden	
CatBoost on knowledge-base truths	−0.031 [−0.043, −0.020] · [−0.042, −0.015]
Read-out trained with modality dropout	−0.026 [−0.036, −0.018] · [−0.036, −0.012]
Visual stratum trained on ground photographs, scoring satellite crops	
Record verdict alone	−0.071 [−0.094, −0.052] · [−0.103, −0.039]

Difference, intervals and notation as in Table G.1.
[a] Abduction with the largest-feasibility rule, a sensitivity not selected on validation; with the nearest threshold of the selected rule, both intervals of the difference from the class-I residual stay above zero.
[b] The comparators were not retrained or retuned, and those on knowledge-base truths keep the truths of the saturating-envelope variant of the remote-sensing stratum they were fitted on (Appendix G.5).

families, and ahead of global means at scoring on every split (+0.147, +0.121 and +0.164). With fields hidden on unseen families, MICE at scoring ranks higher, with differences of −0.055 [−0.081, −0.029] · [−0.088, −0.030], −0.103 [−0.130, −0.078] · [−0.139, −0.084] and −0.113 [−0.140, −0.087] · [−0.155, −0.096] with 10, 20 and 30% of the fields hidden; MICE fills the values that the masked verdict leaves unread, and the report lists the constraints that remain checkable. With MICE imputation enabled, an optional setting of NeSySpec in place of exact masking (Section 4.5), the verdict is on par with the MLP trained with masks given the same imputer, +0.0089 [−0.0039, +0.0230] · [−0.0037, +0.0349] on test with 10% of the fields hidden.

G.4 Record Coherence

NeSySpec's record verdict reaches a test AUROC of 0.979 ± 0.002 on generator labels, from 0.937 on cruise to 1.000 on geometric corruptions by corruption type, each type scored against the 48 clean-labeled test records (16 clean records and their 32 noisy-but-coherent copies; means of 8 seeds), and each of its flags names the violated constraint, the altered field and a repair proposal. Table G.7 groups the detectors of Section 4.4 by the knowledge they receive, with what each returns. On test the

Table G.3: Contrasts behind the body results: record coherence.

Comparator	Difference and intervals
Record coherence (AUROC); NeSySpec returns its verdict, each flag with the violated constraint, field and repair	
Test, learners on the raw fields	
CatBoost	+0.187 [+0.114, +0.259] · [+0.094, +0.335]
TabPFN v2	+0.197 [+0.118, +0.266] · [+0.105, +0.320]
XGBoost	+0.205 [+0.119, +0.280] · [+0.109, +0.334]
LightGBM	+0.231 [+0.136, +0.315] · [+0.119, +0.395]
Test, other detectors	
Rules with pooling	+0.075 [+0.039, +0.115] · [+0.041, +0.132]
Train-fitted rule checker	+0.151 [+0.101, +0.172] · [+0.083, +0.165]
Mahalanobis	+0.186 [+0.132, +0.241] · [+0.136, +0.269]
DL2 method	+0.276 [+0.223, +0.322] · [+0.224, +0.353]
Qwen3-8B	+0.359 [+0.275, +0.437] · [+0.296, +0.427]
Test, learners on the residuals	
MLP	−0.010 [−0.022, +0.002] · [−0.017, +0.002]
TabPFN v2	+0.021 [−0.008, +0.069] · [−0.011, +0.037]
One-class SVM	+0.022 [+0.0003, +0.039] · [−0.0001, +0.043]
LightGBM	+0.023 [+0.0003, +0.046] · [−0.0004, +0.034]
Unseen families	
Train-fitted rule checker	+0.090 [+0.043, +0.121] · [+0.014, +0.109]
DL2 method	+0.296 [+0.244, +0.345] · [+0.241, +0.395]
Qwen3-8B	+0.254 [+0.204, +0.305] · [+0.188, +0.332]
MLP	−0.008 [−0.019, +0.0002] · [−0.017, +0.001]
TabPFN v2, fields and residuals	−0.009 [−0.021, +0.0001] · [−0.019, +0.0002]
Unseen class, other detectors	
DL2 method	+0.355 [+0.312, +0.394] · [+0.326, +0.383]
Qwen3-8B	+0.266 [+0.181, +0.356] · [+0.176, +0.347]
Unseen class, learners on the residuals	
MLP	−0.045 [−0.074, −0.005] · [−0.075, −0.004]
TabPFN v2	−0.050 [−0.080, −0.015] · [−0.083, −0.013]
TabPFN v2, native NaN	−0.051 [−0.075, −0.021] · [−0.077, −0.019]
TabICL	−0.046 [−0.078, −0.011] · [−0.080, −0.009]
TabPFN v2, fields and residuals	−0.044 [−0.075, −0.007] · [−0.079, −0.005]
NeSySpec and the comparators retrained on intervention labels (AUROC), test; NeSySpec (0.807) returns its verdict with the per-constraint degrees	
Residual MLP	−0.053 [−0.084, −0.028] · [−0.105, −0.033]
TabPFN v2, native NaN [a]	−0.078 [−0.114, −0.047] · [−0.119, −0.062]
TabPFN v2, imputed [a]	−0.081 [−0.121, −0.048] · [−0.145, −0.060]
False alarms on the 32 controls at a TPR of 0.90, test [b]; each alarm of NeSySpec names the constraint that raised it or, with the alert, the undeclared critical field	
NeSySpec's verdict	
Mahalanobis	−0.570 [−0.844, −0.384] · [−0.792, −0.385]
Rules with pooling	−0.244 [−0.500, −0.031] · [−0.620, −0.056]
Z3 edit count	−0.244 [−0.491, −0.062] · [−0.544, −0.093]
Fellegi-Holt	−0.244 [−0.482, −0.008] · [−0.569, −0.050]
Residual MLP	+0.078 [0.000, +0.188] · [0.000, +0.113]
NeSySpec with the critical-field alert [c]	
NeSySpec's verdict	+0.062 [0.000, +0.188] · [0.000, +0.250]
Mahalanobis	−0.508 [−0.812, −0.306] · [−0.767, −0.264]
Z3 edit count	−0.181 [−0.444, −0.006] · [−0.492, −0.009]
Rules with pooling	−0.181 [−0.444, +0.062] · [−0.545, +0.111]
Fellegi-Holt	−0.181 [−0.437, +0.033] · [−0.486, +0.008]
Residual MLP	+0.141 [0.000, +0.301] · [0.000, +0.278]
Satellite pairs (AUROC) [d]; NeSySpec returns the fused verdict and names the cross-modal axiom the image contradicts	
NeSySpec's fused verdict, ResNet-50	
AutoGluon-MM, test	+0.165 [+0.137, +0.194] · [+0.132, +0.197]
AutoGluon-MM, unseen sites	+0.110 [+0.090, +0.132] · [+0.098, +0.144]
TIP-style, test	+0.067 [+0.046, +0.086] · [+0.055, +0.081]
TIP-style, unseen sites	+0.048 [+0.032, +0.064] · [+0.035, +0.086]
NeSySpec's multimodal verifier and Qwen3-VL-8B given the rules	
ConvNeXt-T, test	+0.247 [+0.208, +0.286] · [+0.201, +0.283]
ConvNeXt-T, unseen sites	+0.324 [+0.287, +0.362] · [+0.200, +0.356]
ResNet-50, test	+0.241 [+0.202, +0.280] · [+0.196, +0.279]
ViT-B/16, test	+0.242 [+0.204, +0.281] · [+0.197, +0.277]
NeSySpec's fused verdict, test image swaps	
Late fusion trained with swaps	−0.012 [−0.026, −0.0033] · [−0.018, −0.004]
Record verdict alone and Qwen3-VL-8B given the rules	
Test	+0.048 [−0.009, +0.103] · [−0.009, +0.094]
Photographs (AUROC), cross-modal stratum; NeSySpec returns the contradicted axiom with the part map of its evidence	
Qwen3-VL-8B, test	+0.025 [−0.010, +0.058] · [−0.046, +0.098]
InternVL3.5-8B, test	+0.032 [−0.002, +0.067] · [−0.033, +0.086]
Qwen3-VL-32B, test	−0.020 [−0.055, +0.016] · [−0.073, +0.042]
Qwen3-VL-32B, unseen families	−0.094 [−0.131, −0.058] · [−0.154, −0.056]

Difference, intervals and notation as in Table G.1; Table G.8 gives the ablation and the same residuals in other models.
[a] Its preprocessing statistics include the scored rows.
[b] Ties at the threshold interpolated; counting every tie, NeSySpec with the critical-field alert differs by −0.852 from the pooled rules [−0.992, +0.062] · [−0.955, +0.111], Fellegi-Holt [−0.992, +0.031] · [−0.955, 0.000] and the Z3 count [−0.992, −0.039] · [−0.955, −0.067].
[c] NeSySpec's optional critical-field alert: a record is flagged when the verdict or the alert fires (Table G.7, note g).
[d] The comparators were not retrained or retuned, and those on knowledge-base truths keep the truths of the saturating-envelope variant of the remote-sensing stratum they were fitted on (Appendix G.5).

Table G.4: Contrasts behind the body results: deliberate falsification.

Comparator	Difference and intervals
Black-box attacks (evasion); NeSySpec returns its verdict, each flag with the field found by abduction and a repair proposal	
Cover-ups, up to five further edits	
Residual MLP	−0.064 [−0.073, −0.055] · [−0.075, −0.053]
Residual TabPFN v2	−0.059 [−0.071, −0.046] · [−0.072, −0.043]
Residual one-class SVM	−0.062 [−0.077, −0.048] · [−0.081, −0.043]
Raw-field LightGBM	−0.065 [−0.075, −0.055] · [−0.076, −0.055]
Isolation Forest	−0.049 [−0.064, −0.035] · [−0.064, −0.028]
Rules with pooling	+0.004 [−0.004, +0.013] · [−0.005, +0.013]
Mahalanobis	+0.010 [−0.0033, +0.025] · [−0.005, +0.028]
Bare lies	
Rules with pooling	−0.113 [−0.143, −0.084] · [−0.149, −0.064]
Raw-field LightGBM	−0.169 [−0.208, −0.130] · [−0.205, −0.121]
Residual MLP	−0.016 [−0.052, +0.023] · [−0.064, +0.052]
Residual LightGBM	+0.052 [+0.009, +0.096] · [+0.005, +0.115]
Residual TabPFN v2	+0.105 [+0.058, +0.155] · [+0.056, +0.176]
Omissions of up to two fields (evasion); NeSySpec returns its verdict and lists each undeclared field with the constraints it leaves unchecked	
Residual MLP	+0.026 [+0.014, +0.039] · [+0.012, +0.044]
Residual one-class SVM	+0.143 [+0.103, +0.184] · [+0.103, +0.200]
Mahalanobis	+0.119 [+0.071, +0.178] · [+0.069, +0.196]
Omissions of up to two fields, attacker aware of the alert, NeSySpec with the critical-field alert [a]; each flag names the undeclared critical field or the violated constraint	
Comparators with the same alert, all splits	
Residual MLP	−0.027 [−0.064, +0.014] · [−0.073, +0.046]
Residual one-class SVM	+0.019 [−0.035, +0.073] · [−0.058, +0.104]
Residual MLP with the same alert, per split	
Test	−0.077 [−0.117, −0.029] · [−0.119, +0.016]
Unseen families	−0.074 [−0.103, −0.043] · [−0.099, −0.008]
Unseen class	+0.125 [+0.060, +0.197] · [+0.060, +0.197]

Difference, intervals and notation as in Table G.1.
[a] NeSySpec's optional critical-field alert: a record is flagged when the verdict or the alert fires (Table G.7, note g).

Table G.5: Constraint verdicts against per-constraint labels on test.

Reading	Generator-oracle labels	Intervention labels
Differentiable logic: a probability per constraint		
Network with the DL2 loss on the same residuals	0.474/0.661/0.020	0.222/0.426/0.019
same network, thresholds tuned on validation	0.436/0.746/0.056	0.218/0.541/0.056
DL2 method on the declared fields	0.106/0.425/0.140	0.048/0.259/0.141
same method, thresholds tuned on validation	0.053/0.484/0.216	0.015/0.319/0.214
Crisp rules and edits: a crisp violated set		
Crisp edits (Fellegi-Holt, Z3), small classes pooled	0.569/0.829/0.036	0.212/0.505/0.036
Rule checker, per-class bands, small classes not pooled	0.103/0.691/0.243	0.049/0.494/0.252
Residual statistics: a z-score per residual		
Class-standardized residual, threshold on specification validation	0.446/0.714/0.069	0.194/0.480/0.078
same residual, threshold on intervention validation	-	0.138/0.560/0.174
Knowledge in words: a judgment with the constraints it lists		
Language model given the constraints in words	0.113/0.153/0.140	-
same model, read only when it calls a record incoherent	0.235/0.065/0.000	-
NeSySpec, graded degree and held or violated verdict per constraint	**0.504/0.795/0.074**	**0.283/0.575/0.084**

Test split. Generator-oracle labels follow the axioms' relations and were the per-constraint training targets; intervention labels, set from the edit alone, are independent of them. Each cell: exact-set agreement / macro-F1 / false-violation rate (lower is better); - not measured. Intervention labels mark every available constraint whose scope contains a field moved beyond tolerance; their false-violation rate is computed on unaltered records and their noisy-but-coherent copies; on records edited only within tolerance, which these labels also count as unaltered, it is 0.119 for NeSySpec, 0.063 for the DL2-loss network and 0.087 for the crisp edits. The ablation trained with the binary record label only, without per-constraint targets, reaches 0.081/0.634/0.194 with generator-oracle labels (Section 4.8).

record verdict is on par with TabPFN v2 and close to the residual MLP, LightGBM and the one-class SVM on the residuals and ahead of the learners on the raw fields, the rule checkers, the Mahalanobis distance, the DL2 method and the language model (Table G.3). On the unseen class five learners on the residuals rank higher in record-level detection (residual MLP 0.978, TabPFN v2 up to 0.984); there, with NeSySpec's optional rank-then-explain setting, the residual MLP raises the alarm and NeSySpec explains it with the violated constraint, the altered field and a repair, flagging and correctly diagnosing 0.849 ± 0.043 of the corrupted unseen-class records (Table G.6).

Against intervention labels the record verdict reaches 0.701, 0.750 and 0.719 on test, unseen families and the unseen class, without retraining, while the generator oracle itself, with its bands on all 107 aircraft, reaches 0.676, 0.709 and 0.714 on the same labels. These labels count as incoherent the 613 altered rows that break no constraint under those bands, and as coherent the edits within tolerance, some of which do break one. With only the clean records and their noisy copies as negatives, the verdict reaches 0.840, 0.847 and 0.811. On the altered rows that break no constraint, scored against all clean rows, it stays below chance (0.394, 0.449 and 0.439), as the comparators that use the physics do on the pooled held-out splits.

G.5 Records with Images

Without an image, NeSySpec's fused verdict equals the record verdict bit for bit in 9 of 9 satellite cells, against 0 of 207 learned method and cell rows, and its second head with ConvNeXt-T (8 seeds) reaches 0.998 ± 0.001 on test and 0.991 ± 0.003 on unseen sites over all pair types. On the satellite pairs of Section 4.4, the best learned comparators rank the pairs higher after learning from labeled pairs: CatBoost on knowledge-base truths on test swaps, TabPFN on visual posteriors and image embeddings on unseen-site swaps, and late fusions on knowledge-base truths over all pair types on unseen sites (nine runs of three backbones). Learned heads also rank the pairs higher in the contrasts of Tables G.2 and G.3, where NeSySpec's fused verdict differs by -0.012 from the late fusion trained with swaps on test swaps and by -0.031 from CatBoost on knowledge-base truths with 30% of the images hidden. The fused verdict, instead, aggregates NeSySpec's strata without labels and names the contradicted axiom; on photographs, Qwen3-VL-32B ranks unseen-family pairs higher than the cross-modal stratum (-0.094). With 30% of satellite images hidden, the fused verdict (0.947) keeps the exact reduction; the read-out trained with modality dropout ranks the pairs at 0.973 without keeping it. When a visual

Table G.6: Diagnosis, localization and repair.

Constraint diagnosis, top-1 on corrupted records	Test	Unseen families	Unseen class
MLP-SHAP on residuals	0.942	0.895	0.970
Class-standardized residual	0.929	0.955	0.667
LightGBM-SHAP on residuals	0.871	0.856	0.930
Rule diagnosis, rules ranked by exceedance	0.857	0.863	0.815
Z3, first unsatisfiable core	0.825	0.799	0.778
Fellegi-Holt, complete violated set	0.795	0.799	0.670
DL2 method, declared fields	0.443	0.377	0.298
Language model, constraints in words	0.128	0.218	0.163
Chance	0.177	0.166	0.152
NeSySpec, second head (secondary key)	0.940 ± 0.005	0.973	0.778 ± 0.064
NeSySpec (ours), largest-violation key	**0.942**	**0.973**	$\mathbf{0.897 \pm 0.007}$
Field localization, top-1 on single-field corruptions	Test	Unseen families	Unseen class
Minimum cardinality, oracle tie-break	0.629	0.670	0.481
Minimum cardinality, random tie-break	0.367	0.377	0.284
Minimum cardinality, worst tie-break	0.210	0.205	0.156
Largest class-I residual	0.597	0.693	0.119
Class-standardized residual with Ochiai	0.403	0.557	0.300
Chance	0.077	0.076	0.079
NeSySpec (ours), abduction	$\mathbf{0.742 \pm 0.011}$	$\mathbf{0.736 \pm 0.010}$	$\mathbf{0.644 \pm 0.014}$
Repair on test records, lower is better	Altered field	Record	Clean records changed
MICE, class-I localization	0.120	0.161	0.917
MICE, class-I localization and detection gate	0.155	0.162	0.271
Class-I regression, class-I localization	0.160	0.282	1.000
missForest, class-I localization	0.161	0.238	1.000
MICE, Fellegi-Holt localization	0.211	0.352	0.312
Class-I PiShield-style ordering	0.218	0.264	0.250
Fellegi-Holt, all optimal solutions	0.241	0.314	0.250
No repair	0.366	0.366	0.000
NeSySpec (ours), repair proposal	$\mathbf{0.133 \pm 0.002}$	$\mathbf{0.200 \pm 0.002}$	$\mathbf{0.562 \pm 0.033}$
Flagged and correctly diagnosed, corrupted records	Test	Unseen families	Unseen class
Residual-MLP alarm, explained by (with NeSySpec's key: its optional rank-then-explain setting)			
MLP-SHAP	0.910	0.858	0.890
Class-standardized residual	0.895	0.878	0.661
Rule diagnosis, rules ranked by exceedance	0.832	0.806	0.791
Z3	0.800	0.743	0.759
Fellegi-Holt	0.770	0.738	0.648
NeSySpec, second head (secondary key)	0.900 ± 0.011	0.875 ± 0.015	0.732 ± 0.053
NeSySpec (ours), largest-violation key	$\mathbf{0.901 \pm 0.013}$	$\mathbf{0.875 \pm 0.015}$	$\mathbf{0.849 \pm 0.043}$
NeSySpec's own alarm, explained by its largest-violation key			
NeSySpec with the critical-field alert	0.929	0.925 ± 0.016	0.871 ± 0.007
NeSySpec's verdict (ours)	**0.929**	$\mathbf{0.925 \pm 0.016}$	$\mathbf{0.864 \pm 0.008}$

Generator-oracle labels, train prior; comparators as in Table F.1; top-1 rates count tied candidates at their expected value under random tie-breaking. Rows of NeSySpec: mean ± sd over 8 seeds (in the last block 5 seed pairs, the seeds 0 to 4 for which MLP-SHAP was computed), without ± where every seed agrees; comparator rows: mean over their seeds. Every row of NeSySpec reads the same degrees, which also give the held or violated verdict per constraint and the report.
Largest-violation key: the available constraint with the largest violation degree, selected on held-out training classes. Second head: the monotone read-out on the degrees, kept as the secondary key. SHAP: grouped exact Shapley values of the residual model, one group per constraint. Oracle and worst tie-break: ties among the optimal localizations broken toward and away from the altered field. Repair errors, on single-field corruptions: log-space error of the altered field, plus the change of the correct fields for the record; clean records changed: share of the 48 clean and control test records. Class-I PiShield-style ordering: sequential projection onto the crisp constraints, fields in ascending order of their class-I residual. Detection gate: imputation only on the records that the class-I residual flags, a residual outside its 95% prediction interval. Residual-MLP alarm and own alarm: the residual MLP or the verdict at its validation threshold; with the alert, the verdict or NeSySpec's optional critical-field alert.

stratum trained on ground photographs scores satellite crops, the record verdict alone (0.796) is the fallback across sensors without labels: the fused verdict reaches 0.724 and an out-of-domain abstention restores the record level, while the fusion selected with satellite labels reaches 0.786 and a refit trained on satellite labels, a reference with privileged access to them, reaches 0.979. Fitted on specification records and applied without satellite training through the shared axioms R2 to R4 (Appendix D.1), the knowledge base reaches 0.767 ± 0.007 on satellite test annotations and 0.764 ± 0.008 on unseen sites, against 0.796 and 0.812 for the remote-sensing record verdict trained on these annotations; fitted on satellite annotations and applied to specification records, it reaches 0.541 ± 0.003 (8 seeds, no interval); the military roles, absent from the records, keep their prior bands. Comparators on satellite pairs

Table G.7: Detection.

Detector	AUROC			FPR, controls
	Test	Unseen families	Unseen class	
Learners on the physical residuals: score; the network with the DL2 loss also a probability per constraint				
MLP	0.990 ± 0.003	0.999 ± 0.001	0.978 ± 0.011	0.008
TabPFN v2	0.958 ± 0.006	0.999 ± 0.001	0.983 ± 0.006	-
TabPFN v2, native missing values	0.957 ± 0.009	0.999 ± 0.001	0.984 ± 0.006	-
TabPFN v2, raw fields and residuals[a]	0.965 ± 0.007	1.000 ± 0.000	0.978 ± 0.006	-
TabICL	0.953 ± 0.003	0.996 ± 0.001	0.979 ± 0.005	-
LightGBM	0.957 ± 0.004	0.993 ± 0.002	0.904 ± 0.006	-
XGBoost, CatBoost and logistic regression	0.893-0.958	0.988-0.992	0.876-0.930	-
Learners on raw fields and residuals[a]	0.897-0.955	0.975-0.997	0.881-0.905	-
Network with the DL2 loss	0.965 ± 0.021	-	0.921 ± 0.018	-
One-class SVM, label-free	0.958	0.969	0.901	-
Local outlier factor and autoencoder, label-free	0.893-0.940	0.919-0.952	0.462-0.917	-
Crisp rules on the knowledge base: score and crisp violated set				
Rules with hierarchical pooling[c]	0.904	0.933	0.924[b]	0.330[d]
Z3 violated-edit count[c]	0.867	0.868	0.825	0.330[d]
Fellegi-Holt[c]	0.856	0.867	0.757	0.330[d]
Rule checker, train-fitted bands	0.828	0.901	0.924[b]	-
Knowledge in other forms: score with a probability per constraint (DL2 method) or a judgment (language model)				
DL2 method, declared fields	0.703 ± 0.030	-	0.579 ± 0.027	-
Language model (Qwen3-8B), constraints in words	0.620	0.737	0.667	-
Knowledge-free: error score				
Mahalanobis distance	0.793	0.739	0.616	0.656
Learners on raw fields[a]	0.748-0.792	0.834-0.863	0.505-0.641	-
MLP on normalized fields	0.711 ± 0.040	0.713 ± 0.037	0.537 ± 0.021	-
Isolation Forest	0.576 ± 0.022	0.542 ± 0.016	0.525 ± 0.028	-
NeSySpec, Logic Tensor Network: score, verdict per constraint, key, altered field, repair and report				
Saturating-envelope ablation[e]	0.821 ± 0.008	0.883 ± 0.005	0.898 ± 0.004	0.363
Textbook prior (sensitivity)[f]	0.977 ± 0.003	0.988 ± 0.001	0.929 ± 0.002	0.062
NeSySpec, second head	0.976 ± 0.002	0.993 ± 0.002	0.822 ± 0.043	-
NeSySpec with the critical-field alert[g]	-	-	-	0.148
NeSySpec's verdict (ours)	$\mathbf{0.979 \pm 0.002}$	$\mathbf{0.991 \pm 0.002}$	$\mathbf{0.934 \pm 0.001}$	**0.086**

AUROC on generator-oracle labels, mean ± sd over seeds where a method has several; a range gives the lowest and highest mean of the detectors of its row. FPR, controls: false-positive rate on the 32 noisy-but-coherent test controls at a true-positive rate of 0.90, mean over seeds. A cell marked - is not reported; in the alert row, the AUROC is the verdict's (note g).

[a] Raw fields: the declared fields with the availability mask and one-hot categories. The learners on raw fields are LightGBM, XGBoost, CatBoost and TabPFN v2; the learners on raw fields and residuals are LightGBM, XGBoost and CatBoost.

[b] Band-width exceedances. With log-ratio exceedances the rule checker scores 0.748 and the rules with hierarchical pooling 0.720; validation holds no business jet to choose between the two normalizations.

[c] Pooling candidates designed after a diagnosis on test; validation supports them independently. Fellegi-Holt (fewest fields to change, ties broken by the smallest change in log space) and the Z3 count (number of violated edits) use the same pooled bands; their scores do not depend on the exceedance normalization of note b.

[d] Every record that no edit flags scores 0, and these ties straddle the threshold: the printed value interpolates them, and it is 1.000 when every tie counts as a flag; every other row has the same value under both conventions.

[e] Clamped envelope truths in place of the graded predicates that carry the residual magnitude.

[f] Handbook envelopes in place of the bands fitted on the training fleet; its class mapping was written after the per-split audit; values of the verdict alone, without the alert.

[g] NeSySpec's optional critical-field alert: a record is flagged when the verdict or the alert fires. The alert leaves the score unchanged, so the AUROC columns are the verdict's. Among the genuine test records and their copies, the alert fires only on one aircraft of a training class, on its record and on its two noisy-but-coherent copies; 5 of the 6 false alerts on the 46 genuine held-out records fall on the 19 records of the unseen business-jet class, whose alert pool is the turbofan training records.

were not retrained or retuned, and those that read knowledge-base truths (CatBoost and TabPFN on knowledge-base truths, the best fusion and the fusion selected without satellite labels) keep the truths of the saturating-envelope variant of the remote-sensing stratum they were fitted on (the specification-record row of this ablation is in Table G.8). NeSySpec's graded remote-sensing stratum was selected on validation among eight candidate configurations; it trains with per-constraint targets that the end-to-end, fusion and vision-language comparators do not receive, and its validation split serves both early stopping and the choice of configuration.

Table G.8: Ablation study.

Ablation or model	Test AUROC	Test Difference	Unseen families AUROC	Unseen families Difference	Unseen class AUROC	Unseen class Difference
Ablations of NeSySpec on specification records, generator-oracle labels, train prior; intervals [aircraft] · [family]						
Saturating envelopes (bounded envelope truths)	0.821 [+0.073, +0.223] · [+0.022, +0.196]	+0.158	0.883 [+0.031, +0.163] · [+0.014, +0.123]	+0.108	0.898 [+0.0031, +0.062] · [+0.005, +0.067]	+0.035
Generalized mean, $p = 4$, instead of the mean degree	0.983 [−0.020, +0.013] · [−0.008, +0.004]	−0.004	0.980 [−0.004, +0.041] · [−0.004, +0.065]	+0.011	0.940 [−0.016, +0.006] · [−0.015, +0.006]	−0.006
Without in-band grading	0.970 [+0.002, +0.017] · [+0.0003, +0.015]	+0.009	0.977 [+0.0012, +0.039] · [+0.0010, +0.052]	+0.014	0.908 [+0.014, +0.036] · [+0.014, +0.036]	+0.025
Flat banks instead of hierarchical banks	0.965 [+0.004, +0.026] · [+0.0031, +0.023]	+0.015	0.960 [+0.006, +0.065] · [+0.009, +0.090]	+0.031	0.915 [+0.008, +0.030] · [+0.008, +0.031]	+0.019
Binary label alone, without per-constraint targets	0.964 [+0.005, +0.026] · [+0.0028, +0.023]	+0.015	0.974 [+0.0019, +0.046] · [+0.0028, +0.060]	+0.017	0.921 [+0.0024, +0.023] · [+0.0025, +0.024]	+0.012
Textbook prior instead of the train prior	0.977 [−0.004, +0.009] · [−0.009, +0.007]	+0.002	0.988 [−0.003, +0.016] · [−0.0019, +0.010]	+0.003	0.929 [−0.018, +0.027] · [−0.018, +0.025]	+0.004
Knowledge base at its initial values, untrained	0.907 [+0.045, +0.091] · [+0.035, +0.089]	+0.073	0.930 [+0.015, +0.106] · [+0.009, +0.118]	+0.061	0.845 [+0.058, +0.109] · [+0.056, +0.111]	+0.088
Neural predicates, learned monotone functions	0.984 [−0.022, +0.010] · [−0.011, +0.0018]	−0.005	0.991 [−0.008, +0.012] · [−0.0028, +0.017]	+0.0004	0.958 [−0.043, −0.0009] · [−0.044, −0.0018]	−0.024
Learners on the same physical residuals, generator-oracle labels; intervals [aircraft] · [family]						
Monotone model on the residual magnitudes, without LTN	0.966 [−0.0019, +0.029] · [−0.007, +0.023]	+0.013	0.993 [−0.013, +0.006] · [−0.011, +0.006]	−0.002	0.921 [−0.019, +0.042] · [−0.023, +0.042]	+0.012
Network with the DL2 loss on the residuals	0.965 [+0.0028, +0.026] · [+0.004, +0.025]	+0.014	-	-	0.921 [−0.010, +0.039] · [−0.011, +0.039]	+0.013
TabPFN v2	0.958 [−0.008, +0.069] · [−0.011, +0.037]	+0.021	0.999 [−0.020, +0.0003] · [−0.017, +0.0002]	−0.008	0.983 [−0.080, −0.015] · [−0.083, −0.013]	−0.050
MLP	0.990 [−0.022, +0.002] · [−0.017, +0.002]	−0.010	0.999 [−0.0185, +0.0002] · [−0.0170, +0.0014]	−0.0076	0.978 [−0.074, −0.005] · [−0.075, −0.004]	−0.045
NeSySpec's verdict (ours), full model	**0.979**		**0.991**		**0.934**	

Satellite pairs, ablations of NeSySpec's fused verdict; intervals [annotated aircraft] · [site]

Ablation or model	Test AUROC	Test Difference	Unseen sites AUROC	Unseen sites Difference
Single masked quantifier, 30% of images withheld	-	−0.002 [−0.0064, +0.0014] · [−0.0050, +0.0035]	-	−0.008 [−0.0117, −0.0046] · [−0.0086, −0.0045]
Single masked quantifier, all images	0.995 [−0.0097, −0.0002] · [−0.0068, +0.0002]	−0.004	0.992 [−0.0201, −0.0073] · [−0.0147, −0.0058]	−0.013
Kleene-Dienes implication, $\max(1-a, b)$ [21], instead of Reichenbach	0.991 [−0.0000137, 0] · [−0.0000118, 0]	−0.0000032	0.979 [−0.0000347, −0.0000008] · [−0.0000338, 0]	−0.0000158
NeSySpec's fused verdict (ours), all images	**0.991**		**0.979**	

AUROC: mean over 8 seeds, one model for the knowledge base at its initial values and for the monotone model; on satellite pairs nine runs, three backbones with three seeds each; with 30% of images withheld, differences only, from the three ResNet-50 runs with 20 random image masks each, masks that differ from those behind the 30% levels of Section 4.5. Difference: NeSySpec's verdict, on satellite pairs its fused verdict on the same images, minus the row, so a positive difference favors NeSySpec; under it, its 95% cluster-bootstrap intervals at the two levels of the block heading, one per line; - not reported. Values near 0 are printed with more decimals, seven in the Kleene-Dienes cells.
The ablations with $p = 4$, without in-band grading and with flat banks are trained models, and the textbook prior is a sensitivity. The knowledge base at its initial values starts from the train prior and is a single deterministic model. Other detectors are in Table G.7.

G.6 Grounding

On the FGVC pseudo-masks the grounding axiom raises test pointing over the network without a localization term on every backbone, by +0.664 with ConvNeXt-T, +0.699 with ResNet-50 at 512 pixels, +0.615 with ViT token-CAM and +0.677 with ResNet-50 at 448 pixels (8 against 3 seeds), each with an exact seed-bootstrap interval above zero. Figure G.1 gives its pointing levels against the controls (panel a, over its 8 seeds on FGVC-Aircraft and its 3 runs at weight 1 on PASCAL-Part; the location prior, a reference without learning, always points to the grid cell that the part covers most often in other images: the training images for FGVC-Aircraft and the other half of the images for PASCAL-Part) and its comparison with the localization and explanation-supervision methods (panel b, 3 runs per method at the weights selected on validation, where the axiom points to the part in 0.790 of test pseudo-masks). At the same loss weight the auxiliary loss puts about 4.3 times the axiom's

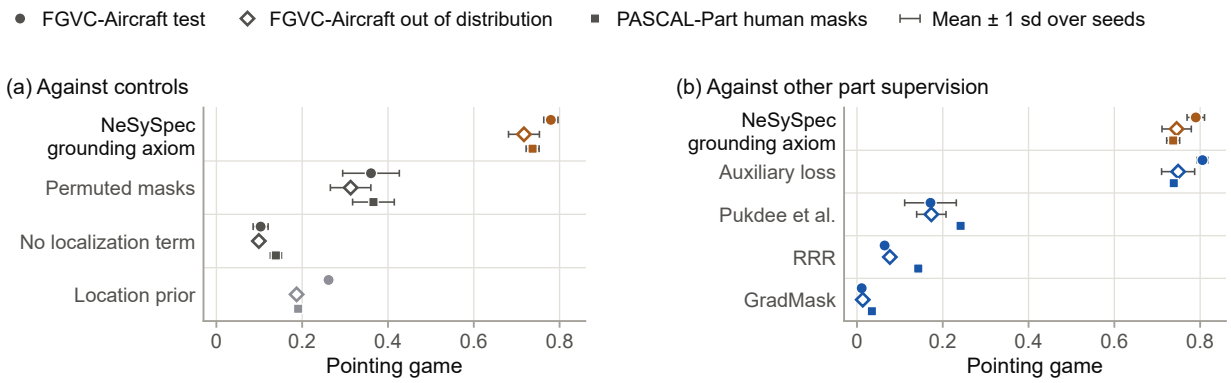

Figure G.1: Grounding of visual parts.

localization weight on each conjunct (Section 3.3), so the two are compared at their weights selected on validation (Appendix H).

On satellite crops (ConvNeXt-T, seed intervals over 8 paired seeds) the axiom keeps attribute and head AUROC on par with the network without a localization term.

G.7 Deliberate Falsification

Protocol. Each of the 736 falsification problems inflates one capacity of a genuine held-out record, the range, the cruise speed, the maximum take-off mass or the ceiling, by 5, 10, 20 or 40% (46 records: the 16 clean test records, 11 of unseen families and 19 of the unseen class). A bare lie submits the inflated record; a cover-up also edits up to five further declared fields within the ranges of the fleet, found by a black-box search that reads only the detector's score; an omission withholds up to two declared fields other than the inflated one. A lie evades a detector when the submitted record scores below the detector's validation threshold, and evasion is the share of evading lies among the problems whose genuine record the detector accepts; the smallest number of further edits that covers a flagged lie is censored at six.

NeSySpec with the critical-field alert, an optional setting, flags a record when the verdict or the alert fires. Under black-box falsification, a lie flagged by NeSySpec's record verdict is covered less often than one flagged by the residual MLP: with up to five further edits 0.936 of cover-ups pass against 1.000 (Table G.4), and 0.247 of its flagged bare lies cannot be covered, against 0.000. The smallest cover of one of its flagged lies has 2.69 further edited fields on average (2.68 with the critical-field alert), against 1.01 to 1.63 for the learned detectors (on the residuals, MLP 1.38, LightGBM 1.37 and TabPFN v2 1.63) and more for the Mahalanobis distance and the pooled rules (3.45 and 3.62; point estimates, censored at six). On the 72 problems that the verdict and the residual MLP flag on every seed, 0.413 of the verdict's flagged lies cannot be covered, against 0.000, with 3.76 edits per cover against 1.58. On omissions of up to two fields and on bare lies some comparators are evaded less often (Table G.4: on omissions the residual MLP +0.026, the residual one-class SVM +0.143 and the Mahalanobis distance +0.119; on bare lies residual LightGBM +0.052 and residual TabPFN v2 +0.105). A withheld field masks the axioms that read it (Proposition 1); with the optional critical-field alert, which names the undeclared field and the constraints it leaves unchecked, NeSySpec is evaded less than the residual MLP with the same alert on unseen families (−0.074) and more on the unseen class (+0.125). On this fleet a rule without the knowledge base (an alert on any undeclared field that the record's class, or its engine type for classes with fewer than five training records, declares in at least 95% of the training records, whether or not an applicable axiom reads it) gives the alert's protection with the same false alerts on the genuine records; NeSySpec's report also names the constraints that each withheld field leaves unchecked.

G.8 Ablation and Other Forms of the Same Knowledge

Training the knowledge base through Real Logic is ahead of the same knowledge base at its initial values on every split (+0.073, +0.061 and +0.088), and the full model is ahead of the saturating-envelope ablation and of the ablation with flat banks on every split; Table G.8 gives each ablation with its difference and both intervals, Section 4.8 states the reading, and the textbook prior, a sensitivity, is on par with the full model on test and the unseen class and close to it on unseen families. On satellite pairs, a single masked quantifier in place of the strata ranks unseen-site pairs higher (differences −0.008 with 30% of the images withheld and −0.013 with all images) and is close to the fused verdict on test, while NeSySpec's strata keep one verdict per evidence source, so the report says whether the record or the image raised the flag.

The knowledge encoded. To measure what each constraint contributes, NeSySpec was retrained with subsets of its 12 specification constraints. A removed constraint is unavailable for every record in training and scoring, while the record labels of training stay unchanged. The record verdict was then scored on intervention labels, which no constraint defines, and on generator labels, which the same constraints define, so that a removal hides the corruptions of its constraint by construction (Figure 10 in Section 4.8, all panels on intervention labels; Table G.9). Over 96 subsets of 1 to 11 constraints, all 12 of sizes 1 and 11 and 8 random ones of each size from 2 to 10, with 3 seeds each, and the full knowledge base, the intervention-label AUROC averaged over test, unseen families and the unseen class has a mean per size that rises from one size to the next, from 0.535 with one constraint to 0.723 with all twelve (Spearman 0.93 [0.90, 0.95]), by steps whose intervals include zero at 5 of the 11 steps. The returns diminish: a logarithmic trend fits better than a linear one (R^2 0.87 against 0.84; on the unseen class alone the two fit equally), the gain is about 0.028 per constraint up to six and 0.008 beyond, and the twelfth still adds a small gain (0.009 [0.003, 0.015]). The AUROC grows linearly with the share of edits that alter a field read by some active constraint (R^2 0.88, against 0.74 for a logarithmic trend), a share that goes from 0.22 with one constraint to all edits with twelve, while on the edits of fields the constraints read it stays between 0.70 and 0.75 at every size: the gain follows the coverage of the edits. The residual MLP on the same subsets gains from the same knowledge (0.540 to 0.706), with no trend in its gap to NeSySpec. Removing one constraint from the full knowledge base lowers the intervention-label AUROC for A5 and A2, slightly for A4, A7 and A8 and for A4 with A4P together, and leaves it on par for A1, A4P, A9 and A10 and close for A3, A6 and A11; removing A1 lowers it on the unseen class and raises it slightly on test, and no removal raises it averaged over the three splits. Stall speed, engine count and fuselage length, the fields that only A6, A9 and A11 read, are never edited alone in the intervention records, and stall speed is never edited at all; A3, A4P and A10 read only fields that other constraints also read, as do A2, A4, A7 and A8, whose removal still lowers the AUROC, slightly for A4, A7 and A8. On generator labels the removal of A1, A2, A4, A4P, A5 or A8 lowers the AUROC, that of A7 or A10 slightly, and that of A3 raises it slightly (0.001); abduction is lower without A1 (0.187), A5 (0.077), A3, A4, A4P, A8 and A11, slightly lower without A2, A7 and A10, and on par without A6 and A9. In Figure 10c the darker line gives the aircraft interval and the lighter the family interval, with filled points where both exclude zero.

G.9 Selected Components and Sensitivity Analyses

Selected components. The diagnosis key was selected on held-out training classes: with each of the six training classes left out in turn, the verifier retrained and every statistic, prior and head recomputed without it, the largest-violation key ranked a violated constraint first in 0.802 of the held-out corrupted records (mean over the six classes), against 0.630 for the second head, with a higher mean on each held-out class (differences from 0.006 to 0.432); on validation records of the seen classes the two keys differ by 0.0008, and the second head is kept as the secondary key. NeSySpec's optional critical-field alert was selected because, on validation, it lowered evasion by an attacker who knows the alert from 0.949 to 0.739 and fired on 2 of the 17 genuine validation records, and because, with each training class held out in turn, it fired on 0.028 of that class's genuine records (mean over six classes; limit 0.25; largest rate 0.091, against 0.263 on the unseen class). Validation selected the abduction rule,

Table G.9: Specification constraints removed from the knowledge base.

	Detection AUROC, intervention labels		Detection AUROC, generator labels		Abduction top-1, intervention labels	
Removed	Drop	Intervals	Drop	Intervals	Drop	Intervals
None (level)	0.723		0.968		0.542	
A5	+0.033	[+0.016, +0.053]	+0.059	[+0.042, +0.079]	+0.077	[+0.054, +0.100]
		[+0.017, +0.060]		[+0.043, +0.078]		[+0.049, +0.098]
A2	+0.022	[+0.011, +0.034]	+0.041	[+0.031, +0.053]	+0.035	[−0.000, +0.071]
		[+0.003, +0.035]		[+0.032, +0.056]		[+0.006, +0.075]
A1	+0.017	[−0.004, +0.039]	+0.136	[+0.120, +0.150]	+0.187	[+0.161, +0.214]
		[−0.006, +0.038]		[+0.119, +0.153]		[+0.164, +0.212]
A4 and A4P	+0.014	[+0.003, +0.026]	+0.032	[+0.021, +0.044]	+0.052	[+0.024, +0.078]
		[+0.001, +0.029]		[+0.022, +0.049]		[+0.023, +0.081]
A4	+0.013	[+0.004, +0.023]	+0.019	[+0.011, +0.029]	+0.034	[+0.015, +0.055]
		[+0.003, +0.025]		[+0.009, +0.030]		[+0.011, +0.054]
A8	+0.007	[+0.000, +0.014]	+0.022	[+0.013, +0.032]	+0.030	[+0.012, +0.050]
		[−0.002, +0.015]		[+0.011, +0.030]		[+0.014, +0.052]
A10	+0.003	[−0.003, +0.008]	+0.007	[+0.002, +0.013]	+0.017	[+0.001, +0.035]
		[−0.005, +0.007]		[+0.002, +0.013]		[+0.000, +0.033]
A7	+0.002	[+0.001, +0.004]	+0.002	[+0.001, +0.004]	+0.011	[+0.003, +0.021]
		[+0.001, +0.004]		[+0.001, +0.004]		[+0.000, +0.021]
A9	+0.001	[−0.003, +0.005]	−0.003	[−0.008, +0.003]	+0.001	[−0.006, +0.009]
		[−0.004, +0.005]		[−0.008, +0.003]		[−0.008, +0.008]
A6	+0.000	[−0.000, +0.001]	−0.000	[−0.001, +0.000]	−0.000	[−0.006, +0.005]
		[−0.000, +0.001]		[−0.001, +0.000]		[−0.006, +0.005]
A11	−0.000	[−0.001, +0.001]	−0.001	[−0.002, +0.001]	+0.032	[+0.021, +0.043]
		[−0.001, +0.002]		[−0.002, +0.000]		[+0.021, +0.044]
A4P	−0.001	[−0.006, +0.005]	+0.017	[+0.010, +0.024]	+0.024	[+0.005, +0.044]
		[−0.006, +0.007]		[+0.011, +0.028]		[+0.006, +0.045]
A3	−0.001	[−0.002, +0.001]	−0.001	[−0.003, −0.000]	+0.010	[+0.004, +0.017]
		[−0.002, +0.001]		[−0.002, −0.000]		[+0.004, +0.015]

Drop: the full knowledge base minus the knowledge base without the constraint, 8 seeds each, paired on the same records; detection averaged over test, unseen families and the unseen class, abduction over the 695 single-field intervention records of the three splits. Intervals: [aircraft] above [family], 95% cluster-bootstrap percentiles, 2,000 draws. Generator labels are defined by the same constraints, so a removal hides the corruptions of its constraint by construction.

the smallest-change rule of Table F.2, and its feasibility threshold 0.9, also chosen in 7 of 8 leave-one-seed-out selections. With threshold 0.8, abduction reaches 0.750 on single-field test corruptions, its readings against minimum-cardinality localization unchanged and still ahead of the Class I residual on test. The largest-feasibility rule of Table G.2, which ranks every field by the lowest verdict it reaches, gives 0.692 on test, 0.747 on unseen families and 0.660 on the unseen class, on par with the Class I residual on test. On intervention labels every reading holds under both alternatives except the margin over the oracle tie-break (+0.112 and +0.095). For a class without training aircraft NeSySpec reads fleet statistics; five alternative reference statistics (engine-type or declared-mass pools with the union or weighted mean of class bands, and the nearest training class), compared on held-out training classes, met no adoption rule, and on the business jets each flags every clean record (1.000 against 0.322 with fleet statistics) with a detection AUROC of 0.836 to 0.906 against 0.934; under the mass pools and the nearest class a withheld field changes the reference class of 16.0, 21.7 and 1.8% of masked held-out records.

Sensitivity analyses. On generator labels, every band alternative that places more of NeSySpec's learned bands inside the training fleet range widened by 12%, [0.88 min, 1.12 max], lowers validation detection by more than 0.005: the learned bands act as tolerances learned from the fleet, 45.7% of them with both edges inside that range. The alternatives project the band edges onto that range after every step, or pull them toward it with a penalty of weight 0.01, 0.1 or 1.0 (eight seeds each). Retrained on intervention labels, the four alternatives hold 68.5, 86.7, 96.5 and 100% of the 46 counted bands per seed (those with a training fleet range) inside the range and stay within the validation margin; the projection reaches 0.804 on test, against 0.807 for NeSySpec retrained on the same labels. The penalty of weight 0.01, the alternative selected on validation, keeps its share at 45.7% and was not adopted; it reaches on generator-label test a detection AUROC of 0.977, on par with 0.979, the same diagnosis top-1 (0.942), a localization top-1 of 0.728 against 0.742 and a false-positive rate on the controls of 0.082 against 0.086, and on the unseen class a detection AUROC slightly lower, 0.932 against 0.934;

with weight 1.0 (53.8% inside) unseen-class detection is lower, 0.924 against 0.934. These alternatives are sensitivity analyses of NeSySpec.

Oracle bands fitted on the training aircraft. With the class bands of the generator oracle fitted on the 44 clean training aircraft alone (business jets, which have no training aircraft, on the band of all training aircraft), and the generator, its seed and the family splits unchanged, the benchmark holds 1,239 instances instead of 1,312. These bands flag 38 of the 63 clean held-out aircraft, 15 of the 16 in test, mostly on A1, A9 and A11; controls must stay clean under the oracle, so the test controls fall from 32 to 4. NeSySpec was retrained with its selected configuration on 8 seeds, as were the residual MLP and TabPFN v2 with their full selection protocol and the DL2-loss network with its selected hyperparameters, whose search was not repeated; the crisp edits and minimum-cardinality localization were recomputed, and the rule diagnosis, TabPFN v2 on intervention labels and the alternative bootstrap streams were not rerun (Tables G.10 and G.11). The validation rule of the crisp edits now selects per-class bands without pooling (validation AUROC 0.919 against 0.915), the bands of the new label oracle itself, with which the crisp edits lead NeSySpec on generator labels in key and macro-F1 on every split; the tables keep the pooled bands of the paper, and on intervention labels NeSySpec is on par with the unpooled crisp edits in macro-F1 on test (-0.006) and ahead on the other two splits. On intervention labels, whose records are the same under both oracles, NeSySpec stays ahead on every split of the DL2-loss network in macro-F1 ($+0.045$ $[+0.012,\ +0.142] \cdot [+0.012,\ +0.163]$ on test), of Fellegi-Holt and Z3 in diagnosis ($+0.126$ and $+0.092$ on test) and of minimum-cardinality localization with random tie-break ($+0.226$ on test); on test it is on par with the oracle tie-break ($+0.079$) and with the crisp edits in macro-F1 (-0.016), and the DL2-loss network has the higher exact-set agreement (-0.149) and fewer false violations. NeSySpec's own test levels there are lower, with more false violations, every paired interval excluding zero: macro-F1 0.488 against 0.575, exact-set agreement 0.070 against 0.283, false-violation rate 0.219 against 0.084, key 0.732 against 0.786 and abduction 0.532 against 0.593, while detection is on par (0.695 against 0.701). On unseen families key, abduction and detection are on par and macro-F1 comparable (0.666 against 0.694), exact-set agreement is lower (0.255 against 0.373) and the false-violation rate slightly higher (0.088 against 0.028); on the unseen class the verdicts, abduction and detection are on par and the key is lower (0.743 against 0.778). Trained on the flags of these bands through its per-constraint targets, NeSySpec learns narrower bands where the oracle's bands narrow. Over the 46 counted bands, the median ratio of its learned widths to those learned under the paper's oracle is 0.903 (seed-bootstrap interval $[0.897,\ 0.911]$), and 64% of the bands narrow, with a moderate rank correlation with the oracle's width ratios over 45 class bands (Spearman 0.43). The narrowing is largest for the thrust bands of A9 and the passenger-density bands of A11 (median 0.71), which carry the largest rises in false violations on the 16 clean test aircraft ($+0.35$ and $+0.33$; 0.061 to 0.208 over all constraints), followed by A2 ($+0.26$), whose learned bands shift. Seven of its eight seeds reached their best epoch by epoch 22, against 80 to 100 for the paper's seeds; the seed whose best epoch is 93 shows a similar narrowing on the bands of A3 and the thrust bands of A9 (0.78 and 0.74) and a smaller one on the passenger-density bands of A11 (0.78). On generator labels the records differ between the two benchmarks: the held-out corrupted records keep the violations their clean origin already has under the narrower bands, and the clean record itself stays labeled clean. There the residual MLP is slightly higher in detection on test (0.950 against 0.910, -0.040 $[-0.075,\ -0.010] \cdot [-0.071,\ +0.001]$) and on unseen families (-0.013), as is TabPFN v2 on unseen families (-0.020), and stays slightly ahead on the unseen class (0.986 against 0.959), where TabPFN v2 is now on par (-0.008). The key reaches 0.944 on test and 0.931 on unseen families (0.973 before), slightly ahead of Fellegi-Holt and Z3 on test ($+0.102$) and on par with them on unseen families ($+0.094$); on the unseen class, where the pooled crisp bands equal those of the label oracle, Fellegi-Holt and Z3 lead (-0.094). The verdicts stay ahead of the DL2-loss network in macro-F1 on every split ($+0.288$ on test), are on par with the crisp edits on test and unseen families and below them on the unseen class (-0.125), with more false violations than the crisp edits on every split. Abduction is now ahead of the oracle tie-break on test ($+0.199$ $[+0.042,\ +0.370] \cdot [+0.052,\ +0.528]$) and stays ahead of random tie-breaking on every split. The repair still lowers the error on the altered field on every split (0.126 against 0.236 without repair on test); as the verifier now rates almost every clean and control

test record as violating, its target is reached on 0.036 of single-field test records (0.555 before) and it proposes a change for 0.988 of the clean and control test records (0.588 on unseen families and 0.707 on the unseen class); the repair rule re-selected on the new validation split (feasibility threshold 0.8 instead of 0.9) reaches the target on 0.190 and changes 0.894 of them. Bands fitted on all 107 aircraft cover the range of the fleet the records describe, which bands fitted on 44 aircraft do not.

Oracle fitted on training and validation aircraft, tolerances chosen out of family. Fitted on the 61 clean training and validation aircraft, the oracle had each of its three tolerances set, by a rule declared before any result, to the smallest value of a grid that flags at most 5% of the clean aircraft when their family is held out in turn: the range margin, the relative widening of each class band ([0.88 min, 1.12 max] at 0.12), becomes 0.50 instead of 0.12 (2 of 61 aircraft flagged, against 24 at 0.12); the equality factor, which multiplies the largest absolute log-residual of the clean records of each class for the identities A1, A5 and A6, becomes 1.25 instead of 1.15 (3 of 60, at the limit of the rule); and the mass-balance margin of A7, subtracted from the smallest clean mass-balance margin, stays at 0.05; with these values 5 of the 61 aircraft are flagged by some tolerance. This oracle flags 14 of the 46 clean test and out-of-distribution aircraft: 11 of the 16 in test, 10 of them narrowbodies on A1, whose class tolerance rests on 3 aircraft of 2 families, 3 business jets on A9, judged by the band fitted on all 61 aircraft, and none of the unseen families; its wider bands let the generator escalate further (median magnitude 0.8 against 0.4), and the benchmark holds 1,260 instances with 10 test controls. Retrained on its labels as above (Tables G.10 and G.11), NeSySpec on intervention labels keeps its false violations at the paper's level on test (0.073 against 0.084, on par), with fewer on the unseen class (0.087 against 0.128), and its exact-set agreement on par on test and unseen families and slightly higher on the unseen class (0.206 against 0.155); its macro-F1 is lower on test (0.537 against 0.575) and unseen families (0.636 against 0.694), its abduction on test (0.491 against 0.593) and unseen families (0.508 against 0.594), and its key on the unseen class (0.735 against 0.778) and slightly on test (0.738 against 0.786), while detection is on par on test and unseen families and close on the unseen class (0.711 against 0.719). Against the comparators on intervention labels it stays ahead on every split of Fellegi-Holt and Z3 in diagnosis and of minimum-cardinality localization with random tie-break, ahead of the DL2-loss network in macro-F1 on unseen families, slightly on test (+0.037) and on par on the unseen class, with the network keeping fewer false violations on test and the unseen class. In macro-F1 it is on par with the crisp edits on test and unseen families and slightly ahead on the unseen class (+0.029), and, against the bands with a 25% margin, small classes pooled, that the validation rule of the crisp edits now selects (validation AUROC 0.943 against 0.942), on par on test and ahead on the other two splits. In abduction it is on par with the oracle tie-break on every split, and against the Class I residual on par on test, lower on unseen families (−0.096) and ahead on the unseen class. On generator labels, with larger corruptions, NeSySpec reaches a detection AUROC of 0.976 on test, slightly below the residual MLP (0.994), and on the unseen class slightly below the residual MLP and TabPFN v2 (0.988 against 0.999 for both); its macro-F1 is on par with the DL2-loss network on test and the unseen class and ahead on unseen families, ahead of the crisp edits on test, slightly ahead on unseen families and below them on the unseen class; its key, 0.955 on test, is ahead of Fellegi-Holt on every split and of Z3 on test and unseen families, and on par with Z3 on the unseen class; abduction is on par with the oracle tie-break on every split and ahead of random tie-breaking; the repair lowers the error on the altered field on every split, reaches its target on 0.570 of single-field test records and changes 0.418 of the clean and control test records; and 109 of the 154 corrupted test records carry a violation their clean origin already has. Under these tolerances the bands widen, the training corruptions grow and training stops earlier (best epochs 15 to 27, against 80 to 100), so their effects on the key and abduction levels are not separated.

Table G.10: Generator oracles: levels of NeSySpec and its comparators.

	Oracle fitted on		
Reading	107 aircraft	44 aircraft	61 aircraft
Benchmark			
Instances	1,312	1,239	1,260
Test controls	32	4	10
Clean test aircraft flagged by the oracle, of 16	0	15	11
Intervention labels, same records: NeSySpec			
Constraint verdicts, macro-F1, test	0.575	0.488	0.537
unseen families	0.694	0.666	0.636
unseen class	0.431	0.444	0.400
Constraint verdicts, exact-set, test	0.283	0.070	0.223
unseen families	0.373	0.255	0.348
unseen class	0.155	0.153	0.206
Constraint verdicts, false-violation rate, test	0.084	0.219	0.073
unseen families	0.028	0.088	0.015
unseen class	0.128	0.120	0.087
Largest-violation key, top-1, test	0.786	0.732	0.738
unseen families	0.828	0.808	0.795
unseen class	0.778	0.743	0.735
Abduction, top-1, test	0.593	0.532	0.491
unseen families	0.594	0.580	0.508
unseen class	0.465	0.453	0.441
Detection AUROC, test	0.701	0.695	0.696
unseen families	0.750	0.747	0.737
unseen class	0.719	0.716	0.711
Intervention labels, test: NeSySpec minus the comparator			
DL2-loss network, macro-F1	+0.149	+0.045	+0.037
DL2-loss network, exact-set	+0.061	−0.149	−0.058
DL2-loss network, false-violation rate	+0.065	+0.194	+0.059
Crisp edits, small classes pooled, macro-F1	+0.071	−0.016	+0.032
Fellegi-Holt, key top-1	+0.180	+0.126	+0.132
Z3, key top-1	+0.146	+0.092	+0.098
Minimum cardinality, random tie-break, abduction	+0.287	+0.226	+0.186
Minimum cardinality, oracle tie-break, abduction	+0.140	+0.079	+0.038
Class I sizing-regression residual, abduction	+0.140	+0.079	+0.038
Generator labels, test and unseen class: different records under each oracle			
Detection AUROC, residual MLP	0.990	0.950	0.994
Detection AUROC, TabPFN v2 on the residuals	0.958	0.932	0.983
Detection AUROC, NeSySpec	**0.979**	**0.910**	**0.976**
Detection AUROC, unseen class, residual MLP	0.978	0.986	0.999
Detection AUROC, unseen class, TabPFN v2 on the residuals	0.983	0.967	0.999
Detection AUROC, unseen class, NeSySpec	**0.934**	**0.959**	**0.988**
Constraint verdicts, DL2-loss network	0.474/0.661/0.020	0.145/0.456/0.022	0.431/0.760/0.012
Constraint verdicts, crisp edits, small classes pooled	0.569/0.829/0.036	0.421/0.664/0.030	0.256/0.662/0.059
Constraint verdicts, NeSySpec	**0.504/0.795/0.074**	**0.176/0.745/0.187**	**0.360/0.773/0.059**
Key top-1, chance	0.177	0.373	0.243
Key top-1, Fellegi-Holt	0.795	0.842	0.733
Key top-1, Z3	0.825	0.842	0.822
Key top-1, NeSySpec	**0.942**	**0.944**	**0.955**
Abduction top-1, minimum cardinality, random tie-break	0.367	0.192	0.423
Abduction top-1, minimum cardinality, oracle tie-break	0.629	0.421	0.590
Abduction top-1, NeSySpec	**0.742**	**0.620**	**0.706**
Repair, log error on the altered field, no repair	0.366	0.236	0.511
Repair, log error on the altered field, NeSySpec	**0.133**	**0.126**	**0.169**
Repair, target reached on single-field records, NeSySpec	0.555	0.036	0.570
Repair, clean and control records changed, NeSySpec	0.562	0.988	0.418

Oracle fitted on all 107 aircraft (the paper), on the 44 training aircraft with the paper's tolerances, and on the 61 training and validation aircraft with tolerances chosen by leaving one family out (range margin 0.50, equality factor 1.25, mass-balance margin 0.05). Learned methods: means over 8 seeds. Intervention labels, set from the edit alone, are the same records under every oracle; generator labels differ between the three benchmarks. Constraint verdicts on generator labels: exact-set agreement / macro-F1 / false-violation rate. On intervention labels only NeSySpec and the DL2-loss network change; the crisp edits, Fellegi-Holt, Z3, minimum cardinality and the Class I residual are the same under every oracle, and since the oracle tie-break and the Class I residual reach the same intervention-label test level (0.453), their differences from NeSySpec coincide. Intervals of the intervention-label rows in Table G.11; those of the differences under the oracle on 107 aircraft in the DL2-loss macro-F1, the key and abduction are in Tables G.1 to G.4 and Appendix G.1.

Table G.11: Generator oracles fitted without test or out-of-distribution aircraft: paired changes and differences with their intervals.

	Oracle on 44 aircraft		Oracle on 61 aircraft	
Reading	Value	Intervals	Value	Intervals
Intervention labels, same records: NeSySpec, change from the oracle on 107 aircraft				
Constraint verdicts, macro-F1, test	−0.087	[−0.173, −0.022]·[−0.146, −0.029]	−0.039	[−0.109, −0.007]·[−0.098, −0.0034]
unseen families	−0.028	[−0.054, +0.017]·[−0.061, +0.003]	−0.058	[−0.105, −0.022]·[−0.116, −0.026]
unseen class	+0.013	[−0.010, +0.041]·[−0.011, +0.039]	−0.032	[−0.059, +0.003]·[−0.060, +0.002]
Constraint verdicts, exact-set, test	−0.213	[−0.309, −0.125]·[−0.279, −0.117]	−0.060	[−0.164, +0.032]·[−0.187, +0.043]
unseen families	−0.118	[−0.226, −0.005]·[−0.188, −0.008]	−0.026	[−0.083, +0.027]·[−0.108, +0.032]
unseen class	−0.002	[−0.040, +0.044]·[−0.041, +0.041]	+0.050	[+0.001, +0.109]·[−0.007, +0.118]
Constraint verdicts, false-violation rate, test	+0.135	[+0.092, +0.178]·[+0.068, +0.168]	−0.011	[−0.044, +0.025]·[−0.049, +0.030]
unseen families	+0.061	[+0.011, +0.111]·[−0.011, +0.099]	−0.013	[−0.032, +0.002]·[−0.042, +0.003]
unseen class	−0.007	[−0.024, +0.010]·[−0.024, +0.010]	−0.041	[−0.059, −0.023]·[−0.061, −0.021]
Largest-violation key, top-1, test	−0.054	[−0.094, −0.019]·[−0.109, −0.023]	−0.048	[−0.091, −0.007]·[−0.101, −0.002]
unseen families	−0.020	[−0.059, +0.016]·[−0.055, +0.045]	−0.033	[−0.063, +0.004]·[−0.060, +0.035]
unseen class	−0.035	[−0.066, −0.004]·[−0.065, −0.009]	−0.043	[−0.072, −0.018]·[−0.070, −0.017]
Abduction, top-1, test	−0.061	[−0.101, −0.024]·[−0.102, −0.030]	−0.102	[−0.151, −0.059]·[−0.173, −0.063]
unseen families	−0.014	[−0.048, +0.021]·[−0.036, +0.043]	−0.086	[−0.136, −0.043]·[−0.129, −0.060]
unseen class	−0.012	[−0.045, +0.025]·[−0.044, +0.025]	−0.024	[−0.077, +0.033]·[−0.077, +0.035]
Detection AUROC, test	−0.006	[−0.022, +0.008]·[−0.026, +0.007]	−0.005	[−0.022, +0.011]·[−0.030, +0.009]
unseen families	−0.003	[−0.031, +0.025]·[−0.025, +0.030]	−0.013	[−0.035, +0.005]·[−0.043, +0.017]
unseen class	−0.003	[−0.016, +0.005]·[−0.015, +0.005]	−0.008	[−0.020, −0.000]·[−0.020, +0.000]
Intervention labels: NeSySpec minus the comparator under each oracle				
DL2-loss network, macro-F1, test	+0.045	[+0.012, +0.142]·[+0.012, +0.163]	+0.037	[+0.002, +0.113]·[+0.010, +0.143]
unseen families	+0.220	[+0.184, +0.277]·[+0.209, +0.267]	+0.129	[+0.094, +0.168]·[+0.101, +0.172]
unseen class	+0.141	[+0.106, +0.185]·[+0.111, +0.182]	+0.000	[−0.029, +0.040]·[−0.028, +0.038]
DL2-loss network, exact-set, test	−0.149	[−0.203, −0.095]·[−0.207, −0.077]	−0.058	[−0.148, +0.026]·[−0.199, +0.050]
unseen families	+0.041	[−0.065, +0.143]·[−0.016, +0.134]	+0.081	[+0.023, +0.133]·[+0.033, +0.152]
unseen class	−0.093	[−0.166, −0.016]·[−0.163, −0.013]	−0.033	[−0.101, +0.037]·[−0.105, +0.034]
DL2-loss network, false-violation rate, test	+0.194	[+0.147, +0.242]·[+0.126, +0.228]	+0.059	[+0.022, +0.100]·[+0.021, +0.101]
unseen families	+0.070	[+0.031, +0.112]·[+0.017, +0.102]	+0.010	[−0.004, +0.028]·[−0.009, +0.032]
unseen class	+0.102	[+0.052, +0.152]·[+0.053, +0.152]	+0.064	[+0.020, +0.109]·[+0.023, +0.112]
Crisp edits, small classes pooled, macro-F1, test	−0.016	[−0.061, +0.053]·[−0.072, +0.045]	+0.032	[−0.025, +0.098]·[−0.036, +0.088]
unseen families	+0.097	[+0.025, +0.178]·[+0.006, +0.177]	+0.067	[−0.007, +0.142]·[−0.039, +0.148]
unseen class	+0.073	[+0.033, +0.103]·[+0.035, +0.103]	+0.029	[+0.001, +0.057]·[+0.002, +0.056]
Fellegi-Holt, key top-1, test	+0.126	[+0.086, +0.170]·[+0.097, +0.193]	+0.132	[+0.084, +0.186]·[+0.092, +0.227]
unseen families	+0.183	[+0.114, +0.256]·[+0.139, +0.280]	+0.170	[+0.103, +0.244]·[+0.118, +0.268]
unseen class	+0.231	[+0.173, +0.292]·[+0.175, +0.298]	+0.223	[+0.159, +0.289]·[+0.161, +0.296]
Z3, key top-1, test	+0.092	[+0.047, +0.153]·[+0.038, +0.164]	+0.098	[+0.046, +0.168]·[+0.040, +0.187]
unseen families	+0.179	[+0.101, +0.279]·[+0.129, +0.250]	+0.166	[+0.093, +0.266]·[+0.101, +0.264]
unseen class	+0.145	[+0.107, +0.186]·[+0.107, +0.186]	+0.137	[+0.094, +0.181]·[+0.095, +0.183]
Minimum cardinality, random tie-break, abduction, test	+0.226	[+0.160, +0.295]·[+0.176, +0.351]	+0.186	[+0.119, +0.256]·[+0.125, +0.302]
unseen families	+0.269	[+0.164, +0.389]·[+0.181, +0.471]	+0.197	[+0.063, +0.336]·[+0.066, +0.402]
unseen class	+0.288	[+0.215, +0.359]·[+0.219, +0.364]	+0.276	[+0.185, +0.363]·[+0.191, +0.371]
Minimum cardinality, oracle tie-break, abduction, test	+0.079	[−0.013, +0.185]·[−0.017, +0.299]	+0.038	[−0.052, +0.142]·[−0.048, +0.248]
unseen families	+0.120	[−0.033, +0.290]·[−0.003, +0.440]	+0.048	[−0.123, +0.230]·[−0.084, +0.350]
unseen class	+0.145	[+0.015, +0.273]·[+0.022, +0.280]	+0.133	[−0.013, +0.277]·[−0.003, +0.286]
Class I sizing-regression residual, abduction, test	+0.079	[+0.022, +0.131]·[+0.016, +0.115]	+0.038	[−0.016, +0.095]·[−0.020, +0.064]
unseen families	−0.024	[−0.098, +0.049]·[−0.071, +0.098]	−0.096	[−0.171, −0.033]·[−0.144, −0.022]
unseen class	+0.362	[+0.291, +0.433]·[+0.292, +0.432]	+0.350	[+0.282, +0.424]·[+0.283, +0.423]

Intervals: [aircraft]·[family], 95% cluster-bootstrap percentiles; the first panel gives the paired change of NeSySpec on the same records, the second the difference at each oracle (Table G.10 gives the levels and the differences under the oracle on 107 aircraft).

G.10 Training Checkpoints

Figure 9 in Section 4.8 shows the training curves of the axioms; two readings complete it. At the frozen checkpoint, restored at the best validation epoch, the mean truth of each specification constraint is lower on corrupted than on clean-labeled records in 46 of 46 split and constraint pairs with available evidence, by as little as 1.0×10^{-5} on the unseen class. With the train prior (8 seeds), the best epoch, selected by early stopping on the validation AUROC (Table F.2), reaches the declared cap of 100 epochs in three of eight seeds.

H Explanation-Supervision Comparison

This appendix gives the results and the protocol of the comparison with explanation supervision summarized in Section 4.6. Under the protocol described at the end of this appendix, the grounding axiom localizes the annotated parts more than RRR [57], GradMask [58] and Pukdee et al. [61], with every cluster-bootstrap interval of the difference above zero on the FGVC pseudo-masks, on PASCAL-Part human masks (levels in Figure G.1b) and on two further human-mask datasets. Against RRR, GradMask and Pukdee et al., in this order, its pointing exceeds theirs by +0.725 [+0.700, +0.750], +0.779 and +0.618 on the FGVC test pseudo-masks, by +0.594 [+0.559, +0.630], +0.702 and +0.496 on PASCAL-Part, by +0.523 [+0.475, +0.571], +0.631 and +0.420 on PartImageNet and by +0.483 [+0.359, +0.595], +0.636 and +0.446 on the 40 items of UDA-Part, with wider intervals on the last two; Pukdee et al. is run at the calibration of the protocol, which flattened their teacher toward the uniform prediction. The intervals of all 27 differences, nine mask configurations by three baselines, lie above zero, also under an independent bootstrap seed. Across the weights 0.3, 1 and 3 the axiom keeps at most one collapsed head, a head whose F1 is at most 0.05, against up to 5.7 on average for the auxiliary loss. At weight 1 it keeps an out-of-distribution accuracy lead of +0.016 [+0.012, +0.020] over the auxiliary loss at its validation weight 0.3; test accuracy, by contrast, is degenerate (a constant predictor scores 0.926).

Protocol. A first set of runs of RRR, GradMask, RES [59] (its RES-L variant, adapted to these part masks) and Pukdee et al. used $\lambda = 1$ and penalties coupled across a batch through batch normalization [94] in training mode. The comparison retrains RRR, GradMask and Pukdee et al. with per-example penalties, the normalization constants held fixed, selects $\lambda \in \{0.1, 1, 10\}$ times a calibration constant on validation with one seed per value, and evaluates three seeds at the selected value. RES-L is not retrained: on these masks its adapted objective does not reward localization, so its pointing does not measure what RES can localize. A second arm of Pukdee et al., added after the results, measures its calibration constant at the first projection step instead of at initialization; it avoids the flattened teacher and weights the constraint about three orders of magnitude less. It points to the part in 0.179 of PASCAL-Part human masks, where the axiom's lead over it is +0.558 [+0.524, +0.591], and the intervals of the axiom's differences from it also lie above zero on the FGVC pseudo-masks and on every other human-mask set. Holding the normalization constants fixed makes each penalty per-example and leaves the network in inference mode unconstrained by it, a gap that a diagnostic measures for RRR.