SetGRPO: Calibrating Grader-Induced Policy Optimization

Jiale Huang[1] Zixu Li[1,†] Zhiwei Chen[1] Zhiheng Fu[1] Jinhe Bi[2] Yupeng Hu[1,✉]
[1]Shandong University [2]National University of Singapore
huangjiale.cs@outlook.com, {lizixu.cs, zivczw, fuzhiheng8}@gmail.com
bijinhe@outlook.com, huyupeng@sdu.edu.cn

Code: https://github.com/ArcadiaDream/SetGRPO

arXiv:2610.36807v1 [cs.LG] 27 Oct 2026

*Rubric-based reinforcement learning uses fine-grained grader feedback to train large language models on open-ended tasks, but grader randomness can alter GRPO updates in ways that score discrepancies alone cannot characterize. Because GRPO jointly normalizes rewards within each response group and combines the resulting training coefficients with response-level gradients, a scoring discrepancy can shift the coefficients of the entire group and induce policy-dependent update changes. Selective execution further affects the direction-reversal risk of retained updates. To address this, we propose **Set**-Valued **G**roup **R**elative **P**olicy **O**ptimization (**SetGRPO**). It comprises three modules: (a) Joint Reward Set Calibration constructs a joint uncertainty set for the training coefficients; (b) Geometry-Guided Safe Recovery maps candidate coefficients into the current policy's local update geometry and determines the largest safe recovery ratio for each response group; and (c) Selective Risk-Gated Optimization selects a global threshold through independent risk calibration to control the conditional direction-reversal risk of accepted updates while preserving the original GRPO objective. Under exchangeability, SetGRPO provides finite-sample marginal coverage for the joint coefficient set and simultaneous risk control across candidate thresholds. On ten open-ended benchmarks spanning three domains, SetGRPO achieves the highest macro average among all compared training methods on both Qwen3.5 and Llama3.2 model families.*

1. Introduction

Nowadays, Rubric-based Reinforcement Learning (Rubric-based RL) has emerged as an important training paradigm for improving the performance of large language models on open-ended generation tasks (Li et al., 2026a; Zhou et al., 2026; Kim et al., 2024; Ye et al., 2024). It uses fine-grained rubrics and LLM judges to evaluate open-ended responses across dimensions such as relevance, completeness, factuality, and style (Li et al., 2026a; Zhou et al., 2026; Kim et al., 2024; Ye et al., 2024; Liu et al., 2023; Zheng et al., 2023). These dimension-level scores are aggregated into response-level rewards that serve as training signals for reinforcement learning (Ouyang et al., 2022; Stiennon et al., 2020), thereby accommodating grader's score discrepancies during training and enabling optimization of open-ended generation tasks. Related work has also explored multimodal fusion, video-language adaptation, and parameter-efficient training for large models (Li et al., 2024; Hu et al., 2025; Li et al., 2026b).

Existing rubric-based RL studies build on GRPO (Shao et al., 2024; Guo et al., 2025) and examine reward over-optimization (Gao et al., 2023; Lambert et al., 2025), LLM-judge consistency and bias (Liu et al., 2023; Zheng et al., 2023; Wang et al., 2024), fine-grained evaluation (Kim et al., 2024; Ye et al., 2024; Liu et al., 2023), rubric constructions and rubric-based training for open-ended tasks (Li et al., 2026a; Zhou et al., 2026). Recent reinforcement-learning studies have also explored how to improve the utility of sampled trajectories through reflection, trajectory reuse, and rollout transformation (Bi et al., 2026a,b). Other work examines selection and self-evaluation mechanisms

† Project Leader: Zixu Li. ✉Corresponding Author: Yupeng Hu.

in multimodal or embodied settings (Aniri et al., 2026a; Bi et al., 2025a,b). Related studies further investigate reasoning, logical behavior, and the gap between model-based preference and real user engagement (Zhang et al., 2026a,b,c). However, these studies largely focus on the construction, reuse, or assessment of training signals and do not characterize how stochastic grader scores propagate through GRPO's within-group normalization and gradient aggregation to alter local policy updates under the current gradient geometry. Specifically, for a given response group, GRPO derives the training coefficients from the shared mean and scale of the group rewards and combines them with the response-level gradients of each response under the current policy to form a local policy update. We therefore conjecture that *a change in a single reward alters multiple training coefficients within the group through the shared statistics, and that different policy snapshots induce different update outcomes*.

To verify this conjecture, we score fixed response groups twice with the same grader and compare the two scoring results pairwise (detailed in Section 3). As shown in Figure 1(a), we first compare score discrepancies with the coefficient residuals obtained after within-group normalization, testing whether shared statistics propagate score discrepancies to the entire group of coefficients. As shown in Figure 1(b), we further compare the local updates of the same response group under different policy snapshots, testing whether the effect of coefficient discrepancies depends on the gradient geometry of the current policy. The results show that shared normalization propagates score discrepancies to multiple coefficients within a group, whereas similar coefficient shifts can induce different update directions under different policy geometries, thereby changing which groups are retained by selective execution. These findings support our conjecture and motivate the following question: *how can we characterize the effect of the training coefficients used by the optimizer on the current policy update while preserving effective training progress and avoiding directional deviation?*

To address this problem, we propose **Set**-Valued **G**roup **R**elative **P**olicy **O**ptimization (**SetGRPO**), which models GRPO's within-group training coefficients as a joint uncertainty set and uses it to select response groups and construct policy updates. Concretely, SetGRPO comprises three components: (1) *Joint Reward Set Calibration*, which uses paired reward discrepancies and independent calibration data to characterize the joint coefficient shifts induced by within-group normalization; (2) *Geometry-Guided Safe Recovery*, which maps candidate coefficients through the current policy's response-level gradients and determines the largest direction-safe recovery ratio; and (3) *Selective Risk-Gated Optimization*, which selects a global threshold on an independent risk-calibration set to control the conditional direction-reversal risk of accepted updates. During training, only groups satisfying this threshold contribute their recovered coefficients to the original GRPO objective, preserving effective updates while limiting overall update risk.

Our main contributions are summarized as follows:

- We first investigate the propagation mechanism of rubric score discrepancies in GRPO: a score discrepancy couples the training coefficients of an entire group through the shared mean and scale, and then jointly with the response-level gradients of the current policy determines the local update, yielding different update magnitudes and directions.
- We formalize this propagation process as a training-coefficient-driven policy-update problem and propose SetGRPO, which conformally calibrates GRPO's joint coefficient space and combines policy-update geometry with selective risk gating to limit direction deviation while preserving effective training progress.
- We systematically evaluate SetGRPO on ten open-ended task benchmarks across three domains with the Qwen3.5 and Llama-3.2 model families. The results show that SetGRPO achieves the best performance among all compared training methods on both model families, confirming its effectiveness in improving the reliability of GRPO training under rubric score discrepancies.

2. Related Work

Reinforcement Learning for Large Language Models and Rubric Feedback. Alignment training for large language models has evolved from preference learning and supervised fine-tuning to human- or AI-feedback-based instruction alignment (Ouyang et al., 2022; Stiennon et al., 2020; Christiano et al., 2017; Ziegler et al., 2019; Bai et al., 2022), giving rise to PPO and reward-model-free DPO (Schulman et al., 2017; Rafailov et al., 2023). For reasoning tasks, verifiers and process supervision provide fine-grained training signals (Cobbe et al., 2021; Uesato et al., 2022; Lightman et al., 2024),while DeepSeekMath and DeepSeek-R1 apply GRPO to large-language-model reasoning (Shao et al., 2024; Guo et al., 2025). Beyond standard single-agent training, recent work has also studied group-agent interaction and long-horizon social simulation in multi-agent settings (Zhang et al., 2026d,e; Song et al., 2026). Related work has also explored composed image retrieval and egocentric reasoning (Huang et al., 2025; Li et al., 2026c).For open-ended tasks without unique reference answers, G-Eval, FLASK,MT-Bench, and Prometheus use rubrics or LLM judges for structured evaluation (Kim et al., 2024; Ye et al., 2024; Liu et al., 2023; Zheng et al., 2023),while related retrieval studies consider semantic enhancement, focus revision,and multi-modification alignment (Huang et al., 2026a; Chen et al., 2025; Li et al., 2026d). RewardBench, RubricHub, and RuscaRL assess reward-model capability or use rubric feedback for reinforcement learning (Li et al., 2026a; Zhou et al., 2026; Lambert et al., 2025). Despite evidence of reward overoptimization and systematic judge bias (Zheng et al., 2023; Gao et al., 2023; Wang et al., 2024), prior work largely treats grader outputs as fixed training signals and does not trace how the score discrepancies propagate through GRPO's group normalization and response-level gradient aggregation to alter the actual policy update (Bi et al., 2026c; Aniri et al., 2026b), while related robustness-oriented studies in multimodal retrieval address noisy supervision and distributional variation through noise mitigation and invariance modeling (Chen et al., 2026a; Huang et al., 2026b; Fu et al., 2026; Li et al., 2026e).

Conformal Prediction and Selective Risk Control. The geometry of reasoning trajectories has also been studied through layerwise trajectory evolution and self-evaluation (Bi et al., 2026c). Conformal prediction exploits exchangeability to construct, from calibration samples, prediction sets with finite-sample coverage guarantees (Shafer and Vovk, 2008; Angelopoulos and Bates, 2023), and has extended from regression and classification coverage to input-dependent uncertainty and general risk functions (Lei et al., 2018; Romano et al., 2019, 2020; Bates et al., 2021; Angelopoulos et al., 2024). Selective prediction controls the risk on an accepted subset through abstention or sample screening (Bi et al., 2025a; Geifman and El-Yaniv, 2019).Beyond this setting, related video and egocentric systems have also studied temporal reranking, scene adaptation, and reliability-aware action composition (Hu et al., 2026; Chen et al., 2026b; Li et al., 2026f).The Clopper–Pearson interval provides a classical one-sided upper confidence bound for finite-sample binomial risk(Clopper and Pearson, 1934). SetGRPO draws on these ideas of calibration and risk control, but shifts the object of calibration to the joint training coefficients actually used by GRPO and the policy updates they induce, and further controls the conditional direction-reversal risk over the set of accepted updates, while related work has also considered calibration and reliability in retrieval and search settings (Li et al., 2026g; Zhang et al., 2026f; Yang et al., 2026).

3. Preliminary

In rubric-based reinforcement learning, grader rewards are normalized within each GRPO response group to produce training coefficients, which are combined with the current policy's response-level gradients to form local updates. We therefore analyze score discrepancies along this reward-to-update chain, first formalizing the mapping and then evaluating its training consequences through paired pilot experiments.

3.1. Rubric-based GRPO

For a problem x, the rollout policy $\pi_{\theta_{\text{old}}}$ generates a group of K responses, over which the grader produces response-level rewards according to the given rubric. Denote the response group, the reward vector, and the nominal training coefficients of GRPO as

$$\mathbf{Y}_x = (y_{x,1}, \ldots, y_{x,K}), \qquad \mathbf{r}_x = (r_{x,1}, \ldots, r_{x,K})^\top, \qquad \hat{\mathbf{a}}_x = T_{\text{train}}(\mathbf{r}_x), \tag{1}$$

where T_{train} denotes the training mapping of the original GRPO that group-centers and scale-normalizes a group of rewards, and $\hat{\mathbf{a}}_x$ denotes the training coefficients induced by the rewards produced by the grader. Because the responses within a group share the same center and scale, each $\hat{a}_{x,i}$ is determined jointly by the rewards of the entire group.

These coefficients then enter the original GRPO objective as response-level advantages. We define the group-level objective as

$$\mathcal{J}_{\text{GRPO}}(\boldsymbol{\theta}; \hat{\mathbf{a}}_x) = \mathbb{E}_x \left[\frac{1}{K} \sum_{i=1}^{K} \frac{1}{|y_{x,i}|} \sum_{t=1}^{|y_{x,i}|} \mathcal{L}_{\text{GRPO}} \left(\omega_{x,i,t}(\boldsymbol{\theta}), \hat{a}_{x,i} \right) \right], \tag{2}$$

where $\omega_{x,i,t}(\boldsymbol{\theta}) = \frac{\pi_{\boldsymbol{\theta}}(y_{x,i,t}|x, y_{x,i,<t})}{\pi_{\boldsymbol{\theta}_{\text{old}}}(y_{x,i,t}|x, y_{x,i,<t})}$ is the token-level probability ratio, and $\mathcal{L}_{\text{GRPO}}(\omega, a)$ denotes the token-level surrogate term retaining the probability ratio and the clipping mechanism.

These training coefficients determine each response's relative weight in the group objective but not the update direction, which also depends on the response-specific gradient directions and their mutual relations under the current policy. Let $\mathbf{g}_{x,i}$ denote the mean response log-probability gradient for $y_{x,i}$. Performing a first-order expansion around the current rollout policy and neglecting the learning rate and the common scaling across groups, the contribution of the coefficient vector $\mathbf{a}$ to the local update of the response group is

$$\mathbf{g}_{x,i} = \nabla_{\boldsymbol{\theta}} \left[\frac{1}{|y_{x,i}|} \sum_{t=1}^{|y_{x,i}|} \log \pi_{\boldsymbol{\theta}} \left(y_{x,i,t} \mid x, y_{x,i,<t} \right) \right], \quad \Delta\boldsymbol{\theta}_x(\mathbf{a}) \propto \sum_{i=1}^{K} a_i \mathbf{g}_{x,i}. \tag{3}$$

3.2. Preliminary Study on Grader-Induced Update Risk

The above equation reveals the propagation chain: *a score discrepancy changes the group's training coefficients, which modifies the local update through the current policy's gradients and ultimately affects which response groups are retained.* To assess this risk and motivate SetGRPO, we conduct pilot experiments asking whether (i) within-group normalization couples scoring errors, (ii) the effect of coefficient shifts depends on the policy state, and (iii) selective execution changes the aggregate direction-reversal risk.

Setup. We run original GRPO on Qwen3.5-9B with the RubricHub training set and collect response groups with their corresponding policy snapshots; training details are given in Section 5. For each group, we freeze the responses and their order, score it twice with the same grader, and map both scores through T_{train}. For the second scoring, we define $\mathbf{r}_x^\star$, $\mathbf{a}_x^\star = T_{\text{train}}(\mathbf{r}_x^\star)$, $\boldsymbol{\delta}_x = \mathbf{a}_x^\star - \hat{\mathbf{a}}_x$. Using the same policy snapshot, we compare the updates induced by the two coefficient sets and analyze coefficient-residual dependence, policy-state dependence, and aggregate risk after group-wise screening.

Does Grader Discrepancy Remain Response-wise Independent? By Eq. 1 and Eq. 2, group normalization uses a shared center and scale. Therefore, as shown in Figure 1(a), since T_{train} normalizes with a center and scale shared over the whole group, a change in the score of a single response

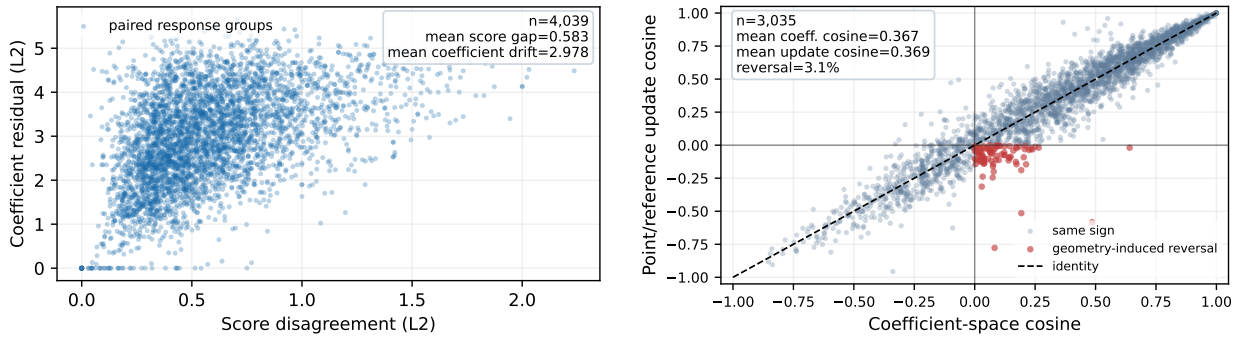

(a) Score disagreement propagates jointly. (b) Policy geometry changes local update alignment.

Figure 1. Illustration of score disagreement and its effect on local policy-update geometry.

affects the entire coefficient vector. Moreover, the corresponding residuals appear as a within-group correlated holistic shift rather than mutually independent response-wise deviations. The uncertainty involved in training should thus be modeled for the entire coefficient vector, rather than for each response separately.

Can Coefficient Drift Alone Characterize Update Safety? Fig. 1(b) applies the coefficients from the two scorings to the response-level gradients of the same response group and compares the local updates under different policy snapshots. The results show that similar coefficient shifts may produce different update magnitudes and directions, and that the updates induced by the two scorings may even have a negative inner product. Distance in coefficient space therefore cannot characterize update safety by itself, since the direction and magnitude of the update also depend on the gradient geometry of the current policy.

Can Group-wise Screening Control Aggregate Reversal Risk? The analyses in Fig. 1 jointly indicate that when the optimizer retains only the response groups satisfying a local criterion, the sample distribution entering optimization changes from all response groups to the acceptance set. Since different screening rules yield different acceptance sets, a per-group criterion only shows that a local condition holds and cannot directly provide control over the aggregate direction-reversal risk on the acceptance set. Risk assessment must therefore be based on conditional statistics over the response groups that are actually accepted.

In summary, group normalization turns score discrepancies into correlated shifts in training coefficients, whose update effects depend on the current policy's response-level gradient geometry. Selective execution further changes the accepted groups and their aggregate direction-reversal risk. SetGRPO therefore jointly models coefficient uncertainty, policy-dependent update geometry, and acceptance-conditioned reversal risk.

4. Methodology

As the main innovation, SetGRPO constructs a joint coefficient set in the training-coefficient space actually used by GRPO and propagates it into the update space of the current policy, so as to recover effective and directionally consistent training updates as much as possible under grader score discrepancies. To address the propagation chain revealed by the pilot experiments in Section 3, SetGRPO consists of three interconnected modules: (a) *Joint Reward Set Calibration, JRSC*, (b) *Geometry-Guided Safe Recovery, GGSR*, and (c) *Selective Risk-Gated Optimization, SRGO*. In the following, we introduce these three modules in turn.

4.1. Joint Reward Set Calibration (JRSC)

As shown in Fig. 1(a), two grader scorings of the same response group produce correlated within-group shifts after the shared training mapping. We therefore propose *Joint Reward Set Calibration*, which models this coupled uncertainty as a joint coefficient set rather than independent response-wise errors. Specifically, for problem x, let $\mathbf{r}_x$ and $\mathbf{r}_x^\star$ denote the rewards from the first and second scorings. Using Eq. 1, $\hat{\mathbf{a}}_x = T_{\text{train}}(\mathbf{r}_x), \quad \mathbf{a}_x^\star = T_{\text{train}}(\mathbf{r}_x^\star)$, where T_{train} maps group rewards to within-group normalized coefficients. We define the paired residual $\boldsymbol{\delta}_x = \mathbf{a}_x^\star - \hat{\mathbf{a}}_x$, which calibrates the coverage radius of the subsequent joint set.

Subsequently, to capture how reward changes jointly perturb the coefficient vector, we linearize T_{train} around the first-scoring rewards $\mathbf{r}_x$ and define $\mathbf{J}_x = \nabla_{\mathbf{r}} T_{\text{train}}(\mathbf{r})|_{\mathbf{r}=\mathbf{r}_x}$. Its j-th column describes how a perturbation in reward j affects all group coefficients, preserving the cross-response coupling induced by shared normalization. The resulting coefficient-uncertainty matrix under unit reward perturbations is

$$\mathbf{S}_x^{\text{coef}} = \mathbf{J}_x \mathbf{J}_x^\top + \lambda \mathbf{I}_K, \qquad \lambda \geq 0, \tag{4}$$

where $\mathbf{J}_x \mathbf{J}_x^\top$ captures the cross-coefficient effects of local reward perturbations and determines set's directional scales, while $\lambda \mathbf{I}_K$ stabilizes inversion and distance computation in near-degenerate cases.

However, $\mathbf{S}_z^{\text{coef}}$ specifies only the shape of the joint coefficient set, not its radius. We therefore calibrate a unified coverage radius from independent paired residuals. For each coverage-calibration group $z \in \mathcal{D}_{\text{cov}}$ [1] , we compute the Mahalanobis score (Mahalanobis, 2018) $\rho_z = \sqrt{\boldsymbol{\delta}_z^\top (\mathbf{S}_z^{\text{coef}})^{-1} \boldsymbol{\delta}_z}$. This score accounts for directional scale differences and within-group correlations induced by reward normalization.

Next, let $n = |\mathcal{D}_{\text{cov}}|$, $\rho_{(j)}$ denote the j-th order statistic, and $\alpha \in (0, 1)$ be the target miss rate. Under exchangeability of the calibration and target groups (i.e., identical data construction, response sampling, and scoring protocols with permutation-invariant joint distributions), we define

$$Q_\alpha = \rho_{(\lceil (n+1)(1-\alpha) \rceil)}, \quad C_x = \left\{ \mathbf{a} \in \mathbb{R}^K : (\mathbf{a} - \hat{\mathbf{a}}_x)^\top (\mathbf{S}_x^{\text{coef}})^{-1} (\mathbf{a} - \hat{\mathbf{a}}_x) \leq Q_\alpha^2 \right\}. \tag{5}$$

Then C_x is an ellipsoidal joint coefficient set centered at $\hat{\mathbf{a}}_x$ that contains the training coefficients of the second scoring with marginal probability at least $1 - \alpha$. If $\lceil (n+1)(1-\alpha) \rceil = n + 1$, no corresponding $(n+1)$-th order statistic exists in the calibration set.

4.2. Geometry-Guided Safe Recovery (GGSR)

JRSC models the discrepancy between two grader scorings as the joint coefficient set C_x. We then propose *Geometry-Guided Safe Recovery (GGSR)*, which maps C_x into the policy-update and reward spaces to identify the largest coefficient-recovery ratio whose update remains directionally consistent with the second scoring under the current policy.

Specifically, we first let $q \in \{\text{full}, \text{reward}\}$ denote the policy-update and reward spaces, respectively, with dimension D_q. Let $\mathcal{S}^q$ be a fixed low-dimensional linear map, $\mathbf{P}_x^q$ a linear transformation, and $\mathbf{h}_x^q$ a coefficient-independent offset. Notably, by full geometry we mean the complete local update geometry under a fixed projected representation, rather than the unprojected original parameter space. From the response-level gradients in Eq. 3, we define

$$\mathbf{G}_x^q = \left[\mathcal{S}^q(\mathbf{g}_{x,1}), \ldots, \mathcal{S}^q(\mathbf{g}_{x,K}) \right] \in \mathbb{R}^{D_q \times K}, \quad \mathbf{u}_x^q(\mathbf{a}) = \mathbf{P}_x^q \left(\mathbf{G}_x^q \mathbf{a} + \mathbf{h}_x^q \right) \in \mathbb{R}^{D_q}, \tag{6}$$

[1] The paired groups are split by problem root into disjoint coverage- and risk-calibration sets. The former calibrates the set radius, while the latter selects the execution threshold; both are fixed before formal training.

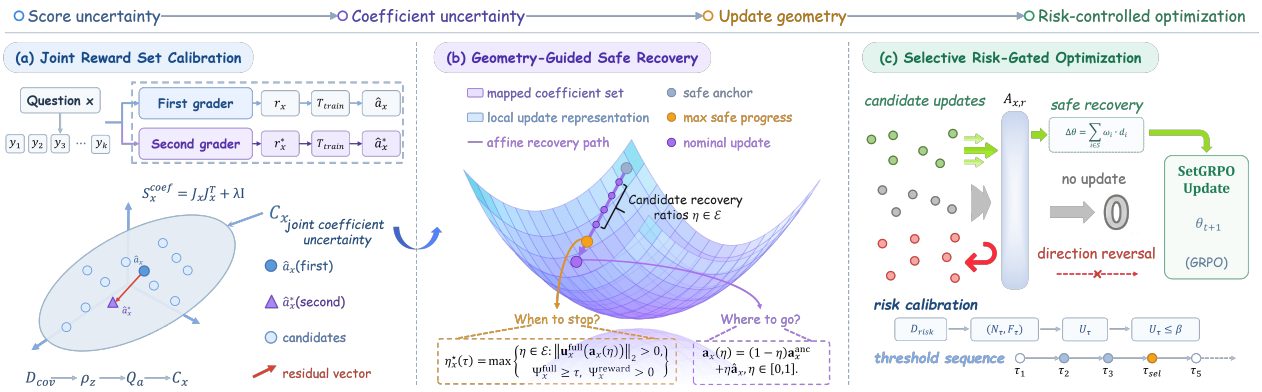

Figure 2. Overview of SetGRPO. *(a) Joint reward-set calibration* uses two grader scores on the same response group to construct a joint coefficient uncertainty set. *(b) Geometry-guided safe recovery* maps candidate coefficients through the current policy geometry and identifies the largest direction-preserving recovery ratio. *(c) Selective risk-gated optimization* selects a global threshold and applies the recovered coefficients only to updates satisfying the calibrated conditional-risk constraint.

where the i-th column of $\mathbf{G}_x^q$ is the projected gradient of response $y_{x,i}$, and $\mathbf{u}_x^q(\mathbf{a})$ is the local update induced by the coefficients $\mathbf{a}$ in the space q.

For a candidate coefficient $\mathbf{a}$, let $\mathbf{d} = \mathbf{u}_x^q(\mathbf{a})$ and define the worst-case projection over the joint coefficient set as $\Psi_x^q(\mathbf{d}) := \min_{\tilde{\mathbf{a}} \in C_x} \mathbf{d}^\top \mathbf{u}_x^q(\tilde{\mathbf{a}})$. This quantity measures the candidate update's worst-case alignment and magnitude under grader uncertainty. If $\mathbf{a}_x^\star \in C_x$, then $\Psi_x^{\text{full}}\left(\mathbf{u}_x^{\text{full}}(\mathbf{a})\right) \le \left\langle \mathbf{u}_x^{\text{full}}(\mathbf{a}), \mathbf{u}_x^{\text{full}}(\mathbf{a}_x^\star) \right\rangle$. Thus, a positive full-space projection rules out direction reversal, while a positive reward-space projection preserves reward-direction consistency. Because C_x is ellipsoidal and $\mathbf{u}_x^q$ is affine in $\mathbf{a}$, the projection has the closed form

$$\Psi_x^q(\mathbf{d}) = \mathbf{d}^\top \mathbf{u}_x^q(\hat{\mathbf{a}}_x) - Q_\alpha \sqrt{\mathbf{d}^\top \mathbf{P}_x^q \mathbf{G}_x^q \mathbf{S}_x^{\text{coef}} (\mathbf{G}_x^q)^\top (\mathbf{P}_x^q)^\top \mathbf{d}}. \tag{7}$$

However, the minimum-projection criterion evaluates a candidate coefficient but does not select one from C_x. We therefore choose the feasible coefficient with the smallest local-update norm as an anchor, denoted as $\mathbf{a}_x^{\text{anc}} \in \arg\min_{\mathbf{a} \in C_x} \frac{1}{2} \|\mathbf{u}_x^{\text{full}}(\mathbf{a})\|_2^2$. Since C_x is convex and centered at $\hat{\mathbf{a}}_x$, the path $\mathbf{a}_x(\eta) = (1-\eta)\mathbf{a}_x^{\text{anc}} + \eta \hat{\mathbf{a}}_x, \quad \eta \in [0,1]$, remains in C_x. Here, η denotes the fraction of training recovered from the conservative anchor toward the first-scoring update. For a predetermined candidate set $\mathcal{E} \subseteq [0,1]$, the maximum acceptable ratio under threshold $\tau \ge 0$ is

$$\eta_x^\star(\tau) = \max \left\{ \eta \in \mathcal{E} : \left\| \mathbf{u}_x^{\text{full}}(\mathbf{a}_x(\eta)) \right\|_2 > 0,\ \Psi_x^{\text{full}}\left(\mathbf{u}_x^{\text{full}}(\mathbf{a}_x(\eta))\right) \ge \tau,\ \Psi_x^{\text{reward}}\left(\mathbf{u}_x^{\text{reward}}(\mathbf{a}_x(\eta))\right) > 0 \right\}. \tag{8}$$

If no η satisfies these conditions, group x produces no update.

4.3. Selective Risk-Gated Optimization (SRGO)

Equation 8 determines the maximum feasible recovery ratio for each response group at threshold τ. Because τ also controls whether a nonzero recovery exists, it induces the acceptance set $\mathcal{A}_\tau = \{x : A_{x,\tau} = 1\}$, where $A_{x,\tau} = 1$ iff the constraints admit a nonzero recovery ratio, and $A_{x,\tau} = 0$ otherwise. The aggregate direction-reversal risk therefore depends on the composition of $\mathcal{A}_\tau$, not only on the group-wise criterion. We thus introduce *Selective Risk-Gated Optimization*, which evaluates acceptance-conditioned risk for candidate thresholds on the independent risk-calibration set $\mathcal{D}_{\text{risk}}$ and selects a threshold that satisfies the risk constraint while retaining as much training signal as possible.

Specifically, for each candidate threshold $\tau \in \mathcal{T}$, we define the direction-reversal event for accepted groups as $Z_{x,\tau} = \mathbb{I}\left[\left\langle \mathbf{u}_x^{\text{full}}\left(\mathbf{a}_x(\eta_x^\star(\tau))\right), \mathbf{u}_x^{\text{full}}(\mathbf{a}_x^\star)\right\rangle < 0\right]$, and its acceptance-conditioned risk as

Table 1. Main results on ten downstream benchmarks. HB: HealthBench; LEM: LLMEval-Med; MQA: MedQA; MCQA: MedMCQA; WB: WritingBench; CW: Creative Writing v3; IFE: IFEval; IFB: IFBench; IB: InfoBench; and MIF: MulDimIF. Our results report the mean across repeated runs, with absolute improvements over the corresponding Base model shown in parentheses.

Model	Medical				Writing		Instruction Following				Macro Avg.
	HB	LEM	MQA	MCQA	WB	CW	IFE	IFB	IB	MIF	
Closed-source Models											
OpenAI-o3	59.80	78.90	96.00	84.70	77.70	81.40	91.60	67.80	89.40	64.00	79.13
GPT-4.1	47.90	71.20	92.40	80.00	69.00	79.00	87.00	37.20	89.40	64.00	71.71
Gemini-3 Pro Preview	49.30	72.70	67.00	62.20	78.50	81.50	94.20	61.20	85.60	66.83	71.90
Open-sourced Models											
Qwen3.5-4B	37.09	26.24	77.77	63.26	50.20	27.60	46.40	17.35	78.22	41.25	46.54
+ GRPO	56.94 (+19.85)	43.29 (+17.05)	79.11 (+1.34)	63.55 (+0.29)	64.92 (+14.72)	29.98 (+2.38)	62.28 (+15.88)	25.09 (+7.74)	80.38 (+2.16)	48.55 (+7.30)	55.41 (+8.87)
+ RuscaRL	65.10 (+28.01)	45.30 (+19.06)	79.60 (+1.83)	63.30 (+0.04)	67.50 (+17.30)	31.80 (+4.20)	65.90 (+19.50)	24.40 (+7.05)	80.74 (+2.52)	47.80 (+6.55)	57.14 (+10.60)
+ RubricHub	67.18 (+30.09)	47.71 (+21.47)	78.06 (+0.29)	**64.22 (+0.96)**	69.62 (+19.42)	**32.80 (+5.20)**	67.88 (+21.48)	26.31 (+8.96)	**82.91 (+4.69)**	49.84 (+8.59)	58.65 (+12.11)
+ SetGRPO (Ours)	**69.64 (+32.55)**	**50.08 (+23.84)**	**80.55 (+2.78)**	63.84 (+0.58)	**72.14 (+21.94)**	32.23 (+4.63)	**70.45 (+24.05)**	**28.97 (+11.62)**	82.32 (+4.10)	**52.35 (+11.10)**	**60.26 (+13.72)**
Qwen3.5-9B	40.94	34.63	76.20	67.03	53.83	32.19	51.39	19.39	80.98	41.17	49.78
+ GRPO	61.06 (+20.12)	51.65 (+17.02)	77.68 (+1.48)	67.10 (+0.07)	68.84 (+15.01)	34.65 (+2.46)	67.61 (+16.22)	26.91 (+7.52)	83.02 (+2.04)	48.33 (+7.16)	58.69 (+8.91)
+ RuscaRL	69.00 (+28.06)	53.50 (+18.87)	78.35 (+2.15)	67.18 (+0.15)	71.70 (+17.87)	36.40 (+4.21)	71.30 (+19.91)	25.80 (+6.41)	82.86 (+1.88)	47.20 (+6.03)	60.33 (+10.55)
+ RubricHub	71.04 (+30.10)	55.62 (+20.99)	**79.86 (+3.66)**	67.28 (+0.25)	73.76 (+19.93)	34.21 (+2.02)	73.38 (+21.99)	28.24 (+8.85)	**85.42 (+4.44)**	**52.33 (+11.16)**	62.11 (+12.33)
+ SetGRPO (Ours)	**73.86 (+32.92)**	**58.15 (+23.52)**	79.32 (+3.12)	**67.34 (+0.31)**	**76.28 (+22.45)**	**36.91 (+4.72)**	**75.94 (+24.55)**	**30.57 (+11.18)**	84.83 (+3.85)	51.87 (+10.70)	**63.51 (+13.73)**
Llama 3.2-3B	14.37	15.96	57.18	48.74	30.63	32.57	71.08	23.67	44.93	35.69	37.48
+ GRPO	21.08 (+6.71)	17.42 (+1.46)	58.66 (+1.48)	49.32 (+0.58)	40.49 (+9.86)	35.41 (+2.84)	77.69 (+6.61)	25.58 (+1.91)	45.14 (+0.21)	76.42 (+40.73)	44.72 (+7.24)
+ RuscaRL	22.84 (+8.47)	17.96 (+2.00)	60.47 (+3.29)	51.05 (+2.31)	42.52 (+11.89)	40.28 (+7.71)	79.58 (+8.50)	27.76 (+4.09)	46.67 (+1.74)	79.35 (+43.66)	46.85 (+9.37)
+ RubricHub	24.61 (+10.24)	18.19 (+2.23)	61.34 (+4.16)	51.27 (+2.53)	47.33 (+16.70)	39.84 (+7.27)	80.93 (+9.85)	28.67 (+5.00)	47.44 (+2.51)	82.08 (+46.39)	48.17 (+10.69)
+ SetGRPO (Ours)	**26.72 (+12.35)**	**18.47 (+2.51)**	**62.18 (+5.00)**	**52.14 (+3.40)**	**45.66 (+15.03)**	**42.51 (+9.94)**	**83.47 (+12.39)**	**30.38 (+6.71)**	**49.46 (+4.53)**	**84.67 (+48.98)**	**49.57 (+12.09)**

$R_\tau = \Pr(Z_{x,\tau} = 1 \mid A_{x,\tau} = 1)$. On the independent risk-calibration set $\mathcal{D}_{\text{risk}}$, let N_τ and F_τ denote the numbers of accepted and direction-reversed groups, respectively. We estimate R_τ using the one-sided Clopper–Pearson upper bound $U_\tau = \text{CPUpper}\left(F_\tau, N_\tau; 1 - \frac{\gamma}{|\mathcal{T}|}\right)$, where β is the target risk level and γ is the family-wise failure probability across candidate thresholds. A threshold is feasible if $N_\tau > 0$ and $U_\tau \le \beta$; among feasible thresholds, we select the one maximizing cumulative recovery Y_τ, formualted as,

$$Y_\tau = \sum_{x \in \mathcal{D}_{\text{risk}}} A_{x,\tau} \eta_x^\star(\tau), \quad \tau_{\text{sel}} = \min\left(\arg\max_{\tau \in \mathcal{T}: N_\tau > 0,\, U_\tau \le \beta} Y_\tau\right). \tag{9}$$

If several thresholds tie, the smaller one is taken; if no threshold satisfies the risk constraint, τ_{sel} is undefined and the risk gating rule is disabled in subsequent training.

Once τ_{sel} is determined, SetGRPO retains the probability ratio, clipping, and token-level surrogate term of the original GRPO and replaces only the training coefficients with,

$$\mathcal{J}_{\text{SetGRPO}}(\boldsymbol{\theta}) = \mathcal{J}_{\text{GRPO}}\left(\boldsymbol{\theta}; \widetilde{\mathbf{a}}_x\right), \qquad \widetilde{\mathbf{a}}_x = \begin{cases} \mathbf{a}_x\left(\eta_x^\star(\tau_{\text{sel}})\right), & A_{x,\tau_{\text{sel}}} = 1, \\ \mathbf{0}_K, & A_{x,\tau_{\text{sel}}} = 0. \end{cases} \tag{10}$$

Accepted response groups use their recovered coefficients, while the remaining groups yield no policy update. We further evaluate the predetermined τ_{sel} across different training checkpoints to test its empirical stability under policy drift.

5. Experiments

5.1. Experimental Setup

Data and evaluation. We perform reinforcement learning training on the RubricHub training set (Li et al., 2026a), whose data covers five task types: Chat, Instruction Following, Medical, Science, and Writing. After training, each model is evaluated on ten downstream benchmarks spanning three domains: healthcare, writing, and instruction following. The healthcare benchmarks include HealthBench, LLMEval-Med, MedQA, and MedMCQA; the writing benchmarks include WritingBench and Creative Writing v3; and the instruction following benchmarks include IFEval, IFBench, InfoBench, and MulDimIF. For benchmarks that require model scoring, we uniformly invoke the GPT-5.6-Luna

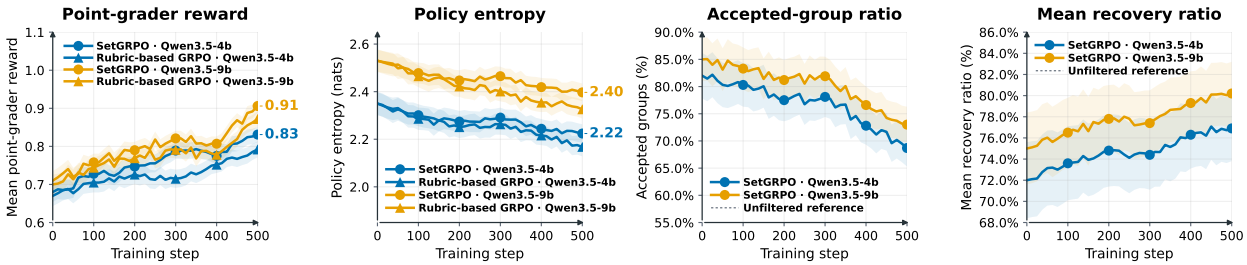

Figure 3. Training dynamics of SetGRPO, including the reward obtained from the first score, policy entropy, accepted-group ratio, and mean recovery ratio over the first 500 training steps.

API. Each benchmark uses its official evaluation metrics, and Macro Avg. is used only to aggregate the results of the ten benchmarks.

Models and baselines. We use Qwen3.5-4B, Qwen3.5-9B and Llama 3.2-3B as backbones and compare Base (the backbone without the training in this study) against four training methods: Rubric-based GRPO, RubricHub (Li et al., 2026a), RuscaRL (Zhou et al., 2026), and SetGRPO. Rubric-based GRPO uses the same rewards obtained from the first scoring as SetGRPO and is the main control for differences in the training objective; RubricHub and RuscaRL are run following their publicly available training procedures, with identical data scale, training budget, and final evaluation protocol.

Training details. For each prompt, we generate $K = 8$ rollout responses with a maximum response length of 4096 tokens. All methods use the same training data, number of training steps, and generation budget, and are evaluated under the same protocol after training. For benchmarks requiring model scoring, all methods use the same judge, prompt, and decoding configuration. The rubric rewards during training are provided by a Qwen3-32B grader. The online rewards of SetGRPO come from the first grader scoring; the second grader scoring is computed only for pre-frozen response groups, used for calibration and excluded from the training coefficients of online groups.

5.2. Main Results

Overall performance. Table 1 reports results on ten benchmarks across three backbones. SetGRPO achieves the highest Macro Avg. on Qwen3.5-4B and Qwen3.5-9B, reaching 60.26 and 63.51, respectively. On Qwen3.5-9B, it improves over RubricHub by 2.82 points on HealthBench and 2.52 on WritingBench, and over RuscaRL by 4.86 on HealthBench and 4.65 on LLMEval-Med. Relative to RubricHub, its average gains across healthcare, writing, and instruction-following tasks are 1.22, 2.61, and 0.96 points, respectively, with consistent improvements across backbones. These results indicate that SetGRPO improves performance across different backbones and task domains by jointly modeling the training-coefficient shifts caused by score discrepancies and screening for actionable updates in combination with the update geometry of the current policy.

Training dynamics. Figure 3 shows how the mean reward from the first scoring, policy entropy, proportion of accepted groups, and mean recovery ratio over accepted groups evolve during training. The mean reward from the first scoring rises steadily in the early phase and saturates later; policy entropy decreases gradually without an abrupt collapse. Compared with unfiltered training, the proportion of accepted groups is lower but remains nonzero in the later phase, and the mean recovery ratio over accepted groups stays within 76%–80%. This indicates that screening excludes response groups that fail the projection condition while nonzero training coefficients still enter the optimization. The steady rise of reward indicates that the retained updates bring training progress, and the stable decline of entropy indicates that the policy does not exhibit marked additional shrinkage.

5.3. Same-Backbone Comparison

In this section, we compare SetGRPO with RubricHub on Qwen3-4B and Qwen3-14B using the same backbone for each comparison. All other experimental settings follow those used for the Qwen3.5 and Llama3.2 experiments, including the data splits, rollout group size, decoding configuration, training steps, generation budget, grader configuration, and evaluation protocol. The backbone-specific implementation details are summarized separately in Tables 3 and 4. Consistent with the main text (detailed in Section 5.2), we report results on ten benchmarks and analyze them across the Medical, Writing, and Instruction Following domains.

On Qwen3-4B, SetGRPO achieves a Macro Avg. of 57.56, improving over RubricHub by 1.50 points. In the Medical domain, SetGRPO improves over RubricHub by 2.48 and 2.45 points on HealthBench and LLMEval-Med, respectively, while the differences on MedQA and MedMCQA are 0.47 and 0.46 points. Thus, the gain in the medical domain mainly comes from open-ended medical evaluation, while performance on multiple-choice evaluation remains close to that of RubricHub. SetGRPO represents the correlated training-coefficient shifts caused by shared normalization of within-group score discrepancies as a joint coefficient set. It then combines this set with the response-level gradients of the current policy to determine whether candidate coefficients induce policy updates consistent with the reference update, thereby reducing erroneous or reversed updates caused by score discrepancies.

Table 3. Training and rollout configuration for Qwen3-4B and Qwen3-14B.

Configuration	**Qwen3-4B**	**Qwen3-14B**
Rollout group size K	8	8
Max prompt length	2048	2048
Max response length	4096	4096
Batch size	64	64
Mini-batch size	32	32
Optimizer	AdamW	AdamW
Learning rate	10^{-6}	10^{-6}
Weight decay	0.01	0.01
KL coefficient	0.001	0.001
Rollout decoding (T / top-p / top-k)	1.0 / 1.0 / 20	1.0 / 1.0 / 20
Training steps	500	500
Checkpoint frequency	10	10

In the Writing domain, SetGRPO improves over RubricHub by 2.46 and 1.71 points on WritingBench and Creative Writing v3, respectively. For writing tasks that jointly impose requirements on content, style, format, and length, as well as open-ended tasks emphasizing creative quality, SetGRPO retains effective score signals related to writing quality by re-

Table 4. Frozen coverage and risk calibration records for Qwen3-4B and Qwen3-14B. ECR denotes empirical conditional risk.

Backbone	Cov.	Risk	Q_α	τ_{sel}	N	F	ECR	U
Qwen3-4B	27,229	27,229	1.32	0.25	19113	1047	0.0548	0.575
Qwen3-14B	27,229	27,229	1.45	0.25	19137	1058	0.0554	0.580

Table 2. Same-backbone comparison between RubricHub and SetGRPO on Qwen3-4B and Qwen3-14B. RubricHub results were reproduced by our implementation using its publicly available training procedure. HB: HealthBench; LEM: LLMEval-Med; MQA: MedQA; MCQA: MedMCQA; WB: WritingBench; CW: Creative Writing v3; IFE: IFEval; IFB: IFBench; IB: InfoBench; and MIF: MulDimIF.

Backbone	**Method**	**Medical**				**Writing**		**Instruction Following**				**Macro Avg.**
		HB	**LEM**	**MQA**	**MCQA**	**WB**	**CW**	**IFE**	**IFB**	**IB**	**MIF**	
Qwen3-4B	RubricHub	64.34	45.12	**78.38**	**61.88**	66.51	28.47	64.21	24.08	**80.21**	47.35	56.06
	SetGRPO (Ours)	**66.82**	**47.57**	77.91	61.42	**68.97**	**30.18**	**66.83**	**26.49**	79.66	**49.72**	**57.56**
Qwen3-14B	Rubrichub	68.44	53.36	**77.53**	**64.76**	70.42	32.71	69.74	25.61	**82.88**	46.79	59.12
	SetGRPO (Ours)	**70.91**	**54.88**	76.94	64.18	**72.83**	**34.62**	**72.21**	**27.94**	82.17	**49.18**	**60.59**

covering executable training progress that satisfies the projection conditions. It therefore improves performance on both evaluations. In the Instruction Following domain, SetGRPO improves over RubricHub by 2.62, 2.41, and 2.37 points on IFEval, IFBench, and MulDimIF, respectively, while the difference on InfoBench is 0.55 points. The gains on the first three evaluations indicate improved adherence to explicit, unseen, and multidimensional constraints. For instructions requiring multiple requirements to be satisfied simultaneously, the method first screens executable updates using the joint coefficient set and the projection conditions in the policy-update space, and then uses independent risk calibration to select a unified threshold. This reduces the damage that direction-reversed updates can cause to multi-constraint behavior.

On Qwen3-14B, SetGRPO improves the Macro Avg. over RubricHub by 1.47 points. After the joint coefficient set and execution threshold are calibrated separately for the two backbones, SetGRPO achieves consistent overall gains. This indicates that the improvement comes from its treatment of score discrepancies, policy updates, and risk screening rather than dependence on a particular backbone. Overall, SetGRPO jointly represents the training-coefficient shifts caused by within-group score discrepancies, uses the response-level gradients of the current policy to screen and recover executable training progress, and applies risk calibration to determine a unified execution threshold, thereby improving benchmark performance across backbones and task domains.

6. Analysis and Discussion

This section reports the conditional direction-reversal risk of accepted updates and analyzes the effect of different design variants and the additional computational overhead on the final results.

6.1. Risk Calibration and Retained Updates

Figure 4 shows the empirical direction-reversal risk and its one-sided Clopper–Pearson upper confidence bound across candidate thresholds, together with the target risk β and the selected threshold $\tau_{\text{sel}} = 0.25$. As the screening condition tightens, both the empirical risk and its upper confidence bound decrease overall. At the selected threshold, the upper confidence bound remains below $\beta = 0.08$.

6.2. Ablation Studies

To isolate the effects of joint coefficient modeling, policy-update geometry, coefficient recovery, and threshold selection, we evaluate three variant groups: (a) an isotropic coefficient set instead of the joint set; (b) variants removing the policy-space projection condition, coefficient recovery, or reward-space projection condition; and (c) a fixed threshold instead of risk-calibrated threshold selection. We also include a matched random-discard baseline that retains the same number of groups as SetGRPO to distinguish geometry-based screening from merely reducing updates.

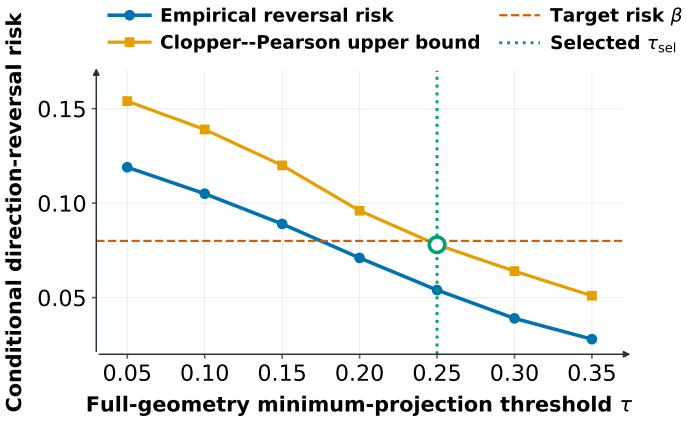

Figure 4. Risk calibration and retained updates at the selected threshold τ_{sel}.

Table 5 shows that all variants underperform the full setting (Macro Avg. 60.26). Replacing the joint coefficient set with a spherical set reduces the Macro Avg. to 56.72 and raises reversal risk to 0.096, indicating that this variant fails to preserve the within-group correlated coefficient shifts. Removing the projection condition of the policy-update space reduces the macro average to 55.18 and raises the risk to 0.118, indicating that screening in coefficient space alone is insufficient to determine the effect of coefficient changes on the current policy update. Keeping only the minimum-norm starting point gives a macro average of 57.04 with the recovery amount dropping to 0.36, indicating that not recovering the training coefficients costs training progress. Removing the minimum projection condition in the reward space gives a macro average of 57.88 with the risk rising to 0.103, indicating that the direction screening in the policy-update space cannot substitute for the positive projection condition in the reward space. The fixed threshold variant reaches 58.21 with a risk 0.042 higher than the full setting, indicating that risk calibration contributes to the selection of the execution threshold. Finally, matched random discard reaches only 54.93 despite a comparable acceptance rate, showing that the gains come from targeted screening rather than simply reducing updates.

Table 5. Ablation study of SetGRPO on Qwen3.5-4B. Δ denotes the change from Full SetGRPO in each metric's native units; upward arrows indicate the preferred direction. CR is normalized cumulative recovery and NUR is the nonzero update rate. Reversal risk, NUR, and CR lie in $[0, 1]$.

D#	Derivative	Macro Avg. ↑	Macro Avg. Δ	Reversal risk ↓	Reversal risk Δ	CR ↑	CR Δ	NUR ↑	NUR Δ
(a) Coefficient-set construction									
1	Spherical set	56.62	-3.64	0.096	+0.049	0.61	-0.19	0.512	-0.190
(b) Local geometry and recovery									
2	Identity geometry	55.08	-5.18	0.118	+0.071	0.51	-0.29	0.438	-0.264
3	Safe-anchor-only	56.94	-3.32	0.061	+0.014	0.36	-0.44	0.318	-0.384
4	Reward-only recovery	57.78	-2.48	0.103	+0.056	0.57	-0.23	0.566	-0.136
(c) Threshold selection and update control									
5	Fixed-threshold	58.11	-2.15	0.089	+0.042	0.60	-0.20	0.548	-0.154
6	Matched random discard	54.83	-5.43	0.154	+0.107	0.49	-0.31	0.702	0.000
SetGRPO (Ours)		**60.26**	**0.00**	**0.047**	**0.000**	**0.80**	**0.000**	**0.702**	**0.000**

6.3. Efficiency Comparison

In Table 6, we compare GRPO and SetGRPO in computational cost and benchmark performance. SetGRPO increases per-step training time from 82 to 103 s, peak VRAM from 61 to 67 GB, and cumulative GPU hours from 22.8 to 28.6, corresponding to increases of 25.6%, 9.8%, and 25.4%, while rollout and grader times remain unchanged at 48 s and 18 s. The overhead mainly comes from Jacobian computation for coefficient-set calibration and per-group search for the largest projection-safe recovery ratio. In return, the ten-benchmark Macro Avg. increases from 55.41 to 60.26 (+4.75), indicating improved performance at the cost of additional computation.

Table 6. Efficiency comparison between GRPO and SetGRPO.

Method	Training-time breakdown					Resource usage		10-Benchmark Macro Avg.
	Single-step time (s)	Rollout time (s)	Grader time (s)	Jacobian time (s)	Coefficient recovery time (s)	Peak VRAM (GB)	Cumulative GPU-hours	
GRPO	82	48	18	0	0	61	22.8	55.41
SetGRPO	103	48	18	13	7	67	28.6	60.26

6.4. Recovery and Selection Diagnostics

We report the main training curves in Figure 3. This section further reports the coefficient norm, full-space and reward-space minimum projections, and cumulative recovery to examine the scale of recovered coefficients, the projection conditions over the joint coefficient set, and the recovery retained during training.

1) Coefficient-Norm Dynamics. Figure 5 shows the norm of the recovered coefficients throughout training. The coefficient norm remains generally stable and changes moderately with the policy state,

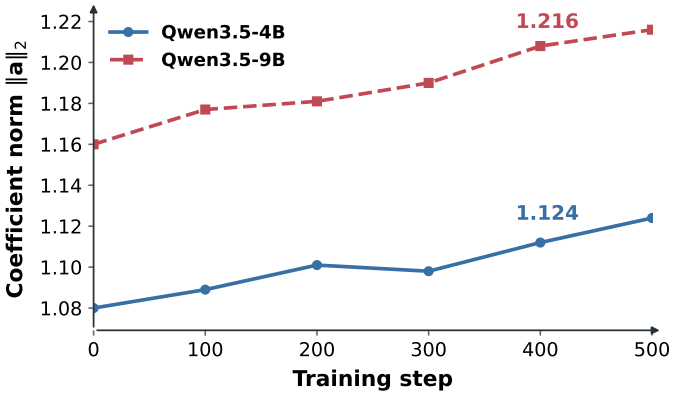

(a) Coefficient-norm trajectory.

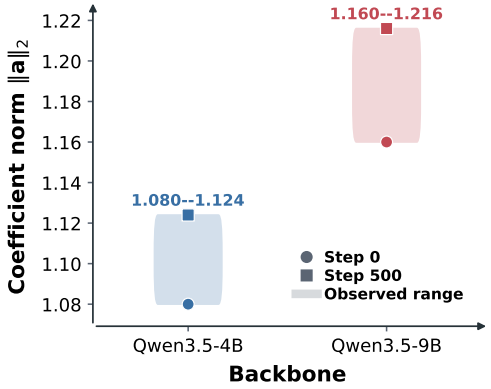

(b) Observed norm range.

Figure 5. Coefficient-norm dynamics during training. The left panel shows the coefficient-norm trajectories of the two backbones, while the right panel summarizes the observed norm range between training steps 0 and 500. The bounded ranges indicate that coefficient recovery does not lead to sustained numerical growth.

without a sustained upward trend. This result is consistent with the recovery range determined by the joint coefficient set and the recovery path from the minimum-update-norm anchor to the point-grader training coefficients. It indicates that the performance gain comes from screening and retaining executable recovery updates rather than continuously enlarging the training coefficients.

2) Minimum-Projection Diagnostics. Figure 6 shows the full-space and reward-space minimum projections. For accepted updates, the full-space minimum projection remains above the selected threshold τ_{sel}, while the reward-space minimum projection remains positive. Response groups that produce no recovery update fail at least one of these conditions. This result shows that screening jointly checks the policy-update projection under the joint coefficient set C_x and the reward-space projection, thereby keeping retained updates consistent with the local updates induced by the reference training coefficients $\mathbf{a}_x^\star$ while maintaining a positive projection in reward space.

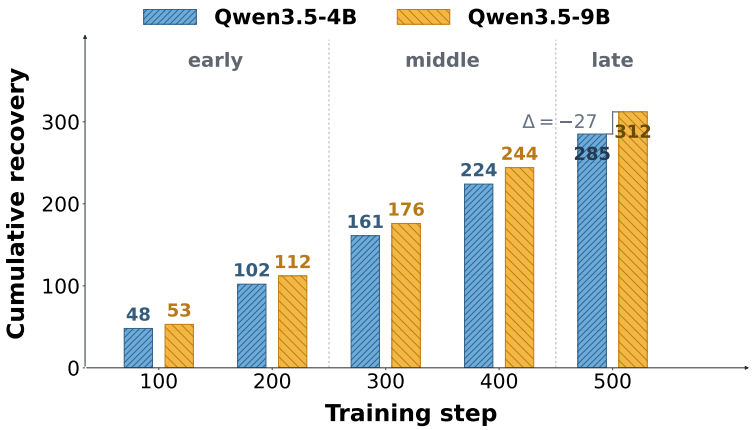

Figure 7. Cumulative recovery over the first 500 training steps for Qwen3.5-4B and Qwen3.5-9B. Shaded regions distinguish the early, middle, and late stages; the bracket reports the final cumulative-recovery gap.

3) Cumulative Recovery. Figure 7 shows cumulative recovery as a function of training steps. Cumulative recovery continues to increase, but its growth rate varies across training stages, reflecting differences in the number of response groups satisfying the recovery conditions and in their maximum recovery ratios. The curve shows that response groups continue to satisfy the recovery conditions and write back nonzero recovered coefficients throughout formal training. Thus, the retained training progress is not concentrated in only a few training steps.

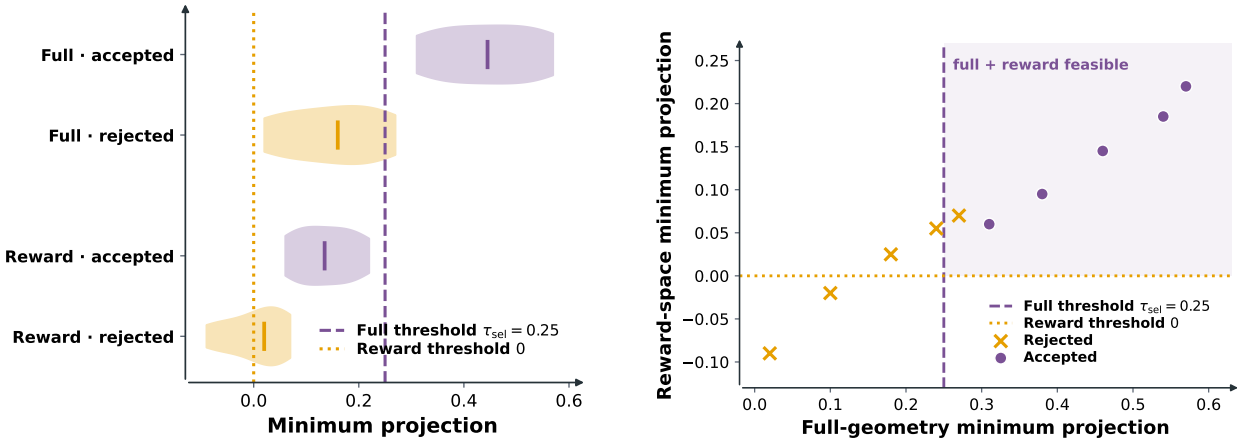

(a) Projection distributions.

(b) Joint projection constraints.

Figure 6. Projection diagnostics for accepted and rejected updates. The left panel shows the distributions of full-geometry and reward-space minimum projections. The right panel shows their joint constraints; the shaded upper-right region satisfies both the full-geometry threshold τ_{sel} and the positive reward-space projection condition.

6.5. Affine-Invariance Sanity Check

Training coefficients are determined by the centering and scale normalization of rewards within each group. To test this property, we apply a positive-affine transformation to frozen response groups and policy snapshots:

$$\mathbf{r}'_x = a\mathbf{r}_x + b\mathbf{1}_K, \qquad a > 0. \tag{11}$$

Because a positive-affine transformation does not change within-group standardized rewards, the training coefficients and joint coefficient sets before and after transformation should remain unchanged. Under the same response-level gradients, $\mathbf{G}^q_x$, and $\mathbf{u}^q_x(\mathbf{a})$, with the calibrated risk rule fixed, we compare the local updates, maximum recovery ratios, and response groups entering training under the two reward representations.

Figure 8 shows the coefficient residuals $(\mathbf{a}'_x - \mathbf{a}_x) \times 10^4$ under four positive-affine transformations. The point clouds are centered around zero, and the within-group means are close to zero, indicating that the standardized training coefficients differ only within numerical tolerance. Table 7 reports the corresponding directional and recovery metrics. The coefficient cosine and update cosine range from 0.9997 to 1.0000 and from 0.9995 to 1.0000, respectively. None of the four transformations produces a decision flip, and the largest difference in recovery ratio is only 0.8×10^{-3}. Thus, the training coefficients, local updates, and recovery ratios remain free of systematic changes under positive-affine transformations, indicating that the screening results are stable to reward translation and positive rescaling.

Table 7. Positive-affine invariance diagnostics.

Transform	Coeff. cosine	Update cosine	Flip (%)	$\lvert\Delta\eta^\star\rvert$ (10^{-3})
$a = 0.5,\ b = -1$	0.9998	0.9996	0.00	0.6
$a = 1,\ b = 0$	1.0000	1.0000	0.00	0.0
$a = 2,\ b = 1$	0.9999	0.9998	0.00	0.4
$a = 3,\ b = -2$	0.9997	0.9995	0.00	0.8

This invariance follows from the elimination of positive-affine transformations by within-group normalization. For $\mathbf{r}'_x = a\mathbf{r}_x + b\mathbf{1}_K$ with $a > 0$, the centered and scale-normalized rewards remain unchanged. Consequently, the training coefficients obtained from Eq. 1 remain unchanged, and the

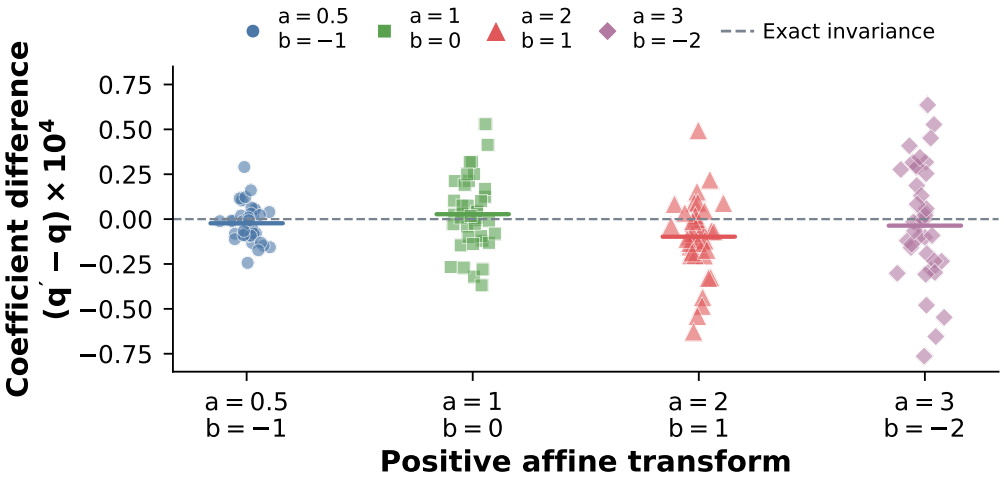

Figure 8. Positive-affine invariance under frozen-policy control. The left plot reports coefficient differences between the original and transformed standardized coefficients. The right table summarizes coefficient cosine, update cosine, decision flips, and the absolute recovery-ratio difference for each transform $\mathbf{r}'_x = a\mathbf{r}_x + b\mathbf{1}_K$ with $a > 0$.

joint coefficient set constructed from the paired coefficient residuals does not change materially. With the response-level gradients, $\mathbf{G}_x^q$, and $\mathbf{u}_x^q(\mathbf{a})$ of the current policy fixed, the updates and their minimum projections in the policy-update space and reward space also remain stable. Along the same recovery path, $\eta_x^\star$ changes only at the level of floating-point error.

This experiment shows that the training coefficients, local updates, minimum projections, and recovery ratios in SetGRPO are determined by the relative structure of within-group rewards. Translating or positively rescaling all rewards in a group does not systematically change these quantities. In contrast, transformations that do not preserve the within-group reward ordering, such as negative rescaling or response-specific non-affine perturbations, may alter the normalized rewards and consequently change the training coefficients, minimum projections, and recovery ratios.

6.6. Empirical Coverage and Risk Selection

Figure 9 reports the coverage performance of the joint coefficient set and the post-screening direction-reversal risk. For each predefined nominal coverage level, we compute Q_α on the coverage calibration groups and compare nominal and empirical coverage on independent risk calibration groups. Empirical coverage varies with the target level and remains close to it overall. Under the final setting, both nominal and empirical coverage are 0.95, indicating that the joint coefficient set constructed from coefficient residuals standardized by $\mathbf{S}_x^{\text{coef}}$ covers the reference training coefficients. This result is consistent with Eq. 5: when the calibration groups and the response groups to be covered satisfy exchangeability, Q_α provides the corresponding finite-sample marginal coverage guarantee.

Figure 9(b) compares the directional relationship between local updates induced by the training coefficients from the two scores for response groups entering and not entering training under τ_{sel}. According to Eq. 9, the direction-reversal event $Z_{x,\tau}$ and its conditional risk R_τ are defined only for response groups with nonzero recovery updates. For groups not entering training, we report only the comparison indicator between the local updates induced by the training coefficients from the two scores:

$$Z_x^{\text{nom}} = \mathbb{I}!\left[\left\langle \mathbf{u}_x^{\text{full}}(\hat{\mathbf{a}}_x), \mathbf{u}_x^{\text{full}}(\mathbf{a}_x^\star)\right\rangle < 0\right]. \tag{12}$$

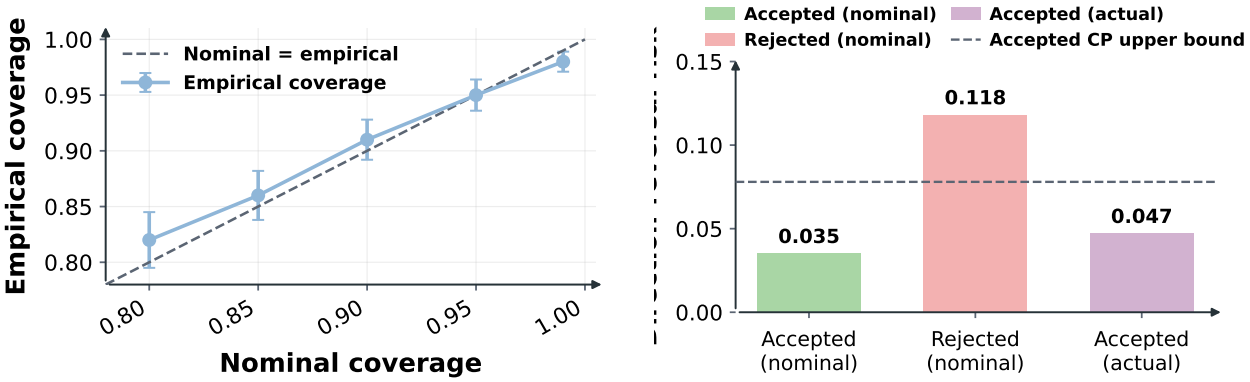

(a) Coverage of the joint coefficient set.

(b) Risk separation at the selected threshold.

Figure 9. Empirical coverage of the joint coefficient set and direction-reversal risk separation at the selected threshold.

In the figure, response groups entering training satisfy $\Pr(Z_x^{\mathrm{nom}} = 1 \mid A_{x,\tau_{\mathrm{sel}}} = 1) = 0.035$, whereas the corresponding proportion is $\Pr(Z_x^{\mathrm{nom}} = 1 \mid A_{x,\tau_{\mathrm{sel}}} = 0) = 0.118$ for groups not entering training. For groups that actually enter training, the empirical direction-reversal risk is 0.047, with a one-sided Clopper–Pearson upper confidence bound of $U_{\tau_{\mathrm{sel}}} = 0.078$, satisfying $U_{\tau_{\mathrm{sel}}} \le \beta$. Thus, the screened training groups exhibit a lower direction-reversal rate. This is because Eq. 8 retains only groups with a nonzero feasible recovery ratio that satisfy $\Psi_x^{\mathrm{full}} \ge \tau_{\mathrm{sel}}$ and $\Psi_x^{\mathrm{reward}} > 0$. The first condition requires the candidate update to maintain sufficient policy-update projection under the joint coefficient set, while the second requires a positive projection in reward space. Therefore, risk-calibrated screening changes the composition of groups entering training and keeps the acceptance-conditioned direction-reversal risk within the target range.

6.7. Checkpoint-wise Validation under Policy Drift

Figure 10 presents checkpoint validation of the fixed calibration results along the training trajectory. During calibration, we determine Q_α and τ_{sel}, and then rerun SRGO on the same response groups at global steps $t \in \{0, 100, 250, 500\}$. At each checkpoint, the projected local updates, Ψ_x^{full}, Ψ_x^{reward}, and recovery ratios are computed solely from the response-level gradients of the current policy. The values of Q_α and τ_{sel} remain unchanged across checkpoints.

As training proceeds, the local update geometry of the current policy changes, but the screening behavior does not exhibit abrupt changes. In Figure 10(a), the accepted-group proportion decreases slightly from the initial to the intermediate checkpoint and then stabilizes. In Figure 10(b), the mean recovery ratio of accepted groups remains positive, indicating that policy changes do not eliminate feasible recovery paths. Meanwhile, the empirical conditional direction-reversal rate of accepted groups remains low, and its one-sided upper confidence bound shows no substantial increase at any checkpoint. Rejected groups consistently exhibit higher direction-reversal rates. These results indicate that policy drift mainly changes the positions of response groups relative to the joint coefficient set, without disrupting the screening relation defined by $\Psi_x^{\mathrm{full}} \ge \tau_{\mathrm{sel}}$ and $\Psi_x^{\mathrm{reward}} > 0$.

This result is consistent with the joint coefficient set and risk-selection threshold fixed during calibration. During formal training, Q_α and τ_{sel} remain fixed, while each checkpoint recomputes Ψ_x^{full}, Ψ_x^{reward}, and the recovery ratios using only the local gradients at that checkpoint. Thus, differences across checkpoints reflect the effect of changing local update geometry on recovery ratios rather than re-estimation of the calibration quantities. Response groups satisfying $\Psi_x^{\mathrm{full}} \ge \tau_{\mathrm{sel}}$ and $\Psi_x^{\mathrm{reward}} > 0$

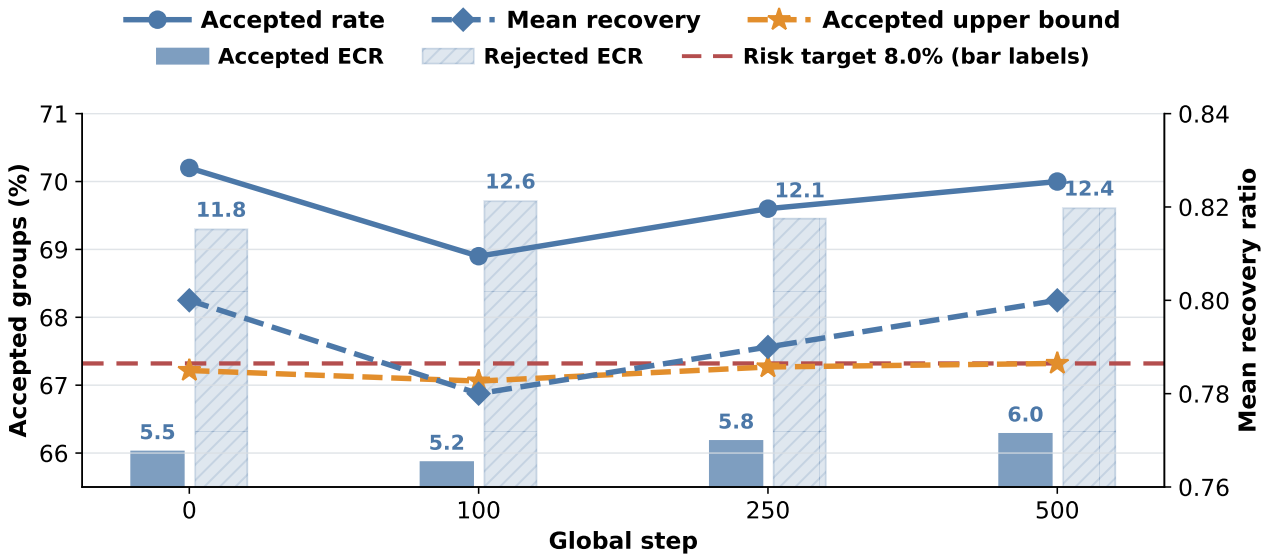

Figure 10. Checkpoint-wise validation of the fixed SRGO rule under policy drift on Qwen3.5-4B. Curves show the accepted-group rate, mean recovery ratio, and accepted-group upper confidence bound; bars show empirical direction-reversal rates for accepted and rejected groups. The red dashed line denotes $\beta = 0.08$.

retain nonzero recovered coefficients, whereas the remaining groups stop recovery and write back zero coefficients. The results show that policy drift changes projection values and recovery ratios, while acceptance and rejection continue to be determined by the same projection conditions. Accepted groups retain lower direction-reversal risk and positive recovery ratios, indicating that the screening rule continues to distinguish recoverable updates from updates that fail the conditions throughout the training trajectory.

7. Representative Cases

Figures 11–13 present three types of response groups: accepted, rejected, and recovery-infeasible, with three representative cases for each type. Each case reports the two scores, the corresponding training coefficients, the minimum projections in the policy-update space and reward space, the recovery ratio $\eta_x^\star$, and the final decision rationale. For the accepted cases, consider WritingBench #4 as an example. Its two minimum projections are 0.84 and 0.81, satisfying $\Psi_x^{\text{full}} \geq \tau_{\text{sel}}$ and $\Psi_x^{\text{reward}} > 0$, with a maximum feasible recovery ratio of $\eta_x^\star = 0.82$. The rejected cases all have nonzero candidate updates but fail at least one projection condition. For example, HealthBench 920ca4e9 has $M^{\text{full}} = 0.29$, satisfying the full-space condition, but $M^{\text{reward}} = -0.06$, which fails the positive reward-space projection condition. Its recovery ratio is therefore $\eta^\star = 0.00$. The recovery-infeasible cases have no nonzero feasible recovery ratio satisfying Eq. 8. For example, IFEval #2035 has $M^{\text{full}} = 0.11$ and $M^{\text{reward}} = -0.03$, neither of which satisfies the corresponding condition, and is therefore marked as infeasible. The two scores do not vary monotonically with the decision outcome across the three types of cases. In contrast, the minimum projections and maximum feasible recovery ratio directly correspond to the screening conditions in Eq. 8.

The decision differences across cases arise because the same score discrepancy induces different local updates under different response-level gradients. SetGRPO represents the joint uncertainty of the resulting training coefficients by C_x, then combines it with the response-level gradients of the current policy to evaluate the minimum projections of candidate updates in the policy-update space

and reward space and determine the maximum recovery ratio that can be written back for each response group. Thus, response groups satisfying the projection conditions retain nonzero training updates, whereas groups that fail these conditions do not write back an update.

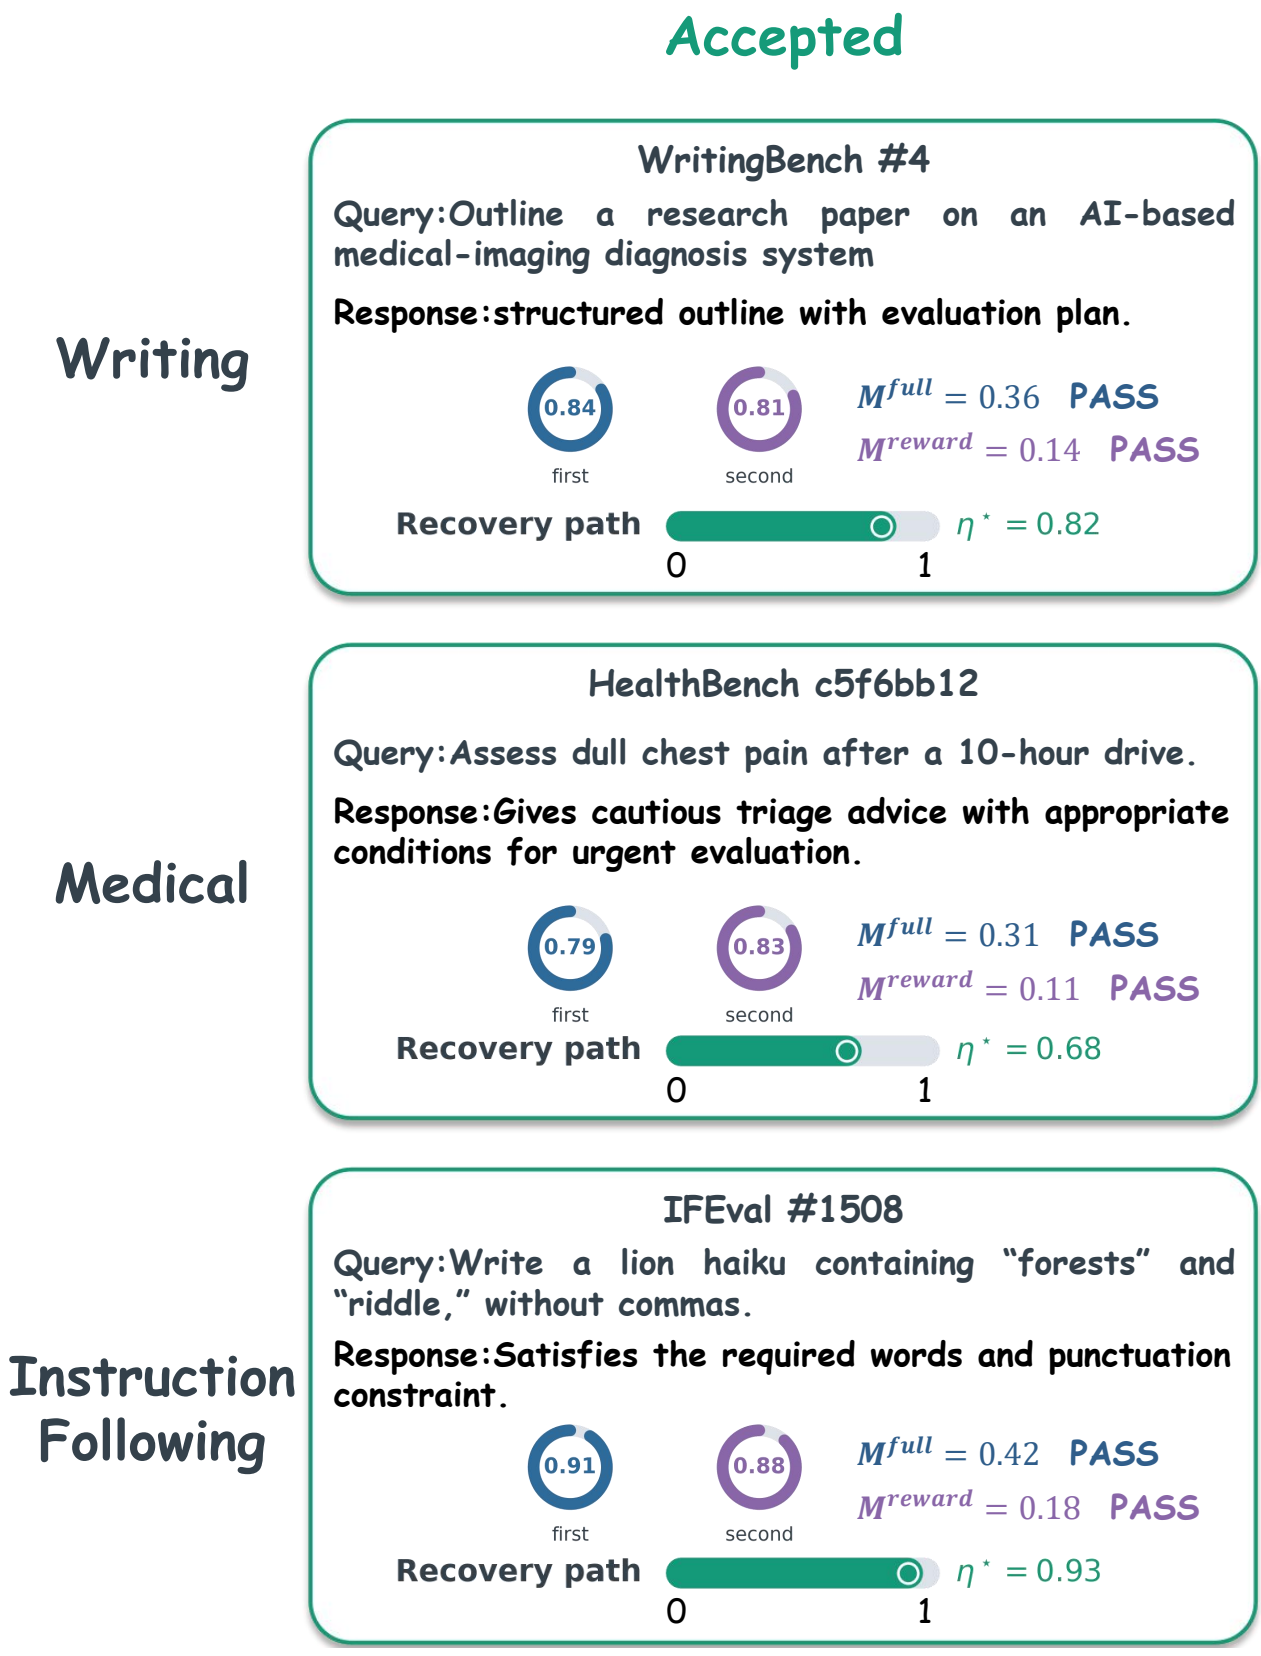

Figure 11. Representative accepted cases. Each case shows the anonymized query–response record, point and reference scores, full-geometry and reward-space projection checks, the recovery path from $\eta = 0$ to $\eta = 1$, the recovered training coefficient $\hat{\mathbf{a}}_x$, and the final decision.

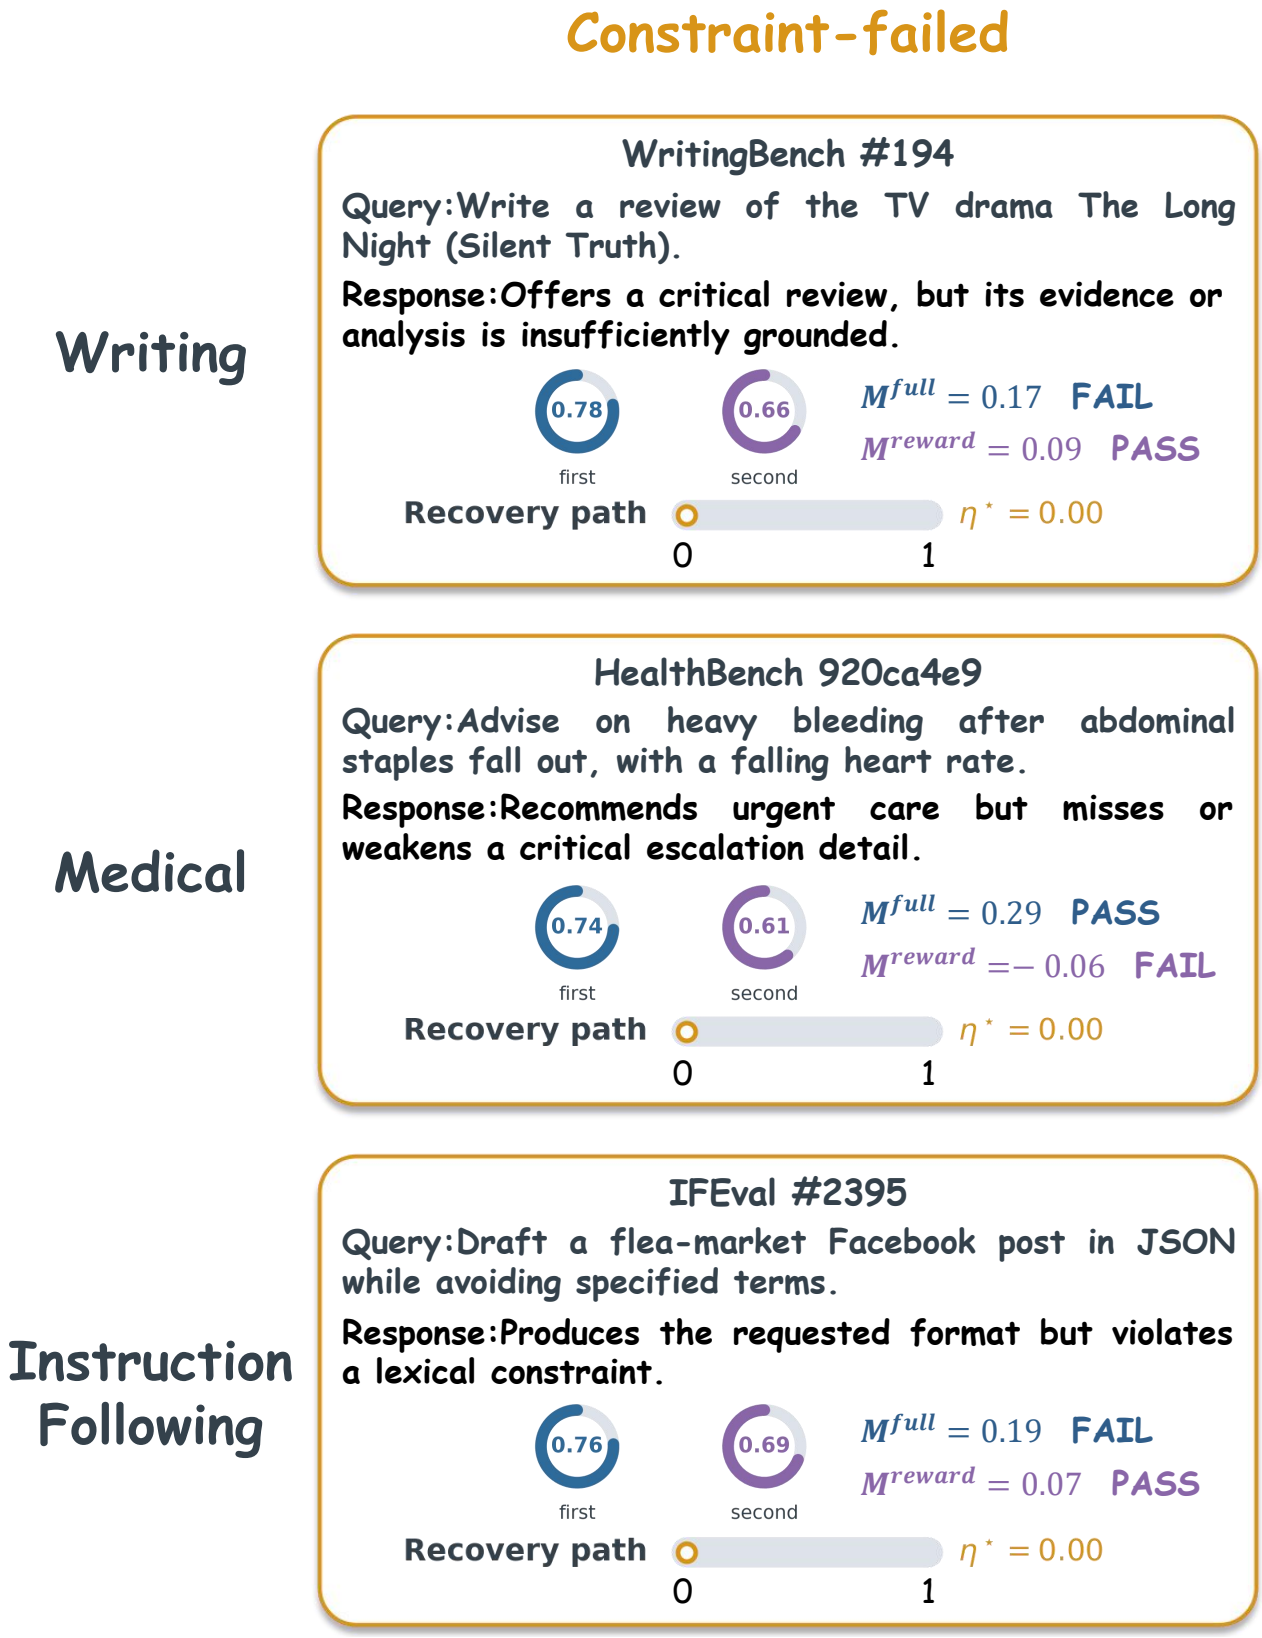

Figure 12. Representative rejected cases. Each case shows the anonymized query–response record, point and reference scores, full-geometry and reward-space projection checks, the recovery path from $\eta = 0$ to $\eta = 1$, the recovered training coefficient $\hat{\mathbf{a}}_x$, and the final decision.

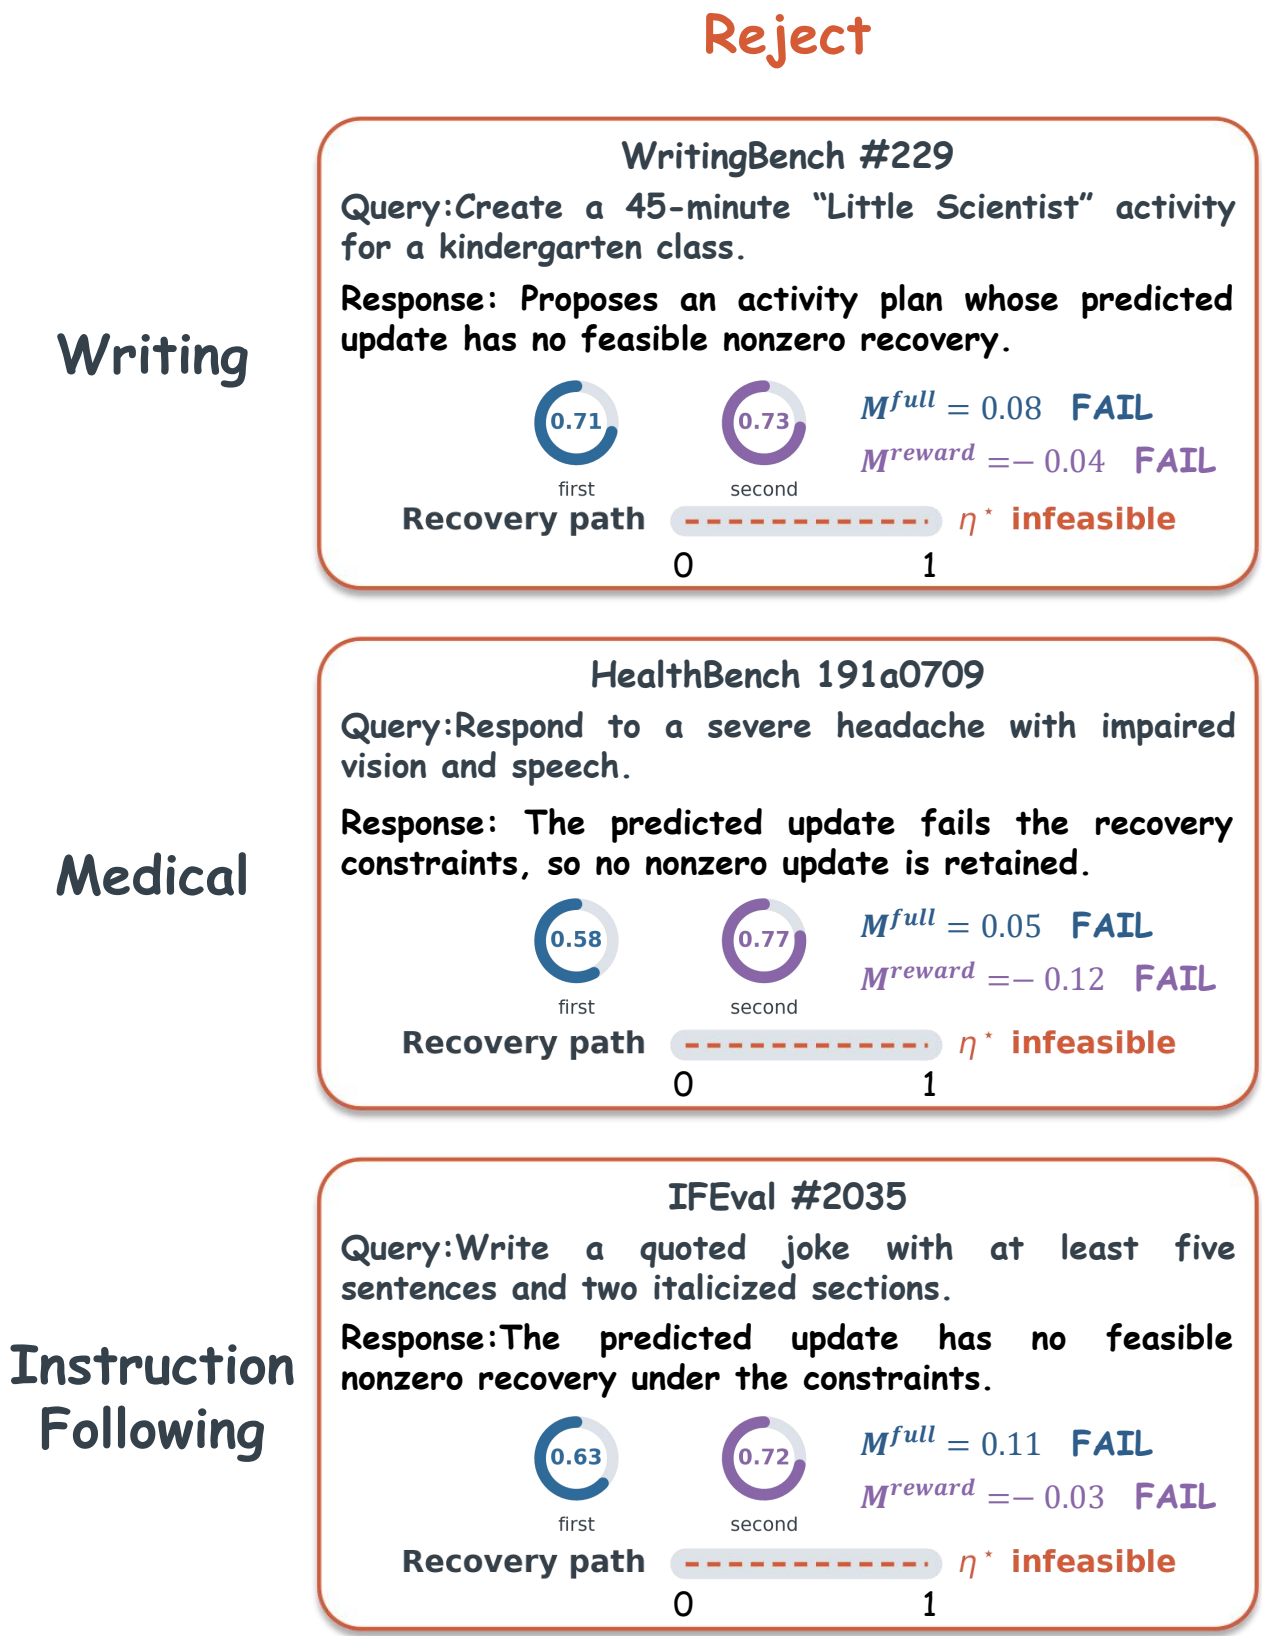

Figure 13. Representative recovery-infeasible cases. Each case shows the anonymized query–response record, point and reference scores, full-geometry and reward-space projection checks, the recovery path from $\eta = 0$ to $\eta = 1$, the recovered training coefficient $\hat{\mathbf{a}}_x$, and the final decision.

8. Conclusion

This paper proposed SetGRPO to address the effect of grader stochasticity on policy updates in rubric-based GRPO. Our pilot experiments showed that within-group normalization in GRPO converted grader stochasticity into correlated shifts in the training coefficients, while the response-level gradients of the current policy determined how these shifts affected local updates. Selective execution further changed the composition of the accepted response groups and their aggregate direction-reversal risk. Based on this observation, SetGRPO constructed a joint uncertainty set in training-coefficient space, calibrated possible coefficient shifts through JRSC, used GGSR to recover acceptable training signals in conjunction with the update geometry of the current policy, and finally calibrated and controlled the conditional direction-reversal risk of accepted updates through SRGO, while preserving the optimization structure of the original GRPO. Across ten open-ended task benchmarks in three domains, SetGRPO achieved the best macro-average performance among all compared training methods on the Qwen-3.5 and Llama-3.2 model families.

Ethics Statement SetGRPO uses paired rubric-grader scores and response-level gradients from the current policy to select and recover reinforcement-learning updates. These components may introduce several risks: grader judgments may be stochastic, biased, or subject to distribution shift; calibration data may contain sensitive information; and the coefficient calibration and selective update mechanisms may amplify artifacts present in the data or grader. The finite-sample guarantee in this work concerns only the conditional direction-reversal risk of accepted updates under the stated exchangeability assumption. It does not guarantee privacy, fairness, factuality, or harmlessness for individual responses. The training, coverage-calibration, and risk-calibration splits are separated by problem root, and downstream evaluation benchmarks are not used for training, calibration, or checkpoint selection. Any deployment should apply data minimization, appropriate authorization and access controls, and anonymization where applicable. It should also include regular audits for grader drift and subgroup disparities, as well as human review for high-risk applications. Updates should be suspended or reviewed when the exchangeability assumption, calibration conditions, or grader reliability are not satisfied.

AI Assistance Statement During the preparation of this manuscript, the authors used generative AI tools only to polish English writing, check grammar, and improve exposition. Generative AI tools were not used to design the method, derive the theory, conduct the experiments, or analyze the results. All AI-assisted text was reviewed by the authors, who take responsibility for the final content of this work.

References

Sunzhu Li, Jiale Zhao, Huimin Ren, Zhenlin Wei, Yang Zhou, Jingwen Yang, Shunyu Liu, Kaike Zhang, and Chen Wei. Rubrichub: A comprehensive and highly discriminative rubric dataset via automated coarse-to-fine generation. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 31320–31344, 2026a.

Yang Zhou, Sunzhu Li, Shunyu Liu, Wenkai Fang, Kongcheng Zhang, Jiale Zhao, Jingwen Yang, Yihe Zhou, Jianwei Lv, Tongya Zheng, Hengtong Lu, Chen Wei, Xie Yan, and Mingli Song. Breaking the exploration bottleneck: Rubric-scaffolded reinforcement learning for open-ended LLM reasoning. In *Forty-third International Conference on Machine Learning*, 2026.

Seungone Kim, Jay Shin, Joel Jang, Shayne Longpre, Hwaran Lee, Sangdoo Yun, Ryan Shin, Sungdong Kim, James Thorne, Minjoon Seo, et al. Prometheus: Inducing fine-grained evaluation capability in language models. In *International Conference on Learning Representations*, volume 2024, pages 29927–29962, 2024.

Seonghyeon Ye, Doyoung Kim, Sungdong Kim, Hyeonbin Hwang, Seungone Kim, Yongrae Jo, James Thorne, Juho Kim, and Minjoon Seo. Flask: Fine-grained language model evaluation based on alignment skill sets. In *International Conference on Learning Representations*, volume 2024, pages 55361–55414, 2024.

Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. G-eval: Nlg evaluation using gpt-4 with better human alignment. In *Proceedings of the 2023 conference on empirical methods in natural language processing*, pages 2511–2522, 2023.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, et al. Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in neural information processing systems*, 36:46595–46623, 2023.

Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022.

Nisan Stiennon, Long Ouyang, Jeffrey Wu, Daniel Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford, Dario Amodei, and Paul F Christiano. Learning to summarize with human feedback. *Advances in neural information processing systems*, 33:3008–3021, 2020.

Wenbing Li, Hang Zhou, Junqing Yu, Zikai Song, and Wei Yang. Coupled mamba: Enhanced multimodal fusion with coupled state space model. *Advances in Neural Information Processing Systems*, 37:59808–59832, 2024.

Yangliu Hu, Zikai Song, Na Feng, Yawei Luo, Junqing Yu, Yi-Ping Phoebe Chen, and Wei Yang. Sf2t: Self-supervised fragment finetuning of video-llms for fine-grained understanding. *arXiv preprint arXiv:2504.07745*, 2025.

Wenbing Li, Zikai Song, Hang Zhou, Junqing Yu, Yunyao Zhang, and Wei Yang. Lora-mixer: Coordinate modular lora experts through serial attention routing. In *International Conference on Learning Representations*, volume 2026, pages 14694–14716, 2026b.

Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, YK Li, Yang Wu, et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.

Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Peiyi Wang, Qihao Zhu, Runxin Xu, Ruoyu Zhang, Shirong Ma, Xiao Bi, et al. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.

Leo Gao, John Schulman, and Jacob Hilton. Scaling laws for reward model overoptimization. In *International conference on machine learning*, pages 10835–10866. PMLR, 2023.

Nathan Lambert, Valentina Pyatkin, Jacob Morrison, LJ Miranda, Bill Yuchen Lin, Khyathi Chandu, Nouha Dziri, Sachin Kumar, Tom Zick, Yejin Choi, et al. Rewardbench: Evaluating reward models for language modeling. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 1755–1797, 2025.

Peiyi Wang, Lei Li, Liang Chen, Zefan Cai, Dawei Zhu, Binghuai Lin, Yunbo Cao, Lingpeng Kong, Qi Liu, Tianyu Liu, et al. Large language models are not fair evaluators. In *Proceedings of the 62nd annual meeting of the association for computational linguistics (volume 1: Long papers)*, pages 9440–9450, 2024.

Jinhe Bi, Chennan Zhou, Zengjie Jin, Aniri, Shuo Lu, Wenke Huang, Hu Cao, Xun Xiao, Zhihong Zhu, Volker Tresp, Fei Shen, Yunpu Ma, and Tat-Seng Chua. Reflectrl: Learning from golden negative trajectories via reflective-to-direct reasoning, 2026a. URL https://arxiv.org/abs/2608.03972.

Jinhe Bi, Aniri, Minglai Yang, Xingcheng Zhou, Wenke Huang, Sikuan Yan, Yujun Wang, Zixuan Cao, Michael Färber, Xun Xiao, Volker Tresp, and Yunpu Ma. EchoRL: Reinforcement learning via rollout echoing. In *Forty-third International Conference on Machine Learning*, 2026b. URL https://openreview.net/forum?id=A6az59SGtF.

Aniri, Chen Yilin, Jinhe Bi, Junfei Guo, Donglai Ran, Xu Bian, Zengjie Jin, Yujun Wang, Yijun Tian, Volker Tresp, Fei Shen, Tat-Seng Chua, and Yunpu Ma. Fabrimae i trust myself? self-evaluating vla action generation with markov attention entropy, 2026a. URL https://arxiv.org/abs/2608.16697.

Jinhe Bi, Yifan Wang, Danqi Yan, Xun Xiao, Artur Hecker, Volker Tresp, and Yunpu Ma. Prism: Self-pruning intrinsic selection method for training-free multimodal data selection. *ArXiv*, abs/2502.12119, 2025a. URL https://api.semanticscholar.org/CorpusID:276421326.

Jinhe Bi, Yujun Wang, Haokun Chen, Xun Xiao, Artur Hecker, Volker Tresp, and Yunpu Ma. LLaVA steering: Visual instruction tuning with 500x fewer parameters through modality linear representation-steering. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar, editors, *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15230–15250, Vienna, Austria, July 2025b. Association for Computational Linguistics. ISBN 979-8-89176-251-0. URL https://aclanthology.org/2025.acl-long.739/.

Xinglang Zhang, Yuanmeng Xiang, Yunyao Zhang, Zeliang Chen, Junqing Yu, and Zikai Song. Too good to be real? diagnosing and reducing the gap between ai preference and real user engagement. *arXiv preprint arXiv:2609.18282*, 2026a.

Yunyao Zhang, Xinglang Zhang, Junxi Sheng, Wenbing Li, Junqing Yu, Yi-Ping Phoebe Chen, Wei Yang, and Zikai Song. Semantic-aware logical reasoning via a semiotic framework, 2026b. URL https://arxiv.org/abs/2509.24765.

Xinglang Zhang, Yunyao Zhang, ZeLiang Chen, Junqing Yu, Wei Yang, and Zikai Song. Logical phase transitions: Understanding collapse in llm logical reasoning, 2026c. URL https://arxiv.org/abs/2601.02902.

Paul F Christiano, Jan Leike, Tom Brown, Miljan Martic, Shane Legg, and Dario Amodei. Deep reinforcement learning from human preferences. *Advances in neural information processing systems*, 30, 2017.

Daniel M Ziegler, Nisan Stiennon, Jeffrey Wu, Tom B Brown, Alec Radford, Dario Amodei, Paul Christiano, and Geoffrey Irving. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593*, 2019.

Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, et al. Constitutional ai: Harmlessness from ai feedback. *arXiv preprint arXiv:2212.08073*, 2022.

John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.

Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. *Advances in neural information processing systems*, 36:53728–53741, 2023.

Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.

Jonathan Uesato, Nate Kushman, Ramana Kumar, Francis Song, Noah Siegel, Lisa Wang, Antonia Creswell, Geoffrey Irving, and Irina Higgins. Solving math word problems with process-and outcome-based feedback. *arXiv preprint arXiv:2211.14275*, 2022.

Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let's verify step by step. In *International Conference on Learning Representations*, volume 2024, pages 39578–39601, 2024.

Yunyao Zhang, Yihao Ai, Zuocheng Ying, Qirui Mi, Junqing Yu, Wei Yang, and Zikai Song. Coupling macro dynamics and micro states for long-horizon social simulation, 2026d. URL https://arxiv.org/abs/2604.05516.

Yunyao Zhang, Zuocheng Ying, Xinglang Zhang, Junqing Yu, Peng Fang, Xu Chen, Wei Yang, and Zikai Song. Intervensim: Intervention-aware social network simulation for opinion dynamics, 2026e. URL https://arxiv.org/abs/2604.06600.

Zikai Song, Xiajie Li, Yunyao Zhang, Xinglang Zhang, Wei Yang, and Junqing Yu. Social intelligence modeling: A comprehensive survey from social perception to social simulation. ResearchGate preprint, 2026. URL https://www.researchgate.net/publication/411123257_Social_Intelligence_Modeling_A_Comprehensive_Survey_from_Social_Perception_to_Social_Simulation. Preprint available on ResearchGate.

Qinlei Huang, Zhiwei Chen, Zixu Li, Chunxiao Wang, Xuemeng Song, Yupeng Hu, and Liqiang Nie. Median: Adaptive intermediate-grained aggregation network for composed image retrieval. In *ICASSP 2025-2025 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 1–5. IEEE, 2025.

Zixu Li, Zhiwei Chen, Zhiheng Fu, Wenbo Wang, Yupeng Hu, Weili Guan, and Liqiang Nie. Omniego-r^2: A routed reasoning framework for the 1st cross-domain egocross challenge at cvpr 2026. *arXiv preprint arXiv:2605.24481*, 2026c.

Jiale Huang, Zixu Li, Zhiwei Chen, Zhiheng Fu, Chunxiao Wang, and Yupeng Hu. Imagine: Adaptive schema-imagery enhanced composition for composed video retrieval. In *Proceedings of the 2026 International Conference on Multimedia Retrieval*, pages 288–297, 2026a.

Zhiwei Chen, Yupeng Hu, Zixu Li, Zhiheng Fu, Xuemeng Song, and Liqiang Nie. Offset: Segmentation-based focus shift revision for composed image retrieval. In *ACM MM*, page 6113–6122, 2025.

Zixu Li, Yupeng Hu, Zhiheng Fu, Zhiwei Chen, Yongqi Li, and Liqiang Nie. Tema: Anchor the image, follow the text for multi-modification composed image retrieval. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 24421–24442, 2026d.

Jinhe Bi, Danqi Yan, Yifan Wang, Wenke Huang, Haokun Chen, Guancheng Wan, Mang Ye, Xun Xiao, Hinrich Schuetze, Volker Tresp, and Yunpu Ma. The geometry of reasoning: Self-evaluation via layerwise trajectory evolution. In *Forty-third International Conference on Machine Learning*, 2026c. URL https://openreview.net/forum?id=WQyrwQwzmK.

Aniri, Jinhe Bi, Peng Liao, Zengjie Jin, Volker Tresp, Fei Shen, Yunpu Ma, and Tat-Seng Chua. Opd-v: Visual on-policy self-distillation with modality balance, 2026b. URL https://arxiv.org/abs/2608.05131.

Zhiwei Chen, Yupeng Hu, Zhiheng Fu, Zixu Li, Jiale Huang, Qinlei Huang, and Yinwei Wei. Intent: Invariance and discrimination-aware noise mitigation for robust composed image retrieval. In *AAAI*, volume 40, pages 20463–20471, 2026a.

Jiale Huang, Zixu Li, Zhiheng Fu, Zhiwei Chen, Qinlei Huang, and Yupeng Hu. Rankvr: Low-rank structure perception and value recalibration for robust composed image retrieval. In *Proceedings of the 2026 International Conference on Multimedia Retrieval*, pages 269–278, 2026b.

Zhiheng Fu, Yupeng Hu, Qianyun Yang, Shiqi Zhang, Zhiwei Chen, and Zixu Li. Air-know: Arbiter-calibrated knowledge-internalizing robust network for composed image retrieval. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 2658–2670, 2026.

Zixu Li, Yupeng Hu, Zhiwei Chen, Mingyu Zhang, Zhiheng Fu, and Liqiang Nie. Conesep: Cone-based robust noise-unlearning compositional network for composed image retrieval. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 16897–16909, 2026e.

Glenn Shafer and Vladimir Vovk. A tutorial on conformal prediction. *Journal of machine learning research*, 9(3), 2008.

Anastasios N Angelopoulos and Stephen Bates. Conformal prediction: A gentle introduction. *Foundations and Trends in Machine Learning*, 16(4):494–591, 2023.

Jing Lei, Max G'Sell, Alessandro Rinaldo, Ryan J Tibshirani, and Larry Wasserman. Distribution-free predictive inference for regression. *Journal of the American Statistical Association*, 113(523): 1094–1111, 2018.

Yaniv Romano, Evan Patterson, and Emmanuel Candes. Conformalized quantile regression. *Advances in neural information processing systems*, 32, 2019.

Yaniv Romano, Matteo Sesia, and Emmanuel Candes. Classification with valid and adaptive coverage. *Advances in neural information processing systems*, 33:3581–3591, 2020.

Stephen Bates, Anastasios Angelopoulos, Lihua Lei, Jitendra Malik, and Michael Jordan. Distribution-free, risk-controlling prediction sets. *Journal of the ACM (JACM)*, 68(6):1–34, 2021.

Anastasios Angelopoulos, Stephen Bates, Adam Fisch, Lihua Lei, and Tal Schuster. Conformal risk control. In *International conference on learning representations*, volume 2024, pages 55198–55218, 2024.

Yonatan Geifman and Ran El-Yaniv. Selectivenet: A deep neural network with an integrated reject option. In *International conference on machine learning*, pages 2151–2159. PMLR, 2019.

Yupeng Hu, Zixu Li, Zhiwei Chen, Qinlei Huang, Zhiheng Fu, Mingzhu Xu, and Liqiang Nie. Refine: Composed video retrieval via shared and differential semantics enhancement. *ACM Transactions on Multimedia Computing, Communications and Applications*, 22(7):1–24, 2026.

Zhiwei Chen, Yupeng Hu, Zixu Li, Zhiheng Fu, Guozhi Qiu, Weili Guan, and Liqiang Nie. Egoadapt: A multi-scene egocentric adaptation method for cvpr 2026 hd-epic vqa challenge. *arXiv preprint arXiv:2605.24500*, 2026b.

Zixu Li, Yupeng Hu, Zhiwei Chen, Shiqi Zhang, Qinlei Huang, Zhiheng Fu, and Yinwei Wei. Habit: Chrono-synergia robust progressive learning framework for composed image retrieval. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pages 6762–6770, 2026f.

Charles J Clopper and Egon S Pearson. The use of confidence or fiducial limits illustrated in the case of the binomial. *Biometrika*, 26(4):404–413, 1934.

Zixu Li, Yupeng Hu, Zhiwei Chen, Qinlei Huang, Guozhi Qiu, Zhiheng Fu, and Meng Liu. Retrack: Evidence-driven dual-stream directional anchor calibration network for composed video retrieval. In *AAAI*, volume 40, pages 23373–23381, 2026g.

Yulun Zhang, Zixu Li, Zhiwei Chen, Zhiheng Fu, Wenbo Wang, Zihang Qiu, Zhilin Wang, Ruxin Wang, and Yupeng Hu. Lightair: Lightweight action inversion and riemannian rectification for text-based person anomaly search. *arXiv preprint arXiv:2608.09152*, 2026f.

Qianyun Yang, Zhiwei Chen, Yupeng Hu, Zixu Li, Zhiheng Fu, and Liqiang Nie. Stable: Efficient hybrid nearest neighbor search via magnitude-uniformity and cardinality-robustness. *IEEE Transactions on Knowledge and Data Engineering*, 2026.

Prasanta Chandra Mahalanobis. On the generalized distance in statistics. *Sankhyā: The Indian Journal of Statistics, Series A (2008-)*, 80:S1–S7, 2018.