arXiv:2610.36763v1 [cs.IR] 27 Oct 2026

The Vocabulary Is Already Here: Adapting Pretrained Interpretable Features for Generative Recommendation

Ziluowen Luo
Central South University
Changsha, Hunan, China
lzlwddl@csu.edu.cn

Ruochen Liu
The Hong Kong University of Science and Technology (Guangzhou)
Guangzhou, China
rliu190@connect.hkust-gz.edu.cn

Jun Yin
Hong Kong Polytechnic University
Hongkong, China
Junmay.yin@connect.polyu.hk

Hao Yan
Central South University
ChangSha, Hunan, China
csuyh1999@csu.edu.cn

Zhengxuan Hu
Queen Mary University of London
London, UK
jp2024213778@qmul.ac.uk

Xinze Shi
Queen Mary University of London
London, UK
jp2024213114@qmul.ac.uk

Chaozhuo Li
Microsoft
China
lichaozhuo1991@gmail.com

Chengqi Zhang
Hong Kong Polytechnic University
Hongkong, China
chengqi.zhang@polyu.edu.hk

Senzhang Wang✉
Central South University
Changsha, China
szwang@csu.edu.cn

Abstract

Generative recommendation (GR) follows the token generation paradigm of language models, predicting the next item through its semantic ID (SID). Existing works commonly fit SID vocabularies to pretrained language model (PLM) representations of item descriptions, incurring additional training costs and making vocabulary coverage dependent on the fitting domains. Alternatively, we observe that pretrained sparse autoencoders (P-SAEs) expose PLM features associated with human-interpretable semantic annotations, allowing feature indices to serve as a natural SID vocabulary. However, reusing this vocabulary poses two challenges. First, *recommendation adaptation* requires preserving item distinctions under compact token budgets and learning to favor the target item over candidates with shared dominant features. Second, *retrieval explanation* requires extending feature interpretability from item content to retrieval decisions. To address these challenges, we propose ***API***, a framework that Adapts the Pretrained Interpretable vocabulary for effective and explainable GR without additional vocabulary learning. Within API, Pretrained Interpretable IDs (PIDs) provide compact item representations that retain activateed feature indices and relative activation strengths for item discrimination. With these PIDs as supervision, the Semantic Prediction Adapter (SPA) learns user preferences from interaction histories through joint feature-index and strength prediction. SPA then supervises agreement with the target PID and encourages the observed next item to outrank its alternatives. Using this score for retrieval, API further explains candidate score differences through feature contributions paired with their pretrained semantic annotations. Extensive experiments show that API achieves competitive performance across backbones and decoding schemes, while preserving item distinctions under the compact token budget and enabling explainable retrieval.

CCS Concepts

• **Information systems** → **Information retrieval**; • **Computing methodologies** → *Natural language processing*.

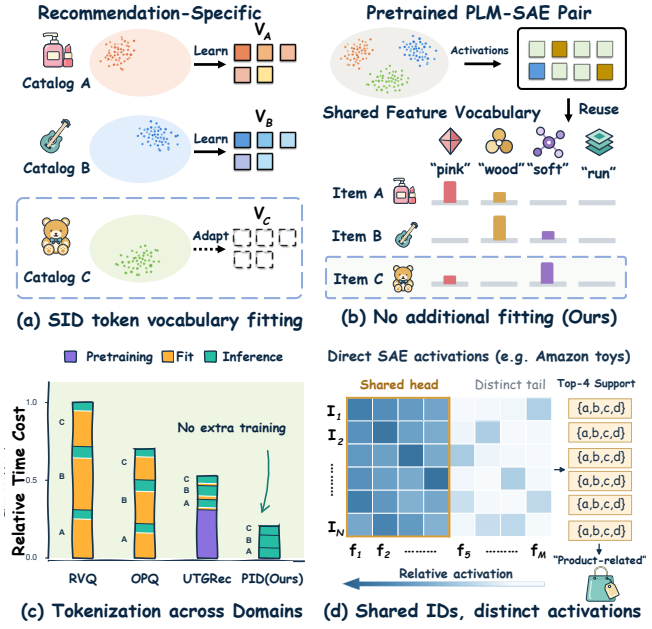

Figure 1: The vocabulary is already here. (a–c) Pretrained SAEs enable vocabulary sharing without extra training. (d) Distinct items can share the same dominant feature indices.

Keywords

Generative recommendation, Interpretability, Vocabulary reuse

Relevance: This work is relevant to the Track of *User modeling, personalization and recommendation*. This work contributes to recommendation via adapting the pretrained vocabulary and provides a practical path for applying interpretability to deployment.

1 Introduction

Generative recommendation systems (GRs) adopts the token generation paradigm of language models (LMs) to retrieve items conditioned on the user's interaction history [12, 13, 16]. A widely used approach represents items as the semantic IDs (SIDs), each consisting of a short sequence of discrete tokens. The mapping from items to SIDs, termed *item tokenization*, affects the GR pipeline from two fundamental perspectives: *item representation* and *prediction supervision* [28]. The resulting SIDs determine which item distinctions are retained and define the token-level predictive targets used to train the recommendation model. Therefore, a central problem in tokenization is how to construct compact SIDs that preserve item distinctions and provide effective supervision for recommendation.

Existing approaches commonly construct SIDs by fitting a special discrete vocabulary over item representations produced by pretrained language models (PLMs) [5, 20, 31]. Inspired by the strong generalizability of unified natural language vocabularies in LMs , recent efforts have explored shared vocabularies for generalizable item tokenization [4, 26]. For example, UTGRec [30] learns a transferable tokenizer through multi-domain pretraining, while UniTok [6] combines shared and specific experts for unified tokenization. Though these advances broaden tokenizer generalizability, they still construct vocabulary through additional fitting on the recommendation-specific data. The semantic coverage of the learned vocabulary therefore depends on the fitting data, while vocabulary construction incurs additional optimization costs. Moreover, fitting a vocabulary to recommendation-specific data does not necessarily yield SID tokens with readily interpretable semantics [2]. These observations raise a natural question: *Is there a free lunch for vocabulary construction in generative recommendation?*

On the basis of recent advances in *mechanistic interpretability*, we answer this question in the affirmative. Specifically, pretrained sparse autoencoders (P-SAEs) already provide feature dictionaries for their paired PLMs, with many features admitting human-interpretable descriptions [3]. For each item, the P-SAE encodes its PLM representation into a sparse activation vector, whose nonzero coordinates identify the expressed features and whose values indicate their strengths. For example, a red T-shirt might activate a feature associated with "red" more strongly than one associated with "blue". As illustrated in Figure 1(b–c), the indices of these features can serve as SID tokens, providing a natural vocabulary that can be reused without additional recommendation-specific vocabulary learning. This vocabulary also carries inspectable semantics that P-SAE features often respond to recognizable semantic patterns in their inputs, allowing the corresponding tokens to indicate which properties are expressed in an item.

However, bringing this vocabulary int effective and explainable GR still requires addressing two key challenges. First, ***recommendation adaptation*** requires preserving item distinctions in both SID construction and the supervision built on these SIDs. As illustrated in Figure 1(d), similar items can share dominant features while differing in their activation strengths and less active features. Representing items only by their most active feature indices can discard these distinctions, while predicting shared indices alone does not ensure that the target item outranks similar candidates. Moreover, while these features make item semantics inspectable, unleashing the potential of the vocabulary requires extending this interpretability to the retrieval. This motivates the second challenge, ***retrieval explanation***, which is to trace feature contributions to items and explain why the model favors one item over another.

To address these challenges, we propose ***API***, a generative recommendation framework that Adapting the Pretrained Interpretable vocabulary for effective and explainable generative recommendation. This adaptation spans two complementary levels: *item representation* through Pretrained interpretable ID (**PID**) construction, and *recommendation learning* through a Semantic Prediction Adapter (**SPA**). PID construction prioritizes strongly activated features while discounting those prevalent across the catalog, producing compact identifiers with accompanying activation evidence. Building on these representations, SPA incorporates feature identities and relative activation strengths into history modeling and aligns feature prediction of the target. To supervise these predictions, SPA combines sequence likelihood and strength agreement into retrieval and introduces preference alignment against competing items. API further traces feature-level contributions to retrieval, linking the semantics of each feature to the predicted evidence that favors a retrieved item. Our contributions are summarized as:

- We propose API, a generative recommendation framework that replaces recommendation-specific vocabulary learning with pretrained concept adaptation, opening a practical path for applying interpretability to practical deployment.
- We introduce PIDs that preserve item distinctions through feature selection and relative strengths, and theoretically characterize their separability under compact token budget.
- We develop SPA to effective and interpretable recommendation, aligning user preferences through item-level supervision and explain rankings through feature contributions.
- Experiments substantiate pretrained vocabulary reuse as a viable basis for GR, showing that API support competitive recommendation and enables explainable retrieval. Code is available at https://anonymous.4open.science/r/API-C1F9/.

2 Related Works

Generative recommendation and item tokenization. Generative recommendation (GR) makes item identifiers part of the prediction space, placing item tokenization at the interface between item representation and model supervision. Language-based approaches include the text-to-text formulation [4] and the learned textual identifiers [22]. For structured identifiers, TIGER constructs semantic IDs (SIDs) through residual quantization [9] of pretrained language model (PLM) representations [20] while SEATER combines balanced tree indexing with contrastive learning [21]. LC-Rec then incorporate collaborative semantics into item indexing [24, 29] and ETEGRec couples tokenizer learning with recommendation through alternating optimization [11]. UTGRec and UniTok further extend vocabulary learning across domains through shared tokenization [6, 30]. Despite these advances, the vocabulary learning still relies on recommendation-specific data which requires additional optimization and its semantic organization remains limited by the training domains. Our work extends pretrained knowledge reuse from continuous representations to the vocabulary units, adapting an existing PLM feature dictionary for recommendation.

Sparse autoencoders and pretrained feature reuse. Motivated by the *mechanistic interpretability* [1, 7], sparse autoencoders (SAEs) decompose language model activations into sparse combinations of dictionary features, exposing interpretable structure within dense representations [16]. Subsequent work scales these dictionaries [3, 23], improves reconstruction under sparsity constraints [19] and releases pretrained SAEs for reuse through resources such as Gemma Scope [10]. In recommendation, SAERec uses SAE-derived intent priors to enrich sequence modeling [25], but does not explore feature-level explanations of generative retrieval. In this work, we explores a practical route from mechanistic interpretability to recommendation systems, reusing a pretrained SAE feature dictionary as a shared vocabulary for GR. This further requires adapting reconstruction-oriented features for recommendation supervision and connecting feature-level readable concepts to item retrieval.

3 Preliminaries

Generative recommendation. Let $\mathcal{I}$ denote an item set, where each item $i \in \mathcal{I}$ is associated with a textual description x_i [14, 15]. Given a user's sequential interaction history $\mathcal{H}_u = (i_1, \ldots, i_T)$, the task is to predict the subsequent item $i^+ \in \mathcal{I}$. Generative recommendation (GR) approaches this task by generating a discrete token sequence conditioned on $\mathcal{H}_u$ and mapping the generating sequence to items. Autoregressive (AR) decoding generates tokens sequentially, conditioned on preceding tokens [20], whereas parallel (PR) decoding predicts multiple token in a single decoding step [5].

Item tokenization. For text-based item tokenization [4], a pretrained language model (PLM) (e.g., GPT-2 [17], T5 [18]) encodes the item description x_i into continuous hidden representations h_i. The tokenizer then converts them into a semantic ID (SID), a compact sequence of discrete tokens drawn from a shared set $\mathcal{V}$, referred to as the *vocabulary*. Formally, the tokenizer τ is defined:

$$\tau : \mathcal{I} \to \mathcal{V}^K, \qquad c_i = \tau(i) = (c_{i,1}, \ldots, c_{i,K}), \tag{1}$$

where K is the length and $c_{i,r} \in \mathcal{V}$ denotes the token at position r. The vocabulary thus determine the information retained in these tokens and the token-level supervision used for recommendation.

Pretrained sparse autoencoders. A pretrained sparse autoencoders (P-SAE) decomposes continuous PLM representations into sparsely activated dictionary features [3]. Given a hidden representation $\mathbf{h}$ from its paired PLM, the encoder E produces a non-negative activation vector $\mathbf{z} = E(\mathbf{h})$. The corresponding reconstruction is

$$\hat{\mathbf{h}} = \mathbf{D}\mathbf{z} + \mathbf{b}_{\text{dec}}, \tag{2}$$

where $\mathbf{D} = [\mathbf{d}_1, \ldots, \mathbf{d}_m]$ is the decoder dictionary and $\mathbf{b}_{\text{dec}}$ is the decoder bias. Each column $\mathbf{d}_j$ represents a feature identified by its index j and the activation z_j indicates how strongly that feature is expressed in $\mathbf{h}$. The associated pretrained semantic annotation m_j provides a human-readable description of the feature (e.g., "red"). These feature indices can thus serve as a reusable token vocabulary, accompanied by activation strengths and semantic annotations.

4 Methodology

In this section, we propose API, a generative recommendation framework comprising Pretrained Interpretable ID (PID) construction and the Semantic Prediction Adapter (SPA). As illustrated in Figure 2, PID construction selects pretrained features and retains their relative activation strengths to mitigate collisions(§ 4.1). Building on this representation, SPA then learns user preferences from interaction histories through PID prediction (§ 4.2), with item-level supervision encouraging the GR model to favor the target item over other candidates (§ 4.3). We further extend explanation from item representations to retrieval decisions, revealing the evidence behind the preference for one candidate over another (§ 4.4).

4.1 Pretrained Interpretable IDs Construction

To construct PIDs that preserve item distinctions, we first identify the pretrained features expressed in each item and quantify their strengths. For an item description sentence x_i, the frozen PLM–SAE pair[1] produces sparse activations $\mathbf{z}_i^{(t)}$ over token t. We then aggregate the activations into an item-level strength for feature j:

$$a_{ij} = \frac{\|\mathbf{d}_j\|_2}{|\mathcal{T}_i|} \sum_{t \in \mathcal{T}_i} z_{i,j}^{(t)}, \tag{3}$$

where $\mathcal{T}_i$ denotes the token positions in x_i, excluding special token. The resulting a_{ij} quantifies the item-level activation strength of feature j, indicating how strongly the item expresses that feature.

However, selecting feature indices solely by activation strength can assign the same token to distinct items. Within a domain, shared dominant features can occupy the token budget, excluding less active features that distinguish individual items. To allow these features to compete for inclusion, we form a candidate pool C_i containing the M features with the highest positive activation strengths. For each feature $j \in C_i$, we define a *selection priority* that favors strong activations while discounting features prevalent:

$$\pi_{ij} = \log(1 + a_{ij}) \, \log \frac{1}{\rho_j}. \tag{4}$$

where ρ_j denotes the smoothed prevalence of feature j, estimated from its occurrence in candidate pools of training-visible items. Lower prevalence can compensate for weaker activation, allowing less active attributes to compete with shared dominant features. We then select the K features with the highest priorities π_{ij} and order their indices by decreasing activation strength a_{ij}. Let $c_{i,r}$ denote the feature index at position r. The resulting sequence $c_i = (c_{i,1}, \ldots, c_{i,K})$ thus forms the token component of the PID.

Although selection by priority mechanism mitigates collisions caused by shared dominant features, a compact token budget still leave semantically similar items with identical token sequences. We observe that the selected features can nevertheless exhibit different relative activation strengths across them. To preserve these differences, we retain the relative activation strengths in a vector $\mathbf{q}_i = (q_{i,1}, \ldots, q_{i,K})$ aligned with the token sequence c_i. Each entry $q_{i,r}$ records normalized log strength of the feature indexed by $c_{i,r}$:

$$q_{i,r} = \frac{\log(1 + a_{i,c_{i,r}})}{\sum_{k=1}^{K} \log(1 + a_{i,c_{i,k}})}. \tag{5}$$

We retain the corresponding pretrained semantic concepts as $\mathbf{m}_i = (m_{c_{i,1}}, \ldots, m_{c_{i,K}})$, associating each feature index with its human-readable meaning. The feature-index sequence, relative activation

[1]We use the pretrained SAE released at https://decoderesearch.github.io/SAELens/latest/pretrained_saes/, together with the PLM and layer used for SAE pretraining.

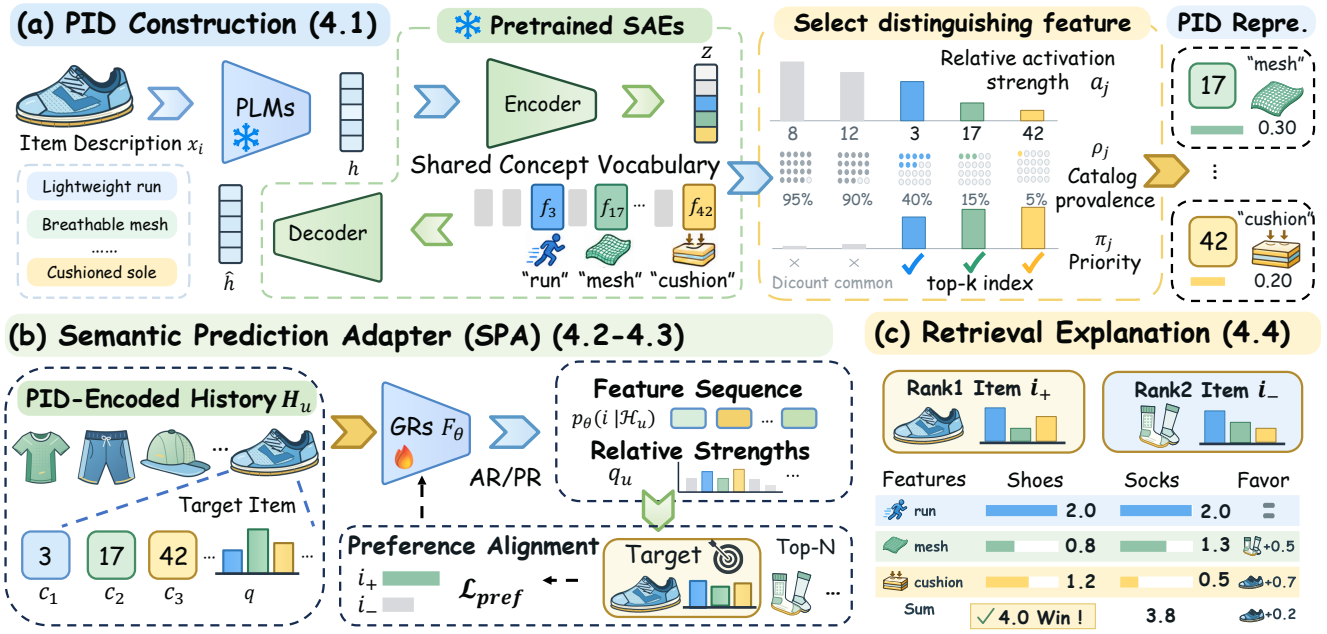

Figure 2: Overview of API. (a) PID construction selects pretrained features and retains their relative activation strengths and semantic annotations. (b) SPA learns user preferences through feature-index and strength prediction with item-level supervision. (c) Candidate score differences are decomposed into semantic feature contributions for retrieval explanation.

strengths, and aligned semantic annotations together form the PID:

$$\mathrm{PID}_i = \big(c_i, \mathbf{q}_i, \mathbf{m}_i\big). \tag{6}$$

By retaining relative activation strengths, PIDs can distinguish items that share the same feature-index sequence. The following proposition further bounds the number of such items.

Proposition 4.1 (Item separability with PID). *For a group of items sharing the same token sequence, let Δ denote the minimum pairwise ℓ_2 distance between their mean PLM representations. If $\Delta > 2\epsilon$, the maximum number $n_{\max}$ of items sharing the same PID satisfies*

$$n_{\max} \le 1 + \left\lfloor \frac{\delta}{\Delta - 2\epsilon} \right\rfloor, \tag{7}$$

where ϵ bounds the reconstruction error using the selected features, and δ bounds the range of $\sum_{r=1}^{K} a_{i,c_{i,r}}$ across items in the group.

Proof. See Appendix B.1. The bound relates collisions to reconstruction error and variation in summed strengths. When $\Delta > 2\epsilon + \delta$,the strengths can distinguish items when their token sequence coincide.

4.2 SPA: PID Prediction

Although PIDs retain relative activation strengths alongside feature, predicting the indices alone does not explicitly model these strengths. The Semantic Prediction Adapter (SPA) therefore incorporates feature-index and activation strengths into history encoding and predicts the next item's token sequence and relative strengths.

Encoding history based PIDs. Though the PID construction selects K features for compact item discrimination, the excluded may still provide evidence of user preferences. To retain this, SPA encodes each historical item using the candidate pool C_i from which c_i is selected, together with the corresponding activation strengths. Specifically, we represent item i by the feature–strength pairs $\mathcal{E}_i = \{(j, a_{ij}) : j \in C_i\}$. The feature encoder Ψ_θ combines trainable feature-index embeddings with activation-strength encodings and aggregates them into a fixed number of tokens per item. Given the interaction history $\mathcal{H}_u = (i_1, \ldots, i_T)$, a recommendation Transformer F_θ processes these tokens in interaction order:

$$\mathbf{H}_u = F_\theta\big(\Psi_\theta(\mathcal{E}_{i_1}), \ldots, \Psi_\theta(\mathcal{E}_{i_T})\big)\,. \tag{8}$$

Predicting target PIDs. Given the history representation $\mathbf{H}_u$, SPA models the next item's feature-index sequence through the conditional distribution $p_\theta(c \mid \mathbf{H}_u)$. SPA also predicts a strength vector $\hat{\mathbf{q}}_u$ from $\mathbf{H}_u$, where $\hat{q}_{u,r}$ estimates $q_{i^+,r}$ at activation-ranked position r. Its entries satisfy $\hat{q}_{u,r} \ge 0$ and $\sum_{r=1}^{K} \hat{q}_{u,r} = 1$. The semantic annotations are obtained by looking up m_j for each predicted feature index j. Its entries are normalized across the K positions, satisfying $\sum_{r=1}^{K} \hat{q}_{u,r} = 1$. In practice, the feature-index sequence can be generated autoregressively [4] or in parallel [5]. The remaining problem is how to favor the next item over competing catalog items.

4.3 SPA: Item-Level Supervision

Building on the feature-index and strength predictions, we use the observed next item's PID to construct item-level supervision for GR training. We combine sequence likelihood and strength agreement into the score to supervise PID prediction and candidates ranking.

Comparing candidate PIDs. For a item i, let $\ell_u(i) = \log p_\theta(c_i \mid \mathbf{H}_u)$ denote the log likelihood of its feature-index sequence. We then define the item score by balancing sequence likelihood $\ell_u(i)$ with

Table 1: Overall recommendation performance. Best and second-best results are shown in bold and underlined, respectively.

	Beauty				Sports				Toys			
Method	R@5	R@10	N@5	N@10	R@5	R@10	N@5	N@10	R@5	R@10	N@5	N@10
GRU4Rec	0.0433	0.0646	0.0304	0.0372	0.0222	0.0371	0.0141	0.0189	0.0384	0.0543	0.0272	0.0324
Bert4Rec	0.0281	0.0434	0.0184	0.0233	0.0129	0.0224	0.0079	0.0110	0.0201	0.0313	0.0133	0.0168
Tiger (2023)	**0.0464**	0.0662	0.0305	0.0369	0.0237	0.0341	0.0156	0.0189	0.0402	0.0584	0.0263	0.0322
LCRec (2024)	0.0228	0.0362	0.0159	0.0197	0.0082	0.0133	0.0073	0.0089	0.0174	0.0285	0.0119	0.0157
RPG (2025)	0.0322	0.0502	0.0210	0.0269	0.0176	0.0266	0.0124	0.0153	0.0306	0.0477	0.0192	0.0247
UTGRec (2026)	0.0193	0.0330	0.0122	0.0167	0.0126	0.0219	0.0080	0.0109	0.0177	0.0304	0.0107	0.0147
Ghost (2026)	0.0388	0.0590	0.0257	0.0322	0.0260	**0.0422**	0.0171	0.0223	0.0333	0.0543	0.021	0.0277
DIGR (2026)	0.0431	0.0605	**0.0307**	0.0362	0.0186	0.0291	0.0129	0.0161	0.0396	0.0571	0.0268	0.0319
TDPM (2026)	0.0418	0.0655	0.0279	0.0355	0.0216	0.0321	0.0140	0.0173	0.0398	0.0591	0.0259	0.0322
API (Ours)	0.0457	**0.0694**	0.0305	**0.0389**	**0.0271**	0.0411	**0.0188**	**0.0239**	**0.0414**	**0.0616**	**0.0291**	**0.0353**

agreement between the predicted and candidate relative strengths:

$$s_\theta(u,i) = \ell_u(i) - \beta D_{\mathrm{H}}(\hat{\mathbf{q}}_u, \mathbf{q}_i), \tag{9}$$

where $\beta \geq 0$ controls the strength penalty. The squared Hellinger distance D_{H} compares the predicted strengths $\hat{\mathbf{q}}_u$ with the candidate's strengths $\mathbf{q}_i$ at corresponding activation-ranked positions:

$$D_{\mathrm{H}}(\hat{\mathbf{q}}_u, \mathbf{q}_i) = 1 - \sum_{r=1}^{K} \sqrt{\hat{q}_{u,r} q_{i,r}}. \tag{10}$$

For these normalized vectors, the distance lies in $[0,1]$ and equals zero exactly when their relative strengths match.

Learning user preferences. For the observed next item i^+, minimizing $-s_\theta(u,i^+)$ supervises its feature-index sequence and relative strengths. However, a better match to the target PID still leave competing items with higher scores. To encourage the target item to outrank competing candidates, we introduce a preference alignment loss over the candidate set $\mathcal{J}_u \subseteq \mathcal{I}$ containing i^+:

$$\mathcal{L}_{\mathrm{pref}}(u,i^+) = \log \sum_{n \in \mathcal{J}_u} \exp\big(s_\theta(u,n) - s_\theta(u,i^+)\big). \tag{11}$$

We then minimize the joint loss over observed interactions:

$$\mathcal{L} = \mathbb{E}_{(u,i^+)}\big[-s_\theta(u,i^+) + \lambda \mathcal{L}_{\mathrm{pref}}(u,i^+)\big], \tag{12}$$

where $\lambda \geq 0$ weights preference alignment. The following proposition shows how the relative strengths in PIDs can help SPA rank the target above the competing items with the same feature sequence.

Proposition 4.2 (Item ranking with SPA). *For a target item i^+ and a competing item i^- sharing the same feature-index sequence, let $\Delta_q = D_{\mathrm{H}}(\mathbf{q}_{i^+}, \mathbf{q}_{i^-}) > 0$ denote their strength separation. For $\beta > 0$, the target's ranking margin satisfies*

$$s_\theta(u,i^+) - s_\theta(u,i^-) \geq \beta\big(\Delta_q - 2\sqrt{\Delta_q \varepsilon_u}\big). \tag{13}$$

where $\hat{\mathbf{q}}_u$ denotes SPA's prediction of the next item's relative strengths given $\mathbf{H}_u$, and $\varepsilon_u = D_{\mathrm{H}}(\hat{\mathbf{q}}_u, \mathbf{q}_{i^+})$ measures its error against the target PID. In particular, $\varepsilon_u < \Delta_q/4$ ensures that SPA ranks i^+ above i^-.

Proof. See Appendix B.2. Under the stated condition, the strength differences retained in PIDs yield a positive ranking margin for the target. At retrieval, SPA ranks items using the same score $s_\theta(u,i)$.

4.4 Retrieval Explanation

To explain why SPA favors one candidate over another, we decompose their margin into contributions from individual interpretable features. We then associate these contributions with their semantic annotations to explain the feature evidence behind the ranking.

Decomposing retrieval scores to pretrained features. To quantify each feature's contribution to the item score, we decompose the sequence log likelihood as $\ell_u(i) = \sum_{r=1}^{K} \ell_{u,r}(i)$. Here, $\ell_{u,r}(i)$ denotes the log probability assigned by SPA to feature index $c_{i,r}$ at position r. Combining this term with the strength agreement at the same position gives the feature contribution:

$$\phi_{u,i}(c_{i,r}) = \ell_{u,r}(i) + \beta\sqrt{\hat{q}_{u,r} q_{i,r}}, \tag{14}$$

where $\phi_{u,i}(c_{i,r})$ measures the contribution of feature $c_{i,r}$ to candidate i's score. For features j absent from c_i, we set $\phi_{u,i}(j) = 0$.

Explaining candidate rankings. To explain the ranking of candidates i and i', we compare the contributions of each feature j to their respective scores. Their score difference decomposes as:

$$s_\theta(u,i) - s_\theta(u,i') = \sum_j \big(\phi_{u,i}(j) - \phi_{u,i'}(j)\big). \tag{15}$$

Each term measures feature j's contribution to this difference, with positive values favoring i and negative values favoring i'. We pair these contributions with the pretrained semantic annotations m_j to show which features favor each candidate and by how much.

5 Experiments

We evaluate pretrained vocabulary reuse and proposed API to generative recommendation through the following research questions:

- **RQ1:** How well does API recommend across decoding strategies?
- **RQ2:** How well can PIDs preserve item distinctions and content?
- **RQ3:** Which supervision signals improve item retrieval?
- **RQ4:** How effectively can the pretrained vocabulary be reused?
- **RQ5:** Can item retain readable meanings with the PID construction and GR model support targeted-feature steering to retrieval?

5.1 Experimental Setup

Benchmark. We use the Beauty, Sports and Toys subsets of the preprocessed Amazon datasets aligned with prior work [4, 8], retaining

Table 2: Recommendation performance on three datasets with T5 and GPT-2 backbones. K counts semantic tokens. R@k and N@k denote Recall@k and NDCG@k. The PID with the SPA constitutes API. For each backbone and decoding strategy, bold and underlined values indicate the best and second-best results, respectively, including ties.

Back.	Generation	Tokenization	Beauty				Sports				Toys			
			R@5	R@10	N@5	N@10	R@5	R@10	N@5	N@10	R@5	R@10	N@5	N@10
T5	AR	**Random**	0.0310	0.0454	0.0216	0.0262	0.0082	0.0111	0.0051	0.0060	0.0188	0.0250	0.0140	0.0161
		RVQ	0.0464	0.0662	0.0305	0.0369	0.0237	0.0341	0.0156	0.0189	0.0402	0.0584	0.0263	0.0322
		RQVAE	0.0402	0.0602	0.0269	0.0334	0.0215	0.0310	0.0139	0.0170	0.0344	0.0492	0.0240	0.0290
		RKMeans	0.0441	0.0667	0.0305	0.0378	0.0217	0.0330	0.0145	0.0181	0.0384	0.0564	0.0256	0.0314
		HiD-VAE	0.0419	0.0644	0.0283	0.0355	0.0215	0.0328	0.0138	0.0175	0.0388	0.0595	0.0257	0.0324
		PID	0.0435	0.0630	0.0299	0.0361	0.0210	0.0346	0.0135	0.0178	0.0401	0.0585	0.0269	0.0328
	+ SPA	**PID**	**0.0476**	**0.0701**	**0.0323**	**0.0396**	**0.0243**	**0.0372**	**0.0157**	**0.0198**	**0.0419**	**0.0607**	**0.0287**	**0.0347**
	PR	**OPQ ($K = 4$)**	0.0290	0.0466	0.0196	0.0256	0.0057	0.0105	0.0034	0.0049	0.0248	0.0373	0.0160	0.0200
		OPQ ($K = 8$)	0.0368	0.0543	0.0247	0.0303	0.0080	0.0134	0.0052	0.0070	0.0320	0.0485	0.0210	0.0263
		OPQ ($K = 16$)	0.0312	0.0484	0.0203	0.0259	0.0114	0.0199	0.0076	0.0102	0.0278	0.0424	0.0171	0.0218
		PID ($K = 4$)	0.0278	0.0402	0.0194	0.0234	0.0049	0.0087	0.0029	0.0041	0.0294	0.0488	0.0182	0.0245
		PID ($K = 8$)	0.0370	0.0541	0.0262	0.0316	0.0094	0.0158	0.0062	0.0083	0.0332	**0.0547**	0.0207	0.0276
		PID ($K = 16$)	0.0340	0.0523	0.0225	0.0284	0.0132	0.0194	0.0086	0.0106	0.0348	0.0509	0.0209	0.0262
	+ SPA	**PID ($K = 4$)**	0.0388	0.0572	0.0264	0.0323	0.0168	0.0267	0.0111	0.0142	0.0352	0.0517	**0.0242**	**0.0295**
		PID ($K = 8$)	**0.0401**	**0.0627**	**0.0274**	**0.0346**	**0.0192**	**0.0289**	**0.0126**	**0.0157**	**0.0354**	0.0517	0.0229	0.0282
		PID ($K = 16$)	0.0394	0.0575	0.0271	0.0329	0.0162	0.0241	0.0106	0.0131	0.0319	0.0462	0.0210	0.0256
GPT2	AR	**Random**	0.0292	0.0401	0.0207	0.0247	0.0070	0.0122	0.0039	0.0056	0.0133	0.0176	0.0094	0.0108
		RVQ	0.0394	0.0594	0.0255	0.0320	0.0197	0.0331	0.0126	0.0169	0.0327	0.0536	0.0206	0.0274
		RQVAE	0.0362	0.0542	0.0242	0.0301	0.0166	0.0271	0.0103	0.0137	0.0296	0.0439	0.0192	0.0238
		RKMeans	0.0423	0.0640	0.0278	0.0348	0.0182	0.0294	0.0110	0.0145	0.0313	0.0508	0.0195	0.0258
		HiD-VAE	0.0398	0.0607	0.0261	0.0329	0.0223	0.0360	0.0141	0.0185	0.0335	0.0567	0.0211	0.0286
		PID	0.0420	0.0663	0.0282	0.0361	0.0190	0.0292	0.0126	0.0159	0.0380	0.0529	0.0265	0.0313
	+ SPA	**PID**	**0.0457**	**0.0694**	**0.0305**	**0.0389**	**0.0271**	**0.0411**	**0.0188**	**0.0239**	**0.0414**	**0.0616**	**0.0291**	**0.0353**
	PR	**OPQ ($K = 4$)**	0.0292	0.0457	0.0197	0.0250	0.0100	0.0175	0.0065	0.0089	0.0267	0.0449	0.0164	0.0222
		OPQ ($K = 8$)	0.0374	0.0549	0.0248	0.0305	0.0161	0.0241	0.0109	0.0135	0.0274	0.0441	0.0165	0.0219
		OPQ ($K = 16$)	0.0322	0.0502	0.0210	0.0269	0.0176	0.0266	0.0124	0.0153	0.0306	0.0477	0.0192	0.0247
		PID ($K = 4$)	0.0267	0.0393	0.0193	0.0234	0.0153	0.0234	0.0101	0.0127	0.0330	0.0483	0.0230	0.0279
		PID ($K = 8$)	0.0379	0.0540	0.0270	0.0322	0.0162	**0.0309**	0.0118	**0.0167**	0.0307	0.0447	0.0214	0.0259
		PID ($K = 16$)	0.0336	0.0525	0.0228	0.0288	0.0179	0.0281	0.0121	0.0156	0.0267	0.0383	0.0189	0.0226
	+ SPA	**PID ($K = 4$)**	**0.0411**	**0.0618**	**0.0277**	**0.0343**	**0.0195**	0.0300	**0.0129**	0.0163	**0.0393**	**0.0559**	**0.0281**	**0.0334**
		PID ($K = 8$)	0.0395	0.0573	0.0274	0.0331	0.0186	0.0294	0.0120	0.0154	0.0384	0.0531	0.0271	0.0318
		PID ($K = 16$)	0.0398	0.0573	**0.0277**	0.0333	0.0161	0.0253	0.0104	0.0133	0.0377	0.0515	0.0262	0.0306

their dataset splits and evaluation protocol. Item descriptions combine the title, brand, category and price fields. Recommendation models are trained on training interactions; PID selection statistics and activation calibration are estimated from training-item contents. Dataset statistics and details are provided in Appendix A.

Compared methods. We organize the main comparisons at two levels: item tokenization in Table 2 and complete recommendation frameworks in Table 1. The tokenization comparison includes Random, OPQ, HiD-VAE and residual quantization baselines (RVQ, RQ-VAE and RKMeans) that aligned with GRID [8]. We then evaluate them with T5- and GPT-2-style backbones, adopting AR and PR implementations based on TIGER [20] and RPG [5], respectively. Detailed baseline and configurations are provided in Appendix A.3.

Implementation. We extract item features using pretrained GPT-2-small [17] and its paired P-SAE with 24,576 labeled dictionary features, keeping both models frozen. The T5-style and GPT-2-style recommendation backbones are trained from random initialization. We retain 32 candidate features per item and use $K \in \{4, 8, 16\}$ tokens. We optimize the recommendation models using Adam with a learning rate of 10^{-3} and weight decay of 10^{-4}. The loss and strength weights are set to $\lambda_{\text{pref}} = 0.1$ and $\beta = 8$, respectively. Checkpoints are selected by Recall@5 on the validation split. Detailed configurations are provided in Appendix A.3.

Evaluation. We report Recall@5/10 and NDCG@5/10 on the complete test split, using the full catalog without sampled negatives. For PID discrimination, *collision rate* measures the fraction of items whose selected feature set is shared by another item. For content preservation, *Overlap@20* measures the average fraction of shared neighbors between representation-based and reference top-20 lists. For steering, we report *feature exposure@10*, the fraction of top-10 recommendations whose PIDs contain the target feature. PR uses exact catalog scoring, while AR uses catalog-constrained beam search with a width of 50. Detailed are provided in Appendix A.1.

Findings for RQ1 (Recommendation Performance)

Finding 1: Effective vocabulary reuse. Pretrained features support competitive GR without additional vocabulary learning.
Finding 2: Benefits of recommendation adaptation. SPA improves over direct PID in all evaluated AR settings and most PR settings.
Finding 3: Competitive framework performance. API achieves the best reported results in nine of twelve dataset–metric comparisons.

5.2 Recommendation Performance (RQ1)

Table 1 compares the complete API framework with sequential and generative recommenders. **API achieves the best reported result in nine of the twelve dataset–metric comparisons**, including all four metrics on Toys. It ranks second on Beauty Recall@5/NDCG@5 and Sports Recall@10, with a tie on Beauty NDCG@5. Table 2 further examines how recommendation adaptation affects the use of the pretrained vocabulary. Direct PID does not uniformly outperform learned tokenizers: on Beauty with T5-PR and $K = 4$, its Recall@5 is 0.0278, compared with 0.0290 for OPQ, while adding SPA raises it to 0.0388. Across the two backbones and three datasets, **SPA improves Recall@5 over direct PID in all six AR settings and 16 of the 18 PR settings**. With T5-AR, the complete API pipeline exceeds the strongest learned-tokenizer baseline by 2.5%–4.2% in Recall@5. Both PR exceptions occur at $K = 16$, while the best API-PR Recall@5 for every backbone–dataset pair is attained at $K = 4$ or $K = 8$. These results also show that **increasing identifier length does not consistently improve recommendation**.

5.3 PID Quality (RQ2)

Table 3 compares feature selection policies within the same top-32 candidate pool, measuring collisions between selected feature sets. **PID improves item discrimination under compact budgets**: at $K = 4$, it reduces the mean collision rate across the three datasets from 80.76% under native selection to 17.09%. The policies coincide at $K = 32$, when the entire candidate pool is retained. Random selection yields even fewer collisions, but Figure 55 shows substantially weaker content alignment at $K = 4$ and $K = 8$. At these budgets, PID achieves the highest Overlap@20 among the compared methods on Beauty and Sports against both title-based and PLM-based references. These comparisons indicate that **low collision rates alone do not establish content preservation**. To further separate these properties, Table 4 permutes complete feature–activation profiles across items while keeping item content fixed. Collision rates remain unchanged, whereas Overlap@20 falls below 0.2% against both references. Finally, **activation-aware neighborhoods show stronger content alignment**: with the selected feature indices fixed at $K = 8$, activation-weighted cosine neighborhoods exceed index-only Jaccard neighborhoods by 1.03%, 2.60% and 4.41 % in PLM-based Overlap@20 on all benchmarks.

Findings for RQ2 (PID Quality)

Finding 1: Compact discrimination. Priority-based feature selection reduces collisions within the same frozen vocabulary.
Finding 2: Discrimination do not guarantee content alignment.

Table 3: Collision-affected items (%) across token budgets.

Dataset	Selection	$K = 4$	$K = 8$	$K = 16$	$K = 24$	$K = 32$
Beauty	Native SAE	77.37	42.72	11.11	2.86	1.23
	Random ID	2.50	0.00	0.00	0.00	1.23
	PID	15.16	3.33	1.26	1.19	1.23
Sports	Native SAE	79.96	51.77	19.98	6.30	3.44
	Random ID	3.67	0.00	0.00	0.00	3.44
	PID	21.48	7.71	3.43	3.20	3.44
Toys	Native SAE	84.95	57.59	17.96	4.68	1.49
	Random ID	2.59	0.00	0.00	0.00	1.49
	PID	14.64	4.54	1.51	1.49	1.49

Table 4: Content alignment under collision-preserving permutation (%). Collision rates (Coll.) remain unchanged.

		Index only				Index + Activation			
		Title		PLM		Title		PLM	
Dataset	Coll.	Orig.	Perm.	Orig.	Perm.	Orig.	Perm.	Orig.	Perm.
Beauty	3.33	24.15	0.19	33.02	0.16	25.54	0.17	34.05	0.14
Sports	7.71	24.83	0.12	37.01	0.16	26.21	0.11	39.61	0.15
Toys	4.54	25.90	0.12	38.81	0.12	27.47	0.15	43.22	0.14

5.4 SPA Supervision (RQ3)

Table 5 examines preference alignment and strength modeling under T5-PR with $K = 8$, fixing the SPA architecture, activation-aware history encoding, feature-index targets, and output ordering. Starting from SID prediction alone, we add preference alignment, strength prediction and scoring, or both. **Preference alignment consistently improves retrieval**, increasing Recall@10 by 13.3%, 45.9%, and 9.4% on Beauty, Sports, and Toys, respectively. Adding strength prediction and scoring alone lowers NDCG@10 on all three datasets. Once preference alignment is included, however, strength modeling further improves NDCG@10 by 4.0%–6.0%. Full SPA achieves the highest Recall@10 and NDCG@10 throughout this ablation. These results suggest that **strength evidence is more effective when coupled with item-level preference alignment**.

Findings for RQ3

Finding 1: Learning through item competition. Catalog-grounded supervision provides the most consistent gains in this ablation.
Finding 2: Complementary strength supervision. SPA provides additional ranking gains when combined with preference alignment.

5.5 Vocabulary Reuse (RQ4)

Figure 3 evaluates learned tokenizers fitted on Beauty or Toys and frozen across all three target catalogs. PID retains the frozen P-SAE dictionary and estimates IDF from target training-item contents; all recommenders are trained on target-domain interactions. On Beauty and Toys, several learned tokenizers incur more collisions and lower Recall@5 when fitted on the other domain. **PID reuses the vocabulary without refitting**, maintaining comparatively low collision rates and competitive Recall@5 across targets: it leads on Toys and Sports under both source configurations, while R-SAE remains ahead on Beauty. Native top-K selection produces substantially more collisions than PID using the same dictionary,

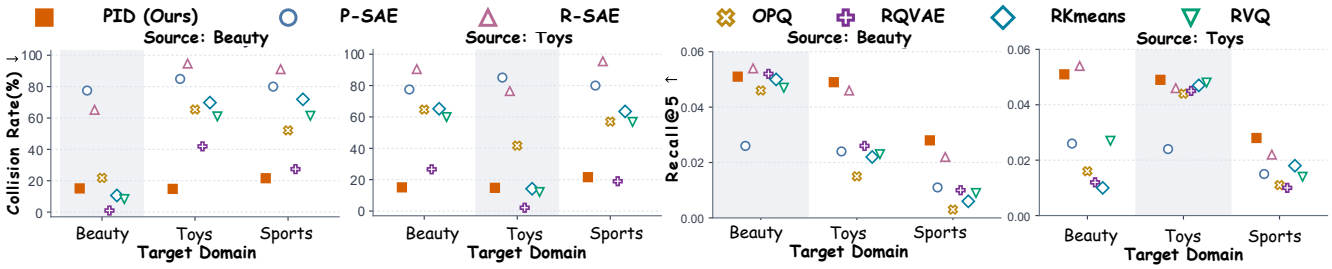

Figure 3: Vocabulary reuse across domains. Collision rates and downstream Recall@5 are reported for each source–target pair; shaded columns mark the source domain of data-trained tokenizers.

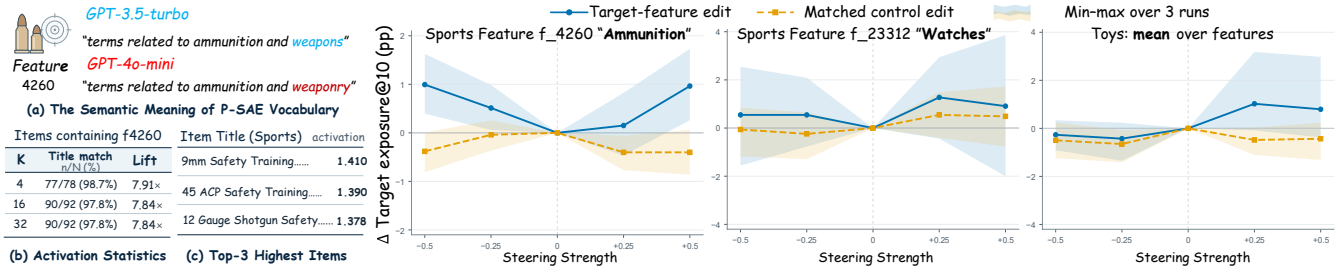

Figure 4: Interpretability and controlled steering with P-SAE features. Left: The labeled f_{4260} by the LLMs, title matches to the meanings and top-activating items. K is the PID budget; lift uses the catalog baseline ($\approx 12.5\%$). Right: changes in target items exposure@10 under activation steering, relative to the non-target. Shaded bands show the min–max range over three runs.

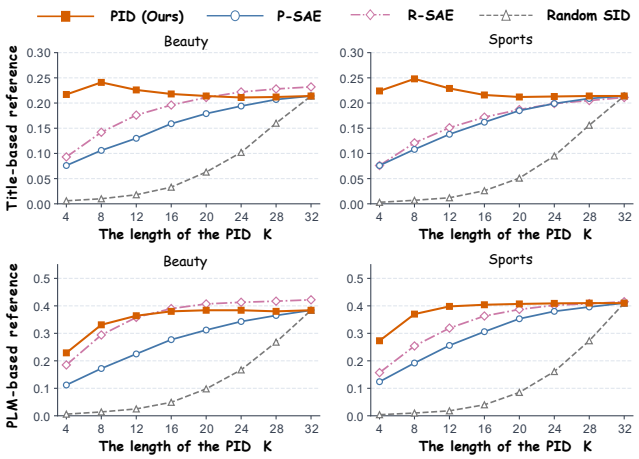

Figure 5: Content alignment across semantic-token budgets, measured by nearest-neighbor overlap@20 against title-based and PLM-based references. Higher values indicate greater agreement with the reference neighborhoods.

highlighting **the importance of catalog-specific selection**. Figure 6(b) extrapolates single-task timings to three equal-workload tokenization tasks, counting UTGRec pretraining once and applicable fitting and processing per task. Treating the pretrained PLM and P-SAE as available resources, cumulative time relative to one GPT-2 encoding task is 3.78× for API, 407.69× for RVQ, and 674.99× for

Table 5: Supervision ablation of SPA under T5-style parallel prediction with $K = 8$ features.

	Beauty		Sports		Toys	
Supervision	R@10	N@10	R@10	N@10	R@10	N@10
Index-only	0.05080	0.02959	0.01980	0.01104	0.04523	0.02518
+ Strength	0.04905	0.02833	0.01893	0.01055	0.04533	0.02447
+ Preference	0.05755	0.03312	0.02888	0.01509	0.04950	0.02656
Full SPA	**0.06265**	**0.03456**	**0.02891**	**0.01569**	**0.05172**	**0.02816**

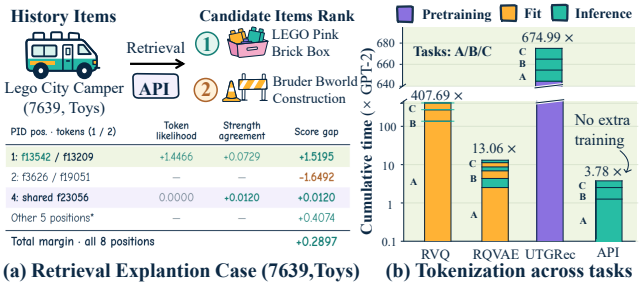

Figure 6: Retrieval explanation and tokenization efficiency. (a) A Toys case showing candidate score-gap contributions grouped by PID position. (b) Projected umulative tokenization time for three tasks, counting UTGRec pretraining once.

UTGRec. **API thus avoids additional vocabulary-fitting costs**, while retaining per-task feature extraction and encoding.

5.6 Retrieval Explanation (RQ5)

Figure 4 illustrates **readable feature semantics** through Sports feature f_{4260}, whose external descriptions refer to ammunition and weapons. Among training items whose $K = 4$ PIDs contain this feature, 77 of 78 match the hunting-and-shooting lexical rule, yielding 98.7% agreement and a 7.91× lift over the rule's prevalence in the training catalog. The highest-activation titles also describe training ammunition. The Toys case in Figure 6 further provides **traceable ranking evidence**: a LEGO brick box outranks Bruder construction accessories by a margin of +0.2897, comprising +0.2993 from token likelihoods and −0.0096 from strength agreement. At PID position 1, the contrast between LEGO-related f_{13542} and f_{13209} favors the LEGO item, whereas position 2 opposes it. Finally, we edit existing history activations while keeping model parameters and candidate PIDs fixed to probe **feature-dependent retrieval responses**. For f_{4260}, a +0.50 shift in standardized log-activation increases exposure@10 by 1.365 percentage points relative to matched control edits on 332 eligible test users (Appendix C.2).

Findings for RQ4

Finding 1: Vocabulary sharing without refitting. Frozen pretrained dictionary supports competitive recommendation across benchmarks.

Finding 2: Adaptation within a shared vocabulary. PID reduces the collisions while preserving the pretrained feature-index identities.

Findings for RQ5

Finding 1: Inspectable ranking evidence. Score decomposition links readable features to evidence supporting or opposing a candidate.

Finding 2: Local steering effects. Activation edits can change exposure, with responses varying across features and doses.

6 Conclusion

We propose API to address catalog-dependent vocabulary learning in generative recommendation by adapting a pretrained interpretable dictionary into a shared vocabulary. Through PID construction and SPA, API connects compact item representations and user preference learning with interpretable retrieval evidence, opening a practical path for applying mechanistic interpretability to recommendation. Learning more general-purpose interpretable vocabularies and fully realizing their potential for recommendation remain promising directions for future research.

LLM Usage. Assist with baseline reproduction, refining grammar, and reviewing mathematical proofs. The authors take full responsibility for the final manuscript and experimental results.

References

[1] Trenton Bricken, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermyn, Tom Conerly, Nick Turner, Cem Anil, Carson Denison, Amanda Askell, Robert Lasenby, Yifan Wu, Shauna Kravec, Nicholas Schiefer, Tim Maxwell, Nicholas Joseph, Zac Hatfield-Dodds, Alex Tamkin, Karina Nguyen, Brayden McLean, Josiah E Burke, Tristan Hume, Shan Carter, Tom Henighan, and Christopher Olah. 2023. Towards Monosemanticity: Decomposing Language Models With Dictionary Learning. *Transformer Circuits Thread* (2023). https://transformer-circuits.pub/2023/monosemantic-features/index.html

[2] Dengzhao Fang, Jingtong Gao, Chengcheng Zhu, Yu Li, Xiangyu Zhao, and Yi Chang. 2025. HiD-VAE: Interpretable Generative Recommendation via Hierarchical and Disentangled Semantic IDs. arXiv:2508.04618 [cs.IR] doi:10.48550/arXiv.2508.04618

[3] Leo Gao, Tom Dupré la Tour, Henk Tillman, Gabriel Goh, Rajan Troll, Alec Radford, Ilya Sutskever, Jan Leike, and Jeffrey Wu. 2025. Scaling and evaluating sparse autoencoders. In *The Thirteenth International Conference on Learning Representations*. https://proceedings.iclr.cc/paper_files/paper/2025/file/42ef3308c230942d223c411adf182c88-Paper-Conference.pdf

[4] Shijie Geng, Shuchang Liu, Zuohui Fu, Yingqiang Ge, and Yongfeng Zhang. 2022. Recommendation as Language Processing (RLP): A Unified Pretrain, Personalized Prompt & Predict Paradigm (P5). In *Proceedings of the 16th ACM Conference on Recommender Systems*. https://arxiv.org/abs/2203.13366

[5] Yupeng Hou, Jiacheng Li, Ashley Shin, Jinsung Jeon, Abhishek Santhanam, Wei Shao, Kaveh Hassani, Ning Yao, and Julian McAuley. 2025. Generating Long Semantic IDs in Parallel for Recommendation. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2*. Association for Computing Machinery. doi:10.1145/3711896.3736979

[6] Yu Hou and Won-Yong Shin. 2026. Tokenize Once, Recommend Anywhere: Unified Item Tokenization for Multi-domain LLM-based Recommendation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 40. 14848–14855. doi:10.1609/aaai.v40i17.38505

[7] Robert Huben, Hoagy Cunningham, Logan Riggs Smith, Aidan Ewart, and Lee Sharkey. 2024. Sparse Autoencoders Find Highly Interpretable Features in Language Models. In *The Twelfth International Conference on Learning Representations*. https://openreview.net/forum?id=F76bwRSLeK

[8] Clark Mingxuan Ju, Liam Collins, Leonardo Neves, Bhuvesh Kumar, Louis Yufeng Wang, Tong Zhao, and Neil Shah. 2025. Generative Recommendation with Semantic IDs: A Practitioner's Handbook. In *Proceedings of the 34th ACM International Conference on Information and Knowledge Management*. Association for Computing Machinery. doi:10.1145/3746252.3761612

[9] Doyup Lee, Chiheon Kim, Saehoon Kim, Minsu Cho, and Wook-Shin Han. 2022. Autoregressive Image Generation using Residual Quantization. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 11513–11522. doi:10.1109/CVPR52688.2022.01123

[10] Tom Lieberum, Senthooran Rajamanoharan, Arthur Conmy, Lewis Smith, Nicolas Sonnerat, Vikrant Varma, Janos Kramar, Anca Dragan, Rohin Shah, and Neel Nanda. 2024. Gemma Scope: Open Sparse Autoencoders Everywhere All At Once on Gemma 2. In *Proceedings of the 7th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*. Association for Computational Linguistics, 278–300. doi:10.18653/v1/2024.blackboxnlp-1.19

[11] Enze Liu, Bowen Zheng, Cheng Ling, Lantao Hu, Han Li, and Wayne Xin Zhao. 2025. Generative Recommender with End-to-End Learnable Item Tokenization. In *Proceedings of the 48th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '25)*. Association for Computing Machinery. doi:10.1145/3726302.3729989

[12] Ruochen Liu, Hao Chen, Yuanchen Bei, Lijia Chen, Qijie Shen, Feiran Huang, Fakhri Karray, and Senzhang Wang. 2026. SetLLM: Set Large Language Model for Cold-Start Item Recommendation. In *Proceedings of the 32nd ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 2*. 3282–3293.

[13] Ruochen Liu, Hao Chen, Yuanchen Bei, Qijie Shen, Fangwei Zhong, Senzhang Wang, and Jianxin Wang. 2024. Fine tuning out-of-vocabulary item recommendation with user sequence imagination. *Advances in Neural Information Processing Systems* 37 (2024), 8930–8955.

[14] Ziluowen Luo, Senzhang Wang, Chaozhuo Li, Jun Yin, Hao Yan, Ming Cheng, Chenxu Wang, Songyang Liu, Litian Zhang, Qiwei Ye, et al. 2026. Distilling Agentic Systems: A Roadmap across Models, Artifacts, and Harnesses. *arXiv preprint arXiv:2609.36630* (2026).

[15] Ziluowen Luo, Jun Yin, Ruochen Liu, Ming Cheng, Shirui Pan, Chengqi Zhang, and Senzhang Wang. 2026. NICE: Scale-Stable Perturbations for Graph Neural Network Explanations via Noise Corruption. *arXiv preprint arXiv:2608.16038* (2026).

[16] Ziluowen Luo, Jun Yin, Hao Yan, Ruochen Liu, Ming Cheng, and Senzhang Wang. 2026. One-Turn Knockout: Traceable and Editable Proxy Unlearning Under Asymmetric Access Constraints. In *Proceedings of the Thirty-Fifth International Joint Conference on Artificial Intelligence, IJCAI-26*. 688–696. Main Track. doi:10.24963/ijcai.2026/78

[17] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. 2019. *Language Models are Unsupervised Multitask Learners*. Technical Report. OpenAI. https://cdn.openai.com/better-language-models/language_models_are_unsupervised_multitask_learners.pdf

[18] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. 2020. Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer. *Journal of Machine Learning Research* 21, 140 (2020), 1–67. https://www.jmlr.org/papers/v21/20-074.html

[19] Senthooran Rajamanoharan, Tom Lieberum, Nicolas Sonnerat, Arthur Conmy, Vikrant Varma, János Kramár, and Neel Nanda. 2024. Jumping Ahead: Improving Reconstruction Fidelity with JumpReLU Sparse Autoencoders. *arXiv preprint arXiv:2407.14435* (2024). arXiv:2407.14435 https://arxiv.org/abs/2407.14435

[20] Shashank Rajput, Nikhil Mehta, Anima Singh, Raghunandan Hulikal Keshavan, Trung Vu, Lukasz Heldt, Lichan Hong, Yi Tay, Vinh Tran, Jonah Samost, Maciej

Kula, Ed Chi, and Maheswaran Sathiamoorthy. 2023. Recommender Systems with Generative Retrieval. In *Advances in Neural Information Processing Systems*, Vol. 36. doi:10.52202/075280-0452

[21] Zihua Si, Zhongxiang Sun, Jiale Chen, Guozhang Chen, Xiaoxue Zang, Kai Zheng, Yang Song, Xiao Zhang, Jun Xu, and Kun Gai. 2024. Generative Retrieval with Semantic Tree-Structured Identifiers and Contrastive Learning. In *Proceedings of the 2024 Annual International ACM SIGIR Conference on Research and Development in Information Retrieval in the Asia Pacific Region*. 154–163. doi:10.1145/3673791.3698408

[22] Juntao Tan, Shuyuan Xu, Wenyue Hua, Yingqiang Ge, Zelong Li, and Yongfeng Zhang. 2024. IDGenRec: LLM-RecSys Alignment with Textual ID Learning. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval*. 355–364. doi:10.1145/3626772.3657821

[23] Adly Templeton, Tom Conerly, Jonathan Marcus, Jack Lindsey, Trenton Bricken, Brian Chen, Adam Pearce, Craig Citro, Emmanuel Ameisen, Andy Jones, Hoagy Cunningham, Nicholas L Turner, Callum McDougall, Monte MacDiarmid, C. Daniel Freeman, Theodore R. Sumers, Edward Rees, Joshua Batson, Adam Jermyn, Shan Carter, Chris Olah, and Tom Henighan. 2024. Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet. *Transformer Circuits Thread* (2024). https://transformer-circuits.pub/2024/scaling-monosemanticity/index.html

[24] Wenjie Wang, Honghui Bao, Xinyu Lin, Jizhi Zhang, Yongqi Li, Fuli Feng, See-Kiong Ng, and Tat-Seng Chua. 2024. Learnable Item Tokenization for Generative Recommendation. In *Proceedings of the 33rd ACM International Conference on Information and Knowledge Management (CIKM '24)*. Association for Computing Machinery. doi:10.1145/3627673.3679569

[25] Jiangnan Xia, Xuansheng Wu, Yu Yang, Xin Wang, and Ninghao Liu. 2026. SAERec: Constructing Fine-grained Interpretable Intents Priors via Sparse Autoencoders for Recommendation. *arXiv preprint arXiv:2606.18897* (2026). arXiv:2606.18897 https://arxiv.org/abs/2606.18897

[26] Linting Xue, Noah Constant, Adam Roberts, Mihir Kale, Rami Al-Rfou, Aditya Siddhant, Aditya Barua, and Colin Raffel. 2021. mT5: A Massively Multilingual Pre-trained Text-to-Text Transformer. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. Association for Computational Linguistics, 483–498. doi:10.18653/v1/2021.naacl-main.41

[27] Jun Yin, Bangguo Zhu, Peng Huo, Ruochen Liu, Hao Chen, Senzhang Wang, Shirui Pan, and Chengqi Zhang. 2026. Echoes in Filter Bubble: Diagnosing and Curing Popularity Bias in Generative Recommenders. *arXiv preprint arXiv:2605.16825* (2026).

[28] Yuanbo Zhao, Ruochen Liu, Senzhang Wang, Jun Yin, Yuxin Dong, Huan Gong, Hao Chen, Shirui Pan, and Chengqi Zhang. 2026. SimGR: Escaping the Pitfalls of Generative Decoding in LLM-based Recommendation. *arXiv preprint arXiv:2602.07847* (2026).

[29] Bowen Zheng, Yupeng Hou, Hongyu Lu, Yu Chen, Wayne Xin Zhao, Ming Chen, and Ji-Rong Wen. 2024. Adapting Large Language Models by Integrating Collaborative Semantics for Recommendation. In *2024 IEEE 40th International Conference on Data Engineering (ICDE)*. 1435–1448. https://arxiv.org/abs/2311.09049

[30] Bowen Zheng, Hongyu Lu, Yu Chen, Wayne Xin Zhao, and Ji-Rong Wen. 2026. Universal item tokenization for transferable generative recommendation. In *Proceedings of the 49th International ACM SIGIR Conference on Research and Development in Information Retrieval*. 2509–2520.

[31] Jieming Zhu, Mengqun Jin, Qijiong Liu, Zexuan Qiu, Zhenhua Dong, and Xiu Li. 2024. CoST: Contrastive Quantization based Semantic Tokenization for Generative Recommendation. In *Proceedings of the 18th ACM Conference on Recommender Systems*. Association for Computing Machinery. doi:10.1145/3640457.3688178

A Experimental Details

A.1 Data and Recommendation Evaluation

We retain the supplied Amazon splits and item orderings, with one validation and one test target per user. Validation uses up to 20 preceding training interactions; test uses up to 20 preceding interactions, including the validation item. The TIGER sampler converts 32 user sequences into at most 512 contiguous next-item windows per micro-batch, each containing up to 29 history items and one target. Evaluation uses unaugmented queries in batches of eight, including the final partial batch.

For target rank r_u, Recall@k and NDCG@k average $\mathbf{1}[r_u \le k]$ and $\mathbf{1}[r_u \le k]/\log_2(r_u + 1)$ over users, respectively, for $k \in \{5, 10\}$; unretrieved targets contribute zero. Evaluation uses item identities, without sampled negatives or seen-item filtering. PR scores the full catalog exactly; AR uses beam search over the catalog trie with width 50. Checkpoints selected by full-validation Recall@5 are reloaded for full-test evaluation.

Table 6: Dataset statistics after pre-processing.

Dataset	Users	Items	Interactions	Train items
Beauty	22,363	12,101	198,502	12,068
Sports	35,598	18,357	296,337	18,284
Toys	19,412	11,924	167,597	11,876

A.2 Pretrained Resources and PID Construction

We extract GPT-2-small [17] with Bloom's SAE through SAELens (`gpt2-small-res-jb`, hook `blocks.8.hook_resid_pre`). Item text concatenates title, brand, category, and price, truncated to 128 tokens. We pool token-level SAE activations excluding special and padding tokens, with decoder-norm scaling as in Eq. (3). Recommendation backbones are randomly initialized.

Each item contributes its top-256 positive features. We estimate $\rho_j = (n_j + 1)/(N_{\text{tr}} + 1)$ from training items, where n_j counts candidate sets containing j and N_{tr} counts training items; compact remapping preserves original feature identities. SPA orders selected features by decreasing activation, breaking ties by index; direct PID-PR uses training-frequency ordering. Ordering is fixed within SPA ablations. A nonsemantic suffix disambiguates shared supports, yielding $\tilde{c}_i = (c_i, \nu_i)$; strength profiles and collision diagnostics exclude it.

History encoding standardizes $\log(1 + a_{ij})$ using featurewise training means μ_j and standard deviations σ_j, smoothed with 16 global pseudo-observations and a standard-deviation floor of 0.05. Values clipped to $[-2, 2]$ interpolate eight learned gain/shift embeddings that modulate feature embeddings before GELU. Four priority-selected anchors incorporate projected mean top-32 context and position embeddings; a fifth encodes the suffix. Layer normalization yields five history tokens per item for every output K. Output strengths normalize selected log activations without history calibration.

A.3 Backbones, Training, and Baselines

Table 7: Recommendation backbones. E/D: encoder/decoder layers. GPT-2 AR and PR are separate capacity cohorts.

Setting	T5 AR/PR	GPT-2 PR	GPT-2 AR
Layers (E/D)	4/4	0/4	0/4
Width / heads	128/6	448/8	128/8
K/V per head	64	56	16
FFN widths	1,024/1,024	1,792	1,024
Activation	ReLU	SiLU	GELU
Dropout	0.15	0.15	0.15

T5 uses relative positions and a three-linear-layer FFN; GPT-2 uses learned positions (512 for PR, history-plus-target capacity for

AR). RPG-style PR baselines flatten SIDs, use position-specific vocabularies and tied input/output embeddings, and apply normalized-dot-product prediction heads at temperature 0.07. T5-SPA-PR uses a noncausal query decoder; GPT-2-SPA-PR appends learned role queries under causal attention. Both predict all roles in one pass without target-token inputs. AR uses shifted teacher forcing and candidate-specific prefixes. The history-only strength prediction is shared across candidates.

The standard recommendation recipe uses seed 45, the gradient accumulation over 16 microbatches, without learning-rate warmup or scheduling. Full validation occurs every 1,600 microbatches; early stopping allows ten unimproved checks, with limits of ten epochs and 320,000 optimizer steps. These limits concern recommendation training, not tokenizer fitting. Equal data protocols do not imply equal parameter counts or early-stopped exposure.

Full PR-SPA uses $(\lambda, \beta) = (1, 8)$ in Eq. (12). The four supervision settings are $(0, 0)$, $(0, 8)$, $(1, 0)$, and $(1, 8)$; all retain activation-aware histories and SID prediction. PR preference loss is positive-item cross-entropy after a temperature-one softmax over all catalog scores. The AR implementation instead uses the positive and four distinct uniform negatives, scoring each under its own prefix; this five-item objective is not exact full-catalog cross-entropy. T5-SPA-AR settings use $(\lambda, \beta) = (0.25, 2)$ on Beauty and $(1, 8)$ on Sports/Toys. At inference, AR normalizes over the full token vocabulary before masking illegal trie edges and applies strengths at complete leaves.

OPQ rotates and product-quantizes mean-pooled PLM residuals; the matched PR recipe uses K subquantizers, 256 centroids each, 25 rotation/PQ iterations, and no additional PCA or whitening. Residual-tokenizer baselines preserve stage-specific code identities: RVQ learns residual codebooks, RKMeans applies residual clustering, and RQ-VAE learns an encoder/decoder around residual quantization. The verified GPT-2 residual-AR assets use four semantic indices plus one suffix. Their original-recipe tokenizers fit all catalog item contents (without held-out interaction labels), with seed 42 and batch size 2,048. RVQ uses per-stage residual normalization, Adagrad at 10^{-3}, commitment weight 0.25, and 3,000 updates; RKMeans uses the same residual normalization, SGD at 0.5, and 30 updates. RQ-VAE uses batch/L2-normalized inputs, a 768–768–256–128–64 encoder and mirrored decoder, no per-stage normalization, and the same Adagrad rate, commitment weight, and update budget as RVQ. Random recommendation codes are distinct from the random-positive-feature selection control used below. The source-transfer recipes in Appendix C.1 form a separate diagnostic cohort, not additional replications of the main-table tokenizer fits.

Baselines. (1) Item tokenization. Random assigns codes independently of item content. OPQ [5] jointly optimizes a rotation and subspace codebooks. RVQ learns successive residual codebooks, whereas R-KMeans clusters successive residuals; we follow their GRID implementations [8]. RQ-VAE [9] learns an encoder and decoder around residual quantization, and HiD-VAE [2] constructs hierarchical, disentangled semantic identifiers. Native P-SAE selects the strongest pretrained features, whereas R-SAE learns a sparse dictionary from item embeddings. The random-positive control samples from the active feature pool and differs from Random item codes. *(2) Recommendation frameworks.* Our tokenizer comparisons use TIGER-style AR [20] and RPG-style PR [5] with T5/GPT-2 backbones: the former predicts SID tokens sequentially, while the latter predicts the code positions in parallel. The complete-framework comparison additionally includes LC-Rec [29], which aligns collaborative semantics with language-model representations; UTGRec [30], which learns a shared tokenizer for transfer across recommendation domains; Ghost [27], which addresses popularity bias through asymmetric unlikelihood optimization and skeleton-founded tokenization. DiGR [28] instead uses SIDs as item representations for direct item-level scoring. GRU4Rec [12] and BERT4Rec [22] provide conventional sequential-recommendation references using recurrent and bidirectional masked-item modeling.

B Proofs and Score Derivatives

B.1 Proof of Proposition 4.1

Proof. Fix a group $\mathcal{G}$ of items sharing the token sequence $c = (c_1, \ldots, c_K)$. Let $\bar{\mathbf{h}}_i = |\mathcal{T}_i|^{-1} \sum_{t \in \mathcal{T}_i} \mathbf{h}_i^{(t)}$ denote the mean PLM representation over the positions used in Eq. (3). We assume nonnegative activations with positive sums $\omega_i = \sum_{r=1}^{K} a_{i,c_r}$, ensuring that $\mathbf{q}_i$ is well-defined.

Write $\mathbf{u}_j = \mathbf{d}_j / \|\mathbf{d}_j\|_2$ for nonzero decoder columns, and set $\mathbf{u}_j = \mathbf{0}$ otherwise. By Eq. (3), the mean SAE reconstruction using only the selected features is

$$\tilde{\mathbf{h}}_i = \mathbf{b}_{\text{dec}} + \sum_{r=1}^{K} \mathbf{u}_{c_r} a_{i,c_r}. \tag{16}$$

The assumptions in the proposition give $\|\bar{\mathbf{h}}_i - \tilde{\mathbf{h}}_i\|_2 \le \epsilon$ for every $i \in \mathcal{G}$, while the sums ω_i lie in an interval of width at most δ.

Consider $n_{\max}$ items sharing the same PID. The case $n_{\max} = 1$ is immediate. For any two distinct items i, i' in this set, $\mathbf{q}_i = \mathbf{q}_{i'}$. Equation (5) therefore implies that

$$\log(1 + a_{i,c_r}) = \alpha \log(1 + a_{i',c_r}) \qquad \text{for all } r$$

for some $\alpha > 0$. Since the log strengths are nonnegative and the logarithm is strictly increasing, all differences $a_{i,c_r} - a_{i',c_r}$ have the same sign. Using $\|\mathbf{u}_{c_r}\|_2 \le 1$ and the triangle inequality, we obtain

$$\begin{aligned} \Delta &\le \|\bar{\mathbf{h}}_i - \bar{\mathbf{h}}_{i'}\|_2 \\ &\le 2\epsilon + \|\tilde{\mathbf{h}}_i - \tilde{\mathbf{h}}_{i'}\|_2 \\ &\le 2\epsilon + \sum_{r=1}^{K} |a_{i,c_r} - a_{i',c_r}| \\ &= 2\epsilon + |\omega_i - \omega_{i'}|. \end{aligned} \tag{17}$$

Thus, any two such items have activation sums separated by at least $\Delta - 2\epsilon > 0$.

Order these sums as $\omega_{(1)} \le \cdots \le \omega_{(n_{\max})}$. Summing the lower bounds on adjacent gaps gives

$$(n_{\max} - 1)(\Delta - 2\epsilon) \le \omega_{(n_{\max})} - \omega_{(1)} \le \delta. \tag{18}$$

Rearranging and using the integrality of $n_{\max}$ yields

$$n_{\max} \le 1 + \left\lfloor \frac{\delta}{\Delta - 2\epsilon} \right\rfloor.$$

In particular, $\delta < \Delta - 2\epsilon$ forces $n_{\max} = 1$, so all items in $\mathcal{G}$ have distinct PIDs. □

B.2 Proof of Proposition 4.2

Proof. Fix a user u, a target item i^+, and a competing item i^-. Let $\hat{\mathbf{q}}_u$ denote SPA's prediction of the next item's relative strengths given $\mathbf{H}_u$, and let $\mathbf{q}_{i^+}$ and $\mathbf{q}_{i^-}$ be the strength vectors recorded in their PIDs.

Write $\mathbf{x} = \sqrt{\hat{\mathbf{q}}_u}$, $\mathbf{y} = \sqrt{\mathbf{q}_{i^+}}$, and $\mathbf{z} = \sqrt{\mathbf{q}_{i^-}}$, with square roots taken elementwise. Since the strength vectors are nonnegative and sum to one, $\|\mathbf{x}\|_2 = \|\mathbf{y}\|_2 = \|\mathbf{z}\|_2 = 1$. The definitions of Δ_q and ε_u therefore give

$$\Delta_q = 1 - \mathbf{y}^\top \mathbf{z} = \frac{1}{2}\|\mathbf{y} - \mathbf{z}\|_2^2, \qquad \varepsilon_u = \frac{1}{2}\|\mathbf{x} - \mathbf{y}\|_2^2. \tag{19}$$

By Eq. (10), the difference between the two strength penalties is

$$D_{\mathrm{H}}(\hat{\mathbf{q}}_u, \mathbf{q}_{i^-}) - D_{\mathrm{H}}(\hat{\mathbf{q}}_u, \mathbf{q}_{i^+}) = \mathbf{x}^\top(\mathbf{y} - \mathbf{z}). \tag{20}$$

Writing $\mathbf{x} = \mathbf{y} + (\mathbf{x} - \mathbf{y})$ and using $\mathbf{y}^\top(\mathbf{y} - \mathbf{z}) = \Delta_q$, the Cauchy–Schwarz inequality yields

$$\mathbf{x}^\top(\mathbf{y} - \mathbf{z}) \geq \Delta_q - \|\mathbf{x} - \mathbf{y}\|_2\|\mathbf{y} - \mathbf{z}\|_2 = \Delta_q - 2\sqrt{\Delta_q \varepsilon_u}. \tag{21}$$

Substituting this bound into Eq. (9) gives

$$s_\theta(u, i^+) - s_\theta(u, i^-) \geq \ell_u(i^+) - \ell_u(i^-) + \beta\big(\Delta_q - 2\sqrt{\Delta_q \varepsilon_u}\big). \tag{22}$$

This inequality holds for general candidate pairs. Under the proposition's assumption $c_{i^+} = c_{i^-}$, their sequence likelihoods are equal, so the bound reduces to Eq. (13).

Finally, $\varepsilon_u < \Delta_q/4$ implies $\Delta_q - 2\sqrt{\Delta_q \varepsilon_u} > 0$. Together with $\beta > 0$, this gives $s_\theta(u, i^+) > s_\theta(u, i^-)$, establishing the claimed ranking. □

C Diagnostic and Interpretation Protocols

C.1 Vocabulary Quality

For a catalog of N items and a suffix-free key $g(i)$, let n_i count items sharing that key. The collision rate is $N^{-1}\sum_i \mathbf{1}[n_i > 1]$. SAE keys are unordered positive-feature supports; quantizer keys retain stage identity. Strengths are excluded, so this statistic measures support ambiguity rather than equality of complete PIDs. Native, priority-based, and random-positive selection share a top-32 candidate pool; random selection samples without replacement.

Content alignment uses 2,000 fixed queries per domain (seed 45), the full catalog excluding each query, and a shared random tie order (seed 46). For representation and reference top-20 neighborhoods $\mathcal{N}_i$ and $\mathcal{N}_i^{\mathrm{ref}}$,

$$\text{Overlap@20} = \frac{1}{20|Q|}\sum_{i \in Q} |\mathcal{N}_i \cap \mathcal{N}_i^{\mathrm{ref}}|. \tag{23}$$

References use cosine similarity of title TF–IDF (lowercase, English stop-word removal, unigrams/bigrams, sublinear TF, smoothed IDF) or mean-pooled PLM residuals. Index-only neighborhoods use support Jaccard; activation-aware neighborhoods use cosine similarity of sparse, decoder-rescaled raw activations, distinct from SPA's normalized log strengths. Permutation moves complete feature–activation profiles across items with content fixed, preserving collision multiplicities.

Reuse diagnostics freeze source-fitted tokenizers across targets, using four semantic tokens and the same 10,000-item subset per target for collision evaluation. R-SAE uses 4,096 features and 3,000 updates; RQ-VAE uses 3,000 updates. RVQ/R-KMeans use globally RMS-scaled residuals without per-stage normalization, taking the first four stages of 32-stage encoders trained for 100 updates per stage. OPQ fits source rotation/codebooks. PID retains the frozen P-SAE dictionary and estimates IDF from target training-item contents. Comparisons match token counts, not bit rates; downstream recommenders are trained on target interactions.

C.2 Semantic Annotations and Steering

Neuronpedia supplies external, automatically generated feature descriptions, unused as model inputs. Title agreement is the lexical-match rate among feature-containing training items; lift divides it by the training-catalog match rate. For f_{4260}, the case-insensitive whole-word rule matches hunt, hunting, rifle, rifles, gun, guns, pistol, shotgun, holster, ammo, ammunition, or scope.

Steering freezes T5-PR Full SPA ($K = 8$) and edits six prespecified features per domain at doses $d \in \{0, \pm 0.25, \pm 0.50\}$. Existing positive history activations are shifted by d in unclipped standardized log coordinates, bounded below by the positive-activation floor, clipped to $[-2, 2]$, and re-encoded; supports, anchors, suffixes, and candidate PIDs stay fixed. Eligible users have the target feature in at least two history items and an anchor responsive in both directions. Within this cohort, three deterministic control features match history frequency, anchors, and activation statistics; per-query scaling matches post-clipping perturbation norms.

Exposure@10 measures the fraction of retrieved items whose PID contains the target feature. Changes from unedited histories are reported in percentage points; Toys equally averages six feature-specific effects. Doses are validation-selected to retain at least 98% Recall@10 on both the full split and matched cohort. Curve bands show min–max over three runs. Paired target-minus-control effects average controls within users; nominal 95% percentile intervals use 500 paired user-bootstrap resamples without correction across 36 feature–direction comparisons. For Sports f_{4260} at $d = 0.50$, the effect is +1.365 points on 332 test users, with interval $[0.652, 2.189]$.

D Additional Results

Vocabulary reuse. Figure 7 extends the source-domain comparison to Sports. Across the three target catalogs, PID has substantially fewer collisions than native P-SAE while using the same frozen dictionary. RQ-VAE attains lower collision rates in this comparison, while PID supports item distinction without extra vocabulary fitting.

Collisions under compact budgets. Figure 8 compares feature selection and quantization at $K \in \{4, 8, 16\}$. At $K = 4$, PID reduces the collision-affected rate from roughly 80% under native selection to 14.64%–21.48% across datasets; the gap narrows as the budget increases. Random-positive selection and RQ-VAE often yield still fewer collisions. Together with the content-alignment results in the main text, this comparison highlights why low collision rates alone are insufficient to assess semantic item representations.

Additional semantic examples. The Beauty feature f_{14730} and Toys feature f_{13542} further connect pretrained descriptions with

Figure 7: Vocabulary reuse from Sports. Collision-affected items (%) are reported; shaded columns mark the source domain used to fit learned tokenizers. PID reuses the frozen P-SAE dictionary with target-catalog IDF.

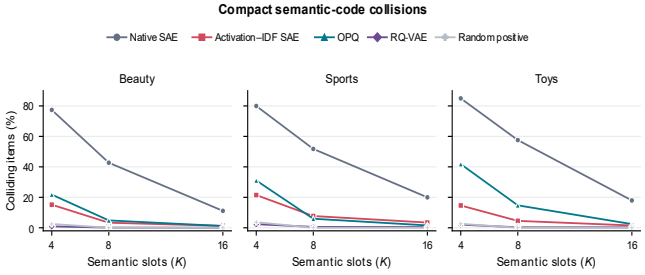

Figure 8: Collision-affected items (%) across semantic-token budgets on Beauty, Sports, and Toys. Activation–IDF SAE denotes PID; random-positive selection samples active SAE features. Strengths and disambiguation suffixes are excluded from collision statistics.

item content. At $K = 4$, title agreement reaches 92.4% for brushes and 99.6% for LEGO, with lifts of 18.31× and 13.77×, respectively. Their top-activating titles exhibit the same patterns, providing complementary evidence for interpreting the selected features.

Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009