arXiv:2610.40249v1 [cs.AI] 29 Oct 2026

Reaction precedents extend retrosynthetic prediction beyond training coverage

Jiaxing Sun[1†], **Xuli Shen**[1†], **Jiang Wu**[1†‡], **Ruiyi Ding**[2†], **Ruikang Zhang**[1,3†], **Xiao An**[1,4†], **Xibao Li**[1†], **Yubin Wang**[1†], **Xingjian Wei**[1†], **Yinfan Wang**[1], **Yu Liu**[5], **Lingli Ge**[6], **Xi Xue**[7,1], **Xiaojian Wang**[7], **Lijun Wu**[1], **Chang-Yu Hsieh**[8], **Yuqiang Li**[1], **Bowen Zhou**[1], **Dahua Lin**[1], **Conghui He**[1] ✉

[1]Shanghai Artificial Intelligence Laboratory, [2]Fudan University, [3]Peking University, [4]Wuhan University, [5]University of Science and Technology of China, [6]Shanghai Jiao Tong University, [7]Institute of Materia Medica, Chinese Academy of Medical Sciences & Peking Union Medical College, [8]Zhejiang University

Reaction precedents guide retrosynthesis, but how far they can compensate for gaps in task-specific training coverage remains unclear. Using 2.71 million reactions, we restricted training and retrieval coverage by transformation or time, then expanded only the retrieval corpus with model parameters fixed. For transformations excluded from training, mean Top-1 accuracy increased from 13.0% to 58.0% on a historical test set and from 4.9% to 23.1% on a temporally held-out test set. For temporal coverage gaps, retrieval expansion enabled a model trained on records before 2017 to match or exceed a non-retrieval model trained on records before 2025 on both test sets. Without task-specific fine-tuning, precedents improved all 14 representative general-purpose models on both test sets and USPTO-50K; the strongest with precedents surpassed all eight representative specialized models. These findings show that inference-time reaction precedents can extend retrosynthetic prediction beyond task-specific training coverage without parameter updates.

Date: October 30, 2026
Equal Contribution (†): Jiaxing Sun, Xuli Shen, Jiang Wu, Ruiyi Ding, Ruikang Zhang, Xiao An, Xibao Li, Yubin Wang, Xingjian Wei
Project Lead (‡): Jiang Wu, wujiang@pjlab.org.cn
Correspondence: Conghui He, heconghui@pjlab.org.cn

DianShi, the romanized form of 点石, evokes the Chinese idiom 点石成金 —literally, "turning stone into gold" —and represents our mission to transform organic-chemistry knowledge scattered across patents, papers, and other literature into structured, AI-ready data and intelligence.

1 Introduction

Scientific prediction draws on both established knowledge and newly available evidence. As AI models are increasingly used for scientific prediction, they should be able to draw on this growing evidence without being retrained for every update[1]. When planning syntheses, chemists complement their acquired expertise by consulting reaction precedents in the literature and databases. This motivates a broader question for scientific AI: to what extent can a trained model expand its predictive scope by drawing on external evidence?

Retrosynthesis—the task of identifying possible precursors to a target molecule—is central to synthesis planning for medicines and materials. Models trained on reaction data can support the search for multistep synthetic routes[4] by proposing precursor candidates for each target molecule[5,6]. Their

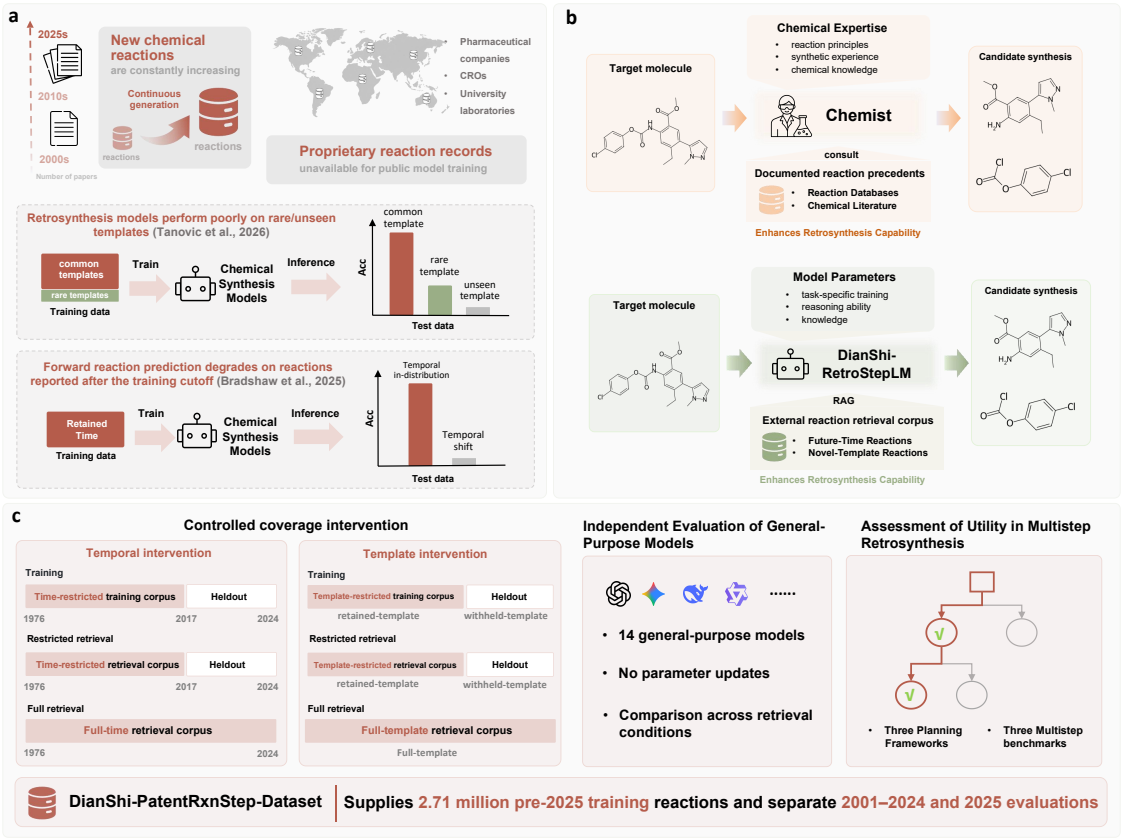

Fig. 1 | Coverage gaps in retrosynthetic prediction and the design of the retrieval-expansion study. **a**, New reactions are continually discovered and reported, so the record available keeps growing while a trained model's coverage stays frozen at its training corpus. A second, independent gap is access: a substantial share of reaction records, held by pharmaceutical companies and contract research organizations, is proprietary and never enters public corpora. Both gaps limit training coverage, and coverage limits prediction: single-step retrosynthesis models are substantially less accurate on rare or absent templates, including models that do not explicitly use templates[2], and forward reaction prediction degrades on reactions from later periods excluded from training[3]. Whether reaction precedents supplied at inference can extend predictive coverage beyond task-specific training without further parameter updates motivates this study. **b**, Reaction knowledge reaches chemists and models in the same two ways. A chemist identifies candidate precursors for a target molecule by combining chemical expertise (reaction principles, synthetic experience and chemical knowledge) with documented reaction precedents drawn from databases and the literature. Likewise, a model can absorb such knowledge into task-specific parameters during training, where it acts together with the model's reasoning ability and stored knowledge, or receive it at inference time as retrieved reaction precedents, which can cover reactions recorded after training or following templates absent from the training corpus. **c**, Study design. Two complementary interventions restrict coverage in both the training and the matching retrieval corpora: a temporal intervention retaining only records before 2017 (Pre2017; 1976–2016 in training), and a template intervention withholding about half of the distinct reaction templates (retained- and withheld-template partitions, S42–S44). Each intervention pairs a restricted retrieval corpus with a full one (full-time, 1976–2024; full-template), while the training corpus remains restricted throughout and evaluation targets are held out. After training, each fixed checkpoint is re-evaluated on the same targets with only the retrieval corpus replaced, keeping model parameters, retrieval procedure and evaluation targets unchanged. Independently, 14 representative general-purpose models, without parameter updates, are compared across a no-precedent condition and three retrieval corpora, and DianShi single-step models are assessed within three planning frameworks (RetroAgent, Retro-R1-modified, DianShi-RetroPlanSystem-v0) on three multistep benchmarks (Pistachio Hard, USPTO-190, PatMol-150). DianShi-PatentRxnStep-Dataset supplies 2.71 million pre-2025 training reactions, with separate 2001–2024 and 2025 evaluation sets.

predictive scope, however, depends on the chemistry represented in training. Evaluations of single-step retrosynthesis models have shown substantially lower accuracy for reactions whose templates are rare or absent in the training data, including in models that do not explicitly use templates[2]. Complementary studies of forward reaction prediction have revealed limitations when models are evaluated on reactions from later periods excluded from training[3]. These coverage gaps have practical relevance: researchers worldwide continually generate new experimental reaction data after a model has been trained, while pharmaceutical companies and contract research organizations (CROs) hold proprietary reaction records that are generally unavailable for public model training. Retrosynthesis therefore provides a concrete setting for scientific AI in which to test whether external evidence can extend a model's predictive scope: training coverage and the evidence supplied at inference can be varied independently, with model parameters held fixed.

Previous studies have demonstrated the value of retrieved reaction information for retrosynthetic prediction[7,8,9,10,11]. Reaction examples have also improved predictions by general-purpose language models without task-specific fine-tuning[12,13]. However, how far reaction precedents can mitigate the predictive limitations associated with gaps in training coverage remains unclear. A complementary question concerns the capabilities of general-purpose models: earlier evaluations reported limited retrosynthetic performance[12], but advances in foundation models motivate reassessing these capabilities and the additional benefit provided by reaction precedents.

To address these questions, we combined controlled retrieval corpus expansion experiments with evaluations of contemporary general-purpose language models. We constructed DianShi-PatentRxnStep-Dataset, comprising 2.71 million training reactions recorded before 2025, with separate evaluation sets for 2001–2024 and 2025. To assess compensation for gaps in task-specific training coverage, we restricted transformation or temporal coverage in both the training corpus and the retrieval corpus. After training, we held model parameters and the retrieval procedure fixed and expanded only the retrieval corpus, comparing predictions on the same targets before and after expansion. We further evaluated 14 representative general-purpose models with and without retrieved reaction precedents, without task-specific fine-tuning, and compared their performance with eight representative released specialized retrosynthesis models (Fig. 1).

With model parameters held fixed, expanding the retrieval corpus improved prediction for transformations absent from the model's fine-tuning data, including those represented by few records in the expanded corpus. It also enabled a model trained on a temporally restricted corpus to match or exceed a non-retrieval model trained on the full corpus, without further parameter updates. Without task-specific fine-tuning, retrieved precedents improved all 14 representative general-purpose models across three test sets, and the strongest general-purpose model with DianShi precedents outperformed all eight representative released specialized models evaluated here. The DianShi single-step models also improved route finding across nine planning settings. Together, these findings show that inference-time access to reaction precedents can extend retrosynthetic prediction beyond task-specific training coverage without further parameter updates. More broadly, they suggest a way for scientific predictors to benefit from an expanding body of external evidence without retraining for every update.

2 Results

2.1 A common data foundation for controlled retrieval corpus expansion

We constructed DianShi-PatentRxnStep-Dataset from DianShi-RxnDB[14], a large-scale, provenance-linked organic reaction-data platform that we previously developed, to study reaction evidence accessed through task-specific training and inference-time retrieval. Following standardization, deduplication and reaction-quality filtering, the dataset comprised 2,712,259 training reactions recorded before 2025 (Train-Full), together with separate validation and test sets. To assess prediction on historical chemistry and under temporal shift, we defined two validation–test pairs: Val-2001-2024 and Test-2001-2024, each

containing 3,798 reactions from 2001 to 2024; and Val-2025 and Test-2025, each containing 1,886 reactions from 2025 and therefore temporally held out from pre-2025 training. The five training, validation and test partitions were disjoint by standardized target product, with additional product-level exclusions against external comparison datasets (Methods; Supplementary Table 1). Train-Full also served as the full retrieval corpus (Supplementary Notes 1 and 2).

We developed two single-step predictors from the Qwen3-4B-Base model[15]. DianShi-RetroStepLM-Direct predicted precursor sets from the target product alone, whereas DianShi-RetroStepLM-RAG, a retrieval-augmented generation (RAG) model, additionally received five retrieved reaction precedents. RetroStepLM-RAG was trained by retrieval-conditioned supervised fine-tuning, in which the target product and its five precedents were supplied together as input, and received precedents in the same format at inference. We refer to these two models hereafter as RetroStepLM-Direct and RetroStepLM-RAG.

To determine whether inference-time evidence could compensate for gaps in task-specific training coverage, we constructed complementary template- and time-restricted training corpora, each with a matching restricted retrieval corpus. The template intervention removed all reactions bearing selected templates from both the task-specific training corpus and the retrieval corpus; the temporal intervention retained only records before 2017 (Pre2017). After training each RetroStepLM-RAG model, we froze its checkpoint and replaced only its restricted retrieval corpus with the full retrieval corpus, keeping all other components and the evaluation targets unchanged. Accuracy changes were then computed on retained- and withheld-template reactions in the template intervention and on Test-2001-2024 and Test-2025 in the temporal intervention. This design isolated the effect of expanding accessible reaction evidence without further parameter updates (Fig. 1; Supplementary Notes 1 and 4).

2.2 Expanding the retrieval corpus extends predictive coverage without further training

We first investigated whether expanding the retrieval corpus could improve prediction for transformations absent from the model's fine-tuning data. We constructed three alternative template partitions using partition seeds 42, 43 and 44, representing different withheld-template assignments rather than independent training replicates. In each partition, approximately half of the distinct local reaction templates represented in Test-2001-2024 and Test-2025 were withheld, and all training reactions bearing a withheld template were removed from both the task-specific training corpus and the retrieval corpus (Methods). The resulting restricted corpora were denoted TemplateRestricted-S42, TemplateRestricted-S43 and TemplateRestricted-S44. With each trained RetroStepLM-RAG checkpoint fixed, replacing only its matching restricted retrieval corpus with the full pre-2025 retrieval corpus improved Top-1 accuracy on withheld-template reactions in every partition (Fig. 2a; Extended Data Fig. 1; Supplementary Note 4). The arithmetic mean of the three partition-level accuracies rose from 13.0% to 58.0% in Test-2001-2024 and from 4.9% to 23.1% in Test-2025. The proportional increase was similar in Test-2001-2024 and Test-2025 (approximately 4.5-fold versus 4.7-fold; partition-level folds of 4.0–4.8 and 4.1–5.7, respectively). In the same paired comparisons, the corresponding mean on retained-template reactions decreased from 69.9% to 68.5% in Test-2001-2024 and from 38.2% to 34.6% in Test-2025, with paired decreases in every partition of 0.6–2.1 and 2.7–4.7 percentage points, respectively. These decreases were modest relative to the gains on withheld-template reactions. Retrieval corpus expansion therefore substantially improved prediction where task-specific training lacked transformation coverage, with similar proportional improvement on historical chemistry in Test-2001-2024 and temporally held-out chemistry in Test-2025.

These benefits were not confined to transformations with abundant precedents. We stratified withheld-template reactions by the frequency of their local reaction template in the full pre-2025 retrieval corpus (1–100, 101–1,000, 1,001–10,000 and $\geq$10,001) and found that corpus expansion improved Top-1 accuracy in all four ranges in both Test-2001-2024 and Test-2025 (Fig. 2b; Extended Data Fig. 2; Supplementary Note 4). For templates represented by 1–100 reactions in the full corpus, mean

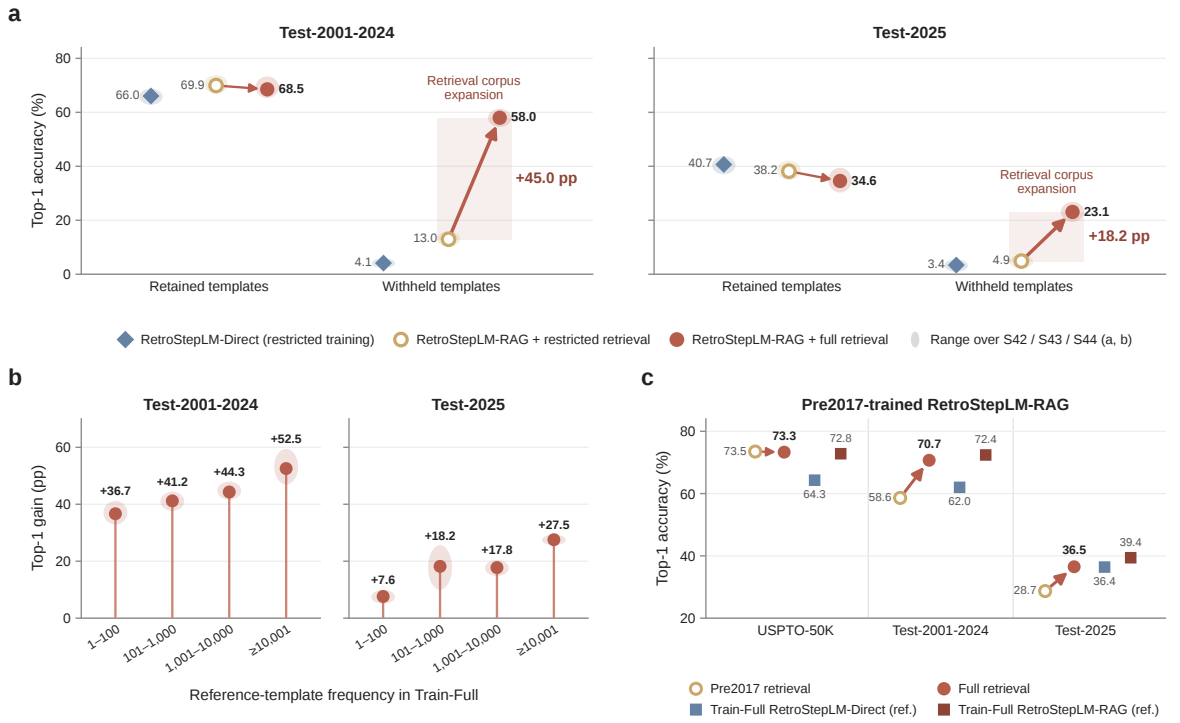

Fig. 2 | Expanding the retrieval corpus extends predictive coverage without further training. **a**, Top-1 accuracy (exact match of the complete predicted precursor set against the recorded reference) on retained- and withheld-template reactions in Test-2001-2024 (left) and Test-2025 (right). All points are arithmetic means over the three template partitions (S42, S43, S44), with light shading spanning the partitions and reactions not pooled across them. The two RetroStepLM-RAG conditions share one fixed checkpoint and differ only in the retrieval corpus; RetroStepLM-Direct is separately trained. **b**, Paired accuracy gain (full minus restricted retrieval; percentage points) on withheld-template reactions across four reference-template frequency ranges in Train-Full (1–100, 101–1,000, 1,001–10,000 and ≥10,001 records). **c**, Top-1 accuracy on USPTO-50K, Test-2001-2024 and Test-2025 (left to right) for a fixed Pre2017-trained RetroStepLM-RAG checkpoint with Pre2017 or full retrieval, together with Train-Full-trained RetroStepLM-Direct and RetroStepLM-RAG as separately trained references. Results and denominators are reported in Supplementary Notes 1 and 4.

partition-level gains were 36.7 percentage points in Test-2001-2024 and 7.6 percentage points in Test-2025 (ranges across partitions: 34.4–39.7 and 6.6–8.4). These results indicate that corpus expansion can benefit transformations with relatively few recorded precedents.

We then examined the complementary temporal intervention: whether access to reaction records beyond the training period could improve prediction without further training. RetroStepLM-RAG was trained on the Pre2017 subset (1.72 million reactions; 63.45% of Train-Full), with precedents retrieved from the same corpus. After training, we fixed the checkpoint and expanded only the retrieval corpus to include the remaining 2017–2024 records in Train-Full. This increased Top-1 accuracy from 58.6% to 70.7% on Test-2001-2024 (+12.1 percentage points) and from 28.7% to 36.5% on Test-2025 (+7.8 percentage points) (Fig. 2c; Extended Data Fig. 3; Supplementary Note 4).

With the expanded retrieval corpus, the Pre2017-trained model exceeded full-corpus RetroStepLM-Direct on Test-2001-2024 (70.7% versus 62.0%) and matched its accuracy on Test-2025 (36.5% versus 36.4%), despite having been trained on only 1.72 million of the 2.71 million reactions. Inference-time access to the additional records thus compensated sufficiently for the temporal training restriction to match or exceed full-corpus RetroStepLM-Direct on both test sets. Even so, accuracy remained below that of

RetroStepLM-RAG trained on Train-Full and evaluated with the full retrieval corpus: 70.7% versus 72.4% on Test-2001-2024 and 36.5% versus 39.4% on Test-2025.

On USPTO-50K[7], which comprises reactions from patents published between 1976 and 2016, retrieval corpus expansion left accuracy essentially unchanged (73.5% versus 73.3%). This limited benefit may reflect the benchmark's temporal coverage, which falls entirely within the Pre2017 training period, whereas expansion added the 2017–2024 records.

2.3 Reaction precedents improve retrosynthesis without task-specific training

We evaluated whether general-purpose models could use reaction precedents without task-specific parameter updates, building on earlier demonstrations that frozen general-purpose models can benefit from retrieved or same-class reaction examples[12,13]. Using a common retrieval-and-prompting framework, we supplied five precedents for each target product to 14 representative general-purpose models and compared their predictions with those generated without retrieved precedents. To assess how the retrieval corpus affected prediction, precedents were drawn from three retrieval corpora: the full DianShi retrieval corpus, supplied by Train-Full; a USPTO retrieval corpus assembled by merging the training sets of USPTO-50K[7], USPTO-MIT[16] and USPTO-Full[17]; and a retrieval corpus constructed from the reactions used to pretrain RXNGraphormer[6], referred to hereafter as the RXNGraphormer corpus (Methods; Supplementary Note 6). Across all 42 model–test-set combinations, Top-1 accuracy was highest when precedents were retrieved from the DianShi corpus, above the USPTO and RXNGraphormer corpora in every case (Fig. 3a; Supplementary Table 15). Precedents from the external corpora were not uniformly better than none: on Test-2025, gemini-3.8-flash reached 18.12% with RXNGraphormer precedents, below its 22.85% without retrieved precedents, whereas DianShi precedents raised accuracy to 35.74%.

The strongest general-purpose model, gemini-3.8-flash with DianShi precedents, achieved Top-1 accuracies of 70.80%, 69.33% and 35.74% on USPTO-50K, Test-2001-2024 and Test-2025, respectively, compared with 42.94%, 27.49% and 22.85% without retrieved precedents. On each test set, gemini-3.8-flash with DianShi precedents exceeded the accuracies of all eight representative released specialized models evaluated here: LocalRetro[18], Graph2Edits[19], R-SMILES[20], EditRetro[21], NAG2G[22], ReactSeq[5], RXNGraphormer[6] and RSGPT[23] (Fig. 3b; Supplementary Note 6). The benefit was not confined to large API-accessed models: Qwen3.5-27B rose from 1.08%, 0.26% and 0.27% without precedents to 36.85%, 36.12% and 13.52% with DianShi precedents. Its accuracy with precedents exceeded that of glm-5.2 without precedents (20.17%, 9.98% and 8.17%), a 744-billion-parameter sparse mixture-of-experts model[24], although with the same precedents glm-5.2 remained higher (53.45%, 52.79% and 21.84%; Supplementary Table 15). Together, these comparisons show that inference-time access to reaction precedents enabled general-purpose models to surpass the evaluated task-specialized predictors without task-specific training.

As a task-specialized reference, the 4B-parameter RetroStepLM-RAG model trained on Train-Full and evaluated with the full retrieval corpus achieved Top-1 accuracies of 72.8%, 72.4% and 39.4% on USPTO-50K, Test-2001-2024 and Test-2025, respectively—the highest among the predictors evaluated here (Supplementary Note 6). It outperformed both full-corpus RetroStepLM-Direct and the strongest general-purpose model supplied with DianShi precedents on all three test sets.

Within a fixed retrieval corpus, the usefulness of precedents also depended on which reactions were selected for each target. We used reference reaction templates to constrain precedent selection in separate diagnostics across four of the 14 representative general-purpose models and the RetroStepLM-RAG model trained on Train-Full. Matching-template exclusion reduced accuracy on all three test sets relative to standard retrieval. On Test-2025, matching-template exclusion fell below the no-precedent baseline for gemini-3.8-flash (19.35% versus 22.85%) and deepseek-v4.1-flash (8.43% versus 15.85%). Reference-template restriction improved prediction for all five models across the three test sets (Extended Data Fig. 4; Supplementary Note 5). These results show that product-structure similarity alone does not ensure useful reaction precedents, so whether an expanded corpus helps also depends on how precedents

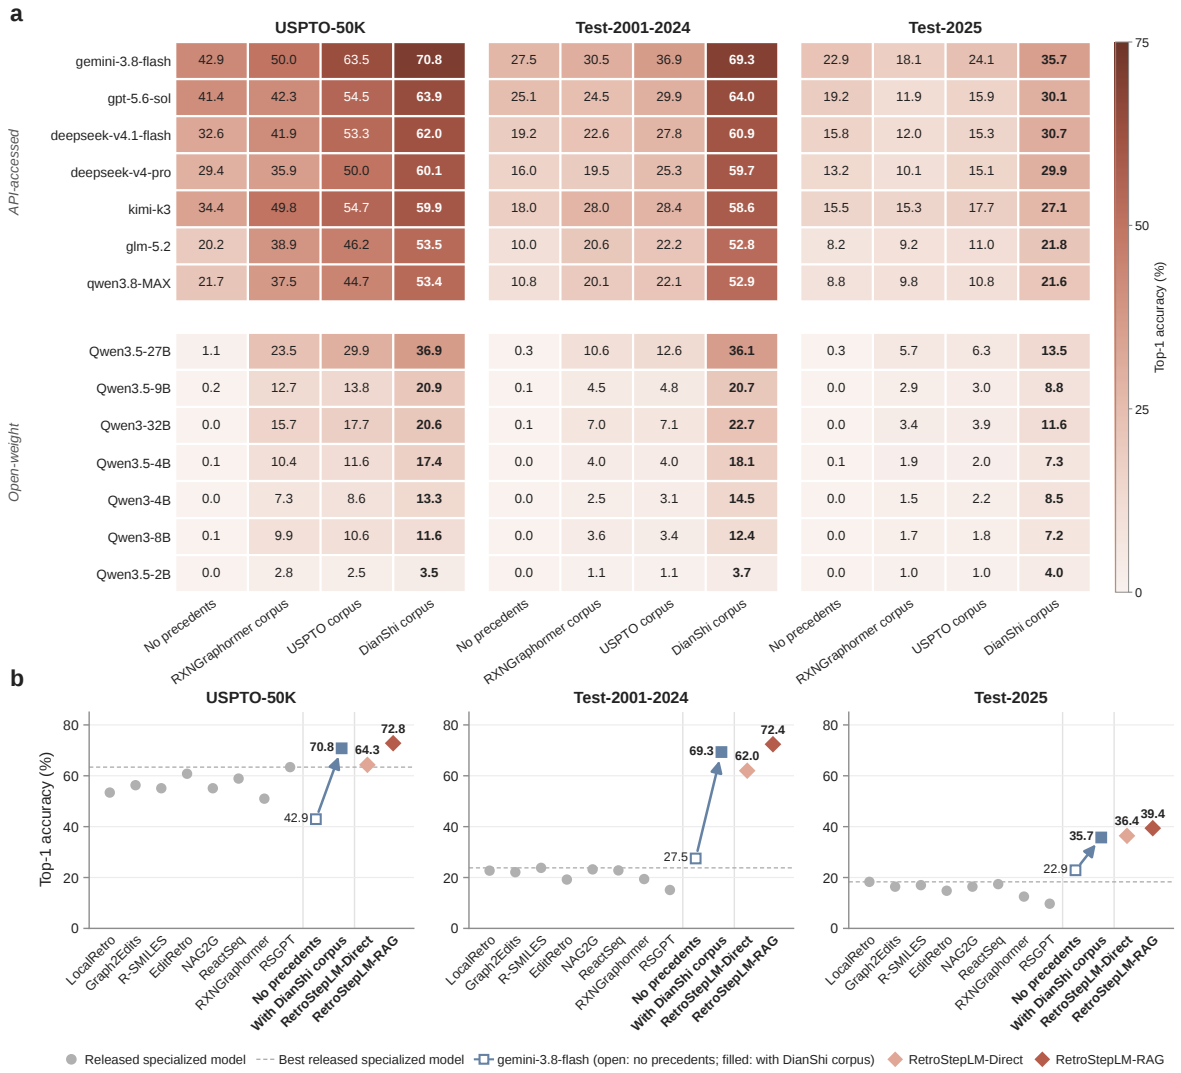

Fig. 3 | Reaction precedents improve retrosynthesis without task-specific training. **a**, Top-1 accuracy of 14 representative general-purpose models on USPTO-50K, Test-2001-2024 and Test-2025 (left to right). Within each test set, columns compare no precedents with five precedents retrieved from the RXNGraphormer, USPTO or DianShi corpus. The DianShi corpus is derived from Train-Full, from which products in the DianShi evaluation partitions are excluded; the USPTO corpus merges the training sets of USPTO-50K, USPTO-MIT and USPTO-Full; the RXNGraphormer corpus is constructed from the reactions used to pretrain RXNGraphormer. **b**, Accuracy of eight representative released specialized models alongside gemini-3.8-flash without precedents (open symbol) or with the DianShi corpus (filled symbol), and of RetroStepLM-Direct and RetroStepLM-RAG trained on Train-Full, on USPTO-50K, Test-2001-2024 and Test-2025 (left to right). Results are reported in Supplementary Note 6.

are selected for each target.

Together, these results show that reaction precedents can support strong retrosynthetic prediction without task-specific training, while the highest accuracy was achieved by combining retrieval-conditioned fine-tuning with inference-time access to reaction precedents. Task-specific training and inference-time evidence thus provided complementary benefits, with the utility of that evidence depending on precedent selection.

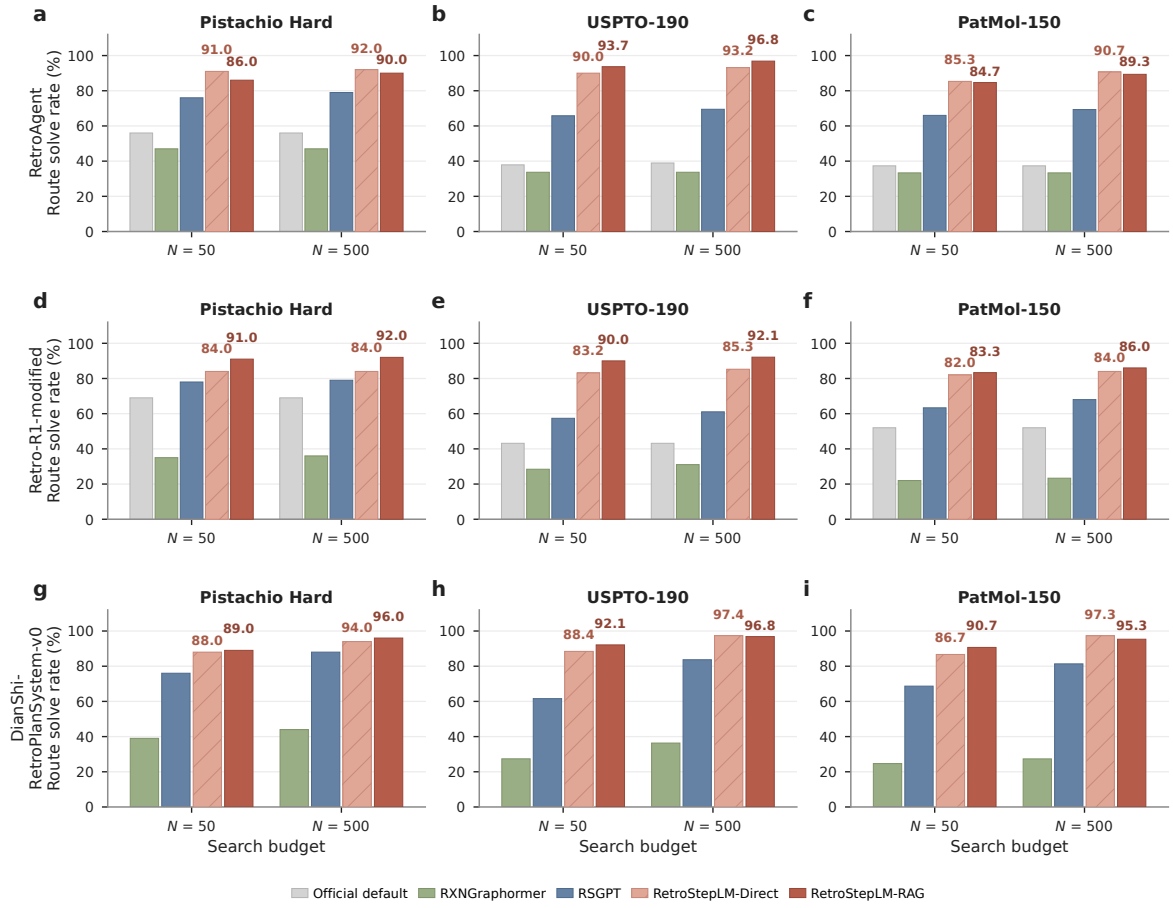

Fig. 4 | DianShi single-step models improve solve rates across planning frameworks. **a–i**, Route solve rates within RetroAgent (**a–c**), Retro-R1-modified (**d–f**) and DianShi-RetroPlanSystem-v0 (**g–i**); columns show Pistachio Hard (100 targets), USPTO-190 (190) and PatMol-150 (150), respectively. Solve rate counts targets for which at least one complete route to the inventory was found. Within each framework, planner, starting-material inventory and search protocol are fixed; the complete budget trajectories are shown in Extended Data Fig. 5. Protocols and complete results are reported in Supplementary Note 7.

2.4 DianShi single-step models improve solve rates across planning frameworks

We next evaluated whether the DianShi single-step models improve downstream route finding. We used three planning frameworks: RetroAgent[25], an LLM-based agentic planner; Retro-R1-modified[26], an adapted Retro-R1 pipeline; and our own planning framework, DianShi-RetroPlanSystem-v0. Each framework was evaluated on three benchmarks: Pistachio Hard[27] (100 targets; evaluated here without the original specified-starting-material constraint), USPTO-190[28] (190 targets) and PatMol-150, a benchmark we constructed from 2025 public patents (150 targets, reference routes of 7–13 steps). Within each framework, we compared single-step models under a fixed planner, starting-material inventory and search protocol, with all models using the same template-based reaction validation[29] and their respective template libraries. We measured solve rate as the fraction of targets for which a complete route to the starting-material inventory was found within the search budget (Fig. 4; Extended Data Fig. 5; Supplementary Note 7).

Both DianShi models achieved higher solve rates than every evaluated non-DianShi model across all nine planner–benchmark settings, at both N=50 and N=500. At N=500, their solve rates ranged from

84.00% to 97.37%, exceeding the strongest non-DianShi model in each setting by 5.00–31.06 percentage points.

RetroStepLM-RAG performed better overall in multistep planning, achieving higher solve rates than RetroStepLM-Direct in seven of the nine settings at N=50 and in five of the nine at N=500, with advantages of 1.00–7.00 and 2.00–8.00 percentage points, respectively. RetroStepLM-Direct achieved higher solve rates in the remaining two and four settings, with advantages of 0.66–5.00 and 0.53–2.00 percentage points. These budget-dependent differences are consistent with the higher single-step Top-1 accuracy of RetroStepLM-RAG on all three test sets, but they also indicate that candidate-set composition and search dynamics influence multistep performance[30,31]. Together, these results show that stronger single-step models can substantially improve multistep planning, while single-step accuracy alone does not determine multistep utility.

3 Discussion

Our controlled interventions with a language model trained through retrieval-conditioned fine-tuning show that retrosynthetic prediction can extend beyond task-specific training coverage. With model parameters held fixed, expanding the inference-time retrieval corpus improved prediction for transformations absent from the model's fine-tuning data. One interpretation is that retrieval-conditioned fine-tuning equips the model not only with knowledge of particular transformations, but also with an ability to use reaction precedents supplied at inference. External evidence can then provide information needed for prediction without requiring each additional transformation to be incorporated through further parameter training[1]. Benefits extended to transformations with few recorded precedents and to records beyond the training period. These two gaps parallel the shift types that prior work constructs as separate out-of-distribution splits, in which either transformation types or target molecules are held out[32]; here the restrictions acted on corpus coverage rather than on evaluation splits, so that the recovered accuracy could be attributed to retrieved evidence alone. Inference-time access to external evidence therefore provides a means of extending the useful scope of a trained predictor. This offers a route for incorporating newly reported reactions and proprietary records held by pharmaceutical companies and CROs into prediction without updating model parameters.

The extent of these gains, however, also depended on which precedents were selected for each target. Structurally similar products can arise from different transformations, so product similarity does not guarantee that the selected precedents share the transformation of the recorded reaction. Diagnostics that use template information from the recorded answer isolate this dependence: restricting precedent selection to matching-template reactions increased reference-set accuracy on all three test sets, whereas excluding such reactions reduced it in every model and test set, and on Test-2025 fell below the no-precedent reference for two general-purpose models. Precedent selection therefore influences agreement with the recorded reaction even when model parameters and the retrieval corpus remain fixed. The expansion did not raise accuracy uniformly either: withheld-template reactions improved substantially, whereas retained-template reactions, whose transformations were already represented in the fine-tuning data, decreased modestly. Whether an expanded corpus improves accuracy therefore depends on whether the additional records contain the transformation a target requires, and on how those records enter the precedent set for that target.

The general-purpose model results show that contemporary foundation models can achieve strong retrosynthetic performance with reaction precedents and without additional task-specific fine-tuning. The strongest general-purpose model with DianShi precedents outperformed all eight representative released specialized models evaluated across the three test sets. These findings suggest that leading foundation models possess substantial chemical knowledge and reasoning capabilities that can be brought to bear on retrosynthesis through access to relevant reaction examples. Performance also depended on the corpus from which precedents were retrieved, highlighting the importance of external evidence alongside model capability. Related work has advanced retrosynthesis mainly by internalizing

reaction knowledge in model parameters—through domain-specific pretraining and supervised training on reaction data[5,6,23,33], or through reinforcement-learning and reasoning-oriented post-training[34,35,36]. Extending either route requires further parameter training. Reaction precedents supplied at inference, by contrast, extended coverage in the present experiments with the checkpoint unchanged, and can complement both approaches for chemical language models. At the same time, the strongest predictor evaluated here combined retrieval-conditioned fine-tuning with inference-time access to reaction precedents, supporting the complementary roles of parameter adaptation and inference-time evidence. This combination parallels how chemists draw on acquired expertise and consult the literature and reaction databases for precedents relevant to a target molecule. External reaction evidence can therefore complement learned capabilities without requiring every additional reaction to be encoded in model parameters.

The planning results demonstrate the practical value of improved single-step models for multistep synthesis planning. Both DianShi models increased solve rates across all nine planner–benchmark settings, showing that their downstream benefits extended across different planning frameworks and target collections. These consistent gains underscore the importance of single-step candidate generation to route finding: the reactions proposed at each expansion determine which synthetic pathways the planner can explore. These gains were not fully captured by single-step Top-1 accuracy. RetroStepLM-RAG performed better overall in multistep planning, consistent with its higher single-step accuracy, but its advantage did not extend to every planning setting. This variation suggests that the utility of candidate reactions during search matters alongside agreement with recorded precursors[30,31]. Exact matching against one recorded precursor set does not capture all chemically plausible alternatives[37,38]. Ultimately, computational route finding must be complemented by experimental validation of the proposed reactions and their conditions[39].

4 Methods

4.1 Reaction dataset construction and partitioning

DianShi-PatentRxnStep-Dataset was derived from DianShi-RxnDB, a provenance-linked organic reaction-data platform that we previously developed[14]. DianShi-RxnDB uses a fully automated pipeline to extract structured reaction records from patent text, images and reaction schemes, primarily from USPTO and EPO organic-synthesis patents published between 1976 and 2025. Each record represents a reported single-step experiment and links to its source patent. We curated these records for single-step retrosynthesis through reaction standardization, deduplication, quality filtering and product-disjoint partitioning (Supplementary Note 2).

Reactant and product SMILES were standardized using RDKit[40]. Reaction-level duplicates were identified using a standardized reactant-set–product-set signature, with atom-mapping numbers removed from the comparison representation and molecular components sorted deterministically. For duplicate records spanning multiple years, the record with the earliest year recorded in DianShi-RxnDB was retained. Quality filtering followed the curation conventions used for single-step retrosynthesis datasets[2] and excluded reactions with multiple products, repeated molecules on either side, structural overlap between reactants and products, more than four reactant components, inorganic products, inconsistent atom mappings or unsuccessful reaction-centre or template extraction. Local reaction templates[18] were grouped across the quality-controlled corpus, and templates represented by fewer than six reactions were removed before temporal partitioning, following the cutoff used in previous template-frequency analyses[2]. Detailed filtering and partitioning procedures are provided in Supplementary Note 2.

For temporal evaluation[3,8], partitions were defined using the year of the source patent recorded in DianShi-RxnDB. Records before 2025 supplied the training and historical evaluation pools, whereas records from 2025 supplied the temporally held-out evaluation pool. Val-2001-2024 and Test-2001-2024 were sampled from 2001–2024 records with stratification by year and template frequency in the full

training corpus; Val-2025 and Test-2025 were sampled from the 2025 pool with stratification by patent type and template frequency. The five partitions—Train-Full, Val-2001-2024, Test-2001-2024, Val-2025 and Test-2025—were constructed jointly and were mutually disjoint by standardized target product[2]. To control overlap with external benchmarks, training candidates were screened against products in the validation and test partitions of USPTO-50K[7], USPTO-Full[17] and USPTO-MIT[16]. DianShi validation and test candidates were additionally screened against products in the training partitions of these USPTO datasets and of the corpus used to pretrain RXNGraphormer[6].

The resulting dataset comprised 2,712,259 pre-2025 training reactions in Train-Full, 3,798 reactions each in Val-2001-2024 and Test-2001-2024, and 1,886 reactions each in Val-2025 and Test-2025. Train-Full also served as the full retrieval corpus. Dataset identifiers and partition roles are summarized in Supplementary Note 1.

4.2 Fixed-parameter retrieval corpus expansion experiments

Training coverage can be incomplete in two distinct ways: the transformations a target requires may be absent from the fine-tuning data, or the available data may predate the chemistry of interest. We therefore imposed two complementary coverage restrictions—one on transformations, one on time—and examined whether reaction precedents retrieved at inference time can compensate for these gaps without further parameter updates.

In every experiment, DianShi-RetroStepLM-RAG was trained on a restricted training corpus, with precedents drawn from the matching restricted retrieval corpus, and the resulting checkpoint was then fixed. Each evaluation target was scored with either the restricted retrieval corpus or the full retrieval corpus. The full retrieval corpus comprised the 2,712,259 reactions in Train-Full, and each restricted retrieval corpus comprised the reactions of its corresponding restricted training corpus. Paired comparisons therefore switched only between the two retrieval indices, keeping the checkpoint, retrieval procedure, precedent serialization, prompting, decoding and scoring unchanged. Retrieval and training procedures are described in Reaction corpora and precedent retrieval and DianShi single-step models and training, respectively.

For the template-coverage intervention, we constructed three alternative template partitions using template-partition seeds 42, 43 and 44. Within each partition, approximately half of the distinct reaction templates represented in each of Test-2001-2024 and Test-2025 were designated as withheld. Templates shared between the two test sets received a common retain-or-withhold assignment. All reactions bearing withheld templates were removed from Train-Full, and the remaining records formed the corresponding TemplateRestricted-S42, TemplateRestricted-S43 or TemplateRestricted-S44 corpus, which served as both the restricted training corpus and the matching restricted retrieval corpus. Withheld templates were therefore absent from both restricted resources, whereas retained evaluation templates had at least one corresponding training reaction. The approximately 50% withholding target referred to distinct templates rather than reaction counts. Test-2001-2024 and Test-2025 were each divided into retained- and withheld-template subsets according to the corresponding partition. The three partitions represented alternative withheld-template assignments rather than independent training replicates.

To assess whether the effects of retrieval corpus expansion extended to transformations with few recorded precedents, we stratified withheld-template evaluation reactions by the number of records bearing their local reaction template[2,18] in Train-Full, which also constituted the full retrieval corpus. The four frequency ranges were 1–100, 101–1,000, 1,001–10,000 and at least 10,001 reactions per template.

For the temporal intervention, the Pre2017 training corpus retained the Train-Full reactions before 2017. RetroStepLM-RAG was trained on this corpus with precedents retrieved from the matching Pre2017 retrieval corpus and, after training, evaluated with either the Pre2017 retrieval corpus or the full retrieval corpus.

RetroStepLM-Direct models trained on the corresponding restricted corpora and RetroStepLM-Direct and RetroStepLM-RAG models trained on Train-Full provided reference configurations. Configuration correspondence is provided in Supplementary Note 1; restricted-corpus sizes and template-partition composition are reported in Supplementary Note 4. Accuracy calculation and cross-partition reporting are described in Single-step evaluation.

4.3 Reaction corpora and precedent retrieval

Retrieval indices were constructed from Train-Full and each restricted training corpus defined above, without additional corpus-level subsampling. Two further retrieval corpora were constructed for the general-purpose model evaluations: a USPTO retrieval corpus assembled by merging the training sets of USPTO-50K[7], USPTO-MIT[16] and USPTO-Full[17], and an RXNGraphormer retrieval corpus built from the RXNGraphormer training corpus[6]. Retrieval from all corpora used the same procedure and index settings, so that differences between corpora reflected their reaction content rather than the retrieval implementation. External retrieval corpora are summarized in Supplementary Note 6.

Standard retrieval used a two-stage structural-similarity search based on the target product alone, without reference reaction-class or template information. Products were represented by extended-connectivity fingerprints (ECFPs)[41]. A binary flat index returned 200 candidate reactions ranked by Hamming distance, which were then reranked by the Tanimoto similarity between their product fingerprints and the query fingerprint. Reference-informed modifications to precedent selection are described in Precedent-relevance diagnostics.

During RetroStepLM-RAG fine-tuning, the query reaction, duplicate reactions and potentially answer-disclosing records from the same patent were excluded from candidate precedents. At each training epoch, five eligible precedents were randomly selected from the ten highest-similarity candidates. At standard inference, RetroStepLM-RAG and the general-purpose models received the five highest-ranked eligible precedents deterministically. Product-level exclusions during dataset construction prevented products in the DianShi evaluation partitions and the USPTO-series validation and test partitions from appearing in the DianShi training corpus or its derived retrieval corpora.

Each precedent contained standardized product and reactant SMILES, together with atom-mapping-derived synthon SMILES and corresponding leaving groups. Five precedents were supplied alongside the target product as reference context for RetroStepLM-RAG and the general-purpose models, using the structured prompt family provided in Supplementary Note 3.

4.4 DianShi single-step models and training

DianShi-RetroStepLM-Direct and DianShi-RetroStepLM-RAG were developed from Qwen3-4B-Base[15] using full-parameter supervised fine-tuning. RetroStepLM-Direct received the target product alone, whereas RetroStepLM-RAG underwent retrieval-conditioned supervised fine-tuning, receiving the target product together with five reaction precedents selected using the retrieval procedure described in Reaction corpora and precedent retrieval. Both models were trained to generate the same structured target sequence.

Synthon- and leaving-group-based decompositions have also been used in structured retrosynthesis models[5,42]. Training used intermediate-label supervision, with synthon and leaving-group labels derived from atom-mapped reactions. Synthons represented product fragments associated with the retrosynthetic disconnection, and corresponding leaving groups provided information for reconstructing the reactants. Target sequences serialized synthon SMILES, leaving groups and the final standardized reactant set in the order `synthons_smiles`, `leaving_groups` and `reactants`, with the final reactant SMILES additionally enclosed in `<answer>` tags. For input x and target sequence $y = (y_1, \ldots, y_T)$, training minimized the autoregressive negative log-likelihood over the complete target sequence,

$$\mathcal{L}(\theta) = -\sum_{t=1}^{T} \log p_\theta(y_t \mid x, y_{<t}),$$

where x comprised the target product alone for RetroStepLM-Direct and the target product together with retrieved precedents for RetroStepLM-RAG.

Full-corpus models used Train-Full, whereas coverage-intervention models used the corresponding Pre2017, TemplateRestricted-S42, TemplateRestricted-S43 or TemplateRestricted-S44 training corpus. For each RetroStepLM-RAG model, training precedents were retrieved from the same record set that supplied its fine-tuning reactions. SMILES were space-segmented into atom- and bond-level tokens following the Molecular Transformer tokenization scheme[43] and encoded using the native Qwen tokenizer without vocabulary expansion. Models were trained for ten epochs with an effective batch size of 128 and a learning rate of 1×10^{-4}, using cosine decay and linear warm-up over the first 5% of training steps. Training used bfloat16 precision; hardware and distributed-training settings are provided in Supplementary Table 4 (Supplementary Note 3).

At inference, both models used beam search to generate ranked precursor-set candidates for Top-1, Top-3, Top-5 and Top-10 evaluation. Only the final reactant set was used for precursor-set scoring, as described in Single-step evaluation. Decoding and candidate-extraction procedures are described in Supplementary Note 3.

4.5 Precedent-relevance diagnostics

We conducted reference-informed retrieval diagnostics to examine how precedent selection affected retrosynthetic prediction. Evaluations used the RetroStepLM-RAG model trained on Train-Full and four of the 14 general-purpose models described below: gemini-3.8-flash, deepseek-v4.1-flash, Qwen3-32B and Qwen3.5-27B. All five models were evaluated on USPTO-50K, Test-2001-2024 and Test-2025. Across precedent-selection conditions, each model used the full DianShi retrieval corpus and received five precedents per query. Model parameters, evaluation targets, prompt structure, decoding and scoring settings were held fixed; only the precedent-selection rule changed.

Standard retrieval, which ranks precedents by target-product structural similarity, was compared with three reference-informed conditions. Under reference-class restriction, retrieval selected only precedents sharing the recorded reference reaction class; under reference-template restriction, it selected only those sharing the reference reaction template. Matching-template exclusion removed reactions bearing the reference template from precedent selection. In the matching-template exclusion condition, retrieval continued down the ranked candidate list to collect five eligible precedents; when fewer than five remained, eligible precedents were cycled to fill the five context positions. Predictions without retrieved precedents provided additional reference conditions for the four general-purpose models. Performance was assessed at Top-1 using the complete precursor-set exact-match metric defined in Single-step evaluation.

Reaction classes and templates were obtained from the recorded reference reactions, making these privileged diagnostics rather than deployable retrieval settings. Diagnostic conditions and results are provided in Supplementary Note 5.

4.6 General-purpose model evaluation and specialized-model comparisons

We evaluated 14 representative general-purpose models with and without retrieved reaction precedents, without task-specific parameter updates, and compared their performance with DianShi models and eight representative released specialized models. These 14 models comprised seven API-accessed models across major providers and seven open-weight Qwen models spanning 2B to 32B parameters. Each model was evaluated on USPTO-50K, Test-2001-2024 and Test-2025 under four conditions: no retrieved

precedents, or five precedents retrieved from the RXNGraphormer retrieval corpus, the USPTO retrieval corpus or the full DianShi retrieval corpus. Model identities are provided in Supplementary Note 6.

General-purpose models produced one final precursor-set prediction per target and were assessed at Top-1 using the common scoring procedure described in Single-step evaluation. API requests used temperature 0, top-p 1.0, a maximum output-token setting of 65,535 and provider-default thinking settings. Incomplete API responses were rerun until a complete response was obtained. General-purpose model prompting and API request settings are described in Supplementary Notes 3 and 6.

For specialized-model comparisons, we evaluated the released weights of eight representative specialized retrosynthesis models covering the main published single-step approaches: LocalRetro[18], Graph2Edits[19], R-SMILES[20], EditRetro[21], NAG2G[22], ReactSeq[5], RXNGraphormer[6] and RSGPT[23], without retraining on DianShi data. Inputs were formatted according to each method's requirements, and Test-2001-2024 and Test-2025 inference used no molecular test-time augmentation. Predictions were assessed using the common precursor-set matching metric described in Single-step evaluation. Specialized-model evaluation protocols and complete comparison results are provided in Supplementary Table 17 (Supplementary Note 6).

4.7 Single-step evaluation

Single-step prediction was evaluated by exact matching of complete precursor sets against the recorded reference. Final precursor-set predictions were extracted from each model's output according to its output format. Predicted and reference molecules were standardized using RDKit and converted to canonical SMILES, with molecular components sorted deterministically before comparison[44]. Component order therefore did not affect correctness. Missing predictions, invalid molecular strings and partial precursor-set matches were counted as incorrect.

For specialized models, Top-k accuracy measured whether at least one of the first k ranked candidates matched the complete reference precursor set. For M evaluation reactions, reference set R_i and candidate $\hat{R}_{ij}$ at rank j, the metric was

$$\mathrm{EM@}k = \frac{1}{M}\sum_{i=1}^{M}\mathbf{1}\left[\exists j \le k : \hat{R}_{ij} = R_i\right], \qquad k \in \{1,3,5,10\}.$$

Candidates were evaluated in their generated rank order. General-purpose models produced one scored precursor set per target and were assessed at Top-1. Accuracy measured agreement with the recorded reference precursor set.

For template-coverage analyses, cross-partition summaries were arithmetic means of the three template-partition-specific accuracies, calculated separately for each test set, retained- or withheld-template subset and, where applicable, precedent-frequency stratum. Evaluation reactions were not pooled across template partitions as independent samples. Reported mean gains were arithmetic means of the three within-partition accuracy differences. Template-subset definitions and reporting conventions are provided in Supplementary Note 1.

4.8 Multistep planning and evaluation

Multistep evaluation used Pistachio Hard[27] (100 targets), USPTO-190[28] (190 targets) and PatMol-150 (150 targets). Pistachio Hard was evaluated without its original specified-starting-material constraint. PatMol-150 was constructed from 2025 public patents in DianShi-RxnDB and contained reference routes of 7–13 steps. After molecular standardization, duplicate targets were removed, and targets were excluded if they appeared among products in Train-Full, targets in either public benchmark or molecules in the starting-material inventory. Benchmark details and the PatMol-150 construction procedure are provided in Supplementary Note 7.

We evaluated three planning frameworks: RetroAgent[25], Retro-R1-modified[26] and DianShi-RetroPlanSystem-v0. All three used Qwen3.5-397B as the planner model and a common RDKit-standardized inventory of 23,081,629 commercially available eMolecules building blocks[45]. Within each framework, the planner model, inventory, Top-10 single-step candidate setting and search protocol were held fixed across single-step model comparisons. Retro-R1-modified adapted the original inference protocol for a general-purpose planner model. DianShi-RetroPlanSystem-v0 represented search as an AND–OR tree, in which a molecule node closed when at least one candidate reaction closed and a reaction node closed only when all required precursor molecules were available or solved. Protocol details, including the tool interface, closure propagation and cycle handling, are provided in Supplementary Note 7.

All evaluations applied a common AOT*-style template-based reaction-validation procedure[29] after single-step prediction, with each model's corresponding template library and rdchiral[46] for template application (Train-Full for the DianShi models; USPTO-derived libraries for RSGPT, RXNGraphormer and the default planning models). Candidates rewritten to regenerate the target product, or any of its components, were discarded. Template-library sources, matching thresholds and fallback settings are provided in Supplementary Note 7.

Performance was assessed as target-level solve rate, defined as the fraction of targets for which at least one complete route, with all terminal precursors in the starting-material inventory, was found within the search budget. Multiple routes for one target counted as one success. Budget N counted newly expanded distinct molecules in RetroAgent, single-step tool calls in Retro-R1-modified, and expanded molecule nodes in DianShi-RetroPlanSystem-v0, where each expansion invoked the single-step model once. Results at $N = 50, 100, 200, 300, 400$ and 500 were budget levels along the same search trajectory. Single-step model comparisons were performed within each framework. Solve-rate trajectories are provided in Supplementary Note 7 and Extended Data Fig. 5.

5 Data availability

The DianShi-curated datasets and derived resources are available at https://github.com/DianShiGroup/DianShi-RetroStepLM and https://huggingface.co/datasets/DianShiGroup/DianShi-RetroPlanSystem-Assets under a Creative Commons Attribution 4.0 International License (CC BY 4.0; https://creativecommons.org/licenses/by/4.0/). They comprise the evaluation partitions DianShi-PatentRxnStep-Test-2001-2024 (3,798 reactions), DianShi-PatentRxnStep-Test-2025 (1,886 reactions), DianShi-PatentRxnStep-Val-2001-2024 (3,798 reactions) and DianShi-PatentRxnStep-Val-2025 (1,886 reactions), the PatMol-150 multistep benchmark (150 targets), the Train-Full template library (the library used for reaction validation) and the template-coverage cutoff scheme (the per-partition retained and withheld template assignments for seeds 42, 43 and 44, together with the generation procedure).

6 Code availability

The DianShi-RetroStepLM-Direct checkpoint (https://huggingface.co/DianShiGroup/DianShi-RetroStepLM-Direct) and the DianShi-RetroPlanSystem-v0 multistep search framework (https://github.com/DianShiGroup/DianShi-RetroStepLM) are available under the Apache License 2.0 (https://www.apache.org/licenses/LICENSE-2.0).

7 Acknowledgments

This project was supported by Shanghai Artificial Intelligence Laboratory.

References

1. Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, volume 33, pages 9459–9474, 2020.

2. S. Tanovic, E. Wieczorek, and F. Duarte. An exploration of dataset bias in single-step retrosynthesis prediction. *Digital Discovery*, 5:793–802, 2026. doi: 10.1039/d5dd00358j.

3. J. Bradshaw, A. Zhang, B. Mahjour, D. E. Graff, M. H. S. Segler, and C. W. Coley. Challenging reaction prediction models to generalize to novel chemistry. *ACS Central Science*, 11:539–549, 2025. doi: 10.1021/acscentsci.5c00055.

4. Marwin H. S. Segler, Mike Preuss, and Mark P. Waller. Planning chemical syntheses with deep neural networks and symbolic AI. *Nature*, 555:604–610, 2018.

5. J. Xiong et al. Bridging chemistry and artificial intelligence by a reaction description language. *Nature Machine Intelligence*, 7:782–793, 2025. doi: 10.1038/s42256-025-01032-8.

6. L.-C. Xu, M.-J. Tang, J. An, F. Cao, and Y. Qi. A unified pre-trained deep learning framework for cross-task reaction performance prediction and synthesis planning. *Nature Machine Intelligence*, 7:1561–1571, 2025. doi: 10.1038/s42256-025-01098-4.

7. Connor W. Coley, Luke Rogers, William H. Green, and Klavs F. Jensen. Computer-assisted retrosynthesis based on molecular similarity. *ACS Central Science*, 3:1237–1245, 2017. doi: 10.1021/acscentsci.7b00355.

8. Y. Qian, Z. Li, Z. Tu, C. W. Coley, and R. Barzilay. Predictive chemistry augmented with text retrieval. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, 2023. doi: 10.18653/v1/2023.emnlp-main.784.

9. L. Fang, J. Li, M. Zhao, L. Tan, and J.-G. Lou. Single-step retrosynthesis prediction by leveraging commonly preserved substructures. *Nature Communications*, 14:2446, 2023. doi: 10.1038/s41467-023-37969-w.

10. A. Qiao, Z. Wang, J. Rao, Y. Yang, and Z. Wei. Advancing retrosynthesis with retrieval-augmented graph generation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 20004–20013, 2025. doi: 10.1609/aaai.v39i19.34203.

11. S. Xie et al. Retrosynthesis prediction with local template retrieval. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pages 5330–5338, 2023. doi: 10.1609/aaai.v37i4.25664.

12. T. Guo et al. What can large language models do in chemistry? a comprehensive benchmark on eight tasks, 2023. NeurIPS 2023 Datasets and Benchmarks Track.

13. A. K. Hassen et al. Atom-anchored LLMs speak chemistry: A retrosynthesis demonstration, 2025.

14. Y. Wang et al. DianShi-RxnDB: A large-scale, fine-grained organic reaction data platform built via a fully automated pipeline for researchers and ai agents, 2026.

15. A. Yang et al. Qwen3 technical report, 2025.

16. Wengong Jin, Connor W. Coley, Regina Barzilay, and Tommi Jaakkola. Predicting organic reaction outcomes with Weisfeiler–Lehman network. In *Advances in Neural Information Processing Systems*, volume 30, 2017.

17. Hanjun Dai, Chengtao Li, Connor W. Coley, Bo Dai, and Le Song. Retrosynthesis prediction with conditional graph logic network. In *Advances in Neural Information Processing Systems*, volume 32, 2019.

18. Shuan Chen and Yousung Jung. Deep retrosynthetic reaction prediction using local reactivity and global attention. *JACS Au*, 1:1612–1620, 2021. doi: 10.1021/jacsau.1c00246.

19. Weihe Zhong, Ziduo Yang, and Calvin Yu-Chian Chen. Retrosynthesis prediction using an end-to-end graph generative architecture for molecular graph editing. *Nature Communications*, 14:3009, 2023. doi: 10.1038/s41467-023-38851-5.

20. Z. Zhong et al. Root-aligned SMILES: a tight representation for chemical reaction prediction. *Chemical Science*, 13:9023–9034, 2022. doi: 10.1039/d2sc02763a.

21. Y. Han et al. Retrosynthesis prediction with an iterative string editing model. *Nature Communications*, 15:6404, 2024. doi: 10.1038/s41467-024-50617-1.

22. L. Yao et al. Node-aligned graph-to-graph: Elevating template-free deep learning approaches in single-step retrosynthesis. *JACS Au*, 4:992–1003, 2024. doi: 10.1021/jacsau.3c00737.

23. Y. Deng et al. RSGPT: a generative transformer model for retrosynthesis planning pre-trained on ten billion datapoints. *Nature Communications*, 16:7012, 2025. doi: 10.1038/s41467-025-62308-6.

24. Z.ai. GLM-5.2 model card. https://huggingface.co/zai-org/GLM-5.2, 2026.

25. Y. Zhu et al. RetroAgent: Harnessing LLMs to search over structured memory for agentic retrosynthesis planning, 2026.

26. Wei Liu, Jiangtao Feng, Hongli Yu, Yuxuan Song, Yuqiang Li, Shufei Zhang, Lei Bai, Wei-Ying Ma, and Hao Zhou. Retro-R1: LLM-based agentic retrosynthesis. In *Advances in Neural Information Processing Systems*, volume 38, pages 78799–78827, 2025. doi: 10.52202/085713-2378.

27. K. Yu et al. Double-ended synthesis planning with goal-constrained bidirectional search. In *Advances in Neural Information Processing Systems*, volume 37, pages 112919–112949, 2024.

28. Binghong Chen, Chengtao Li, Hanjun Dai, and Le Song. Retro*: Learning retrosynthetic planning with neural guided A* search. In *Proceedings of the 37th International Conference on Machine Learning*, pages 1608–1616, 2020.

29. Xiaozhuang Song, Xuanhao Pan, Xinjian Zhao, Hangting Ye, Shufei Zhang, Jian Tang, and Tianshu Yu. AOT*: Efficient synthesis planning via LLM-empowered AND–OR tree search, 2025. arXiv:2509.20988v1; doi: 10.48550/arXiv.2509.20988.

30. P. Torren-Peraire et al. Models matter: the impact of single-step retrosynthesis on synthesis planning. *Digital Discovery*, 3:558–572, 2024. doi: 10.1039/D3DD00252G.

31. J. Choe, H. Kim, Y. T. Chok, M. Gim, and J. Kang. Retrosynthetic crosstalk between single-step reaction and multi-step planning. *Journal of Cheminformatics*, 17:130, 2025. doi: 10.1186/s13321-025-01088-z.

32. Y. Yu et al. RetroOOD: Understanding out-of-distribution generalization in retrosynthesis prediction. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 374–382, 2024. doi: 10.1609/aaai.v38i1.27791.

33. Yu Zhang et al. Large language models to accelerate organic chemistry synthesis. *Nature Machine Intelligence*, 7:1010–1022, 2025.

34. Shuaike Zhang et al. RetroDFM-R: reasoning-driven retrosynthesis prediction with large language models via reinforcement learning, 2026. Preprint at https://arxiv.org/abs/2507.17448v2.

35. Hyunseung Ko et al. RetroReasoner: a reasoning LLM for strategic retrosynthesis prediction, 2026. Preprint at https://arxiv.org/abs/2603.12666v3.

36. X. Li, S. Wang, Y. Lin, and Y. Wu. Retro-Expert: collaborative reasoning for interpretable retrosynthesis. In *Proceedings of the 43rd International Conference on Machine Learning (ICML)*, pages 70831–70866. PMLR, 2026.

37. K. Maziarz et al. Chemist-aligned retrosynthesis by ensembling diverse inductive bias models. *Nature*, 2026. doi: 10.1038/s41586-026-11160-9.

38. B. Zagribelnyy et al. When single answer is not enough: Rethinking single-step retrosynthesis benchmarks for LLMs, 2026.

39. B. Mikulak-Klucznik et al. Computational planning of the synthesis of complex natural products. *Nature*, 588:83–88, 2020. doi: 10.1038/s41586-020-2850-3.

40. RDKit. RDKit: Open-source cheminformatics. https://www.rdkit.org, 2025. Version 2025.9.5.

41. David Rogers and Mathew Hahn. Extended-connectivity fingerprints. *Journal of Chemical Information and Modeling*, 50:742–754, 2010. doi: 10.1021/ci100050t.

42. Peng-Cheng Zhao et al. Single-step retrosynthesis prediction via multitask graph representation learning. *Nature Communications*, 16:814, 2025.

43. Philippe Schwaller, Teodoro Laino, Théophile Gaudin, Peter Bolgar, Christopher A. Hunter, Costas Bekas, and Andrea Passerini. Molecular Transformer: A model for uncertainty-calibrated chemical reaction prediction. *ACS Central Science*, 5:1572–1583, 2019. doi: 10.1021/acscentsci.9b00576.

44. K. Maziarz et al. Re-evaluating retrosynthesis algorithms with Syntheseus. *Faraday Discussions*, 256: 568–586, 2025. doi: 10.1039/D4FD00093E. Associated preprint: arXiv:2310.19796v3.

45. eMolecules. eMolecules building-block inventory. https://www.emolecules.com, 2025.

46. Connor W. Coley, William H. Green, and Klavs F. Jensen. RDChiral: An RDKit wrapper for handling stereochemistry in retrosynthetic template extraction and application. *Journal of Chemical Information and Modeling*, 59:2529–2537, 2019. doi: 10.1021/acs.jcim.9b00286.

47. Samuel Genheden and Esben Bjerrum. PaRoutes: towards a framework for benchmarking retrosynthesis route predictions. *Digital Discovery*, 1:527–539, 2022. doi: 10.1039/D2DD00015F.

8 Extended Data for *Reaction precedents extend retrosynthetic prediction beyond training coverage*

8.1 Extended Data Fig. 1 | Top-k accuracy under template restriction and retrieval corpus expansion

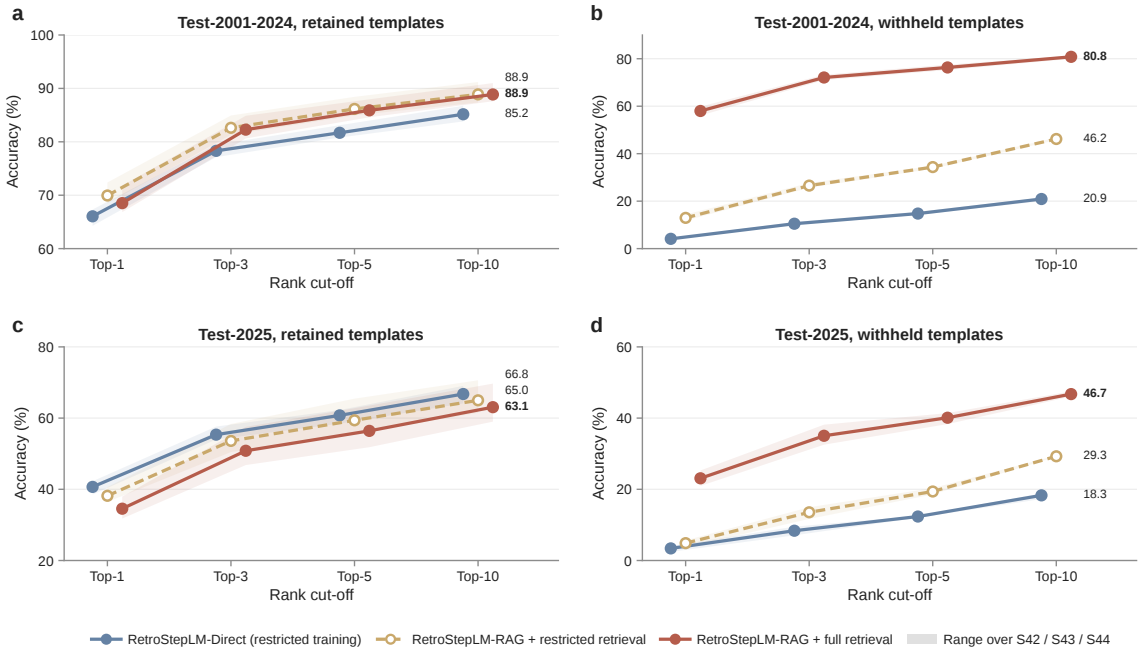

a–**d**, Complete precursor-set exact-match accuracy at the Top-1, Top-3, Top-5 and Top-10 rank cut-offs on Test-2001-2024 (**a**, retained; **b**, withheld templates) and Test-2025 (**c**, retained; **d**, withheld templates). Points are arithmetic means over the three template partitions (S42, S43, S44), with light shading spanning the partitions and reactions not pooled across them. Results are reported in Supplementary Note 4.

8.2 Extended Data Fig. 2 | Prediction accuracy across precedent-frequency ranges

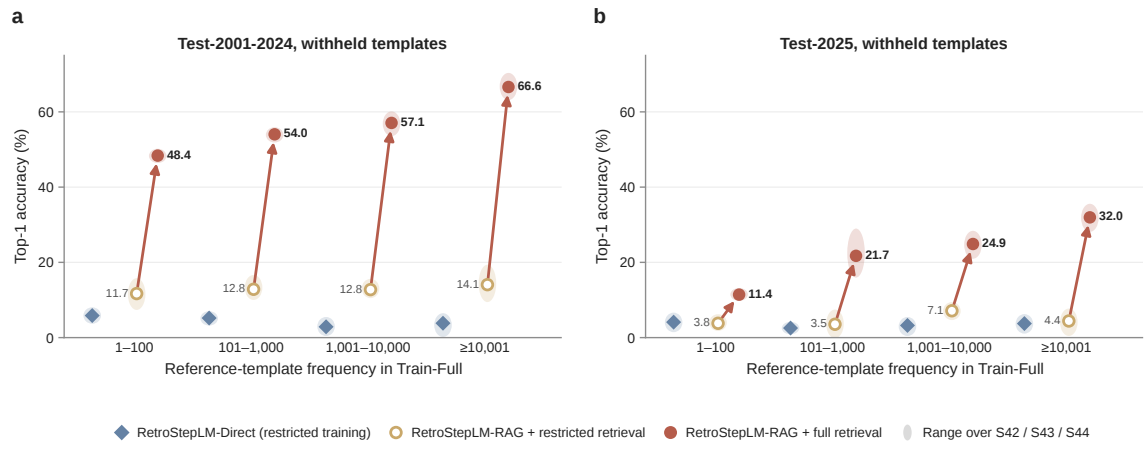

a,b, Top-1 accuracy on withheld-template reactions in Test-2001-2024 (**a**) and Test-2025 (**b**), stratified by the frequency of the recorded local reaction template in Train-Full: 1–100, 101–1,000, 1,001–10,000 and ≥10,001 records. Points are arithmetic means over the three template partitions (S42, S43, S44), with light shading spanning the partitions and reactions not pooled across them. Template counts, reaction denominators and results are reported in Supplementary Note 4.

8.3 Extended Data Fig. 3 | Top-k accuracy under temporal restriction and retrieval corpus expansion

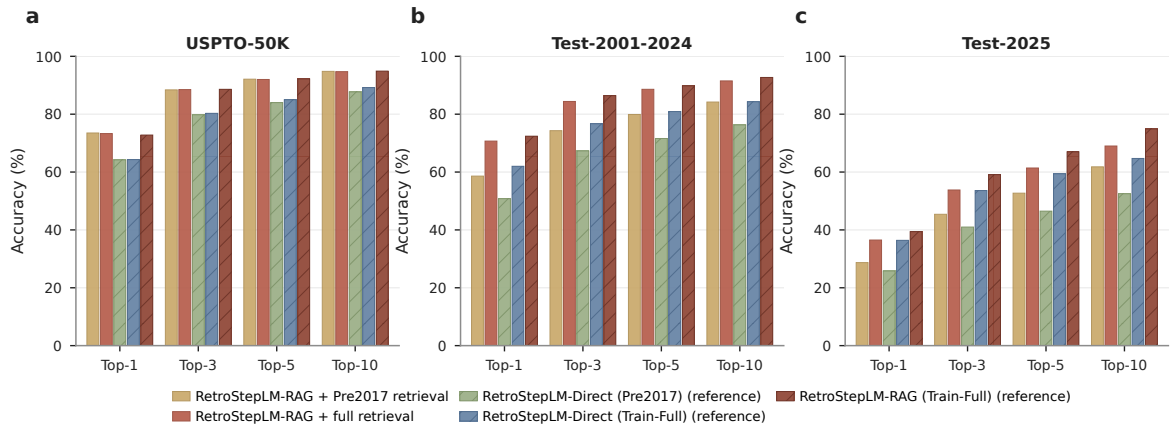

a–**c**, Complete precursor-set exact-match accuracy at the Top-1, Top-3, Top-5 and Top-10 rank cut-offs on USPTO-50K (**a**), Test-2001-2024 (**b**) and Test-2025 (**c**). Complete results are reported in Supplementary Note 4.

8.4 Extended Data Fig. 4 | Prediction accuracy under different precedent-selection conditions

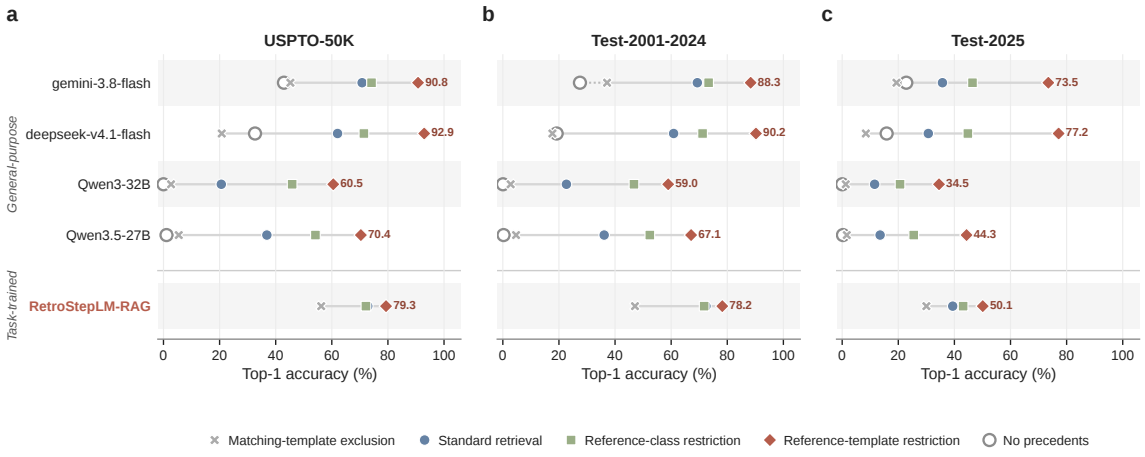

a–**c**, Top-1 accuracy on USPTO-50K (**a**), Test-2001-2024 (**b**) and Test-2025 (**c**) for four representative general-purpose models and RetroStepLM-RAG trained on Train-Full. Conditions use the full DianShi retrieval corpus with standard retrieval, reference-class restriction, reference-template restriction or

matching-template exclusion; grey circles mark the no-precedent condition for the general-purpose models. Reference classes and templates derive from the recorded answer, making these privileged diagnostics. The corpus and model remain fixed across selection conditions. Protocols and complete results are reported in Supplementary Note 5.

8.5 Extended Data Fig. 5 | Multistep solve rates across search budgets

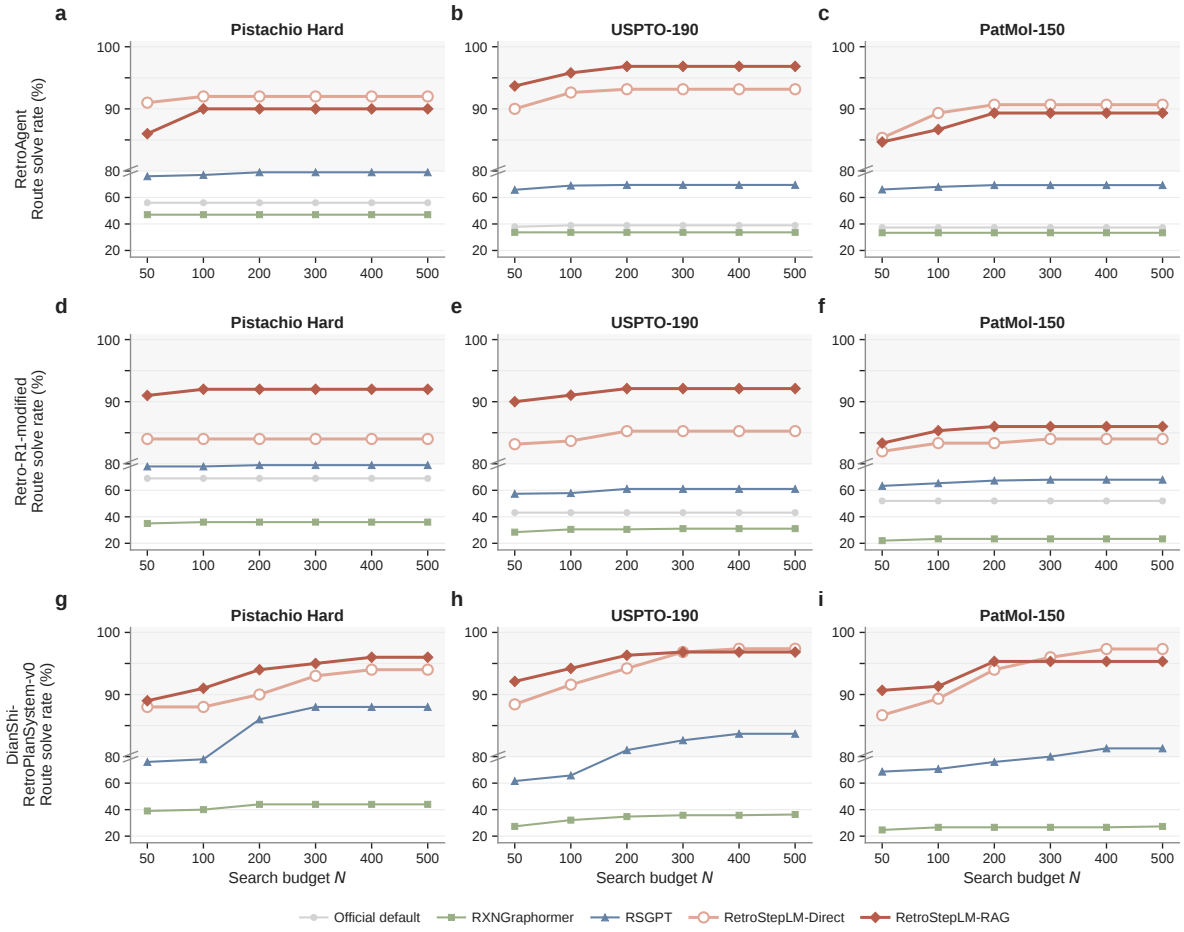

a–**i**, Route solve rates within RetroAgent (**a**–**c**), Retro-R1-modified (**d**–**f**) and DianShi-RetroPlanSystem-v0 (**g**–**i**) on Pistachio Hard (100 targets), USPTO-190 (190) and PatMol-150 (150). Within each framework, planner, starting-material inventory and search protocol are fixed. Protocols and complete results are reported in Supplementary Note 7.

9 Supplementary Information for *Reaction precedents extend retrosynthetic prediction beyond training coverage*

9.1 Supplementary Note 1. Dataset nomenclature and experimental configurations

1.1 Dataset partitions

DianShi-PatentRxnStep-Dataset comprises five base partitions, summarized in Supplementary Table 1. Years refer to the year of the source patent recorded in DianShi-RxnDB[14]. The five partitions are mutually disjoint by standardized target product. Their membership remains fixed across training and retrieval configurations.

Supplementary Table 1 | Dataset partitions.

Partition	Dataset identifier	Recorded years	Reactions	Role
Train-Full	`DianShi-PatentRxnStep-Train-Full`	1976–2024	2,712,259	Full training partition; also supplies the full retrieval corpus
Val-2001-2024	`DianShi-PatentRxnStep-Val-2001-2024`	2001–2024	3,798	Validation on historical records
Test-2001-2024	`DianShi-PatentRxnStep-Test-2001-2024`	2001–2024	3,798	Evaluation on historical records
Val-2025	`DianShi-PatentRxnStep-Val-2025`	2025	1,886	Temporally held-out validation
Test-2025	`DianShi-PatentRxnStep-Test-2025`	2025	1,886	Temporally held-out evaluation

The short partition names are used throughout this Supplementary Information. *Full* denotes the final training partition after the study's quality filtering and product-level exclusions, rather than the original database or the complete construction pool. Dataset construction, filtering counts and partitioning procedures are described in Supplementary Note 2.

1.2 Training and retrieval configurations

We distinguish the **training corpus**, which supplies reactions for task-specific parameter updates; the **checkpoint**, which contains the resulting model parameters; the **retrieval corpus**, which contains reaction records available for precedent selection; the **retrieval index**, which implements search over that corpus; and the **retrieved precedents**, which are the selected examples supplied to the model for a particular query. The **restricted retrieval corpus** matches a coverage-restricted training corpus, whereas the **full retrieval corpus** comprises all Train-Full reactions. *Restricted* and *full* qualify corpus coverage only: the retrieval procedure is identical across coverage conditions, and procedure-level constraints on precedent selection are described in Supplementary Note 5. A record set can serve both training and retrieval roles without becoming a separate dataset.

Three coverage labels identify the DianShi training and retrieval resources used in the coverage experiments:

- **Full** refers to Train-Full.
- **Pre2017** refers to the subset of Train-Full containing records before 2017, identified as `DianShi-PatentRxnStep-Train-Pre2017`.
- **TemplateRestricted-S42**, **TemplateRestricted-S43** and **TemplateRestricted-S44** refer to subsets of Train-Full obtained by removing all reactions bearing the templates withheld in the corresponding template partition. Their identifiers are `DianShi-PatentRxnStep-Train-TemplateRestricted-S42`, `DianShi-PatentRxnStep-Train-TemplateRestricted-S43` and `DianShi-PatentRxnStep-Train-TemplateRestricted-S44`.

For each restricted RetroStepLM-RAG configuration, the restricted record set supplied both the fine-tuning reactions and the matching restricted retrieval corpus. After training, the same checkpoint was evaluated with either that restricted retrieval corpus or the full retrieval corpus.

1.3 Template-defined evaluation subsets and reporting

The template intervention uses three alternative template partitions, denoted S42, S43 and S44. Within each partition, approximately half of the distinct reaction templates represented in the two base test sets were designated as withheld. Templates shared between the two test sets received the same assignment.

Each base test set was divided into **retained-template** and **withheld-template** subsets according to the corresponding partition. Retained templates had at least one reaction in the restricted training corpus and its matching restricted retrieval corpus; withheld templates had none. Temporal and transformation-based shifts represent distinct evaluation axes[32].

Subset display names combine the evaluation window, template status and partition, for example, *2001–2024 / Withheld (partition seed 42)*. The corresponding dataset identifier is `DianShi-PatentRxnStep-Test-2001-2024-TemplateWithheld-S42`; retained subsets use `TemplateRetained`.

Within each partition, the retained- and withheld-template subsets are mutually exclusive and together comprise the corresponding base test set.

Single-step accuracy measures exact agreement with the complete standardized reference precursor set, following the scoring procedure in Methods. Accuracy within a subset uses the number of evaluation reactions in that subset as its denominator. Cross-partition summaries are arithmetic means of the three partition-specific accuracies, calculated separately for each test set, template status and, where applicable, precedent-frequency stratum. Evaluation reactions are not pooled across partitions as independent samples. A reported mean gain is the arithmetic mean of the three within-partition accuracy differences.

1.4 Configuration correspondence

Supplementary Table 2 summarizes the coverage experiments and their reference configurations. Coverage labels refer to the record sets defined in Section 1.2. RetroStepLM-Direct receives no retrieved precedents. For RetroStepLM-RAG, the retrieval corpus used during training matches its training corpus; the inference retrieval corpus is specified separately.

Supplementary Table 2 | Configuration correspondence.

Model	Training corpus	Inference retrieval corpus	Comparison role
RetroStepLM-Direct	`TemplateRestricted-S{s}`	None	Separately trained template-restricted reference
RetroStepLM-RAG	`TemplateRestricted-S{s}`	`TemplateRestricted-S{s}`	Restricted-retrieval condition
Same RetroStepLM-RAG checkpoint	`TemplateRestricted-S{s}`	`Full`	Expanded-retrieval condition
RetroStepLM-Direct	`Pre2017`	None	Separately trained temporally restricted reference
RetroStepLM-RAG	`Pre2017`	`Pre2017`	Restricted-retrieval condition
Same RetroStepLM-RAG checkpoint	`Pre2017`	`Full`	Expanded-retrieval condition
RetroStepLM-Direct	`Full`	None	Separately trained full-corpus reference
RetroStepLM-RAG	`Full`	`Full`	Separately trained full-corpus reference

Here, $s \in \{42, 43, 44\}$; the three template-restricted rows apply separately to each partition. Within each paired RetroStepLM-RAG comparison, the checkpoint, evaluation targets, retrieval procedure, prece-

dent serialization, prompting, decoding and scoring remain fixed. Only the retrieval corpus and its corresponding index change. Standard inference supplies five retrieved precedents per query.

The full-corpus RetroStepLM-RAG reference has separately trained parameters. It is therefore distinct from a restricted-training checkpoint evaluated with the full retrieval corpus. General-purpose model configurations and additional adaptation and retrieval-transfer experiments are described in Supplementary Notes 6 and 8, respectively.

9.2 Supplementary Note 2. Reaction dataset construction and partitioning

2.1 Reaction standardization and deduplication

DianShi-PatentRxnStep-Dataset was derived from DianShi-RxnDB[14], a provenance-linked organic reaction-data platform that we previously developed. DianShi-RxnDB uses a fully automated pipeline to extract structured reaction records from patent text, images and reaction schemes, primarily from USPTO and EPO organic-synthesis patents published between 1976 and 2025. Each record represents a reported single-step experiment and links to its source patent.

Reactant and product SMILES were standardized using RDKit[40]. Reaction-level duplicates were identified using a standardized reactant-set–product-set signature. Atom-mapping numbers were removed from the comparison representation, and molecular components on each side were sorted deterministically, so differences in component order or mapping-number assignment did not produce distinct signatures. Records with identical signatures were treated as duplicates. When duplicate records spanned multiple years, the record with the earliest year recorded in DianShi-RxnDB was retained.

2.2 Reaction-quality filtering and template processing

Reaction-quality filtering excluded records with multiple products, repeated molecules on either side, structural overlap between reactants and products, more than four reactant components or inorganic products. Records were also excluded when atom mappings were inconsistent or reaction-centre or template extraction was unsuccessful.

Local reaction templates[18] were grouped across the quality-controlled corpus. Templates represented by fewer than six reactions were removed before temporal partitioning, following the cutoff used in previous template-frequency analyses[2]. This threshold applied to template frequency in the corpus at that processing stage.

The principal processing and partitioning rules are summarized in Supplementary Table 3.

Supplementary Table 3 | Processing and partitioning rules.

Processing component	Rule	Scope
Molecular standardization	Standardize reactant and product SMILES using RDKit	Reaction records entering dataset construction
Reaction deduplication	Compare standardized reactant-set–product-set signatures after removing mapping numbers and sorting components; retain the earliest recorded year for cross-year duplicates	Duplicate reaction records
Structural quality filtering	Exclude multiple products, repeated molecules, reactant–product overlap, more than four reactant components and inorganic products	Deduplicated records
Mapping and extraction filtering	Exclude inconsistent atom mappings and unsuccessful reaction-centre or template extraction	Records undergoing quality control

Processing component	Rule	Scope
Template-frequency filtering	Remove local reaction templates represented by fewer than six reactions	Quality-controlled corpus, before temporal partitioning
External benchmark exclusions	Compare standardized target products against the specified external dataset partitions	Training and evaluation candidates, as detailed in Section 2.3
Internal partition separation	Require standardized target products to be mutually disjoint across all five base partitions	Final training, validation and test partitions

2.3 Temporal partitioning and product-level exclusions

For temporal evaluation[3,8], partitions were defined using the year of the source patent recorded in DianShi-RxnDB. Records before 2025 supplied the training and historical evaluation pools, whereas records from 2025 supplied the temporally held-out evaluation pool.

Val-2001-2024 and Test-2001-2024 were sampled from 2001–2024 records, with stratification by year and template frequency in the full training corpus. Val-2025 and Test-2025 were sampled from the 2025 pool, with stratification by patent type and template frequency. The five base partitions were constructed jointly and were mutually disjoint by standardized target product[2]. Consequently, overlapping year ranges between training and historical evaluation did not entail shared target products.

Product-level exclusions controlled overlap with external comparison datasets. Training candidates were screened against products in the validation and test partitions of USPTO-50K[7] and USPTO-Full[17], and USPTO-MIT[16]. DianShi validation and test candidates were additionally screened against products in the training partitions of these USPTO datasets and of the corpus used to pretrain RXNGraphormer[6]. These exclusions operated on standardized target products irrespective of the associated precursor set.

2.4 Final dataset composition

The resulting dataset comprised 2,712,259 reactions in Train-Full, 3,798 reactions each in Val-2001-2024 and Test-2001-2024, and 1,886 reactions each in Val-2025 and Test-2025. Dataset identifiers and partition roles are summarized in Supplementary Note 1.

Train-Full contains records spanning 1976–2024 and supplies all reaction records in the full DianShi retrieval corpus, without additional corpus-level subsampling. The restricted training and retrieval corpora used in the coverage experiments were derived from this fixed partition, as described in Supplementary Note 4.

9.3 Supplementary Note 3. Precedent retrieval, intermediate-label supervision and model training

3.1 Retrieval implementation and precedent selection

Retrieval used the target product alone, without reference reaction-class or template information. Products in each retrieval corpus were represented by extended-connectivity fingerprints (ECFPs)[41]. A binary flat index returned 200 candidate reactions ranked by Hamming distance to the query fingerprint. These candidates were then reranked by the Tanimoto similarity between their product fingerprints and the query fingerprint.

During RetroStepLM-RAG fine-tuning, the query reaction, duplicate reactions and potentially answer-disclosing records from the same patent were excluded from candidate precedents. At each training epoch, five eligible precedents were randomly selected from the ten highest-similarity candidates. At

standard inference, RetroStepLM-RAG received the five highest-ranked eligible precedents deterministically. The general-purpose model evaluations used the same standard inference selection procedure.

Indices were constructed from the full training partition and each restricted training corpus without additional corpus-level subsampling. Within each retrieval corpus expansion comparison, the checkpoint and retrieval procedure remained unchanged; inference switched between indices built from the corresponding restricted and full record sets. Configuration correspondence is provided in Supplementary Note 1, and the interventions are detailed in Supplementary Note 4. Reference-informed modifications to precedent selection are described separately in Supplementary Note 5.

3.2 Precedent representation and model inputs

Each retrieved precedent contained standardized product and reactant SMILES, together with synthon SMILES and corresponding leaving groups derived from the atom-mapped reaction. These fields represented the precedent reaction and were supplied as reference context.

RetroStepLM-Direct received only the query product. RetroStepLM-RAG received the query product together with five retrieved precedents. The query reaction's reference reactants, synthons and leaving groups were used to construct training targets and were not included in the query input. At inference, these fields were generated by the model.

In the prompt records below, `{products}` denotes the query product and `{rag_examples}` denotes the retrieved-precedent block. The placeholders `{synthons_smiles}`, `{leaving_groups}` and `{reactants}` occur in the supervised target and denote labels derived from the query reaction.

3.3 Intermediate-label supervision and prompt templates

Both DianShi models used intermediate-label supervision. Synthon- and leaving-group-based decompositions have also been used in structured retrosynthesis models[5,42]. Atom-mapped reactions supplied labels comprising synthons associated with the retrosynthetic disconnection, corresponding leaving groups and the final standardized reactant set. The target sequence presented these fields in the order `synthons_smiles`, `leaving_groups` and `reactants`, with the final reactant string additionally enclosed in `<answer>` tags.

The structured prompt text and supervised target are shown separately below. The `<think>` tags delimit the serialized intermediate fields.

RetroStepLM-Direct input

```
---SYSTEM---
You are an expert chemist specializing in retrosynthetic analysis. Given the SMILES representation of a
    product molecule, your task is to predict the SMILES of the reactants required to synthesize it.
Input:
The product SMILES.
Task requirements (step-by-step reasoning is required):
1. Infer the synthons based on the structure of the product.
2. Identify the corresponding leaving groups associated with these synthons.
3. Reconstruct the most plausible reactant molecules by combining the synthons with appropriate leaving
    groups.
You must explicitly show your reasoning step by step for each of the above stages.
Please output the reasoning process, and write the final reactant SMILES strictly in the following
    format:
<answer>REACTANT_SMILES</answer>
If multiple reactants are required, separate their SMILES using a period (.).
---USER---
The product SMILES is {products}.
```

RetroStepLM-Direct supervised target

```
<think>The product SMILES is {products}. By analyzing the product SMILES, the synthons involved in this
    reaction are identified as {synthons_smiles}.
Based on the product structure and the inferred synthons, the corresponding leaving groups for this
    reaction are deduced to be {leaving_groups}.
By combining the synthons with the identified leaving groups, the final reactant(s) are inferred as:
    {reactants}</think>
<answer>{reactants}</answer>
```

RetroStepLM-RAG input text

```
You are an expert chemist specializing in retrosynthetic analysis. Given the SMILES representation of a
    product molecule, your task is to predict the SMILES of the reactants required to synthesize it.
    The input consists of a product SMILES and multiple referenced reactions. Task requirements
    (step-by-step reasoning is required):
1. Infer the synthons based on the structure of the product.
2. Identify the corresponding leaving groups associated with these synthons.
3. Reconstruct the most plausible reactant molecules by combining the synthons with appropriate leaving
    groups.
You must explicitly show your reasoning step by step for each of the above stages.
Please output the reasoning process, and write the final reactant SMILES strictly in the following
    format:
<answer>REACTANT_SMILES</answer>
If multiple reactants are required, separate their SMILES using a period (.).

The product SMILES is {products}. Below are similar reactions retrieved from a reaction database for
    reference:
{rag_examples}
```

RetroStepLM-RAG supervised target

```
<think>The product SMILES is {products}. By analyzing the product SMILES and multiple referenced
    reactions, the synthons involved in this reaction are identified as {synthons_smiles}.
Based on the product structure and the inferred synthons, the corresponding leaving groups for this
    reaction are deduced to be {leaving_groups}.
By combining the synthons with the identified leaving groups, the final reactant(s) are inferred as:
    {reactants}</think>
<answer>{reactants}</answer>
```

Training minimized the autoregressive negative log-likelihood over the complete supervised target sequence, as defined in Methods. Only the final reactant set contributed to precursor-set accuracy. General-purpose model prompting and model-specific call settings are described in Supplementary Note 6.

3.4 Model training

DianShi-RetroStepLM-Direct and DianShi-RetroStepLM-RAG were developed from Qwen3-4B-Base[15] using full-parameter supervised fine-tuning; for RetroStepLM-RAG, this took the form of retrieval-conditioned supervised fine-tuning, with the target product and five retrieved precedents supplied as input. SMILES were space-segmented into atom- and bond-level tokens following the Molecular Transformer tokenization scheme[43] and encoded using the native Qwen tokenizer without vocabulary expansion.

Full-corpus models used Train-Full. Coverage-intervention models used their corresponding Pre2017 or TemplateRestricted-S{s} training corpus, where $s \in \{42, 43, 44\}$. For each RetroStepLM-RAG model, training precedents were retrieved from the same record set that supplied its fine-tuning reactions.

Training used 32 H200 GPUs across four nodes, with bfloat16 precision, FlashAttention, DeepSpeed ZeRO-3 and gradient checkpointing. The recorded training settings are summarized in Supplementary Table 4.

Supplementary Table 4 | Training settings.

Item	Setting
Base model	Qwen3-4B-Base
Parameter updates	Full-parameter supervised fine-tuning
Supervision	Synthon–leaving-group intermediate labels
Hardware	32 H200 GPUs across four nodes
Precision	bfloat16
Training implementation	FlashAttention; DeepSpeed ZeRO-3; gradient checkpointing
Per-device batch size	4
Gradient accumulation	1 step
Effective batch size	128
Training duration	10 epochs
Learning rate	1×10^{-4}
Learning-rate schedule	Cosine decay with linear warm-up over the first 5% of training steps
Tokenization	Molecular Transformer-style SMILES spacing followed by the native Qwen tokenizer; no vocabulary expansion
RetroStepLM-RAG training precedents	Random five from the ten highest-similarity eligible candidates per epoch
RetroStepLM-RAG standard inference precedents	Deterministic Top-5

Additional adaptation experiments and training-time precedent-sampling ablations are described in Supplementary Note 8. Complete single-step performance comparisons are reported in Supplementary Note 6.

3.5 Decoding and candidate extraction

At inference, both DianShi models used beam search to generate ranked precursor-set candidates for Top-1, Top-3, Top-5 and Top-10 evaluation. The final reactant string enclosed in `<answer>` tags was extracted from each generated sequence.

Predicted and reference precursor sets were standardized using RDKit, converted to canonical SMILES and compared after deterministic component sorting[44]. Component order therefore did not affect correctness. Candidates were evaluated in their generated rank order. Missing predictions, invalid molecular strings and partial precursor-set matches were counted as incorrect under the common scoring procedure defined in Methods.

9.4 Supplementary Note 4. Fixed-parameter retrieval corpus expansion

4.1 Template-coverage intervention

We constructed three alternative template partitions using partition seeds 42, 43 and 44. Within each partition, approximately half of the distinct local reaction templates represented in Test-2001-2024 and Test-2025 were designated as withheld. Templates shared between the two test sets received the same retain-or-withhold assignment. The withholding target referred to distinct templates rather than reaction counts.

All reactions bearing withheld templates were removed from Train-Full. The remaining records served as both the TemplateRestricted-S{s} training corpus and its matching restricted retrieval corpus, where $s \in \{42, 43, 44\}$. Withheld evaluation templates therefore had zero occurrence in both restricted resources, whereas retained evaluation templates had at least one corresponding training reaction. Restricted-corpus sizes are summarized in Supplementary Table 5.

Supplementary Table 5 | Template-restricted training and retrieval corpora.

Record set	Reactions retained	Percentage of Train-Full (%)	Reactions removed
Train-Full	2,712,259	100.00	0
TemplateRestricted-S42	1,804,926	66.55	907,333
TemplateRestricted-S43	1,673,416	61.70	1,038,843
TemplateRestricted-S44	1,564,454	57.68	1,147,805

For each partition, RetroStepLM-Direct and RetroStepLM-RAG were trained on the restricted corpus. RetroStepLM-RAG training used precedents retrieved from the matching restricted retrieval corpus. After training, its checkpoint was fixed and evaluated with either that restricted retrieval corpus or the full retrieval corpus. The two conditions used identical evaluation targets and execution settings apart from the retrieval corpus and corresponding index. RetroStepLM-Direct provided a separately trained reference.

Each base test set was divided into retained- and withheld-template subsets according to the corresponding partition. These assignments remained unchanged after retrieval corpus expansion. Subset sizes are reported in Supplementary Table 6.

Supplementary Table 6 | Template-defined evaluation subset sizes. Entries are numbers of evaluation reactions.

Partition	2001–2024 / Retained	2001–2024 / Withheld	2025 / Retained	2025 / Withheld
S42	1,977	1,821	970	916
S43	1,842	1,956	937	949
S44	1,797	2,001	878	1,008

Within each partition, the retained and withheld subsets contained 671 distinct templates each for Test-2001-2024, and 426 and 427 distinct templates, respectively, for Test-2025.

Supplementary Table 7 reports complete precursor-set exact-match accuracy. All models in a given partition were trained on its TemplateRestricted-S{s} corpus. *Restricted retrieval* and *full retrieval* denote evaluations of the same RetroStepLM-RAG checkpoint.

Supplementary Table 7 | Complete template-intervention results (%).

a, Test-2001-2024 / Retained.

Partition	Configuration	Top-1	Top-3	Top-5	Top-10
S42	RetroStepLM-Direct	67.32	79.46	83.11	86.75
S42	RetroStepLM-RAG, restricted retrieval	72.38	84.93	88.42	91.20
S42	RetroStepLM-RAG, full retrieval	71.78	84.88	87.96	91.00
S43	RetroStepLM-Direct	66.56	78.34	81.16	84.91
S43	RetroStepLM-RAG, restricted retrieval	69.00	81.65	84.96	87.35
S43	RetroStepLM-RAG, full retrieval	66.94	81.32	85.12	88.00
S44	RetroStepLM-Direct	64.27	77.13	80.80	83.81
S44	RetroStepLM-RAG, restricted retrieval	68.45	81.30	85.09	88.09
S44	RetroStepLM-RAG, full retrieval	66.83	80.63	84.59	87.59

b, Test-2001-2024 / Withheld.

Partition	Configuration	Top-1	Top-3	Top-5	Top-10
S42	RetroStepLM-Direct	4.39	11.26	15.54	21.25
S42	RetroStepLM-RAG, restricted retrieval	11.81	25.65	33.66	46.07
S42	RetroStepLM-RAG, full retrieval	56.07	71.00	75.23	79.96
S43	RetroStepLM-Direct	4.24	10.12	14.57	21.47
S43	RetroStepLM-RAG, restricted retrieval	14.37	27.66	34.92	45.76
S43	RetroStepLM-RAG, full retrieval	58.08	72.03	76.23	81.19
S44	RetroStepLM-Direct	3.75	10.09	14.19	19.99
S44	RetroStepLM-RAG, restricted retrieval	12.69	26.34	34.48	46.73
S44	RetroStepLM-RAG, full retrieval	59.82	73.26	77.51	81.26

c, Test-2025 / Retained.

Partition	Configuration	Top-1	Top-3	Top-5	Top-10
S42	RetroStepLM-Direct	41.24	57.63	62.27	67.11
S42	RetroStepLM-RAG, restricted retrieval	40.62	58.35	65.46	70.62
S42	RetroStepLM-RAG, full retrieval	37.94	57.53	63.51	69.69
S43	RetroStepLM-Direct	38.53	54.43	61.58	69.05
S43	RetroStepLM-RAG, restricted retrieval	37.46	51.97	57.10	62.86
S43	RetroStepLM-RAG, full retrieval	33.94	48.13	53.90	60.51
S44	RetroStepLM-Direct	42.26	53.99	58.43	64.12
S44	RetroStepLM-RAG, restricted retrieval	36.45	50.46	55.58	61.50
S44	RetroStepLM-RAG, full retrieval	31.78	46.81	51.82	59.00

d, Test-2025 / Withheld.

Partition	Configuration	Top-1	Top-3	Top-5	Top-10
S42	RetroStepLM-Direct	3.71	8.52	12.55	18.78
S42	RetroStepLM-RAG, restricted retrieval	3.71	11.68	18.56	29.15
S42	RetroStepLM-RAG, full retrieval	20.96	32.53	38.32	46.07
S43	RetroStepLM-Direct	2.42	7.06	12.01	18.86
S43	RetroStepLM-RAG, restricted retrieval	5.27	14.96	20.34	29.50
S43	RetroStepLM-RAG, full retrieval	25.29	38.15	41.41	47.52
S44	RetroStepLM-Direct	4.07	9.52	12.50	17.26
S44	RetroStepLM-RAG, restricted retrieval	5.65	13.99	19.25	29.17
S44	RetroStepLM-RAG, full retrieval	23.02	34.42	40.58	46.63

Retrieval corpus expansion improved withheld-template Top-1 accuracy in every partition. Mean partition-level accuracy increased from 13.0% to 58.0% on Test-2001-2024 and from 4.9% to 23.1% on Test-2025. Retained-template accuracy decreased in every partition, with corresponding means changing from 69.9% to 68.5% and from 38.2% to 34.6%.

4.2 Precedent-frequency stratification

We stratified withheld-template evaluation reactions by the number of records bearing their local reaction template in Train-Full, which also constituted the full retrieval corpus. Frequencies were calculated before template withholding. The four frequency ranges were 1–100, 101–1,000, 1,001–10,000 and at least 10,001 reactions per template. Supplementary Table 8 reports distinct-template counts, evaluation-reaction counts and Top-1 accuracy for each partition and stratum. Evaluation-reaction

counts provide the accuracy denominators. All three columns within a row used models trained on the corresponding TemplateRestricted-S{s} corpus; the two RetroStepLM-RAG columns share a checkpoint.

Supplementary Table 8 | Withheld-template composition and frequency-stratified Top-1 accuracy (%).

a, Test-2001-2024 / Withheld.

Reactions per template in full corpus	Partition	Distinct templates	Evaluation reactions	RetroStepLM-Direct	RetroStepLM-RAG, restricted retrieval	RetroStepLM-RAG, full retrieval
1–100	S42	290	305	6.56	8.85	48.52
1–100	S43	290	309	5.18	11.97	47.90
1–100	S44	290	308	5.84	14.29	48.70
101–1,000	S42	269	446	4.71	11.88	53.36
101–1,000	S43	269	452	5.53	11.50	54.42
101–1,000	S44	269	429	5.36	15.15	54.31
1,001–10,000	S42	92	592	2.70	11.99	55.07
1,001–10,000	S43	92	574	1.92	14.11	57.49
1,001–10,000	S44	92	551	3.99	12.16	58.62
≥10,001	S42	20	478	4.81	13.39	64.64
≥10,001	S43	20	621	4.99	17.87	66.34
≥10,001	S44	20	713	1.68	10.94	68.86

b, Test-2025 / Withheld.

Reactions per template in full corpus	Partition	Distinct templates	Evaluation reactions	RetroStepLM-Direct	RetroStepLM-RAG, restricted retrieval	RetroStepLM-RAG, full retrieval
1–100	S42	188	212	4.25	4.72	11.32
1–100	S43	188	204	2.94	3.43	11.27
1–100	S44	188	214	5.14	3.27	11.68
101–1,000	S42	147	222	2.70	1.35	20.27
101–1,000	S43	147	211	2.37	3.32	27.49
101–1,000	S44	147	235	2.55	5.96	17.45
1,001–10,000	S42	72	276	3.99	6.16	22.46
1,001–10,000	S43	72	261	1.92	7.28	25.29
1,001–10,000	S44	72	264	3.79	7.95	26.89
≥10,001	S42	20	206	3.88	1.94	29.61
≥10,001	S43	20	273	2.56	6.23	34.07
≥10,001	S44	20	295	4.75	5.08	32.20

Expansion improved Top-1 accuracy in every partition and frequency range in both test sets. For templates represented by 1–100 reactions in the full corpus, the mean partition-level gain was 36.7 percentage points on Test-2001-2024 and 7.6 percentage points on Test-2025.

4.3 Temporal-coverage intervention

The Pre2017 training corpus retained 1,720,962 reactions before 2017, representing 63.45% of Train-Full. The remaining 991,297 reactions were excluded from both task-specific training and the Pre2017 retrieval corpus.

RetroStepLM-RAG was trained using Pre2017 reactions and precedents retrieved from the matching corpus. After training, its checkpoint was fixed and evaluated with either the Pre2017 retrieval corpus or the full retrieval corpus.

Supplementary Table 9 reports complete precursor-set exact-match accuracy. The two Pre2017-trained RetroStepLM-RAG rows share the same checkpoint. Pre2017-trained RetroStepLM-Direct and the full-corpus models are separately trained references.

Supplementary Table 9 | Complete temporal-intervention results (%).

a, USPTO-50K.

Model	Training corpus	Inference retrieval corpus	Top-1	Top-3	Top-5	Top-10
RetroStepLM-Direct	Pre2017	None	64.2	79.8	84.0	87.7
RetroStepLM-Direct	Full	None	64.3	80.3	85.0	89.2
RetroStepLM-RAG	Pre2017	Pre2017	73.5	88.4	92.1	94.8
RetroStepLM-RAG	Pre2017	Full	73.3	88.5	92.0	94.7
RetroStepLM-RAG	Full	Full	72.8	88.6	92.2	94.9

b, Test-2001-2024.

Model	Training corpus	Inference retrieval corpus	Top-1	Top-3	Top-5	Top-10
RetroStepLM-Direct	Pre2017	None	50.8	67.3	71.6	76.3
RetroStepLM-Direct	Full	None	62.0	76.7	80.9	84.3
RetroStepLM-RAG	Pre2017	Pre2017	58.6	74.3	79.9	84.2
RetroStepLM-RAG	Pre2017	Full	70.7	84.4	88.6	91.5
RetroStepLM-RAG	Full	Full	72.4	86.4	89.8	92.7

c, Test-2025.

Model	Training corpus	Inference retrieval corpus	Top-1	Top-3	Top-5	Top-10
RetroStepLM-Direct	Pre2017	None	25.8	41.0	46.5	52.5
RetroStepLM-Direct	Full	None	36.4	53.6	59.4	64.7
RetroStepLM-RAG	Pre2017	Pre2017	28.7	45.4	52.7	61.8
RetroStepLM-RAG	Pre2017	Full	36.5	53.8	61.4	69.0
RetroStepLM-RAG	Full	Full	39.4	59.1	67.0	74.9

Retrieval corpus expansion increased Top-1 accuracy by 12.1 percentage points on Test-2001-2024 and 7.8 percentage points on Test-2025. USPTO-50K accuracy remained essentially unchanged, from 73.5% to 73.3%.

With full retrieval, the Pre2017-trained RetroStepLM-RAG checkpoint exceeded full-corpus RetroStepLM-Direct on Test-2001-2024 and achieved a similar Top-1 accuracy on Test-2025. It remained below the separately trained full-corpus RetroStepLM-RAG reference on both test sets.

Expansion changed corpus size, recency and composition simultaneously; these comparisons do not separate their contributions.

9.5 Supplementary Note 5. Precedent-relevance diagnostics

5.1 Models and evaluation settings

We conducted reference-informed diagnostics to examine how precedent selection affected retrosynthetic prediction when the model and retrieval corpus were fixed. Evaluations used DianShi-RetroStepLM-RAG trained on Train-Full and four representative general-purpose models: gemini-3.8-flash, deepseek-v4.1-flash, Qwen3-32B and Qwen3.5-27B. All five models were evaluated on USPTO-50K, Test-2001-2024 and Test-2025.

Each retrieval condition used the full DianShi retrieval corpus and supplied five precedents per query. Within each model, parameters, evaluation targets, prompt structure, decoding and scoring settings remained fixed; only the precedent-selection rule changed.

5.2 Precedent-selection conditions

Standard retrieval was compared with three reference-informed selection conditions. Reaction classes and local reaction templates were obtained from the recorded reference reactions (Supplementary Table 10).

Supplementary Table 10 | Precedent-selection rules used in the diagnostics.

Condition	Selection rule
Standard retrieval	Select precedents by target-product structural similarity, without reference reaction-class or template constraints
Reference-class restriction	Select precedents sharing the recorded reference reaction class
Reference-template restriction	Select precedents sharing the local reaction template of the reference reaction
Matching-template exclusion	Exclude precedents bearing the local reaction template of the reference reaction

Standard retrieval followed the procedure described in Supplementary Note 3. In the matching-template exclusion condition, retrieval continued down the ranked candidate list to collect five eligible precedents. When fewer than five remained, eligible precedents were cycled to fill the five context positions. Thus, maintaining five positions did not necessarily imply five distinct precedents.

The reference-informed conditions used information derived from the query reaction's recorded answer. They were therefore privileged diagnostics rather than retrieval settings assumed available for prospective prediction.

5.3 Diagnostic results

Supplementary Table 11 reports Top-1 complete precursor-set exact-match accuracy and corresponds to Extended Data Fig. 4. Conditions are defined in Supplementary Table 10. A dash indicates that no corresponding result is reported.

Supplementary Table 11 | Prediction accuracy under different precedent-selection conditions (%).

a, USPTO-50K.

Model	No precedents	Standard retrieval	Reference-class restriction	Reference-template restriction	Matching-template exclusion
gemini-3.8-flash	42.94	70.80	74.16	90.75	45.24
deepseek-v4.1-flash	32.63	62.05	71.44	92.93	20.77
Qwen3-32B	0.00	20.61	45.86	60.52	2.70
Qwen3.5-27B	1.08	36.85	54.16	70.36	5.45
RetroStepLM-RAG	—	72.80	72.16	79.33	56.22

b, Test-2001-2024.

Model	No precedents	Standard retrieval	Reference-class restriction	Reference-template restriction	Matching-template exclusion
gemini-3.8-flash	27.49	69.33	73.35	88.34	37.15
deepseek-v4.1-flash	19.19	60.87	71.20	90.23	17.67
Qwen3-32B	0.05	22.67	46.74	58.95	2.79
Qwen3.5-27B	0.26	36.12	52.40	67.11	4.71

Model	No precedents	Standard retrieval	Reference-class restriction	Reference-template restriction	Matching-template exclusion
RetroStepLM-RAG	—	72.40	71.77	78.23	47.10

c, Test-2025.

Model	No precedents	Standard retrieval	Reference-class restriction	Reference-template restriction	Matching-template exclusion
gemini-3.8-flash	22.85	35.74	46.45	73.49	19.35
deepseek-v4.1-flash	15.85	30.70	44.80	77.15	8.43
Qwen3-32B	0.00	11.56	20.62	34.52	1.27
Qwen3.5-27B	0.27	13.52	25.50	44.33	1.64
RetroStepLM-RAG	—	39.40	43.11	50.11	30.01

RetroStepLM-RAG, developed in this work, received retrieved precedents during both training and inference; the general-purpose models received precedents only at inference, without task-specific parameter updates.

Reference-template restriction improved accuracy relative to standard retrieval for all five models on all three test sets. Matching-template exclusion reduced accuracy in every model–test-set combination. On Test-2025, the exclusion condition fell below the no-precedent reference for gemini-3.8-flash (19.35% versus 22.85%) and deepseek-v4.1-flash (8.43% versus 15.85%). For deepseek-v4.1-flash, the exclusion condition also fell below the no-precedent reference on USPTO-50K and Test-2001-2024.

Reference-class restriction improved accuracy for all four general-purpose models on all three test sets. Its effect on RetroStepLM-RAG differed: accuracy decreased slightly on USPTO-50K and Test-2001-2024 and increased on Test-2025. Reference-class restriction therefore did not provide a uniform improvement across the evaluated models.

5.4 Interpretation and scope

These diagnostics show that the utility of retrieved precedents depends on their selection. Structurally similar products can arise from different transformations, and standard retrieval does not guarantee that the selected precedents share the reference transformation.

The reference-informed conditions nevertheless have a limited interpretation. Reference classes and templates are derived from the recorded answer and are generally unavailable for prospective prediction. Their accuracies describe performance under privileged selection conditions rather than deployable retrieval performance.

9.6 Supplementary Note 6. General-purpose model evaluation and specialized-model comparisons

6.1 General-purpose models and inference settings

We evaluated 14 representative general-purpose models without task-specific parameter updates: seven API-accessed models across major providers and seven open-weight Qwen models spanning 2B to 32B parameters. Each model was evaluated on USPTO-50K, Test-2001-2024 and Test-2025, both without retrieved precedents and with precedents drawn from each of three retrieval corpora. The evaluated models are listed in Supplementary Table 12.

Supplementary Table 12 | General-purpose models.

Evaluation group	Model identifiers
API-accessed models	glm-5.2; qwen3.8-MAX; kimi-k3; deepseek-v4-pro; deepseek-v4.1-flash; gpt-5.6-sol; gemini-3.8-flash
Open-weight Qwen models	Qwen3-4B; Qwen3-8B; Qwen3-32B; Qwen3.5-2B; Qwen3.5-4B; Qwen3.5-9B; Qwen3.5-27B

Models used structured sequence-to-sequence prompting to generate one final precursor-set prediction per target. Retrieved conditions supplied five reaction precedents using the representation and structured prompt family described in Supplementary Note 3.

API requests used the settings in Supplementary Table 13.

Supplementary Table 13 | API request settings.

Item	Setting
Temperature	0
Top-p	1.0
Maximum output tokens	65,535
Thinking	Provider-default settings
Incomplete-response handling	Rerun the API request until a complete response was obtained
Scored output	One final precursor set per target
Reported metric	Top-1 complete precursor-set exact match

Only the final precursor set was scored. Predicted and reference molecules were standardized and compared using the common procedure in Methods. Missing predictions, invalid molecular strings and partial precursor-set matches were counted as incorrect.

6.2 Retrieval corpora and no-precedent condition

Each model was evaluated under four conditions, summarized in Supplementary Table 14.

Supplementary Table 14 | Retrieval-resource conditions.

Condition	Retrieval record set	Reactions	Precedents per query
NoPrecedents	None	0	0
RXNGraphormer corpus	`RXNGraphormer-RetrievalCorpus`, constructed from the RXNGraphormer training corpus	6,591,946	5
USPTO corpus	`USPTO-RetrievalCorpus`, assembled by merging the training sets of USPTO-50K, USPTO-MIT and USPTO-Full	1,217,980	5
DianShi corpus	`DianShi-PatentRxnStep-Train-Full`	2,712,259	5

Reaction counts for the RXNGraphormer and USPTO retrieval corpora are given after deduplication. The three retrieval corpora differ in source, scale, temporal coverage and curation; the comparison therefore evaluates retrieval sources as a whole.

6.3 General-purpose model results

Supplementary Table 15 reports Top-1 complete precursor-set exact-match accuracy. Column headings identify the retrieval conditions defined in Section 6.2.

Supplementary Table 15 | General-purpose model accuracy across retrieval conditions (%).

a, USPTO-50K.

Model	NoPrecedents	RXNGraphormer corpus	USPTO corpus	DianShi corpus
glm-5.2	20.17	38.89	46.24	53.45
qwen3.8-MAX	21.69	37.55	44.70	53.43
kimi-k3	34.39	49.83	54.70	59.94
deepseek-v4-pro	29.38	35.85	49.97	60.14
deepseek-v4.1-flash	32.63	41.92	53.33	62.05
gpt-5.6-sol	41.42	42.34	54.54	63.89
gemini-3.8-flash	42.94	49.99	63.45	70.80
Qwen3-4B	0.00	7.29	8.61	13.32
Qwen3-8B	0.14	9.89	10.63	11.60
Qwen3-32B	0.00	15.74	17.70	20.61
Qwen3.5-2B	0.00	2.80	2.48	3.50
Qwen3.5-4B	0.06	10.43	11.58	17.38
Qwen3.5-9B	0.16	12.68	13.82	20.87
Qwen3.5-27B	1.08	23.47	29.86	36.85

b, Test-2001-2024.

Model	NoPrecedents	RXNGraphormer corpus	USPTO corpus	DianShi corpus
glm-5.2	9.98	20.64	22.22	52.79
qwen3.8-MAX	10.80	20.06	22.12	52.87
kimi-k3	17.98	27.96	28.36	58.58
deepseek-v4-pro	16.03	19.48	25.28	59.69
deepseek-v4.1-flash	19.19	22.56	27.80	60.87
gpt-5.6-sol	25.09	24.54	29.91	63.95
gemini-3.8-flash	27.49	30.46	36.89	69.33
Qwen3-4B	0.00	2.55	3.08	14.53
Qwen3-8B	0.03	3.58	3.42	12.43
Qwen3-32B	0.05	7.03	7.14	22.67
Qwen3.5-2B	0.00	1.08	1.05	3.69
Qwen3.5-4B	0.00	4.03	4.00	18.11
Qwen3.5-9B	0.05	4.48	4.82	20.67
Qwen3.5-27B	0.26	10.56	12.64	36.12

c, Test-2025.

Model	NoPrecedents	RXNGraphormer corpus	USPTO corpus	DianShi corpus
glm-5.2	8.17	9.17	11.03	21.84
qwen3.8-MAX	8.75	9.76	10.82	21.58
kimi-k3	15.48	15.32	17.66	27.09
deepseek-v4-pro	13.20	10.13	15.06	29.90
deepseek-v4.1-flash	15.85	11.98	15.32	30.70
gpt-5.6-sol	19.19	11.93	15.91	30.12
gemini-3.8-flash	22.85	18.12	24.13	35.74

Model	NoPrecedents	RXNGraphormer corpus	USPTO corpus	DianShi corpus
Qwen3-4B	0.00	1.48	2.17	8.48
Qwen3-8B	0.00	1.70	1.75	7.16
Qwen3-32B	0.00	3.39	3.87	11.56
Qwen3.5-2B	0.00	1.01	1.01	4.03
Qwen3.5-4B	0.05	1.91	2.01	7.26
Qwen3.5-9B	0.00	2.86	2.97	8.85
Qwen3.5-27B	0.27	5.67	6.26	13.52

DianShi retrieval produced the highest accuracy among the four evaluated conditions in all 42 model–test-set combinations. It also improved every model relative to its no-precedent condition on each test set. Improvements from the external retrieval corpora were less uniform. For example, on Test-2025, gemini-3.8-flash achieved 18.12% with RXNGraphormer retrieval compared with 22.85% without precedents, whereas DianShi retrieval increased accuracy to 35.74%.

The strongest general-purpose model was gemini-3.8-flash with DianShi retrieval, achieving 70.80%, 69.33% and 35.74% on USPTO-50K, Test-2001-2024 and Test-2025, respectively.

6.4 Specialized-model evaluation protocol

We evaluated eight representative released specialized models: LocalRetro[18], Graph2Edits[19], R-SMILES[20], EditRetro[21], NAG2G[22], ReactSeq[5], RXNGraphormer[6] and RSGPT[23].

The evaluated models span template-based (LocalRetro), graph-editing (Graph2Edits, NAG2G), sequence-based (R-SMILES, EditRetro, ReactSeq) and pretrained generative (RXNGraphormer, RSGPT) approaches. These models were used without retraining on DianShi data. Inputs were formatted according to each method's requirements. Inference on Test-2001-2024 and Test-2025 used no molecular test-time augmentation.

Ranked predictions were evaluated using the common complete precursor-set matching procedure in Methods. Top-k accuracy measured whether at least one of the first k ranked candidates matched the complete standardized reference precursor set. Candidates were assessed in their generated rank order. The evaluation protocol is summarized in Supplementary Table 16.

Supplementary Table 16 | Released specialized-model evaluation protocol.

Component	Protocol
Model parameters	Released weights
Retraining on DianShi data	None
Input formatting	According to each method's requirements
Molecular test-time augmentation on Test-2001-2024 and Test-2025	None
Scoring	Common standardized complete precursor-set exact match
Reported ranks	Top-1, Top-3, Top-5 and Top-10
Training-data and compute budget	Not matched across models

6.5 Complete single-step comparison

Supplementary Table 17 reports complete Top-k results for the eight representative released specialized models and the two DianShi models trained on Train-Full. RetroStepLM-Direct received the target product alone; RetroStepLM-RAG additionally received five precedents retrieved from the full DianShi retrieval

corpus. Training and decoding procedures for the DianShi models are described in Supplementary Note 3.

The strongest general-purpose model is included as a Top-1 reference. Dashes in its Top-3, Top-5 and Top-10 columns indicate that these ranks were not evaluated because only one final precursor set was scored per target. Additional USPTO-50K adaptation experiments are reported separately in Supplementary Note 8.

Supplementary Table 17 | Complete single-step performance (%).

a, USPTO-50K.

Model	Top-1	Top-3	Top-5	Top-10
LocalRetro	53.4	77.5	85.9	92.4
Graph2Edits	56.3	79.2	86.2	91.0
R-SMILES	55.1	77.3	83.4	89.4
EditRetro	60.8	80.6	86.0	90.3
NAG2G	55.1	76.9	83.4	89.9
ReactSeq	58.9	80.5	86.4	91.4
RXNGraphormer	51.0	69.0	74.2	79.2
RSGPT	63.4	84.2	89.2	93.0
gemini-3.8-flash + DianShi retrieval	70.80	—	—	—
DianShi-RetroStepLM-Direct	64.3	80.3	85.0	89.2
DianShi-RetroStepLM-RAG	72.8	88.6	92.2	94.9

b, Test-2001-2024.

Model	Top-1	Top-3	Top-5	Top-10
LocalRetro	22.7	35.4	40.8	45.7
Graph2Edits	22.1	33.9	38.8	42.7
R-SMILES	23.8	36.0	41.0	45.4
EditRetro	19.2	29.9	31.8	32.2
NAG2G	23.2	36.0	40.7	45.9
ReactSeq	22.8	35.4	39.7	43.9
RXNGraphormer	19.4	26.9	28.9	29.7
RSGPT	15.1	24.1	28.1	32.0
gemini-3.8-flash + DianShi retrieval	69.33	—	—	—
DianShi-RetroStepLM-Direct	62.0	76.7	80.9	84.3
DianShi-RetroStepLM-RAG	72.4	86.4	89.8	92.7

c, Test-2025.

Model	Top-1	Top-3	Top-5	Top-10
LocalRetro	18.3	29.1	33.7	38.0
Graph2Edits	16.4	26.9	31.3	36.4
R-SMILES	17.0	28.3	32.8	37.6
EditRetro	14.8	22.5	23.8	24.0
NAG2G	16.4	27.8	33.1	38.0
ReactSeq	17.4	27.6	32.3	37.0
RXNGraphormer	12.5	18.8	20.3	21.4
RSGPT	9.7	15.5	18.3	22.3
gemini-3.8-flash + DianShi retrieval	35.74	—	—	—

Model	Top-1	Top-3	Top-5	Top-10
DianShi-RetroStepLM-Direct	36.4	53.6	59.4	64.7
DianShi-RetroStepLM-RAG	39.4	59.1	67.0	74.9

On each test set, gemini-3.8-flash with DianShi retrieval exceeded the Top-1 accuracy of all eight representative released specialized models evaluated here. Full-corpus RetroStepLM-RAG achieved the highest Top-1 accuracy among the predictors reported in this note, combining retrieval-conditioned fine-tuning with inference-time reaction precedents.

9.7 Supplementary Note 7. Multistep benchmarks and planning evaluation

7.1 Multistep benchmarks

Multistep evaluation used three benchmarks: Pistachio Hard[27], USPTO-190[28] and PatMol-150. Pistachio Hard and USPTO-190 are public synthesis-planning benchmarks. PatMol-150 was constructed for this study from 2025 public patents in DianShi-RxnDB. Benchmark composition is summarized in Supplementary Table 18.

Supplementary Table 18 | Multistep benchmark composition.

Benchmark	Source	Targets	Average steps
Pistachio Hard	Public benchmark derived from Pistachio reaction data	100	7.2[27]
USPTO-190	Public benchmark introduced with Retro*	190	6.0[27]
PatMol-150	This study; 2025 public patents	150	8.17

The minimum reference-route length for PatMol-150 is 7 steps.

Pistachio Hard was evaluated without its original specified-starting-material constraint. All three benchmarks were evaluated against the common starting-material inventory described in Section 7.3.

7.2 PatMol-150 construction

PatMol-150 was constructed through single-step reaction cleaning, route-tree assembly, applicant-source screening and final target selection.

Single-step reaction extraction and cleaning. Construction started from 1,300,500 single-step reactions extracted from 2025 public patents. Records were retained only when their SMILES parsed successfully. Filtering removed reactions with repeated molecules within either side, overlap between reactant and product sets, multiple products, product atom mappings absent from the reactants, more than four reactant components or purely inorganic products. Within each patent, duplicate reactions differing only in component order were merged. These steps retained 846,230 reactions. Extraction of local reaction templates, RDChiral[46] representations and synthons was then performed, and records with unsuccessful extraction were discarded, leaving 835,875 reactions.

Route-tree assembly. Following the PaRoutes approach[47], reactions within the same patent Example, or title group, were connected when the product of one reaction appeared as a reactant in another. AND–OR synthesis trees were constructed, cyclic paths were removed, and trees whose root molecule had already appeared in pre-2025 patent data were discarded. Cross-patent connections were subsequently made, and routes were graded by reaction completeness and template commonness. This route-construction stage yielded 24,760 routes covering 2,923 patents.

Applicant-source screening. Candidate targets were deduplicated internally and against products in Train-Full. Applicant screening used the U.S. News 2025–2026 Best Global Universities top 100 for

universities and the 2026 Fortune Global 500 pharmaceutical list and its intersection with the Fortune Global 500 healthcare list for pharmaceutical and healthcare organisations. Applicant identities were checked using Google Patents.

Final target selection. For each patent, routes leading to the final target molecule were retained. Route length was defined as the maximum reaction depth from the target to a leaf node, $d_{\max}$, with a minimum of seven steps. Targets were standardized and deduplicated, then excluded if they appeared among products in Train-Full, targets in USPTO-190 or Pistachio Hard, or molecules in the starting-material inventory.

The final benchmark contained 150 target molecules and 150 patent reference routes. Reference-route depths ranged from 7 to 13 steps, with a mean of 8.17 and a median of 8. Forty-one routes, or 27.3%, had depths of at least nine steps. The depth distribution is reported in Supplementary Table 19.

Supplementary Table 19 | PatMol-150 reference-route depth distribution.

Maximum reaction depth	Routes	Percentage (%)
7	63	42.0
8	46	30.7
9	17	11.3
10	9	6.0
11	9	6.0
12	3	2.0
13	3	2.0
Total	150	100.0

7.3 Starting-material inventory and planning frameworks

All evaluations used a common RDKit-standardized inventory of 23,081,629 commercially available eMolecules building blocks[45]. Inventory membership defined terminal precursor availability during search.

We evaluated RetroAgent[25], Retro-R1-modified[26] and DianShi-RetroPlanSystem-v0. All three used Qwen3.5-397B as the planner model. Within each framework, the planner, inventory, Top-10 single-step candidate setting and search protocol were fixed across single-step model comparisons.

RetroAgent. Evaluations used the RetroAgent planning framework with its default single-step model or one of the alternative single-step predictors listed in Section 7.6. The planner model and framework-specific search protocol remained fixed across these comparisons.

Retro-R1-modified. The original inference protocol was adapted for execution by a general-purpose planner. The system prompt explicitly specified tool syntax and two action types: invoking the single-step model and selecting a candidate reaction. It also provided reaction-selection and search guidance. The complete interaction history was retained, and format reminders were issued when the planner produced no tool call. Following branch failure, the environment traversed untried sibling candidates without requiring an explicit planner-issued backtracking action. This adapted protocol was fixed across single-step model comparisons.

DianShi-RetroPlanSystem-v0. Search was represented as an AND–OR tree. A molecule node closed when at least one candidate reaction closed; a reaction node closed only when all required precursor molecules were available in the inventory or solved through further expansion. The planner selected an expandable molecule, and a deterministic scheduler validated actions, invoked the single-step predictor, inserted candidates, checked inventory membership and propagated closure or failure. Invalid molecular strings and duplicate precursor sets were removed. Candidates containing the current molecule or an ancestor were blocked from cyclic expansion.

7.4 Reaction validation and fallback

The generative single-step models (RetroStepLM-Direct/RetroStepLM-RAG, RSGPT and RXNGraphormer) were filtered with a common AOT*-style template-based reaction-validation procedure[29] after single-step prediction. Candidate reactions were checked using each model's corresponding template library, with template-based rewriting where applicable.

Templates were preselected using product PatternFP fingerprints and a Tversky similarity threshold of 1 ($\alpha = 0$, $\beta = 1$). Product- and reactant-side SMARTS were required to match one-to-one. In our adaptation, we applied up to 256 matching templates with rdchiral to rewrite each candidate. If these applications yielded no acceptable candidate, the procedure considered the 1,000 highest reactant-side fingerprint matches among the product-preselected templates, ranked the resulting reactant sets with the AOT*-style reward, and retained the best; the unsolved penalty was applied only on this fallback path, with $\gamma = 0.1$. Rewritten outputs in which the target product, or any of its components, appeared as a reactant were discarded. These settings are summarized in Supplementary Table 20.

Supplementary Table 20 | Reaction-validation settings.

Component	Setting
Template preselection	Product PatternFP; Tversky similarity 1 ($\alpha = 0$, $\beta = 1$)
Structural matching	One-to-one SMARTS matching on product and reactant sides
Initial template application	Up to 256 matching templates rewritten with rdchiral
Template-search fallback	1,000 highest reactant-side fingerprint matches among product-preselected templates
Fallback ranking	AOT*-style reward on reactant sets; best retained; $\gamma = 0.1$
Unsolved penalty	Applied only on the template-search fallback path
Identity-output handling	Discard outputs in which the target product, or any of its components, appears as a reactant

The validation procedure was shared, whereas template libraries depended on the single-step model. Template-library sources are listed in Supplementary Table 21.

Supplementary Table 21 | Template libraries used in planning.

Single-step model	Template-library source	Entries
RetroStepLM-Direct and RetroStepLM-RAG	Train-Full	412,027
RSGPT and RXNGraphormer	DESP USPTO library	270,794
RetroAgent default model	DESP USPTO library	270,794
Retro-R1 default model	Retro* USPTO library	381,302

Because the Retro-R1 default model used a template library of similar size (381,302 versus 412,027 entries) yet solved 49 percentage points fewer USPTO-190 targets within the same framework, library size alone does not account for the observed differences.

7.5 Search budgets and solve-rate calculation

Performance was assessed as target-level solve rate, defined as the fraction of targets for which at least one complete route, with all terminal precursors in the starting-material inventory, was found within the search budget. Multiple routes for one target counted as one success.

Search budget N had a framework-specific definition (Supplementary Table 22).

Supplementary Table 22 | Search-budget units by framework.

Framework	Budget unit
RetroAgent	Newly expanded distinct molecules
Retro-R1-modified	Single-step tool calls
DianShi-RetroPlanSystem-v0	Expanded molecule nodes; each expansion invoked the single-step model once

Results at N = 50, 100, 200, 300, 400 and 500 were budget levels along the same search trajectory for each target and single-step model. They were not independent reruns or averages across repeated searches. Equal numerical budgets did not imply equal computational cost across frameworks, so single-step model comparisons were made within each framework.

Solve rate measured computational route finding under the stated inventory, validation and search protocol. It did not establish reaction conditions, selectivity, yield or experimental executability.

7.6 Complete planning results

Supplementary Table 23 reports solve rates at all six budget levels and corresponds to Extended Data Fig. 5. *Default* denotes the respective framework's default single-step model. No default-model result is included for DianShi-RetroPlanSystem-v0.

Supplementary Table 23 | Multistep solve rates across search budgets (%).

a, RetroAgent.

Benchmark	Single-step model	$N = 50$	$N = 100$	$N = 200$	$N = 300$	$N = 400$	$N = 500$
Pistachio Hard	Default	56.00	56.00	56.00	56.00	56.00	56.00
Pistachio Hard	RXNGraphormer	47.00	47.00	47.00	47.00	47.00	47.00
Pistachio Hard	RSGPT	76.00	77.00	79.00	79.00	79.00	79.00
Pistachio Hard	RetroStepLM-Direct	91.00	92.00	92.00	92.00	92.00	92.00
Pistachio Hard	RetroStepLM-RAG	86.00	90.00	90.00	90.00	90.00	90.00
USPTO-190	Default	37.89	38.95	38.95	38.95	38.95	38.95
USPTO-190	RXNGraphormer	33.68	33.68	33.68	33.68	33.68	33.68
USPTO-190	RSGPT	65.79	68.95	69.47	69.47	69.47	69.47
USPTO-190	RetroStepLM-Direct	90.00	92.63	93.16	93.16	93.16	93.16
USPTO-190	RetroStepLM-RAG	93.68	95.79	96.84	96.84	96.84	96.84
PatMol-150	Default	37.33	37.33	37.33	37.33	37.33	37.33
PatMol-150	RXNGraphormer	33.33	33.33	33.33	33.33	33.33	33.33
PatMol-150	RSGPT	66.00	68.00	69.33	69.33	69.33	69.33
PatMol-150	RetroStepLM-Direct	85.33	89.33	90.67	90.67	90.67	90.67
PatMol-150	RetroStepLM-RAG	84.67	86.67	89.33	89.33	89.33	89.33

b, Retro-R1-modified.

Benchmark	Single-step model	$N = 50$	$N = 100$	$N = 200$	$N = 300$	$N = 400$	$N = 500$
Pistachio Hard	Default	69.00	69.00	69.00	69.00	69.00	69.00
Pistachio Hard	RXNGraphormer	35.00	36.00	36.00	36.00	36.00	36.00
Pistachio Hard	RSGPT	78.00	78.00	79.00	79.00	79.00	79.00
Pistachio Hard	RetroStepLM-Direct	84.00	84.00	84.00	84.00	84.00	84.00
Pistachio Hard	RetroStepLM-RAG	91.00	92.00	92.00	92.00	92.00	92.00
USPTO-190	Default	43.16	43.16	43.16	43.16	43.16	43.16
USPTO-190	RXNGraphormer	28.42	30.53	30.53	31.05	31.05	31.05
USPTO-190	RSGPT	57.37	57.89	61.05	61.05	61.05	61.05
USPTO-190	RetroStepLM-Direct	83.16	83.68	85.26	85.26	85.26	85.26

Benchmark	Single-step model	$N = 50$	$N = 100$	$N = 200$	$N = 300$	$N = 400$	$N = 500$
USPTO-190	RetroStepLM-RAG	90.00	91.05	92.11	92.11	92.11	92.11
PatMol-150	Default	52.00	52.00	52.00	52.00	52.00	52.00
PatMol-150	RXNGraphormer	22.00	23.33	23.33	23.33	23.33	23.33
PatMol-150	RSGPT	63.33	65.33	67.33	68.00	68.00	68.00
PatMol-150	RetroStepLM-Direct	82.00	83.33	83.33	84.00	84.00	84.00
PatMol-150	RetroStepLM-RAG	83.33	85.33	86.00	86.00	86.00	86.00

c, DianShi-RetroPlanSystem-v0.

Benchmark	Single-step model	$N = 50$	$N = 100$	$N = 200$	$N = 300$	$N = 400$	$N = 500$
Pistachio Hard	RXNGraphormer	39.00	40.00	44.00	44.00	44.00	44.00
Pistachio Hard	RSGPT	76.00	78.00	86.00	88.00	88.00	88.00
Pistachio Hard	RetroStepLM-Direct	88.00	88.00	90.00	93.00	94.00	94.00
Pistachio Hard	RetroStepLM-RAG	89.00	91.00	94.00	95.00	96.00	96.00
USPTO-190	RXNGraphormer	27.37	32.11	34.74	35.79	35.79	36.32
USPTO-190	RSGPT	61.58	65.79	81.05	82.63	83.68	83.68
USPTO-190	RetroStepLM-Direct	88.42	91.58	94.21	96.84	97.37	97.37
USPTO-190	RetroStepLM-RAG	92.11	94.21	96.32	96.84	96.84	96.84
PatMol-150	RXNGraphormer	24.67	26.67	26.67	26.67	26.67	27.33
PatMol-150	RSGPT	68.67	70.67	76.00	80.00	81.33	81.33
PatMol-150	RetroStepLM-Direct	86.67	89.33	94.00	96.00	97.33	97.33
PatMol-150	RetroStepLM-RAG	90.67	91.33	95.33	95.33	95.33	95.33

RetroStepLM-Direct and RetroStepLM-RAG were developed in this work.

Both DianShi models exceeded every evaluated non-DianShi model across all nine framework–benchmark settings at both $N = 50$ and $N = 500$. At $N = 500$, DianShi solve rates ranged from 84.00% to 97.37%. RetroStepLM-Direct solved more targets than RetroStepLM-RAG in four of the nine settings at $N = 500$ and in two of the nine at $N = 50$, despite RetroStepLM-RAG's higher single-step Top-1 accuracy.

9.8 Supplementary Note 8. Additional ablations and retrieval corpus transfer

8.1 Intermediate-label and retrieval ablations

We evaluated combinations of intermediate-label supervision and retrieval conditioning using Qwen3-4B as the common base model. Four models received task-specific training on Train-Full, with intermediate labels and retrieval enabled or omitted. Qwen3-4B without task-specific training or retrieved precedents provided an additional reference.

Models with intermediate labels included the synthon and leaving-group fields described in Supplementary Note 3 in their training targets. Models without intermediate labels omitted these fields. Retrieval-enabled models used precedents from Train-Full during training and inference; models without retrieval received the query product alone. The models with intermediate labels, without and with retrieval, correspond to the full-corpus RetroStepLM-Direct and RetroStepLM-RAG models, respectively.

Supplementary Table 24 reports complete precursor-set exact-match accuracy. The four trained models form a comparison of intermediate-label and retrieval combinations; the untrained reference additionally differs in whether task-specific parameter updates were performed.

Supplementary Table 24 | Intermediate-label and retrieval ablation results (%).

a, USPTO-50K.

Training corpus	Intermediate labels	Retrieval	Top-1	Top-3	Top-5	Top-10
None	No	No	0.0	0.0	0.0	0.0
Train-Full	No	No	61.49	78.37	82.98	86.74
Train-Full	Yes	No	64.3	80.3	85.0	89.2
Train-Full	No	Yes	71.70	86.84	90.71	94.19
Train-Full	Yes	Yes	72.8	88.6	92.2	94.9

b, Test-2001-2024.

Training corpus	Intermediate labels	Retrieval	Top-1	Top-3	Top-5	Top-10
None	No	No	0.0	0.0	0.0	0.0
Train-Full	No	No	58.29	73.85	78.07	82.10
Train-Full	Yes	No	62.0	76.7	80.9	84.3
Train-Full	No	Yes	70.06	84.81	88.47	91.73
Train-Full	Yes	Yes	72.4	86.4	89.8	92.7

c, Test-2025.

Training corpus	Intermediate labels	Retrieval	Top-1	Top-3	Top-5	Top-10
None	No	No	0.0	0.0	0.0	0.0
Train-Full	No	No	35.63	51.48	57.26	63.26
Train-Full	Yes	No	36.4	53.6	59.4	64.7
Train-Full	No	Yes	34.36	54.19	61.40	69.19
Train-Full	Yes	Yes	39.4	59.1	67.0	74.9

Including intermediate labels increased accuracy at every reported rank in both retrieved and non-retrieved trained models. The combination of intermediate labels and retrieval achieved the highest accuracy among these models on all three test sets. Retrieval alone did not improve every metric: without intermediate labels, Test-2025 Top-1 accuracy decreased from 35.63% to 34.36%, although Top-3, Top-5 and Top-10 accuracy increased.

8.2 Training-time precedent sampling

We compared two training-time precedent-selection strategies for full-corpus RetroStepLM-RAG: deterministic selection of the five highest-similarity eligible precedents, and random selection of five precedents from the ten highest-similarity eligible candidates at each epoch. Both models used Train-Full, intermediate labels and deterministic Top-5 retrieval at inference. The comparison changed training-time precedent selection while retaining the common training and evaluation protocol.

These were separately trained models. Their comparison therefore differs from the fixed-checkpoint inference interventions in Supplementary Note 4. Results are reported in Supplementary Table 25.

Supplementary Table 25 | Training-time precedent-sampling results (%).

a, USPTO-50K.

Training-time selection	Top-1	Top-3	Top-5	Top-10
Deterministic Tanimoto Top-5	71.66	87.20	91.33	94.21
Random five from Tanimoto Top-10	72.8	88.6	92.2	94.9

b, Test-2001-2024.

Training-time selection	Top-1	Top-3	Top-5	Top-10
Deterministic Tanimoto Top-5	70.19	84.99	88.84	92.18
Random five from Tanimoto Top-10	72.4	86.4	89.8	92.7

c, Test-2025.

Training-time selection	Top-1	Top-3	Top-5	Top-10
Deterministic Tanimoto Top-5	36.80	54.88	62.25	70.15
Random five from Tanimoto Top-10	39.4	59.1	67.0	74.9

Random selection from the Top-10 candidates produced higher accuracy at every reported rank on all three test sets. This strategy was used for the main RetroStepLM-RAG model. The results describe the evaluated training runs and do not estimate variability across repeated training.

8.3 Additional USPTO-50K adaptation

We additionally evaluated a two-stage RetroStepLM-Direct model. Starting from the model trained on Train-Full, a second training stage adapted the model on USPTO-50K-train. Both the original and adapted models predicted from the query product without retrieved precedents.

This model differs from both the unadapted full-corpus RetroStepLM-Direct model and the USPTO-only RetroStepLM-RAG model described in Section 8.4. Its training history includes DianShi training followed by USPTO-50K adaptation. Results are reported in Supplementary Table 26.

Supplementary Table 26 | RetroStepLM-Direct performance before and after USPTO-50K adaptation (%).

a, USPTO-50K.

Training history	Top-1	Top-3	Top-5	Top-10
Train-Full	64.3	80.3	85.0	89.2
Train-Full → USPTO-50K-train	67.9	86.0	90.9	95.0

b, Test-2001-2024.

Training history	Top-1	Top-3	Top-5	Top-10
Train-Full	62.0	76.7	80.9	84.3
Train-Full → USPTO-50K-train	57.4	76.2	81.7	86.5

c, Test-2025.

Training history	Top-1	Top-3	Top-5	Top-10
Train-Full	36.4	53.6	59.4	64.7
Train-Full → USPTO-50K-train	34.6	52.9	60.9	69.5

Adaptation increased USPTO-50K accuracy at all reported ranks. On both DianShi test sets, Top-1 and Top-3 accuracy decreased, whereas Top-5 and Top-10 accuracy increased. The effect of adaptation therefore depended on both the evaluation set and candidate rank.

8.4 Retrieval corpus transfer for a USPTO-only RetroStepLM-RAG model

A separate RetroStepLM-RAG model was initialized from Qwen3-4B and trained only on USPTO-50K-train using intermediate labels and retrieval-conditioned fine-tuning. USPTO-50K-train supplied both the supervised reactions and training precedents. This model did not undergo task-specific training on Train-Full.

After training, the checkpoint was fixed and evaluated with three inference retrieval corpora: USPTO-50K-train, the DianShi Pre2017 corpus and Train-Full. The model's task-specific training history remained unchanged across these conditions.

USPTO-50K-train here denotes the individual benchmark's training partition. It is distinct from `USPTO-RetrievalCorpus` in Supplementary Note 6, which merges the training sets of USPTO-50K, USPTO-MIT and USPTO-Full. Results are reported in Supplementary Table 27.

Supplementary Table 27 | Retrieval corpus transfer for the USPTO-only RetroStepLM-RAG checkpoint (%).

a, USPTO-50K.

Inference retrieval corpus	Top-1	Top-3	Top-5	Top-10
USPTO-50K-train	53.88	73.66	80.27	85.98
Pre2017	64.81	81.85	86.90	91.13
Full	64.57	81.33	87.00	91.31

b, Test-2001-2024.

Inference retrieval corpus	Top-1	Top-3	Top-5	Top-10
USPTO-50K-train	20.75	31.49	35.49	40.44
Pre2017	37.52	49.87	53.98	58.48
Full	44.92	57.82	61.90	66.14

c, Test-2025.

Inference retrieval corpus	Top-1	Top-3	Top-5	Top-10
USPTO-50K-train	14.48	23.17	27.68	31.97
Pre2017	16.44	26.56	31.28	36.00
Full	21.16	32.24	37.01	42.05

Both DianShi retrieval corpora improved accuracy relative to retrieval from USPTO-50K-train at every reported rank on all three test sets. Full retrieval further improved all reported metrics on Test-2001-2024 and Test-2025 relative to Pre2017 retrieval. On USPTO-50K, the differences between the two DianShi corpora were small and varied by rank.

These results support the utility of updating inference-time retrieval resources when the available corpus differs from the model's task-specific training corpus.