arXiv:2611.03377v1 [cs.AI] 2 Nov 2026

# Same Outputs, Different Scores: Auditing the Unchecked Graders of Multimodal LLM Benchmarks

Ahmad Jarjis Hasan
Independent Researcher
ahmad.hasan2k8@gmail.com

Jonathan Lwowski
Nebius
jonathan.lwowski@nebius.com

## Abstract

*Every multimodal-LLM multiple-choice benchmark ends with a* scoring rule *that converts a model's response, often free-form text rather than a clean option letter, into a right/wrong verdict, buried in its code. Unlike the dataset, this rule is chosen without justification and left unexamined, yet it sets the published number. We reproduce the official scorers of eight peer-reviewed benchmarks on five open vision-language models (VLMs) that fit a single 16 GB GPU, and introduce a blinded human-plus-LLM-judge adjudication (Cohen's $\kappa = 0.956$) that recovers what the model actually answered. The scoring rule alone moves accuracy by up to* ***22 points*** *on identical outputs, yet* exactly 0 *wherever a model emits clean option letters, so it hides on strong-model leaderboards. It is the fallback and not the rambling that does the damage: the benchmark with the* highest *non-compliance we measure has an artifact of exactly zero, because its scorer abstains instead of guessing; a reasoning model with non-compliance* 0.000 *still swings* 9.5 *points with which regex match of its trace the parser keeps. No standard rule is reliable: snapping a non-answer to its closest option credits answers never given yet buries correct ones, abstention discards answers stated in words, and likelihood scoring follows answer-token priors, so a single grade conflates capability with luck rather than measuring either. We release all code, labels, and the protocol, and argue benchmarks should report a non-compliance rate and at least two rules, not one unexamined grade.*

## 1. Introduction

Consider a vision-language model asked whether an industrial part is defective, with options "(A) Yes" and "(B) No." It replies: *"The object is a nail."* The model never answered the question. Yet MMAD [15], an ICLR'25 benchmark, scores this response **correct**: its grader calls `difflib.get_close_matches(response, options, cutoff=0.0)`, which with a zero threshold snaps any string to the lexically nearest option: here "Yes," by character overlap, not base rate. An abstaining grader scores the same response *wrong*. A likelihood grader returns whatever the model's logits prefer. A human reading it records that the model gave *no answer*. Four graders, four verdicts, one output (Fig. 1). The grade came from the grader, not the model.

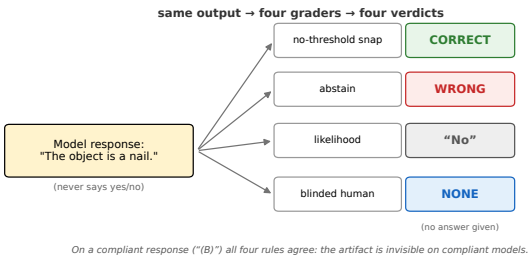


Figure 1. **The grade comes from the grader, not the model.** A single rambling response receives contradictory verdicts under four scoring rules shipped by published benchmarks; none matches what a blinded human says the model answered. On a compliant response the rules agree, which is why the artifact is invisible on leaderboards of strong, format-following models.

This is not a quirk of one benchmark. Every multimodal-LLM multiple-choice benchmark must decide what to do when a model's free-form text does not cleanly name an option, and that decision, a few lines at the bottom of a scoring script, is made differently across the field: a no-threshold snap (MMAD [15], OmniMedVQA [14], PMC-VQA [36]), a random fallback (MMMU [33], ScienceQA [20]), a default class (POPE [18]), abstention, or likelihood scoring. These benchmarks cleared top-tier peer review and are cited widely, but reviewers judged the *datasets*; the grader was an

© 2027 IEEE. Personal use of this material is permitted. Permission from IEEE must be obtained for all other uses, in any current or future media, including reprinting/republishing this material for advertising or promotional purposes, creating new collective works, for resale or redistribution to servers or lists, or reuse of any copyrighted component of this work in other works. Accepted at WACV 2027.

unexamined implementation detail that produced numbers which looked reasonable, so it went unopened. This paper opens it.

**Why the obvious fix is not enough.** "Just parse the letter and mark the rest wrong" (abstention) is one choice among several, and it is not obviously right: a model that answers "No." in words, or reasons correctly and is then cut off before emitting a letter, is marked wrong. Conversely, the no-threshold snap rewards rambles that never commit. The rules are not interchangeable, and which one a benchmark ships changes its leaderboard. We scope this audit to the *locally deployable* regime (open models of $\leq$13B on one consumer GPU), the binding constraint for the applications these benchmarks target: privacy-bound medical VQA and air-gapped industrial inspection, where a frontier API model is not an option. Where a model emits clean letters the rules coincide, so the choice is invisible there; in the regime practitioners are confined to, the scores that guide deployment inherit it.

**What prior work leaves open.** A line of work studies the *format sensitivity* of multimodal-LLM scoring: first-token logits versus decoded text on answers the model *did* give in parseable form [23, 37], and whether models should be allowed to abstain on unanswerable inputs [24]. Standardized harnesses do not dissolve the problem: VLMEvalKit [9], which powers the OpenVLM leaderboard, falls back to `random.choice` once its extraction retries fail and to a fixed sentinel under its parse-only policy: the two idioms we audit, shipped as defaults. The closest prior work anchors answer *extraction* to human judgment for text LLMs and does annotate non-answers, refusals and truncation among them [25, App. A.3]; what it does not do is *grade* them, recording a miss instead, and it varies prompt format at a fixed extractor rather than the scorer at fixed outputs. We audit what the released scorer *credits* when *no* answer is extractable, a different contrast, and find 21–22 points. The reason ours can be this large is mechanical. Text LLMs mostly emit a parseable letter; on the open VLMs practitioners can actually deploy, the fallback fires on 29–37% of items. Multimodal is where the fallback stops being a rounding error and starts setting the number, which is why this audit had to be done here, in the spirit of evaluation reality-checks in other subfields [22, 26].

**Scope of evidence.** That the parse-failure rule *varies and is unreported* across the field is a code audit of eight peer-reviewed benchmarks (Table 1), six of which manufacture a verdict and two of which do not. That *no rule reliably recovers what the model answered* is a stronger claim, demonstrated where we have human ground truth: binary MMAD and four-option medical PMC-VQA. We anchor it on binary MMAD by design, not by limitation: yes/no is where blinded readers agree almost perfectly on what a model said (inter-annotator $\kappa = 0.91$), so it is where the grader's fabrication is *certifiable* rather than asserted. Breadth is carried by the field-wide survey, proof by the binary anchor.

**Contributions.**

- **An audit of released graders** (Sec. 3, 5). We reproduce the official scoring code of eight peer-reviewed benchmarks on five open 16 GB-class VLMs; the parse-failure rule alone shifts accuracy by up to **22 points** (Fig. 2) on identical outputs and exactly 0 on format-compliant models; the artifact is invisible until a model rambles.
- **A human-certified adjudication protocol** (Sec. 4, 5). A blinded human plus LLM-judge procedure recovers what the model actually answered ($\kappa = 0.956$, released), and shows *no* standard rule is reliable, in both directions: the snap credits answers never given on one model yet *deflates* another's correct answers, abstention discards answers given, and likelihood scoring is biased by answer-token priors.
- **A practice recommendation and release** (Sec. 6). We release all code, labels, and the protocol,[1] and recommend benchmarks report a non-compliance rate and at least two scoring rules rather than one unexamined grade.

## 2. Related Work

**Format sensitivity of multimodal-LLM scoring.** Lyu et al. [23] report large disagreements between a model's first-token logit and its decoded free-text answer on multiple-choice questions; others study option-order and prompt robustness and the generative-versus-discriminative gap [37]; Rosenthal et al. [27] further show that semantically-neutral prompt-format changes alone reorder MLLM leaderboards, as does prompt format in text LLMs [28]; Alzahrani et al. [2] perturb the benchmark and the answer-selection method together and move leaderboard rank by up to eight positions. All vary the *input*, or analyze answers produced in *parseable* form. We hold the outputs byte-identical and vary only the parse-*failure* rule, anchoring each to an external human label rather than to another automated rule.

**Answer extraction and standardized evaluation.** A parallel line improves how a *parseable* answer is recovered or standardized: trained extractors that pinpoint the option a model chose [32], standardized prompt-and-parse protocols [12], reproducibility practices for evaluation harnesses [6], and, for text LLMs, evidence that

[1] https://github.com/ahmadhasan2k8/grader-audit

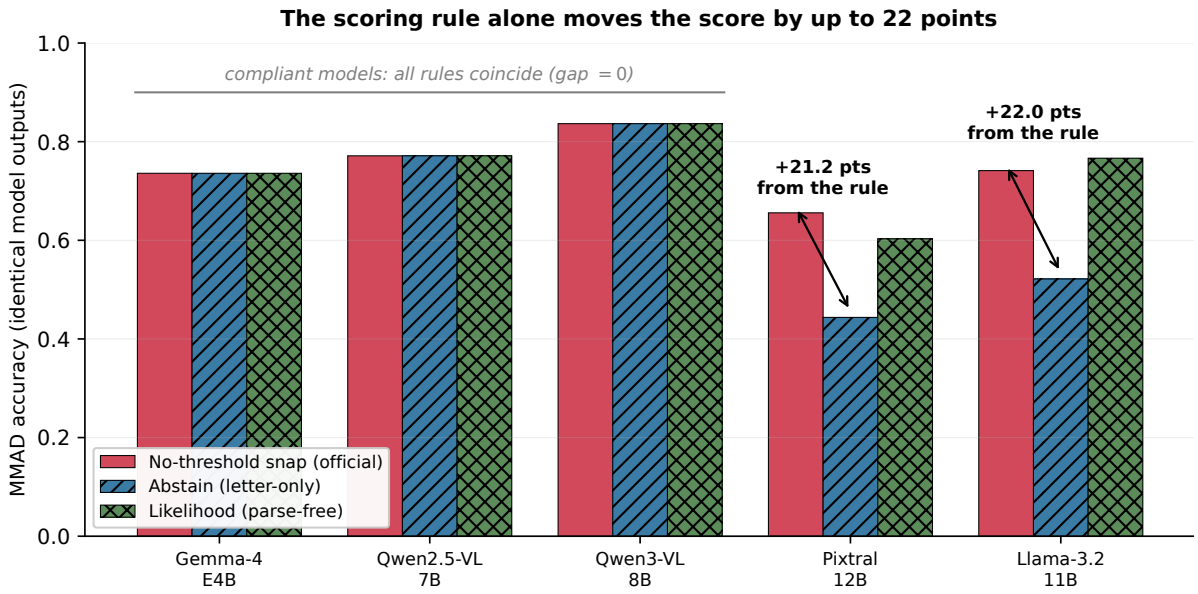


Figure 2. **The scoring rule alone moves the score by up to 22 points.** MMAD accuracy on *identical* stored model outputs under three parse-failure scoring rules. The three format-compliant models score the same under every rule (gap = 0); on the two models that ramble, the no-threshold closest-option snap inflates accuracy by 21–22 points over abstention, with the likelihood rule different again. Exact on all 1,670 items, deterministic given the outputs.

answer-extraction choices disagree with human judgment on MCQA [25], and that the extraction algorithm alone shifts reasoning-model results [16]; in the multimodal setting, MMBench [19] adds an LLM choice-extractor and circular evaluation inside a new benchmark. A stronger response is to abandon the format outright: Chandak et al. [7] show that free-form *answer matching*, graded against a reference by a capable model, reaches human-level agreement where multiple choice does not, and argue evaluation should move off discriminative MCQ. That is a prescription for the benchmarks of the future; ours is a diagnosis of the ones in use today. These make the parser better, the protocol uniform, the benchmark harder, or the format different, but none audits, for existing multimodal benchmarks, what the scorer does when the parser *fails*. We target exactly that residual: the parse-failure fallback, measured against what the model actually answered.

**Abstention and unanswerable inputs.** A separate line argues multimodal LLMs should be allowed to abstain and that benchmarks mismeasure when they force a choice [24]. We share the premise but apply it to the *grader*: the parse-failure fallback is a forced choice imposed by the grading code, independent of whether the model tried to abstain.

**Evaluation reality-checks.** A recurring genre re-examines a trusted evaluation practice and finds it unreliable: GAN metrics under equal compute budgets [22], deep metric-learning gains under fair comparison [26]. The genre is now active for LLM and multimodal evaluation and well represented at datasets-and-benchmarks venues: judge reliability [38], benchmark robustness and format noise [31], multimodal harness reality-checks [35], lexical matching that underestimates generative models against human judgment [17], grading logic that over-credits agents [39], and, closest to us in modality, dataset validity for LVLMs [8] (visual necessity and leakage). We are the reality-check for the *scoring rule* of multimodal-LLM multiple-choice evaluation. Unlike a general reproduction, we hold each benchmark's own official parser fixed and change only the parse-failure fallback, which attributes the measured gap to the scoring rule alone.

**Benchmarks studied.** We audit MMAD [15], MMMU [33], ScienceQA [20], OmniMedVQA [14], PMC-VQA [36], and POPE [18], and use MathVista [21] and MathVision [29] as the cases where an extraction step and an abstaining rule, respectively, defuse the same idiom. All are peer-reviewed.

## 3. The Graders

Let a multiple-choice item have $k$ options and gold answer $g$; the model emits free text $r$. A scoring rule is a function $s(r) \to \{1, \ldots, k\} \cup \{\perp\}$, where $\perp$ means abstain (scored

wrong). The rule has a *parser* that tries to extract a clean option from $r$ and a *fallback* for when the parser fails. Both embed unreported decisions: where the parser looks and how it breaks ties among several candidate letters, and what the scorer does when it finds nothing. Our headline measurements isolate the fallback, which is where these benchmarks differ most and which is almost never justified or reported; Sec. 5.3 shows the parser side is not inert either, and that it is the side the field's move to reasoning models activates. Table 1 lists the released idiom for each benchmark we audit; it is a *code* audit verified at pinned commits, and we quantify the effect (Sec. 5) for the five benchmarks whose official prompt yields a clean reproduction, treating ScienceQA, whose released regex accepts exactly one answer format that our models rarely produce (Sec. 5.3), as a code-presence finding only.

**Three families.** **No-threshold snap** maps $r$ to the lexically closest option with no similarity cutoff, so any string yields an option; it is *directional*, concentrating verdicts on whichever option text is the common shape of an answer (e.g. "Yes" in anomaly detection). **Random fallback** returns a uniform option, undirected, $\approx 1/k$ correct per firing. **Default class** returns a fixed label. Against these, two rules do not manufacture a verdict from an unparseable answer: **abstention** (score wrong) and **likelihood** (never parse text; pick the option the model assigns highest likelihood; unconditional, never gated by parsing, specified in the supplement). A natural reading is that the first three are "bad" and the last two are "safe." Section 5 shows the truth is messier: abstention and likelihood each fail in their own way, and which rule is least wrong is model-dependent.

**The field already edits these rules silently.** Two of the released scorers we audit have changed a scoring decision without reporting it as one. MMAD's thirteen released harnesses share one answer parser but index its output inconsistently: eleven keep the *last* matching letter, two keep the *first*, a difference worth nothing on the models it was built for and 9.5 accuracy points on a reasoning model (Sec. 5.3). MMMU parses the entire response; its successor MMMU-Pro [34] restricts the search to the text after the last ``Answer:'' (`evaluate.py`, whose docstring cites the very MMMU file it diverges from), a change announced in neither paper. Scoring rules are revised in practice, just not in print.

## 4. Setup

**Models and reproduction.** We use five current open vision-language models that each run 4-bit on a single 16 GB GPU: Gemma-4-E4B [10], Qwen2.5-VL-7B [3], Qwen3-VL-8B [4], Pixtral-12B [1], and Llama-3.2-11B-Vision [11]. For each benchmark we clone the official repository at a pinned commit and call its released scoring function on our models' outputs; we do not reimplement the scorer. Non-compliance is not an artifact of under-specified prompting: MMAD's released prompt already ends "Answer with the option's letter from the given choices directly," and Pixtral still fails to emit a letter on $489/1{,}670$ items ($29\%$), once replying "Here's the option's letter from the choices directly." We generate every response with greedy decoding (temperature 0), so the stored outputs are deterministic and carry no sampling variance; this is what licenses our "same outputs, every rule" design, since all rule differences are then attributable to the scorer alone. Models are loaded in 4-bit (bitsandbytes NF4, `bfloat16` compute) with a generation budget of 16 tokens on the letter-answer benchmarks and up to 768 on the chain-of-thought ones. ScienceQA is the one benchmark whose non-compliance is a *format* mismatch rather than a failure to answer (Sec. 5.3), so we treat its magnitude as a code-presence finding only (Sec. 3). Elsewhere non-compliance is a property of the models, not the budget: quadrupling it to 64 tokens (Tab. 2) leaves MMAD non-compliance nearly flat and the snap artifact large, because these models describe the image or answer in words rather than emit a letter. As a validity check we score our stored outputs twice, once with MMAD's *own* released harness (`caculate_accuracy_mmad`) and once with our wrapper, and their per-item verdicts agree to four decimals on four of five models (its printed headline further macro-averages the normal/anomaly split; we report plain accuracy). The fifth, Pixtral, differs by 0.2 points because the harness keeps the *last* regex match where our wrapper keeps the first: immaterial here, since only Pixtral ever emits two option letters (19 of 8,350 responses) and the artifact is identical under both conventions, but the same unreported choice that decides 9.5 points on a reasoning model (Sec. 5.3).

**The adjudication protocol.** To anchor the rules to ground truth we adjudicate what the model actually answered. For each item we build a *blinded* record containing only the question, the options, and the model's full response; the gold answer and every rule's verdict are withheld. A reader assigns one label: a letter (the model committed to that option, by letter or by stating its text), NONE (it never committed: it described, hedged, refused, or ran out), or MULTIPLE (contradictory). The join from saved responses to questions is verified by matching the gold sequence (it aborts on mismatch).

The 130 MMAD items we adjudicate comprise 110 non-compliant cases (where the rules disagree) and 20 compliant controls (55 and 10 per model), drawn from the two models that ramble (Pixtral, Llama). **Two** human annota-

| Benchmark | Venue | Parse-failure fallback (released code) |
|---|---|---|
| MMAD [15] | ICLR'25 | no-threshold snap (`get_close_matches`, cutoff 0) |
| OmniMedVQA [14] | CVPR'24 | no-threshold snap (`SequenceMatcher`) |
| PMC-VQA [36] | Comm. Med.'24 | no-threshold snap (`SequenceMatcher`) |
| MMMU [33] | CVPR'24 | `random.choice` over options |
| ScienceQA [20] | NeurIPS'22 | `random.choice` over options |
| POPE [18] | EMNLP'23 | default class ("yes") |
| MathVista [21] | ICLR'24 | snap, *preceded by* an LLM extraction step |
| MathVision [29] | NeurIPS'24 D&B | abstain: no fallback, unparseable scored wrong |

Table 1. **The parse-failure fallback is an unreported design choice that varies across peer-reviewed benchmarks.** We read each released scorer to recover the rule; none of these papers reports the accuracy its fallback is responsible for, which is the quantity we measure. The bottom two rows are the counterexamples that make the choice visible: an extraction buffer and an abstaining rule each defuse the idiom (Sec. 5.3).

Table 2. **The artifact is robust to the generation budget.** Quadrupling the MMAD budget from 16 to 64 tokens leaves both models' non-compliance above 0.26 and the snap−abstain artifact above 17 points; the rambling is a model behavior, not a truncation artifact. Full 1,670-item set.

| | Non-compliance | | Snap−abstain | |
|---|---|---|---|---|
| Model | 16 | 64 | 16 | 64 |
| Pixtral-12B | 0.29 | 0.27 | +0.212 | +0.183 |
| Llama-3.2-11B | 0.37 | 0.33 | +0.220 | +0.171 |

Table 3. **Adjudication agreement.** The LLM judge tracks the human anchor, and the rubric is stable across two independent humans. On medical PMC-VQA the raw $\kappa$ is a notch lower, but almost every disagreement is a NONE-vs-MULTIPLE boundary call on a ramble; collapsing those two non-committal labels into the one distinction the decomposition uses raises $\kappa$ to 0.95.

| Comparison | Benchmark | $k$ | $n$ | Cohen's $\kappa$ |
|---|---|---|---|---|
| Judge vs. human | MMAD | 2 | 130 | 0.956 |
| Human vs. human | MMAD | 2 | 40 | 0.910 |
| Judge vs. human | PMC-VQA | 4 | 60 | 0.839 |

tors, one author and one annotator naive to the hypothesis, label them independently; on a 40-item overlap they agree on 38 (95%, inter-human Cohen's $\kappa_{\text{HH}} = 0.91$), establishing that the letter/NONE/MULTIPLE rubric is well-posed and not an artifact of one reader. A blinded LLM judge (Claude Opus 4.8, temperature 0, same rubric) labels the same items so it can scale beyond the human subset, and agrees with the human label on 127 of 130 (97.7%, *judge-vs-human* Cohen's $\kappa = 0.956$); the three disagreements are boundary cases (a hedge, a contradictory echo, a buried commitment). We release the manifest, both label sets, the rubric, and the code; the judge is a frontier model used only as a research-time labeling aid on this released data, not a deployment-time dependency, and anchoring the decomposition to the human rather than the judge leaves it unchanged (Sec. 5.3). A rule's accuracy is then decomposed against the human label into *over-credit* (it scores correct on an item the model did not answer with the gold option) and *under-credit* (it scores wrong on an item the model did answer correctly). We extend the same protocol to a second domain, medical PMC-VQA ($k \approx 4$, Sec. 5.3); Table 3 summarizes agreement across both.

# 5. Results

## 5.1. The scoring rule alone moves the score

On identical model outputs, swapping only the parse-failure fallback changes accuracy by up to **22 percentage points** (Fig. 2); every "point" in this paper is a percentage point of accuracy. On MMAD the no-threshold snap scores Pixtral-12B at 0.656 and Llama-3.2-11B at 0.741, while abstaining scores them at 0.444 and 0.522, a +0.212 (Wilson 95% CI $[0.193, 0.232]$) and +0.220 (95% CI $[0.201, 0.240]$) gap on the full $n$=1,670 set. On the three models that emit clean letters the gap is *exactly* 0: the artifact is invisible until a model rambles, and it rides on each paper's specific models' compliance. The effect is smaller but present elsewhere (MMMU +0.04 on the model that rambles most, Gemma, and +0.09 at its largest, Llama's; OmniMedVQA up to +0.08; POPE ≈ +0.01, an easy task where almost everything is compliant); Table 4 gives every benchmark–model cell we measure. The magnitude is set by the fallback *type* and the option count, not by task difficulty: at matched non-compliance (MMAD Pixtral 0.29, MMMU Gemma 0.27) the swing is 3–5× larger on the *easier* binary task (+0.212 vs +0.040–0.071), because a directional two-option snap concentrates on the high-base-rate option while a $k$-way random fallback disperses at $\approx 1/k$, so the biggest artifact

lands on the simplest benchmark. A 3-point shift looks minor beside MMAD's 22, but on a competitive benchmark it is often the entire gap between the top model and the fifth-ranked one, so an unreported fallback can silently invalidate a comparative claim, and the cost falls on the specialist models these benchmarks exist to serve.

### 5.2. No rule recovers ground truth

Anchored to the human label ($\kappa = 0.956$, Sec. 4), each rule fails differently. On the 110 non-compliant MMAD items, the model *committed* to an option only 23% of the time, yet the snap scored 72% "correct": it credits answers the model never gave. Abstention has the opposite failure: on the 17% of those items where the model *did* state the gold answer in words ("No."), the letter-only rule marks it wrong. Likelihood scoring fails a third way (Sec. 5.3). There is no neutral default: the snap over-credits, abstention under-credits, and the gap between them on MMAD (0.21) is a gap between two flawed rules, not a measurement of the model. The 22-point figure needs no adjudication: it is exact on all 1,670 items; adjudication supports only the stronger claim that neither rule recovers what the model said. The two rules also *span* the scores a reporting policy can produce on the full set: $[0.444, 0.656]$ for Pixtral, $[0.522, 0.741]$ for Llama. We report that as a policy-sensitivity interval, not a bound on adjudicated accuracy: the snap under-credits even on MMAD (Table 5), the lower end assumes a parsed letter *is* the committed answer, which Sec. 5.3(v) shows can fail, and on PMC-VQA the ordering inverts. A gallery of real responses with each rule's verdict and the blinded human label is in the supplement.

### 5.3. The error has no safe direction, and compliance does not fix it

Five further results show no rule is a safe default. **(i) The same scorer adds 22 points to one model and removes 19 from another** (Fig. 3)**.** PMC-VQA's matcher is *unconditional* (it runs on every response), so besides filling parse failures it can override an answer the parser already read. Separating the two on identical outputs, the fallback alone adds $+0.097$ to Pixtral and $+0.020$ to Llama, while the *override* subtracts 0.210 from Llama, mapping onto a different option the answer its own parser reads as *gold* on $153/600$ items (25.5% of the benchmark). The aggregate $-0.190$ is the sum, and the larger term is not the fallback, a distinction our other benchmarks never force, because their matchers run only when the parser fails. Anchored to the human label the fallback is bidirectional too. A 60-item audit (judge–human $\kappa = 0.84$, 0.95 on the commit-vs-no-commit split this turns on) confirms a real per-item effect: the snap *under-credits* Llama on 31% of items, marking wrong medical answers it actually gave, while over-crediting Pixtral on 10%; the split is unchanged

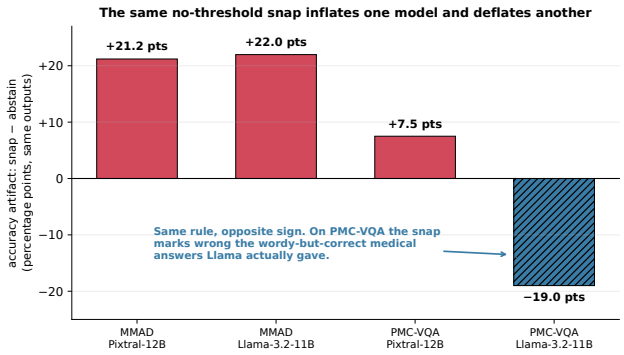


Figure 3. **The official scorer adds 22 points to one model and removes 19 from another.** Signed accuracy swing (official scorer − clean parse) on our two human-anchored benchmarks. On MMAD this is *entirely* the parse-failure fallback: the two rules agree on every compliant item, so $+0.212/+0.220$ is attributable to the fallback alone. On PMC-VQA the matcher is unconditional and the total decomposes into a fallback term ($+0.097$ Pixtral, $+0.020$ Llama) and an override term ($-0.022$, $-0.210$); Llama's $-0.190$ is dominated by the latter (Sec. 5.3). Per-item mechanism in Tab. 5.

whether unaudited items anchor to the human or the judge (0.29→0.31). "Unreliable," not "optimistic." Because the sign of the error flips with the model's verbosity and the dataset's option format, the grader acts as a *confound*: it injects variance tied to response length rather than capability, so one unexamined grade measures neither cleanly. A weak probe agrees: on items where a human found no single committed option, the model's own option likelihood prefers gold at chance for Pixtral ($26/52$) but above chance for Llama ($24/34$, two-sided $p$=0.024), so even the *character* of the snap's unearned credit is model-dependent; the single-token logit is itself token-prior biased [13, 30] (ii), so we do not read this as a capability claim. **(ii) The principled rule can be the most wrong: 0.403 against the snap's 0.565.** Parse-free likelihood is not a clean ground truth: on OmniMedVQA, Llama's option-letter logit collapses onto the token "A" on $459/600$ items while gold is "A" only 29% of the time, giving accuracy 0.403, *below* the snap's 0.565. Any single rule, including the principled one, can be the most wrong. **(iii) An extraction buffer defuses the idiom: +10 points collapses to +0.9.** MathVista runs an LLM extraction step before the same no-threshold snap; scoring the snap on raw responses yields a spurious +10-point artifact, but the faithful pipeline (extraction first, reproduced with a local model) collapses it to +0.9 points, evidence that the fix is a competent parse, not a different fallback. **(iv) One character of unreported regex is worth 29.5 points.** ScienceQA is the purest case, and not for the reason we first assumed: its released regex demands the literal `The answer is X.`, and the mod-

Table 4. **The parse-failure artifact across every benchmark and model we audit.** Each cell is the accuracy swing attributable to the parse-failure rule alone, with the rate at which that rule fires (non-compliance) in parentheses, computed on *identical* stored outputs under each benchmark's own released scorer. "Non-compliance" is in each case the rate at which that benchmark's own scorer reaches its fallback. Because each cell isolates the fallback, no artifact can exceed its firing rate, and none does. MMMU and ScienceQA cells are one seeded draw of their own `random.choice` fallback, whose expectation reorders MMMU's models (supplement); MMAD's snap is deterministic. Two regimes appear. **Firing:** the fallback manufactures a verdict and moves the score. **Defused:** an extraction buffer (MathVista) or an abstaining rule (MathVision) means the scorer never fabricates. MathVision is the control that separates prevalence from harm: it has the *highest* non-compliance of any benchmark here (0.54–0.69) and an artifact of *exactly* zero, because its scorer abstains instead of guessing; zero artifact shows the fallback adds no credit, not that extraction is faithful. Rambling is not the bug; the fallback is.

| Benchmark | Regime | Gemma-4-E4B | Qwen2.5-VL-7B | Qwen3-VL-8B | Pixtral-12B | Llama-3.2-11B |
|---|---|---|---|---|---|---|
| MMAD | firing | +0.000 (0.00) | +0.000 (0.00) | +0.000 (0.00) | +0.212 (0.29) | +0.220 (0.37) |
| MMMU | firing | +0.040 (0.27) | +0.007 (0.04) | +0.047 (0.17) | +0.027 (0.16) | +0.087 (0.21) |
| OmniMedVQA | firing | +0.005 (0.01) | +0.013 (0.015) | +0.010 (0.01) | +0.083 (0.39) | +0.018 (0.06) |
| PMC-VQA‡ | firing | +0.015 (0.04) | +0.000 (0.00) | +0.002 (0.01) | +0.097 (0.61) | +0.020 (0.12) |
| POPE | firing | +0.017 (0.05) | +0.000 (0.00) | +0.003 (0.003) | +0.025 (0.05) | +0.003 (0.003) |
| ScienceQA† | firing | +0.062 (0.20) | +0.070 (0.23) | +0.027 (0.07) | +0.118 (0.34) | +0.303 (0.84) |
| MathVista | defused | +0.004 (0.03) | +0.011 (0.03) | +0.004 (0.01) | +0.017 (0.04) | +0.007 (0.04) |
| MathVision | defused | +0.000 (0.61) | +0.000 (0.56) | +0.000 (0.61) | +0.000 (0.69) | +0.000 (0.54) |

†ScienceQA's released regex requires the literal string `The answer is X.`; the models overwhelmingly answer in a near-miss format instead (Sec. 5.3), so the magnitude is exact on these outputs but the *firing rate* would differ under a different prompt. ‡PMC-VQA's matcher is unconditional: besides the fallback tabulated here, it also *overrides* answers its own parser reads cleanly, a second and larger failure on one model (Sec. 5.3).

Table 5. **The snap's error is bidirectional and model-dependent.** Human-anchored over- vs. under-credit on adjudicated items (MMAD anchored to the human label, $n$=65; medical PMC-VQA to the certified judge, $n$=48). "Gave gold" counts adjudicated items committed to the gold option. On binary MMAD the snap over-credits heavily; on medical PMC-VQA the same rule *under-credits* Llama on 15/48 items, 11 of them items where the fallback itself fired. Wilson 95% CIs: MMAD Pixtral over-credit 62% [49, 72], PMC-VQA Llama under-credit 31% [20, 45]. Within PMC-VQA Llama the direction itself is significant (15 under-credits against 4 over-credits, exact binomial $p$=0.019), so the sign flip is not a small-sample artifact.

| Benchmark | Model | $n$ | Gave gold | Over | Under |
|---|---|---|---|---|---|
| MMAD | Pixtral-12B | 65 | 11 | 40 (62%) | 0 (0%) |
| MMAD | Llama-3.2-11B | 65 | 23 | 25 (38%) | 4 (6%) |
| PMC-VQA | Pixtral-12B | 48 | 4 | 5 (10%) | 2 (4%) |
| PMC-VQA | Llama-3.2-11B | 48 | 17 | 4 (8%) | 15 (31%) |

els answer `The answer is (X).` instead. Of Llama's 507 parse failures, 460 contain that parenthesised form and *none* is truncated (median response 110 characters). The scorer therefore discards answers that are present and correct, and randomly fabricates in their place: accepting both spellings (abstaining on the residual failures) raises Llama from 0.448 to 0.743 on identical stored outputs (≈0.80 keeping that fallback instead). A single character of unreported regex is worth 29.5 points to one model, which is the thesis of this paper in its least ambiguous form. **(v) Non-compliance 0.000, and the grade still moves 9.5 points.** Every scorer we audit applies its parser to the *entire* response: harmless when models emitted two characters, not once they reason first. On MMAD, Qwen3-VL-8B-Thinking is fully format-compliant: non-compliance 0.000, the parse-failure fallback fires *zero* times, the artifact this paper measures is exactly 0, and it closes its `</think>` block within budget on 99.8% of 1,670 items. Yet its trace restates MMAD's option list before answering, so *which* regex match the scorer keeps now decides the grade. MMAD's eleven majority harnesses keep the last and recover the model's stated answer on 1666/1666 items; the two that keep the first disagree with it on 15.9% and score identical outputs at 0.689 rather than 0.784, a 9.5-point gap (category-clustered 95% CI [6.4, 13.1], 15 MVTec [5] categories). The benchmark ships both numbers and reports neither choice. This is not a capability claim: correctly scored, the reasoning model (0.784) still trails its instruction-tuned sibling (0.837) on these items. It is the sharper form of our recommendation: a non-compliance rate of 0 certifies that nothing was fabricated, not that the grade is safe, a pattern Jo et al. [16] document for text LLMs and we find in shipped multimodal scoring code. MathVision [29] makes the same point from the opposite direction, abstaining rather than snapping at the highest non-compliance we measure (Tab. 4): what converts rambling into a corrupted score is the fallback, not the difficulty.

**What we do not claim.** Holding each benchmark's parser fixed and varying only the fallback, we find no robust leaderboard reordering: the strongest model stays top-ranked, and a top-1 change on MMMU under abstention is not significant at our sample size ($p = 0.58$). Where the released *parser* is itself the variable the picture differs: ac-

cepting ScienceQA's near-miss format moves Llama from last to fourth, which is one reason we report ScienceQA as a code-presence finding rather than folding it into the headline. The contribution is the unreliability of the grade, not a corrected ranking: stable ranks are the weaker guarantee, and the same ordering sits on a 21-point spread in the number practitioners deploy by and later papers must match. Nor does a shared harness dissolve it: VLMEvalKit ships the same random fallback (Sec. 1), and the thirteen harnesses inside MMAD alone disagree with each other by 9.5 points (item v).

## 6. Recommendation

No single scoring rule reflects what a model knew or said, so a benchmark should not hide its choice of rule inside a scoring script and report one number. We recommend three low-cost practices.

1. **Report the non-compliance rate** (the fraction of outputs on which the parse-failure fallback fires) alongside accuracy. It is one number, and it is the single statistic that exposes how much of a score the grader manufactured; when it is large, the headline accuracy is not comparable across papers. It is necessary but not sufficient, and reasoning models are where the difference starts to matter: a model can be perfectly compliant (non-compliance 0, the fallback never firing) while its grade still turns on where in the response the parser looks (Sec. 5.3). Report the rate, and state the parser's scope and tie-break alongside it.
2. **Report at least two rules.** Publish accuracy under both an abstaining rule (crediting only parseable letters) and the closest-option snap, rather than a single grade under either alone. On MMAD, where adjudication shows the snap rarely under-credits, the spread approximates the range of defensible scores; on PMC-VQA it inverts, the snap scoring Llama *below* the abstaining rule. Unconditionally, the spread is the rule-sensitivity a single number hides, and a spread that inverts is itself the warning. Likelihood scoring is a useful parse-free cross-check, not a safe default: it too can be the most-wrong rule (Sec. 5.3).
3. **Adjudicate when it matters.** Our blinded human-plus-LLM-judge protocol ($\kappa = 0.956$) recovers what the model actually answered for roughly one person-day per benchmark, and the certified judge can then label the full non-compliant set automatically, at no further human cost. We release the **adjudication corpus**: one stored generation per item for all eight benchmarks and five models, 130 blinded MMAD adjudications (40 dual-annotated by a reader naive to the hypothesis, $\kappa_{\mathrm{HH}} = 0.91$), 60 blinded PMC-VQA adjudications, and the certified judge's labels on the audited items, with the manifest builder, blinded MCQ tool, judge, and analysis code, under permissive licenses (MIT / CC-BY-4.0), with a datasheet and one-command containerized reproduction, so any benchmark author can audit what their grader credits.

MathVista already shows the remedy works in the wild: an LLM extraction step placed before the snap collapses a spurious $+10$-point artifact to $+0.9$ (Sec. 5.3). Our protocol operationalizes and certifies that step rather than inventing one.

None of these is novel machinery: abstention, likelihood, and LLM-as-judge are all in use. The contribution is the evidence that the choice between them is consequential and currently unexamined, and a certified, released procedure to make it.

## 7. Limitations

We state the boundaries plainly. (i) The human-anchored ground truth is strongest on *binary* MMAD, the easiest case for both the artifact and the adjudication. We extend the certified judge to medical PMC-VQA ($k{\approx}4$), where a 60-item audit gives judge–human $\kappa = 0.84$ overall and 0.95 on the commit-vs-no-commit distinction the decomposition uses (residual disagreement is almost entirely NONE-vs-MULTIPLE on rambles, treated identically by the decomposition), so the cross-benchmark claim does not rest on a single anchor. (ii) Outside MMAD and code-presence-only ScienceQA, the certified fallback magnitudes are typically 1–5 points, reaching 8–10 on the models that ramble; PMC-VQA's override term reaches $-21$ and ScienceQA's parser effect 29.5; the contribution is the unreliability of the grade, not its size. (iii) Varying only the fallback, we find no robust ranking change; varying the released *parser* does move ranks (Sec. 5.3), reported but not built on. (iv) The single-token option logit is a weak probe of "latent knowledge"; we use it only to show the snap's credit is model-dependent, and report where it is itself confounded (Sec. 5.3). (v) We test five open 7–12B models 4-bit on a 16 GB GPU; we measure neither proprietary nor much larger models; an 8-bit control does not restore compliance (supplement), and the artifact is a property of the scorer *given* the outputs. Compliance is not an escape route: Sec. 5.3(v) reports a model with non-compliance 0.000 whose grade still moves 9.5 points, so the failure does not retire as models strengthen: it migrates from the fallback to the parser.

## 8. Conclusion

A benchmark score is read as a measurement of the model. But when a model rambles, as the locally deployable models these benchmarks target routinely do, part of that score is manufactured by a grading choice no paper wrote down. Scores are therefore not comparable across papers that

chose differently, and the remedy is nearly free: state the rule, report how often it fires, publish more than one number. We give a certified procedure for checking what a grader credits, and hope the next benchmark affords the last lines of its scoring code the scrutiny it affords its data: the dataset is not the only thing worth reviewing.

## References


[1] Pravesh Agrawal, Szymon Antoniak, Emma Bou Hanna, Baptiste Bout, Devendra Chaplot, Jessica Chudnovsky, Diogo Costa, Baudouin De Monicault, Saurabh Garg, Theophile Gervet, et al. Pixtral 12B. *arXiv preprint arXiv:2410.07073*, 2024. 4

[2] Norah Alzahrani, Hisham Abdullah Alyahya, Yazeed Alnumay, Sultan AlRashed, Shaykhah Alsubaie, Yousef Almushayqih, Faisal Mirza, Nouf Alotaibi, Nora Al-Twairesh, Areeb Alowisheq, M Saiful Bari, and Haidar Khan. When benchmarks are targets: Revealing the sensitivity of large language model leaderboards. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 13787–13805, 2024. 2

[3] Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibo Song, Kai Dang, Peng Wang, Shijie Wang, Jun Tang, et al. Qwen2.5-VL technical report. *arXiv preprint arXiv:2502.13923*, 2025. 4

[4] Shuai Bai et al. Qwen3-VL technical report. *arXiv preprint arXiv:2511.21631*, 2025. 4

[5] Paul Bergmann, Michael Fauser, David Sattlegger, and Carsten Steger. MVTec AD – a comprehensive real-world dataset for unsupervised anomaly detection. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 9584–9592, 2019. 7

[6] Stella Biderman, Hailey Schoelkopf, Lintang Sutawika, Leo Gao, Jonathan Tow, et al. Lessons from the trenches on reproducible evaluation of language models. *arXiv preprint arXiv:2405.14782*, 2024. 2

[7] Nikhil Chandak, Shashwat Goel, Ameya Prabhu, Moritz Hardt, and Jonas Geiping. Answer matching outperforms multiple choice for language model evaluation. *arXiv preprint arXiv:2507.02856*, 2025. 3

[8] Lin Chen, Jinsong Li, Xiaoyi Dong, Pan Zhang, Yuhang Zang, Zehui Chen, Haodong Duan, Jiaqi Wang, Yu Qiao, Dahua Lin, and Feng Zhao. Are we on the right way for evaluating large vision-language models? In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. 3

[9] Haodong Duan, Junming Yang, Yuxuan Qiao, Xinyu Fang, Lin Chen, Yuan Liu, Xiaoyi Dong, Yuhang Zang, Pan Zhang, Jiaqi Wang, Yu Qiao, Dahua Lin, and Kai Chen. VLMEvalKit: An open-source toolkit for evaluating large multi-modality models. In *Proceedings of the 32nd ACM International Conference on Multimedia (MM)*, 2024. 2

[10] Gemma Team. Gemma 4 technical report. *arXiv preprint arXiv:2607.02770*, 2026. 4

[11] Aaron Grattafiori et al. The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024. 4

[12] Yuling Gu, Oyvind Tafjord, Bailey Kuehl, Dany Haddad, Jesse Dodge, and Hannaneh Hajishirzi. OLMES: A standard for language model evaluations. In *Findings of the Association for Computational Linguistics: NAACL*, 2025. 2

[13] Ari Holtzman, Peter West, Vered Shwartz, Yejin Choi, and Luke Zettlemoyer. Surface form competition: Why the highest probability answer isn't always right. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 7038–7051, 2021. 6

[14] Yutao Hu, Tianbin Li, Quanfeng Lu, Wenqi Shao, Junjun He, Yu Qiao, and Ping Luo. OmniMedVQA: A new large-scale comprehensive evaluation benchmark for medical LVLM. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 1, 3, 5

[15] Xi Jiang, Jian Li, Hanqiu Deng, Yong Liu, Bin-Bin Gao, Yifeng Zhou, Jialin Li, Chengjie Wang, and Feng Zheng. MMAD: A comprehensive benchmark for multimodal large language models in industrial anomaly detection. In *International Conference on Learning Representations (ICLR)*, 2025. 1, 3, 5

[16] Hwiyeol Jo, Joosung Lee, Jaehong Lee, Sang-Woo Lee, Joonsuk Park, and Kang Min Yoo. Finding answers in thought matters: Revisiting evaluation on large language models with reasoning. *arXiv preprint arXiv:2510.14773*, 2025. 3, 7

[17] Ehsan Kamalloo, Nouha Dziri, Charles L. A. Clarke, and Davood Rafiei. Evaluating open-domain question answering in the era of large language models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 5591–5606, 2023. 3

[18] Yifan Li, Yifan Du, Kun Zhou, Jinpeng Wang, Wayne Xin Zhao, and Ji-Rong Wen. Evaluating object hallucination in large vision-language models. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023. 1, 3, 5

[19] Yuan Liu, Haodong Duan, Yuanhan Zhang, Bo Li, Songyang Zhang, Wangbo Zhao, Yike Yuan, Jiaqi Wang, Conghui He, Ziwei Liu, Kai Chen, and Dahua Lin. MMBench: Is your multi-modal model an all-around player? In *European Conference on Computer Vision (ECCV)*, 2024. 3

[20] Pan Lu, Swaroop Mishra, Tony Xia, Liang Qiu, Kai-Wei Chang, Song-Chun Zhu, Oyvind Tafjord, Peter Clark, and Ashwin Kalyan. Learn to explain: Multimodal reasoning via thought chains for science question answering. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. 1, 3, 5

[21] Pan Lu, Hritik Bansal, Tony Xia, Jiacheng Liu, Chunyuan Li, Hannaneh Hajishirzi, Hao Cheng, Kai-Wei Chang, Michel Galley, and Jianfeng Gao. MathVista: Evaluating mathematical reasoning of foundation models in visual contexts. In *International Conference on Learning Representations (ICLR)*, 2024. 3, 5

[22] Mario Lucic, Karol Kurach, Marcin Michalski, Sylvain Gelly, and Olivier Bousquet. Are GANs created equal? a large-scale study. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2018. 2, 3

[23] Chenyang Lyu, Minghao Wu, and Alham Fikri Aji. Beyond probabilities: Unveiling the misalignment in evaluating large language models. In *Workshop on Towards Knowledgeable Language Models (KnowLLM), ACL*, 2024. 2

[24] Atsuyuki Miyai, Jingkang Yang, Jingyang Zhang, Yifei Ming, Qing Yu, Go Irie, Yixuan Li, Hai Li, Ziwei Liu, and Kiyoharu Aizawa. Unsolvable problem detection: Robust understanding evaluation for large multimodal models. In *Annual Meeting of the Association for Computational Linguistics (ACL)*, 2025. 2, 3

[25] Francesco Maria Molfese, Luca Moroni, Luca Gioffré, Alessandro Scirè, Simone Conia, and Roberto Navigli. Right answer, wrong score: Uncovering the inconsistencies of LLM evaluation in multiple-choice question answering. In *Findings of the Association for Computational Linguistics: ACL*, 2025. 2, 3

[26] Kevin Musgrave, Serge Belongie, and Ser-Nam Lim. A metric learning reality check. In *European Conference on Computer Vision (ECCV)*, 2020. 2, 3

[27] Fabio Rosenthal, Sebastian Schmidt, Thorsten Graf, Thorsten Bagdonat, Stephan Günnemann, and Leo Schwinn. Unexplored flaws in multiple-choice VQA make benchmarking unreliable. In *Findings of the Association for Computational Linguistics: EMNLP*, 2026. 2

[28] Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. Quantifying language models' sensitivity to spurious features in prompt design. In *International Conference on Learning Representations (ICLR)*, 2024. 2

[29] Ke Wang, Junting Pan, Weikang Shi, Zimu Lu, Houxing Ren, Aojun Zhou, Mingjie Zhan, and Hongsheng Li. Measuring multimodal mathematical reasoning with MATH-Vision dataset. In *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2024. 3, 5, 7

[30] Xinpeng Wang, Bolei Ma, Chengzhi Hu, Leon Weber-Genzel, Paul Röttger, Frauke Kreuter, Dirk Hovy, and Barbara Plank. "my answer is c": First-token probabilities do not match text answers in instruction-tuned language models. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 7407–7416, 2024. 6

[31] Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming Ren, Aaran Arulraj, Xuan He, Ziyan Jiang, Tianle Li, Max Ku, Kai Wang, Alex Zhuang, Rongqi Fan, Xiang Yue, and Wenhu Chen. MMLU-pro: A more robust and challenging multi-task language understanding benchmark. In *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2024. 3

[32] Qingchen Yu, Zifan Zheng, Shichao Song, Zhiyu Li, Feiyu Xiong, Bo Tang, and Ding Chen. xFinder: Large language models as automated evaluators for reliable evaluation. In *International Conference on Learning Representations (ICLR)*, 2025. 2

[33] Xiang Yue, Yuansheng Ni, Kai Zhang, Tianyu Zheng, Ruoqi Liu, Ge Zhang, Samuel Stevens, Dongfu Jiang, Weiming Ren, Yuxuan Sun, et al. MMMU: A massive multi-discipline multimodal understanding and reasoning benchmark for expert AGI. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 1, 3, 5

[34] Xiang Yue, Tianyu Zheng, Yuansheng Ni, Yubo Wang, Kai Zhang, Shengbang Tong, Yuxuan Sun, Botao Yu, Ge Zhang, Huan Sun, Yu Su, Wenhu Chen, and Graham Neubig. MMMU-Pro: A more robust multi-discipline multimodal understanding benchmark. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 15134–15186, 2025. 4

[35] Kaichen Zhang, Bo Li, Peiyuan Zhang, Fanyi Pu, Joshua Adrian Cahyono, Kairui Hu, Shuai Liu, Yuanhan Zhang, Jingkang Yang, Chunyuan Li, and Ziwei Liu. LMMs-Eval: Reality check on the evaluation of large multimodal models. In *Findings of the Association for Computational Linguistics: NAACL 2025*, 2025. 3

[36] Xiaoman Zhang, Chaoyi Wu, Ziheng Zhao, Weixiong Lin, Ya Zhang, Yanfeng Wang, and Weidi Xie. Development of a large-scale medical visual question-answering dataset. *Communications Medicine*, 4:277, 2024. 1, 3, 5

[37] Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. Large language models are not robust multiple choice selectors. In *International Conference on Learning Representations (ICLR)*, 2024. 2

[38] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, Hao Zhang, Joseph Gonzalez, and Ion Stoica. Judging LLM-as-a-judge with MT-bench and chatbot arena. In *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2023. 3

[39] Yuxuan Zhu, Tengjun Jin, Yada Pruksachatkun, Andy Zhang, Shu Liu, Sasha Cui, Sayash Kapoor, Shayne Longpre, Kevin Meng, Rebecca Weiss, Fazl Barez, Rahul Gupta, Jwala Dhamala, Jacob Merizian, Mario Giulianelli, Harry Coppock, Cozmin Ududec, Antony Kellermann, Jasjeet Sekhon, Jacob Steinhardt, Sarah Schwettmann, Arvind Narayanan, Matei A. Zaharia, Ion Stoica, Percy Liang, and Daniel Kang. Establishing best practices in building rigorous agentic benchmarks. In *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2025. 3