arXiv:2611.04819v1 [cs.AI] 3 Nov 2026

Reconstructing Fraudulent Strategies in Online Payments via Interpretable Inverse Reinforcement Learning

Huilin He[1,2,†], Weiyu Han[2], Linbo Shao[1], Ruobing Zuo[3,2,†],
Xiao Tan[2,*], Shuai Chen[2], and Dawei Cheng[1,*]
[1] Tongji University, [2] Ant International, [3] East China Normal University
{huilin3,2634102,dcheng}@tongji.edu.cn,{hanweiyu.hwy,alex.tx,shuai.cs}@ant-intl.com,rbzuo@stu.ecnu.edu.cn

Abstract

Payment fraud causes financial losses and undermines users' trust in web payment services. Although fraud detection is a critical defense for payment security, complex and diverse fraudulent behaviors can still evade detection. Generating fraud samples for data augmentation can broaden detectors' coverage of different fraud patterns. However, existing methods matching the distribution of transaction attributes alone do not ensure that generated sequences preserve the sequential relationships among transaction choices, payment feedback, and subsequent behavioral adjustments. To this end, we propose PaySeqGen, an interpretable inverse reinforcement learning framework that reconstructs fraud behavior policies from fraudsters' observed payment trajectories. We design a mechanism for constructing candidate actions to support plausible payment sequence generation while enabling exploration. The framework also learns interpretable behavioral rewards that capture fraudsters' preferences across different transaction contexts. By alternating between reward learning and policy optimization, the framework reconstructs policies that generate new fraud sequences for detector augmentation and learns interpretable rewards supporting sequential fraud-pattern analysis in real trajectories. Experiments in real-world payment scenarios at Ant International, a leading global digital payment provider, show that PaySeqGen outperforms state-of-the-art generative augmentation methods. Reward-based trajectory analysis also identifies fraud patterns of practical value, helping analyst locating suspicious segments and examine multistep fraud tactics. Evaluation in the deployed system of online payment further confirms its effectiveness.

1 Introduction

Online payment systems underpin e-commerce and digital services, making their security a critical concern for users and merchants [2]. Payment fraud causes financial losses for users, merchants, and payment providers, raises risk management and dispute resolution costs, damages platform reputations, and undermines trust in online payment services [8, 22]. Fraud detection therefore serves as a key line of defense [5]. However, fraudulent behavior takes diverse forms across different payment contexts, and limited training data may not adequately capture complex fraud patterns and their variations. This limited coverage constrains detectors' ability to learn and reliably recognize such patterns in practice [7, 30].

Generating additional payment sequences from observed fraudulent behavior provides a way to expand this coverage by introducing behavioral variations into detector training. Existing sequence generation methods typically model transaction attribute distributions and temporal dependencies [26, 46, 48, 49]. However, matching observed sequences alone does not necessarily ensure that generated decisions remain coherent as the history evolves and new feedback becomes available. In web payment settings, each payment attempt updates the transaction history, and its outcome can influence subsequent decisions about payment amounts, intervals, and card and device use, as illustrated in Figure 1. Fraud sequence generation therefore calls for explicit modeling of this interaction process: selecting the next payment action based on the current history, incorporating the resulting payment feedback into that history, and making subsequent decisions that reflect adaptive fraud strategies while respecting payment constraints.

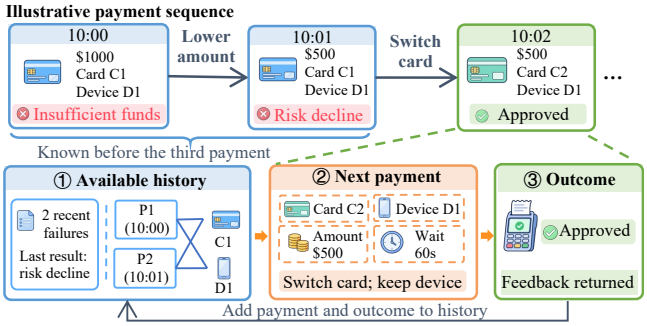

Figure 1: An illustrative online fraud payment sequence and its decision context shaped by feedback.

Such feedback-driven sequential decision-making is naturally modeled using reinforcement learning (RL), which learns policies by maximizing expected cumulative reward and has also been applied to sequence generation [20, 47]. In our setting, however, the reward is unknown. Handcrafting rewards requires specifying in advance which payment behaviors should be encouraged. These predefined preferences may not align with those reflected in real demonstrations and can bias the learned policy. Inverse reinforcement learning (IRL) [9, 11, 50] offers an alternative by inferring rewards that explain demonstrated behavior. In our setting, observed payment trajectories can serve as demonstrations for learning a behavioral reward and reconstructing a policy that generates new payment sequences. A decomposable reward representation further exposes the scoring contributions of individual behavioral factors, providing interpretable evidence for analyzing sequential fraud patterns. Prior work on rumor detection has used IRL to reconstruct graph-structure attack policies [21], providing a methodological precedent for this approach. However, using interpretable IRL to reconstruct fraudulent strategies in online payments remains underexplored.

* Corresponding Author.
† Work done during internships at Ant International.

Applying this approach to web payments presents two challenges. First, candidate action construction must balance behavioral coherence and exploration. Decisions to reuse or switch cards and devices must account for transaction history and payment feedback, as must decisions about payment amounts, intervals, and other attributes. Generating these fields independently can disrupt behavioral coherence [20]. Meanwhile, the combinatorial space of payment actions is too large to enumerate. The challenge is therefore to cover historically supported behavioral combinations with a finite candidate set while allowing exploration beyond combinations observed in demonstrations. Second, rewards must capture context-dependent behavioral preferences [15]. The same action, such as switching cards or increasing a payment amount, may reflect different behavioral patterns depending on preceding transactions and payment outcomes. Scoring action attributes in isolation makes it difficult to distinguish context-dependent preferences.

To address these challenges, we propose PaySeqGen, an interpretable inverse reinforcement learning framework for reconstructing payment fraud policies from observed trajectories. At each step, the state summarizes payment history and feedback, and the action specifies the next payment. To balance behavioral coherence and exploration within a finite action set, we combine history-guided sampling with exploratory sampling. History-guided sampling draws joint changes from historical transitions matching contexts and feedback, adapting them to current history. Exploratory sampling recombines payment components under resource and time constraints. To capture context-dependent behavioral preferences, we learn a linear behavioral reward over base features and feedback–action interactions, allowing action scores to vary with preceding payment feedback. We alternate reward learning with policy optimization to reconstruct payment fraud strategies, using a fixed payment feedback interface adapted from the platform's risk-control system to provide outcomes for generated payments. After training, the reconstructed policy generates fraud payment sequences to augment detector training. The learned interpretable reward helps characterize sequential fraud patterns in observed trajectories.

Our contributions are summarized as follows:

- We propose PaySeqGen, an interpretable IRL framework for reconstructing fraudulent payment strategies from observed trajectories, supporting both sequence generation for detector augmentation and reward-guided analysis of sequential fraud patterns.
- We design a candidate construction mechanism that combines history-guided and exploratory sampling to balance behavioral coherence and exploration, and a decomposable linear reward with feedback–action interactions to capture context-dependent behavioral preferences.
- We evaluate PaySeqGen on large-scale data from a real-world online payment system operated by Ant International[1], a leading global digital payment and financial technology provider. Offline evaluations and online A/B tests demonstrate PaySeqGen's augmentation benefits. Reward-guided analysis also provides traceable evidence for fraud investigation. These results demonstrate PaySeqGen's practical value in strengthening web payment security.

[1]https://www.ant-intl.com/en/

2 Preliminaries

2.1 Reinforcement Learning and Actor-Critic

In reinforcement learning, a policy $\pi_\omega(a_t \mid s_t)$ selects actions at each step and interacts with environment dynamics P to generate a trajectory $\tau = (s_1, a_1, \ldots, s_T, a_T)$. Policy learning maximizes the expected cumulative reward:

$$J(\omega) = \mathbb{E}_{\tau\sim\pi_\omega,P}\left[\sum_{t=1}^{T} r(s_t, a_t)\right]. \tag{1}$$

Here, ω denotes the policy parameters, P is the environment transition distribution, τ is a trajectory containing T decision steps, and r is the reward function. Policy updates based on sampled returns can have high variance [38]. We adopt Actor–Critic [16] to provide a learned value baseline for these updates. The Actor π_ω outputs action probabilities, while the Critic $V_\psi(s_t)$ estimates the expected return from s_t under the current policy, with ψ denoting the Critic parameters. Let $\widehat{G}_t$ denote a return estimate based on the selected action and subsequent interactions. The estimated advantage is defined as:

$$\widehat{A}_t = \widehat{G}_t - V_\psi(s_t). \tag{2}$$

The advantage measures how much better the selected action performs relative to the current policy's average behavior at the same state. The Actor uses this signal to favor actions with better-than-expected returns, while the Critic is trained to predict the return targets $\widehat{G}_t$ computed from sampled trajectories.

2.2 Maximum Entropy IRL

Inverse reinforcement learning (IRL) assumes that expert behavior is optimal or near-optimal under an unknown reward function and aims to recover a reward that explains the observed behavior. To account for variability in expert behavior, maximum entropy IRL models a stochastic policy that favors higher cumulative rewards while retaining behavioral diversity [50]. Let $r_\theta(s, a)$ denote a parameterized reward function and π_θ the stochastic policy induced by this reward under the maximum-entropy model. Given an expert demonstration set $\mathcal{D}$, reward learning can be formulated as maximizing the log-likelihood of the demonstrated actions:

$$\theta^* = \arg\max_\theta \sum_{\tau\in\mathcal{D}} \sum_t \log \pi_\theta(a_t^E \mid s_t^E), \tag{3}$$

where (s_t^E, a_t^E) is a demonstrated state–action pair in trajectory τ. This objective adjusts the reward parameters so that the induced policy assigns higher probabilities to expert actions, thereby encoding demonstrated behavioral preferences as explicit reward scores. Our framework follows this perspective through sampled feature-expectation matching and alternating reward and policy updates.

3 Methodology

3.1 Overview

As shown in Figure 2, PaySeqGen learns an interpretable behavioral reward and reconstructs a fraud payment policy from real demonstrations. The policy selects the next payment based on the current history, while the reward evaluates whether this choice reflects the behavioral preferences in the demonstrations. To support these roles, candidate construction balances behavioral coherence

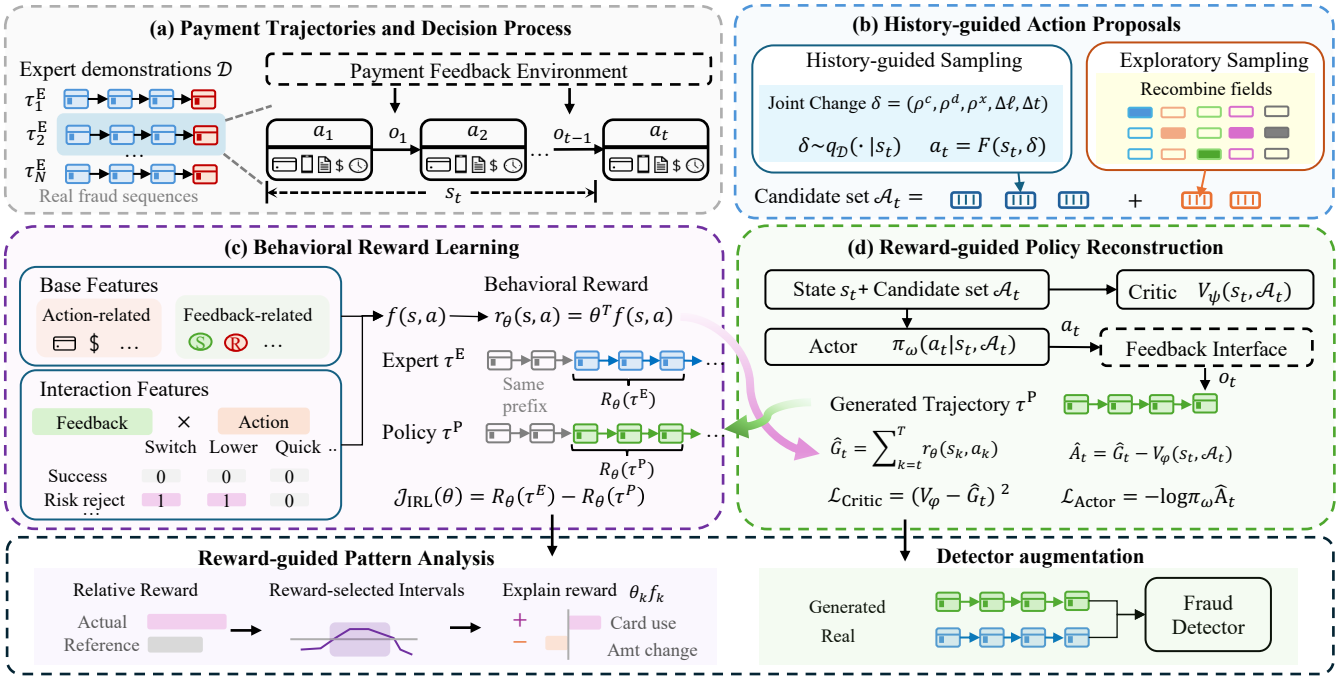

Figure 2: Overview of PaySeqGen. (a) Model payment trajectories as a sequential decision process. (b) Construct candidate actions through history-guided sampling and exploratory sampling. (c) Learn behavioral rewards with feedback–action interactions, supporting fraud-pattern analysis and guiding fraud policy reconstruction. (d) Optimize the policy using learned rewards to generate payment sequences for detector augmentation.

and exploration, while behavioral reward learning captures how action preferences vary with prior feedback. Within this process, reward learning and policy reconstruction alternate: differences between demonstrated and generated behavior inform reward updates, which in turn guide policy optimization. After training, the reconstructed policy generates sequences for detector augmentation, while the learned reward supports fraud-pattern analysis.

3.2 Payment Trajectories and Decision Process

Let a real trajectory be $\tau = \{(a_t, o_t)\}_{t=1}^{T}$, where a_t is a payment action and o_t is the feedback returned after execution. We represent the next payment event choice as an action:

$$a_t = (c_t, d_t, x_t, m_t, \Delta t_t) \tag{4}$$

where c_t and d_t denote the card and device, respectively; x_t contains card attributes and payment-context attributes, such as country and terminal information; m_t is the payment amount; and Δt_t is the waiting time since the previous payment. The state s_t is represented as the ordered history of payments and their observed feedback:

$$s_t = \big((a_1, o_1), \ldots, (a_{t-1}, o_{t-1})\big). \tag{5}$$

Real payment transactions receive feedback from the payment system, whereas generated actions require simulated feedback, since underlying business states, such as account balances, are not fully observable. We therefore use a payment feedback interface adapted from the platform's risk-control system with a multiclass output head. This interface approximates the distribution of payment outcomes, including success, rejection, and other failure categories, and remains fixed during reward and policy learning:

$$\begin{aligned} o_t &\sim P_{\text{env}}(\cdot \mid s_t, a_t), \\ s_{t+1} &= \text{Update}(s_t, a_t, o_t). \end{aligned} \tag{6}$$

Here, P_{env} is the conditional distribution of payment feedback, and Update appends the current payment and its feedback to the history. Feedback for the current payment becomes available only after execution and is used in the next decision.

3.3 History-guided Action Proposals

The action candidate set $\mathcal{A}_t$ defines the payment choices available to the policy at step t, shaping both sequence coherence and the scope of exploration. Constructing this set is challenging because payment attributes form a large combinatorial space, while independent field sampling can disrupt their dependencies. We therefore combine joint candidates guided by real payment histories with constrained random exploration, preserving observed attribute dependencies while expanding the behavioral combinations available to the policy. We first construct history-guided candidates from joint changes observed between consecutive payments in the demonstration set $\mathcal{D}$. Each change is represented as:

$$\delta_i = (\rho_i^c, \rho_i^d, \rho_i^x, \Delta\ell_i, \Delta t_i), \tag{7}$$

where ρ_i^c and ρ_i^d describe card and device usage, such as retaining the current resource, returning to a previously used resource, or introducing a new one. ρ_i^x records whether the payment context is retained or replaced, together with the corresponding attribute

template. $\Delta\ell_i$ is the change in log amount $\log(1+m)$, where m is the payment amount, and Δt_i is the interval between payments. These components are extracted from the same pair of consecutive payments and sampled jointly to preserve co-occurring behavioral changes. At state s_t, we sample K_h joint changes by matching the latest payment feedback, amount bin, and payment context, and adapt them to the current history:

$$\begin{aligned} \delta_{t,j} &\sim q_{\mathcal{D}}(\cdot \mid s_t), \\ a^{\text{hist}}_{t,j} &= F(s_t, \delta_{t,j}), \qquad j = 1, \ldots, K_h. \end{aligned} \tag{8}$$

Here, $q_{\mathcal{D}}$ samples joint changes from demonstrations under matching payment conditions. The mapping F adapts the sampled changes to the current resources, attributes, and time budget. The resulting candidates form $\mathcal{A}^{\text{hist}}_t$. For exploration, we draw K_e proposals $\tilde{a}_{t,j} \sim q_{\text{exp}}(\cdot \mid s_t)$ and construct candidates as:

$$\begin{aligned} q_{\text{exp}}(\tilde{a} \mid s_t) &= \prod_{u \in \{c,d,x,m,\Delta\}} q_u(\tilde{a}_u \mid s_t), \\ a^{\text{exp}}_{t,j} &= F_{\text{exp}}(s_t, \tilde{a}_{t,j}), \qquad j = 1, \ldots, K_e. \end{aligned} \tag{9}$$

Here, q_u samples a payment component without historical transition matching, and F_{exp} enforces resource-consistency and time-budget constraints. These candidates form $\mathcal{A}^{\text{exp}}_t$. The final candidate list is $\mathcal{A}_t = \text{Concat}(\mathcal{A}^{\text{hist}}_t, \mathcal{A}^{\text{exp}}_t)$.

3.4 Behavioral Reward Learning

3.4.1 Decomposable Behavioral Rewards. To make learned behavioral preferences interpretable, we define a linear reward over a set of fixed behavioral features:

$$\begin{aligned} f(s,a) &= [b(s,a); \chi(s,a)], \\ r_\theta(s,a) &= \theta^\top f(s,a). \end{aligned} \tag{10}$$

Here, $b(s,a)$ summarizes amounts, payment intervals, resource usage, and historical feedback. Preferences for the same action may vary with prior payment feedback. We therefore construct the interaction vector $\chi(s,a)$ from products of feedback-related and action-related features. For example,

$$\begin{aligned} \chi_{\text{reject,switch}}(s_t, a_t) = \mathbf{1}[o_{t-1} = \text{risk rejection}] \\ \cdot \mathbf{1}[c_t \neq c_{t-1},\ d_t = d_{t-1}] \end{aligned} \tag{11}$$

captures switching cards while retaining the device after a risk rejection. The feature definitions are fixed and shared by the policy and reward. The policy uses them for nonlinear scoring, whereas the reward learns their linear weights. Each contribution $\theta_k f_k(s,a)$ can therefore be traced to a specific behavioral factor, supporting subsequent fraud-pattern analysis.

3.4.2 Reward Optimization. We use an observed payment prefix as a shared starting history to compare expert and policy behavior under the same initial conditions. The expert branch replays the recorded continuation, while the policy branch generates an alternative. The gap between their behavioral feature expectations provides the signal for updating the reward weights. For the $L = T - p$ decisions following a prefix of p payments, let $\bar{f}(\tau) = L^{-1} \sum_{t=p+1}^{T} f(s_t, a_t)$ denote the average feature vector over the continuation. The expert and policy feature expectations are:

$$\begin{aligned} \mu^E &= \mathbb{E}_{\tau \sim \mathcal{D}}[\bar{f}(\tau)] \\ \mu^\pi &= \mathbb{E}_{\tau \sim \pi_\omega, P_{\text{env}}}[\bar{f}(\tau)], \end{aligned} \tag{12}$$

where $\mathcal{D}$ is the demonstration set. Following feature-expectation matching in IRL [1, 50], we hold the sampled trajectories fixed and update the reward by minimizing:

$$\mathcal{L}_{\text{IRL}}(\theta) = \theta^\top(\mu^\pi - \mu^E) + \frac{\lambda}{2}\|\theta\|_2^2. \tag{13}$$

We optimize the reward weights by gradient descent, with each weight clipped to $[-w_{\max}, w_{\max}]$. The regularization term discourages large weights, while clipping imposes an explicit bound on each learned reward coefficient. The feature-matching term increases weights for features underrepresented in generated trajectories and decreases those overrepresented relative to the expert demonstrations. The updated reward then guides the next policy optimization step. Through alternating reward and policy updates, the learned weights identify which behavioral factors and feedback–action combinations to encourage or discourage to reproduce demonstrated feature statistics, expressing context-dependent preferences as explicit, decomposable scores.

3.5 Reward-guided Policy Reconstruction

3.5.1 Policy representation and initialization. Given the candidate list $\mathcal{A}_t$, the policy selects the next payment from its K candidates. For each candidate $a_{t,j}$, we combine the behavioral features defined above with learned attribute encodings into a vector: $\mathbf{h}_{t,j} = [f(s_t, a_{t,j}); e_a(a_{t,j}); e_s(s_t)]$. Here, semicolons denote concatenation, e_a encodes candidate attributes, and e_s mean-pools learned attribute embeddings over the payment history. The Actor scores all candidates with the same multilayer perceptron (MLP) and normalizes their scores into selection probabilities:

$$\pi_\omega(j \mid s_t, \mathcal{A}_t) = \frac{\exp\big(\text{MLP}_\omega(\mathbf{h}_{t,j})\big)}{\sum_{k=1}^{K} \exp\big(\text{MLP}_\omega(\mathbf{h}_{t,k})\big)}. \tag{14}$$

The Critic estimates the expected remaining return under the current policy as $V_\psi(s_t, \mathcal{A}_t) = \text{MLP}_\psi\left(K^{-1}\sum_{j=1}^{K} \mathbf{h}_{t,j}\right)$. Conditioning the Critic on the candidate list allows its baseline to account for the payment choices available at each step. The Actor and Critic share the attribute encoders and use separate MLP heads. Their parameter sets, ω and ψ, include the shared encoder parameters. Before autonomous generation, we initialize the Actor through behavior cloning: the demonstrated next payment is included in each candidate set, and the policy maximizes its log-probability under the corresponding real history. This provides an initial payment-selection policy for subsequent reinforcement learning.

3.5.2 Reward-guided policy optimization. With the current behavioral reward r_θ fixed, the policy samples the next payment from $\mathcal{A}_t$. The payment feedback interface provides an outcome, which is added to the history together with the selected action. We add an entropy bonus to encourage exploration among candidates and normalize by the number of generated steps:

$$\tilde{r}_t = \frac{r_\theta(s_t, a_t) - \alpha \log \pi_\omega(j_t \mid s_t, \mathcal{A}_t)}{L}, \tag{15}$$

where L is the number of generated decision steps following the initial prefix, and α controls the strength of entropy regularization. We use generalized advantage estimation (GAE) [33] to estimate $\widehat{A}_t$ by combining sampled rewards with Critic predictions over multiple horizons. This balances variability in sampled returns

against value-estimation bias, helping stabilize policy updates. The corresponding return targets $\widehat{G}_t$ are computed using Eq. (2). We then minimize the joint loss:

$$\mathcal{L}_{\mathrm{RL}}(\omega, \psi) = -\mathbb{E}_t \left[\log \pi_\omega(j_t \mid s_t, \mathcal{A}_t)\, \widehat{A}_t \right] + \beta_V \mathbb{E}_t \left[\left(V_\psi(s_t, \mathcal{A}_t) - \widehat{G}_t \right)^2 \right], \tag{16}$$

where $\mathbb{E}_t$ averages over generated decision steps in the sampled batch, and β_V weights the Critic loss. The advantages and return targets are held fixed during each update. The first term increases the probability of choices with better-than-expected multistep returns, while the second improves the Critic's return estimates.

3.6 Alternating Training and Applications

We initialize the policy by behavior cloning on demonstrations and set the reward weights to zero. Training then alternates between two steps: updating the reward using feature-expectation differences between expert and current policy trajectories, and fixing the reward while sampling fresh trajectories for policy optimization. Each policy update changes the generated distribution, which in turn informs the next reward update.

3.6.1 Fraud sequence generation and detector augmentation. After training, we use the reconstructed fraud policy to generate synthetic fraudulent payment sequences for detector augmentation. Each sequence begins with an initial payment assembled from training-data components and uses fresh user and resource identifiers to avoid adding synthetic payments to real users' existing histories. The policy then generates subsequent payments using feedback from the payment feedback interface. The final payment serves as the fraud-labeled target, with preceding payments providing its historical context. The synthetic sequences are combined with real training samples for detector augmentation.

3.6.2 Reward-guided Fraud Pattern Analysis. With the learned reward fixed, we replay real trajectories to identify informative behavioral intervals for fraud pattern analysis. At each historical state s_t, we reuse the candidate construction mechanism from policy generation to obtain $\mathcal{A}_t$. We define the relative reward Δr_t as the difference between the observed action's reward and the mean reward of the candidate actions:

$$\Delta r_t = r_\theta(s_t, a_t^E) - \frac{1}{|\mathcal{A}_t|} \sum_{a \in \mathcal{A}_t} r_\theta(s_t, a). \tag{17}$$

Here, a_t^E is the observed payment action. A larger Δr_t indicates a stronger learned preference for this choice relative to the candidate alternatives under the same history. Since the reward is learned from fraud-related demonstrations, predominantly positive scores across consecutive payments provide evidence of sustained alignment with the modeled behavioral preferences. We therefore select such intervals as candidate segments for fraud-pattern analysis. Signed feature contributions identify which factors raise or lower the observed actions' rewards relative to the candidate mean. Together with recorded amounts, resource changes, and payment feedback, these contributions help characterize the sequential behavior and contextual preferences within each selected interval.

4 Experiments

4.1 Experiment Setup

4.1.1 Datasets. We conduct experiments on real-world data collected from a large online payment platform operated by Ant International. Our dataset covers approximately 1.29 million users and 26.28 million distinct payment events across 3.46 million labeled detection examples. Each example corresponds to a user's payment sequence ending in a target payment, and its binary label indicates whether that target payment is fraudulent. Target payments span March 2025 to January 2026. The records include payment amounts and timing, card and device usage, and available payment feedback. All detectors share the same chronological training, validation, and test splits. Expert demonstrations are drawn from training examples with fraudulent targets; action-proposal templates and sampling statistics are also derived exclusively from the training split.

4.1.2 Target Models. We evaluate the utility of generated samples for augmenting four fraud detectors with different architectures. Our primary detector, referred to as RT, is a Transformer-based [39] payment fraud detection model deployed in our production risk control system. Given the sequential nature of our task, we also include TabFormer [28], which models payment sequences using a hierarchical Transformer architecture at the field and payment-event levels, with masked-attribute pretraining followed by fine-tuning for fraud detection. We also evaluate GraphSAGE [13], motivated by the widespread use of graph neural networks in fraud detection [6, 19]. This allows us to assess whether the benefits of generated samples extend to detectors that model relationships among payments. Finally, we include XGBoost [4], a well-established baseline in production fraud detection systems, to assess whether the benefits of generated samples extend to tree-based fraud detectors.

4.1.3 Generation Baselines. We compare PaySeqGen with two categories of generation baselines. Distribution-based methods include CTGAN [44] for conditional adversarial generation of mixed-type attributes, Diffusion-TS [48] for diffusion-based sequence generation with temporal decomposition, and TabDiff [37] for joint diffusion over continuous and discrete fields. Policy-based methods include Behavior Cloning (BC) [27] and MoE-BiEntIRL [21]. BC learns a payment policy through supervised imitation of expert actions. It shares PaySeqGen's policy architecture and candidate construction mechanism but is trained independently by maximizing demonstrated-action likelihood, without IRL or RL updates. MoE-BiEntIRL combines IRL with a mixture-of-experts formulation. We additionally include a *Real-fraud replay* control, which adds an equal number of resampled real fraud examples. This control helps distinguish the benefits of generating behavioral variations from the effect of simply increasing the weight of fraud examples during detector training. All generation methods use the same training data sources and generation lengths.

4.1.4 Evaluation Metrics. We evaluate detector augmentation using the area under the ROC curve (AUC), Recall@2%, and the Kolmogorov–Smirnov (KS) statistic. AUC measures the detector's ability to rank fraudulent payments above legitimate ones. Recall@2 measures the proportion of all fraudulent samples captured among the top 2% of samples ranked by predicted risk. KS measures the

Table 1: Performance changes across four fraud detectors under different augmentation methods. Values are percentage-point differences from training without augmentation. Boldface and underlining indicate the best and second-best results per column.

Detectors		RT			TabFormer			GraphSAGE			XGBoost		
Metrics		ΔAUC	ΔRecall@2	ΔKS	ΔAUC	ΔRecall@2	ΔKS	ΔAUC	ΔRecall@2	ΔKS	ΔAUC	ΔRecall@2	ΔKS
Control	Fraud replay	-0.046	+1.094	+0.166	-0.092	-1.679	-0.042	+0.036	+1.098	+0.416	**+0.109**	**+0.325**	**+0.074**
Baselines	BC	-0.141	-0.360	+0.055	+0.089	-2.972	-0.097	+0.681	+3.375	+0.975	-0.195	-1.266	-0.354
	CTGAN	-0.036	+1.559	-0.169	+0.148	-2.218	+0.170	+0.066	+1.370	+0.934	-0.025	-0.193	-0.304
	Diffusion-TS	-0.189	+0.871	-0.610	-0.042	-1.923	-0.423	+0.343	+1.454	+1.056	-0.095	-0.976	-0.232
	TabDiff	+0.169	+0.769	+0.108	+0.198	-2.000	+0.295	-0.251	+0.111	-0.517	-0.093	-0.103	-0.628
	MoE-BiEntIRL	-0.051	+1.036	-0.042	+0.149	-2.274	-0.130	+0.360	+1.951	+0.697	-0.030	-0.004	-0.428
Ours	FULL	**+0.393**	**+1.621**	**+0.218**	**+0.328**	**-1.613**	**+0.451**	**+0.878**	**+4.561**	**+1.686**	+0.062	-0.047	-0.032
Ablations	w/o coh.	+0.175	+1.171	+0.065	+0.089	-1.700	+0.186	+0.654	+1.426	+1.332	+0.034	-0.411	-0.054
	w/o expl.	+0.272	+0.568	-0.071	+0.100	-1.872	+0.006	+0.596	+1.165	+1.042	+0.042	-0.704	+0.092
	w/o inter.	-0.257	+0.653	-0.151	+0.240	-1.929	+0.343	+0.588	+1.250	+1.454	-0.022	-0.231	-0.187

maximum difference between the true positive rate and the false positive rate across all thresholds, reflecting class separation. Higher values indicate better performance for all three metrics.

4.1.5 Implementation Details. Experiments are conducted on a server using CUDA 12.4 for acceleration and PyTorch 2.4.0 [29] for training and evaluation. For each generator and random seed, we augment the detector training set with synthetic fraud samples equivalent to approximately 2.5% of its original size. Original training examples and the validation and test sets remain unchanged. PaySeqGen constructs 64 action candidates per step, comprising 48 history-guided candidates and 16 exploration candidates. We use a policy hidden dimension of 128 and a training batch size of 256. The policy and reward learning rates are 3×10^{-4} and 0.03, respectively. For generalized advantage estimation, we set the discount factor to $\gamma = 1$ and the smoothing parameter to $\lambda_{\text{GAE}} = 0.95$. Baselines and detectors follow their original implementations or reported settings, with adaptations to payment sequences.

4.2 Improvement of Detection Performance with Generated Samples

Table 1 shows that PaySeqGen achieves the best or second-best results across all four detectors. It ranks first on all three metrics for RT, TabFormer, and GraphSAGE, highlighting the value of its generated sequences for detectors that capture sequential or relational dependencies. On our production detector RT, PaySeqGen improves AUC by 0.393 percentage points over training without augmentation, exceeding the strongest generation baseline by 0.224 points. It also achieves the largest Recall@2 gain of 1.621 points, together with a KS increase of 0.218 points. Compared with real-fraud replay, PaySeqGen provides an additional 0.527 points of Recall@2, supporting the utility of generated samples beyond increasing the proportion of fraud examples in training. Large transaction volumes in web payment systems make it important to prioritize high-risk payments for further assessment. The Recall@2 gain therefore indicates that PaySeqGen helps identify more fraud among these prioritized payments, helping platforms allocate limited risk-control resources more effectively. The benefits are also evident in graph-based detection. On GraphSAGE, the Recall@2 gain reaches 4.561 points, exceeding the strongest generation baseline on this metric by 1.186 points. The accompanying AUC and KS improvements extend the evidence for useful augmentation to a detector that models relationships among payments. On XGBoost, PaySeqGen ranks second to real-fraud replay on AUC and KS, while outperforming all generation baselines, suggesting that its generated sequences also capture useful attribute-level patterns for fraud detection. Together, these findings support our fraud-policy reconstruction as a practical way to convert observed fraud behavior into reusable training resources for large-scale payment risk control, strengthening online payment security.

4.3 Ablation Study

To assess the candidate construction mechanism and behavioral reward design, we compare three variants: *w/o coh.* removes history conditioning from transition sampling while retaining joint payment changes; *w/o expl.* replaces exploratory candidates with history-guided candidates while keeping the candidate count unchanged; and *w/o inter.* removes feedback–action interactions from the reward while retaining all policy inputs. These comparisons assess how each component contributes to downstream detection performance. All three ablations generally weaken performance relative to the full model, although their relative effects vary across detectors and metrics. Removing history conditioning reduces Recall@2 by 0.450 percentage points on RT and 3.135 points on GraphSAGE, suggesting that aligning payment proposals with their preceding histories is important for generating useful fraud examples. Removing exploration produces the largest or second-largest Recall@2 decline on every detector, including the largest decrease on RT (1.053 points). This is consistent with its role in broadening behavioral coverage for detector training. Removing interaction reward leads to the largest AUC decrease on three of the four detectors, suggesting that modeling contextual preferences is important for reconstructing behaviors that support accurate risk ranking. Together, these findings support combining history-guided sampling, exploratory sampling, and feedback–action interaction reward to generate useful training sequences for payment fraud detection.

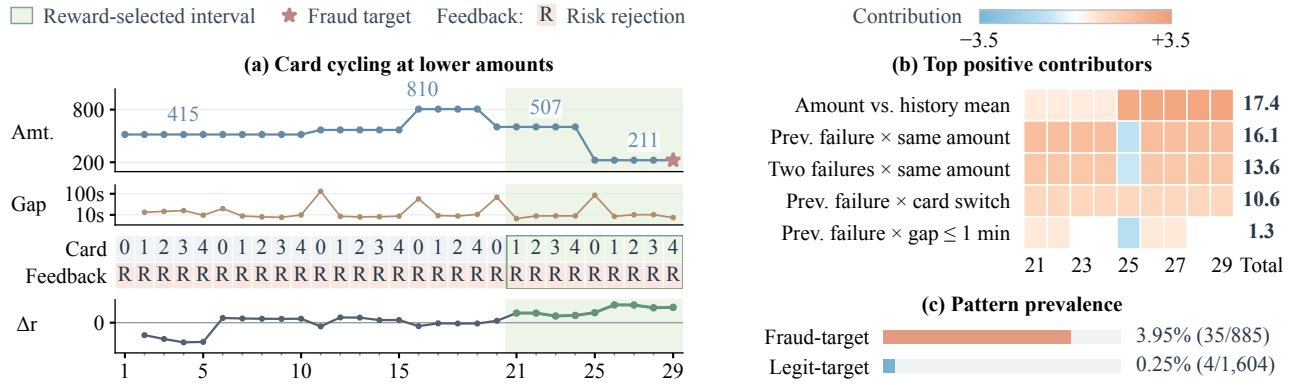

Figure 3: Reward-guided fraud-pattern analysis. (a) Payment history and relative rewards, with the selected interval shaded and the fraudulent target marked by a star. (b) Features with the largest positive contributions to relative reward within the selected interval. (c) Prevalence of the corresponding pattern among real fraud-target and legitimate-target trajectories.

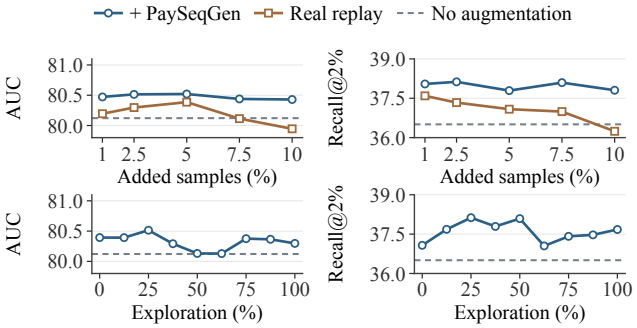

Figure 4: Sensitivity Analysis with respect to the ratio of augmented samples and the exploration candidate ratio.

4.4 Parameter Sensitivity

We further investigate the effects of augmentation size and the proportion of exploratory candidates. As shown in Figure 4, PaySeqGen achieves higher AUC than real-fraud replay across the tested augmentation ratios, supporting the utility of generated samples beyond oversampling existing fraud examples. Its AUC increases as the augmentation ratio rises from 1% to 5%, whereas Recall@2 peaks at 2.5%. Increasing the ratio beyond 5% provides no consistent additional benefit. The 2.5% setting thus combines the highest observed Recall@2 with near-peak AUC and a modest increase in training data, illustrating the value of a limited augmentation budget. For candidate construction, a 25% exploration ratio achieves the highest AUC and Recall@2 among the tested settings. Performance fluctuates at higher ratios, suggesting that increasing exploration alone does not consistently improve detector augmentation. This may reflect a trade-off between behavioral coverage and coherence: insufficient exploration may restrict the diversity of generated behaviors, while excessive exploration may introduce candidates that are less consistent with the observed payment history, adding noise to the generated sequences. These results support combining history-guided and exploratory candidates to generate useful training sequences for payment fraud detection.

4.5 Interpretability Analysis

We use a real payment trajectory with a fraudulent target to illustrate reward-guided fraud-pattern analysis. In Figure 3, consecutive positive relative rewards (Eq. (17)) highlight an interval in which the observed actions receive higher scores than the candidate mean under the same history. Within this interval, the payment records show rapid retries across multiple cards at a fixed amount, followed by further retries at a lower fixed amount. The heatmap presents top positive contributions to relative reward, including amounts relative to the historical mean and interactions between preceding failures and same-amount retries, card switching, or short payment intervals. Together, the payment records and reward contributions provide traceable evidence for interpreting how payment behavior evolves following repeated rejections. To examine the pattern beyond this case, we further assess the prevalence of such sequence pattern using the reward-derived rule (specified in Appendix E). We apply this rule to all eligible trajectories in the validation cohort without reward-based preselection. It matches 3.95% of fraud-target trajectories and 0.25% of legitimate-target trajectories. The higher prevalence in fraud-target trajectories suggests that the pattern is associated with fraudulent payment behavior. These findings demonstrate how PaySeqGen's interpretable behavioral cues can be used to help analysts focus their review of complex payment histories in large-scale online payment systems.

4.6 Industrial Deployment and A/B Test

PaySeqGen is integrated into Ant International's production online payment risk-control system to strengthen existing fraud detectors through data augmentation. As shown in Figure 5, risk services first perform rule and blacklist checks, followed by feature extraction and model-based risk scoring. Risk services log payment records and outcomes in the database. Model services use these records and fraud labels to reconstruct the policy and generate a synthetic sample pool. Sequences generated by PaySeqGen are combined with real payment records to retrain and update the detectors, which are deployed to risk services for subsequent payment screening. Behavioral cues derived from the learned reward are presented in the analyst console to assist analysts in fraud risk analysis.

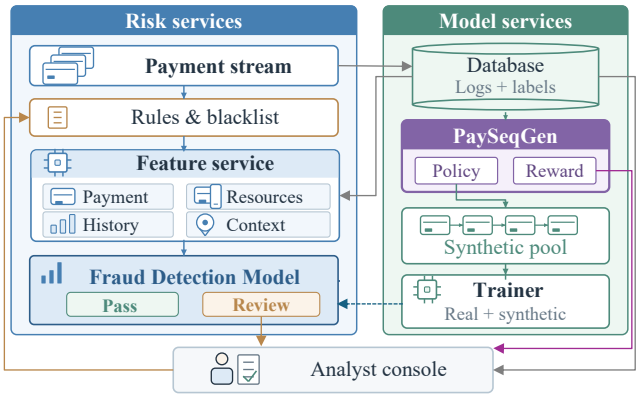

Figure 5: Schematic overview of online payment risk-control pipeline integrating risk services and model services.

Table 2: Results of online A/B tests.

Method	AUC	Recall@2	KS
Model-1	86.02	40.66	54.69
Model-1 (PaySeqGen)	86.23	43.85	55.45
Model-2	85.90	41.79	55.03
Model-2 (PaySeqGen)	86.16	44.63	56.09

Following deployment, we evaluate the resulting augmentation through online A/B tests on two production detectors within a segment of payment traffic over a 30-day period. We track both detection metrics and operational outcomes throughout the evaluation period, including additional fraud identification and the associated reduction in losses. As shown in Table 2, PaySeqGen augmentation improves all three metrics for both detectors. For Model-1 and Model-2, AUC increases by 0.21 and 0.26 percentage points, KS by 0.76 and 1.06 points, and Recall@2 by 3.19 and 2.84 points, respectively. The recall gains indicate that more fraudulent payments are captured at the same 2% screening rate. Beyond these metric improvements, the results show that, compared with the system without augmentation, PaySeqGen augmentation helps identify 250 additional users involved in fraud and detect 393 additional fraudulent payments, preventing an additional US$9,738 in losses. These results demonstrate the practical value of PaySeqGen for improving fraud detection performance and reducing financial losses in online payment services.

5 Related Work

5.1 Generative Data Augmentation

Generative data augmentation operates at both the record and sequence levels to provide additional synthetic samples for model training. Record-level generative models synthesize samples by learning dependencies among heterogeneous attributes [18, 34, 37, 44, 49], while fraud-oriented approaches use adversarial or diffusion models to augment minority-class transactions [23, 43]. Beyond individual transactions, fraud patterns also unfold across successive payment attempts, where temporal dependencies and evolving choices carry information that isolated records cannot preserve. This motivates sequence-level generation that captures temporal structure [31, 36, 46, 48], with extensions incorporating contextual conditions [25] or counterfactual augmentation of minority-class sequences [41]. Banksformer further specializes autoregressive generation for financial transaction sequences [26]. However, these approaches do not explicitly model the action–feedback loop through which payment outcomes guide subsequent decisions. PaySeqGen instead reconstructs a policy from real demonstrations to generate payment sequences through this interaction process.

5.2 Policy Learning for Sequence Generation

Modeling this interaction requires specifying or learning policies. AMLworld simulates financial activity through empirically calibrated agent mechanisms [3]. To learn from observed behavior, behavior cloning and its extensions imitate expert choices and behavioral styles [27, 45]. Beyond direct imitation, adversarial methods guide policy learning and sequence generation with discriminator rewards [14, 47]. For fraud generation, Fraud-RLA uses PPO with task-specific rewards and learns action covariance to model attribute dependencies [20]. For payment fraud reconstruction, the behavioral reward is unknown, and predefined objectives may not reflect demonstrated preferences. IRL instead infers rewards that explain demonstrations and guide policy learning [9, 11, 15, 32, 40, 50]. For example, MoE-BiEntIRL reconstructs graph-attack policies with interpretable rewards over graph and node features to generate adversarial samples [21]. However, using IRL to reconstruct fraudulent payment strategies for sequence generation remains underexplored. This requires constructing coherent payment actions from interdependent attributes and learning behavioral preferences conditioned on preceding payment outcomes. PaySeqGen therefore constructs a finite candidate list through history-guided sampling of observed joint payment changes and exploratory sampling. Its decomposable behavioral reward models feedback–action interactions, allowing the same action to score differently after different outcomes. Together, these designs support interpretable reconstruction of feedback-driven fraud strategies, providing additional training sequences and traceable evidence for online payment risk analysis.

6 Conclusion

In this paper, we propose PaySeqGen, an interpretable inverse reinforcement learning framework for reconstructing fraudulent payment strategies from observed online payment sequences. Specifically, our candidate construction mechanism combines history-guided and exploratory sampling to balance behavioral coherence and exploration. We further learn a decomposable linear behavioral reward that captures context-dependent preferences through feedback–action interactions. Through alternating reward learning and policy optimization, PaySeqGen translates observed behavior into additional training sequences and interpretable evidence for sequential fraud-pattern analysis. Experiments on large-scale real-world payment data demonstrate the utility of its generated sequences across multiple detector architectures, while a case study illustrates how reward contributions support the interpretation of multistep fraud behavior. Online A/B tests in Ant International's large-scale payment system further show Recall@2 gains of 3.19 and 2.84 percentage points for two production fraud detectors. These findings highlight PaySeqGen's practical value for strengthening fraud detection and supporting the analysis of complex fraud patterns in real-world web payment services.

References

[1] Pieter Abbeel and Andrew Y Ng. 2004. Apprenticeship learning via inverse reinforcement learning. In *Proceedings of the twenty-first international conference on Machine learning*. 1.

[2] L. Alfalahi and H. Nobanee. 2024. Conceptual Building of Sustainable Economic Growth and Corporate Bankruptcy. *Social Science Electronic Publishing* (2024). doi:10.2139/ssrn.3472409 [Online; accessed 2024-12-12].

[3] Erik Altman, Jovan Blanuša, Luc von Niederhäusern, Beni Egressy, Andreea Anghel, and Kubilay Atasu. 2023. Realistic Synthetic Financial Transactions for Anti-Money Laundering Models. In *Advances in Neural Information Processing Systems*, A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine (Eds.), Vol. 36. Curran Associates, Inc., 29851–29874.

[4] Tianqi Chen and Carlos Guestrin. 2016. Xgboost: A scalable tree boosting system. In *Proceedings of the 22nd acm sigkdd international conference on knowledge discovery and data mining*. 785–794.

[5] Dawei Cheng, Sheng Xiang, Chencheng Shang, Yiyi Zhang, Fangzhou Yang, and Liqing Zhang. 2020. Spatio-temporal attention-based neural network for credit card fraud detection. In *Proceedings of the AAAI conference on artificial intelligence*, Vol. 34. 362–369.

[6] Dawei Cheng, Yao Zou, Sheng Xiang, and Changjun Jiang. 2025. Graph neural networks for financial fraud detection: a review. *Frontiers of Computer Science* 19, 9 (2025), 199609.

[7] Annie D'souza, M Swetha, and Sunita Sarawagi. 2025. Synthetic tabular data generation for imbalanced classification: The surprising effectiveness of an overlap class. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 39. 16127–16134.

[8] Federal Trade Commission. 2025. *Consumer sentinel network data book 2024*. Technical Report. Federal Trade Commission.

[9] Chelsea Finn, Sergey Levine, and Pieter Abbeel. 2016. Guided Cost Learning: Deep Inverse Optimal Control via Policy Optimization. In *Proceedings of The 33rd International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 48)*, Maria Florina Balcan and Kilian Q. Weinberger (Eds.). PMLR, New York, New York, USA, 49–58.

[10] Peyré Gabriel and Cuturi Marco. 2019. Computational optimal transport with applications to data sciences. *Foundations and trends® in machine learning* 11, 5-6 (2019), 355–607.

[11] Divyansh Garg, Shuvam Chakraborty, Chris Cundy, Jiaming Song, and Stefano Ermon. 2021. IQ-Learn: Inverse soft-Q Learning for Imitation. In *Thirty-Fifth Conference on Neural Information Processing Systems*.

[12] Arthur Gretton, Karsten M Borgwardt, Malte J Rasch, Bernhard Schölkopf, and Alexander Smola. 2012. A kernel two-sample test. *The journal of machine learning research* 13 (2012), 723–773.

[13] Will Hamilton, Zhitao Ying, and Jure Leskovec. 2017. Inductive representation learning on large graphs. *Advances in Neural Information Processing Systems* 30 (2017).

[14] Jonathan Ho and Stefano Ermon. 2016. Generative Adversarial Imitation Learning. In *Advances in Neural Information Processing Systems*, D. Lee, M. Sugiyama, U. Luxburg, I. Guyon, and R. Garnett (Eds.), Vol. 29. Curran Associates, Inc.

[15] Jingyang Ke, Feiyang Wu, Jiyi Wang, Jeffrey Markowitz, and Anqi Wu. 2025. Inverse Reinforcement Learning with Switching Rewards and History Dependency for Characterizing Animal Behaviors. In *Proceedings of the 42nd International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 267)*, Aarti Singh, Maryam Fazel, Daniel Hsu, Simon Lacoste-Julien, Felix Berkenkamp, Tegan Maharaj, Kiri Wagstaff, and Jerry Zhu (Eds.). PMLR, 29608–29625.

[16] Vijay Konda and John Tsitsiklis. 1999. Actor-critic algorithms. *Advances in neural information processing systems* 12 (1999).

[17] Jianhua Lin. 1991. Divergence measures based on the Shannon entropy. *IEEE Transactions on Information theory* 37, 1 (1991), 145–151.

[18] Xiaofeng Lin, Chenheng Xu, Matthew Yang, and Guang Cheng. 2025. CTSyn: A Foundation Model for Cross Tabular Data Generation. In *The Thirteenth International Conference on Learning Representations*.

[19] Xin Liu, Xiyuan Chen, Chenglong Wu, Xuan Zong, Jun Zhou, and Dawei Cheng. 2026. Empowering Credit Risk Detection in Weixin Pay with Billion-Scale Deep Graph Learning. In *Proceedings of the 32nd ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 2*. 7679–7690.

[20] Daniele Lunghi, Yannick Molinghen, Alkis Simitsis, Tom Lenaerts, and Gianluca Bontempi. 2026. Fraud-RLA: A Reinforcement Learning Adversarial Attack Against Credit Card Fraud Detection. *IEEE Transactions on Dependable and Secure Computing* 23, 3 (2026), 7468–7481.

[21] Yuefei Lyu, Chaozhuo Li, Sihong Xie, and Xi Zhang. 2024. Enhancing robustness of graph neural networks on social media with explainable inverse reinforcement learning. *Advances in Neural Information Processing Systems* 37 (2024), 31736–31758.

[22] Sadaf R. Máté, J. Oláh, et al. 2019. The Effects of Accountability, Governance Capital, and Legal Origin on Reported Frauds. *Technological and Economic Development of Economy* (2019). doi:10.3846/tede.2019.10717

[23] Chandrahas Mishra and Ashish Kumar Mishra. 2026. Diffusion-based minority augmentation and confidence-calibrated tabular learning for imbalanced rare-event classification. *Information Systems* 140 (2026), 102716.

[24] Muhammad Ferjad Naeem, Seong Joon Oh, Youngjung Uh, Yunjey Choi, and Jaejun Yoo. 2020. Reliable fidelity and diversity metrics for generative models. In *International conference on machine learning*. PMLR, 7176–7185.

[25] Sai Shankar Narasimhan, Shubhankar Agarwal, Oguzhan Akcin, Sujay Sanghavi, and Sandeep P. Chinchali. 2024. Time Weaver: A Conditional Time Series Generation Model. In *Proceedings of the 41st International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 235)*, Ruslan Salakhutdinov, Zico Kolter, Katherine Heller, Adrian Weller, Nuria Oliver, Jonathan Scarlett, and Felix Berkenkamp (Eds.). PMLR, 37293–37320.

[26] Kyle Nickerson, Terrence Tricco, Antonina Kolokolova, Farzaneh Shoeleh, Charles Robertson, John Hawkin, and Ting Hu. 2022. Banksformer: A Deep Generative Model for Synthetic Transaction Sequences. In *Machine Learning and Knowledge Discovery in Databases: European Conference, ECML PKDD 2022, Grenoble, France, September 19–23, 2022, Proceedings, Part VI*. Springer-Verlag, Berlin, Heidelberg, 121–136.

[27] Takayuki Osa, Joni Pajarinen, Gerhard Neumann, J Andrew Bagnell, Pieter Abbeel, and Jan Peters. 2018. An algorithmic perspective on imitation learning. *Foundations and Trends® in Robotics* 7, 1-2 (2018), 1–179.

[28] Inkit Padhi, Yair Schiff, Igor Melnyk, Mattia Rigotti, Youssef Mroueh, Pierre Dognin, Jerret Ross, Ravi Nair, and Erik Altman. 2021. Tabular transformers for modeling multivariate time series. In *ICASSP 2021-2021 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 3565–3569.

[29] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, et al. 2019. Pytorch: An imperative style, high-performance deep learning library. *Advances in neural information processing systems* 32 (2019).

[30] Yiran Qiao, Ningtao Wang, Yuncong Gao, Yang Yang, Xing Fu, Weiqiang Wang, and Xiang Ao. 2025. Online fraud detection via test-time retrieval-based representation enrichment. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 39. 12470–12478.

[31] Cécile Rousseau, Tobia Boschi, Giandomenico Cornacchia, Dhaval Salwala, Alessandra Pascale, and Juan Moreno. 2025. Forging Time Series with Language: A Large Language Model Approach to Synthetic Data Generation. In *Advances in Neural Information Processing Systems*, D. Belgrave, C. Zhang, H. Lin, R. Pascanu, P. Koniusz, M. Ghassemi, and N. Chen (Eds.), Vol. 38, Main Conference. Curran Associates, Inc., 78402–78432.

[32] Silvia Sapora, R Devon Hjelm, Omar Attia, Alexander Toshev, and Bogdan Mazoure. 2026. GRACE: A Language Model Framework for Explainable Inverse Reinforcement Learning. In *International Conference on Learning Representations*, C. Vondrick, B. Hariharan, C. Raffel, L. Pinto, D. Yang, and A. Faust (Eds.), Vol. 2026. 15000–15027.

[33] John Schulman, Philipp Moritz, Sergey Levine, Michael Jordan, and Pieter Abbeel. 2015. High-dimensional continuous control using generalized advantage estimation. *arXiv preprint arXiv:1506.02438* (2015).

[34] Aditya Shankar, Yuandou Wang, Rihan Hai, and Lydia Y. Chen. 2026. Harpoon: Generalised Manifold Guidance for Conditional Tabular Diffusion. In *The Fourteenth International Conference on Learning Representations*.

[35] Linbo Shao, Huilin He, Yating Lou, and Dawei Cheng. 2026. Behavior-Grounded Semantic Enrichment for Financial Fraud Modeling and Reasoning. arXiv:2609.34211 [cs.AI]

[36] Lifeng Shen, Kai Syun Hou, Weiyu Chen, and James Kwok. 2026. Latent-to-Data Cascaded Diffusion Models for Unconditional Time Series Generation. In *The Fourteenth International Conference on Learning Representations*.

[37] Juntong Shi, Minkai Xu, Harper Hua, Hengrui Zhang, Stefano Ermon, and Jure Leskovec. 2025. Tabdiff: a mixed-type diffusion model for tabular data generation. In *International Conference on Learning Representations*, Vol. 2025. 37353–37375.

[38] Richard S Sutton, David McAllester, Satinder Singh, and Yishay Mansour. 1999. Policy gradient methods for reinforcement learning with function approximation. *Advances in neural information processing systems* 12 (1999).

[39] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. *Advances in neural information processing systems* 30 (2017).

[40] Zhenglin Wan, Xingrui Yu, David Mark Bossens, Yueming Lyu, Qing Guo, Flint Xiaofeng Fan, Yew-Soon Ong, and Ivor W. Tsang. 2025. Diversifying policy behaviors via extrinsic behavior curiosity. In *Proceedings of the 42nd International Conference on Machine Learning* (Vancouver, Canada) *(ICML'25)*. JMLR.org, Article 2460, 20 pages.

[41] Lei Wang, Shanshan Huang, Chunyuan Zheng, Jun Liao, Xiaofei Zhu, Haoxuan Li, and Li Liu. 2025. Mitigating Data Imbalance in Time Series Classification Based on Counterfactual Minority Samples Augmentation. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2* (Toronto ON, Canada) *(KDD '25)*. Association for Computing Machinery, New York, NY, USA, 2962–2973.

[42] Sheng Xiang, Mingzhi Zhu, Dawei Cheng, Enxia Li, Ruihui Zhao, Yi Ouyang, Ling Chen, and Yefeng Zheng. 2023. Semi-supervised credit card fraud detection

via attribute-driven graph representation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 37. 14557–14565.
[43] Yu Xie, Junkai Shan, Lifei Wei, Jiamin Yao, and MengChu Zhou. 2025. GAN-Based Hybrid Sampling Method for Transaction Fraud Detection. *IEEE Transactions on Knowledge and Data Engineering* 37, 10 (2025), 5905–5918.
[44] Lei Xu, Maria Skoularidou, Alfredo Cuesta-Infante, and Kalyan Veeramachaneni. 2019. Modeling tabular data using conditional gan. *Advances in neural information processing systems* 32 (2019).
[45] Hanlin Yang, Jian Yao, Weiming Liu, Qing Wang, Hanmin Qin, Kong hansheng, Kirk Tang, Jiechao Xiong, Chao Yu, Kai Li, Junliang Xing, Hongwu Chen, Juchao Zhuo, QIANG FU, Yang Wei, and Haobo Fu. 2025. Diverse Policies Recovering via Pointwise Mutual Information Weighted Imitation Learning. In *The Thirteenth International Conference on Learning Representations*.
[46] Jinsung Yoon, Daniel Jarrett, and Mihaela van der Schaar. 2019. Time-series Generative Adversarial Networks. In *Advances in Neural Information Processing Systems*, H. Wallach, H. Larochelle, A. Beygelzimer, F. d'Alché-Buc, E. Fox, and R. Garnett (Eds.), Vol. 32. Curran Associates, Inc.
[47] Lantao Yu, Weinan Zhang, Jun Wang, and Yong Yu. 2017. SeqGAN: sequence generative adversarial nets with policy gradient. In *Proceedings of the Thirty-First AAAI Conference on Artificial Intelligence* (San Francisco, California, USA) *(AAAI'17)*. AAAI Press, 2852–2858.
[48] Xinyu Yuan and Yan Qiao. 2024. Diffusion-ts: Interpretable diffusion for general time series generation. In *International Conference on Learning Representations*, Vol. 2024. 41582–41610.
[49] Hengrui Zhang, Jiani Zhang, Zhengyuan Shen, Balasubramaniam Srinivasan, Xiao Qin, Christos Faloutsos, Huzefa Rangwala, and George Karypis. 2024. Mixed-Type Tabular Data Synthesis with Score-based Diffusion in Latent Space. In *The Twelfth International Conference on Learning Representations*.
[50] Brian D. Ziebart, Andrew Maas, J. Andrew Bagnell, and Anind K. Dey. 2008. Maximum entropy inverse reinforcement learning. In *Proceedings of the 23rd National Conference on Artificial Intelligence - Volume 3* (Chicago, Illinois) *(AAAI'08)*. AAAI Press, 1433–1438.

A Pseudocode of PaySeqGen

Algorithm 1 summarizes alternating reward and policy learning. Expert and generated continuations share the same observed prefix. We write $\tilde{f}(\tau) = L^{-1} \sum_{t=p+1}^{T} f(s_t, a_t)$ for the average continuation features, where $L = T - p$. Rollouts are applied to each prefix in the sampled minibatch, and η_θ denotes the reward learning rate. The payment feedback interface and candidate sampling rules remain fixed throughout training.

Algorithm 1 PaySeqGen training

Input: Demonstrations $\mathcal{D}$; fixed feedback interface P_{env}; candidate counts K_h, K_e; policy updates per reward update U
Output: Policy π_ω and reward r_θ

Precompute joint changes between consecutive payments in $\mathcal{D}$
Initialize Actor–Critic parameters ω, ψ; set $\theta \leftarrow \mathbf{0}$
Initialize π_ω by behavior cloning with demonstrated actions included in candidate lists
for each training round **do**
  Sample expert trajectories $\mathcal{E}$ and initial prefixes
  Replay expert continuations with recorded feedback
  $\mathcal{P} \leftarrow \textsc{Rollout}(\pi_\omega, \text{prefixes})$
  $\mu^E \leftarrow |\mathcal{E}|^{-1} \sum_{\tau \in \mathcal{E}} \tilde{f}(\tau)$
  $\mu^\pi \leftarrow |\mathcal{P}|^{-1} \sum_{\tau \in \mathcal{P}} \tilde{f}(\tau)$
  $g_\theta \leftarrow \mu^E - \mu^\pi - \lambda\theta$
  $\theta \leftarrow \text{clip}(\theta + \eta_\theta g_\theta, -w_{\max}, w_{\max})$
  for $u = 1$ **to** U **do**
    Sample fresh $\mathcal{P} \leftarrow \textsc{Rollout}(\pi_\omega, \text{prefixes})$
    Compute $\tilde{r}_t$ using Eq. (15), with θ fixed
    Compute $\widehat{A}_t, \widehat{G}_t$ using Eq. (18)
    Update ω, ψ using Eq. (16), with $\widehat{A}_t, \widehat{G}_t$ detached
  end for
end for
return π_ω and $r_\theta(s, a) = \theta^\top f(s, a)$

function $\textsc{Rollout}(\pi_\omega, \text{prefixes})$
  for each remaining payment step **do**
    Construct K_h joint history-guided actions using Eq. (8)
    Add K_e constrained exploration actions to form $\mathcal{A}_t$
    Sample $j_t \sim \pi_\omega(\cdot \mid s_t, \mathcal{A}_t)$; set $a_t \leftarrow a_{t,j_t}$
    Sample $o_t \sim P_{\text{env}}(\cdot \mid s_t, a_t)$
    Record the decision; $s_{t+1} \leftarrow \text{Update}(s_t, a_t, o_t)$
  end for
  return generated continuations and decision records
end function

Advantage and return estimation. For a completed rollout, abbreviate the Critic prediction as $V_t = V_\psi(s_t, \mathcal{A}_t)$. With terminal values $V_{T+1} = \widehat{A}_{T+1} = 0$, GAE computes the following quantities backward:

$$\begin{aligned} e_t &= \tilde{r}_t + \gamma V_{t+1} - V_t, \\ \widehat{A}_t &= e_t + \gamma\lambda_{\text{GAE}}\widehat{A}_{t+1}, \\ \widehat{G}_t &= \widehat{A}_t + V_t. \end{aligned} \tag{18}$$

Here, e_t is the temporal-difference residual. The discount factor γ controls the weight of future rewards, while λ_{GAE} controls how strongly later residuals contribute to the advantage. Both targets are detached during optimization. Entropy is already included in $\tilde{r}_t$ and is not added again to the loss.

B Details of Baselines

We compare PaySeqGen with five generation baselines and a real-fraud replay control, and evaluate their augmentation effects on four target detectors.

B.1 Generation Baselines

The distribution-based baselines model the joint structure of payment records or sequences:

- **CTGAN** [44] uses conditional adversarial learning to capture dependencies among continuous and categorical attributes. We adapt it to generate payment continuations conditioned on an observed prefix.
- **Diffusion-TS** [48] generates multivariate time series through diffusion, combining temporal decomposition with reconstruction-based learning to capture sequence structure. We use conditional infilling to generate payment continuations.
- **TabDiff** [37] models numerical and categorical attributes through a unified diffusion framework with learnable noise schedules. We apply it to mixed-type payment sequences while keeping observed prefix fields fixed.

The policy-based baselines learn how to select successive payment actions:

- **BC** [27] learns a policy through supervised imitation of demonstrated actions. It uses PaySeqGen's policy architecture and generation candidate mechanism, but is trained

independently without IRL reward learning or subsequent RL optimization.
- **MoE-BiEntIRL** [21] combines mixture-of-experts modeling with IRL to reconstruct attack policies and explain behavioral preferences. We adapt its graph-specific action representation to complete payment actions.

We additionally include **Real-fraud replay**, which appends an equal number of resampled real fraud-target examples. This control distinguishes the benefits of generating behavioral variations from increasing the training weight of observed fraud examples.

B.2 Target Models

We use four detectors to assess augmentation across sequence, graph, and tree architectures:

- **RT**, our production Transformer-based payment fraud detector, serves as the primary target model. It models historical payment sequences to assess fraud risk.
- **TabFormer** [28] models payment sequences using hierarchical Transformers at the field and payment-event levels. It performs masked-attribute pretraining followed by fine-tuning for fraud detection.
- **GraphSAGE** [13] learns node representations by aggregating neighborhood information. We construct temporal, same-card, and same-device relations within payment histories to capture relationships among payments.
- **XGBoost** [4] is a well-established tree-based baseline in production fraud detection. It predicts risk from summary features of the payment history.

C Results on Public Dataset

We additionally evaluate on MS-FFSD [35, 42], a public financial transaction dataset with semantic enrichment and synthesized timestamps. The released data contain 77,881 payment records from 30,346 users. We split events chronologically in proportions of approximately 60%/20%/20%, keeping equal timestamps in the same split. Only explicitly legitimate or fraudulent targets contribute to detector supervision; unlabeled events may appear in the preceding history. This yields 16,965 training, 6,519 validation, and 6,159 test targets. Detectors use at most 19 strictly preceding payments, without target attributes or historical fraud labels, keeping the same with main experiment. Generator training is restricted to users with more than 20 training-period payments, yielding 3,662 complete windows, including 109 fraud-target windows. Each method adds 2,090 fraud-target examples, approximately 12% of the original detector training set. MS-FFSD provides merchant, location, transaction type, province, amount, and time, but lacks payment outcomes and reliable card/device identifiers. We therefore omit the payment feedback interface and feedback–action interactions and use 34 fixed features describing amounts, intervals, and merchant/location usage. The adapted policy still selects from 48 history-guided and 16 exploration candidates.

Table 3 reports changes relative to the unaugmented detectors. Many augmentation methods yield negative gains on this dataset. For generators, this may be partly due to the limited availability of long payment histories and the small number of fraud-target windows available for learning sequential behavioral patterns. Despite these constraints and the absence of feedback-related components, PaySeqGen remains competitive, achieving the best or second-best AUC and KS on RT, TabFormer, and GraphSAGE. These results support PaySeqGen's adaptability to financial datasets with limited behavioral information, even without payment feedback or feedback–action interactions.

D Complexity analysis

Let n be the maximum visible history length, K the candidate count, and N the number of training payment transitions. With fixed feature dimensions, network sizes, relation types, and path depth, history processing and policy scoring cost $O(n^2 + Kn)$. After preprocessing, indexed retrieval and candidate adaptation cost $C_{\text{cand}} = O(K_h \log N + Kn)$, with fixed fallback and resampling limits. The feedback interface caches its Transformer history representation and scores candidates through a shared prediction head, giving $C_{\text{env}} = O(n^2 + K)$. The total per-decision cost is therefore

$$C_{\text{step}} = O(n^2 + Kn) + C_{\text{cand}} + C_{\text{env}}. \tag{19}$$

For S decision steps processed during training, including expert replay and policy rollouts, the total cost is $O(SC_{\text{step}})$ after preprocessing and initialization. Generating an L-step continuation costs $O(LC_{\text{step}})$. Reward updates require only linear aggregation of precomputed features, avoiding the expert-responsibility and gating updates required by MoE-BiEntIRL [21]. Thus, reward learning preserves the asymptotic order of the underlying Actor–Critic computation. Training and generation take place during detector preparation and add no generation overhead to real-time payment scoring, supporting practical integration into online payment risk-control systems.

E Details of Fraud-Pattern Prevalence

Reward-derived rule. We derive a behavioral rule from the interval highlighted by PaySeqGen's relative rewards in Figure 3 and the corresponding feature contributions. The contributions link same-amount retries and card switching to preceding failures, while the payment records reveal two consecutive stages with lower amounts in the second stage. We encode this evidence as a deterministic window-matching rule. For a payment interval I, let $\ell_i = \log(1+m_i)$ and $N_{\text{sw}}(I)$ count changes between consecutive cards. We define the stage-validity predicate:

$$\text{Valid}(I) = [N_{\text{sw}}(I) \geq 2] \wedge \left[\max_{i\in I} \ell_i - \min_{i\in I} \ell_i \leq \epsilon\right], \tag{20}$$

where $\epsilon = 10^{-4}$. We scan windows of 6–12 consecutive payments, excluding windows with missing card identifiers. Each window is partitioned into two consecutive stages containing at least three payments each. For a window W of n locally indexed payments, each split $k = 3, \ldots, n-3$ defines $I_1 = \{1, \ldots, k\}$ and $I_2 = \{k+1, \ldots, n\}$. The matching predicate is:

$$\begin{aligned}\text{Match}_\kappa(W,k) = \text{Valid}(I_1) \wedge \text{Valid}(I_2) \wedge \left[\min_{i\in I_1} m_i > 0\right] \\ \wedge \left[\max_{i\in I_2} m_i \leq (1-\kappa) \min_{i\in I_1} m_i\right] \wedge \bigwedge_{i=1}^{n-1} [o_i \in \mathcal{O}_{\text{fail}}],\end{aligned} \tag{21}$$

Table 3: Augmentation results on public credit card fraud detection dataset MS-FFSD.

Augmentation	RT			TabFormer			GraphSAGE			XGBoost		
	ΔAUC	ΔRecall@2	ΔKS	ΔAUC	ΔRecall@2	ΔKS	ΔAUC	ΔRecall@2	ΔKS	ΔAUC	ΔRecall@2	ΔKS
Real-fraud replay	-3.874	-0.654	-2.781	+0.249	-0.292	+0.159	-1.221	**-0.905**	-0.327	-0.874	-0.028	-0.242
BC	-11.970	-0.710	-10.376	-1.583	-0.306	-1.531	-0.439	-1.184	+0.708	-1.175	+0.084	-1.019
CTGAN	-12.406	-1.058	-11.004	-4.368	-0.794	-3.842	-6.084	-1.351	-5.278	-0.404	**+0.292**	-0.234
Diffusion-TS	-4.191	-0.947	-2.809	+0.449	-0.111	+0.559	-0.548	-1.253	+0.258	-0.894	+0.014	-0.401
TabDiff	**-2.361**	-0.961	-3.202	-4.600	-0.473	-4.566	-3.620	-1.392	-3.241	**-0.337**	+0.265	-0.150
MoE-BiEntIRL	-9.098	-0.654	-8.137	+0.157	**-0.097**	+0.201	-0.399	-1.003	+0.960	-0.609	0.000	**-0.119**
PaySeqGen	-2.528	**-0.432**	**-1.748**	**+1.011**	-0.223	**+1.132**	**+2.194**	-1.267	**+2.825**	-0.704	-0.097	-0.486

Table 4: Pattern prevalence under different minimum amount reductions. Entries show matched/eligible trajectories and occurrence rates.

Reduction	Fraud-target	Legitimate-target
10%	43/885 (4.86%)	5/1,604 (0.31%)
20%	35/885 (3.95%)	4/1,604 (0.25%)
30%	31/885 (3.50%)	3/1,604 (0.19%)

where O_{fail} contains failure outcomes and $\kappa = 0.2$ specifies the minimum amount reduction. A trajectory matches if this predicate holds for any window and split. The final outcome within each window is unused, and no additional constraints are imposed on device identity, failure subtype, or payment intervals.

Prevalence analysis. The matching criteria and thresholds were fixed before cohort-wide counting. Table 4 reports prevalence under minimum amount reductions of 10%, 20%, and 30%. The pattern is consistently more prevalent among fraud-target trajectories, suggesting that the observed association is not specific to the default 20% threshold. These results provide exploratory evidence that the reward-highlighted pattern recurs beyond the illustrated case and is associated with fraudulent targets.

F Quality of Generated Fraud Samples

We compare PaySeqGen with generation baselines on behavioral fidelity and coverage. We derive fixed trajectory descriptors from payment attributes and outcomes, summarizing local behavior, feedback-dependent transitions, and card/device reuse structure.

- **Local W1.** We average the 1-Wasserstein distances [10] between real and generated distributions of trajectory-level statistics, with each distance normalized using its training-data scale.
- **Conditional and two-step JS.** Conditional Jensen–Shannon (JS) divergence [17] compares distributions of discretized action categories conditioned on preceding feedback, weighted by real feedback frequencies. Two-step JS compares consecutive feedback–resource-change patterns separately for cards and devices.
- **Path MMD^2.** Within each trajectory, directed edges connect payments to later payments sharing a card or device. We compare distributions of normalized path counts for paths with one to three edges using squared maximum mean discrepancy with a Gaussian kernel [12].
- **Coverage.** Following Naeem et al. [24], coverage is the fraction of real trajectory representations whose neighborhoods contain at least one generated representation. Representations combine the descriptors above. Neighborhood radii are distances to the kth nearest other real representation, with $k = 5$.

Table 5: Behavioral fidelity and coverage of generated fraud samples on the proprietary dataset. Values are means; bold denotes the best value in each column.

Method	Local W1 ↓	Cond. JS ↓	2-step JS ↓	Path MMD^2 ↓	Cov. (%) ↑
BC	**0.316**	**0.180**	0.486	**0.0313**	**24.44**
CTGAN	1.462	0.380	0.349	0.2352	1.39
Diffusion-TS	0.613	0.299	0.284	0.0574	10.21
TabDiff	1.086	0.335	0.240	0.0373	0.98
MoE-BiEntIRL	0.730	0.451	0.472	0.1407	4.97
PaySeqGen	0.560	0.369	**0.176**	0.0540	18.12

Lower discrepancy scores indicate closer agreement with real behavioral statistics, while higher coverage indicates that more real trajectories have nearby generated samples in the descriptor space.

Table 5 shows that PaySeqGen achieves the lowest two-step JS, second-lowest Local W1, and second-highest coverage. Two-step JS evaluates consecutive feedback–resource-change patterns, making it particularly relevant to reconstructing sequential payment behavior. PaySeqGen's advantage on this metric suggests that it better preserves these short behavioral patterns, providing evidence for its effectiveness in modeling feedback-dependent decisions. Performance nevertheless varies across metrics, with BC achieving the best results on the other four measures. These metrics assess agreement with selected reference statistics and neighborhood coverage in the chosen descriptor space. Such agreement alone does not establish whether generated samples introduce useful behavioral variations for detector training. We therefore use downstream detection performance as the primary evaluation criterion and report these metrics as complementary diagnostics. Together, the results provide supporting evidence of generated-sample fidelity, particularly in consecutive behavioral patterns relevant to fraud-strategy reconstruction.