ZeroRAG: Adaptive Semantic Peer Routing for Decentralized Retrieval-Augmented Generation

Petru Neague
Delft University of Technology
Delft, The Netherlands
p.m.neague@tudelft.nl

Quinten Stokkink
Delft University of Technology
Delft, The Netherlands
q.a.stokkink@tudelft.nl

Martijn de Vos
EPFL
Lausanne, Switzerland
martijn.devos@epfl.ch

Rafael Pires
EPFL
Lausanne, Switzerland
rafael.pires@epfl.ch

Mathis Randl
EPFL
Lausanne, Switzerland
mathis.randl@epfl.ch

Anne-Marie Kermarrec
EPFL
Lausanne, Switzerland
anne-marie.kermarrec@epfl.ch

Johan Pouwelse
Delft University of Technology
Delft, The Netherlands
j.a.pouwelse@tudelft.nl

arXiv:2611.07928v1 [cs.DC] 5 Nov 2026

Abstract

Retrieval-augmented generation (RAG) is a standard approach to enhance the reliability of large language model (LLM) outputs by retrieving documents relevant to a query and supplying them in the LLM prompt. Today's RAG systems collect, index, and serve these documents from centralized stores. In practice, however, documents are owned by organizations, devices, and individuals who cannot or will not surrender custody of them to a central party, *e.g.*, because of privacy, copyright or regulatory compliance. Existing modern decentralized retrieval techniques either reintroduce a coordinating authority, or rely on a global data structure that requires continuous network-wide coordination (*e.g.*, DHTs). We present ZeroRAG, a decentralized routing layer for document retrieval in RAG that performs semantic query routing across an overlay of peers without global data structures or centralized orchestration. Documents remain with the peers that own them, and each peer keeps only lightweight local state: a bounded set of neighbor advertisements and a cache of recently routed queries. Using these, ZeroRAG steers each query through a best-first walk that progressively rewires the overlay toward semantically relevant neighbors, so recall improves over time as the network self-organizes. We evaluate ZeroRAG using three real-world workloads and two decentralized baselines. We find that ZeroRAG outperforms baselines in terms of recall and required network messages, approaches the recall and LLM answer accuracy of a centralized document store, self-heals after catastrophic peer failures, and scales to 100 000 peers.

CCS Concepts

• **Information systems** → **Peer-to-peer retrieval**; **Top-k retrieval in databases**; • **Networks** → *Peer-to-peer networks.*

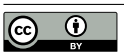

This work is licensed under a Creative Commons Attribution 4.0 International License.
Middleware '26, Tarragona, Spain
© 2026 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-2621-7/2026/12
https://doi.org/10.1145/3801927.3858412

Keywords

Decentralized retrieval, Retrieval-augmented generation (RAG), Approximate nearest neighbor search

ACM Reference Format:
Petru Neague, Quinten Stokkink, Martijn de Vos, Rafael Pires, Mathis Randl, Anne-Marie Kermarrec, and Johan Pouwelse. 2026. ZeroRAG: Adaptive Semantic Peer Routing for Decentralized Retrieval-Augmented Generation. In *27th International Middleware Conference (Middleware '26), December 14–18, 2026, Tarragona, Spain.* ACM, New York, NY, USA, 14 pages. https://doi.org/10.1145/3801927.3858412

1 Introduction

Large language models (LLMs) have transformed natural language processing because of their strong capabilities across tasks such as text generation, translation, and summarization [10, 56]. One of the most popular applications of LLMs is question answering [64]. In this domain, a fundamental challenge is to ensure the reliability of LLM-generated answers. However, LLMs are prone to hallucinations where they confidently generate false or misleading information [4]. This makes it difficult to trust the outputs of LLMs without double-checking factuality against other sources of truth [27]. These reliability concerns are particularly problematic in high-stakes settings, such as healthcare, finance, or law, where inconsistent or fabricated answers can have serious real-world consequences [24].

RAG is a popular approach to improving LLM reliability without retraining or fine-tuning the model [14, 28]. Rather than relying solely on parametric knowledge embedded in the model itself, RAG first retrieves relevant documents from an external source based on the user's query, then includes them in the prompt before generating a response. Both queries and documents are typically represented as high-dimensional embedding vectors capturing semantic meaning, and retrieval involves finding the embeddings nearest to the query, a process known as nearest neighbor search (NNS) [42]. This allows the LLM to ground its answers in up-to-date or domain-specific information, improving factual accuracy and consistency.

Existing RAG systems typically assume that documents are collected, indexed, and served from a centralized store under unified administrative control [14, 28, 62]. While this simplifies the RAG pipeline and results in low latency for document retrieval, it is not always a viable or desirable architecture. Data may be inherently distributed across organizations, devices, or individuals who cannot or will not surrender custody of their documents to a central party, whether for reasons of privacy, regulatory compliance, data ownership, or operational boundaries [19, 26]. A hospital network, for instance, cannot pool patient records into a shared index without additional access control [46, 55]. Similarly, a federation of research institutions may each hold proprietary datasets they are unwilling to freely share with other parties. Beyond ownership concerns, centralizing large, heterogeneous, and continuously evolving document collections introduces infrastructure costs and creates a single point of failure [17, 60, 61]. Despite these limitations, no general-purpose decentralized retrieval infrastructure exists for RAG: existing systems either assume centralized control, or distribute retrieval only within tightly managed federated or cloud settings that still require some coordinating authority.

We present ZeroRAG, a decentralized routing layer that enables document retrieval by routing queries across a decentralized overlay, without requiring global data structures or centralized orchestration. Documents remain with the peers that own them and peers themselves are in control whether they want to share documents with the party requesting them. Retrieval with ZeroRAG only uses local coordination. Each peer in ZeroRAG maintains only lightweight local state: compact summaries of the content exposed by neighboring peers and a bounded cache of previously-routed queries, and uses these to guide incoming queries to peers likely to hold documents relevant to that query. This is a key departure from prior decentralized nearest-neighbor systems, which route based only on what peers *hold*: ZeroRAG additionally exploits what peers have recently *asked*, referring query issuers toward peers whose recent query history is semantically aligned, not only their local inventory. For example, REMINDIN' [54] only learns from past responses and INGA [33] routes toward previous query issuers using explicit query topics. This design leaves owners in control of which documents they share, while peers route queries using document and query embeddings and peer referrals.

The key technical challenge ZeroRAG solves is effective routing and overlay management without global knowledge or coordination. Rather than consulting a centralized directory or broadcasting queries to all peers, a query in ZeroRAG involves a best-first walk over the overlay that adapts as queries flow through the network while also progressively rewiring peer connections toward semantically relevant neighbors. The result is a retrieval layer for RAG whose recall improves over time as the overlay organizes, scales gracefully with the number of peers and documents in terms of communication and computation overhead, and remains resilient to peer churn, without requiring network-wide maintenance.

We implement ZeroRAG and evaluate its efficiency and effectiveness for decentralized document retrieval in RAG pipelines, using three real-world workloads (AOL-IA [36, 43], Amazon Electronics [3], and PubMedQA [25]) and two decentralized baselines (Semantica [41] and DHT-LSH). Our experimental evaluation demonstrates that ZeroRAG consistently outperforms other decentralized baselines while requiring no globally maintained structure. On the AOL-IA workload, ZeroRAG achieves this at one tenth of the message overhead of the strongest baseline (DHT-LSH). ZeroRAG approaches the recall and end-to-end RAG answer accuracy of a centralized document store, without any global index, and 99% of queries locate the best supporting document in at most 13 peer contacts. Beyond static performance, our ablation studies show that ZeroRAG self-heals after catastrophic failure in which 50% of all peers crash simultaneously, quickly adapts as new topics are introduced, and scales to networks of up to one hundred thousand peers and documents. We make the following contributions:

- We introduce ZeroRAG, a decentralized routing layer for RAG that enables document retrieval by routing queries across a decentralized overlay without labeled data, shared data structures, or centralized orchestration (Section 3).
- We evaluate three real-world datasets, comparing against centralized and decentralized baselines on recall, RAG answer accuracy, message efficiency, resilience to peer churn, and topic introduction (Section 4). We show that ZeroRAG matches or exceeds the performance of decentralized alternatives while scaling gracefully to large peer populations.

2 Background and Problem Description

Our aim is to construct a decentralized document retrieval layer for RAG pipelines. We first describe the standard RAG workflow and then discuss the technical challenges that this work addresses.

2.1 Centralized RAG Workflows

Retrieval-augmented generation (RAG) is a standard and popular approach to improve the factual reliability of LLM outputs by supplying external context relevant to the user query directly in the prompt [28]. Whereas conventional language models depend entirely on parametric knowledge encoded during pre-training, RAG consults an external data source at inference time, yielding more accurate, up-to-date responses and reducing hallucinations without requiring costly retraining or finetuning. We illustrate the RAG workflow in Figure 1, using an example where a user asks a medical question, and explain the individual operations next.

The pipeline operates in two phases. In the offline phase, documents, videos, or other source material are preprocessed: each item is divided into chunks, and every chunk is encoded into a high-dimensional vector by an embedding model [45] ❶, with the resulting embeddings stored in a vector database ❷. For presentation clarity, we will refer to chunks as documents throughout this work. In the online phase, a user query ❸ is embedded and forwarded to the retriever ❹, which locates the most semantically similar stored vectors ❺ and fetches the corresponding documents from the vector database ❻. The original query and the retrieved context are then jointly passed to the LLM, which produces a response grounded in that evidence ❼, yielding a more accurate and context-aware answer than generation without retrieval.

The vector database is the central component enabling retrieval. Rather than matching exact keywords, it uses similarity search over high-dimensional embedding spaces. Prior to insertion, each document is encoded by a pre-trained machine learning (ML) model,

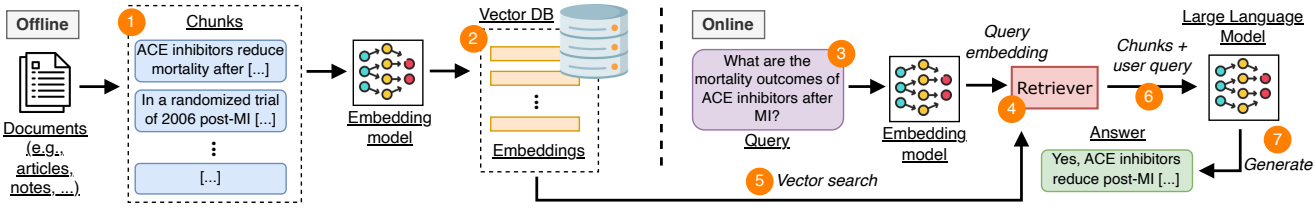

Figure 1: The standard RAG pipeline: offline indexing (left) encodes documents into a vector database; online retrieval (right) matches a query to relevant chunks before LLM answer generation.

typically a sentence transformer [45]. Encodings allow the database to capture semantic relationships between queries and documents, supporting retrieval beyond what keyword-based indices permit [35]. Incoming queries are embedded in the same space, and the database returns the nearest neighbors using vector distance. For tractability on large corpora, Approximate Nearest Neighbor (ANN) indexing techniques are used to accelerate this search [29].

2.2 Data Ownership and Availability

Our objective is to build a decentralized mechanism for the retrieval of relevant documents. A vast body of works exists surrounding peer-to-peer technologies and information retrieval [47]. Our work challenges two commonplace assumptions of these technologies and RAG: (1) that data can be freely ingested regardless of copyright or access control, and (2) that data must be globally available.

Copyright and access control. Both peer-to-peer retrieval systems and RAG architectures typically assume that information is generally available; RAG systems in particular oftentimes expect to freely ingest any corpus. We argue that this assumption will collapse under scrutiny when it comes to the intended access control or copyright law [48]. LLMs have been shown to be trained on scraped websites, whereas their terms of service prohibit this, and some even on pirated material, completely circumventing access control [37]. There have also been various court cases, that were lost by the companies behind the technology, where ML algorithms generated copyrighted material verbatim [5]. Strict access control is therefore a hard requirement in RAG systems, particularly in domains where data sovereignty is key, like healthcare or finance. By adhering to purely decentralized interactions, ZeroRAG is compatible with existing access control mechanisms and peers can refuse or allow other peers to access their information at their discretion.

Global data availability. Major decentralized systems, *e.g.*, Bitcoin [40], BitTorrent [44], and IPFS [11], are built around the assumption that data should be available everywhere, and they solve for this by replicating data across nodes. As previously discussed, replication is not always legal or permissible: a hospital, for instance, cannot freely copy patient records to arbitrary peers. Data has to stay with the peer that owns it and peers make no guarantees about availability or access. Rather than solving the replication problem, ZeroRAG routes queries to find data where it already resides, using only local queries and data set descriptions, without replicas of the underlying documents in the network. This comes at a cost: because ZeroRAG maintains no global index, it cannot offer the routing-time guarantees that a system like Kademlia [39] provides, and retrieval is therefore best-effort by design.

2.3 Towards Decentralized RAG

Our routing protocol should, given a user query, locate non-local owners of target or relevant documents to that query under bounded coordination and without centralized control. Forwarding queries to retrieve documents is a promising approach to decentralizing RAG workflows. This allows peers to perform access control on three levels: identity-based (or attribute-based) access control, query-based access control, and view-based (or purpose-based, or context-aware) access control. Our innovation therefore is in the *Retriever* component in the RAG workflow (step 4 in Figure 1), which still receives results from the local vector database, but also combines these with results from (remote) peers in an overlay network. This introduces two new challenges: (1) ZeroRAG must quickly find and query peers in the network that hold relevant documents to a given query, and (2) ZeroRAG must adapt to changes in the available documents over time, *e.g.*, due to peers joining and leaving or updating their local stores, causing semantic drift when the topology becomes misaligned with the current distribution of documents and queries.

Challenge 1: locality-aware peer sampling. Simply querying all peers in a decentralized overlay ensures maximum recall in terms of document retrieval, but cannot work on Internet-scale networks or data volumes. Instead, each peer should maintain connections to a small set of semantically proximate peers, such that relevant documents can be retrieved from currently connected neighbors alone, without sacrificing recall. This is the goal of *locality* optimization: steering peer connections toward neighbors whose document collections are likely to satisfy incoming queries [57]. For example, in a network of hospitals with different specializations, two pediatric hospitals are more likely to find similar cases in each other's documents than in those of hospitals with unrelated specializations. However, in semantically homogeneous settings, such as a network of *only* pediatric hospitals, all peers hold similar documents, so any neighbor is as good as any other and a locality-aware overlay brings no benefit over random connections. Therefore, we assume a heterogeneous setting. However, locality optimization cannot be treated as a one-time configuration: peers may leave the network, and the documents they hold may change over time, shifting what constitutes a relevant neighborhood. The first challenge addressed by ZeroRAG is *maintaining temporally adaptive peer connections that remain semantically relevant* as the network evolves.

Challenge 2: connection reconfiguration. Existing decentralized retrieval systems that achieve efficient lookup typically do so by maintaining a globally shared index structure, such as a semantic trie [41] or a distributed hash table (DHT) [39]. These structures require continuous network-wide reorganization as peers join and

leave and as document collections evolve. In our setting, where coordination must remain strictly local, maintaining such a global structure is not an option. Thus, the second challenge that ZeroRAG takes on is to *reconfigure neighborhood connections directly* as queries and documents evolve. Our approach is comparable to how border gateway protocol (BGP) routes packets without a global directory: each router maintains only local routing tables that are updated incrementally as traffic flows through the network, and efficient paths emerge from these local updates alone.

3 ZeroRAG Architecture and Design

ZeroRAG is a decentralized retrieval layer that routes queries to peers that hold relevant documents. In this section, we first discuss the actors and assumptions, and then peers' local algorithms for neighborhood maintenance and query answering.

3.1 System Model

Our system model focuses on query routing; we empirically evaluate the accuracy of the downstream RAG answer, but do not model the quality of the citation or the reduction of hallucinations.

Network and communication. ZeroRAG coordinates decentralized document retrieval by maintaining a permissioned decentralized overlay network and reconfiguring connections between peers when queries are issued. The incoming and outgoing connections by each peer form a directed graph. Peers joining the ZeroRAG network can bootstrap into the network by connecting (outgoing) to several random peers, which can be obtained by using existing peer sampling services [18, 23]. Threats targeted at the network level, like Sybil [13] and Eclipse [21] attacks, are beyond the scope of this paper, but we note that peer sampling mechanisms can overcome these after a cold-boot period of several minutes [52].

ZeroRAG only uses primitive request-response messages (presented shortly), does not maintain shared data structures, or require synchronization steps between peers. Therefore, messages may be asynchronous or drop. Likewise, peer connections can drop at any time, *e.g.*, when a peer goes offline or crashes. We assume the peer sampling mechanism can detect dropped connections and that this can be implemented with minimal overhead using a ping-pong scheme, message-reception acknowledgments and timeouts. Beyond this, we do not investigate Byzantine behavior of peers, like attempts to bias others' search results or black hole attacks [1].

Documents and embeddings. In ZeroRAG, peers exchange embeddings of documents and queries. The embeddings are used for indexing and network reconfiguration. The transfer of the documents corresponding to those embeddings is outside ZeroRAG. Exchanged embeddings need to be stored by peers and each peer operates a local vector database for this. We assume that peers use the same embedding model ϕ_e, so embeddings produced by different peers remain comparable. As embeddings are local, routing does not depend on centrally assigned labels or a global topic taxonomy, which was the common approach in the late 2000's [32, 50].

Peers are expected to accrue documents with attribution over time, *i.e.*, a verifiable record of which peer a document originates from, as they exchange them after ZeroRAG matches their embeddings to queries. The attribution itself is out of the scope of ZeroRAG and can be implemented with a digital signature. Within the scope of the paper we are interested in routing, not in the manner in which this history is built. Therefore, we shortcut the emergent effect of collecting documents by using a `HISTORY()` remote procedure call, which gathers the history of accrued documents from neighbors *in memory*. When considering deployment of ZeroRAG, this can (and should) be replaced with a performant storage layer.

3.2 High-Level Description

ZeroRAG is built around the insight that if a peer's neighbors hold semantically similar content, then a query is best served by greedily walking toward the most similar peer at each step. To this end, whenever a peer i connects to another peer j, j sends a set of *advertisements* to i. These advertisements are compact summaries, in the form of document embedding vectors, of the content j locally holds and wants to advertise for routing. This allows owners to retain their documents and decide whether and what to disclose, which is not the same as privacy-preserving retrieval: advertisements, query embeddings, and issuer referrals may reveal information even when documents are withheld. Now, some peer p aims to obtain the top-k documents that are relevant to a query q. To facilitate this search, peer p maintains a *frontier*, which is a ranked list of candidate peers to visit, ordered by their estimated relevance to the query. The best-first walk proceeds by repeatedly popping the highest-scoring peer from the frontier, sending it the query, and collecting its response: a set of relevant local document embeddings together with new *referrals* to that peer's own neighbors. These referrals are scored against the query and put back into the frontier, extending the search into previously unvisited parts of the network. Meanwhile, peer p keeps track of the document embeddings that it has seen. After a fixed number of steps, p will inspect the these embeddings and compare their relevance against query q. Peer p then retrieves the top-k relevant documents from the peers that offer them.

ZeroRAG continuously reconnects the overlay itself as queries are issued. After each query completes, the querier adds a directed edge toward whichever peer returned the most relevant result, evicting the oldest existing connection if the neighbor list is full. This implicit feedback loop causes the overlay to self-organize around semantic affinity: peers that frequently serve useful results for a given topic accumulate more incoming connections, and future queries on that topic are naturally steered toward them earlier in the walk. As the overlay network reconnects, peers are increasingly likely to quickly find the documents they need, without any global coordination or shared data structure. In addition, each peer keeps a bounded routed-query memory: a cache of embedding vectors from queries it has recently responded to, which were issued by other peers. When scoring an incoming query, a peer consults this memory alongside its neighbor advertisements, and can refer the querier not only toward peers with relevant documents but also toward peers that have recently asked similar questions, and are thus likely to be connected to, or themselves hold, relevant content.

Illustrative example. Figure 2 (left) shows the progression of a document search, issued by peer a ❶. Peer a iteratively sends search queries to peers in its frontier and receives replies with relevant documents and referrals. We show a document search where four peers are queried by a (❷- ❺). After the search completes, peer a adds peer i as neighbor since it has the most relevant documents ❻.

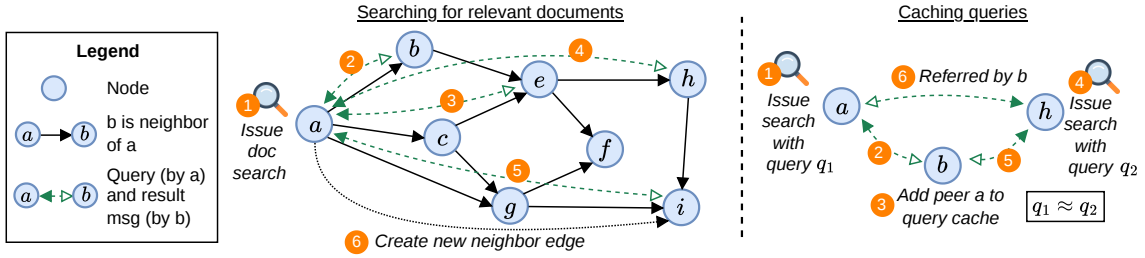

Figure 2: *(left)* Example of 4 query steps of a document searched by peer a with query q. After the search concludes, peer a adds peer i, which has the most relevant document to a's query, as neighbor. *(right)* The dynamics of query caching in ZeroRAG.

Figure 2 (right) shows query caching. Peer a issues a search with query q_1 and queries peer b (❶,❷). Peer b adds peer a and the incoming search query to its local query cache ❸. When another peer h later issues a query q_2, which is similar to q_1 (*i.e.*, their embeddings are close) and interacts with peer b, b refers peer h to peer a as a is likely to have found relevant documents for q_2.

3.3 Technical Description

We describe the concrete steps in ZeroRAG through the roles of a peer u that aims to retrieve k documents relevant to a query q (with embedding e_q) by repeatedly sending QUERY messages to candidate peers in the frontier, and a peer p that receives such a QUERY message. These roles are purely explanatory: any peer can fulfill either role, even at the same time. Algorithm 1 explains the search procedure from the perspective of peer u and Algorithm 2 shows how peer p handles an incoming query.

Local state. To facilitate routing, every peer i stores a local state σ_i, given by the 4-tuple (D_i, N_i, A_i, M_i). The documents stored in the local database of peer i are represented by D_i. The neighbors of i are given by N_i, which is a FIFO list with limited capacity δ_n. Insertion and eviction from this list is based on the most recent highest-scoring query. The collection of advertisements collected by i, of its neighbors, is given by A_i. Lastly, the collection of the most recent δ_m received pairs (u, e_q), as a result of acting as the role of the queried peer p, form the chronological FIFO list M_i. The local state is designed for the simplest effective implementation. FIFO avoids maintaining estimates of each neighbor's usefulness: eviction needs only insertion order [57]. Maximum scoring allows one close match to guide routing, despite unrelated vectors [53]. One neighbor per query limits connection changes and advertisement exchanges [51].

Both algorithms 1 and 2 use *scores* of relevance. These scores are simply the cosine similarity of two embeddings, or a score of minus infinity if the two have not been compared. ZeroRAG attempts to localize documents with the highest score in the network.

Retrieving documents. A peer u invokes the `retrieveDocuments` procedure, given in Algorithm 1, to retrieve up to k of the most-relevant documents to the query q. The remainder of the algorithm will use the embedding e_q of q (line 2) to query neighbors of u. If the required documents are not found locally, the peer constructs a frontier F of uncontacted neighbors, and successive neighbors of neighbors. It initially consists of pairs of currently-connected neighbors N_u of u and scores of minus infinity to signify that they have not been queried (line 5). During the following best-first-search loop (line 7), F may already receive scores for identities from previously queried peers. The search loop successively visits any (there may be multiple with the same score) most high-scoring peer x in the frontier, up to C peers, inserting x into the visited peer set Q (line 7-13). Peer u then sends a QUERY message to peer x, containing the query embedding e_q and the value of k (line 14). Peer u awaits a response QUERYRESULT message from x, containing the top remote results and top referrals of x (line 15). The document results are merged into the set R, which is then trimmed down to the top-k (line 16). Returned identities and their scores are inserted into F, or the identity's score is updated if it was already present, while the best-known score and its (potentially unvisited) associated identity is saved in B (lines 17-22). After the search loop is finished, the network connections are reconfigured (line 24) and the documents are retrieved, by their embedding, from the peer that stores them and returned (line 25 and 26).

Neighborhood reconfiguration. As a result of a query and its retrieved document embeddings R and best identity-score pair B, the current neighbors of u are updated using the `updateNeighbors` procedure, given in Algorithm 1 (line 28). The identity a of the peer that has the highest-scoring document embedding in R is retrieved (line 31). Otherwise, if R was empty, the identity of the peer a that was advertised as the best score is retrieved (line 33), if there was one. If an identity a is available, it is connected to and if that causes the current number of neighbors to exceed δ_n the oldest neighbor is removed from the list (lines 35-40).

Query handling. We next explain how peer p handles an incoming QUERY message, also see Algorithm 2. First, peer p queries its local vector database of to retrieve document embeddings H based on e_q (line 2). Then, p forms a pool of routable, known embeddings E by all known pairs of document embeddings and query embeddings of neighbors, coupled to the identity of the neighbor identifier v that sent them (line 3). The scores S of identities are initialized with minus infinity (line 4) and then set to the highest score over the elements in E that are associated with it (lines 5-9). The top r peers are then taken from S (line 10) to be returned as the referrals $Refs$ later. The pair of the querying peer u and its query embedding e_q is stored in the routed-query memory M_p (line 11), to support future routing of queries similar to u. The oldest entry in M_p is evicted if the number of these tuples exceeds δ_m size (lines 12-14). Finally, p sends a QUERYRESULT, with the local top-k embeddings in H and the top-r referrals in $Refs$ (line 15).

Algorithm 1 Retrieving top-k documents for query q of peer u.

Require: C: contact budget, r: referral limit, δ_n: max amount of neighbors each peer maintains
Require: Persistent state $\sigma_u = (D_u, N_u, A_u, M_u)$ ▷ local documents D_u, outgoing neighbors N_u, neighbor advertisements A_u, and routed-query memory M_u.
1: **procedure** RETRIEVEDOCUMENTS(q: query, k: result limit)
2: $e_q \leftarrow \phi_e(q)$ ▷ generate embedding of q
3: $Q \leftarrow \emptyset$ ▷ peers already contacted
4: $R \leftarrow \emptyset$ ▷ document results found so far
5: $F \leftarrow \{(u, -\infty)\}$ ▷ frontier of peers to contact
6: $B \leftarrow \emptyset$ ▷ best referral score seen for each peer
7: **while** $F \neq \emptyset$ **and** $|Q| < C$ **do**
8: $(x, _) \leftarrow \arg\max_{(v,s)\in F} s$ ▷ peer with highest score
9: F.REMOVEPEER(x)
10: **if** $x \in Q$ **then**
11: **continue**
12: **end if**
13: $Q \leftarrow Q \cup \{x\}$
14: **send** QUERY(e_q, k) message to peer x
15: **receive** QUERYRESULT(H_x, $Refs_x$) message from peer x
16: $R \leftarrow$ TOPK($R \cup H_x, k$) ▷ keep the best document results
17: **for all** $(v, s) \in Refs_x$ **do**
18: **if** $v \notin Q$ **and** ($B[v]$ is undefined **or** $s > B[v]$) **then**
19: $B[v] \leftarrow s$
20: F.UPDATE((v, s))
21: **end if**
22: **end for**
23: **end while**
24: UPDATENEIGHBORS(R, B) ▷ Overlay maintenance
25: $D \leftarrow$ OBTAINDOCUMENTS(R) ▷ retrieve docs from docrefs
26: **return** D
27: **end procedure**
28: **procedure** UPDATENEIGHBORS(R: found docrefs, B: ref scores)
29: $a \leftarrow \perp$
30: **if** $R \neq \emptyset$ **then** ▷ We found at least one document
31: $a \leftarrow$ peer returning the highest-scoring document in R
32: **else if** $B \neq \emptyset$ **then**
33: $a \leftarrow \arg\max_v B[v]$
34: **end if**
35: **if** $a \neq \perp$ **and** $a \neq u$ **then**
36: N_u.ADD(a)
37: **if** $|N_u| > \delta_n$ **then**
38: N_u.EVICT() ▷ Evict oldest neighbor
39: **end if**
40: **end if**
41: **end procedure**

Algorithm 2 Query handling at peer p.

Require: r: referral limit, δ_m: memory size limit
Require: Persistent state $\sigma_p = (D_p, N_p, A_p, M_p)$
1: **On receive** QUERY(e_q, k) message **from** peer u
2: $H \leftarrow D_p$.SEARCH(e_q, k) ▷ local document search
3: $E \leftarrow A_p \cup M_p$ ▷ advertisements and past routed queries
4: $S[v] \leftarrow -\infty$ **for** every $(v, _) \in E$
5: **for all** $(v, z) \in E$ **do** ▷ compute best scores for peers
6: **if** $v \neq u$ **then**
7: $S[v] \leftarrow \max\{S[v], \langle e_q, z\rangle\}$
8: **end if**
9: **end for**
10: $Refs \leftarrow$ TOPK($\{(v, S[v]) : S[v] > -\infty\}$, r)
11: M_p.ADD((u, e_q))
12: **if** $|M_p| > \delta_m$ **then**
13: M_p.EVICT() ▷ Evict oldest routed-query memory item
14: **end if**
15: **send** QUERYRESULT(H, $Refs$) message to peer u

Retrieval parallelization. To reduce end-to-end retrieval latency, it is possible to parallelize Algorithm 1. Instead of performing a best-first search, a breadth-first search can send and await up to f QUERY and QUERYRESULT messages in parallel. As requests and responses can be asynchronous, the number of visited peers that have responded might then overshoot the contact budget C by the node fanout of the connection graph. In the remainder of this paper, we use the sequential version of Algorithm 1 to avoid additional variance in our metrics, which do not include wall-clock time.

Cleanup. Our algorithms do not show the response of ZeroRAG to churn, as signaled by the peer sampling. When peer p leaves, peer u removes p from N_u and deletes entries associated with p from A_u and M_u. Documents in D_u do not need cleanup and other peers in the network will now start being routed to u instead of p.

ZeroRAG only has a QUERY and a QUERYRESULT message that can be dropped. The dropping of either of these can be remedied by waiting until some timeout when receiving the response message in Algorithm 1 (line 15). If the timeout is triggered, the query results are an empty list of documents and an empty list of referrals.

Scalability and topology convergence. With fixed-size encodings, each query sends at most C requests and C replies, carrying $O(C((k+1)d+r))$ bytes for d-dimensional embeddings, k results and r referrals. If each neighbor advertises at most a vectors, removed when that neighbor is evicted, persistent routing state occupies $O(\delta_n + d(\delta_n a + \delta_m))$ bytes, where δ_n and δ_m cap neighbors and cached queries. Each search adds at most one connection, transferring $O(ad)$ advertisement bytes. The total routing state and traffic for one query per peer are $O(n)$ for n peers, excluding document transfer, retries, and background maintenance. Regarding topology convergence, stability can only be achieved if the number of connected peers matches the number of interests. At capacity δ_n, adding a new neighbor for the current query evicts the oldest connection, potentially displacing a peer useful to other interests; in this configuration therefore the topology cannot feasibly converge.

4 Experimental Evaluation

We implement ZeroRAG as routing middleware between a network overlay and local document stores. Our experimental evaluation answers the following five questions:

(1) What is the performance of ZeroRAG on real-world data sets in terms of recall and accuracy (Section 4.2)?
(2) How does the end-to-end accuracy of RAG change when using ZeroRAG (Section 4.3)?
(3) How resilient is ZeroRAG to network-wide failure (Section 4.4)?
(4) What is the efficacy of ZeroRAG's inputs and what is the effect of their ablation (Section 4.5)?
(5) How quickly does ZeroRAG make queries routeable when documents with new topics are introduced (Section 4.6)?

4.1 Experimental Setup

We next describe our experimental setup, including our workloads, baselines and implementation details.

4.1.1 Workloads. We evaluate ZeroRAG and baselines using three real-world workloads: AOL-IA [36, 43], Amazon Electronics [3], and PubMedQA [25]. The three workloads cover complementary retrieval settings: single-target and multi-relevant retrieval, real and synthetic user behavior, and general-domain, e-commerce, and biomedical query distributions. We additionally use PubMedQA for an end-to-end RAG evaluation with ZeroRAG, extending our evaluation beyond routing performance alone.

AOL-IA provides a query-and-click stream with user ids, query text, clicked documents, and timestamps. The full workload contains 657 426 users, 428 157 documents and 2 594 705 total queries. We adapt the workload to our setting by selecting the 10 000 users with the most queries, together performing 2 171 460 queries. Prior to running this workload, we populate the local document stores with the documents that the user has clicked. During our evaluation, we then replay these events in a randomized order by having peers search for the document they are interested in (*e.g.*, $k = 1$), without them looking into their local database first. The single-target retrieval objective (find the exact document) enables us to test the effectiveness of query routing in ZeroRAG with a binary objective score (whether the document is found or not).

Amazon Electronics is an adaptation of the personalized product-search benchmark introduced by Ai et al. [3]. This benchmark contains 63 001 product descriptions, and each product contains, besides some basic metadata, a category path (*e.g.*, "Camera, Photo → Digital Camera Lenses"). The workload authors synthetically constructed search queries by taking each product's category path and flattening it into a query string, stripping stopwords and duplicates. For example, "Camera, Photo → Digital Camera Lenses" becomes the query "photo digital camera lenses." Thus, queries are categorical and fairly repetitive across users. The queries were then mapped to users using the reviews of purchased products: if a reviewer purchased a product associated with a category-derived query, a corresponding reviewer–query pair is made.

In our evaluation, each reviewer in the Amazon Electronics benchmark becomes a peer in the network. Each reviewer–query pair becomes one simulation event, *i.e.*, one search query issued by one peer. The relevant products for that reviewer–query pair become the documents that peer is searching for. The retrieval objective for ZeroRAG is to recover as many relevant products as possible from other peers, rather than to recover a single clicked item. Those documents are distributed across other peers in the network, based on which peers reviewed those products. Because queries come from product categories rather than individual user intent, the same query string will naturally recur across many different reviewers. This means the experiment ends up simulating a setting of repeated, overlapping product-search demand, where many peers are asking semantically similar things at different times.

PubMedQA is a biomedical question-answering dataset built from abstracts in the PubMed database [25]. We adopt the discrete-answer subset of PubMedQA with 1000 items, in which each item contains a biomedical question, an abstract, and a categorical answer (`yes`, `no`, or `maybe`). To increase query-level semantic variety, we supplement each original question with five rephrased queries generated using the GPT-5-nano model, resulting in six query variants (the original question and five rephrasings).

To assign PubMed documents to peers, we use Medical Subject Headings (MeSH) terms, assigned to PubMed articles, as biomedical topic labels. We sample a MeSH-level preference vector $\boldsymbol{\theta}_u$ from a Dirichlet($\boldsymbol{\alpha}$) distribution, where $\boldsymbol{\theta}_u$ defines user u's preference distribution over all 23 904 MeSH topics. This produces user profiles that are biased toward a small number of topics. We fix $\boldsymbol{\alpha} = 7\times10^{-5}$, where a peer gets on average 15.9 documents on the first topic, 6.6 on the second topic, 3.4 on the third topic and so on. Given a target number of documents t, we sample the number of documents assigned to each MeSH topic according to $\boldsymbol{\theta}_u$. We fix $t = 30$ unless otherwise stated. We then uniformly sample documents from each selected MeSH topic, and for each document uniformly choose one of its six query variants. Under this user-model, most peers are interested in 1-3 main topics with the interest dropping off for other topics. Each document of a peer i also carries one randomly selected query variant. This (peer, document, query) triple defines a simulation event: peer i issues the query, and the retrieval objective is to locate the corresponding document on another peer.Each peer issues exactly as many queries as it holds documents. Unless otherwise stated, experiments with the PubMedQA workload contain 50 000 simulated peers and 30 queries per peer, leading to 1.5 million query-document events, issued in random sequence.

4.1.2 Implementation. The experiments were simulated by following a timeline of 'events', one event representing a peer sending a query which aims to retrieve one or more target documents. We fix the contact budget $C = 50$, meaning that a peer contacts at most 50 peers during document retrieval. We used the all-mpnet-base-v2 embedding model for embedding generation of documents and queries [49]. For embedding generation of documents in the AOL-IA workload, we concatenated the `title`, `body`, `text`, and `url` values from the data set using the string "\n\n". For documents in the Amazon Electronics workload, we concatenated their `title`, `description`, `categories`, `brand`, and `price` in a similar fashion.

The overlay topology is implemented as a global directed connection graph, where each graph node represents a peer. Message passing between nodes is implemented using remote procedure calls. Document embeddings are kept in a single unified memory and the local storage of peers are references in that memory.

Figures 3, 4, 6, and 7(b) and all ablation tables use five runs per configuration: a randomized starting topology for each run of ZeroRAG, and order of the peers in the tree for Semantica. For matching windows, bins, or table summaries, we compute the five-run mean $\bar{x}$ and sample SD in Figures 4 and 7(b), and the tables. Figures 3 and 6 instead show the mean with 95% Student-t confidence intervals; DHT-LSH was not repeated because its hash functions are fixed, retrieval does not depend on randomized peer connections, and unrepeated baselines have no uncertainty bars.

4.1.3 Hardware. All simulations were performed using 12 computing nodes which are equipped with two AMD EPYC 7282 processors (16 cores per socket, 2.8 GHz), for 32 physical cores in total.

4.1.4 Baselines. We evaluate ZeroRAG against two related baselines that together occupy different parts in the design space.

Semantica [41] is the closest prior decentralized semantic retrieval system and organizes peers with a globally maintained semantic trie. We use the source code of Semantica as published in the author's Github repository. For the hierarchical navigable small world (HNSW) implementation we utilize hnswlib [38].

DHT-LSH is a baseline that pairs a DHT (*e.g.*, Kademlia [39]) with locality-sensitive hashing (LSH) schemes to provide ANN search in fully decentralized settings. DHT-LSH relies on p-stable hash functions that map embedding vectors into discrete buckets such that semantically similar vectors collide into the same bucket with high probability [12]. To improve recall, DHT-LSH maintains L independent hash tables, each using a composition of k such functions with bin width w. At indexing time, each document embedding is hashed across all L tables, and the resulting bucket identifiers are used as keys in the DHT, which distributes storage of bucket contents across peers. At query time, the query embedding is hashed through the same L tables, yielding L bucket keys, each of which requires a separate DHT lookup to retrieve candidate documents. We adopt the suggested hyperparameter values from the original LSH paper [12], fixing $k = 10$. We adopted $w = 3$ instead the of the $w = 4$ used by the authors because $w = 4$ placed 44.7% of the AOL-IA documents in the single largest bucket.

4.1.5 Metrics. Our primary success metric is *recall*: the percentage of relevant documents retrieved from the entire set available in the network, excluding documents already held locally by the querying peer. For AOL-IA, where each query has a single target document, this reduces to a binary success rate. For Amazon Electronics, where a query may have hundreds of relevant documents, we instead report the *fraction of available relevant documents retrieved* within the routing budget. For PubMedQA we additionally measure end-to-end RAG accuracy: the fraction of questions answered correctly when the retrieved top-k documents are supplied as context. This assesses whether routing quality translates into downstream generation quality. For brevity, in the remainder of this section, we will use the notation *recall@k* for the recall when the retrieved top-k documents are supplied as context. Unless stated otherwise, we report recall averaged over batches of 1000 queries to smooth per-query variance and surface trends in routing performance as the overlay evolves over time.

4.2 Effectiveness of Document Retrieval

We first evaluate the effectiveness of document retrieval with different configurations of ZeroRAG, and the Semantica and DHT-LSH baselines, using the AOL-IA and Amazon Electronics workloads. For DHT-LSH, we vary the number of buckets $L \in \{10, 20, 50\}$.

4.2.1 AOL-IA workload. We issue the search queries in the AOL-IA workload, in randomized order, and measure recall@1, which indicates whether the document has been retrieved or not. Figure 3 shows average recall@1 for each query batch, for all evaluated systems. ZeroRAG obtains an average recall@1 of 0.91 for the last query batch. In contrast, this same number for Semantica is 0.65 and significantly lower. We attribute this to the fact that peers in Semantica have a single (averaged) embedding that is used to advertise their documents, whereas in ZeroRAG peers share more than one non-averaged embedding. DHT-LSH with $L = 50$ approaches ZeroRAG with an average recall@1 of 0.80, while DHT-LSH with $L = 20$ and $L = 10$ reach only 0.56 and 0.41, respectively. However, the drawback of LSH methods is high communication and computation cost. Every bucket check requires traversing the DHT: the overhead in our 10 000-peer network, for instance, is that each search query requires approximately $50 \cdot log_2(10000) \approx 664$ return messages (to another peer and back). In contrast, ZeroRAG only requires 50 return messages per search query for similar performance. Furthermore, the largest DHT bucket for $L = 50$ holds 25% of all documents, therefore every query routed to that bucket has to be compared against all of them.

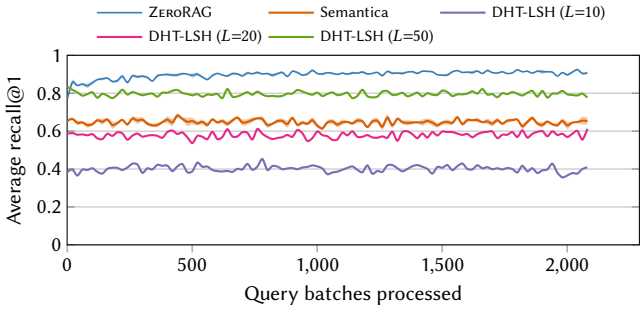

Figure 3: Average recall@1 for AOL-IA query batches, for ZeroRAG and different baseline configurations.

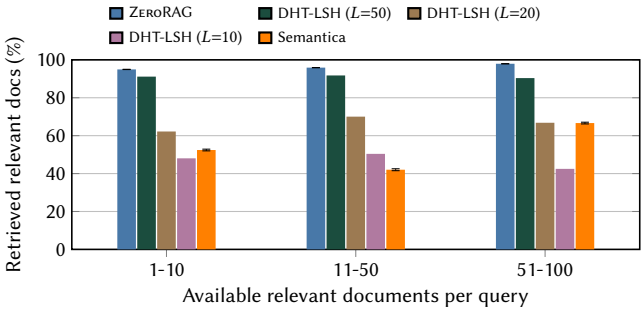

Figure 4: Percentage of retrieved results, binned by the number available results for a query, for the Amazon workload.

Takeaway (Figure 3). Using the AOL-IA workload, ZeroRAG shows superior retrieval effectiveness compared to DHT-LSH and Semantica, at a lower message complexity.

4.2.2 Amazon Electronics workload. We use the same setup as in the experiment in Section 4.2.1 but run the Amazon Electronics workload. However, queries in this workload may have many relevant documents rather than just a single result. Therefore, recall@1 is not well defined and we address this by considering different values of k when computing recall@k for three different bin sizes, based on how many relevant documents there are for each query: if

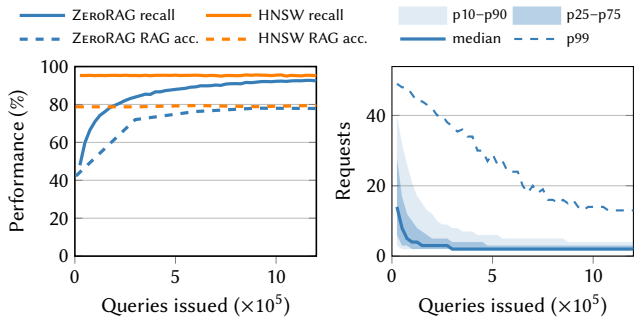

(a) Recall and RAG accuracy

(b) Requests until the first relevant document is found

Figure 5: (a) The PubMedQA end-to-end RAG performance, using ZeroRAG and HNSW. (b) the distribution of number of ZeroRAG requests before the first relevant document.

there are 1-10 available relevant documents, we use recall@10, recall@50 for queries with 11-50 available documents, and recall@100 for queries with 51-100 available documents.

Figure 4 shows the fraction of retrieved relevant documents for ZeroRAG and baselines, grouped by the number of available documents for each query. ZeroRAG retrieves 94.8-98.9% of available relevant products across all three bins. DHT-LSH benefits from broader hash-table coverage but falls behind ZeroRAG: with w=3.0, k=10, L=50 it reaches 90.3-91.7%, with L=20 being 62.1-70.0%, while a value of L=10 drops it to 42.5-50.3%. Semantica does match the performance of DHT-LSH L=20 in the largest bin (66.7%) but falls behind again in the smaller and middle bins (41.7-52.9%).

Takeaway (Figure 4). In terms of recall, ZeroRAG consistently outperforms DHT-LSH and Semantica on the Amazon Electronics workload, which has varying numbers of results per query instead of one best result.

4.3 End-to-End RAG Performance

We now compare the performance of document retrieval with ZeroRAG in an end-to-end RAG pipeline, using the PubMedQA workload. In this experiment, peers retrieve documents using ZeroRAG and then use an LLM (`Phi3.5-mini-instruct` [2]) to generate an answer. We consider the generated answer to be successful if it matches the (yes, no, or maybe) ground truth answer provided by PubMedQA. We use k = 5 for document retrieval.

Figure 5(a) shows the document retrieval recall and RAG accuracy for ZeroRAG and when using a centralized HNSW index, which resembles an upper-bound on performance. The solid and dashed lines represent recall and accuracy, respectively. Both the recall and accuracy of ZeroRAG increase as peers issue searches, showing that ZeroRAG successfully adapts its routing to improve RAG performance. The recall and accuracy of the central HNSW database remains constant across the experiment since it is a static data store. At the end of the experiment, the HNSW baseline reaches 97.3% recall with an accuracy of 79%, while ZeroRAG reaches 93.0% with 77% accuracy. The reason for the maximum recall of 97.3%, even with access to global data and the centralized index, is the imperfect embedding of queries and their (GPT5-nano) rephrasings: it could be that an embedding of a document relevant to a query is not in the top-5 closest embeddings. This shows that ZeroRAG gets performance relatively close to that of a centralized baseline.

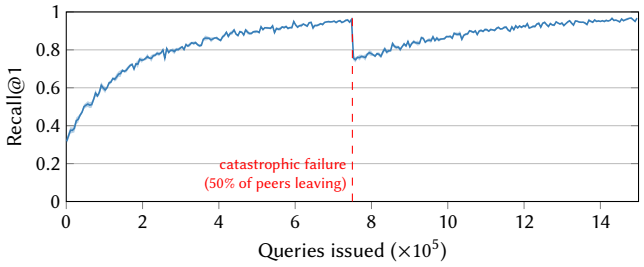

Figure 6: Average recall@1 per processed query batch in ZeroRAG with the PubMedQA workload, when there is a catastrophic failure event (dashed red line).

ZeroRAG can act as a progressive retrieval mechanism: a peer may contact only a few other peers for a fast, approximate result, or contact more peers to improve retrieval quality at the cost of additional messages. We next measure how many request messages (see Section 3) are needed until ZeroRAG finds the closest document that exists in the network as the overlay organizes. Figure 5(b) shows this information, grouped by percentile range and including the median. We observe that as the overlay self-organizes over time, the number of query messages required to find the best document decreases. In the final state, 99% of queries require no more than 13 query messages to find the best supporting document.

On average, the centralized HNSW baseline processes its longest-duration queries in 24 ms/query, whereas ZeroRAG requires 379 ms/query (excluding network latency). This difference is the expected cost of decentralization: a centralized index can exploit global knowledge and optimized single-machine data structures, whereas ZeroRAG achieves comparable recall through purely local interactions across an overlay of peers, without any central authority, global index, or centralized data stores.

Takeaway (Figure 5). End-to-end RAG performance with ZeroRAG approaches the recall and accuracy of the centralized HNSW baseline. For 99% of queries, ZeroRAG can find the best document in a network in just 13 requests.

4.4 Resilience to Catastrophic Peer Failure

We now analyze ZeroRAG's resilience under a catastrophic failure event, in which 50% of all peers go offline simultaneously at the midpoint of the experiment and are gradually replaced by new peers over the remainder of the run. For this experiment, we use the PubMedQA workload and we compute the average recall@1 for each batch of 1000 queries performed. Figure 6 shows these results and we annotate the catastrophic failure event with a vertical dashed line. Recall drops sharply at the moment of failure because

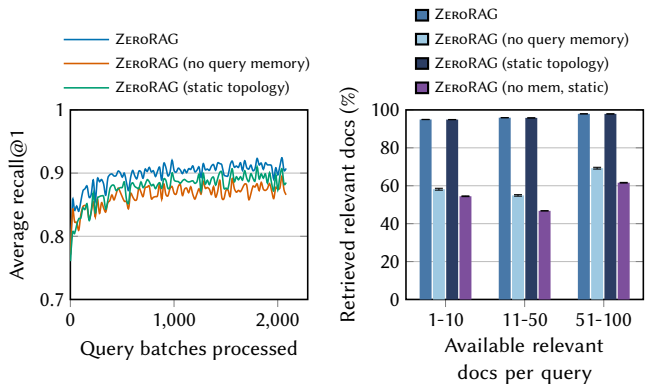

(a) AOL-IA (b) Amazon Electronics

Figure 7: Ablation of query memory and adaptive topology on the (a) AOL-IA and (b) Amazon Electronics workloads.

the learned paths and routed-query memories of the departed peers are lost. Nevertheless, ZeroRAG recovers to 95.5% recall in the final 10% of the run. This recovery happens entirely through local interactions: replacement peers accumulate routed-query memory and adaptive edges as they receive search queries, rebuilding useful routing state. After the catastrophic failure, peers require 15.1 requests per query to find the relevant document, which is higher compared to the same number before the catastrophic event (8.9): the neighborhoods of newly-joined peers are not yet optimal.

Takeaway (Figure 6). ZeroRAG is robust against catastrophic failure where routing dynamics get disrupted significantly, and self-heals entirely through local interactions as replacement peers rebuild routing state from incoming queries.

4.5 Ablation Study

We analyze the effect of key parameters and components on routing performance. Unless stated otherwise, experiments use the PubMedQA workload and results are reported over the final 10% of queries, by which point the overlay has converged. We report *average recall@1*, measuring the fraction of target documents successfully retrieved, and *average requests required*, measuring the number of QUERY messages sent until the target document is found. Together, these capture both routing effectiveness and efficiency.

4.5.1 The effect of query cache and topology adaptation. Figure 7 isolates the effect of ZeroRAG's two routing inputs by disabling routed-query memory, adaptive topology formation, or both, and comparing against the full system on both the AOL-IA and Amazon Electronics workloads. Routed-query memory is the dominant contributor on both workloads: on AOL-IA its removal lowers average recall@1 from 0.89 to 0.85, while on the more demanding multi-relevant Amazon Electronics coverage task it has a much larger effect. Disabling topology adaptation alone leaves retrieval close to the full system, and removing both inputs degrades it the most, confirming that the two mechanisms are complementary.

Table 1: The effect of query cache size.

Max. cache size (δ_m)	Avg. Recall@1	Avg. requests required
0	(77.860 ± 0.177)%	16.796 ± 0.067
10	(85.834 ± 0.120)%	12.757 ± 0.041
100	(97.612 ± 0.037)%	5.900 ± 0.017
1,000	(99.652 ± 0.007)%	3.448 ± 0.006
10,000	(99.652 ± 0.007)%	3.448 ± 0.006

Table 2: The effect of network size.

Network size (peers)	Avg. Recall@1	Avg. requests required
1,000	(99.365 ± 0.272)%	3.209 ± 0.111
5,000	(98.276 ± 0.066)%	4.330 ± 0.059
10,000	(98.378 ± 0.051)%	4.656 ± 0.015
50,000	(99.652 ± 0.007)%	3.448 ± 0.006
100,000	(99.913 ± 0.005)%	5.724 ± 0.011

Table 3: The effect of document count per peer.

Docs. per peer (t)	Avg. Recall@1	Avg. requests required
10	(98.665 ± 0.018)%	5.940 ± 0.023
20	(99.292 ± 0.016)%	5.545 ± 0.017
30	(99.667 ± 0.007)%	5.361 ± 0.012
40	(99.836 ± 0.006)%	5.276 ± 0.009
50	(99.919 ± 0.004)%	5.222 ± 0.007

4.5.2 The effect of query routing cache size (δ_m). We vary the size of the query routing cache (δ_m) and report the average recall@1 and number of requests required to obtain a document for the last 10% of the queries (of the PubMedQA workload). Table 1 shows that δ_m is a key contributor to the routing effectiveness of ZeroRAG. With no routed-query memory ($\delta_m = 0$), ZeroRAG obtains 77.9% recall and requires 16.8 requests to find the target document. Storing 10 routed queries ($\delta_m = 10$) improves recall to 85.8% and reduces the average required requests to 12.75. Increasing δ_m to 100 improves recall to 97.6% and the request count drops to 5.9. Even higher values lead to diminishing returns for recall and required requests.

4.5.3 Scalability. We evaluate the routing effectiveness and efficiency as the number of peers in the network increases. We vary the network size from 1000 to 100 000 peers and analyze the average recall@1 and number of requests required to obtain a document for the last 10% of the queries. Because of our setup, there could be a situation where a peer attempts to retrieve a document that is not advertised by any other peer in the network. We exclude these queries when computing recall@1, so recall is measured conditional on the target document being advertised by at least one remote peer. Table 2 shows that average recall@1 remains high (above 98%), even as the network size grows to 100 000 peers. This indicates that, once a relevant document exists in the network, ZeroRAG can locate it with high probability even as the peer population grows by two orders of magnitude. The average request count is non-monotonic as the network size increases. Furthermore, the average number of requests required to obtain the target document remains low.

Table 4: The effect of neighborhood size. Five seeds.

Max. neighbors (δ_n)	Avg. Recall@1	Avg. requests required
3	(98.614 ± 0.022)%	6.198 ± 0.020
5	(99.080 ± 0.012)%	5.360 ± 0.009
10	(99.417 ± 0.012)%	4.431 ± 0.010
25	(99.606 ± 0.004)%	3.702 ± 0.010
50	(99.644 ± 0.010)%	3.505 ± 0.005
75	(99.666 ± 0.008)%	3.388 ± 0.004
100	(99.675 ± 0.012)%	3.305 ± 0.006

Table 5: The effect of parallel searches.

Fanout	Rounds	Avg. Recall@1	Avg. rounds required
3	10	(99.374 ± 0.010)%	1.718 ± 0.002
6	5	(99.203 ± 0.012)%	1.508 ± 0.001
10	3	(97.456 ± 0.040)%	1.408 ± 0.001

4.5.4 Documents per peer. We evaluate the routing effectiveness and efficiency as we assign t more PubMedQA documents to individual peers. We show average recall@1 and requests required in Table 3. Even for $t = 10$, recall reaches 98.66% and increases to 99.91% for $t = 50$. As documents are more available, the number of requests to obtain them decreases from 5.94 for $t = 10$ to 5.22 for $t = 50$: a decrease of 11.5%. We conclude that ZeroRAG is effective even with small local document stores, while larger stores substantially improve coverage and reduce message overhead.

4.5.5 Neighborhood size (δ_n). We evaluate routing effectiveness and efficiency as we vary the number of neighbors (δ_n) each peer is connected to, see Table 4. Varying δ_n from 3 to 100 has a moderate effect on recall: 98.61% for $\delta_n = 3$ compared to 99.7% for $\delta_n = 100$. Increasing δ_n shows diminishing returns: a small neighborhood size is sufficient to obtain high recall as the network converges. In contrast, the average number of requests shows a sharper decrease: 6.2 for $\delta_n = 3$ to 3.30 for $\delta_n = 100$, a decrease of 46.5%. Larger values of δ_n primarily reduce the number of requests by giving each peer more immediate semantic routing options. This result suggests that moderate neighborhood sizes (*e.g.*, 50) are reasonable default choices, striking a good balance between recall and efficiency.

4.5.6 Parallel requests. We evaluate how parallelizing search affects recall and the number of rounds to find a target document. A peer sends f QUERY and QUERYRESULT in parallel for a fixed number of rounds. We fix the contact budget $C = 30$ and redistribute it across fanout and number of rounds: 3 requests per round for up to 10 rounds, 6 requests per round for 5 rounds, and 10 requests per round for 3 rounds. Table 5 reports the recall and average rounds over the final 10% of queries. The $f = 3$ configuration reaches 99.37% recall and requires 1.72 rounds to find the target document, while increasing f to 6 we get almost the same recall, 99.20%. With $f = 10$, recall drops very slightly to 97.46%. A query typically requires several referral steps to reach the right semantic neighborhood. Consequently, parallelization is most beneficial in low-traffic networks where the added message volume is acceptable, as it reduces wall-clock latency without sacrificing routing depth.

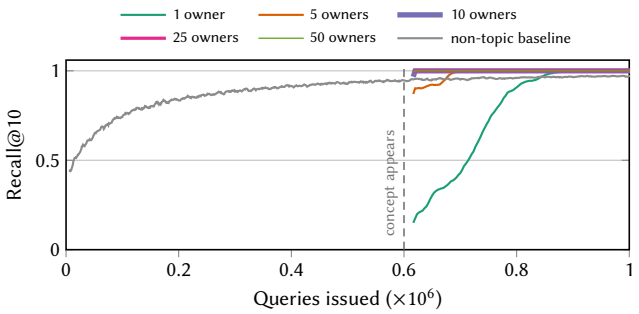

Figure 8: Introduction of a new topic at the 60% mark (vertical line). The lines of owners ≥ 10 overlap.

4.6 Spreading New Concepts

Our last experiment shows how fast ZeroRAG spreads a new topic in a network, analyzing how recall evolves for topics which appear anew, and are very sparsely available, but suddenly become popular with all peers. To do so, we omit a single topic from the PubMedQA workload and re-introduce it at the 60% mark into the experiment. We chose a random topic ("Peptostreptococcus") that has 10 relevant documents, and we vary how many peers own it initially. All peers then continuously send queries to find documents on our chosen (initially omitted) topic. A document starts appearing in other peers' local data stores only once a query successfully reaches it.

Figure 8 shows the rolling recall@10 for both the topic's queries per initial owner count and the average for other topics as a gray line. Seeding the topic at more owners lets recall recover faster once the concept reappears: with 10 or more owners the documents are immediately reachable across the network and the recall becomes 1.0 almost instantly, causing the lines to overlap. With less initial owners the documents also spread less fast, with the single-owner case being the slowest, climbing back only as the documents propagate through the overlay once queries begin to reach them.

Takeaway (Figure 8). ZeroRAG adapts to newly introduced, initially sparse topics through local interactions alone, with its speed governed by how widely the new content is initially held.

5 Related Work

We discuss works on RAG and nearest-neighbor search, federated and distributed RAG that scale retrieval without removing a coordinating authority, and decentralized peer-to-peer retrieval and routing, which is closest to our setting. Table 6 gives an overview.

Retrieval-augmented generation. RAG grounds LLM generation in external documents retrieved at query time, improving factual accuracy without retraining [14, 28, 62]. Retrieval is typically cast as nearest-neighbor search over high-dimensional embeddings: queries and documents are embedded into a shared space, and the documents closest to the query are returned [42]. Because exact search is expensive at scale, practical systems rely on ANN indexes such as HNSW graphs [38] or LSH [12, 29]. These techniques assume documents are collected and indexed under a single

administrative domain, so the index can be built globally. ZeroRAG targets the opposite: when documents remain with their owners and no global index can be constructed.

Federated and distributed RAG. Federated approaches route or merge results across independently managed collections: RAGRoute selects which remote sources to query for a given request [19], FeB4RAG benchmarks federated search for RAG [58], and pFedRAG personalizes retrieval models across federated clients [20]. Others target the privacy or efficiency of the retrieval: FedVSE provides privacy-preserving vector search over federated databases [15], PRAG offers end-to-end privacy-preserving RAG [30], and Proximity caches semantically similar queries to avoid repeated vector-database lookups [7]. A parallel line scales the retrieval substrate itself: DRAGON distributes RAG retrieval across local and cloud resources [31], DGRAG organizes graph-based retrieval over edge-cloud systems [63], and HARMONY and SPIRE build scalable distributed vector databases for high-throughput ANN search [60, 61]. These systems improve scalability, privacy, or cost, but assume a federated, cloud, or managed deployment in which a coordinating component still orchestrates retrieval. ZeroRAG instead routes each query through peer-local state alone.

Decentralized retrieval and routing. The closest systems to ZeroRAG use a decentralized overlay and a local or global state. One family imposes a shared logical structure that must be maintained as peers and documents change. Semantica organizes peers into a global LLM-guided semantic tree overlay [41], while DHT designs such as Kademlia and LSH-based distributed indexes place keys in a global address space for approximate lookup [12, 39]. Blockchain-based designs decentralize control through shared ledger state: D-RAG coordinates decentralized RAG through a blockchain [6], and a related system records a reliability score for each document source on a blockchain so that the trustworthiness of retrieved content cannot be tampered with [34]. These approaches remain decentralized but trade a central authority for global coordination.

A second family keeps routing state local but depends on explicit symbolic descriptions. REMINDIN' routes semantic queries using social metaphors and observations of previously successful peers [54]; DRAG/TARW performs topic-aware random walks given global knowledge of the available topics [59]; and SHIMI organizes decentralized agent memory through an explicit semantic hierarchy [22]. These systems share ZeroRAG's local-state philosophy, but rely on topic labels, expertise categories, or maintained hierarchies rather than routing directly from unlabeled embeddings.

A third, longer line of decentralized recommendation and nearest-neighbor work anticipates our overlay mechanics. WHATSUP disseminates news recommendations through an epidemic protocol that clusters users by interest without any central server [8], and a follow-up obfuscation-based protocol makes such profile-based collaborative filtering privacy-preserving [9]. Most directly, Cluster-and-Conquer constructs a k-nearest-neighbor graph by starting from a random overlay and greedily improving it through neighbor-of-neighbor exploration [16], the same locality-exploiting principle ZeroRAG uses to rewire its overlay. These build on classical decentralized overlay techniques such as gossip-based peer sampling [23], semantic-proximity content search [57], and the broader literature on robust peer-to-peer search [47]. ZeroRAG routes from unlabeled query and document embeddings and, crucially, additionally remembers recently routed queries, letting peers steer toward neighbors with semantically similar request histories.

Table 6: Whether various retrieval and routing systems have ("Decentr.") decentralized governance, use ("Unlab.") unlabeled data, or use a global, local, or no ("N/A") data structure.

System	Decentr.	Unlab.	Data Structure
ZeroRAG [this work]	✓	✓	**Local**
Semantica [41]	✓	✓	Global
LSH/Kademlia [12, 39]	✓	✓	Global
REMINDIN' [54]	✓		Local
SHIMI [22]	✓		Local
Proximity [7]		✓	N/A
DRAG (TARW) [59]	✓		Local
pFedRAG [20]		✓	N/A
DRAGON [31]		✓	N/A
D-RAG (Blockchain) [6]	✓		Global
HARMONY/SPIRE [60, 61]		✓	N/A

Overall, prior work occupies only part of the design space. Shared-structure methods are decentralized and often unlabeled, but require global maintenance; some local routing methods reduce coordination, but rely on labeled or structured semantics; and federated vector-search systems improve scalability, but are not fully decentralized. To our knowledge, ZeroRAG is unique in treating retrieval as a querier-based local routing problem: it avoids the global maintenance that Semantica and Kademlia require, while remaining fully decentralized, unlike centralized or federated alternatives.

6 Summary

We presented ZeroRAG, a decentralized routing layer for RAG that locates relevant documents across an overlay of peers without any central authority, labeled data, or global index. Each peer keeps only a bounded set of neighbor advertisements and a small cache of recently routed queries, routing each query through a best-first walk that continuously reconfigures the overlay toward semantically relevant neighbors. Our evaluation on the AOL-IA, Amazon Electronics, and PubMedQA workloads shows that ZeroRAG consistently outperforms Semantica and DHT-LSH, approaches the recall of a centralized HNSW index in an end-to-end RAG pipeline (93.0% vs. 97.3% at comparable accuracy). ZeroRAG recovers to 95.5% recall after losing half the network, scales to 100 000 peers, and absorbs new topics through local interactions alone. These results indicate that decentralization is a practical foundation for RAG, even when documents remain with the peers that own them.

Availability

The implementation of this work is available at https://github.com/pneague/ZeroRAG-System-Simulation.

Acknowledgments

This work has been funded by the Swiss National Science Foundation, under the project "FRIDAY: Frugal, Privacy-Aware and Practical Decentralized Learning", SNSF grant number 10001796, and by NWO/TKI grant BLOCK.2019.004

References

[1] Ashraf Abdelhamid, Mahmoud Said Elsayed, Anca D Jurcut, and Marianne A Azer. 2023. A lightweight anomaly detection system for black hole attack. *Electronics* 12, 6 (2023), 1294.

[2] Marah Abdin, Jyoti Aneja, Hany Awadalla, Ahmed Awadallah, Ammar Ahmad Awan, Nguyen Bach, Amit Bahree, Arash Bakhtiari, Jianmin Bao, Harkirat Behl, Alon Benhaim, Misha Bilenko, Johan Bjorck, Sébastien Bubeck, Martin Cai, Qin Cai, Vishrav Chaudhary, Dong Chen, Dongdong Chen, Weizhu Chen, Yen-Chun Chen, Yi-Ling Chen, Hao Cheng, Parul Chopra, Xiyang Dai, Matthew Dixon, Ronen Eldan, Victor Fragoso, Jianfeng Gao, Mei Gao, Min Gao, Amit Garg, Allie Del Giorno, Abhishek Goswami, Suriya Gunasekar, Emman Haider, Junheng Hao, Russell J. Hewett, Wenxiang Hu, Jamie Huynh, Dan Iter, Sam Ade Jacobs, Mojan Javaheripi, Xin Jin, Nikos Karampatziakis, Piero Kauffmann, Mahoud Khademi, Dongwoo Kim, Young Jin Kim, Lev Kurilenko, James R. Lee, Yin Tat Lee, Yuanzhi Li, Yunsheng Li, Chen Liang, Lars Liden, Xihui Lin, Zeqi Lin, Ce Liu, Liyuan Liu, Mengchen Liu, Weishung Liu, Xiaodong Liu, Chong Luo, Piyush Madan, Ali Mahmoudzadeh, David Majercak, Matt Mazzola, Caio César Teodoro Mendes, Arindam Mitra, Hardik Modi, Anh Nguyen, Brandon Norick, Barun Patra, Daniel Perez-Becker, Thomas Portet, Reid Pryzant, Heyang Qin, Marko Radmilac, Liliang Ren, Gustavo de Rosa, Corby Rosset, Sambudha Roy, Olatunji Ruwase, Olli Saarikivi, Amin Saied, Adil Salim, Michael Santacroce, Shital Shah, Ning Shang, Hiteshi Sharma, Yelong Shen, Swadheen Shukla, Xia Song, Masahiro Tanaka, Andrea Tupini, Praneetha Vaddamanu, Chunyu Wang, Guanhua Wang, Lijuan Wang, Shuohang Wang, Xin Wang, Yu Wang, Rachel Ward, Wen Wen, Philipp Witte, Haiping Wu, Xiaoxia Wu, Michael Wyatt, Bin Xiao, Can Xu, Jiahang Xu, Weijian Xu, Jilong Xue, Sonali Yadav, Fan Yang, Jianwei Yang, Yifan Yang, Ziyi Yang, Donghan Yu, Lu Yuan, Chenruidong Zhang, Cyril Zhang, Jianwen Zhang, Li Lyna Zhang, Yi Zhang, Yue Zhang, Yunan Zhang, and Xiren Zhou. 2024. Phi-3 Technical Report: A Highly Capable Language Model Locally on Your Phone. (2024). arXiv:2404.14219 [cs.CL]

[3] Qingyao Ai, Yongfeng Zhang, Keping Bi, Xu Chen, and W. Bruce Croft. 2017. Learning a Hierarchical Embedding Model for Personalized Product Search. In *Proceedings of the 40th International ACM SIGIR Conference on Research and Development in Information Retrieval* (Shinjuku, Tokyo, Japan) *(SIGIR '17)*. Association for Computing Machinery, New York, NY, USA, 645–654. doi:10.1145/3077136.3080813

[4] Aisha Alansari and Hamzah Luqman. 2025. Large language models hallucination: A comprehensive survey. arXiv:2510.06265 [cs.CL]

[5] Isha Amjad. 2026. AI Copyright Infringement: Navigating the Legal Risks of AI-Generated Content. *Trends in Intellectual Property Research* 4 (02 2026), 24–28. doi:10.69971/tipr.4.1.2026.104

[6] Tessa Andersen, Ayanna Avalos, Gaby Dagher, and Min Long. 2025. D-RAG: A Privacy-Preserving Framework for Decentralized RAG using Blockchain. (02 2025), 183–198. doi:10.5121/csit.2025.150417

[7] Shai Bergman, Anne-Marie Kermarrec, Diana Petrescu, Rafael Pires, Mathis Randl, Martijn de Vos, and Ji Zhang. 2025. Leveraging Approximate Caching for Faster Retrieval-Augmented Generation. In *Proceedings of the 26th International Middleware Conference* (Vanderbilt University, Nashville, TN, USA) *(Middleware '25)*. Association for Computing Machinery, New York, NY, USA, 340–353. doi:10.1145/3721462.3770776

[8] Antoine Boutet, Davide Frey, Rachid Guerraoui, Arnaud Jégou, and Anne-Marie Kermarrec. 2013. WHATSUP: A Decentralized Instant News Recommender. In *2013 IEEE 27th International Symposium on Parallel and Distributed Processing*. 741–752. doi:10.1109/IPDPS.2013.47

[9] Antoine Boutet, Davide Frey, Rachid Guerraoui, Arnaud Jégou, and Anne-Marie Kermarrec. 2015. Privacy-Preserving Distributed Collaborative Filtering. *Computing* 8593 (03 2015). doi:10.1007/s00607-015-0451-z

[10] Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. Language models are few-shot learners. In *Proceedings of the 34th International Conference on Neural Information Processing Systems* (Vancouver, BC, Canada) *(NeurIPS '20)*. Curran Associates Inc., Red Hook, NY, USA, Article 159, 25 pages.

[11] Erik Daniel and Florian Tschorsch. 2022. IPFS and friends: A qualitative comparison of next generation peer-to-peer data networks. *IEEE Communications Surveys & Tutorials* 24, 1 (2022), 31–52.

[12] Mayur Datar, Nicole Immorlica, Piotr Indyk, and Vahab S. Mirrokni. 2004. Locality-sensitive hashing scheme based on p-stable distributions. In *Proceedings of the Twentieth Annual Symposium on Computational Geometry* (Brooklyn, New York, USA) *(SCG '04)*. Association for Computing Machinery, New York, NY, USA, 253–262. doi:10.1145/997817.997857

[13] John R Douceur. 2002. The sybil attack. In *International workshop on peer-to-peer systems*. Springer, 251–260.

[14] Wenqi Fan, Yujuan Ding, Liangbo Ning, Shijie Wang, Hengyun Li, Dawei Yin, Tat-Seng Chua, and Qing Li. 2024. A survey on rag meeting llms: Towards retrieval-augmented large language models. In *Proceedings of the 30th ACM SIGKDD conference on knowledge discovery and data mining*. 6491–6501.

[15] Zeheng Fan, Yuxiang Zeng, Zhuanglin Zheng, and Yongxin Tong. 2025. FedVSE: A Privacy-Preserving and Efficient Vector Search Engine for Federated Databases. *Proc. VLDB Endow.* 18, 12 (Aug. 2025), 5371–5374. doi:10.14778/3750601.3750674

[16] George Giakkoupis, Anne-Marie Kermarrec, Olivier Ruas, and François Taïani. 2021. Cluster-and-Conquer: When Randomness Meets Graph Locality. In *2021 IEEE 37th International Conference on Data Engineering (ICDE)*. 2027–2032. doi:10.1109/ICDE51399.2021.00195

[17] Seth Gilbert and Nancy Lynch. 2002. Brewer's conjecture and the feasibility of consistent, available, partition-tolerant web services. *Acm Sigact News* 33, 2 (2002), 51–59.

[18] Rachid Guerraoui, Anne-Marie Kermarrec, Anastasiia Kucherenko, Rafael Pinot, and Martijn de Vos. 2024. Peerswap: A peer-sampler with randomness guarantees. In *2024 43rd International Symposium on Reliable Distributed Systems (SRDS)*. IEEE, 271–281.

[19] Rachid Guerraoui, Anne-Marie Kermarrec, Diana Petrescu, Rafael Pires, Mathis Randl, and Martijn de Vos. 2025. Efficient Federated Search for Retrieval-Augmented Generation. In *Proceedings of the 5th Workshop on Machine Learning and Systems* (World Trade Center, Rotterdam, Netherlands) *(EuroMLSys '25)*. Association for Computing Machinery, New York, NY, USA, 74–81. doi:10.1145/3721146.3721942

[20] Hangyu He, Xin Yuan, Kai Wu, Ren Ping Liu, and Wei Ni. 2025. pFedRAG: A Personalized Federated Retrieval-Augmented Generation System with Depth-Adaptive Tiered Embedding Tuning. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, Christos Christodoulopoulos, Tanmoy Chakraborty, Carolyn Rose, and Violet Peng (Eds.). Association for Computational Linguistics, Suzhou, China, 14255–14268. doi:10.18653/v1/2025.findings-emnlp.769

[21] Ethan Heilman, Alison Kendler, Aviv Zohar, and Sharon Goldberg. 2015. Eclipse attacks on {Bitcoin's}{peer-to-peer} network. In *24th USENIX security symposium (USENIX security 15)*. 129–144.

[22] Tooraj Helmi. 2025. Decentralizing AI Memory: SHIMI, a Semantic Hierarchical Memory Index for Scalable Agent Reasoning. arXiv:2504.06135 [cs.AI]

[23] Márk Jelasity, Spyros Voulgaris, Rachid Guerraoui, Anne-Marie Kermarrec, and Maarten Van Steen. 2007. Gossip-based peer sampling. *ACM Transactions on Computer Systems (TOCS)* 25, 3 (2007), 8–es.

[24] Ziwei Ji, Tiezheng Yu, Yan Xu, Nayeon Lee, Etsuko Ishii, and Pascale Fung. 2023. Towards mitigating LLM hallucination via self reflection. In *Findings of the Association for Computational Linguistics: EMNLP 2023*. 1827–1843.

[25] Qiao Jin, Bhuwan Dhingra, Zhengping Liu, William Cohen, and Xinghua Lu. 2019. Pubmedqa: A dataset for biomedical research question answering. In *Proceedings of the 2019 conference on empirical methods in natural language processing and the 9th international joint conference on natural language processing (EMNLP-IJCNLP)*. 2567–2577.

[26] Kevin Kuo, Chhavi Yadav, and Virginia Smith. 2025. Research in Collaborative Learning Does Not Serve Cross-Silo Federated Learning in Practice. arXiv:2510.12595 [cs.LG]

[27] Yoonjoo Lee, Kihoon Son, Tae Soo Kim, Jisu Kim, John Joon Young Chung, Eytan Adar, and Juho Kim. 2024. One vs. many: Comprehending accurate information from multiple erroneous and inconsistent ai generations. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*. 2518–2531.

[28] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Proceedings of the 34th International Conference on Neural Information Processing Systems* (Vancouver, BC, Canada) *(NeurIPS '20)*. Curran Associates Inc., Red Hook, NY, USA, Article 793, 16 pages.

[29] Wen Li, Ying Zhang, Yifang Sun, Wei Wang, Mingjie Li, Wenjie Zhang, and Xuemin Lin. 2020. Approximate Nearest Neighbor Search on High Dimensional Data — Experiments, Analyses, and Improvement . *IEEE Transactions on Knowledge & Data Engineering* 32, 08 (Aug. 2020), 1475–1488. doi:10.1109/TKDE.2019.2909204

[30] Zhijun Li, Minghui Xu, Huayi Qi, Wenxuan Yu, Tingchuang Zhang, Qiao Zhang, GuangYong Shang, Zhen Ma, and Xiuzhen Cheng. 2026. PRAG: End-to-End Privacy-Preserving Retrieval-Augmented Generation. arXiv:2604.26525 [cs.CR]

[31] Shangyu Liu, Zhenzhe Zheng, Xiaoyao Huang, Fan Wu, Guihai Chen, and Jie Wu. 2025. Efficient Distributed Retrieval-Augmented Generation for Enhancing Language Model Performance. arXiv:2504.11197 [cs.LG]

[32] Alexander Loser, Kai Schubert, and Frederik Zimmer. 2004. Taxonomy-based routing overlays in p2p networks. In *Proceedings. International Database Engineering and Applications Symposium, 2004. IDEAS'04.* IEEE, 407–412.

[33] Alexander Löser, Steffen Staab, and Christoph Tempich. 2005. Semantic Methods for P2P Query Routing. In *Multiagent System Technologies*, Torsten Eymann, Franziska Klügl, Winfried Lamersdorf, Matthias Klusch, and Michael N. Huhns (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 15–26.

[34] Yining Lu, Wenyi Tang, Max Johnson, Taeho Jung, and Meng Jiang. 2025. A Decentralized Retrieval Augmented Generation System with Source Reliabilities Secured on Blockchain. arXiv:2511.07577 [cs.CR]

[35] Le Ma, Ran Zhang, Yikun Han, Shirui Yu, Zaitian Wang, Zhiyuan Ning, Jinghan Zhang, Ping Xu, Pengjiang Li, Ziyue Qiao, Wei Ju, Chong Chen, Dongjie Wang, Kunpeng Liu, Pengyang Wang, Pengfei Wang, Yanjie Fu, Chunjiang Liu, Yuanchun Zhou, and Chang-Tien Lu. 2026. A Comprehensive Survey on Vector Database: Storage and Retrieval Technique, Challenge. (2026). arXiv:2310.11703 [cs.DB]

[36] Sean MacAvaney, Craig Macdonald, and Iadh Ounis. 2022. Reproducing Personalised Session Search Over the AOL Query Log. In *Advances in Information Retrieval: 44th European Conference on IR Research, ECIR 2022, Stavanger, Norway, April 10–14, 2022, Proceedings, Part I* (Stavanger, Norway). Springer-Verlag, Berlin, Heidelberg, 627–640. doi:10.1007/978-3-030-99736-6_42

[37] Pratyush Maini, Hengrui Jia, Nicolas Papernot, and Adam Dziedzic. 2024. LLM Dataset Inference: Did you train on my dataset? (2024). arXiv:2406.06443 [cs.LG]

[38] Yu A. Malkov and D. A. Yashunin. 2020. Efficient and Robust Approximate Nearest Neighbor Search Using Hierarchical Navigable Small World Graphs. *IEEE Trans. Pattern Anal. Mach. Intell.* 42, 4 (April 2020), 824–836. doi:10.1109/TPAMI.2018.2889473

[39] Petar Maymounkov and David Mazières. 2002. Kademlia: A Peer-to-Peer Information System Based on the XOR Metric. In *Revised Papers from the First International Workshop on Peer-to-Peer Systems (IPTPS '01)*. Springer-Verlag, Berlin, Heidelberg, 53–65.

[40] Satoshi Nakamoto. 2008. Bitcoin: A peer-to-peer electronic cash system. (2008).

[41] Petru Neague, Quinten Stokkink, Naman Goel, and Johan Pouwelse. 2025. Semantica: Decentralized Search using a LLM-Guided Semantic Tree Overlay. arXiv:2502.10151 [cs.IR]

[42] Sameer A Nene and Shree K Nayar. 1997. A simple algorithm for nearest neighbor search in high dimensions. *IEEE Transactions on pattern analysis and machine intelligence* 19, 9 (1997), 989–1003.

[43] Greg Pass, Abdur Chowdhury, and Cayley Torgeson. 2006. A picture of search. In *Proceedings of the 1st International Conference on Scalable Information Systems* (Hong Kong) *(InfoScale '06)*. Association for Computing Machinery, New York, NY, USA, 1–es. doi:10.1145/1146847.1146848

[44] J.A. Pouwelse, Pawel Garbacki, D. Epema, and Henk Sips. 2005. The Bittorrent P2P File-Sharing System: Measurements and Analysis. *Peer-to-Peer Systems IV* 3640, 205–216. doi:10.1007/11558989_19

[45] Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. 3973–3983. doi:10.18653/v1/D19-1410

[46] Nicola Rieke, Jonny Hancox, Wenqi Li, Fausto Milletari, Holger R Roth, Shadi Albarqouni, Spyridon Bakas, Mathieu N Galtier, Bennett A Landman, Klaus Maier-Hein, et al. 2020. The future of digital health with federated learning. *NPJ digital medicine* 3, 1 (2020), 119.

[47] John Risson and Tim Moors. 2007. *Survey of research towards robust peer-to-peer networks: Search methods.* Technical Report.

[48] Eleonora Rosati. 2025. Infringing AI: Liability for AI-Generated Outputs under International, EU, and UK Copyright Law. *European Journal of Risk Regulation* 16, 2 (2025), 603–627. doi:10.1017/err.2024.72

[49] Sentence Transformers. [n. d.]. all-mpnet-base-v2. https://huggingface.co/sentence-transformers/all-mpnet-base-v2.

[50] Dou Shen, Jian-Tao Sun, Qiang Yang, and Zheng Chen. 2006. Building bridges for web query classification. In *Proceedings of the 29th annual international ACM SIGIR conference on Research and development in information retrieval*. 131–138.

[51] K. Sripanidkulchai, B. Maggs, and H. Zhang. 2003. Efficient content location using interest-based locality in peer-to-peer systems. In *IEEE INFOCOM 2003. Twenty-second Annual Joint Conference of the IEEE Computer and Communications Societies (IEEE Cat. No.03CH37428)*, Vol. 3. 2166–2176 vol.3. doi:10.1109/INFCOM.2003.1209237

[52] Quinten Stokkink, Can Umut Ileri, Dick Epema, and Johan Pouwelse. 2023. Web3 Sybil avoidance using network latency. *Computer Networks* 227 (2023), 109701.

[53] Chunqiang Tang, Zhichen Xu, and Sandhya Dwarkadas. 2003. Peer-to-peer information retrieval using self-organizing semantic overlay networks. In *Proceedings of the 2003 Conference on Applications, Technologies, Architectures, and Protocols for Computer Communications* (Karlsruhe, Germany) *(SIGCOMM '03)*. Association for Computing Machinery, New York, NY, USA, 175–186. doi:10.1145/863955.863976

[54] Christoph Tempich, Steffen Staab, and Adrian Wranik. 2004. Remindin': semantic query routing in peer-to-peer networks based on social metaphors. In *Proceedings of the 13th International Conference on World Wide Web* (New York, NY, USA) *(WWW '04)*. Association for Computing Machinery, New York, NY, USA, 640–649. doi:10.1145/988672.988759

[55] U.S. Department of Health & Human Services. 1996. Health Insurance Portability and Accountability Act of 1996. https://www.hhs.gov/hipaa Public Law 104-191, 110 Stat. 1936.

[56] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. *Advances in neural information processing systems* 30 (2017).

[57] Spyros Voulgaris, A-M Kermarrec, and Laurent Massoulié. 2004. Exploiting semantic proximity in peer-to-peer content searching. In *Proceedings. 10th IEEE International Workshop on Future Trends of Distributed Computing Systems, 2004. FTDCS 2004.* IEEE, 238–243.

[58] Shuai Wang, Ekaterina Khramtsova, Shengyao Zhuang, and Guido Zuccon. 2024. FeB4RAG: Evaluating Federated Search in the Context of Retrieval Augmented Generation. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval* (Washington DC, USA) *(SIGIR '24)*. Association for Computing Machinery, New York, NY, USA, 763–773. doi:10.1145/3626772.3657853

[59] Chenhao Xu, Longxiang Gao, Yuan Miao, and Xi Zheng. 2025. Distributed Retrieval-Augmented Generation. arXiv:2505.00443 [cs.DC]

[60] Qian Xu, Feng Zhang, Chengxi Li, Lei Cao, Zheng Chen, Jidong Zhai, and Xiaoyong Du. 2025. HARMONY: A Scalable Distributed Vector Database for High-Throughput Approximate Nearest Neighbor Search. *Proc. ACM Manag. Data* 3, 4, Article 249 (Sept. 2025), 28 pages. doi:10.1145/3749167

[61] Yuming Xu, Qianxi Zhang, Qi Chen, Baotong Lu, Menghao Li, Philip Adams, Mingqin Li, Zengzhong Li, Jing Liu, Cheng Li, and Fan Yang. 2025. Scalable Distributed Vector Search via Accuracy Preserving Index Construction. arXiv:2512.17264 [cs.DC]

[62] Penghao Zhao, Hailin Zhang, Qinhan Yu, Zhengren Wang, Yunteng Geng, Fangcheng Fu, Ling Yang, Wentao Zhang, Jie Jiang, and Bin Cui. 2026. Retrieval-augmented generation for ai-generated content: A survey. *Data Science and Engineering* (2026), 1–29.

[63] Wenqing Zhou, Yuxuan Yan, and Qianqian Yang. 2026. DGRAG: Distributed Graph-based Retrieval-Augmented Generation in Edge-Cloud Systems. arXiv:2505.19847 [cs.AI]

[64] Yuchen Zhuang, Yue Yu, Kuan Wang, Haotian Sun, and Chao Zhang. 2023. ToolQA: a dataset for LLM question answering with external tools. In *Proceedings of the 37th International Conference on Neural Information Processing Systems* (New Orleans, LA, USA) *(NeurIPS '23)*. Curran Associates Inc., Red Hook, NY, USA, Article 2180, 27 pages.