arXiv:2611.12088v1 [cs.LG] 9 Nov 2026

Zero-Search Reinforcement Learning for Slay the Spire 2 with Failure-Driven Restarts

Xiangqin Chen
University of Waterloo
xiangqin.chen@uwaterloo.ca

Zhengxuan Xie
University of Waterloo
z94xie@uwaterloo.ca

November 2026

Abstract

Slay the Spire 2 (STS2) is a deck-building roguelike in which a winning run at the highest difficulty takes several hundred decisions, and most losses happen in late boss fights that a learner seldom reaches. We train agents that choose every action with one forward pass of a 6.6M-parameter entity transformer, without lookahead. Training and evaluation run the game's own compiled assembly (v0.107.1) in a headless .NET host. The host replays a run exactly from its seed and actions, so training episodes can restart from states the learner has visited, and a failure-driven curriculum allocates these restarts in proportion to where complete runs are lost. On 12,800 held-out games at Ascension 10 (A10), the released policies win 8.05% averaged over the five characters (Ironclad 15.6%, Regent 12.5%), and AlphaSpire's released networks used without search win 0.55% on a common panel. Longer A10 training raised Ironclad from 12.4% to 15.6% and Defect from 3.0% to 8.1% on these games. At Ascension 0 the released policies win 80.2%, level with the 80.3% that the tree-search agent Spire Mentor reports, while Mentor leads by 5 to 8 points from A4 upward. Run in our host on the same 2,000 A10 games, Mentor's search wins 8.8% against 8.2% for our policies, which need 21–27 ms per decision on one CPU core. GPT-6 Astra at maximum reasoning won none of 11 A10 games through our text interface. In a controlled study of 18 training runs from six random initializations, failure-driven allocation exceeds uniform restart allocation on each of three Regent initializations, by 3.6 percentage points on average, while the Silent effect is small and changes sign. Code, trained policies, an in-game mod and the records behind every table are at https://github.com/HalfVulpes/sts2ai.

1 Introduction

A run of Slay the Spire 2 crosses three acts of procedurally generated maps. Each floor brings a combat, event, shop, treasure or rest site, and the player builds a deck, collects relics and potions, and chooses a route along the way. Ascension levels add cumulative modifiers, and the highest level in the current version, Ascension 10 (A10), ends the run with two consecutive boss fights. A deck assembled poorly in act one can make the final fights unwinnable, so combat skill alone does not secure a win. The decisions that settle a run are also rare in the training data. In an early development measurement at Ascension 0, the act-3 boss accounted for 7% of an Ironclad policy's decisions and 59% of its losses.

Current STS2 agents reach their strongest results with large computation at decision time. Spire Mentor combines a learned run-level policy with thousands of tree-search simulations in every combat decision [8]. We ask how far a learned policy goes when it chooses each action with a single

network evaluation, and which training changes help at the late bottleneck. A policy of this kind runs inside the game on a CPU, and with the simulator included it plays a winning A10 game in about 30 seconds on two CPU cores.

Our approach rests on the game's own code. A headless host loads the installed game assembly and exposes every in-game choice as a decision, and it reproduces a run exactly from its seed and the sequence of chosen actions. This determinism lets a worker save the action prefix that led to an intermediate state and later restart a training episode from it. We use this to give the learner practice where it fails. A failure map tracks where complete runs from the ordinary start are lost, and restart episodes are allocated across twelve situations (act start and the three kinds of combat in each act) in proportion to that map.

This paper makes four contributions.

1. A training and evaluation platform that runs the STS2 v0.107.1 assembly headlessly, with a structured observation encoder shared by the trainer and an in-game mod, exact replay of runs, and an evaluation protocol on the game's canonical seed strings (Section 3).
2. Zero-search policies for all five characters, evaluated on 12,800 held-out A10 games and at four lower difficulties, with the code, the policies and a playable mod released (Sections 4 and 6).
3. A comparison of zero-search policies with Monte Carlo tree search and language models on the same game version, including Spire Mentor run in our host on identical seeds and GPT-6 Astra through a headless text interface (Section 7).
4. A controlled 18-run study of failure-driven restart allocation from independent random initializations, which finds a consistent gain for one character and none for the other (Section 8).

2 Related work

Restart curricula. Starting episodes from intermediate states is an established way to help reinforcement learning with long horizons. Reverse Curriculum Generation grows a set of start states outward from the goal [6], and Backplay starts episodes at late states of a demonstration and moves the start point backward during training [10]. Go-Explore returns to promising archived states before exploring further [5]. Jump-Start RL lets a guide policy produce the first part of each episode for an exploration policy [15]. Prioritized Level Replay revisits training levels in proportion to an estimate of their learning potential [7], and Per-Context Reverse Curricula adjusts the amount of assistance separately for each context from success feedback [4]. Our restart states come from the learner's own past trajectories, including failed ones, so a later state need not be an easier start. The novel element is the outer allocation rule that spreads restarts across situations by where complete runs fail. Our controlled study isolates this rule. It does not compare against faithful reimplementations of the methods above.

Agents for Slay the Spire. For STS2, AlphaSpire releases a policy-and-search system with networks for run-level and combat decisions that guide a tree search [3]. Spire Mentor trains a run-level policy with PPO and plays combat with Monte Carlo tree search at several thousand simulations per decision; it reports win rates for all five characters at five ascension levels [8]. The sts2-rl-research project trains separate combat and choice agents without search and reports 26.6% for Ironclad at Ascension 2 [12]. Epoch AI's SpireClimb evaluates frontier language models on STS2 at A10 from screenshots with access to the game wiki [14]. Its published data do not link per-run logs, so its individual runs cannot be checked. In 8 Defect games that we ran with the same model in a similar setting, Astra won none, but we recorded neither their seeds nor their screenshots

(Section 7). AgenticSTS studies bounded memory for language-model agents on STS2, mainly with Silent at Ascension 0 [2]. SpireBench evaluates language models on the original Slay the Spire at Ascension 20 [13].

Evaluation. We follow recommendations to report uncertainty over independent training runs separately from uncertainty over evaluation episodes, and to fix evaluation protocols before testing [1, 9].

3 Environment and evaluation protocol

3.1 Headless host

STS2 is a Godot/.NET application. The host loads the installed game assembly into a .NET 9 process with a minimal Godot compatibility layer and the game's test mode. It advances the asynchronous game logic on a deterministic synchronization context and requests decisions through the entry points the user interface calls. Patches remove input/output side effects, restore real-game behavior where test mode takes shortcuts, and make accesses to interface nodes null-safe; a static rewriter covers about 500 methods that touch interface nodes. Any exception raised by the game fails the episode, and evaluations report such failures. The host advances prompts with exactly one option and gold rewards automatically, and plays the Crystal Sphere minigame with a fixed local heuristic. The policy, the text interface and every comparison in this paper share these conventions. The research version is the Linux build of v0.107.1 with all content unlocked; the repository records the hashes of all eleven game assemblies, and an importer verifies them on a user's own installation.

A run is reproduced from its seed and the sequence of chosen action indices. Restart training relies on this. Before a restarted episode continues, the worker compares a fingerprint of the encoded observation at the reconstructed state with the stored fingerprint and treats a mismatch as an error. In the final calibration run, 3,765 restart episodes completed without a mismatch. Because the fingerprint covers only the encoded observation, hidden game variables are checked only indirectly, through later observations.

3.2 Observation and actions

Each decision is encoded as up to 192 typed entity tokens: a global run-state token, the player, enemies with their intents, powers, cards in hand and the play pile, multisets of the draw, discard, exhaust and deck piles, orbs, potions, relics, the act bosses, offers such as shop items, rewards, event options and map nodes, and a summary of the run history. A token carries five categorical identifiers and 32 numeric features. Draw piles are encoded as multisets without order, and the observation contains no random-number state. The encoding includes the internal identifier of each enemy's next move as well as its rendered intent, so the observation contains slightly more than a human sees on screen. Each legal action is a token with its own features and pointers to up to two entity tokens, such as a card and its target. Typically 5 to 40 actions are legal, at most 128.

3.3 Seeds, panels and statistics

The game turns a ten-character seed string from a base-34 alphabet without I and O into its random streams. Our evaluations use these canonical strings and the game's own hashing of the act-selection

stream, so a seed means the same in our host as in the native game's seed entry. Earlier development evaluations used a different numeric mapping; we keep them separate and label them as legacy where they appear. Every evaluation panel is drawn uniformly at random with a recorded generator seed before the evaluated policies are trained or chosen. Training samples seeds uniformly from the same range, excluding every evaluation panel; a panel used to choose a checkpoint is never reused to report that checkpoint's result.

We report win rates with Wilson 95% intervals, which describe sampling variation of a fixed policy over games. Comparisons of two policies on the same seeds use the exact McNemar test on the discordant games, with a Bonferroni or Holm correction stated for each family of comparisons. For training methods, the unit of replication is the independent training run, and game-level intervals do not measure variation between training runs.

4 Agent

4.1 Network

Token embeddings sum the five identifier embeddings with a projection of the numeric features, and eight learned memory tokens are appended. Six pre-norm transformer blocks (width 256, eight heads, feed-forward width 1024) encode the set. The state summary concatenates the output of the global token with attention pooling over all tokens. The policy scores each legal action with a three-layer MLP on the action embedding, a projection of the entity tokens it points to, and the state summary, followed by a softmax over legal actions. A value trunk feeds four heads: discounted return, probability of winning the run, floors remaining, and end-of-combat HP fraction; the last three are auxiliary targets. With the vocabulary of 2,650 entries the network has 6,568,453 parameters. At evaluation the agent takes the highest-scoring legal action of the exponential moving average (EMA) of the weights. The same encoder and network run in C# inside the game as a mod with autoplay, single-stepping and an overlay of the policy.

4.2 Training

We train with clipped PPO [11], using separate entropy coefficients for combat and non-combat decisions, $\gamma = 0.9998$, GAE $\lambda = 0.99$, 65,536 samples per iteration, two epochs and minibatches of 4,096. The reward gives +0.01 per floor, +0.02 per elite, +0.1 per boss, +1 for a win and −0.1 for a death, plus a potential-based term on the HP fraction. We make no claim that this shaped objective has the same optimum as the win rate. Episodes end as losses after 30 identical non-combat actions, 150 combat turns or 300 actions in a turn, which removes reward loops. A native collector multiplexes several hundred simulator processes through shared memory and batches their inference on the GPU.

The policies start from an exploratory development history on two GPU servers. Training began with all characters sharing one network, continued with single-character specialists, an ascension curriculum and the restart curriculum of Section 5, and finished with a learning rate of 10^{-4} and training directly at A10 for Ironclad and Regent. We call the final checkpoints of that history the base policies. They are EMA weights after 0.93B (Defect), 1.14B (Silent), 1.17B (Necrobinder), 1.66B (Ironclad) and 1.68B (Regent) collected decisions, each matches a historical EMA tensor exactly, and they were chosen on legacy development seeds. The released Regent and Necrobinder policies are their base checkpoints. The released Ironclad, Defect and Silent policies continued from their base weights for 1,800 updates (118M collected decisions) at A10 with the failure-driven

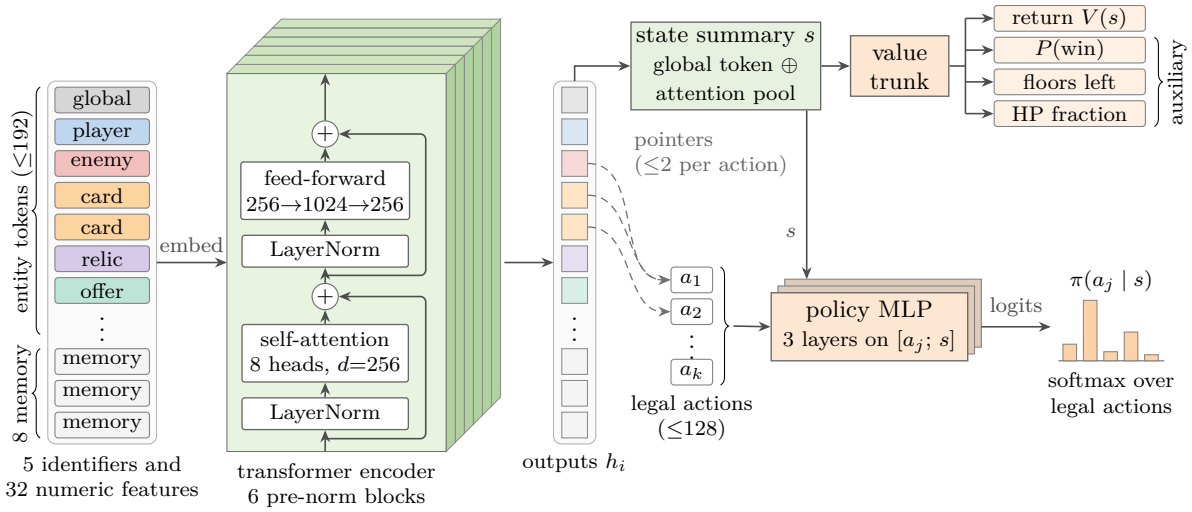

Figure 1: Policy and value network (6,568,453 parameters). Typed entity tokens and eight learned memory tokens pass through six pre-norm transformer blocks; the front block shows their layers. The value trunk reads a state summary of the encoder outputs. The same MLP scores every legal action from its own features, the encoder outputs of the entities it points to (dashed) and the state summary.

curriculum and a learning rate of 10^{-4}, under a protocol frozen before training that named each run's final EMA as its only candidate (Section 9). Every canonical panel in this paper was drawn after the base checkpoints were chosen and is excluded from the training seeds of these runs.

5 Failure-driven restart curriculum

Let ρ_0 be the distribution of ordinary run starts. The objective is the win rate from these starts,

$$J_0(\theta) = \mathbb{E}_{s_0 \sim \rho_0,\, \pi_\theta}\left[\mathbf{1}\{\text{run won}\}\right]. \tag{1}$$

Restart episodes begin at intermediate states and therefore change the distribution of training starts. The PPO probability ratio corrects for the change of policy and leaves this change uncorrected, so a restart update is a heuristic step for J_0. We always evaluate on complete runs from ρ_0.

Situations. Restart states fall into twelve situations: the start of each act, and the first decision of a monster, elite or boss combat in each act. Each worker keeps up to 16 recent states per situation in a first-in-first-out buffer. A state is stored as the game seed, character, ascension and the action prefix that reaches it.

Failure map. Every ten PPO iterations, the losses of fresh single-player runs since the last update form an attribution vector $D_t \in \mathbb{R}^{12}_{\geq 0}$. A loss in combat credits 0.65 to that combat's situation and 0.35 to the start of its act, which stands for the deck-building choices made since then. A loss

outside combat credits the start of its act. The map and the situation distribution are

$$L_t = 0.5\, L_{t-1} + D_t, \tag{2}$$

$$p_t(s) = \frac{L_t(s)/\sum_j L_t(j) + 0.02}{1 + 12 \cdot 0.02}. \tag{3}$$

The distribution starts uniform and keeps its previous value when an update has no new losses. It measures where losses occur, and it is not normalized by how often a situation is visited.

Episode resets. At each reset a worker starts a restart episode with probability 0.5. It samples a situation from p_t restricted to the situations for which it holds a stored state,

$$p_t(s \mid A_w) = \frac{\mathbf{1}\{s \in A_w\}\, p_t(s)}{\sum_{j \in A_w} p_t(j)}, \tag{4}$$

which keeps the actual restart rate equal to the requested rate when some buffers are empty. Within the situation, a state with n_i restarts and ℓ_i losses is sampled with weight

$$w_i = 0.1 + 4q_i(1 - q_i), \qquad q_i = \frac{\ell_i + 1}{n_i + 2}, \tag{5}$$

which favors states whose outcome is uncertain. The worker replays the prefix without querying the learner and then plays on with the current policy. Replayed decisions cost simulator time and do not enter the PPO update. Figure 2 summarizes the procedure.

Failure-driven restart curriculum (one learner, W workers)

Worker w, at every episode reset:

1. With probability 0.5 and a nonempty set A_w of situations with stored states, draw $s \sim p_t(\cdot \mid A_w)$, draw a stored state i of s with weight w_i, replay its action prefix from its seed, and check the observation fingerprint. Otherwise start a fresh run.
2. Play with the current policy. On entering any of the 12 situations in a fresh run, store (seed, action prefix, fingerprint) in that situation's 16-slot buffer.
3. At the end of a restart episode, update n_i and ℓ_i of the state it started from.

Learner, every PPO iteration: update on the collected decisions (replayed prefixes excluded). Every ten iterations, attribute the losses of fresh runs to situations (D_t), update L_t and p_t, and send p_t to the workers.

Figure 2: The restart procedure. The uniform control replaces p_t with the uniform distribution over the same 12 situations and keeps everything else.

Losses are a heuristic for value. Consider two necessary stages that succeed independently with $p_1 = 0.6$ and $p_2 = 0.5$, so $J_0 = 0.3$. Losses fall on the first stage with mass 0.4 and on the second with mass $0.6 \cdot 0.5 = 0.3$, so the failure map favors stage one. If training raises either stage's success probability by the same amount, the gain in J_0 is $\partial J_0/\partial p_1 = p_2 = 0.5$ for stage one and $\partial J_0/\partial p_2 = p_1 = 0.6$ for stage two, and stage two is the better use of practice. Loss counts can therefore point practice to the wrong stage, and the benefit of the rule has to be measured.

6 Released policies

Table 1 pools three disjoint held-out A10 panels: the 512-game panel of the external comparison (Section 7) and the two 1,024-game test panels of the selection and continuation studies (Section 9),

where the base policies serve as the baseline. All 12,800 games of the released policies completed with no engine error. Ironclad wins about one run in six, Regent one in eight and Defect one in twelve; Silent and Necrobinder win about 2%. On these games the continued Ironclad and Defect policies win clearly more often than their base policies (269 against 187 seeds won by only one of the two, exact McNemar $p < 10^{-3}$, and 186 against 56, $p < 10^{-16}$), while Silent's rise from 47 to 61 wins is not significant (54 against 40, $p = .18$). The five-character mean rises from 6.29% to 8.05%.

Table 1: Released zero-search policies at A10 on 2,560 held-out canonical games per character (three disjoint panels of 512, 1,024 and 1,024 games). Base: the policy before the release continuation, on the same games; Regent and Necrobinder are released unchanged. Bold marks the higher count, arrows the change in win rate; paired exact McNemar p. Intervals are Wilson 95%.

Character	Base	Released	Win rate (%)	95% interval (%)	Change (pp)	p
Ironclad	317	**399**	15.59	14.23–17.04	↑3.20	$< 10^{-3}$
Silent	47	**61**	2.38	1.86–3.05	↑0.55	.18
Defect	78	**208**	8.12	7.13–9.25	↑5.08	$< 10^{-16}$
Necrobinder	42 (unchanged)		1.64	1.22–2.21		
Regent	321 (unchanged)		12.54	11.31–13.88		
Mean of five	805	**1031**	8.05		↑1.77	

Table 2 reports the same policies at Ascension 0, 4, 8 and 9 on 512 common seeds per level, which match the levels that Spire Mentor reports. All 10,240 games completed without engine error. The five-character mean falls from 80.2% at A0 to 71.5% at A4, 51.8% at A8 and 22.8% at A9, then to 8.1% at A10. Ironclad and Regent keep about two fifths of their A9 rate at A10 (39.3% to 15.6% and 30.1% to 12.5%), where the run continues into a second final boss. Silent and Necrobinder are already below 13% at A9. Against their base policies on the same seeds, the continued Defect policy wins more often at A4, A8 and A9 (all $p < .01$), the continued Ironclad policy wins more often at A8 (370 against 329 of 512, $p = .002$) and less often at A0 (460 against 479, $p = .027$), and Silent changes significantly at no level.

Table 2: Released zero-search policies across difficulty, win rate in percent. A0–A9: 512 common canonical seeds per level; A10: the 2,560 games per character of Table 1. Bold marks the best character at each level; arrows mark a significant change of a continued policy against its base policy on the same seeds (exact McNemar, $p < .05$).

Character	A0	A4	A8	A9	A10
Ironclad	**89.8**↓	**88.3**	**72.3**↑	**39.3**	**15.6**↑
Silent	67.6	54.3	36.1	12.3	2.4
Defect	83.8	77.0↑	60.0↑	26.6↑	8.1↑
Necrobinder	72.7	59.4	34.2	5.9	1.6
Regent	87.1	78.5	56.4	30.1	12.5
Mean of five	80.2	71.5	51.8	22.8	8.1

A10 adds a second final boss directly after the first. In a development sweep of an earlier Ironclad policy on 100 legacy-protocol seeds, the second final boss ended 25 of the 100 A10 runs, the largest single cause, and the act-1 boss ended 19. This failure profile motivated training directly at A10 (Section 9).

7 Comparison across agent classes

Table 3 places the released policies next to the agents we ran ourselves and the zero-search and tree-search STS2 A10 results reported by others, grouped by the computation they spend on a decision. The agents also differ in what they observe.

Table 3: STS2 v0.107.1, Ascension 10, all five characters. "Measured" rows are our runs; "reported" rows are taken from the cited sources, with their own protocols and seed panels.

Agent	Computation per decision	Source	Games	Mean win (%)
Ours (released)	one network evaluation, 21–27 ms on one CPU core	measured	12,800	**8.05**
AlphaSpire networks without search [3]	one network evaluation	measured	2,560	0.55
Spire Mentor [8]	8,000 MCTS simulations per combat decision, about 112 ms on one core	reported	2,500	14.0
Spire Mentor, in our host	4,000–16,000 (Ironclad) or 8,000 simulations per combat decision	measured	2,000	8.8
GPT-6 Astra, our text interface	maximum reasoning, median 13.7 s	measured	11	0.0
DeepSeek V4.1 Flash 1017, our text interface	maximum reasoning, median 29 s	measured	12	0.0

Zero-search policies. We ran AlphaSpire v0.2.1 with its released checkpoints in its native greedy mode for both combat and run-level decisions, and required its counters of searched decisions and search steps to stay at zero in every game. On a common panel of 512 seeds per character, our policies win 208 games and AlphaSpire's networks win 14, all for Ironclad (Table 4). The two systems run in their own engines: AlphaSpire uses its own simulator in its custom mode, and our host uses Standard mode with all content unlocked. AlphaSpire's checkpoints are trained as priors for search, and the non-Ironclad networks often end turns early when used alone. In 61 Ironclad games, including 4 wins, AlphaSpire's action head fell back to a uniform prior 227 times when the number of actions exceeded its axis. All these games are counted. The table therefore compares two released artifacts used without search. It does not measure AlphaSpire's search system. The sts2-rl-research agents also act without search, but they report only Ironclad at Ascension 2 (26.6%) [12], so they do not enter the A10 table. In our host, the released Ironclad policy wins 88.3% at the harder Ascension 4.

Table 4: Released policies on a common panel of 512 canonical A10 seeds per character, both used without search, wins of 512; bold marks the higher count. Wilson 95% intervals in percent. All 5,120 games completed without engine error.

Character	Ours	Interval	AlphaSpire	Interval
Ironclad	**75**	11.85–17.97	14	1.64–4.54
Silent	**9**	0.93–3.31	0	0.00–0.74
Defect	**47**	6.97–11.99	0	0.00–0.74
Necrobinder	**9**	0.93–3.31	0	0.00–0.74
Regent	**68**	10.61–16.50	0	0.00–0.74

Tree search. Spire Mentor's rows use 8,000 simulations per combat decision with 500 runs per character at A10, measured in October 2026 [8]: Ironclad 27.2%, Defect 22.0%, Silent 9.0%, Regent 6.2% and Necrobinder 5.6%. The authors note that these rows were measured on an engine version that granted rewards after the first of the two final bosses, a bug in A10's favor that they later fixed. Their newer Ironclad policy wins 30.0% of 400 runs on the corrected engine at 4,000 simulations per decision (16,000 in boss fights). On different engines and seed panels, our zero-search Regent policy (12.5%) wins more often than Mentor's searched Regent (6.2%), and Mentor leads on the other four characters, most clearly on Defect. Figure 3 extends the comparison to the lower ascensions. At A0 the two agents are level on average (80.2% against 80.3%), and Mentor leads by 4.8 points at A4, 7.8 at A8, 6.2 at A9 and 5.9 at A10. Our Regent policy leads Mentor's at every level, and our Necrobinder policy leads at A0.

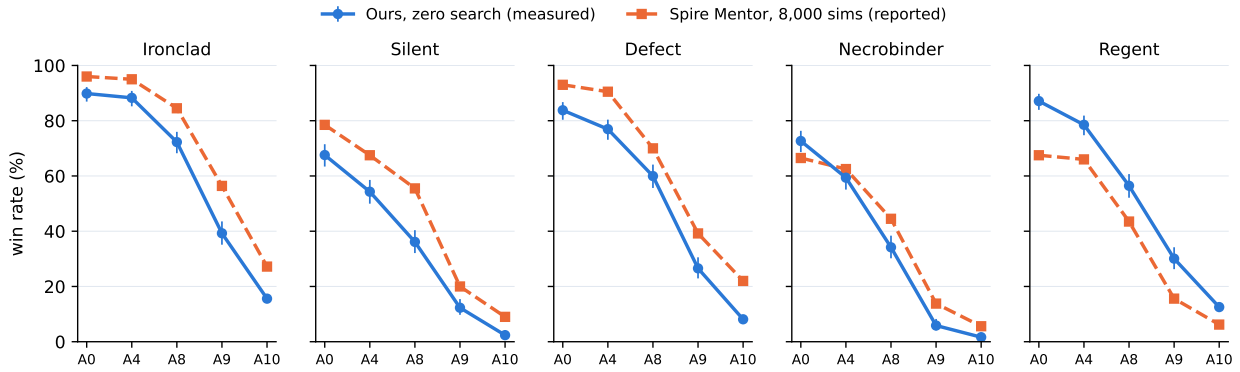

Figure 3: Win rate by ascension. Our released zero-search policies (measured; 512 games per point at A0–A9 and 2,560 at A10, Wilson 95% intervals) and Spire Mentor at 8,000 simulations per combat decision (reported; 200 runs per point at A0–A8 and 500 at A9–A10, no intervals given). Engines and seed panels differ, and Mentor notes that its A10 values come from an engine version that favored A10.

Spire Mentor in our host. To remove the engine and panel differences, we ran Spire Mentor inside our host on 400 common canonical A10 seeds per character, the first 400 of a panel on which our policies were also evaluated. Mentor's own released run-level networks make every non-combat decision through its real-game adapter (map, card rewards, reward screens, shops, rest sites, events, treasure and deck edits), and its search plays every combat move at its published budgets: 4,000 simulations per decision for Ironclad (16,000 in boss fights) and 8,000 for the other characters. The search sees only what a player sees. Four details differ from Mentor's own real-game driver: our policy makes the card choices that some cards open during combat (8 to 39 per run, depending on the character), the search tree is rebuilt from the board at every decision, Mentor's additional Ironclad search settings are not applied, and one rare event is skipped. All runs finished without error, and no non-combat decision fell back to a random choice or to our policy. Table 5 gives the results. On identical games Mentor's search averages 8.8% over the five characters, against the 14.0% it reports in its own engine, and our zero-search policies average 8.2%. Search leads clearly only on Silent (6.8% against 2.3%, $p = .004$). It also leads on Ironclad (16.8% against 13.3%, $p = .13$) and Defect (12.5% against 9.0%, $p = .098$), but not significantly; Necrobinder is level (3.5% against 2.5%, $p = .52$); and our Regent policy wins three times as often as Mentor's (13.8% against 4.5%). Mentor spends three to five minutes of search per run; our policy spends 21–27 ms per decision.

Table 5: Spire Mentor (search) and our released policies (zero search) in the same host on the same 400 canonical A10 seeds per character. Wins of 400; bold marks the higher count; Wilson 95% intervals in percent; paired exact McNemar test on the seeds won by only one agent.

Character	Mentor	Interval	Ours	Interval	Only Mentor / only ours (p)
Ironclad	**67**	13.4–20.7	53	10.3–16.9	44 / 30 (.13)
Regent	18	2.9–7.0	**55**	10.7–17.5	16 / 53 (<.001)
Silent	**27**	4.7–9.6	9	1.2–4.2	27 / 9 (.004)
Defect	**50**	9.6–16.1	36	6.6–12.2	38 / 24 (.098)
Necrobinder	**14**	2.1–5.8	10	1.4–4.5	13 / 9 (.52)
Mean of five	**8.8**		8.2		

Swapping the two layers. Both agents split a run into combat moves and everything else (map, card rewards, shops, rest sites, events). On the same 400 seeds we crossed the two layers for the base Ironclad and Regent policies, pairing our policy's non-combat decisions with Mentor's search in combat, and the reverse (Table 6). Each agent's non-combat play suits its own combat play. With our policy in combat, Mentor's non-combat decisions collapse to 22 and 2 wins, while our policy's own win 41 and 55 (pooled, 90 against 18 seeds won by only one arm, $p < 10^{-10}$). Mentor's search playing the decks our policy builds is the strongest combination for Regent, with 70 wins against 55 for our policy alone and 18 for Mentor alone. Over both characters, search on our decks gains 83 seeds and loses 57 against our policy's own combat play ($p = .034$), about the three points that search at the final bosses alone adds (Section 9). Our policies' gap to search therefore lies in combat; their run-level decisions hold up against Mentor's.

Table 6: Crossing non-combat decisions and combat play on the same 400 A10 seeds, wins of 400, for the base Ironclad and Regent policies. The diagonal cells are the two agents unchanged; bold marks the best combination per character.

Character	Non-combat decisions	Mentor's search in combat	Our policy in combat
Ironclad	Mentor	**67**	22
	ours	52	41
Regent	Mentor	18	2
	ours	**70**	55

Language models. We ran GPT-6 Astra at its maximum reasoning setting on up to 25 A10 games, five per character, drawn at random from the 512-seed panel of Table 4 before any game started. The model receives a structured text rendering of the player-visible state, including card, relic and power descriptions, enemy intents and the unordered draw pile, together with numbered legal actions. It has no tools, wiki or rollouts. Each request carries the current state, a strategic note of at most 1,200 characters that the model maintains, and its four most recent action notes. States with a single legal action advance automatically. Because of the cost, an amendment frozen before any game had ended replaced the fixed 25 games with a Wald sequential test of $p = 0.347$ against $p = 0.10$ with error rates of 5% in both directions. Games are judged in their fixed launch order, using only the longest prefix of finished games, since losses end earlier than wins. The test stops for the lower rate after 0 wins in 10 games, 1 in 15 or 2 in 19, and for the higher rate after 3 wins with at most 2 losses or 4 with at most 6. Astra lost all 11 games it played, including the first 10 in launch order, which ended the cohort (one-sided 95% upper bound on its win rate 24%);

a cost cap of ten games, set while the last games were running, would have ended it at the same point. Astra played well into the run, reaching the final bosses in 5 of 11 games and a mean floor of 34.5, against 32.8 for our released policy on the same seeds, which won one of them. A game took a median of 395 model decisions and 2.2 hours of model time; the cohort used 58.8M input and 2.55M output tokens. This result concerns our interface, with bounded context, text instead of screenshots, and no wiki.

We also ran a smaller check from screenshots. On 8 November 2026, OpenAI's Codex agent with GPT-6 Astra at its maximum reasoning setting opened the live game, the same version as above, on a separate virtual machine and played Defect at A10 through computer use, with access to the game wiki and unrestricted network access. It lost all 8 games, after which we stopped for cost (about 5,000 Codex credits per game). These games had no frozen protocol. We did not record their seeds, and Codex trimmed the screenshot turns from its session history, so the games cannot be reconstructed. We therefore report them as an unverified observation and leave them out of Table 3 and of every statistical statement.

DeepSeek V4.1 Flash 1017, served on our own GPUs, played the same 25 seeds through the same interface and runner. We stopped this cohort after 12 finished games, all lost, because a win had become unlikely. Six games stopped mid-run and six unstarted games are left out of the count, and one game ended when the model reached its 40,000-token output limit three times at the same decision without answering. Its runs ended early: the mean floor was 12.4, against 26.6 for our released policy on the same 12 seeds (one-sided 95% upper bound on its win rate 22%). We release every decision of both text-interface cohorts' games with the frozen protocols: the chosen actions, the model's stated reasons and notes, latencies and token counts. Replaying the recorded actions in our host reproduces the logged outcome of all 23 finished games. The release leaves out the rendered observations because they contain the game's own text; replaying a game with a copy of the game regenerates them exactly.

Computation per decision. We timed the released network in the C# runtime that the mod uses, on 500 recorded observations, pinned to one CPU core. A decision takes 21–27 ms on one core of a Ryzen Threadripper 9980X and 62–67 ms on one core of a Xeon CPU Max 9480; with all threads available, it takes 10 ms on the Threadripper. For each released policy the C# outputs match PyTorch on all 500 observations, with the same greedy action every time and logits within 7.3×10^{-5}. A winning A10 run has a median of about 750 decisions, so it needs about 20 seconds of single-core inference. Spire Mentor reports 14 ms per 1,000 simulations on one thread of a different CPU, about 112 ms per combat decision at 8,000 simulations and 224 ms at the 16,000 its newest Ironclad uses in boss fights. Its Ironclad A10 runs took 162 s of CPU each in its own benchmark on a shared host. Its search also relies on a complete reimplementation of the game kept in step with the live game, while our policy needs only the observation. GPT-6 Astra in our text interface took a median of 13.7 s per decision, with about 10,400 input and 320 output tokens, and the cohort cost about 70 USD per hour (mix of credits and subscription) at eight concurrent games. DeepSeek V4.1 Flash on our own servers took a median of 29 s. One network evaluation per decision therefore costs several times less than tree search on a CPU core and about three orders of magnitude less time than a language model. At that cost it reaches 8.1% at A10, far above AlphaSpire's networks used without search and the language models we ran, and below tree search.

8 Controlled study of restart allocation

8.1 Design

We compare three training methods from scratch: fresh starts only, restarts with uniform situation allocation, and restarts with failure-driven allocation. Both restart methods use the same within-situation sampling, buffers and restart probability, so the comparison of the two isolates the outer allocation rule. The study covers Regent and Silent at Ascension 0 with three independent random initializations per character, giving 18 runs. Within each character–initialization block, the three methods start from identical weights with a new optimizer. Table 7 lists the frozen design.

Table 7: Design of the independent-start study, frozen before training.

Factor	Setting
Methods	Fresh only; uniform situations + uncertainty; failure map + uncertainty
Task	Regent and Silent, A0, single player, no ascension curriculum
Budget	763 updates of 65,536 rows; 50,003,968 rows per run
PPO	learning rate 2×10^{-4}; $\gamma = 0.9998$; $\lambda = 0.99$; two epochs; minibatch 4,096
Workers	256 simulator processes
Development curves	EMA at updates 153, 306, 459, 611 and 763 on 256 development games
Final test	EMA at update 763 on 1,024 held-out games

The budget fixes the number of optimizer rows. Restart runs also replay action prefixes, which costs simulator time outside this budget (Table 9), so the study compares methods at equal optimization samples, with total simulator work recorded separately. All arms restart their worker processes and prefix buffers at the same five segment boundaries, while model, optimizer, EMA, sampler statistics and learner random state carry over. During validation we found that the first worker implementation could give consecutive combats the same identifier when no non-combat decision separated them, which mislabeled the auxiliary HP target and restart points for the two A10 final bosses. We corrected the worker, verified five complete A10 regression runs with identical actions and observations, discarded the affected batch of 59.1M rows, and ran all 18 runs anew from the original six initializations.

8.2 Results

All 18 final policies played the same 1,024 held-out games, 18,432 in total, with no engine error. Table 8 gives every training block.

For Regent, failure-driven allocation wins 10.55% on average, against 6.97% for both the fresh and the uniform arm. It exceeds uniform allocation in all three blocks, by +3.58 points on average (range +2.44 to +4.39). For Silent, the three arms average 5.57%, 5.21% and 5.60%; the failure-minus-uniform difference averages +0.39 points with blocks of both signs. Variation between initializations is large for Silent, where fresh training ranges from 28 to 114 wins. Three blocks per character show a consistent direction for Regent and no effect for Silent. They do not establish the size of the Regent effect for other initializations, other characters or higher difficulty. These short-budget policies win far less than the released ones, which collected at least twenty times more decisions.

Table 8: Final A0 wins out of 1,024 common held-out games. Each row is one random initialization shared by the three methods; bold marks the best method of the row. The last column is the paired difference in win rate, in percentage points.

Character	Initialization	Fresh	Uniform	Failure	Failure − uniform
Regent	41001	74	54	**99**	↑4.39
Regent	41002	60	54	**94**	↑3.91
Regent	41003	80	106	**131**	↑2.44
Silent	52001	**114**	72	86	↑1.37
Silent	52002	28	**53**	40	↓1.27
Silent	52003	29	35	**46**	↑1.07

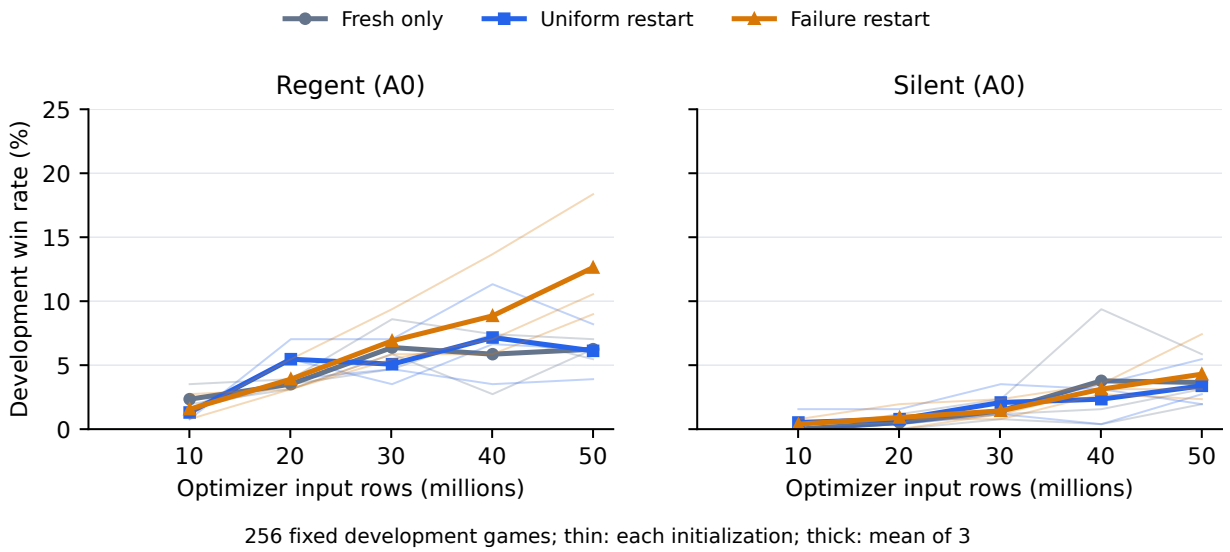

Figure 4: Development curves for all 18 runs. Thin lines are single initializations and thick lines the means of three; each point uses 256 development games. Development games are separate from the final test and were not used to choose checkpoints.

Figure 6 shows where restart episodes started. Uniform allocation gives each situation a similar share; later situations receive somewhat less because workers hold stored states for them less often. Failure-driven allocation moves practice toward act starts and the act-1 and act-2 bosses and away from act-3 monster and elite fights, in a similar pattern for both characters. One visible difference is the act-3 boss, which failure allocation practices more than uniform allocation for Regent (8.5% against 7.2%) and less for Silent (5.9% against 6.4%). This describes three runs per arm and does not establish a cause of the different outcomes.

Restart arms replay about half as many prefix decisions as they collect. Requested and actual restart counts are equal in every run because of the availability masking. Failure-driven allocation replays 10% fewer prefix decisions than uniform allocation while winning more for Regent, so its Regent gain does not come from extra simulator work.

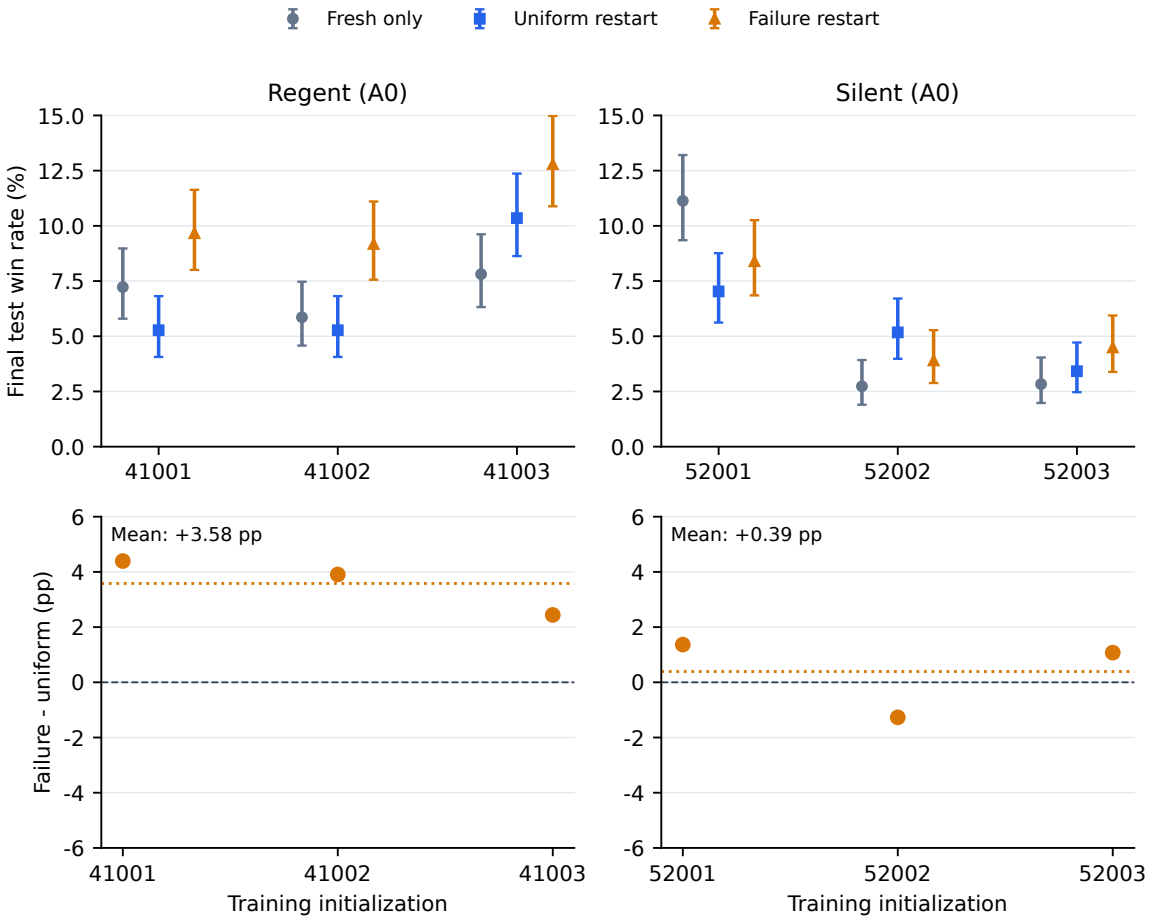

Figure 5: Top: final policies with Wilson 95% intervals over test games. Bottom: the three independent block differences per character and their means.

9 Improving the A10 policies

We made several attempts to raise the base A10 policies, each with a frozen protocol and fresh test seeds. Only longer A10 training produced confirmed gains, and its last run gave the released Ironclad, Defect and Silent policies.

Checkpoint selection. A fixed pool of 21 historical checkpoints, including the five base ones, played 256 development games per character at A10. The best candidate per character then played 1,024 new test games against the base policy. The five-character mean moved from 6.41% to 6.66%, with a 95% bootstrap interval over common seeds of -0.53 to $+1.04$ points, and no paired comparison approached significance (Table 10).

Short continuation. Each base policy trained for 5.05M more rows at A10 with the failure-driven restart curriculum and a learning rate of 3×10^{-5}. A candidate replaced the base policy only if it won strictly more of 256 development games. On 1,024 test games per character the mean moved

Table 9: Recorded training work, summed over the six runs of each method. Decisions in millions. Trainer hours are summed over segments on mixed hardware and do not form a speed comparison.

Quantity	Fresh	Uniform	Failure
Optimizer rows (M)	300.024	300.024	300.024
Collected decisions (M)	300.638	300.618	300.621
Replayed prefix decisions (M)	0	149.819	134.631
Restart episodes	0	591,062	577,858
Completed episodes	895,131	1,187,028	1,161,006
Summed trainer hours	15.17	17.41	17.88

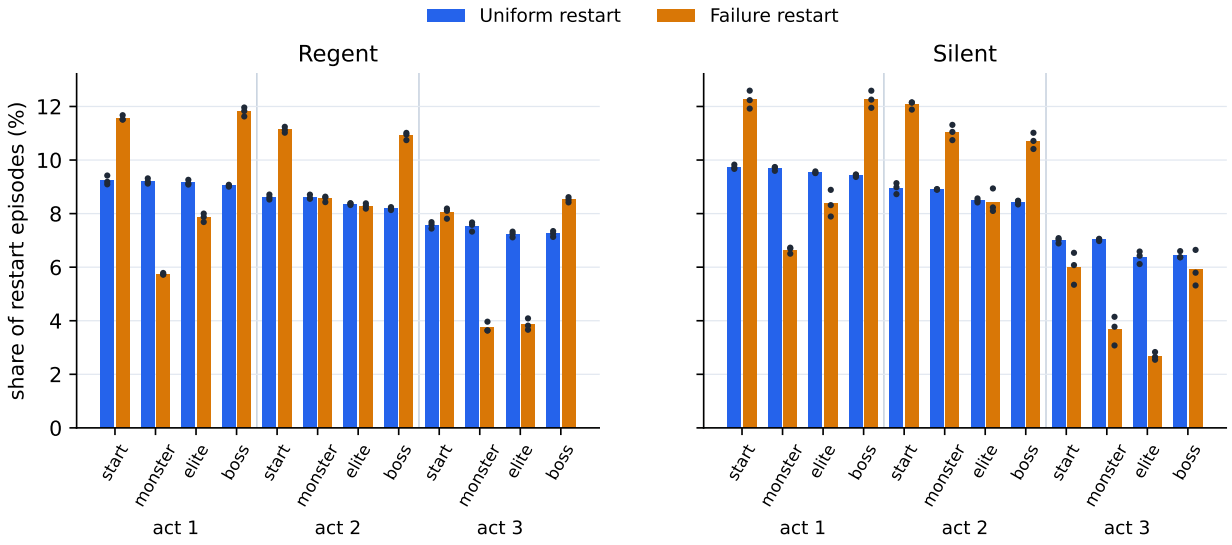

Figure 6: Share of restart episodes by starting situation over the whole run. Bars are means over the three initializations; dots are single runs.

from 6.09% to 6.45% (bootstrap interval −0.31 to +1.02 points; Table 11). Regent's gain of 2.93 points had an unadjusted paired $p = 0.018$ and $p = 0.088$ after correction for five comparisons. A separately frozen confirmation on 2,048 new seeds gave 251 wins for the base Regent policy and 252 for the candidate (153 games gained, 152 lost, exact $p = 1.0$), so the gain did not replicate. Ironclad's development gain from 26 to 41 wins reversed on the test panel.

Longer continuation. Ironclad and Regent continued from their base weights for 600 updates (39.3M rows) at A10 with the recipe that ended their development history. On a 1,024-seed panel the Ironclad candidate won 138 games against the base policy's 99 (paired $p = .003$), a comparison outside the frozen readout of that experiment. A confirmation frozen before any game, on 2,048 new seeds, gave 268 against 247 for Ironclad ($p = .28$) and 305 against 266 for Regent ($p = .047$, Holm-adjusted .093). Neither passes; the first panel happened to be one on which the base Ironclad policy did poorly (9.7% against 12.1% on the confirmation seeds). We do not pool the panels.

Second-boss practice. A curriculum aimed at the final double boss restarted training episodes from second-boss situations stratified by entry HP, added a per-fight outcome reward and repeated

Table 10: Checkpoint selection at A10, wins out of 1,024 test games; bold marks the higher count. Paired exact p-values; all equal 1 after Bonferroni correction over five characters.

Character	Base	Selected	Difference (pp)	p
Ironclad	122	**124**	↑0.20	.939
Silent	20	**22**	↑0.20	.860
Defect	36	**39**	↑0.29	.795
Necrobinder	15	15	0.00	1.000
Regent	135	**141**	↑0.59	.705

Table 11: Short A10 continuation. Development: base → candidate, wins of 256. Test: base → selected, wins of 1,024. Silent and Defect kept the base policy. Bold marks the higher count.

Character	Development	Test	Difference (pp)
Silent	**4** → 1	19 → 19	0.00
Defect	**9** → 4	29 → 29	0.00
Necrobinder	3 → **5**	**18** → 13	↓0.49
Ironclad	26 → **41**	**128** → 121	↓0.68
Regent	32 → **35**	118 → **148**	↑2.93

each practiced fight four times. Against the plain continuation above, with the same updates, seeds and 1,024 test games, Ironclad was unchanged (141 against 138, $p = .88$) and Regent fell (102 against 138; 65 against 101 discordant, Holm-adjusted $p = .013$).

Search at the final bosses, and imitating it. We replayed our policy's own runs to the first final boss and let Spire Mentor's search (8,000 simulations) play both final bosses from those entries: 251 Ironclad and 170 Regent entries from 400 seeds each. The search won 66 and 72 runs where our policy won 51 and 61 (pooled, 52 runs won only by search against 26 only by our policy, $p = .004$), and it won the second boss more often from identical boss-2 states (19 against 7 and 20 against 5). Searching only the final bosses would thus raise these A10 rates by about three points (3.8 for Ironclad, 2.8 for Regent). We then distilled the search's choices into the network, training on our policy's own boss states with a target of 0.8 on the search's pick plus 0.2 on its visit shares and a KL anchor to the base policy on full-run states. Agreement with the search's pick on held-out states rose from 45.7% to 50.0% (Ironclad) and from 40.1% to 45.6% (Regent), while the greedy move changed in 31% and 41% of decisions. From the same boss-1 entries the distilled policies won 60 and 62 runs against 74 and 74 for the base ones (pooled 29 against 55, $p = .006$). An earlier variant trained on visit shares alone also lost (63 against 74, Ironclad). The search's visit distributions are nearly flat, and the network could not predict its picks from its own observation even where the search was confident, so imitation added noise to otherwise coherent play.

Long A10 training for the weakest characters. The A10 recipe that raised Ironclad and Regent during development combined single-character training at A10, the failure-driven curriculum and a learning rate of 10^{-4}, over several hundred million decisions. Silent, Defect and Necrobinder never received it. We trained each of them from its base weights for 201.3M optimizer rows (3,072 updates) at A10 with this recipe. The frozen protocol names the final EMA as the only candidate and compares it with the base policy on 1,024 new test seeds per character, with an exact McNemar test and Holm correction over the three characters; it also checks A0 and A9 on the 512 sweep

seeds. Defect rose from 31 to 80 wins (3.0% to 7.8%; 71 seeds won only by the new policy against 22, Holm-adjusted $p < 10^{-5}$) and Silent from 16 to 31 (1.6% to 3.0%; 27 against 12, $p = .047$); Necrobinder's rise from 13 to 20 is not significant ($p = .28$) (Table 12). All three reach 6 to 8 floors deeper on average. At A9 the new policies win more often for all three characters, clearly so for Defect (125 against 81 of 512, $p < .001$). Across difficulty, the trained Defect policy also wins more often at A4 (393 against 358 of 512, $p = .009$) and A8 (282 against 235, $p < .001$), while the trained Silent and Necrobinder policies are level with the base ones at A4 and A8 (all $p > .3$). At A0 Silent and Defect hold their rates, while Necrobinder falls from 372 to 345 of 512 ($p = .044$, uncorrected), so for that character long A10 training traded some low-difficulty play for A10 and A9 play. Defect's gain came from a recipe that the base Ironclad and Regent policies had already received.

A second round of the same recipe for Defect, frozen before training and started from the first round's final EMA (another 3,072 updates), did not raise it further. On a second set of 1,024 new A10 test seeds it won 88 games against 82 for the first-round policy (8.6% against 8.0%; 55 seeds won only by the second round against 49, $p = .62$). At A0 and A9 on the 512 sweep seeds it won 422 against 400 ($p = .059$) and 135 against 125 ($p = .45$). The first-round policy's 8.0% on these fresh seeds replicates its 7.8% on the first test panel.

Release continuation. The last run of this recipe produced the released Ironclad, Defect and Silent policies. Each continued from its base weights for 1,800 updates (118M rows) at A10 with the failure-driven curriculum and a learning rate of 10^{-4}, as the control arm of a frozen 12-run study whose other arms we will report separately; the protocol named the final EMA as the only candidate. On that study's 2,048 new A10 test seeds per character, the continued policies won 313, 176 and 54 games against 234, 53 and 33 for the base policies (seeds won only by the continued / only by the base policy: Ironclad 216 / 137, $p < 10^{-4}$; Defect 158 / 35, $p < 10^{-19}$; Silent 44 / 23, $p = .014$). These comparisons lie outside the study's frozen readout, which compares training methods with this control, and an independent replicate of the same arm reached 330 Ironclad and 156 Defect wins. We adopted the continued policies as the release after these results; Table 1 then measured them, under a plan frozen before any game, on 2,560 held-out games per character that took no part in that choice. There Ironclad won 399 games against 317 for its base policy (269 seeds won only by the continued policy against 187, $p < 10^{-3}$), Defect 208 against 78 (186 against 56, $p < 10^{-16}$) and Silent 61 against 47 (54 against 40, $p = .18$): the gains confirm for Ironclad and Defect, and Silent's remains uncertain. This run is three times as long as the longer continuation above (1,800 against 600 updates), which did not confirm for Ironclad.

Table 12: Long A10 training (201.3M rows) for the three weakest characters. Wins of 1,024 new A10 test seeds; "discordant" counts the seeds won only by the trained / only by the base policy (exact McNemar, Holm over three characters). Mean floor reached, and wins of 512 on the A0 and A9 sweep panels, base → trained; bold marks the better value. The released Silent and Defect policies come from the release continuation instead.

Character	Base	Trained	Discordant	Holm p	Floor	A0	A9
Silent	16	**31**	27 / 12	.047	24.7→**32.3**	335→**345**	53→**66**
Defect	31	**80**	71 / 22	$< 10^{-5}$	32.4→**38.8**	**407**→400	81→**125**
Necrobinder	13	**20**	19 / 12	.28	24.1→**30.0**	**372**→345	30→**46**

10 Discussion and limitations

A single network evaluation per decision carries the released policies at A10 to about one win in six for Ironclad, one in eight for Regent and one in twelve for Defect, and to one in forty or fewer for Silent and Necrobinder. In our host tree search leads clearly for Silent, leads without a significant difference for Ironclad and Defect, is level for Necrobinder and is behind for Regent. Search can correct a combat line at the moment of play, and the final double boss rewards exact combat sequencing over many turns, so this split is consistent with where our policies lose. The comparison has limits. Mentor's reported numbers come from its own engine, seed panels and observation contract; our in-host run removes the engine and panel differences but not every detail of its real-game driver. Our observation includes the enemies' internal next-move identifier. The host patches parts of the game for headless execution, and behavioral equivalence with the live game is tested on selected cases only.

Why do A10 rates move so little? Three measurements point the same way. First, the win is a product of stages: in our replay of the base Ironclad policy on 400 seeds, 63% of runs reach the first final boss, 64% of those win it and 32% of those win the second, so a five-point gain at either of the first two stages moves the run rate by about one point (two points at the second boss), and a one-point change takes thousands of paired games to detect. Second, the decisive errors come early: for 90 of 125 final-boss losses (72%), none of 12 replays that re-decided all of act 3 won, and for 72 of them (58%) no replay of the boss fights with the same dice won either, so the deck built in acts 1 and 2 could not take both bosses. Third, the policy barely conditions on difficulty. On 139,744 decisions with more than one legal action from its own A10 runs, setting the ascension input from A10 to A0 changes the base Ironclad policy's greedy action in 5.9% of them (Regent 3.4%; shop decisions change most, 15.0% for Ironclad and 10.5% for Regent), while its win-probability head moves from 0.16 to 0.80 (Regent 0.24 to 0.70). The win-probability head separates the levels, and the policy hardly does. The development history trained mostly below A10, through an ascension curriculum, with A10-only training at the end for Ironclad and Regent and none for the other three, so a strategy that wins nine runs in ten at A0 may be reused almost unchanged where it wins one in eight.

The restart study isolates one design choice and shows a character-dependent effect with three initializations per character. It equalizes optimizer rows and records the extra prefix replay separately. It does not compare against faithful implementations of related curricula, and it was run at A0 with short budgets. The base policies come from an exploratory development history whose individual steps were not run as controlled experiments; only the final continuation of three of them ran under a frozen protocol.

11 Conclusion

Running the game's own code headlessly makes exact replay available, and exact replay makes restart curricula practical in a long-horizon card game. Zero-search policies trained this way reach 8.1% at STS2 A10 over five characters, 15.6% for Ironclad and 12.5% for Regent, the highest rates we know of for agents that use one network evaluation per decision. Failure-driven restart allocation helped Regent consistently in a controlled study and did not help Silent. Closing the gap to search at A10 will likely require stronger combat play in the final fights, for example by distilling search into the policy. Code, protocols, evaluation records and the five policies are released (Appendix A).

A Reproduction and artifacts

The release has three parts. The repository https://github.com/HalfVulpes/sts2ai holds the source code: the headless host, the trainer with its restart curricula, the in-game mod and the tools below. Its release page holds the playable mod with the five released policies, each policy's exported network for quick evaluation, and their PyTorch weights. A records archive, https://drive.google.com/drive/folders/1OP9IzFIYPuyR97eg1XZru6bpZDSh87Oz, linked from the repository, holds the frozen protocols, per-game results, the training logs of the released runs, the text-interface game logs and the records behind the Discussion.

1. Import a local v0.107.1 installation with `tools/import_game.py`, which verifies all eleven assembly hashes, then build the host and generate the vocabulary locally. The release contains no game files.
2. Run the source tests without game files; integration tests additionally need the local game.
3. Each study has a frozen protocol in the records archive's `protocols/` folder, with its seed panels, initial tensor hashes and decision rules. Per-game results, commands, hardware records and model hashes are in `records/`, and training update logs in `training_logs/`; the archive's index maps every table and figure to its files.
4. Evaluate a policy greedily with `scripts/evaluate.py` or the host's `play` command, which take the exported policy, character, ascension and canonical seed strings; every game counts, and engine errors are reported. `scripts/train.py` prints and runs the trainer command of each reported recipe.

Original code is released under GPL-2.0-only. The repository holds the algorithms and methods needed to train the policies from scratch, including the restart curricula of the longer continuation, second-boss practice and the release continuation; it leaves out platform-specific job tooling. The Spire Mentor adapter of our host is not part of the release; the protocols and per-game records of the studies that used it are in the archive.

B Hyperparameters

Table 13: PPO and curriculum settings shared by all runs in this paper unless a study states otherwise.

Parameter	Value	Parameter	Value
γ / GAE λ	0.9998 / 0.99	samples per iteration	65,536
clip (policy / value)	0.2 / 0.2	epochs, minibatch	2, 4,096
weight decay	0.01	gradient-norm clip	1.0
entropy (combat / other)	0.003 / 0.01	value / auxiliary coefficient	0.5 / 0.2
EMA decay per step	0.998	restart probability	0.5
situations	12	stored states per situation	16 per worker
failure-map round, decay	10 iterations, 0.5	combat credit, floor	0.65, 2%

C Continuation pilot

Before the independent-start study, nine short runs continued one historical Regent checkpoint for about two million decisions each, three per method with three continuation seeds. On the same 200

legacy development seeds the initial policy won 140. Table 14 lists every run. Failure minus uniform was +6.0, −3.5 and +4.0 points; all nine runs share one pretrained source, so the pilot measures adaptation of one policy.

Table 14: Regent A0 continuation pilot, wins of 200 common legacy development games.

Continuation seed	Fresh	Uniform	Failure	Failure − uniform
101	151	141	**153**	↑6.0 pp
202	132	**144**	137	↓3.5 pp
303	130	126	**134**	↑4.0 pp

D Text interface and earlier language-model diagnostics

The text interface runs the same host and exposes three operations: read the state, choose a numbered action and store a note. It offers no reset, lookahead or hidden state, and every action is logged. A known rendering defect prints some energy icons as resource paths inside card text. Before the Astra cohort, GPT-6.1 Terra at high reasoning (one level below maximum) played one game per character under the same bounded-context protocol and lost all five, ending on floors 17, 42, 12, 17 and 14; the released RL policies won one of the same five seeds. An earlier single Ironclad game switched controllers midway and is excluded from all comparisons.

E Legacy development evaluations

Two development evaluations of the released policies on legacy numeric seeds, 50 games per character and level at A0, A9 and A10, gave a five-character mean of 4.0% at A10 under both GPU and C# inference, with A0 between 66% and 92% per character. Greedy actions matched across the two backends on 32 recorded observations per model, and four complete Ironclad A10 runs matched in outcome, floor and decision count. These numbers use a different seed distribution and are not pooled with the canonical panels.

References

[1] Rishabh Agarwal, Max Schwarzer, Pablo Samuel Castro, Aaron Courville, and Marc G. Bellemare. Deep Reinforcement Learning at the Edge of the Statistical Precipice. NeurIPS, 2021. https://arxiv.org/abs/2108.13264.

[2] Xiangchen Cheng et al. AgenticSTS: A Bounded-Memory Testbed for Long-Horizon LLM Agents. 2026. https://arxiv.org/abs/2607.02255.

[3] AlphaSpire2. AlphaSpire, release v0.2.1. Accessed November 6, 2026. https://github.com/AlphaSpire2/alphaspire/releases/tag/v0.2.1.

[4] Siddharth Aphale and Ayushman Singh. Per-Context Reverse Curricula for Sparse-Reward Reinforcement Learning. Version 2, October 2026. https://arxiv.org/abs/2607.26417v2.

[5] Adrien Ecoffet, Joost Huizinga, Joel Lehman, Kenneth O. Stanley, and Jeff Clune. First return, then explore. Nature 590:580–586, 2021. https://arxiv.org/abs/2004.12919.

[6] Carlos Florensa, David Held, Markus Wulfmeier, Michael Zhang, and Pieter Abbeel. Reverse Curriculum Generation for Reinforcement Learning. CoRL, 2017. https://arxiv.org/abs/1707.05300.

[7] Minqi Jiang, Edward Grefenstette, and Tim Rocktäschel. Prioritized Level Replay. ICML, 2021. https://arxiv.org/abs/2010.03934.

[8] Thomas Friedel. Spire Mentor. GitHub repository, README accessed November 7, 2026. https://github.com/tfriedel/spire-mentor.

[9] Andrew Patterson, Samuel Neumann, Martha White, and Adam White. Empirical Design in Reinforcement Learning. 2023. https://arxiv.org/abs/2304.01315.

[10] Cinjon Resnick, Roberta Raileanu, Sanyam Kapoor, Alexander Peysakhovich, Kyunghyun Cho, and Joan Bruna. Backplay: "Man muss immer umkehren". 2018. https://arxiv.org/abs/1807.06919.

[11] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal Policy Optimization Algorithms. 2017. https://arxiv.org/abs/1707.06347.

[12] shuchi-lab. sts2-rl-research. GitHub repository, accessed November 7, 2026. https://github.com/shuchi-lab/sts2-rl-research.

[13] Scott C. Lowe and Sander W. Keemink. SpireBench: Evaluating Long-Horizon Strategic Play of Agentic Models with Slay the Spire. 2026. https://arxiv.org/abs/2610.12596.

[14] Epoch AI. SpireClimb: benchmark, methodology and data. Accessed November 7, 2026. https://epoch.ai/benchmarks/spireclimb.

[15] Ikechukwu Uchendu et al. Jump-Start Reinforcement Learning. ICML, PMLR 202:34556–34583, 2023. https://proceedings.mlr.press/v202/uchendu23a.html.